

Chapter 1

Likelihood Analysis of SCR Models

In this book we mainly focus on Bayesian analysis of spatial capture-recapture models. And, in the previous chapters we learned how to fit some basic spatial capture-recapture models using a Bayesian formulation of the models analyzed in WinBUGS. Despite our focus on Bayesian analysis, it is instructive to develop the basic conceptual and methodological ideas behind classical analysis based on likelihood methods. In fact, simple SCR models can be analyzed fairly easily using classical likelihood methods. This has been the approach taken by Borchers and Efford (2008), Dawson and Efford (2009) and related papers.

In this chapter we provide some conceptual and technical footing for likelihood-based analysis of spatial capture-recapture models. We recognized earlier (Section xxxx) that SCR models are versions of binomial (or other) GLMs with random effects i.e., GLMMs, which are routinely analyzed by likelihood methods. In particular, likelihood analysis is based on the integrated likelihood in which the random effects are removed by integration from the likelihood. In SCR models, the random effect, s , i.e., the 2-dimensional coordinate, is a bivariate random effect. In this chapter, we show that it is straightforward to compute the maximum likelihood estimates (MLE) for SCR models by integrated likelihood. We develop the MLE framework using R, and we also provide a basic introduction to an R package `secr` (Efford 2011) which is based on the stand-alone package `DENSITY` (Efford et al. 2004). To set the context we analyze the SCR model here when N is known because, in that case, it is precisely a GLMM and does not pose any difficulty at all. We generalize the model to allow for unknown N using both conventional ideas based on the joint likelihood (e.g., Borchers book XYZ) and also using a formulation based on data augmentation. We consider likelihood analysis of SCR models in the context of the wolverine camera trapping study (Magoun et al. XYZ) we analyzed in previous chapters

30 to compare/contrast the results.

31 1.1 Likelihood analysis

32 We noted in Chapter 4 that, with N known, the basic SCR model is a type
 33 of binomial regression with a random effect. For such models we can easily
 34 maximize the integrated likelihood numerically. It is worth analyzing this really
 35 simple case in order to understand the underlying mechanics and basic concepts
 36 which are directly relevant to the manner in which many capture-recapture
 37 models are classically analyzed, such as model Mh, and individual covariate
 38 models (see Chapter 6 from Royle and Dorazio 2008).

39 The observation model specified conditional on \mathbf{s}_i is

$$y[i, j] | s[i] \sim \text{Bin}(K, p[i, j])$$

40 where p_{ij} is a function of the individual-specific random effect, \mathbf{s}_i , which has
 41 a $\text{Unif}(\mathcal{S})$ prior distribution. The conditional-on- \mathbf{s} likelihood for the encounter
 42 history data for individual i is the product of J such terms:

$$L(\mathbf{y}_i | \mathbf{s}_i) = \prod_j \text{Bin}(K, p_{ij})$$

43 We note that this assumes that encounter of individual i in each trap is in-
 44 dependent of encounter in every other trap, conditional on \mathbf{s}_i . The so-called
 45 “marginal likelihood” is computed by removing \mathbf{s}_i , by integration, from the
 46 conditional-on- \mathbf{s} likelihood. That is, we compute:

$$f(y[i, j]) = \int_{\mathcal{S}} f(\mathbf{y}_i | s[i]) g(s[i]) ds[i]$$

47 The joint likelihood for all N individuals, assuming independence of encounters
 48 among individuals, is the product of N such terms:

$$\prod_i f(y[i, j])$$

49 We emphasize that two independence assumptions are explicit in this devel-
 50 opment: independence of trap-specific encounters within individuals and also
 51 independence among individuals. In particular, this would only be valid when
 52 individuals are not physically restrained or removed upon capture, and when
 53 traps do not fill up.

54 The key operation for computing the likelihood is solving a 2-dimensional
 55 integration problem. There are some general purpose R packages that imple-
 56 ment various multi-dimensional integration routines including *adapt* (REF) and
 57 *Rcuba* (REF XYZ). In practice, we won't rely on these extraneous R packages
 58 but instead will use perhaps less efficient methods in which we replace the inte-
 59 gral with a summation over an equal area mesh of points on the state-space \mathcal{S}

and explicitly evaluate the integrand at each point. Let $u=1,2,\dots,nG$ index a grid of points where the area of grid cell u is $a[u] = a$ for a regular grid where each cell has the same area. In this case, the trapezoidal rule for approximating the integral yields

$$f(y[i, j]) = \sum_{u=1}^{nG} f(\mathbf{y}_i|u).g(u) * area$$

Where a is the constant area of the grid cells used to approximate the continuous region. This is a general expression that could be used for approximating the integral for any arbitrary bivariate distribution $g(u)$. In the present context note that $g(u) = (1/area(S)) = 1/[nG*a]$ and thus the grid-cell area a cancels in the above expression and we have

$$f(y[i, j]) = \sum_{u=1}^{nG} f(\mathbf{y}_i|u)(1/nG)$$

Which not surprisingly is the same answer we get if S were inherently discrete, having nG unique values with equal probabilities $1/nG$. Then the marginal probability of $y[I, j]$, i.e., by the Law of Total probability, is precisely that last expression.

1.1.1 Implementation (simulated data)

Here we will illustrate how to carry-out this integration and optimization based on the integrated likelihood using simulated data set (i.e., following that from Chapter 4, section XYZ).

From before, we simulated data for 100 individuals and a 25 trap array layed out in a 5 x 5 grid of unit spacing. The specific encounter model is the logit model given above and we used this code to simulate data used in subsequent analyses. The 100 activity centers were simulated on a state-space defined by a 8 x 8 square within which the trap array was centered (thus the trap array is buffered by 2 units). Therefore, the density of individuals in this system is fixed at 100/64.

```
# create 5 x 5 grid of trap locations with unit spacing
traplocs<- cbind(sort(rep(1:5,5)),rep(1:5,5))
Dmat<-e2dist(traplocs,traplocs)
ntraps<-nrow(traplocs)

# define state-space of point process. (i.e., where animals live).
# "delta" just adds a fixed buffer to the outer extent of the traps.

ssbuffer<-2
Xl<-min(traplocs[,1] - ssbuffer)
Xu<-max(traplocs[,1] + ssbuffer)
Yl<-min(traplocs[,2] - ssbuffer)
Yu<-max(traplocs[,2] + ssbuffer)

N<-100 # population size
```

```

99 K<- 20      # number nights of effort
100
101 sx<-runif(N,Xl,Xu)      # simulate activity centers
102 sy<-runif(N,Yl,Yu)
103 S<-cbind(sx,sy)
104 D<- e2dist(S,traplocs)  # distance of each individual from each trap
105
106 alpha0<- -2.5          # define parameters of encounter probability
107 sigma<- 0.5            #
108 alpha1<- 1/(2*sigma*sigma)
109 probcap<- plogis(alpha0 - alpha1*D)    # probability of encounter
110
111 # now generate the encounters of every individual in every trap
112 y<-matrix(NA,nrow=N,ncol=ntraps)
113 for(i in 1:nrow(y)){
114     y[i,]<-rbinom(ntraps,K,probcap[i,])
115 }

```

Now we need to define the integration grid, say G, which we do with the following set of R commands (delta is the grid spacing):

```

118 delta<- .2
119 xg<-seq(Xl+delta/2,Xu-delta/2,by=delta)
120 yg<-seq(Yl+delta/2,Yu-delta/2,by=delta)
121 npix<-length(xg)  # assumes xg and yg same dimension here
122 area<- (Xu-Xl)*(Yu-Yl)/((npix)*(npix)) # dont need area here
123 G<-cbind(rep(xg,npix),sort(rep(yg,npix)))
124 nG<-nrow(G)

```

In this case, the integration grid is set up as a grid with spacing 0.2 which produces a 20 x 20 grid of points for evaluating the integrand if the state-space buffer is set at 2.

We next create an R function that defines the likelihood as a function of the data objects y and X which were created above but, in general, you would read these files into R, e.g., as done in the previous chapter for the analysis using WinBUGS. In addition to these data objects, we need to have defined the various quantities associated with the integration grid created just previously – G, nG, and area. Instead of worrying about making all of these objects and keeping track of them we just put that code above into the likelihood function and pass delta as an additional (optional) argument and a few other things that we need such as the boundary of the state-space over which the integration (summation) is being done.

Here is one reasonably useful variation of a function for estimation based on the integrated likelihood:

```

140
141 intlik1<-function(parm,y=y,delta=.2,X=traplocs,ssbuffer=2){

```

```

142
143 Xl<-min(X[,1])-ssbuffer
144 Xu<-max(X[,1])+ ssbuffer
145 Yu<-max(X[,2])+ ssbuffer
146 Yl<-min(X[,2])-ssbuffer
147
148 xg<-seq(Xl+delta/2,Xu-delta/2,,length=npix)
149 yg<-seq(Yl+delta/2,Yu-delta/2,,length=npix)
150 npix<-length(xg)
151 area<- (Xu-Xl)*(Yu-Yl)/((npix)*(npix))
152 G<-cbind(rep(xg,npix),sort(rep(yg,npix)))
153 nG<-nrow(G)
154 D<- e2dist(X,G)
155
156 alpha0<-parm[1]
157 alpha1<-parm[2]
158 probcap<- plogis(alpha0-alpha1*D)
159 Pm<-matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
160 n0<-sum(apply(y,1,sum)==0) # all zero encounter histories
161 ymat<-y[apply(y,1,sum)>0,] # encounter histories with at least 1 detection
162 ymat<-rbind(ymat,rep(0,ncol(ymat)))
163 lik.marg<-rep(NA,nrow(ymat))
164 for(i in 1:nrow(ymat)){
165   Pm[1:length(Pm)]<- (dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],log=TRUE))
166   lik.cond<- exp(colSums(Pm))
167   lik.marg[i]<- sum( lik.cond*(1/nG))
168 }
169 nv<-c(rep(1,length(lik.marg)-1),n0)
170 -1*( sum(nv*log(lik.marg)) )
171 }
172

```

173 Remarks about the likelihood function: 1. The function accepts as input the
 174 encounter history matrix, y , the trap locations, X , and the state-space buffer.
 175 This allows us to vary the state-space buffer and easily evaluate the sensitivity
 176 of the MLE to the size of the state-space. 2. We have a peculiar handling of
 177 the encounter history matrix y . In particular, we remove the all-zero encounter
 178 histories from the matrix and tack-on a single all-zero encounter history as the
 179 last row which then gets weighted by the number of such encounter histories
 180 (n_0). This is a bit long-winded and strictly unnecessary when N is known, but we
 181 did it this way because the extension to the unknown- N case is now transparent
 182 (as we demonstrate in the following section). 3. The matrix P_m holds the
 183 log-likelihood contributions of each encounter frequency for each possible state-
 184 space location of the individual. 4. The log contributions are summed up and
 185 the result exponentiated on the next line, producing lik.cond , the conditional-
 186 on- s likelihood. (formula XXX above). 5. The marginal likelihood (lik.marg)

187 sums up the conditional elements weighted by $\Pr(s)$ (formula XXX above). 6.
 188 This function assumes that K , the number of replicates, is constant for each
 189 trap. Further, it assumes that the state-space is a square. As an exercise,
 190 consider resolving these two issues by generalizing the code.

191 Here is the R command for maximizing the likelihood and saving the results
 192 into an object called `frog`. The output is a list of the following structure and
 193 these specific estimates are produced using the data set provided in the Online
 194 Supplement):

```
195 # should take 15-30 seconds
196 > starting.values <- c(a0=-2, a1=1) # just to make things clear
197 > frog<-nlm(intlik1,starting.values,y=y,delta=.1,X=traplocs,ssbuffer=2,hessian=TRUE)
198
199 Warning message:
200 In nlm(intlik1, c(-2, 1), y = y, delta = 0.1, X = traplocs, ssbuffer = 2, :
201   NA/Inf replaced by maximum positive value
202
203 > frog
204 $minimum
205 [1] 182.6882
206
207 $estimate
208 [1] -2.623456  1.658335
209
210 $gradient
211 [1] -1.867945e-06  1.587044e-06
212
213 $hessian
214           [,1]      [,2]
215 [1,]  69.16351 -72.13976
216 [2,] -72.13976 108.22528
217
218 $code
219 [1] 1
220
221 $iterations
222 [1] 12
223
```

224 Details about this output can be found on the help page for `nlm`. We note
 225 briefly that `frog$minimum` is the negative log-likelihood value at the MLEs, which are stored in the `frog$estimate`
 226 component of the list. The hessian is the observed Fisher information matrix,
 227 which can be inverted to obtain the variance-covariance matrix using the com-
 228 mands:

```
229 solve(frog$hessian)
```

230 It is worth drawing attention to the fact that the estimates are different than the
231 Bayesian estimates reported in the previous chapter (section XYZ)!!! How can
232 that be?! There are many reasons for this. First Bayesian inference is based on
233 the posterior distribution and it is not generally the case that the MLE should
234 correspond to any particular value of the posterior distribution. If the prior
235 distributions in a Bayesian analysis are uniform, then the mode of the posterior
236 is the MLE but note that Bayesians almost always report posterior means and
237 so there will typically be a discrepancy there. Secondly, we have implemented
238 an approximation to the integral here and there might be a slight bit of error
239 induced by that. We should evaluate that now..

240 To compute the integrated likelihood we used a discrete representation of
241 the state-space so that the integral could be approximated as a summation
242 over possible values of \mathbf{s} with each value being weighted by its probability of
243 occurring, which is $1/nG$ under the assumption that \mathbf{s} is uniform on the state-
244 space S . In chapter 4 we used a discrete state-space in developing a Bayesian
245 analysis of the model in order to be able to modify the state-space in a flexible
246 manner. Bayesian analysis requires simulation of the point process conditional
247 on the observations, and this can be a difficult task when the state-space is
248 continuous but has irregular geometry. Conversely, if the state-space is a regular
249 polygon then Bayesian analysis by MCMC is possibly more efficient with a
250 continuous state-space. We emphasize that the state-space is a part of the
251 model. In some cases there wont be a natural choice of state space beyond some
252 large rectangle containing the trap grid and, in such cases, for regular detection
253 functions the estimate of density is invariant to the size of the state-space (i.e.,
254 the buffer) as long as it is sufficiently large. However if there are good reasons
255 to restrict the state-space, it will tend to have an influence on the likelihood
256 and hence AIC and so forth. As an illustration, lets do that by changing the
257 state space here Use my polygon clipping stuff.

258 In summary, we note that, for the basic SCR model, integrated likelihood is
259 a really easy calculation when N is known. Even for N unknown it is not too
260 difficult, and we will do that shortly. However, if you can solve the known- N
261 problem then you should be able to do a real analysis, for example by considering
262 different values of N and computing the results for each value and then making
263 a plot of the log-likelihood or AIC and choosing the value of N that produces
264 the best log likelihood or AIC. As a homework problem we suggest that the
265 reader take the code given above and try to estimate N without modifying the
266 code by just repeatedly calling that code for different values of N and trying to
267 deduce the best value. Nevertheless, we will formalize the unknown- N problem
268 shortly. We note that the software package DENSITY (Efford et al. 2004)
269 implements certain types of SCR models using integrated likelihood methods.
270 DENSITY has been made into an R package called secr (Efford 2011) and we
271 provide an analysis of some data using secr shortly along with a discussion of
272 its capabilities.

273 1.2 MLE when N is Unknown

274 Here we build on the previous introduction to integrated likelihood but we
 275 consider now the case in which N is unknown. We will see that adapting the
 276 analysis based on the N -known model is really straightforward for the more
 277 general problem. The main distinction is that we don't observe the all-zero
 278 encounter history so we have to make sure we compute the probability for that
 279 encounter history which we do by tacking a row of zeros onto the encounter
 280 history matrix. In addition, we include the number of such all-zero encounter
 281 histories as an unknown parameter of the model. Call that unknown quantity
 282 n_0 . In addition, we have to be sure to include a combinatorial term to account
 283 for the fact that of the n observed individuals there are N choose n ways to
 284 realize a sample of size n . The combinatorial term involves the unknown n_0 and
 285 thus it must be included in the likelihood (we have already done that in the
 286 previous likelihood construction even though we didn't need it).

287 To summarize, when N is unknown, the n observed encounter histories have
 288 a multinomial distribution with probabilities $\pi(i)$ and sample size N^1 . The last
 289 cell the zero cell is computed by carrying out the integral in expression XYZ
 290 above for the all-zero encounter history and we have to account for the fact that
 291 there are $n_0 = N - n$ such encounter histories.

292 To analyze a specific case, we'll read in our fake data set (simulated using the
 293 parameters given above). To set some things up in our workspace we do this:

```
294 simSCR0.fn( blah blah blah)
295
296 > y<-read.csv("ind_by_trap.csv")[, -1]
297 > y<-as.matrix(y)
298 >
299 > X<-read.csv("traplocs.csv")[, -1]
300 > X<-as.matrix(X)
301 >
302
```

303 Recall that these data were generated with $N=100$, on an 8×8 unit state-
 304 space representing the trap locations (X) buffered by 2 units. As before, the
 305 likelihood is defined in the R workspace as an R function which takes an argu-
 306 ment being the unknown parameters of the model and additional arguments as
 307 prescribed. In particular, as before, we provide the encounter history matrix y ,
 308 the trap locations traplocs , the spacing of the integration grid (delta) and the
 309 state-space buffer. Here is the new likelihood function:

```
310 intlik2<-function(parm,y=y,delta=.3,X=traplocs,ssbuffer=2){
311
312   Xl<-min(X[,1]) -ssbuffer
```

¹Maybe you could show an alternative simulation script to generate data using the `rmulti-`
`nom` function. This would make it a little more clear for people


```

313 Xu<-max(X[,1])+ ssbuffer
314 Yu<-max(X[,2])+ ssbuffer
315 Yl<-min(X[,2])- ssbuffer
316
317 #delta<- (Xu-Xl)/npix
318 xg<-seq(Xl+delta/2,Xu-delta/2,delta)
319 yg<-seq(Yl+delta/2,Yu-delta/2,delta)
320 npix.x<-length(xg)
321 npix.y<-length(yg)
322 area<- (Xu-Xl)*(Yu-Yl)/((npix.x)*(npix.y))
323 G<-cbind(rep(xg,npix.y),sort(rep(yg,npix.x)))
324 nG<-nrow(G)
325 D<- e2dist(X,G)
326
327 alpha0<-parm[1]
328 alpha1<-parm[2]
329 n0<-exp(parm[3])
330 probcap<- plogis(alpha0-alpha1*D)
331 Pm<-matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
332 ymat<-rbind(y,rep(0,ncol(y)))
333
334 lik.marg<-rep(NA,nrow(ymat))
335 for(i in 1:nrow(ymat)){
336 Pm[1:length(Pm)]<- (dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],log=TRUE))
337 lik.cond<- exp(t(Pm)%*%rep(1,nrow(probcap)))
338 lik.marg[i]<- sum( lik.cond*(1/nG) )
339 }
340 nv<-c(rep(1,length(lik.marg)-1),n0)
341 part1<- lgamma(nrow(y)+n0+1) - lgamma(n0+1)
342 part2<- sum(nv*log(lik.marg))
343 -1*(part1+ part2)
344 }

```

345 To execute this function for the data that we just read into the workspace,
 346 we execute the following command (saving the result in our friend frog).

```

347 > frog<-nlm(intlik2,c(-2.5,2,log(4)),hessian=TRUE,y=y,X=traplocs,delta=.2,ssbuffer=2)

```

348 This results in the usual output, including the parameter estimates, the
 349 gradient, and the numerical Hessian which is useful for obtaining asymptotic
 350 standard errors (omitted here).

```

351 > frog
352 $minimum
353 [1] 181.4657
354
355 $estimate

```

```
356 [1] -2.629342  1.790849  3.997268
```

```
357
```

```
358 [ . Additional output deleted . ]
```

```
359     The estimate of population size for the state-space (using the default state-
360 space buffer) is
```

```
361 > nrow(y)+exp(3.997)
```

```
362 [1] 113.4346
```

```
363     Which differs from the data-generating truth (N=100) as we might expect.
```

364 1.2.1 Exercises

```
365 1. Run the analysis with different state-space buffers and comment on the result.
```

```
366 2. Conduct a brief simulation study using this code by simulating 100 data
367 sets and obtain the MLEs for each data set. Do things seem to be working as
368 you expect?
```

```
369 3. Further extensions: It should be straightforward to generalize the in-
370 tegrated likelihood function to accommodate many different situations. For
371 examples, if we want to include more covariates in the model we can just add
372 stuff to the probcap there, and add the relevant parameters to the argument
373 that gets passed to the intlik function. For the simulated data, make up a co-
374 variate by generating a Bernoulli covariate (trap type perhaps baited or not
375 baited) randomly and try to modify the likelihood to accommodate that.
```

```
376 4. We would probably be interested in devising the integrated likelihood for
377 the full 3-d encounter history array so that we could include temporally varying
378 covariates. This is not difficult but naturally will slow down the execution
379 substantially. The interested reader should try to expand the capabilities of
380 this basic R function.
```

381 1.2.2 Integrated Likelihood using the model under data 382 augmentation

```
383 Note that this likelihood analysis is based on the standard likelihood in which
384 N (or n0) is an explicit parameter. This is usually called the joint likelihood or
385 unconditional likelihood. We could also express the joint likelihood using data
386 augmentation, replacing the parameter N with psi. See Royle et al. (2007);
387 Royle and Dorazio (2008, distance sampling example), etc.. Briefly, we note
388 that the likelihood under data augmentation looks like a zero-inflated binomial
389 mixture precisely as an occupancy type model (see Royle 2006). We think the
390 interested reader could adapt the material from Royle and Dorazio (2008) with
391 the R code given above for the likelihood and implement the likelihood analysis
392 based on the model under data augmentation. Despite that we can carry-out
393 likelihood analysis of models under data augmentation, we primarily advocate
394 data augmentation for Bayesian analysis.
```

1.2.3 Extensions and Classical model selection and assessment

There are other types of covariates of interest: behavioral response, sex-specificity of parameters and all of these things. Some of these can be added directly to the likelihood if the covariate is fixed and known for all individuals captured or not. This excludes most covariates but it does include behavioral response. Sex-specificity is more difficult since sex is not known for uncaptured individuals. Trap-specific covariates such as trap type or status, or time-specific covariates such as date, are relatively easy to deal with (we leave these as exercises). We apply these various models in Chapter XXXX. To analyze such models, we do Bayesian analysis of the joint likelihood facilitated by the use of data augmentation. For covariates that are not fixed and known for all individuals, it is hard to do MLE for these based on the joint likelihood as we have developed above. Instead what people normally do is use what is colloquially referred to as the Huggins-Alho type model which is one of the approaches taken in the software package secr (Efford XYZ; see chapter XYZ).

In most analyses, one is interested in choosing from among various potential models. A good thing about classical analysis based on likelihood is we can do rote application of AIC without thinking about anything. With distance as a covariate (e.g., distance sampling) this is usually applied to some arbitrary selection of distance functions. We don't recommend this. Given there is hardly ever (if at all) a rational science-based reason for choosing some particular distance function we believe that this standard approach will invariably lead to over-fitting. The fact that AIC is easy to compute does not mean that it should be abused in such fashions. Further discussion is made in chapters XYZ.

Goodness-of-fit In many analyses based on likelihood methods it is possible to cook-up fit statistics for which asymptotic distributions are known. In general, however, applied statisticians tend to adopt bootstrapping based on heuristically appealing fit statistics. An omnibus global GoF statistic is not so obvious but we can apply bootstrapping principles to SCR models directly which we discuss in chapter XYZ. Bayesian goodness-of-fit is almost always addressed with Bayesian p-values or some other posterior predictive check (REF XXX). Thus the approach whether Bayesian or classical is the same. We identify a fit statistic, we do a bootstrap (classical) or a Bayesian p-value. Royle et al. (2011) decomposed the fit problem into separate evaluations of the CSR hypothesis and the encounter process model. We discuss all of this in Chapter XYZ.

1.3 Likelihood analysis of the wolverine camera trapping data

Here we compute the MLEs for the wolverine data using an expanded version of the function we developed in the previous section. To accommodate that each trap might be operational a variable number of nights, we provided an

additional argument to the likelihood function (allowing for a vector K), which requires also a modification to the construction of the likelihood (see Online Supplement). In addition, we had to accommodate that the state-space is a general rectangle, and we included a line in the code to compute the state-space area which we apply below for computing density. The more general function (`intlik3`) is given in the online supplement (shall we provide it here?).

The data were read into our R session and manipulated using the following commands. Note that we use the utility R function `SCR23darray.fn` which we defined in section XYZ.XYZ.

```

> wcaps<-source("wcaps.R")$value
> wtraps<-source("wtraps.R")$value
> K.wolv<-apply(wtraps[,4:ncol(wtraps)],1,sum)
>
> xx<-SCR23darray.fn(wcaps,ntraps=37,nperiods=165)
> y.wolv<- apply(xx,c(1,3),sum)
> traplocs.wolv<-wtraps[,2:3]
> traplocs.wolv<-traplocs.wolv/10000
>
> nlm(lik,c(-1.5,1.2,log(4)),hessian=TRUE,y=y.wolv,K=K.wolv,X=traplocs.wolv,
delta=.1,ssbuffer=2)$estimate
[1] -1.646692  3.128712  3.761974

```

We obtained the MLEs for a state-space buffer of 2 (standardized units) and for integration grid with spacing $\delta = .3, .2, .1, .05$. The MLES for these 4 cases are as follows:

```

> nlm(intlik3,c(-1.5,1.2,log(4)),hessian=TRUE,y=y.wolv,K=K.wolv, X=traplocs.wolv,delta=.3,ssbuffer=2)$estimate
[1] -1.654535  3.108126  3.723244
> nlm(lik,c(-1.5,1.2,log(4)),hessian=TRUE,,delta=.2,ssbuffer=2)$estimate
[1] -1.643023  3.133985  3.749195
> nlm(lik,c(-1.5,1.2,log(4)),hessian=TRUE,,delta=.1,ssbuffer=2)$estimate
[1] -1.646692  3.128712  3.761974
> nlm(lik,c(-1.5,1.2,log(4)),hessian=TRUE,,delta=.05,ssbuffer=2)$estimate
[1] -1.647191  3.128308  3.769455

```

We see the results change only slightly as the fineness of the integration grid increases. Conversely, the runtime on the platform of the day for the 4 cases was approximately 19s, 40s, 169s, and 723s which as we have suggested before could probably be regarded as relative to each other, across platforms, for gaging the decrease in speed as the fineness of the integration grid increases.

Next we studied the effect of the state-space buffer on the MLEs, using a fixed $\delta = 1$ for all analyses. We used state-space buffers of 1 to 4 units stepped by .5. This produced the following results, given here are the state-space buffer, area of the state-space, the MLE of N for the prescribed state-space and the corresponding MLE of density:

REDO ANALYSES

	ssbuff	Ass	Nhat	Dhat
[1,]	1.0	66.98212	42.18630	0.6298144
[2,]	1.5	84.36242	52.12473	0.6178667
[3,]	2.0	103.74272	64.03329	0.6172317
[4,]	2.5	125.12302	77.36792	0.6183348
[5,]	3.0	148.50332	91.99139	0.6194568
[6,]	3.5	173.88362	107.87487	0.6203855
[7,]	4.0	201.26392	125.01359	0.6211426

The estimates of D stabilize rapidly² and the results suggest that wolverine density is around 0.62 individuals per 100 square km (recall that a unit is 10 x 10 km). This is about 6.2 individuals per thousand square km which compares with XYZ reported in Royle et al. (2011) based on a clipped state-space as described in section XYZ.

In a later chapter analysis by MLE is done with an irregular state-space. Therefore we will have to extend the R function again to accept as possible input a pre-defined state-space grid.

1.3.1 Exercises

1. Compute the 95% confidence interval for wolverine density, somehow.
2. Compute the AIC of this model and modify `intlik3` to consider alternative link functions (at least one additional) and compare the AIC of the different models and the estimates. Comment. [should we do that here?].

1.4 Program DENSITY and the R package secr

DENSITY is a software program developed by Efford et al. (2004) for fitting spatial capture-recapture models based mostly on classical maximum likelihood

²not very convincing

estimation and related inference methods. Efford (2011) has also released an R package named `secr`, that contains many of the functions within `DENSITY` but also incorporates new models and features. Here, we will focus on `secr` as it will continue to be developed, contains more functionality and is based in R. To install and run models in `secr`, you must download the package and load it in R.

```
>install.packages(secr)
>library(secr)
```

`SECR` allows the user to simulate data and fit a suite of models with various detection functions and covariate responses. `SECR` uses the standard R model specification framework using tildes. E.g., the model command is `secr.fit` and is generally written as

```
>secr.fit(capturedata, model = list(D~1, g0~1, sigma~1), buffer = 20000)
```

where we have `g0 1` indicating the intercept model. To include covariates, this would be written as `g0 b` where `b` is a behavioral covariate. Possible predictors for detection probability include both pre-defined variables (e.g., `t` and `b` corresponding to time and behavior), and user-defined covariates of several kinds. The discussion of covariates is developed in chapter XX(8).

Before we can fit the models, the data must first be entered into `SECR`. Two input files are required: trap layout (location and identification information for each trap) and capture data (e.g., sampling session, animal identification, trap day, and trap location). `SECR` requires that you specify the trap type, the two most common for camera trapping/hair snares are proximity detectors and count detectors. The ‘proximity’ detector type allows, at most, one detection of each individual at a particular detector on any occasion. The count detector designation allows repeat encounters of each individual at a particular detector on any occasion. There are other detector types that one can select such as: ‘polygon’ detector type which allows for a trap to be a sampled polygon, e.g., scat surveys, and ‘signal’ detector which allows for traps that have a strength indicator, e.g., acoustic arrays. The detector types `single` and `multi` can be confusing as `multi` seems like it would appropriate for something like a camera trap, but instead these two designations refer to traps that retain individuals, thus precluding the ability for animals to be captured in other traps during the sampling occasion. The `single` type indicates trap that can only catch one animal at a time, while `multi` indicates traps that may catch more than one animal at a time. For a full review of the detector types, one should look at the help manual, which can be accessed in R after installing the `SECR` package by using the command:

```
>RShowDoc("secr-manual", package = "secr")
```

As with all of the `scr` models, `SECR` fits a detection function relating the probability of detection to the distance of a detector from an individual activity center.

SECR allows the user to specify one of a variety of detection functions including the commonly used half-normal, hazard rate, and exponential. There are 12 different functions, but some are only available for simulating data, and one should take caution when using different detection functions as the interpretation of the parameters, such as sigma, may not be consistent across formulations. The different detection functions are defined in the secr manual and can be found by calling the help function for the detection function:

```
> ?detectfn
```

It is useful to note that secr requires the buffer distance to be defined in meters and density will be returned as number of animals per hectare. Thus to make comparisons between secr and other models, we will often have to convert the density to the same units. Also, note that sigma is returned in units of meters.

³

1.4.1 Analysis using secr package

To demonstrate the use of the secr package, we will show how to do the same analysis on the wolverine study as shown in section 4.6. To use the secr package, the data need to be formatted in a similar but slightly different manner than we use in WinBUGS. After installing the secr package, we first have to read in the trap locations and other related information, such as if the trap is operational during a sampling occasion. The secr package reads in the trap data through a command called read.traps, which requires the detector type as input. The detector type is important because it will determine the likelihood that secr will use to fit the model. Here, we have selected proximity since individuals are captured at most once in each trap during each sampling occasion.

```
>traps= read.csv(wtraps.csv)
>colnames(traps)[1:3]<- c("trapID","x", "y") #name the first 3 columns
# to match the secr nomenclature

>trapfile <- read.traps(data = traps, detector = "proximity")
```

After reading in the data, we now need to create the encounter matrix or array. The secr package does this through the use of the make.caphist command, where we provide the capture histories in raw data format (each line contains the session, identification number, occasion, and trap id for only 1 individual). This is the format that was shown in the data input file wcaps, and we only need a line or two to organize the data into the order that the make.caphist command wants. In creating the capture history, we provide also the trapfile with the trap information, and the format (e.g., here fmt= trapID) so that secr

³One question: SECR only ever reports sigma. What exactly is sigma? It is a scale parameter of a detection function and all detection functions have a scale parameter. But in what sense is this sigma parameter related to home range diameter? Efford doesn't explain this, does he? In some sections in chapter 4 or possibly 6 we get into this issue.

596 knows how to match the encounters to the trap, and finally, we provide the
 597 number of occasions.

```
598 >wolv.dat <- wcaps[,c(2, 3, 1)] #NEED TO UPDATE THIS WHEN I GET THE FILES, I JUST GU
599 >wolv.dat <- cbind(rep(1, dim(wolv.dat)[1], wolv.dat)
600 >colnames(wolv.dat) <- c("Session", "ID", "Occasion", "trapID")
601
602 >wolvcapt=make.caphist(wolv.dat, trapfile, fmt = "trapID", noccasions = 165)
```

603 Calling the secr.fit command, will run the model. We are using the basic
 604 model (SCR0), so we do not need to make any specifications in the command
 605 line except for the providing the buffer size (in m). To specify different models,
 606 you can change the default D 1, g0 1, sigma 1, which the interested reader
 607 can do with very little difficulty.

```
608 > wolv.secr=secr.fit(wolvcapt, model = list(D~1, g0~1, sigma~1), buffer = 20000)
609
610 >wolv.secr
611
612 secr.fit( caphist = wolvcapt, buffer = 20000, binomN = 1 )
613 secr 2.0.0, 18:26:39 05 Jul 2011
614
615 Detector type      proximity
616 Detector number    37
617 Average spacing    4415.693 m
618 x-range            593498 652294 m
619 y-range            6296796 6361803 m
620 N animals          : 21
621 N detections       : 115
622 N occasions        : 165
623 Mask area          : 1037069 ha
624
625 Model              : D~1 g0~1 sigma~1
626 Fixed (real)       : none
627 Detection fn       : halfnormal
628 Distribution        : poisson
629 N parameters       : 3
630 Log likelihood     : -746.754
631 AIC                 : 1499.508
632 AICc               : 1500.920
633
634 Beta parameters (coefficients)
635      beta      SE.beta      lcl      ucl
636 D      -9.749576 0.23027860 -10.200913 -9.298238
637 g0     -4.275736 0.15846104 -4.586313 -3.965158
638 sigma  8.699202 0.07868944  8.544973  8.853430
639
```



```

640 Variance-covariance matrix of beta parameters
641           D           g0          sigma
642 D      0.053028233  0.000546922 -0.005226926
643 g0      0.000546922  0.025109900 -0.005885213
644 sigma -0.005226926 -0.005885213  0.006192027
645
646 Fitted (real) parameters evaluated at base levels of covariates
647           link      estimate SE.estimate          lcl          ucl
648 D      log 5.831941e-05 1.360973e-05 3.713638e-05 9.158548e-05
649 g0     logit 1.371121e-02 2.142902e-03 1.008756e-02 1.861207e-02
650 sigma  log 5.998124e+03 4.727205e+02 5.140849e+03 6.998355e+03

```

651 Under the fitted (real) parameters, we find D, the density, given in units
652 of individuals/hectare (1 hectare = 100 m²). To convert this into individ-
653 uals/1000km², we multiply by 100000, thus our density estimate is 5.83 individ-
654 uals/1000 km². Sigma is given in units of meters, to convert to kilometers, we
655 divide by 1000, which puts sigma at 5.99 km. Both of these estimates are very
656 similar to those provided in section 4.6 for the buffer size equal to 20 km. As
657 an exercise, run this analysis for 30 and 40 km buffers and compare those found
658 in section 4.6 under WinBUGS. NOTE: The secr.fit will return a warning when
659 the buffer size appears to be too small. This is useful particularly with the
660 different units being used between programs and packages.

661 1.5 Summary and Outlook

662 In this chapter, we showed that classical analysis of SCR models based on like-
663 lihood methods is a relatively simple proposition. Analysis is based on the
664 so-called integrated likelihood in which the individual activity centers (random
665 effects) are removed from the conditional-on-s likelihood by integration. We
666 showed how to construct the integrated likelihood and fit some simple models
667 in the R programming language. In addition, likelihood analysis for some broad
668 classes of SCR models can be accomplished in the software package DENSITY
669 (and other packages such as ADMB) or in the equivalent R library secr which
670 we provided an illustration of here. In later chapters we provide more detailed
671 analyses of SCR data using the secr package.

672 To compute the integrated likelihood we have to precisely describe the state-
673 space of the underlying point process. In practice, this leads to a buffer around
674 the trap array. We note that this is not really a buffer strip in the sense of
675 Wilson et al. (XYZ) which is a feature of the analysis but it is somewhat more
676 general here. In particular, it establishes the support of the integrand which
677 we generally require to compute any integral. It might be that the integrand
678 itself is finite even if the support is infinity but that may or may not be the
679 case depending on the choice of detection function. As a practical matter then,
680 it will typically be the case that, while estimates of N increase with the size of
681 the buffer, estimates of density stabilize. This is not a feature of the classical
682 methods based on using model M0 or model Mh and buffering the trap array.

Why or why not use likelihood inference exclusively? For certain specific models, it is probably more computationally efficient to produce MLEs. However, WinBUGS is extremely flexible in terms of describing models, although it sometimes can be quite slow. We can devise models in WinBUGS easily that we cannot fit in secr. E.g., random individual effects of various types (see next chapter), we can handle missing covariates in complete generality and seamlessly, and impose arbitrary distributions on random variables. Moreover, models can easily be adapted to include auxiliary data types. For example, we might have camera trapping and genetic data and we can describe the models directly in WinBUGS and fit a joint model. For the MLE we have to write a custom new piece of code for each model or hope someone has done it for us. Later we consider open population models which are straightforward to develop in WinBUGS but, so far, there is no available platform for doing MLE although we imagine one could develop this. . Another thing that is more conceptual here is non-CSR point processes (see chapter XYZ) and generating predictions of how many individuals have home range centers in any particular polygon. Basic benefits of Bayesian analysis have been discussed elsewhere (Chapter 2? BPA book? Link and Barker?) and we believe these are compelling. On the other hand, likelihood analysis makes it easy to do model-selection by AIC. Goodness-of-fit is probably no more difficult or easy under either paradigm (see next chapter?).

In summary, basic SCR models are easy to implement by either likelihood or Bayesian methods but we feel that the typical user will realize much more flexibility in model development using existing platforms for Bayesian analysis. While these tend to be slow (sometimes excruciatingly slow), this will probably not be an impediment in most problems, especially at some near point in the future. Since we spent a lot of time here talking about specific technical details on how to implement likelihood analysis of SCR models, we provided a corresponding treatment in the next chapter on how to devise MCMC algorithms for SCR models. This is a bit more tedious and requires more coding, but is not technically challenging (except perhaps to develop highly efficient algorithms which we don't excel at).