

Chapter 1

Likelihood Analysis of SCR Models

X

In this book we mainly focus on Bayesian analysis of spatial capture-recapture models. And, in the previous chapters we learned how to fit some basic spatial capture-recapture models using a Bayesian formulation of the models analyzed in WinBUGS. Despite our focus on Bayesian analysis, it is instructive to develop the basic conceptual and methodological ideas behind classical analysis based on likelihood methods. In fact, simple SCR models can be analyzed fairly easily using classical likelihood methods. This has been the approach taken by Borchers and Efford (2008), Dawson and Efford (2009) and related papers.

In this chapter we provide some conceptual and technical footing for likelihood-based analysis of spatial capture-recapture models. We recognized earlier (Section xxxx) that SCR models are versions of binomial (or other) GLMs with random effects i.e., GLMMs, which are routinely analyzed by likelihood methods. In particular, likelihood analysis is based on the integrated likelihood in which the random effects are removed by integration from the likelihood. In SCR models, the random effect, \mathbf{s} , i.e., the 2-dimensional coordinate, is a bivariate random effect. In this chapter, we show that it is straightforward to compute the maximum likelihood estimates (MLE) for SCR models by integrated likelihood. We develop the MLE framework using R, and we also provide a basic introduction to an R package `secr` (Efford 2011) which is based on the stand-alone package `DENSITY` (Efford et al. 2004). To set the context we analyze the SCR model here when N is known because, in that case, it is precisely a GLMM and does not pose any difficulty at all. We generalize the model to allow for unknown N using both conventional ideas based on the “joint likelihood” (e.g., Borchers et al., 2002) and also using a formulation based on data augmentation. We consider likelihood analysis of SCR models in the context of the wolverine

30 camera trapping study (Magoun et al., 2011) we analyzed in previous chapters
 31 to compare/contrast the results.

32 1.1 Likelihood analysis

33 We noted in Chapter 4 that, with N known, the basic SCR model is a type
 34 of binomial regression with a random effect. For such models we can easily
 35 obtain maximum likelihood estimators of model parameters based on integrated
 36 likelihood. The integrated likelihood is based on the marginal distribution of the
 37 data y in which the random effects are removed by integration. Conceptually,
 38 our model is a specification of the conditional-on- \mathbf{s} model $[y|\mathbf{s}, \theta]$ and we have a
 39 “prior distribution” for \mathbf{s} , say $[\mathbf{s}]$, and the marginal distribution of the data y is

$$[y|\theta] = \int_{\mathbf{s}} [y|\mathbf{s}, \theta][\mathbf{s}]d\mathbf{s}.$$

40 When viewed as a function of θ for purposes of estimation, the marginal dis-
 41 tribution $[y|\theta]$ is often referred to as the *integrated likelihood*.

42 It is worth analyzing the simplest SCR model with known- N in order to
 43 understand the underlying mechanics and basic concepts. These are directly
 44 relevant to the manner in which many capture-recapture models are classically
 45 analyzed, such as model Mh, and individual covariate models (see chapt. 6 from
 46 Royle and Dorazio (2008)). To develop integrated likelihood for SCR models,
 47 we first identify the conditional likelihood.

48 The observation model for each encounter observation y_{ij} , specified condi-
 49 tional on \mathbf{s}_i , is

$$y_{ij}|\mathbf{s}_i \sim \text{Bin}(K, p_{\theta}(\mathbf{x}_j, \mathbf{s}_i)) \quad (1.1)$$

50 where we have indicated the dependence of p_{ij} on \mathbf{s} and parameters θ explicitly.
 51 For the random effect we have $\mathbf{s}_i \sim \text{Unif}(\mathcal{S})$. The joint distribution of the data
 52 for individual i is the product of J such terms (i.e., contributions from each of
 53 J traps).

$$[\mathbf{y}_i|\mathbf{s}_i, \theta] = \prod_j \text{Bin}(K, p_{\theta}(\mathbf{x}_j, \mathbf{s}_i))$$

54 We note that this assumes that encounter of individual i in each trap is indepen-
 55 dent of encounter in every other trap, conditional on \mathbf{s}_i , this is the fundamental
 56 property of SCR0 or “multi-catch” traps.

57 The so-called “marginal likelihood” is computed by removing \mathbf{s}_i , by integra-
 58 tion, from the conditional-on- \mathbf{s} likelihood. That is, we compute:

$$\mathcal{L}(\theta|\mathbf{y}_i) = \int_{\mathcal{S}} \mathcal{L}(\theta|\mathbf{y}_i|\mathbf{s}_i)g(\mathbf{s}_i)d\mathbf{s}_i$$

59 The joint likelihood for all N individuals, assuming independence of encounters
 60 among individuals, is the product of N such terms:

$$\prod_i f(y[i, j])$$

61 We emphasize that two independence assumptions are explicit in this devel-
 62 opment: independence of trap-specific encounters within individuals and also
 63 independence among individuals. In particular, this would only be valid when
 64 individuals are not physically restrained or removed upon capture, and when
 65 traps do not fill up.

66 The key operation for computing the likelihood is solving a 2-dimensional in-
 67 tegration problem. There are some general purpose **R** packages that implement
 68 a number of multi-dimensional integration routines including **adapt** (REF) and
 69 **Rcuba** (REF XYZ). In practice, we wont rely on these extraneous **R** packages
 70 but instead will use perhaps less efficient methods in which we replace the inte-
 71 gral with a summation over an equal area mesh of points on the state-space \mathcal{S}
 72 and explicitly evaluate the integrand at each point. Let $u = 1, 2, \dots, nG$ index
 73 a grid of nG points where the area of grid cell u is $a(u) \equiv a$ (for a regular grid).
 74 In this case, the trapezoidal rule for approximating the integral yields

$$f(y_{ij}) = \sum_{u=1}^{nG} f(\mathbf{y}_i|u)g(u) * area$$

75 This is a general expression that could be used for approximating the integral
 76 for any arbitrary bivariate distribution $g(u)$. In the present context note that
 77 $g(u) = (1/area(\mathcal{S})) = 1/[nG * a]$ and thus the grid-cell area cancels in the above
 78 expression and we have

$$f(y_{ij}) = \sum_{u=1}^{nG} f(\mathbf{y}_i|u)(1/nG)$$

79 Which not surprisingly is the same answer we get if \mathcal{S} were inherently discrete,
 80 having nG unique values with equal probabilities $1/nG$. Then the marginal
 81 probability of y_{ij} , i.e., by the Law of Total probability, is precisely that last
 82 expression.

83 1.1.1 Implementation (simulated data)

84 Here we will illustrate how to carryout this integration and optimization based
 85 on the integrated likelihood using simulated data (i.e., following that from Chap-
 86 ter 4). Using **simSCR0.fn** we simulate data for 100 individuals and a 25 trap
 87 array layed out in a 5×5 grid of unit spacing. The specific encounter model is
 88 the half-normal model. The 100 activity centers were simulated on a state-space
 89 defined by a 8×8 square within which the trap array was centered (thus the
 90 trap array is buffered by 2 units). Therefore, the density of individuals in this
 91 system is fixed at $100/64$.

92 In the following set of R commands we generate the data and then harvest
 93 the required data objects:

```
94 data<-simSCR0.fn(discard0=FALSE,sd=2013)
95 y<-data$Y
```

```

96 traplocs<-data$traplocs
97 nind<-nrow(y)
98 X<-data$traplocs
99 J<-nrow(X)
100 K<-data$K
101 Xl<-data$xlim[1]
102 Yl<-data$ylim[1]
103 Xu<-data$xlim[2]
104 Yu<-data$ylim[2]

```

Now we need to define the integration grid, say **G**, which we do with the following set of **R** commands (here, `delta` is the grid spacing):

```

107 delta<- .2
108 xg<-seq(Xl+delta/2,Xu-delta/2,by=delta)
109 yg<-seq(Yl+delta/2,Yu-delta/2,by=delta)
110 npix<-length(xg) # assumes xg and yg same dimension here
111 area<- (Xu-Xl)*(Yu-Yl)/((npix)*(npix)) # dont need area for anything
112 G<-cbind(rep(xg,npix),sort(rep(yg,npix)))
113 nG<-nrow(G)

```

In this case, the integration grid is set up as a grid with spacing $\delta = 0.2$ which produces a 40×40 grid of points for evaluating the integrand if the state-space buffer is set at 2.

We next create an **R** function that defines the likelihood as a function of the data objects *y* and *X* which were created above but, in general, you would read these files into **R**, e.g., from a .csv file. In addition to these data objects, we need to have defined the various quantities associated with the integration grid **G** and *nG*. However, instead of worrying about making all of these objects and keeping track of them we just put that code above into the likelihood function and pass δ as an additional (optional) argument and a few other things that we need such as the boundary of the state-space over which the integration (summation) is being done. Here is one reasonably useful variation of a function for estimation based on the integrated likelihood:

```

127
128 intlik1<-function(parm,y=y,delta=.2,X=traplocs,ssbuffer=2){
129
130   Xl<-min(X[,1]) - ssbuffer
131   Xu<-max(X[,1]) + ssbuffer
132   Yu<-max(X[,2]) + ssbuffer
133   Yl<-min(X[,2]) - ssbuffer
134
135   xg<-seq(Xl+delta/2,Xu-delta/2,,length=npix)
136   yg<-seq(Yl+delta/2,Yu-delta/2,,length=npix)
137   npix<-length(xg)
138
139   G<-cbind(rep(xg,npix),sort(rep(yg,npix)))

```

```

140 nG<-nrow(G)
141 D<- e2dist(X,G)
142
143 alpha0<-parm[1]
144 alpha1<-parm[2]
145 probcap<- plogis(alpha0)*exp(-alpha1*D*D)
146 Pm<-matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
147           # all zero encounter histories
148 n0<-sum(apply(y,1,sum)==0)
149           # encounter histories with at least 1 detection
150 ymat<-y[apply(y,1,sum)>0,]
151 ymat<-rbind(ymat,rep(0,ncol(ymat)))
152 lik.marg<-rep(NA,nrow(ymat))
153 for(i in 1:nrow(ymat)){
154   Pm[1:length(Pm)]<- (dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],log=TRUE))
155   lik.cond<- exp(colSums(Pm))
156   lik.marg[i]<- sum( lik.cond*(1/nG))
157 }
158 nv<-c(rep(1,length(lik.marg)-1),n0)
159 -1*( sum(nv*log(lik.marg)) )
160 }

```

The function accepts as input the encounter history matrix, y , the trap locations, X , and the state-space buffer. This allows us to vary the state-space buffer and easily evaluate the sensitivity of the MLE to the size of the state-space. Note that we have a peculiar handling of the encounter history matrix y . In particular, we remove the all-zero encounter histories from the matrix and tack-on a single all-zero encounter history as the last row which then gets weighted by the number of such encounter histories ($n0$). This is a bit long-winded and strictly unnecessary when N is known, but we did it this way because the extension to the unknown- N case is now transparent (as we demonstrate in the following section). The matrix Pm holds the log-likelihood contributions of each encounter frequency for each possible state-space location of the individual. The log contributions are summed up and the result exponentiated on the next line, producing `lik.cond`, the conditional-on- s likelihood (Eq. 1.1 above). The marginal likelihood (`lik.marg`) sums up the conditional elements weighted by $\Pr(s)$ (formula XXX above). Finally, this function assumes that K , the number of replicates, is constant for each trap. Further, it assumes that the state-space is a square. As an exercise, consider resolving these two issues by generalizing the code.

Here is the **R** command for maximizing the likelihood and saving the results into an object called `frog`. The output is a list of the following structure and these specific estimates are produced using the simulated data set:

```

182 # should take 15-30 seconds
183
184 > starting.values <- c(-2, 2)
185 > frog<-nlm(intlik1,starting.values,y=y,delta=.1,X=traplocs,ssbuffer=2,hessian=TRUE)

```

```

186 > frog
187
188 $minimum
189 [1] 297.1896
190
191 $estimate
192 [1] -2.504824  2.373343
193
194 $gradient
195 [1] -2.069654e-05  1.968754e-05
196
197 $hessian
198           [,1]      [,2]
199 [1,]  48.67898 -19.25750
200 [2,] -19.25750  13.34114
201
202 $code
203 [1] 1
204
205 $iterations
206 [1] 11

```

207 Details about this output can be found on the help page for `nlm`. We note
208 briefly that `frog$minimum` is the negative log-likelihood value at the MLEs,
209 which are stored in the `frog$estimate` component of the list. The hessian is
210 the observed Fisher information matrix, which can be inverted to obtain the
211 variance-covariance matrix using the commands:

```

212 solve(frog$hessian)

```

213 It is worth drawing attention to the fact that the estimates are different
214 than the Bayesian estimates reported in the previous chapter (section XYZ)!!!
215 How can that be?! There are several reasons for this. First Bayesian inference
216 is based on the posterior distribution and it is not generally the case that the
217 MLE should correspond to any particular value of the posterior distribution. If
218 the prior distributions in a Bayesian analysis are uniform, then the mode of the
219 posterior is the MLE, but note that Bayesians almost always report posterior
220 means and so there will typically be a discrepancy there. Secondly, we have
221 implemented an approximation to the integral here and there might be a slight
222 bit of error induced by that. We will evaluate that shortly. Third, the Bayesian
223 analysis by MCMC is subject to some amount of Monte Carlo error which the
224 analyst should always be aware of in practical situations. All of these different
225 explanations are likely responsible for some of the discrepancy. Accounting
226 for these, as a practical matter, we see general consistency between the two
227 estimates.

228 To compute the integrated likelihood we used a discrete representation of
229 the state-space so that the integral could be approximated as a summation
230 over possible values of `s` with each value being weighted by its probability of

occurring, which is $1/nG$ under the assumption that s is uniform on the state-space \mathcal{S} . In chapter 4 we used a discrete state-space in developing a Bayesian analysis of the model in order to be able to modify the state-space in a flexible manner. Bayesian analysis requires simulation of the point process conditional on the observations, and this can be a difficult task when the state-space is continuous but has irregular geometry. Conversely, if the state-space is a regular polygon then Bayesian analysis by MCMC is possibly more efficient with a continuous state-space. We emphasize that the state-space is a part of the model. In some cases there wont be a natural choice of state space beyond “some large rectangle containing the trap grid” and, in such cases, for regular detection functions the estimate of density is invariant to the size of the state-space (i.e., the buffer) as long as it is sufficiently large. However if there are good reasons to restrict the state-space, it will tend to have an influence on the likelihood and hence AIC and so forth. As an illustration, lets do that by changing the state space here.....Use my polygon clipping stuff

In summary, we note that, for the basic SCR model, integrated likelihood is a really easy calculation when N is known. Even for N unknown it is not too difficult, and we will do that shortly. However, if you can solve the known- N problem then you should be able to do a real analysis, for example by considering different values of N and computing the results for each value and then making a plot of the log-likelihood or AIC and choosing the value of N that produces the best log likelihood or AIC. As a homework problem we suggest that the reader take the code given above and try to estimate N without modifying the code by just repeatedly calling that code for different values of N and trying to deduce the best value. Nevertheless, we will formalize the unknown- N problem shortly. We note that the software package **DENSITY** (?) implements certain types of SCR models using integrated likelihood methods. **DENSITY** has been made into an **R** package called **secr** (?) and we provide an analysis of some data using **secr** shortly along with a discussion of its capabilities.

1.2 MLE when N is Unknown

Here we build on the previous introduction to integrated likelihood but we consider now the case in which N is unknown. We will see that adapting the analysis based on the N -known model is really straightforward for the more general problem. The main distinction is that we dont observe the all-zero encounter history so we have to make sure we compute the probability for that encounter history which we do by tacking a row of zeros onto the encounter history matrix. In addition, we include the number of such all-zero encounter histories as an unknown parameter of the model. Call that unknown quantity n_0 . In addition, we have to be sure to include a combinatorial term to account for the fact that of the n observed individuals there are $\binom{N}{n}$ ways to realize a sample of size n . The combinatorial term involves the unknown n_0 and thus it must be included in the likelihood (we have already done that in the previous likelihood construction even though we didnt need it).

274 To summarize, when N is unknown, the n observed encounter histories have
 275 a multinomial distribution with probabilities $\pi(i)$ and sample size N^1 . The last
 276 cell the zero cell is computed by carrying out the integral in expression XYZ
 277 above for the all-zero encounter history and we have to account for the fact that
 278 there are $n_0 = N - n$ such encounter histories.

279 To analyze a specific case, we'll read in our fake data set (simulated using the
 280 parameters given above). To set some things up in our workspace we do this:

```
281 simSCR0.fn( blah blah blah)
282
283 > y<-read.csv("ind_by_trap.csv")[, -1]
284 > y<-as.matrix(y)
285 >
286 > X<-read.csv("traplocs.csv")[, -1]
287 > X<-as.matrix(X)
```

288 Recall that these data were generated with $N=100$, on an 8×8 unit state-
 289 space representing the trap locations (X) buffered by 2 units. As before, the
 290 likelihood is defined in the R workspace as an R function which takes an argu-
 291 ment being the unknown parameters of the model and additional arguments as
 292 prescribed. In particular, as before, we provide the encounter history matrix y ,
 293 the trap locations traplocs , the spacing of the integration grid (δ) and the
 294 state-space buffer. Here is the new likelihood function:

```
295 intlik2<-function(parm,y=y,delta=.3,X=traplocs,ssbuffer=2){
296
297   Xl<-min(X[,1]) -ssbuffer
298   Xu<-max(X[,1]) + ssbuffer
299   Yu<-max(X[,2]) + ssbuffer
300   Yl<-min(X[,2]) - ssbuffer
301
302   #delta<- (Xu-Xl)/npix
303   xg<-seq(Xl+delta/2,Xu-delta/2,delta)
304   yg<-seq(Yl+delta/2,Yu-delta/2,delta)
305   npix.x<-length(xg)
306   npix.y<-length(yg)
307   area<- (Xu-Xl)*(Yu-Yl)/((npix.x)*(npix.y))
308   G<-cbind(rep(xg,npix.y),sort(rep(yg,npix.x)))
309   nG<-nrow(G)
310   D<- e2dist(X,G)
311
312   alpha0<-parm[1]
313   alpha1<-parm[2]
314   n0<-exp(parm[3])
```

¹Maybe you could show an alternative simulation script to generate data using the `rmulti-`
`nom` function. This would make it a little more clear for people


```

315 probcap<- plogis(alpha0-alpha1*D)
316 Pm<-matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
317 ymat<-rbind(y,rep(0,ncol(y)))
318
319 lik.marg<-rep(NA,nrow(ymat))
320 for(i in 1:nrow(ymat)){
321   Pm[1:length(Pm)]<- (dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],log=TRUE))
322   lik.cond<- exp(t(Pm)%*%rep(1,nrow(probcap)))
323   lik.marg[i]<- sum( lik.cond*(1/nG) )
324 }
325 nv<-c(rep(1,length(lik.marg)-1),n0)
326 part1<- lgamma(nrow(y)+n0+1) - lgamma(n0+1)
327 part2<- sum(nv*log(lik.marg))
328   -1*(part1+ part2)
329 }

```

330 To execute this function for the data that we just read into the workspace,
 331 we execute the following command (saving the result in our friend `frog`).

```

332 > frog<-nlm(intlik2,c(-2.5,2,log(4)),hessian=TRUE,y=y,X=traplocs,delta=.2,ssbuffer=2)

```

333 This results in the usual output, including the parameter estimates, the
 334 gradient, and the numerical Hessian which is useful for obtaining asymptotic
 335 standard errors (omitted here).

```

336 > frog
337 $minimum
338 [1] 181.4657
339
340 $estimate
341 [1] -2.629342  1.790849  3.997268
342
343 [. Additional output deleted .]

```

344 The estimate of population size for the state-space (using the default state-
 345 space buffer) is

```

346 > nrow(y)+exp(3.997)
347 [1] 113.4346

```

348 Which differs from the data-generating truth ($N=100$) as we might expect.

349 1.2.1 Exercises

- 350 1. Run the analysis with different state-space buffers and comment on the result.
- 351 2. Conduct a brief simulation study using this code by simulating 100 data
- 352 sets and obtain the MLEs for each data set. Do things seem to be working as
- 353 you expect?

3. Further extensions: It should be straightforward to generalize the integrated likelihood function to accommodate many different situations. For examples, if we want to include more covariates in the model we can just add stuff to the probcap there, and add the relevant parameters to the argument that gets passed to the intlik function. For the simulated data, make up a covariate by generating a Bernoulli covariate (trap type perhaps baited or not baited) randomly and try to modify the likelihood to accommodate that.

4. We would probably be interested in devising the integrated likelihood for the full 3-d encounter history array so that we could include temporally varying covariates. This is not difficult but naturally will slow down the execution substantially. The interested reader should try to expand the capabilities of this basic R function.

1.2.2 Integrated Likelihood using the model under data augmentation

Note that this likelihood analysis is based on the standard likelihood in which N (or n_0) is an explicit parameter. This is usually called the joint likelihood or unconditional likelihood. We could also express the joint likelihood using data augmentation, replacing the parameter N with ψ . See Royle et al. (2007); Royle and Dorazio (2008, distance sampling example), etc.. Briefly, we note that the likelihood under data augmentation looks like a zero-inflated binomial mixture precisely as an occupancy type model (see Royle 2006). We think the interested reader could adapt the material from Royle and Dorazio (2008) with the R code given above for the likelihood and implement the likelihood analysis based on the model under data augmentation. Despite that we can carry-out likelihood analysis of models under data augmentation, we primarily advocate data augmentation for Bayesian analysis.

1.2.3 Extensions and Classical model selection and assessment

There are other types of covariates of interest: behavioral response, sex-specificity of parameters and all of these things. Some of these can be added directly to the likelihood if the covariate is fixed and known for all individuals captured or not. This excludes most covariates but it does include behavioral response. Sex-specificity is more difficult since sex is not known for uncaptured individuals. Trap-specific covariates such as trap type or status, or time-specific covariates such as date, are relatively easy to deal with (we leave these as exercises). We apply these various models in Chapter XXXX. To analyze such models, we do Bayesian analysis of the joint likelihood facilitated by the use of data augmentation. For covariates that are not fixed and known for all individuals, it is hard to do MLE for these based on the joint likelihood as we have developed above. Instead what people normally do is use what is colloquially referred to as the Huggins-Alho type model which is one of the approaches taken in the software package secr (Efford XYZ; see chapter XYZ). .

In most analyses, one is interested in choosing from among various potential models. A good thing about classical analysis based on likelihood is we can do rote application of AIC without thinking about anything. With distance as a covariate (e.g., distance sampling) this is usually applied to some arbitrary selection of distance functions. We don't recommend this. Given there is hardly ever (if at all) a rational science-based reason for choosing some particular distance function we believe that this standard approach will invariably lead to over-fitting. The fact that AIC is easy to compute does not mean that it should be abused in such fashions. Further discussion is made in chapters XYZ.

Goodness-of-fit In many analyses based on likelihood methods it is possible to cook-up fit statistics for which asymptotic distributions are known. In general, however, applied statisticians tend to adopt bootstrapping based on heuristically appealing fit statistics. An omnibus global GoF statistic is not so obvious but we can apply bootstrapping principles to SCR models directly which we discuss in chapter XYZ. Bayesian goodness-of-fit is almost always addressed with Bayesian p-values or some other posterior predictive check (REF XXX). Thus the approach whether Bayesian or classical is the same. We identify a fit statistic, we do a bootstrap (classical) or a Bayesian p-value. Royle et al. (2011) decomposed the fit problem into separate evaluations of the CSR hypothesis and the encounter process model. We discuss all of this in Chapter XYZ.

1.3 Likelihood analysis of the wolverine camera trapping data

Here we compute the MLEs for the wolverine data using an expanded version of the function we developed in the previous section. To accommodate that each trap might be operational a variable number of nights, we provided an additional argument to the likelihood function (allowing for a vector K), which requires also a modification to the construction of the likelihood (see Online Supplement). In addition, we had to accommodate that the state-space is a general rectangle, and we included a line in the code to compute the state-space area which we apply below for computing density. The more general function (intlik3) is given in the online supplement (shall we provide it here?).

The data were read into our R session and manipulated using the following commands. Note that we use the utility R function SCR23darray.fn which we defined in section XYZ.XYZ.

```
> wcaps<-source("wcaps.R")$value
> wtraps<-source("wtraps.R")$value
> K.wolv<-apply(wtraps[,4:ncol(wtraps)],1,sum)
>
> xx<-SCR23darray.fn(wcaps,ntraps=37,nperiods=165)
> y.wolv<- apply(xx,c(1,3),sum)
> traplocs.wolv<-wtraps[,2:3]
```

```

438 > traplocs.wolv<-traplocs.wolv/10000
439 >
440 > nlm(lik,c(-1.5,1.2,log(4)),hessian=TRUE,y=y.wolv,K=K.wolv,X=traplocs.wolv,
441 delta=.1,ssbuffer=2)$estimate
442
443 [1] -1.646692  3.128712  3.761974
444

```

We obtained the MLEs for a state-space buffer of 2 (standardized units) and for integration grid with spacing $\delta = .3, .2, .1, .05$. The MLES for these 4 cases are as follows:

```

448
449 > nlm(intlik3,c(-1.5,1.2,log(4)),hessian=TRUE,y=y.wolv,K=K.wolv, X=traplocs.wolv,delta=.3,ssbuffer=2)$estimate
450
451 [1] -1.654535  3.108126  3.723244
452
453 > nlm(lik,c(-1.5,1.2,log(4)),hessian=TRUE,,delta=.2,ssbuffer=2)$estimate
454
455 [1] -1.643023  3.133985  3.749195
456
457 > nlm(lik,c(-1.5,1.2,log(4)),hessian=TRUE,,delta=.1,ssbuffer=2)$estimate
458
459 [1] -1.646692  3.128712  3.761974
460
461 > nlm(lik,c(-1.5,1.2,log(4)),hessian=TRUE,,delta=.05,ssbuffer=2)$estimate
462
463 [1] -1.647191  3.128308  3.769455
464

```

We see the results change only slightly as the fineness of the integration grid increases. Conversely, the runtime on the platform of the day for the 4 cases was approximately 19s, 40s, 169s, and 723s which as we have suggested before could probably be regarded as relative to each other, across platforms, for gaging the decrease in speed as the fineness of the integration grid increases.

Next we studied the effect of the state-space buffer on the MLEs, using a fixed $\delta=.1$ for all analyses. We used state-space buffers of 1 to 4 units stepped by .5. This produced the following results, given here are the state-space buffer, area of the state-space, the MLE of N for the prescribed state-space and the corresponding MLE of density:

```

475
476 ### REDO ANALYSES
477
478      ssbuff      Ass      Nhat      Dhat
479 [1,]      1.0  66.98212  42.18630  0.6298144
480 [2,]      1.5  84.36242  52.12473  0.6178667

```

```

481 [3,]    2.0 103.74272  64.03329 0.6172317
482 [4,]    2.5 125.12302  77.36792 0.6183348
483 [5,]    3.0 148.50332  91.99139 0.6194568
484 [6,]    3.5 173.88362 107.87487 0.6203855
485 [7,]    4.0 201.26392 125.01359 0.6211426
486

```

487 The estimates of D stabilize rapidly² and the results suggest that wolverine
 488 density is around 0.62 individuals per 100 square km (recall that a unit is 10 x
 489 10 km). This is about 6.2 individuals per thousand square km which compares
 490 with XYZ reported in Royle et al. (2011) based on a clipped state-space as
 491 described in section XYZ.

492 In a later chapter analysis by MLE is done with an irregular state-space.
 493 Therefore we will have to extend the R function again to accept as possible
 494 input a pre-defined state-space grid.

495 1.3.1 Exercises

- 496 1. Compute the 95% confidence interval for wolverine density, somehow.
- 497 2. Compute the AIC of this model and modify `intlik3` to consider alternative
 498 link functions (at least one additional) and compare the AIC of the different
 499 models and the estimates. Comment. [should we do that here?].

500 1.4 Program DENSITY and the R package secr

501 DENSITY is a software program developed by Efford et al. (2004) for fitting
 502 spatial capture-recapture models based mostly on classical maximum likelihood
 503 estimation and related inference methods. Efford (2011) has also released an
 504 R package named `secr`, that contains many of the functions within DENSITY
 505 but also incorporates new models and features. Here, we will focus on `secr` as
 506 it will continue to be developed, contains more functionality and is based in R.
 507 To install and run models in `secr`, you must download the package and load it in
 508 R.

```

509 >install.packages(secr)
510 >library(secr)

```

511 SECR allows the user to simulate data and fit a suite of models with various
 512 detection functions and covariate responses. SECR uses the standard R model
 513 specification framework using tildes. E.g., the model command is `secr.fit` and is
 514 generally written as

```

515 >secr.fit(capturedata, model = list(D~1, g0~1, sigma~1), buffer = 20000)

```

²not very convincing

where we have `g0 1` indicating the intercept model. To include covariates, this would be written as `g0 b` where `b` is a behavioral covariate. Possible predictors for detection probability include both pre-defined variables (e.g., `t` and `b` corresponding to time and behavior), and user-defined covariates of several kinds. The discussion of covariates is developed in chapter XX(8).

Before we can fit the models, the data must first be entered into SECR. Two input files are required: trap layout (location and identification information for each trap) and capture data (e.g., sampling session, animal identification, trap day, and trap location). SECR requires that you specify the trap type, the two most common for camera trapping/hair snares are proximity detectors and count detectors. The ‘proximity’ detector type allows, at most, one detection of each individual at a particular detector on any occasion. The count detector designation allows repeat encounters of each individual at a particular detector on any occasion. There are other detector types that one can select such as: ‘polygon’ detector type which allows for a trap to be a sampled polygon, e.g., scat surveys, and ‘signal’ detector which allows for traps that have a strength indicator, e.g., acoustic arrays. The detector types `single` and `multi` can be confusing as `multi` seems like it would appropriate for something like a camera trap, but instead these two designations refer to traps that retain individuals, thus precluding the ability for animals to be captured in other traps during the sampling occasion. The `single` type indicates trap that can only catch one animal at a time, while `multi` indicates traps that may catch more than one animal at a time. For a full review of the detector types, one should look at the help manual, which can be accessed in R after installing the SECR package by using the command:

```
>RShowDoc("secr-manual", package = "secr")
```

As with all of the `scr` models, SECR fits a detection function relating the probability of detection to the distance of a detector from an individual activity center. SECR allows the user to specify one of a variety of detection functions including the commonly used half-normal, hazard rate, and exponential. There are 12 different functions, but some are only available for simulating data, and one should take caution when using different detection functions as the interpretation of the parameters, such as `sigma`, may not be consistent across formulations. The different detection functions are defined in the `secr` manual and can be found by calling the help function for the detection function:

```
> ?detectfn
```

It is useful to note that `secr` requires the buffer distance to be defined in meters and density will be returned as number of animals per hectare. Thus to make comparisons between `secr` and other models, we will often have to convert the density to the same units. Also, note that `sigma` is returned in units of meters.

3

³One question: SECR only ever reports `sigma`. What exactly is `sigma`? It is a scale parameter of a detection function and all detection functions have a scale parameter. But in

1.4.1 Analysis using secr package

To demonstrate the use of the secr package, we will show how to do the same analysis on the wolverine study as shown in section 4.6. To use the secr package, the data need to be formatted in a similar but slightly different manner than we use in WinBUGS. After installing the secr package, we first have to read in the trap locations and other related information, such as if the trap is operational during a sampling occasion. The secr package reads in the trap data through a command called `read.traps`, which requires the detector type as input. The detector type is important because it will determine the likelihood that secr will use to fit the model. Here, we have selected proximity since individuals are captured at most once in each trap during each sampling occasion.

```
>traps= read.csv(wtraps.csv)
>colnames(traps)[1:3]<- c("trapID","x", "y") #name the first 3 columns
# to match the secr nomenclature
>trapfile <- read.traps(data = traps, detector = "proximity")
```

After reading in the data, we now need to create the encounter matrix or array. The secr package does this through the use of the `make.caphist` command, where we provide the capture histories in raw data format (each line contains the session, identification number, occasion, and trap id for only 1 individual). This is the format that was shown in the data input file `wcaps`, and we only need a line or two to organize the data into the order that the `make.caphist` command wants. In creating the capture history, we provide also the trapfile with the trap information, and the format (e.g., here `fmt= trapID`) so that secr knows how to match the encounters to the trap, and finally, we provide the number of occasions.

```
>wolv.dat <- wcaps[,c(2, 3, 1)] #NEED TO UPDATE THIS WHEN I GET THE FILES, I JUST GUESSED AT THE
>wolv.dat <- cbind(rep(1, dim(wolv.dat)[1], wolv.dat)
>colnames(wolv.dat) <- c("Session", "ID", "Occasion", "trapID")
>wolvcapt=make.caphist(wolv.dat, trapfile, fmt = "trapID", noccasions = 165)
```

Calling the `secr.fit` command, will run the model. We are using the basic model (SCR0), so we do not need to make any specifications in the command line except for the providing the buffer size (in m). To specify different models, you can change the default `D 1`, `g0 1`, `sigma 1`, which the interested reader can do with very little difficulty.

```
> wolv.secr=secr.fit(wolvcapt, model = list(D~1, g0~1, sigma~1), buffer = 20000)
>wolv.secr
```

what sense is this sigma parameter related to home range diameter? Efford doesnt explain this, does he? In some sections in chapter 4 or possibly 6 we get into this issue.

```

596
597 secr.fit( capthist = wolvcapt, buffer = 20000, binomN = 1 )
598 secr 2.0.0, 18:26:39 05 Jul 2011
599
600 Detector type      proximity
601 Detector number   37
602 Average spacing   4415.693 m
603 x-range           593498 652294 m
604 y-range           6296796 6361803 m
605 N animals         : 21
606 N detections      : 115
607 N occasions       : 165
608 Mask area         : 1037069 ha
609
610 Model              : D~1 g0~1 sigma~1
611 Fixed (real)      : none
612 Detection fn       : halfnormal
613 Distribution       : poisson
614 N parameters       : 3
615 Log likelihood     : -746.754
616 AIC                : 1499.508
617 AICc               : 1500.920
618
619 Beta parameters (coefficients)
620           beta      SE.beta      lcl      ucl
621 D      -9.749576 0.23027860 -10.200913 -9.298238
622 g0     -4.275736 0.15846104  -4.586313 -3.965158
623 sigma  8.699202 0.07868944   8.544973  8.853430
624
625 Variance-covariance matrix of beta parameters
626           D           g0          sigma
627 D      0.053028233  0.000546922 -0.005226926
628 g0      0.000546922  0.025109900 -0.005885213
629 sigma -0.005226926 -0.005885213  0.006192027
630
631 Fitted (real) parameters evaluated at base levels of covariates
632           link      estimate SE.estimate      lcl      ucl
633 D      log 5.831941e-05 1.360973e-05 3.713638e-05 9.158548e-05
634 g0     logit 1.371121e-02 2.142902e-03 1.008756e-02 1.861207e-02
635 sigma  log 5.998124e+03 4.727205e+02 5.140849e+03 6.998355e+03

```

Under the fitted (real) parameters, we find D, the density, given in units of individuals/hectare (1 hectare = 100 m²). To convert this into individuals/1000km², we multiply by 100000, thus our density estimate is 5.83 individuals/1000 km². Sigma is given in units of meters, to convert to kilometers, we divide by 1000, which puts sigma at 5.99 km. Both of these estimates are very

similar to those provided in section 4.6 for the buffer size equal to 20 km. As an exercise, run this analysis for 30 and 40 km buffers and compare those found in section 4.6 under WinBUGS. NOTE: The `secr.fit` will return a warning when the buffer size appears to be too small. This is useful particularly with the different units being used between programs and packages.

1.5 Summary and Outlook

In this chapter, we showed that classical analysis of SCR models based on likelihood methods is a relatively simple proposition. Analysis is based on the so-called integrated likelihood in which the individual activity centers (random effects) are removed from the conditional-on-s likelihood by integration. We showed how to construct the integrated likelihood and fit some simple models in the R programming language. In addition, likelihood analysis for some broad classes of SCR models can be accomplished in the software package `DENSITY` (and other packages such as `ADMB`) or in the equivalent R library `secr` which we provided an illustration of here. In later chapters we provide more detailed analyses of SCR data using the `secr` package.

To compute the integrated likelihood we have to precisely describe the state-space of the underlying point process. In practice, this leads to a buffer around the trap array. We note that this is not really a buffer strip in the sense of Wilson et al. (XYZ) which is a feature of the analysis but it is somewhat more general here. In particular, it establishes the support of the integrand which we generally require to compute any integral. It might be that the integrand itself is finite even if the support is infinity but that may or may not be the case depending on the choice of detection function. As a practical matter then, it will typically be the case that, while estimates of N increase with the size of the buffer, estimates of density stabilize. This is not a feature of the classical methods based on using model M_0 or model M_h and buffering the trap array.

Why or why not use likelihood inference exclusively? For certain specific models, it is probably more computationally efficient to produce MLEs. However, WinBUGS is extremely flexible in terms of describing models, although it sometimes can be quite slow. We can devise models in WinBUGS easily that we cannot fit in `secr`. E.g., random individual effects of various types (see next chapter), we can handle missing covariates in complete generality and seamlessly, and impose arbitrary distributions on random variables. Moreover, models can easily be adapted to include auxiliary data types. For example, we might have camera trapping and genetic data and we can describe the models directly in WinBUGS and fit a joint model. For the MLE we have to write a custom new piece of code for each model or hope someone has done it for us. Later we consider open population models which are straightforward to develop in WinBUGS but, so far, there is no available platform for doing MLE although we imagine one could develop this. . Another thing that is more conceptual here is non-CSR point processes (see chapter XYZ) and generating predictions of how many individuals have home range centers in any particular polygon.

684 Basic benefits of Bayesian analysis have been discussed elsewhere (Chapter 2?
685 BPA book? Link and Barker?) and we believe these are compelling. On the
686 other hand, likelihood analysis makes it easy to do model-selection by AIC.
687 Goodness-of-fit is probably no more difficult or easy under either paradigm (see
688 next chapter?).

689 In summary, basic SCR models are easy to implement by either likelihood
690 or Bayesian methods but we feel that the typical user will realize much more
691 flexibility in model development using existing platforms for Bayesian analysis.
692 While these tend to be slow (sometimes excruciatingly slow), this will probably
693 not be an impediment in most problems, especially at some near point in the
694 future. Since we spent a lot of time here talking about specific technical details
695 on how to implement likelihood analysis of SCR models, we provided a corre-
696 sponding treatment in the next chapter on how to devise MCMC algorithms
697 for SCR models. This is a bit more tedious and requires more coding, but is
698 not technically challenging (except perhaps to develop highly efficient algorithms
699 which we don't excel at).

Bibliography

- Borchers, D. L., Buckland, S. T., and Zucchini, W. (2002), *Estimating animal abundance: closed populations*, vol. 13, Springer Verlag.
- Magoun, A. J., Long, C. D., Schwartz, M. K., Pilgrim, K. L., Lowell, R. E., and Valkenburg, P. (2011), “Integrating motion-detection cameras and hair snags for wolverine identification,” *The Journal of Wildlife Management*, 75, 731–739.
- Royle, J. and Dorazio, R. (2008), *Hierarchical modeling and inference in ecology: the analysis of data from populations, metapopulations and communities*, Academic Press.