

LIKELIHOOD ANALYSIS OF SPATIAL CAPTURE-RECAPTURE MODELS

We have so far mainly focused on Bayesian analysis of spatial capture-recapture models. And, in the previous chapters we learned how to fit some basic spatial capture-recapture models using a Bayesian formulation of the models analyzed in **BUGS** engines including **WinBUGS** and **JAGS**. Despite our focus on Bayesian analysis, it is instructive to develop the basic concepts and ideas behind classical analysis based on likelihood methods and frequentist inference for SCR models. We recognized earlier (Chapt. 5) that SCR models are versions of binomial (or other) GLMs, but with random effects (i.e., GLMMs). Throughout statistics, such models are routinely analyzed by likelihood methods. In particular, likelihood analysis is based on the integrated or marginal likelihood in which the random effects are removed, by integration, from the conditional-on- \mathbf{s} likelihood (\mathbf{s} being the individual activity center). This has been the approach taken by Borchers and Efford (2008); Dawson and Efford (2009) and related papers. Therefore, in this chapter, we provide some conceptual and technical foundation for likelihood-based analysis of spatial capture-recapture models.

We will show here that it is straightforward to compute the maximum likelihood estimates (MLE) for SCR models by integrated likelihood. We develop the MLE framework using **R**, and we also provide a basic introduction to the **R** package **secr** (Efford, 2011) which does likelihood analysis of SCR models (see also the stand-alone program **DENSITY** (Efford et al., 2004)). To set the context for likelihood analysis of SCR models, we first analyze the SCR model when N is known because, in that case, analysis is no different at all than a standard GLMM. We generalize the model to allow for unknown N using both conventional ideas based on the “full likelihood” (e.g., Borchers et al., 2002) and also using a formulation based on data augmentation. We obtain the MLEs for the SCR model from the wolverine camera trapping study (Magoun et al., 2011) analyzed in previous chapters to compare/contrast the results.

6.1 MLE WITH KNOWN N

We noted in Chapt. 5 that, with N known, the basic SCR model is a type of binomial model with a random effect. For such models we can obtain maximum likelihood estimators of model parameters based on integrated likelihood. The integrated likelihood is based on the marginal distribution of the data y in which the random effects are removed by integration from the conditional-on- \mathbf{s} distribution of the observations. See Chapt. 2 for a review of marginal, conditional and joint distributions. Conceptually, any SCR model begins with a specification of the conditional-on- \mathbf{s} model $[y|\mathbf{s}, \boldsymbol{\alpha}]$ and we have a “prior distribution” for \mathbf{s} , say $[\mathbf{s}]$. Then, the marginal distribution of the data y is

$$[y|\boldsymbol{\alpha}] = \int_{\mathcal{S}} [y|\mathbf{s}, \boldsymbol{\alpha}][\mathbf{s}]d\mathbf{s}.$$

When viewed as a function of $\boldsymbol{\alpha}$ for purposes of estimation, the marginal distribution $[y|\boldsymbol{\alpha}]$ is often referred to as the *integrated likelihood*.

It is worth analyzing the simplest SCR model with known- N in order to understand the underlying mechanics and basic concepts. These are directly relevant to the manner in which many capture-recapture models are classically analyzed, such as model M_h , and individual covariate models (see Chapt. 4).

To develop the integrated likelihood for SCR models, we first identify the conditional-on- \mathbf{s} likelihood. The observation model for each encounter observation y_{ij} , for individual i and trap j , specified conditional on \mathbf{s}_i , is

$$y_{ij}|\mathbf{s}_i \sim \text{Binomial}(K, p_{ij}(\mathbf{x}_j, \mathbf{s}_i)) \quad (6.1.1)$$

where we have indicated the dependence of encounter probability, p_{ij} , on \mathbf{s} and parameters $\boldsymbol{\alpha}$ explicitly. For example, p_{ij} might be the Gaussian model given by

$$p_{ij} = \text{logit}^{-1}(\alpha_0) \exp(-\alpha_1 \|\mathbf{x}_j - \mathbf{s}_i\|^2)$$

where $\alpha_1 = 1/(2\sigma^2)$. The joint distribution of the data for individual i is the product of J such terms (i.e., contributions from each of J traps).

$$[\mathbf{y}_i|\mathbf{s}_i, \boldsymbol{\alpha}] = \prod_{j=1}^J \text{Binomial}(K, p_{ij}(\mathbf{x}_j, \mathbf{s}_i))$$

We note this assumes that encounter of individual i in each trap is independent of encounter in every other trap, conditional on \mathbf{s}_i . This is the fundamental property of the basic model SCR0. The marginal likelihood is computed by removing \mathbf{s}_i , by integration from the conditional-on- \mathbf{s} likelihood, so we compute:

$$[\mathbf{y}_i|\boldsymbol{\alpha}] = \int_{\mathcal{S}} [\mathbf{y}_i|\mathbf{s}_i, \boldsymbol{\alpha}][\mathbf{s}_i]d\mathbf{s}_i$$

In most SCR models, $[\mathbf{s}] = 1/A(\mathcal{S})$ where $A(\mathcal{S})$ is the area of the prescribed state-space \mathcal{S} (but see Chapt. 11 for alternative specifications of $[\mathbf{s}]$).

The joint likelihood for all N individuals, assuming independence of encounters among individuals, is the product of N such terms:

$$\mathcal{L}(\boldsymbol{\alpha}|\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_N) = \prod_{i=1}^N [\mathbf{y}_i|\boldsymbol{\alpha}]$$

We emphasize that two independence assumptions are explicit in this development: independence of trap-specific encounters within individuals and also independence among individuals. In particular, this would only be valid when individuals are not physically restrained or removed upon capture, and when traps do not “fill up.”

The key operation for computing the likelihood is solving a 2-dimensional integration problem. There are some general purpose **R** packages that implement a number of multi-dimensional integration routines including **adapt** (Genz et al., 2007) and **R2cuba** (Hahn et al., 2010). In practice, we won’t rely on these extraneous **R** packages (except see Chapt. 11 for an application of **R2cuba**) but instead will use perhaps less efficient methods in which we replace the integral with a summation over an equal area mesh of points on the state-space \mathcal{S} and explicitly evaluate the integrand at each point. We invoke the rectangular rule for integration here¹ in which we evaluate the integrand on a regular grid of points of equal area and compute the average of the integrand over that grid of points. Let $u = 1, 2, \dots, nG$ index a grid of nG points, \mathbf{s}_u , where the area of grid cells is constant, say A . In this case, the integrand, i.e., the marginal pmf of \mathbf{y}_i , is approximated by

$$[\mathbf{y}_i|\boldsymbol{\alpha}] = \frac{1}{nG} \sum_{u=1}^{nG} [\mathbf{y}_i|\mathbf{s}_u, \boldsymbol{\alpha}] \quad (6.1.2)$$

This is a specific case of the general expression that could be used for approximating the integral for any arbitrary distribution $[\mathbf{s}]$. The general case is

$$[\mathbf{y}|\boldsymbol{\alpha}] = \frac{A(\mathcal{S})}{nG} \sum_{u=1}^{nG} [y|\mathbf{s}_u, \boldsymbol{\alpha}][\mathbf{s}_u]$$

Under the uniformity assumption, $[\mathbf{s}] = 1/A(\mathcal{S})$ and thus the grid-cell area cancels in the above expression to yield Eq. 6.1.2. The rectangular rule for integration can be seen as an application of the Law of Total Probability for a discrete random variable \mathbf{s} , having nG unique values with equal probabilities $1/nG$.

6.1.1 Implementation (simulated data)

Here we will illustrate how to carry out this integration and optimization based on the integrated likelihood using simulated data (i.e., see Sec. 5.5). Using **simSCR0** we simulate data for 100 individuals and an array of 25 traps laid out in a 5×5 grid of traps having unit spacing. The specific encounter model is the Gaussian model. The 100 activity centers were simulated on a state-space defined by an 8×8 square within which the trap array was centered (thus the trap array is buffered by 2 units). Therefore, the density of individuals in this system is fixed at $100/64$. In the following set of **R** commands we generate the data and then harvest the required data objects:

```
## simulate a complete data set (perfect detection)
> data <- simSCR0(discard0=FALSE, rnd=2013)
## extract the objects that we need for analysis
> y <- data$Y
```

¹e.g., http://en.wikipedia.org/wiki/Rectangle_method

```

5790 > traplocs <- data$traplocs
5791 > nind <- nrow(y) ## in this case nind=N
5792 > J <- nrow(traplocs)
5793 > K <- data$K
5794 > xlim <- data$xlim
5795 > ylim <- data$ylim

```

Now, we need to define the integration grid, say **G**, which we do with the following set of **R** commands (here, **delta** is the grid spacing):

```

5798 > delta <- .2
5799 > xg <- seq(xlim[1]+delta/2,xlim[2]-delta/2,by=delta)
5800 > yg <- seq(ylim[1]+delta/2,ylim[2]-delta/2,by=delta)
5801 > npix <- length(xg) # valid for square state-space only
5802 > G <- cbind(rep(xg,npix),sort(rep(yg,npix)))
5803 > nG <- nrow(G)

```

In this case, the integration grid is set up as a grid with spacing $\delta = 0.2$ which produces, for our example, a 40×40 grid of points for evaluating the integrand if the state-space buffer is set at 2. We note that the integration grid is set-up here to correspond exactly to the state-space used in simulating the data. However, in practice, we wouldn't know this, and our estimate of N (for the unknown case, see below) would be sensitive to choice of the extent of the integration grid. As we've discussed previously, density, which is N standardized by the area of the state-space, will not be so sensitive in most cases.

We are now ready to compute the conditional-on-**s** likelihood and carry out the marginalization described by Eq. 6.1.2. We need to do this by defining an **R** function that computes the likelihood for the integration grid, as a function of the data objects **y** and **traplocs** which were created above. However, it is a bit untidy to store the grid information in your workspace, and define the likelihood function in a way that depends on these things that exist in your workspace. Therefore, we build the **R** function so that it computes the integration grid *within* the function, thereby avoiding potential problems if our trapping grid locations change, or if we want to modify the state-space buffer easily. We therefore define the function, called **intlik1**, to which we pass the data objects and other information necessary to compute the marginal likelihood. This function is available in the **scrbook** package (use **?intlik1** at the **R** prompt). The code is reproduced here:

```

5822 intlik1 <- function(parm,y=X,X=traplocs, delta=.2, ssbuffer=2){
5823
5824   Xl <- min(X[,1]) - ssbuffer ## These lines of code are setting up the
5825   Xu <- max(X[,1]) + ssbuffer ## support for the integration which is
5826   Yu <- max(X[,2]) + ssbuffer ## the same as the state-space of "s"
5827   Yl <- min(X[,2]) - ssbuffer
5828   xg <- seq(Xl+delta/2,Xu-delta/2,,length=npix)
5829   yg <- seq(Yl+delta/2,Yu-delta/2,,length=npix)
5830   npix<- length(xg)
5831
5832   G <- cbind(rep(xg,npix),sort(rep(yg,npix)))

```

```

5833   nG <- nrow(G)
5834   D <- e2dist(X,G)
5835
5836   alpha0 <- parm[1]
5837   alpha1 <- exp(parm[2]) # alpha1 restricted to be positive here
5838
5839   probcap <- plogis(alpha0)*exp(-alpha1*D*D)
5840   Pm <- matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
5841           # Frequency of all-zero encounter histories
5842   n0 <- sum(apply(y,1,sum)==0)
5843           # Encounter histories with at least 1 detection
5844   ymat <- y[apply(y,1,sum)>0,]
5845   ymat <- rbind(ymat,rep(0,ncol(ymat)))
5846   lik.marg <- rep(NA,nrow(ymat))
5847
5848   for(i in 1:nrow(ymat)){
5849       ## Next line: log conditional likelihood for ALL possible values of s
5850       Pm[1:length(Pm)] <- dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],
5851                                   log=TRUE)
5852       ## Next line: sum the log conditional likelihoods, exp() result
5853       ## same as taking the product
5854       lik.cond <- exp(colSums(Pm))
5855       ## Take the average value == computing marginal
5856       lik.marg[i] <- sum( lik.cond*(1/nG))
5857   }
5858   ## n0 = number of all-0 encounter histories
5859   nv <- c(rep(1,length(lik.marg)-1),n0)
5860   return( -1*( sum(nv*log(lik.marg)) ) )
5861 }

```

5862 We emphasize that this function (and subsequent) are not meant to be general-purpose
5863 routines for solving all of your SCR problems but, rather, they are meant for illustrative
5864 purposes – so you can see how the integrated likelihood is constructed and how we connect
5865 it to data and other information that is needed.

5866 The function `intlik1` accepts as input the encounter history matrix, `y`, the trap loca-
5867 tions, `X`, and the state-space buffer. This allows us to vary the state-space buffer and easily
5868 evaluate the sensitivity of the MLE to the size of the state-space. Note that we have a
5869 peculiar handling of the encounter history matrix `y`. In particular, we remove the all-zero
5870 encounter histories from the matrix and tack-on a single all-zero encounter history as the
5871 last row which then gets weighted by the number of such encounter histories (`n0`). This is
5872 a bit long-winded and strictly unnecessary when N is known, but we did it this way be-
5873 cause the extension to the unknown- N case is now transparent (as we demonstrate in the
5874 following section). The matrix `Pm` holds the log-likelihood contributions of each encounter
5875 frequency for each possible state-space location of the individual. The log contributions
5876 are summed up and the result exponentiated on the next line, producing `lik.cond`, the
5877 conditional-on-`s` likelihood (Eq. 6.1.1 above). The marginal likelihood (`lik.marg`) sums
5878 up the conditional elements weighted by the probabilities [`s`] (Eq. 6.1.2 above).

5879 This is a fairly primitive function which doesn't allow much flexibility in the data
 5880 structure. For example, it assumes that K , the number of replicates, is constant for each
 5881 trap. Further, it assumes that the state-space is a square. We generalize this to some
 5882 extent later in this chapter.

5883 Here is the **R** command for maximizing the likelihood using **nlm** (the function **optim**
 5884 could also be used) and saving the results into an object called **frog**. The output is a list
 5885 of the following structure and these specific estimates are produced using the simulated
 5886 data set:

```
5887 # should take 15-30 seconds
5888
5889 > starts <- c(-2,2)
5890 > frog <- nlm(intlik1,starts,y=y,X=traplocs,delta=.1,ssbuffer=2,hessian=TRUE)
5891 > frog
5892
5893 $minimum
5894 [1] 297.1896
5895
5896 $estimate
5897 [1] -2.504824  2.373343
5898
5899 $gradient
5900 [1] -2.069654e-05  1.968754e-05
5901
5902 $hessian
5903           [,1]      [,2]
5904 [1,]  48.67898 -19.25750
5905 [2,] -19.25750  13.34114
5906
5907 $code
5908 [1] 1
5909
5910 $iterations
5911 [1] 11
```

5912 Details about this output can be found on the help page for **nlm**. We note briefly that
 5913 **frog\$minimum** is the negative log-likelihood value at the MLEs, which are stored in the
 5914 **frog\$estimate** component of the list. The order of the parameters is as they are defined
 5915 in the likelihood function so, in this case, the first element (value = -2.504824) is the
 5916 logit transform of p_0 and the second element (value = 2.373343) is the value of α_1 the
 5917 “coefficient” on distance-squared. The Hessian is the observed Fisher information matrix,
 5918 which can be inverted to obtain the variance-covariance matrix using the command:

```
5919 > solve(frog$hessian)
```

5920 It is worth drawing attention to the fact that the estimates are slightly different than
 5921 the Bayesian estimates reported previously in Sec. 5.6. There are several reasons for this.
 5922 First Bayesian inference is based on the posterior distribution and it is not generally the

case that the MLE should correspond to any particular value of the posterior distribution. If the prior distributions in a Bayesian analysis are uniform, then the (multivariate) mode of the posterior is the MLE, but note Bayesians almost always report posterior *means* and so there will typically be a discrepancy there. Secondly, we have implemented an approximation to the integral here and there might be a slight bit of error induced by that. We will evaluate that shortly. Third, the Bayesian analysis by MCMC is itself subject to some amount of Monte Carlo error which the analyst should always be aware of in practical situations. All of these different explanations are likely responsible for some of the discrepancy. Accounting for these, we see general consistency between the two estimates.

In summary, for the basic SCR model, computing the integrated likelihood is a simple task when N is known. Even for N unknown it is not too difficult, and we will do that shortly. However, if you can solve the known- N problem then you should be able to do a real analysis, for example by considering different values of N and computing the results for each value and then making a plot of the log-likelihood or AIC and choosing the value of N that produces the best log-likelihood or AIC. As a homework problem we suggest that you can take the code given above and try to estimate N without modifying the code by just repeatedly applying it for different values of N in attempt to deduce the best value. We will formalize the unknown- N problem next.

6.2 MLE WHEN N IS UNKNOWN

Here we build on the previous introduction to integrated likelihood but we consider now the case in which N is unknown. We will see that adapting the analysis based on the known- N model is straightforward for the more general problem. The main distinction is that we don't observe the all-zero encounter history so we have to make sure we compute the probability for that encounter history, which we do by tacking a row of zeros onto the encounter history matrix. In addition, we include the number of such all-zero encounter histories (that is, the number of individuals *not* encountered) as an unknown parameter of the model. Call that unknown quantity n_0 , so that $N = n_0 + n$ where n is the number of unique individuals encountered. We will usually parameterize the likelihood in terms of n_0 because optimization over a parameter space in which $\log(n_0)$ is unconstrained is preferred to a parameter space in which N must be constrained $N \geq n$. With n_0 unknown, we have to be sure to include a combinatorial term to account for the fact that, of the n observed individuals, there are $\binom{N}{n}$ ways to realize a sample of size n . The combinatorial term involves the unknown n_0 and thus it must be included in the likelihood. In evaluating the log-likelihood, we have to compute terms such as the log-factorial, $\log(N!) = \log((n_0 + n)!)$. We do this in **R** by making use of the log-gamma function (`lgamma`) and the identity

$$\log(N!) = \text{lgamma}(N + 1).$$

Therefore, to compute the likelihood, we require the following 3 components: (1) The marginal probability of each \mathbf{y}_i as before,

$$[\mathbf{y}_i | \boldsymbol{\alpha}] = \int_S [\mathbf{y}_i | \mathbf{s}_i, \boldsymbol{\alpha}] [\mathbf{s}_i] d\mathbf{s}_i.$$

5960 (2) We compute the probability of an all-0 encounter history:

$$\pi_0 = [\mathbf{y} = \mathbf{0} | \boldsymbol{\alpha}] = \int_{\mathcal{S}} \text{Binomial}(\mathbf{0} | \mathbf{s}_i, \boldsymbol{\alpha}) [\mathbf{s}_i] d\mathbf{s}_i$$

5961 (3) The combinatorial term: $\binom{N}{n}$. Then, the marginal likelihood has this form:

$$\mathcal{L}(\boldsymbol{\alpha}, n_0 | \mathbf{y}) = \frac{N!}{n!n_0!} \left\{ \prod_{i=1}^n [\mathbf{y}_i | \boldsymbol{\alpha}] \right\} \pi_0^{n_0}. \quad (6.2.1)$$

5962 This is discussed in Borchers and Efford (2008, p. 379) as the conditional-on- N form of the
5963 likelihood – we also call it the “binomial form” of the likelihood because of its appearance.

5964 Operationally, things proceed much as before: We compute the marginal probability
5965 of each observed \mathbf{y}_i , i.e., by removing the latent \mathbf{s}_i by integration. In addition, we com-
5966 pute the marginal probability of the “all-zero” encounter history \mathbf{y}_{n+1} , and make sure to
5967 weight it n_0 times. We accomplish this by “padding” the data set with a single encounter
5968 history having $y_{n+1,j} = 0$ for all traps $j = 1, 2, \dots, J$. Then we be sure to include the
5969 combinatorial term in the likelihood or log-likelihood computation. We demonstrate this
5970 shortly. To analyze a specific case, we’ll simulate our fake data set (simulated using the
5971 parameters given above). To set some things up in our workspace we do this:

```
5972 ## Obtain a simulated data set
5973 > data <- simSCRO(discard0=TRUE, rnd=2013)
5974
5975 ## Extract the items we need for analysis
5976 > y <- data$Y
5977 > nind <- nrow(y)
5978 > traplocs <- data$traplocs
5979 > J <- nrow(traplocs)
5980 > K <- data$K
```

5981 Recall that these data are simulated by default with $N = 100$, on an 8×8 unit state-
5982 space representing the trap locations buffered by 2 units, although you can modify the
5983 simulation script easily.

5984 As before, the likelihood is defined in the **R** workspace as an **R** function, `intlik2`,
5985 which takes an argument being the unknown parameters of the model and additional
5986 arguments as prescribed. In particular, we provide the encounter history matrix **y**, the
5987 trap locations `traplocs`, the spacing of the integration grid (argument `delta`) and the
5988 state-space buffer. Here is the new likelihood function:

```
5989 intlik2 <- function(parm,y=y,X=traplocs,delta=.3,ssbuffer=2){
5990
5991   Xl <- min(X[,1]) - ssbuffer
5992   Xu <- max(X[,1]) + ssbuffer
5993   Yu <- max(X[,2]) + ssbuffer
5994   Yl <- min(X[,2]) - ssbuffer
5995
5996   xg <- seq(Xl+delta/2,Xu-delta/2,delta)
```



```

5997   yg <- seq(Yl+delta/2,Yu-delta/2,delta)
5998   npix.x <- length(xg)
5999   npix.y <- plength(yg)
6000   area <- (Xu-Xl)*(Yu-Yl)/((npix.x)*(npix.y))
6001   G <- cbind(rep(xg,npix.y),sort(rep(yg,npix.x)))
6002   nG <- nrow(G)
6003   D <- e2dist(X,G)
6004   # extract the parameters from the input vector
6005   alpha0 <- parm[1]
6006   alpha1 <- exp(parm[2])
6007   n0 <- exp(parm[3]) # note parm[3] lives on the real line
6008   probcap <- plogis(alpha0)*exp(-alpha1*D*D)
6009   Pm <- matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
6010   ymat <- rbind(y,rep(0,ncol(y)))
6011
6012   lik.marg <- rep(NA,nrow(ymat))
6013   for(i in 1:nrow(ymat)){
6014     Pm[1:length(Pm)] <- (dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],
6015                               log=TRUE))
6016     lik.cond <- exp(colSums(Pm))
6017     lik.marg[i] <- sum( lik.cond*(1/nG) )
6018   }
6019   nv <- c(rep(1,length(lik.marg)-1),n0)
6020   ## part1 here is the combinatorial term.
6021   ## math: log(factorial(N)) = lgamma(N+1)
6022   part1 <- lgamma(nrow(y)+n0+1) - lgamma(n0+1)
6023   part2 <- sum(nv*log(lik.marg))
6024   return( -1*(part1+ part2) )
6025 }

```

6026 To execute this function for the data that we created with `simSCR0`, we execute the
6027 following command (saving the result in our friend `frog`). This results in the usual output,
6028 including the parameter estimates, the gradient, and the numerical Hessian which is useful
6029 for obtaining asymptotic standard errors (see below):

```

6030 > starts <- c(-2.5,0,4)
6031 > frog <- nlm(intlik2,starts,hessian=TRUE,y=y,X=traplocs,delta=.2,ssbuffer=2)
6032
6033 Warning message:
6034 In nlm(intlik2, starts, hessian = TRUE, y = y, X = traplocs, delta = 0.2, :
6035   NA/Inf replaced by maximum positive value
6036
6037 > frog
6038 $minimum
6039 [1] 113.5004
6040
6041 $estimate

```

```
6042 [1] -2.538333  0.902807  4.232810
```

```
6043
```

```
6044 [... additional output deleted ...]
```

6045 Executing `nlm` here usually produces one or more **R** warnings due to numerical calculations
 6046 happening on extremely small or large numbers (calculation of p near the edge of the
 6047 state-space), and they also happen if a poor parameterization is used which produces
 6048 evaluations of the objective function beyond the boundary of the parameter space (e.g.,
 6049 $n_0 < 0$). Such numerical warnings can often be minimized or avoided altogether by picking
 6050 judicious starting values of parameters or properly transforming or scaling the parameters
 6051 but, in general, they can be ignored. You will see from the `nlm` output that the algorithm
 6052 performed satisfactory in minimizing the objective function. The estimate of population
 6053 size, \hat{N} , for the state-space (using the default state-space buffer) is

```
6054 > Nhat <- nrow(y) + exp(4.2328)   ### This is n + MLE of n0
```

```
6055 > Nhat
```

```
6056 [1] 110.9099
```

6057 Which differs from the data-generating value ($N = 100$), as we might expect for a single
 6058 realization. We usually will present an estimate of uncertainty associated with this MLE
 6059 which we can obtain by inverting the Hessian. Note that $\text{Var}(\hat{N}) = n + \text{Var}(\hat{n}_0)$. Since
 6060 we have parameterized the model in terms of $\log(n_0)$ we use the delta method² described
 6061 in Williams et al. (2002, Appendix F4) (see also Ver Hoef, 2012) to obtain the variance
 6062 on the scale of n_0 as follows:

```
6063 > (exp(4.2328)^2)*solve(frog$hessian)[3,3]
```

```
6064 [1] 260.2033
```

```
6065
```

```
6066 > sqrt(260)
```

```
6067 [1] 16.12452
```

6068 Therefore, the asymptotic “Wald-type” confidence interval for N is $110.91 \pm 1.96 \times 16.125 =$
 6069 $(79.305, 142.515)$. To report this in terms of density, we scale appropriately by the area
 6070 of the prescribed state-space which is 64 units of area (i.e., an 8×8 square). Our MLE
 6071 of D is $\hat{D} = 110.91/64 = 1.733$ individuals per square unit. To get the standard error
 6072 for \hat{D} we need to divide the SE for \hat{N} by the area of the state-space, and so $\text{SE}(\hat{D}) =$
 6073 $(1/64) \times 16.12452 = 0.252$.

6074 6.2.1 Integrated likelihood under data augmentation

6075 The likelihood analysis developed in the previous sections is based on the likelihood in
 6076 which N (or n_0) is an explicit parameter. This is usually called the “full likelihood” or
 6077 sometimes “unconditional likelihood” (Borchers et al., 2002) because it is the likelihood
 6078 for all individuals in the population, not just those which have been captured, i.e., not that
 6079 which is *conditional on capture*. It is also possible to express an alternative unconditional

² We found a good set of notes on the delta approximation on Dr. David Patterson’s ST549 notes: <http://www.math.umd.edu/patterson/549/Delta.pdf>

likelihood using data augmentation, replacing the parameter N with ψ (e.g., see Sec. 7.1.6 Royle and Dorazio, 2008, for an example). We don't go into detail here, but we note that the likelihood under data augmentation is a zero-inflated binomial mixture – precisely an occupancy type model (Royle, 2006). Thus, while it is possible to carry out likelihood analysis of models under data augmentation, we primarily advocate data augmentation for Bayesian analysis.

6.2.2 Extensions

We have only considered basic SCR models with no additional covariates. However, in practice, we are interested in covariate effects including “behavioral response”, sex-specificity of parameters, and potentially others. Some of these can be added directly to the likelihood if the covariate is fixed and known for all individuals captured or not. An example is a behavioral response, which amounts to having a covariate $x_{ik} = 1$ if individual i was captured prior to occasion k and $x_{ik} = 0$ otherwise. For uncaptured individuals, $x_{ik} = 0$ for all k . Royle et al. (2011b) called this a global behavioral response because the covariate is defined for all traps, no matter the trap in which an individual was captured. We could also define a *local* behavioral response which occurs at the level of the trap, i.e., $x_{ijk} = 1$ if individual i was captured in trap j prior to occasion k , etc... Trap-specific covariates such as trap type or status, or time-specific covariates such as date, are easily accommodated as well. As an example, Kéry et al. (2010) develop a model for the European wildcat *Felis silvestris* in which traps are either baited or not (a trap-specific covariate with only 2 values), and also encounter probability varies over time in the form of a quadratic seasonal response. We consider models with behavioral response or fixed covariates in Chapt. 7. The integrated likelihood routines we provided above can be modified directly for such cases, which we leave to the interested reader to investigate.

Sex-specificity is more difficult to deal with since sex is not known for uncaptured individuals (and sometimes not even for all captured individuals). To analyze such models, we do Bayesian analysis of the joint likelihood using data augmentation (Gardner et al., 2010b; Russell et al., 2012), discussed further in Chapt. 7. For such covariates (i.e., that are not fixed and known for all individuals), it is somewhat more challenging to do MLE based on the joint likelihood as we have developed above. Instead it is more conventional to use what is colloquially referred to as the “Huggins-Alho” type model which is one of the approaches taken in the software package `secr` (Efford, 2011). We introduce the `secr` package in Sec. 6.5 below.

6.3 CLASSICAL MODEL SELECTION AND ASSESSMENT

In most analyses, one is interested in choosing from among various potential models, or ranking models, or something else to do with assessing the relative merits of a set of models. A good thing about classical analysis based on likelihood is we can apply Akaike Information Criterion (AIC) methods (Burnham and Anderson, 2002) without difficulty. AIC is convenient for assessing the relative merits of these different models although if there are only a few models it is not objectionable to use hypothesis tests or confidence intervals to determine importance of effects. A second model selection context has to do with choosing among various detection models, although, as a general rule, we don't

recommend this application of model selection. This is because there is hardly ever (if at all) a rational subject-matter based reason motivating specific distance functions. As a result, we believe that doing too much model selection will invariably lead to over-fitting and thus over-statement of precision. This is the main reason that we haven't loaded you down with a basket of models for detection probability so far, although we discuss many possibilities in Chapt. 7.

Goodness-of-fit or model-checking – For many standard capture-recapture models, it is possible to identify goodness-of-fit statistics based on the multinomial likelihood, (Cooch and White, 2006, Chapt. 5), and evaluate model adequacy using formal statistical tests. Similar strategies can be applied to SCR models using expected cell-frequencies based on the marginal distribution of the observations. Also, because computing MLEs is somewhat more efficient in many cases compared to Bayesian analysis, it is sometimes feasible to use bootstrap methods. At the present time, we don't know of any applications of goodness-of-fit testing for SCR models based on likelihood inference, although we discuss the use of Bayesian p-values for assessing model fit in Chapt. 8. An important practical problem in trying to evaluate goodness-of-fit is that, in realistic sample sizes, fit tests often lack the power to detect departures from the model under consideration and so they may not be generally useful in practice.

6.4 LIKELIHOOD ANALYSIS OF THE WOLVERINE CAMERA TRAPPING DATA

Here we compute the MLEs for the wolverine data using an expanded version of the function we developed in the previous section. To accommodate that each trap might be operational a variable number of nights, we provided an additional argument to the likelihood function (allowing for a vector $\mathbf{K} = (K_1, \dots, K_J)$), which requires also a modification to the construction of the likelihood. In addition, we accommodate the state-space is a general rectangle, and we included a line in the code to compute the state-space area which we apply below for computing density. The more general function (`intlik3`) is given in the **R** package `scrbook`. Incidentally, this function also returns the area of the state-space for a given set of parameter values, as an attribute to the function value, which will be used in converting \hat{N} to \hat{D} . To use this function to obtain the MLEs for the wolverine camera trap study, we execute the following commands (note: these are in the help file and will execute if you type `example(intlik3)`:

```
> library(scrbook)
> data(wolverine)

> traps <- wolverine$wtraps
> traplocs <- traps[,2:3]/10000
> K.wolv <- apply(traps[,4:ncol(traps)],1,sum)

> y3d <- SCR23darray(wolverine$wcaps,traps)
> y2d <- apply(y3d,c(1,2),sum)

> starts <- c(-1.5,0,3)
```

```

6163 > wolv <- nlm(intlik3,starts,hessian=TRUE,y=y2d,K=K.wolv,X=traplocs,
6164               delta=.2,ssbuffer=2)
6165
6166 > wolv
6167 $minimum
6168 [1] 220.4313
6169
6170 $estimate
6171 [1] -2.8176120  0.2269395  3.5836875
6172
6173 [.... output deleted ....]
    
```

Of course we're interested in obtaining an estimate of population size for the prescribed state-space, or density, and associated measures of uncertainty which we do using the delta method (Williams et al., 2002, Appendix F4). To do all of that we need to manipulate the output of `nlm` since we have our estimate in terms of $\log(n_0)$. We execute the following commands:

```

6179 > wolv <- nlm(intlik3,starts,hessian=TRUE,y=y2d,K=K.wolv,X=traplocs,delta=.2,
6180               ssbuffer=2)
6181 > Nhat <- nrow(y2d)+exp(wolv$estimate[3])
6182 > area <- attr(intlik3(starts,y=y2d,K=K.wolv,X=traplocs,delta=.2,ssbuffer=2),
6183               "SSarea")
6184 > Dhat <- Nhat/area
6185
6186 > Dhat
6187 [1] 0.5494947
6188
6189 > SE <- (1/area)*exp(wolv$estimate[3])*sqrt(solve(wolv$hessian)[3,3])
6190
6191 > SE
6192 [1] 0.1087073
    
```

Our estimate of density is 0.55 individuals per “standardized unit” which is 100 km², because we divided UTM coordinates by 10000. So this is about 5.5 individuals per 1000 km², with a SE of around 1.09 individuals. This compares closely with 5.77 reported in Sec. 5.9 based on Bayesian analysis of the model.

6.4.1 Sensitivity to integration grid and state-space buffer

The effect of approximating the integral by a discrete mesh of points is that it induces some numerical error in evaluation of the integral and, further, that error increases as the coarseness of the mesh increases. To evaluate the effect (or sensitivity) of the integration grid spacing, we obtained the MLEs for a state-space buffer of 2 (standardized units) and for integration grid with spacing $\delta = .3, .2, .1, .05$. The MLEs for these 4 cases including the relative runtime are given in Table 6.1. We see the results change only slightly as the integration grid changes. Conversely, the runtime on the platform of the day for the 4 cases

increases rapidly. These runtimes could be regarded in relative terms, across platforms, for gaging the decrease in speed as the fineness of the integration grid increases.

Table 6.1. Runtime and MLEs for different integration grid resolutions for the wolverine camera trapping data.

| δ | Estimates | | | |
|----------|-------------|------------------|------------------|-----------------------|
| | runtime (s) | $\hat{\alpha}_0$ | $\hat{\alpha}_1$ | $\widehat{\log(n_0)}$ |
| 0.30 | 9.9 | -2.819786 | 1.258468 | 3.569731 |
| 0.20 | 32.3 | -2.817610 | 1.254757 | 3.583690 |
| 0.10 | 115.1 | -2.817570 | 1.255112 | 3.599040 |
| 0.05 | 407.3 | -2.817559 | 1.255281 | 3.607158 |

We studied the effect of the state-space buffer on the MLEs, using a fixed $\delta = .2$ for all analyses. We used state-space buffers of 1 to 4 units stepped by .5. As we can see (Table 6.2), the estimates of D stabilize rapidly and the incremental difference is within the numerical error associated with approximating the integral.

Table 6.2. Results of the effect of the state-space buffer on the MLE. Given here are the state-space buffer, area of the state-space (area), the MLE of N (\hat{N}) for the prescribed state-space and the corresponding MLE of density (\hat{D}).

| Buffer | Area | \hat{N} | \hat{D} |
|--------|-----------|-----------|-----------|
| 1.0 | 66.98212 | 37.73338 | 0.5633352 |
| 1.5 | 84.36242 | 46.21008 | 0.5477567 |
| 2.0 | 103.74272 | 57.00617 | 0.5494956 |
| 2.5 | 125.12302 | 69.03616 | 0.5517463 |
| 3.0 | 148.50332 | 82.17550 | 0.5533580 |
| 3.5 | 173.88362 | 96.44018 | 0.5546249 |
| 4.0 | 201.26392 | 111.83524 | 0.5556646 |

6.4.2 Using a habitat mask (Restricted state-space)

In Sec. 5.10 we used a discrete representation of the state-space in order to have control over its extent and shape. This makes it easy to do things like clip out non-habitat, or create a *habitat mask* which defines suitable habitat. Clearly that formulation of the model is relevant to the calculation of the marginal likelihood in the sense that the discrete state-space is equivalent to the integration grid. Thus, for example, we could easily compute the MLE of parameters under some model with a restricted state-space merely by creating the required state-space at whatever grid resolution is desired, and then inputting that state-space into the likelihood function above, instead of computing it within the function. We can easily create an explicit state-space grid for integration from arbitrary polygons or GIS shapefiles which we demonstrate here. Our approach is to create the integration grid (or state-space grid) outside of the likelihood evaluation, and then determine which points of the grid lie in the polygon defined by the shapefile using functions in the **R** packages **sp** and **maptools**. For each point in the state-space grid (object **G** in the code below which is

assumed to exist), we determine whether it is inside the polygon³, identifying such points with a value of `mask=1` and `mask=0` for points that are *not* in the polygon. We load the shapefile which originates by an application of the `readShapeSpatial` function. We have saved the result into an **R** data object called `SSp` which is in the `scrbook` package. Here are the **R** commands for doing this (see the helpfile `?intlik4`):

```

6230 > library(maptools)
6231 > library(sp)
6232 > library(scrbook)
6233
6234 ##### If we have the .shp file in place, we would use this command:
6235 ##### SSp <- readShapeSpatial('Sim_Polygon.shp')
6236 ##### The object SSp is in data(fakeshapefile)
6237 > data(fakeshapefile)
6238 > Pcoord <- SpatialPoints(G)
6239 > PinPoly <- over(Pcoord,SSp) ### determine if each point is in polygon
6240 > mask <- as.numeric(!is.na(PinPoly[,1])) ## convert to binary 0/1
6241 > G <- G[mask==1,]
    
```

We created the function `intlik4` which accepts the integration grid as an explicit argument, and this function is also available in the package `scrbook`.

We apply this modification to the wolverine camera trapping study. Royle et al. (2011b) created 2, 4 and 8 km state-space grids so as to remove “non-habitat” (mostly ocean, bays, and large lakes). We previously analyzed the model using **JAGS** and **WinBUGS** in Chapt. 5. To set up the wolverine data and fit the model using maximum likelihood we execute the following commands:

```

6249 > library(scrbook)
6250 > data(wolverine)
6251
6252 > traps <- wolverine$wtraps
6253 > traplocs <- traps[,2:3]/10000
6254 > K.wolv <- apply(traps[,4:ncol(traps)],1,sum)
6255
6256 > y3d <- SCR23darray(wolverine$wcaps,traps)
6257 > y2d <- apply(y3d,c(1,2),sum)
6258 > G <- wolverine$grid2/10000
6259
6260 > starts <- c(-1.5,0,3)
6261 > wolv <- nlm(intlik4, starts, y=y2d, K=K.wolv, X=traplocs, G=G)
    
```

³We perform this check using the `over` function. This function takes as its second argument (among others) an object of the class “SpatialPolygons” or “SpatialPolygonsDataFrame”, which can hold additional information for each polygon, and the output value of the function differs slightly for these two classes: if using a “SpatialPolygons” object, the function returns a vector of length equal to the number of points (e.g., in the example above), but if using a “SpatialPolygonsDataFrame” it returns a data frame (e.g., see Sec. 17.5 in Chapt. 17). If you use the `over` function, make sure you know the class of your second argument so that when processing the function output you index it correctly.

Table 6.3. Maximum likelihood estimates (MLEs) and asymptotic standard errors (SE) for the wolverine camera trapping data using 2, 4 and 8 km state-space grids.

| Grid | α_0 | α_1 | $\log(n_0)$ | N | SE | D(1000) | SE |
|------|------------|------------|-------------|-------|-------|---------|------|
| 2 | -3.00 | 1.27 | 4.11 | 81.98 | 16.31 | 8.31 | 1.65 |
| 4 | -2.99 | 1.34 | 4.16 | 84.88 | 16.76 | 8.57 | 1.69 |
| 8 | -3.05 | 1.08 | 4.06 | 78.89 | 15.31 | 7.85 | 1.52 |

```

6262
6263 > wolv
6264
6265 $minimum
6266 [1] 225.8355
6267
6268 $estimate
6269 [1] -2.9955424  0.2350885  4.1104757
6270
6271 [... some output deleted ...]

```

Next we convert the parameter estimates to estimates of total population size for the prescribed state-space, and then obtain an estimate of density (per 1000 km²) using the area computed as the number of pixels in the state-space grid, G , multiplied by the area per grid cell. In the present case (the calculation above) we used a state-space grid with 2 km \times 2 km pixels. Finally, we compute a standard errors using the delta approximation:

```

6277 > area <- nrow(G)*4
6278 # Nhat = n (observed) + MLE of n0 (not observed)
6279 > Nhat <- 21 + exp(wolv$estimate[3])
6280 > SE <- exp(wolv$estimate[3])*sqrt(solve(wolv$hessian)[3,3])
6281 > D <- (Nhat/(nrow(G)*area))*1000
6282 > SE.D <- (SE/(nrow(G)*area))*1000

```

We did this for each the 2 km, 4 km and 8 km state-space grids which produced the estimates summarized in Table 6.3. These estimates compare with the 8.6 (2 km grid) and 8.2 (8 km grid) reported in Royle et al. (2011b) based on a clipped state-space as described in Sec. 5.10.

6.5 DENSITY AND THE R PACKAGE SECR

DENSITY is a software program developed by Efford (2004) for fitting spatial capture-recapture models based mostly on classical maximum likelihood estimation and related inference methods. Efford (2011) has also released an **R** package called **secr**, that contains much of the functionality of **DENSITY** but also incorporates new models and features. Here, we briefly introduce the **secr** package which we prefer to use over **DENSITY**, because it allows us to remain in the **R** environment for data processing and summarization. We provide a brief introduction to **secr** and some of its capabilities here, and we also use

it for doing some analysis in other parts of this book. We believe that **secr** will be sufficient for many (if not most) of the SCR problems that one might encounter. It provides a flexible analysis platform, with a large number of summary features, and “publication ready” output. Its user-interface is clean and intuitive to **R** users, and it has been stable, efficient and reliable in the (fairly extensive) evaluations that we have done.

To install and run models in **secr**, you must download the package and load it in **R**.

```
> install.packages("secr")
> library(secr)
```

secr allows the user to simulate data and fit a suite of models with various detection functions and covariate responses. It also contains a number of helpful constructor functions for creating objects of the proper class that are recognized by other **secr** functions. We provide a brief overview of the capabilities here, but the **secr** help manual can be accessed with the command:

```
> RShowDoc("secr-manual", package = "secr")
```

We note that **secr** has many capabilities that we will not cover or do so only sparingly. We encourage you to read through the manual, the extensive documentation, and the vignettes, in order to get a better understanding of what the package is capable of. We also cover certain capabilities of **secr** in other chapters.

The main model-fitting function in **secr** is called **secr.fit**, which makes use of the standard **R** model specification framework with tildes. As an example, the equivalent of the basic model SCR0 is fitted as follows:

```
> secr.fit(capturedata, model = list(D ~ 1, g0 ~ 1, sigma ~ 1),
  buffer = 20000)
```

where **capturedata** is the object created by **secr** containing the encounter history data and the trap information, and the model expression **g0~1** indicates the intercept-only (i.e., constant) model. Note that we use p_0 for the baseline encounter probability parameter, which is g_0 in **secr** notation. A number of possible models for encounter probability can be fitted including both pre-defined variables (e.g., **t** and **b** corresponding to “time” and “behavior”), and user-defined covariates of several kinds. For example, to include a global behavioral response, this would be written as **g0~b**. The discussion of this (global versus local trap-specific behavioral response) and other covariates is developed more in Chapt. 7. We can also model covariates on density in **secr**, which we discuss in Chapt. 11. It is important to note that **secr** requires the buffer distance to be defined in meters and density will be returned as number of animals per hectare. Thus to make comparisons between **secr** and output from other programs, we will often have to convert the density to the same units.

Before we can fit the models, the data must first be packaged properly for **secr**. We require data files that contain two types of information: trap layout (location and identification information for each trap), which is equivalent to the trap deployment file (TDF) described in Sec. 5.9 and the capture data file containing sampling *session*, animal identification, trap occasion, and trap location, equivalent in information content to the encounter data file (EDF). Sample session can be thought of as primary period identifier

in a robust design like framework – it could represent a yearly sample or multiple sample periods within a year, each of them producing data on a closed population. We discuss “multi-session” models in more detail below, in Sec. 6.5.4 and Chapt. 14.

There are three important constructor functions that help package-up your data for use in `secr`: `read.traps`, `make.caphist` and `read.mask`. We provide a brief description of each here, but apply them to our wolverine camera trapping data in the next section:

(1) `read.traps`: This function points to an external file *or* **R** data object containing the trap coordinates, and other information, and also requires specification of the type of encounter devices (described in the next section). A typical application of this function looks like the following, invoking the `data=` option when there is an existing **R** object containing the trap information:

```
> trapfile <- read.traps(data=traps, detector="proximity")
```

(2) `make.caphist`: This function takes the EDF and combines it with trap information, and the number of sampling occasions. A typical application looks like this:

```
> capturedata <- make.caphist(enc.data, trapfile, fmt="trapID",
                             noccasions=165)
```

See `?make.caphist` for definition of distinct file formats. Specifying `fmt = trapID` is equivalent to our EDF format.

(3) `read.mask`: If there is a habitat mask available (as described in sec. 6.4.2), then this function will organize it so that `secr.fit` knows what to do with it. The function accepts either an external file name (see `?read.mask` for details of the structure) or a $nG \times 2$ **R** object, say `mask.coords`, containing the coordinates of the mask. A typical application looks like the following:

```
> grid <- read.mask(data=mask.coords)
```

These constructor functions produce output that can then be used in the fitting of models using `secr.fit`.

6.5.1 Encounter device types and detection models

The `secr` package requires that you specify the type of encounter device. Instead of describing models by their statistical distribution (Bernoulli, Poisson, etc.), `secr` uses certain operational classifications of detector types including ‘proximity’, ‘multi’, ‘single’, ‘polygon’ and ‘signal’. For camera trapping/hair snares we might consider ‘proximity’ detectors or ‘count’ detectors. The ‘proximity’ detector type allows, at most, one detection of each individual at a particular detector on any occasion (i.e., it is equivalent to what we call the Bernoulli or binomial encounter process model, or model SCR0). The ‘count’ detector designation allows repeat encounters of each individual at a particular detector on any occasion. There are other detector types that one can select such as: ‘polygon’ detector type which allows for a trap to be a sampled polygon (Royle and Young, 2008) which we discuss further in Chapt. 15, and ‘signal’ detector which allows for traps that have a strength indicator, e.g., acoustic arrays (Dawson and Efford, 2009). The detector types ‘single’ and ‘multi’ refer to traps that retain individuals, thus precluding the ability for animals to be captured in other traps during the sampling occasion. The ‘single’ type

6377 indicates trap that can only catch one animal at a time (single-catch traps), while 'multi'
 6378 indicates traps that may catch more than one animal at a time (multi-catch). These are
 6379 both variations of the multinomial encounter models described in Chapt. 9.

6380 As with all SCR models, **secr** fits an encounter probability model ("detection function"
 6381 in **secr** terminology relating the probability of encounter to the distance of a detector from
 6382 an individual activity center. **secr** allows the user to specify one of a variety of detection
 6383 functions including the commonly used half-normal ("Gaussian"), hazard rate ("Gaussian
 6384 hazard"), and (negative) exponential models. There are 12 different functions as of version
 6385 2.3.1 (see Table 7.1 in Chapt. 7), but some are only available for simulating data. The
 6386 different detection functions are defined in the **secr** manual and can be found by calling
 6387 the help function for the detection function:

```
6388 > ?detectfn
```

6389 Most of the detection functions available in **secr** contain some kind of a scale parameter
 6390 which is usually labeled σ . The units of this parameter default to meters in the **secr**
 6391 output. We caution that the meaning of this parameter depends on the specific detection
 6392 model being used, and it should not be directly compared as a measure of home-range size
 6393 across models. Instead, as we noted in Sec. 5.4 most encounter probability models imply
 6394 a model of space-usage and fitted encounter models should be converted to a common
 6395 currency such as "area used."

6396 6.5.2 Analysis using the **secr** package

6397 To demonstrate the use of the **secr** package, we will show how to do the same analysis on
 6398 the wolverine study as shown in Sec. 5.9. To use the **secr** package, the data need to be
 6399 formatted in a similar but slightly different manner than we use in **WinBUGS**.

6400 For example, in Sec. 5.9 we introduced a standard data format for the encounter data
 6401 file (EDF) and trap deployment file (TDF). The EDF shares the same format as that used
 6402 by the **secr** package with 1 row for every encounter observation and 4 columns representing
 6403 trap session ('Session'), individual identity ('ID'), sample occasion ('Occasion'), and trap
 6404 identity ('trapID'). For a standard closed population study that takes place during a single
 6405 season, the 'Session' column in our case is all 1's, to indicate a single primary sampling
 6406 occasion. In addition to providing the encounter data file (EDF), we must tell **secr** infor-
 6407 mation about the traps, which is formatted as a matrix with column labels 'trapID', 'x' and
 6408 'y', the last two being the coordinates of each trap, with additional columns representing
 6409 the operational state of each trap during each occasion (1=operational, 0=not).

6410 We demonstrate these differences now by walking through an analysis of the wolverine
 6411 camera trapping data using **secr**. To read in the trap locations and other related infor-
 6412 mation, we make use of the constructor function **read.traps** which also requires that we
 6413 specify the detector type. The detector type is important because it will determine the
 6414 likelihood that **secr** will use to fit the model. Here, we have selected "proximity" which
 6415 corresponds to the Bernoulli encounter model in which individuals are captured at most
 6416 once in each trap during each sampling occasion:

```
6417 > library(secr)
6418 > library(scrbook)
```

```

6419 > data(wolverine)
6420
6421 > traps <- as.matrix(wolverine$wtraps)
6422 > dimnames(traps) <- list(NULL,c("trapID","x","y",paste("day",1:165,sep="")))
6423 > traps1 <- as.data.frame(traps[,1:3])
6424 > trapfile1 <- read.traps(data=traps1,detector="proximity")

```

Here we note that trap coordinates are extracted from the wolverine data but we do *not* scale them. This is because **secr** defaults to coordinate scaling of meters which is the extant scaling of the wolverine trap coordinates. Note that we add a 'trapID' column to the trap coordinates and provide appropriate column labels to the 'traps' matrix. An important aspect of the wolverine study is that while the camera traps were operated over a 165 day period, each trap was operational during only a portion of that period. We need to provide the trap operation information which is contained in the columns to the right of the trap coordinates in our standard trap deployment file (TDF). Unfortunately, this is less easy to do in **secr**⁴, which requires an external file with a single long string of 1's and 0's indicating the days in which each trap was operational (1) or not (0). The **read.traps** function will not allow for this information on trap operation if the data exists as an **R** object – instead, we can create this external file and then read it back in with **read.traps** using these commands:

```

6438 > hold <- rep(NA,nrow(traps))
6439 > for(i in 1:nrow(traps)){
6440 >   hold[i] <- paste(traps[i,4:ncol(traps)],collapse="")
6441 > }
6442 > traps1 <- cbind(traps[,1:3], "usage"=hold)
6443
6444 > write.table(traps1, "traps.txt", row.names=FALSE, col.names=FALSE)
6445 > trapfile2 <- read.traps("traps.txt",detector="proximity")

```

These operations can be accomplished using the function **scr2secr** which is provided in the **R** package **scrbook**.

After reading in the trap data, we now need to create the encounter matrix or array using the **make.caphist** command, where we provide the capture histories in EDF format, which is the existing format of the data input file **wcaps**. In creating the capture history, we provide also the trapfile created previously, the format (e.g., here EDF format is **fmt= 'trapID'**), and finally, we provide the number of occasions.

```

6453 #
6454 # Grab the encounter data file and format it:
6455 #
6456 wolv.dat <- wolverine$wcaps
6457 dimnames(wolv.dat) <- list(NULL,c("Session","ID","Occasion","trapID"))
6458 wolv.dat <- as.data.frame(wolv.dat)
6459 wolvcapt2 <- make.caphist(wolv.dat,trapfile2,fmt="trapID",noccasions=165)

```

⁴as of v. 2.3.1

We also set up a habitat mask using the 2×2 km grid which we used previously in the analysis of the wolverine data and then pass the relevant objects to `secr.fit` as follows:

```
#
# Grab the habitat mask (2 x 2 km) and format it:
#
gr2 <- (as.matrix(wolverine$grid2))
dimnames(gr2) <- list(NULL,c("x","y"))
gr2 <- read.mask(data=gr2)
#
# To fit the model we use secr.fit:
#
wolv.secr2 <- secr.fit(wolvcapt2,model=list(D ~ 1, g0 ~ 1, sigma ~ 1),
                      buffer=20000,mask=gr2)
```

We are using the “proximity detector” model (SCR0), so we do not need to make any specifications in the command line because we have specified the detector type using the constructor function `read.traps`, except to provide the buffer size (in meters). To specify different models, you can change the default model `D~1`, `g0~1`, `sigma~1`. We provide all of these commands and additional analyses in the `scrbook` package with the function called `secr_wolverine`. Printing the output object produces the following (slightly edited):

```
> wolv.secr2
secr 2.3.1, 15:52:45 29 Aug 2012

Detector type      proximity
Detector number    37
Average spacing    4415.693 m
x-range            593498 652294 m
y-range            6296796 6361803 m
N animals          : 21
N detections       : 115
N occasions        : 165
Mask area          : 987828.1 ha

Model              : D ~ 1 g0 ~ 1 sigma ~ 1
Fixed (real)       : none
Detection fn       : halfnormal
Distribution        : poisson
N parameters       : 3
Log likelihood     : -602.9207
AIC                : 1211.841
AICc               : 1213.253

Beta parameters (coefficients)
      beta      SE.beta      lcl      ucl
```

```

6504 D      -9.390124 0.22636698 -9.833795 -8.946452
6505 g0     -2.995611 0.16891982 -3.326688 -2.664535
6506 sigma  8.745547 0.07664648  8.595323  8.895772
6507
6508 Variance-covariance matrix of beta parameters
6509                D                g0                sigma
6510 D      0.0512420110 -0.0004113326 -0.003945371
6511 g0    -0.0004113326  0.0285339045 -0.006269477
6512 sigma -0.0039453711 -0.0062694767  0.005874683
6513
6514 Fitted (real) parameters evaluated at base levels of covariates
6515      link      estimate SE.estimate      lcl      ucl
6516 D      log 8.354513e-05 1.915674e-05 5.360894e-05 1.301982e-04
6517 g0     logit 4.762453e-02 7.661601e-03 3.466689e-02 6.509881e-02
6518 sigma  log 6.282651e+03 4.822512e+02 5.406315e+03 7.301037e+03

```

The object returned by `secr.fit` provides extensive default output when printed. Much of this is basic descriptive information about the model, the traps, or the encounter data. We focus here on the parameter estimates. Under the fitted (real) parameters, we find D , the density, given in units of individuals/hectare (1 hectare = 10000 m^2). To convert this into individuals/1000 km^2 , we multiply by 100000, thus our density estimate is 8.35 individuals/1000 km^2 . The parameter σ is given in units of meters, and so this corresponds to 6.283 km. Both of these estimates are very similar to those obtained in our likelihood analysis summarized in Table 6.3 which, for the 2×2 km grid, we obtained $\hat{D} = 8.31$ with a SE of $100000 \times 1.915674e - 05 = 1.9156$ and, accounting for the scale difference (1 unit = 10000 m in the previous analysis), $\hat{\sigma} = \sqrt{1/(2\hat{\alpha}_1)} * 10000 = 6.289$ km. The difference in the MLE between Table 6.3 and those produced by `secr` could be due to subtle differences in internal tuning of optimization algorithms, starting values or other numerical settings. In addition, the likelihood is based on a Poisson prior for N (see the next section). On the other hand, the SE is slightly larger based on `secr` which is due to a subtle difference in the interpretation of D under the `secr` model (See below).

6.5.3 Likelihood analysis in the `secr` package

The `secr` package does likelihood analysis of SCR models for most classes of models as developed by Borchers and Efford (2008). Their formulation deviates slightly from the binomial form we presented in Sec. 6.2 above (though Borchers and Efford (2008) also mention the binomial form). Specifically, the likelihood that `secr` implements is that based on removing N from the likelihood by integrating the binomial likelihood (Eq. 6.2.1 above) over a Poisson prior for N – what we will call the *Poisson-integrated likelihood* as opposed to the conditional-on- N (*binomial-form*) considered previously.

To develop the Poisson-integrated likelihood we compute the marginal probability of each \mathbf{y}_i and the probability of an all-0 encounter history, π_0 , as before, to arrive at the marginal likelihood in the binomial-form:

$$\mathcal{L}(\boldsymbol{\alpha}, n_0 | \mathbf{y}) = \frac{N!}{n!n_0!} \left\{ \prod_i [y_i | \boldsymbol{\alpha}] \right\} \pi_0^{n_0}$$

Now, what Borchers and Efford (2008) do is assume that $N \sim \text{Poisson}(\Lambda)$ and they do a further level of marginalization over this prior distribution:

$$\sum_{n_0=0}^{\infty} \frac{N!}{n!n_0!} \left\{ \prod_i [y_i | \alpha] \right\} \pi_0^{n_0} \frac{\exp(-\Lambda) \Lambda^N}{N!}$$

In Chapt. 11 we write $\Lambda = \mu ||\mathcal{S}||$ where $||\mathcal{S}||$ is the area of the state-space, and μ is the density (“intensity”) of the point process. Carrying out the summation above produces exactly this marginal likelihood:

$$\mathcal{L}_2(\alpha, \Lambda | \mathbf{y}) = \left\{ \prod_i [y_i | \alpha] \right\} \Lambda^n \exp(-\Lambda(1 - \pi_0))$$

which is Eq. 2 of Borchers and Efford (2008) except for notational differences. It also resembles the binomial-form of the likelihood in Eq. 6.2.1 with $\Lambda^n \exp(-\Lambda\pi_0)$ replacing the combinatorial term and the $\pi_0^{n_0}$ term. We emphasize there are two marginalizations going on here: (1) the integration to remove the latent variables \mathbf{s} ; and, (2) summation to remove the parameter N . We provide a function for computing this in the `scrbook` package called `intlik3Poisson`. The help file for that function shows how to conduct a small simulation study to compare the MLE under the Poisson-integrated likelihood with that from the binomial form.

The essential distinction between our MLE and Borchers and Efford as implemented in `seccr` is whether you keep N in the model or remove it by integration over a Poisson prior. If you have prescribed a state-space explicitly with a sufficiently large buffer, then we imagine there should be hardly any difference at all between the MLEs obtained by either the Poisson-integrated likelihood or the binomial-form of the likelihood which retains N . There is a subtle distinction in the sense that under the binomial form, we estimate the realized population size N for the state-space whereas, for the Poisson-integrated form we estimate the *prior* expected value which would apply to a hypothetical new study of a similar population (see Sec. 5.7.3).

Both models (likelihoods) assume \mathbf{s} is uniformly distributed over space, but for the binomial model we make no additional assumption about N whereas we assume N is Poisson using the formulation in `seccr` from (Borchers and Efford, 2008). Using data augmentation we could do a similar kind of integration but integrate N over a binomial (M, ψ) prior – which we referred to as the binomial-integrated likelihood in Sec. 4.2.4. So obviously the two approaches (data augmentation and Poisson-integrated likelihood) are approximately the same as M gets large. However, doing a Bayesian analysis by MCMC, we obtain an estimate of both N , the *realized population size*, and the parameter controlling its expected value ψ which are, in fact, both identifiable from the data even using likelihood analysis (Royle et al., 2007). That said we can integrate N out completely and just estimate ψ as we noted in Sec. 6.2.1 above.

6.5.4 Multi-session models in `seccr`

In practice we will often deal with SCR data that have some meaningful stratification or group structure. For example, we might conduct mist-netting of birds on K consecutive days, repeated, say, T times during a year, or perhaps over T years. Or we might collect

data from R distinct trapping grids. In these cases, we have T or R groups which we might reasonably regard as being samples of independent populations. While the groups might be distinct sites, year, or periods within years, they could also be other biological groups such as sex or age. Conveniently, **secr** fits a specific model for stratified populations – referred to as *multi-session* models. These models build on the Poisson assumption which underlies the integrated likelihood used in **secr** (as described in the previous section). To understand the technical framework, let N_g be the population size of group g and *assume*

$$N_g \sim \text{Poisson}(\Lambda_g).$$

Naturally, we model group-specific covariates on Λ_g :

$$\log(\Lambda_g) = \beta_0 + \beta_1 z_g$$

where z_g is some group-specific covariate such as a categorical index to the group, or a trend variable, or a spatial covariate, such as treatment effect or habitat structure, if the groups represent spatial units. Under this model, we can marginalize *all* N_g parameters out of the likelihood to concentrate the likelihood on the parameters β_0 and β_1 precisely as discussed in the previous section. This Poisson hierarchical model is the basis of the multi-session models in **secr**.

To implement a multi-session model (or stratified population model) in **secr**, we provide the relevant stratification information in the ‘Session’ variable of the input encounter data file (EDF). If ‘Session’ has multiple values then a “multi-session” object is created by default and session-specific variables can be described in the model. For example, if the session has 2 values for males and females then we have sex-specific densities, and baseline encounter probability p_0 (g_0 in **secr**) by just doing this (see Chapt. 8 for the **R** code to set this up):

```
> out <- secr.fit(capdata, model=list(D ~ session, g0 ~ session, sigma~ 1),
buffer=20000)
```

More detailed analysis is given in Sec. 8.1 where we fit a number of different models and apply methods of model selection to obtain model-averaged estimates of density.

We can also easily implement stratified population models in the various **BUGS** engines using data augmentation (Converse and Royle, 2012; Royle and Converse, in review) which we discuss, with examples, in Chapt. 14.

6.5.5 Some additional capabilities of **secr**

The **secr** package has capabilities to do a complete analysis of SCR data sets, including model fitting, selection, and many summary analyses. In the previous sections, we’ve given a basic overview, and we do more in later chapters of this book. Here we mention a few of these other capabilities that you should know about as you use **secr**. Of course, you should skim through the associated documentation (`?secr`) to see more of what’s available.

Alternative observation models

secr fits a wide range of alternative observation models besides the Bernoulli encounter model, including multinomial encounter models for “multi-catch” and “single catch” traps, models for sound attenuation from acoustic detection devices, and many others. We discuss many of these other methods in Chapt. 9 and elsewhere in the book.

Summary statistics

secr provides a useful default summary of the data, but it also has summary statistics about animal movement including mean-maximum distance moved (the function **MMDM**). For example, see the help page `?MMDM` which lists a number of other summary functions which take a **capthist** object:

```
> moves(capthist)
> dbar(capthist)
> RPSV(capthist)
> MMDM(capthist, min.recapt = 1, full = FALSE)
> ARL(capthist, min.recapt = 1, plt = FALSE, full = FALSE)
```

The function **moves** returns the observed distances moved, **dbar** returns the average distance moved, **RPSV** produces a measure of dispersion about the home-range center, and **ARL** gives the *Asymptotic Range Length* which is the asymptote of an exponential model fit to the observed range length vs. the number of detections of each individual (Jett and Nichols, 1987).

State-space buffer

secr will produce a warning if the state-space buffer is chosen too small. For example, in fitting the wolverine data as in Sec. 6.5.2 but with a 1000 m buffer, and we see the following warning message:

```
Warning message:
In secr.fit(wolvcapt2, model=list(D ~ 1, g0 ~ 1, sigma ~ 1), buffer=1000):
  predicted relative bias exceeds 0.01 with buffer = 1000
```

This should cause you to contemplate modifying the state-space buffer if that is a reasonable thing to do in the specific application.

Model selection and averaging

secr does likelihood ratio tests to compare nested models using the function **LR.test**. You can create model selection tables based on AIC or AICc, using the function **AIC**, and obtain model-averaged parameter estimates using the function **model.average** (See Chapt. 8 for examples).

Population closure test

secr has a population closure test with the function **closure.test** which implements the tests of Stanley and Burnham (1999) or Otis et al. (1978). The function is used like this: **closure.test(object, SB = FALSE)**. Here **object** is a **capthist** object and **SB** is a logical variable that, if **TRUE**, produces the Stanley and Burnham (1999) test.

Density mapping and effective sample area

secr produces likelihood versions of the various summaries of posterior density and effective sample area that we discussed in Chapt. 5. For example, while **secr** reports estimates of the expected value of N or density directly in the summary output from fitting a model, you can use the function **region.N** to produce estimates of N for any given region. In addition, **secr** has functions for creating maps of detection contours for individuals traps, or for the entire trap array. See the function **pdot.contour**, and also **fxi.contour** for

computing the 2-dimensional pdf of the locations of one or more individual activity centers (as in Sec. 5.11.3). In the context of likelihood analysis, estimation of a random effect \mathbf{s} is based on a plug-in application of Bayes' Rule. When \mathbf{s} has a uniform distribution, and we use a discrete evaluation of the integral, it can be computed simply by renormalizing the likelihood:

$$[\mathbf{s}|\mathbf{y}, \theta] = \frac{[\mathbf{y}|\mathbf{s}, \theta]}{\sum_{\mathbf{s}} [\mathbf{y}|\mathbf{s}, \theta]}.$$

Any of the `intlik` functions given previously in this chapter can be easily modified to return the posterior distribution of \mathbf{s} for any, or all, individuals, or an individual that is not encountered.

Effective sample area (see Sec. 5.12) can be calculated in `secr` using the functions `esa` and `esa.plot`.

Covariate models

`secr` has many capabilities for modeling covariates. It has a number of built-in models that allow certain covariates on encounter probability, which we cover to a large extent in Chapt. 7, and also see Chapt. 8 for more examples. `secr` also allows covariates to be built into the density model (see Chapt. 11). It has some built in response surface models, allowing for the fitting of linear or quadratic response surfaces. This is done by modifying the density model in `secr.fit`. For example, $D \sim 1$ is a constant density surface, and $D \sim \mathbf{x} + \mathbf{y}$ fits a linear response surface, etc.. See the manual `secr-densitysurfaces.pdf` for the details.

There are a number of ways to model your own "custom" covariates (as opposed to pre-specified models). One way is to use the `addCovariates` function and supply it a `mask` or `traps` object along with some "spatialdata." Or, if you have covariates at each trap location then it will extrapolate to all points on the habitat mask. There's also a method by which the user can create a function of geographic coordinates, `userDfn`, which seems to provide additional flexibility, although we haven't used this method. There is a handy function `predictDsurface` for producing density maps under the specified model for density.

6.6 SUMMARY AND OUTLOOK

In this chapter, we discussed basic concepts related to classical analysis of SCR models based on likelihood methods. Analysis is based on the so-called integrated or marginal likelihood in which the individual activity centers (random effects) are removed from the conditional-on- \mathbf{s} likelihood by integration. We showed how to construct the integrated likelihood and fit some simple models in the **R** programming language. In addition, likelihood analysis for some broad classes of SCR models can be accomplished using the **R** library `secr` (Efford, 2011) which we provided a brief introduction to. In later chapters we provide more detailed analyses of SCR data using likelihood methods and the `secr` package.

Why or why not use likelihood inference exclusively? For certain specific models, it is may be more computationally efficient to produce MLEs (for an example see Chapt. 12). And, likelihood analysis makes it easy to do model-selection by AIC and compute standard errors or confidence intervals. However, **BUGS** is extremely flexible in terms of describing models and we can devise models in the **BUGS** language easily that we

cannot fit in **secr**. For example, in Chapt 16 we consider open population models which are straightforward to develop in **BUGS** but, so far, there is no available platform for doing MLE of such models. We can also fit models in **BUGS** that accommodate missing covariates in complete generality (e.g., unobserved sex of individuals), and we can adopt SCR models to include auxiliary data types. For example, we might have camera trapping and genetic data and we can describe the models directly in **BUGS** and fit a joint model (Gopalaswamy et al., 2012b). To do maximum likelihood estimation, we have to write a custom new piece of code for each model⁵ or hope someone has done it for us. You should have some capability to develop your own MLE routines with the tools we provided in this chapter.

⁵Although we may be able to handle multiple survey methods together in **secr** using the multi-session models.