

¹ Chapter 1

² Introduction

³ Chapter 2

⁴ GLMS and WinBUGS

⁵ Chapter 3

⁶ Closed population models

⁷ Chapter 4

⁸ Fully Spatial
⁹ capture-recapture models

¹⁰ **Chapter 5**

¹¹ **Other observation models**

¹² **Chapter 6**

¹³ **MCMC details**

14 Chapter 7

15 MCMC Details

16 7.1 Introduction

17 In this chapter we will dive a little deeper into Markov chain Monte Carlo
18 (MCMC) sampling. We will construct custom MCMC samplers in R, starting
19 with easy-to-code GLMs and GLMMs and moving on to simple SCR mod-
20 els. We will also demonstrate some tricks and simple extensions to the 'spatial
21 null model'. Finally, we will illustrate some alternative ready-to-use software
22 packages for MCMC sampling. We will NOT provide exhaustive background
23 information on the theory and justification of MCMC sampling there are entire
24 books dedicated to that subject and we refer you to ? and ?. Rather we aim to
25 provide you with enough background and technical know-how to start building
26 your own MCMC samplers for SCR models in R.

27 7.1.1 Why build your own MCMC algorithm?

28 The standard program we have used so far to run MCMC analyses is WinBUGS
29 (?). The wonderful thing about WinBUGS is that it will automatically use the
30 most appropriate and efficient form of MCMC sampling for the model specified
31 by the user.

32 The fact that we have such a Swiss Army knife type of MCMC machine begs
33 the question: Why would anyone want to build their own MCMC algorithm?
34 For one, there are a limited number of distributions and functions implemented
35 in WinBUGS. While OpenBUGS provides more options, some more complex
36 models may be impossible to build within these programs. A very simple exam-
37 ple from spatial capture-recapture that can give you a headache in WinBUGS
38 is when your state-space is an irregular-shaped polygon, rather than an ideal
39 rectangle that can be characterized by four pairs of coordinates. It is easy to
40 restrict activity centers to any arbitrary polygon in R using an ESRI shapefile
41 (and we will show you an example in a little bit), but you cannot use a shape

file in a BUGS model.

Sometimes implementing an MCMC algorithm in R may be faster than in WinBUGS - especially if you want to run simulation studies where you have hundreds or more simulated data sets, several years' worth of data or other large models, this can be a big advantage.

Finally, building your own MCMC algorithm is a great exercise to understand how MCMC sampling works. So while using the BUGS language requires you to understand the structure of your model, building an MCMC algorithm requires you to think about the relationship between your data, priors and posteriors, and how these can be efficiently analyzed and characterized. Not to mention that, if you are an R junkie, it can actually be fun. However, if you don't think you will ever sit down and write your own MCMC sampler, consider skipping this chapter - apart from coding it will not cover anything SCR-related that is not covered by other, more model-oriented chapters as well.

7.2 MCMC and posterior distributions

As mentioned in Chapter 2, MCMC is a class of simulation methods for drawing (correlated) random numbers from a target distribution, which in Bayesian inference is the posterior distribution. As a reminder, the posterior distribution is a probability distribution for an unknown parameter, say θ , given a set of observed data and its prior probability distribution (the probability distribution we assign to a parameter before we observe data). The great benefit of computing the posterior distribution of θ is that it can be used to make probability statements about θ , such as the probability that θ is equal to some value, or the probability that θ falls within some range of values. As an example, suppose we conducted a Bayesian analysis to estimate detection probability of some species at a study site (p), and we obtained a posterior distribution of $\text{beta}(20,10)$ for the parameter p . The following R commands demonstrate how we make inferences based upon summaries of the posterior distribution. Fig 1 shows the posterior along with the summary statistics.

```
> (post.median <- qbeta(0.5, 20, 10))
[1] 0.6704151
> (post.95ci <- qbeta(c(0.025, 0.975), 20, 10))
[1] 0.4916766 0.8206164
```

Thus, we can state that there is a 95% probability that θ lies between 0.49 and 0.82.

The posterior distribution summarizes all we know about a parameter and thus, is the central object of interest in Bayesian analysis. Unfortunately, in many if not most practical applications, it is nearly impossible to directly compute the posterior. Recall Bayes theorem:

$$p(\theta|y) = p(y|\theta) * p(\theta)/p(y), \quad (7.1)$$




Figure 7.1: Probability density plot of a hypothetical posterior distribution of $\text{beta}(20,10)$; dashed lines indicate mean and upper and lower 95% interval

81 where θ is the parameter of interest, y is the observed data, $p(\theta|y)$ is the poste-
 82 rior, $p(y|\theta)$ the likelihood of the data conditional on θ , $p(\theta)$ the prior probability
 83 of θ , and, finally, $p(y)$ is the marginal probability of the data, which can also be
 84 written as

$$p(y) = \int p(y|\theta) * p(\theta) d\theta$$

85 This marginal probability is a normalizing constant that ensures that the
 86 posterior integrates to 1. You read in Chapter 2 that this integral is often hard
 87 or impossible to evaluate, unless you are dealing with a really simple model. For
 88 example, consider that you have a Normal model, with a set of n observations,
 89 y that come from a Normal distribution:

$$y \sim \text{Normal}(\mu, \sigma),$$

90 where σ is known and our objective is to obtain an estimate of μ using Bayesian
 91 statistics. To fully specify the model in a Bayesian framework, we first have
 92 to define a prior distribution for μ . Recall from Chapter 2 that for certain
 93 data models, certain priors lead to conjugacy i.e. if you choose the right prior
 94 for your parameter, your posterior distribution will be of a known parametric
 95 form. The conjugate prior for the mean of a normal model is also a Normal
 96 distribution:

$$\mu \sim \text{Normal}(\mu_0, \sigma_0^2)$$

97 If μ_0 and σ_0^2 are fixed, the posterior for μ has the following form (for the algebraic
 98 proof, see XXX):

$$\mu|y \sim \text{Normal}(\mu_n, \sigma_n^2) \quad (7.2)$$

99 where

$$\mu_n = (sig^2/sig^2 + n * sig0^2) * mu0 + (n * sig0^2/sig^2 + n * sig0^2) * y - bar$$

100 And

$$sign^2 = sig^2 * sig0^2 / (sig^2 + n * sig0^2)$$

101 We can directly obtain estimates of interest from this Normal posterior distri-
 102 bution, such as the mean μ -hat and its variance; we do not need to apply
 103 MCMC, since we can recognize the posterior as a parametric distribution, in-
 104 cluding the normalizing constant $p(y)$. But generally we will be interested in
 105 more complex models with several, say n , parameters. In this case, computing
 106 $p(y)$ from Eq. 7.1 requires n -dimensional integration, which is can be difficult
 107 or impossible. Thus, the posterior distribution in generally only known up to a
 108 constant of proportionality:

$$p(\theta|y)proptop(y|\theta) * p(\theta)$$

109 The power of MCMC is that it allows us to approximate the posterior using
 110 simulation without evaluating the high dimensional integrals and to directly
 111 sample from the posterior, even when the posterior distribution is unknown!
 112 The price is that MCMC is computationally expensive. Although MCMC first
 113 appeared in the scientific literature in 1949 (?), widespread use did not occur
 114 until the 1980s when computational power and speed increased (?). It is safe
 115 to say that the advent of practical MCMC methods is the primary reason why
 116 Bayesian inference has become so popular during the past three decades. In
 117 a nutshell, MCMC lets us generate sequential draws of θ (the parameter(s)
 118 of interest) from distributions approximating the unknown posterior over T
 119 iterations. The distribution of the draw at t depends on the value drawn at $t-1$;
 120 hence, the draws from a Markov chain.¹ As T goes to infinity, the Markov
 121 chain converges to the desired distribution in our case the posterior distribution
 122 for $\theta|y$. Thus, once the Markov chain has reached its stationary distribution,
 123 the generated samples can be used to characterize the posterior distribution,
 124 $p(\theta|y)$, and point estimates of θ , its standard error and confidence bounds,
 125 can be obtained directly from this approximation of the posterior. In practice,
 126 although we know that a Markov chain will eventually converge, we can only
 127 generate a limited number of samples a process that depending on the model
 128 can be quite time consuming. Assessing whether our Markov chain has indeed
 129 converged is an important part of MCMC sampling and we will speak about
 130 some common diagnostics in Section XX.

131 7.3 Types of MCMC sampling

132 There are several MCMC algorithms, the most popular being Gibbs sampling
 133 and Metropolis-Hastings sampling. We will be dealing with these two classes in

¹In case you are not familiar with Markov chains, for t random samples $\theta(1), \dots, \theta(t)$ from a Markov chain the distribution of $\theta(t)$ depends only on the most recent value, $\theta(t-1)$.

more detail and use them to construct the MCMC algorithms for SCR models. Also, we will briefly review alternative techniques that are applicable in some situations.

7.3.1 Gibbs sampling

Gibbs sampling was named after the physicist J.W. Gibbs by ², who applied the algorithm to a Gibbs distribution ². The roots of Gibbs sampling can be traced back to work of ², and it is actually closely related to Metropolis sampling (see Chapter 11.5 in ², for the link between the two samplers). We will focus on the technical aspects of this algorithm, but if you find yourself hungry for more background, ² provide a more in-depth introduction to the Gibbs sampler.

In Chapter 2 you already heard about the basic principles of Gibbs sampling³. But as a refresher, let's go back to our simple example from above to understand the motivation and functioning of Gibbs sampling. Recall that for a Normal model with known variance and a Normal prior for μ , the posterior distribution of $\mu|y$ is also Normal. Conversely, with a fixed (known) μ , but unknown variance, the conjugate prior for σ^2 is an Inverse-Gamma distribution with shape and scale parameters a and b :

$$\sigma^2 \sim \text{Inv-Gamma}(a, b),$$

With fixed a and b , the posterior $p(\text{sig}|\mu, y)$ is also an Inverse Gamma distribution, namely:

$$\text{sig}|\mu, y \sim \text{InvGamma}(an, bn), \quad (7.3)$$

where $an = n/2 + a$ and $bn = 1/2\sigma(y_i - \mu)^2 + b$ However, what if we know neither μ nor sig , which is probably the more common case? The joint posterior distribution of μ and sig now has the general structure

$$p(\mu, \text{sig}|y) = \frac{p(y|\mu) * p(\mu) * p(\text{sig})}{\int p(y|\mu) * p(\mu) * p(\text{sig}) d\mu d\text{sig}}$$

Or

$$p(\mu, \text{sig}|y) \propto p(y|\mu) * p(\mu) * p(\text{sig})$$

This cannot easily be reduced to a distribution we recognize. However, we can condition μ on sig (i.e., we treat sig as fixed) and remove all terms from the joint posterior distribution that do not involve μ to construct the full conditional distribution,

$$p(\mu|\text{sig}, y) \propto p(y|\mu) * p(\mu)$$

The full conditional of μ again takes the form of the Normal distribution shown in Eq. ??; similarly, $p(\text{sig}|\mu, y)$ takes the form of the Inverse Gamma

²a distribution from physics we are not going to worry about, since it has no immediate connection with Gibbs sampling other than giving its name

³maybe we should think out chapter 2 and concentrate that material here?

163 distribution shown in Eq. 7.3 both distribution we can easily sample from.
 164 And this is precisely what we do when using Gibbs sampling we break down
 165 high-dimensional problems into convenient one-dimensional problems by con-
 166 structing the full conditional distributions for each model parameter separately;
 167 and we sample from these full conditionals, which, if we choose conjugate priors,
 168 are known parametric distributions. Let's put the concept of Gibbs sampling
 169 into the MCMC framework of generating successive samples, using our simple
 170 Normal model with unknown μ and σ and conjugate priors as an example.
 171 These are the steps you need to build a Gibbs sampler:

172 **Step 0:** Begin with some initial values for θ , $\theta(0)$. In our example, we have to
 173 specify initial values for μ and σ , for example by drawing a random number
 174 from some uniform distribution, or by setting them close to what we think they
 175 might be. (Note: This step is required in any MCMC sampling chains have to
 176 start from somewhere. We will get back to these technical details a little later.)

177 **Step 1:** Draw $\theta_1(1)$ from the conditional distribution $p(\theta_1(1) | \theta_2(0), \dots, \theta_d(0))$
 178 Here, θ_1 is μ , which we draw from the Normal distribution in Eq. 7.3 using
 179 $\sigma(0)$ as value for σ .

180 **Step 2:** Draw $\theta_2(1)$ from the conditional distribution $p(\theta_2(1) | \theta_1(1), \theta_3(0), \dots, \theta_d(0))$
 181 Here, θ_2 is σ , which we draw from the Inverse Gamma distribution of
 182 Eq. 7.3, using $\mu(1)$ as value for μ ...

183 **Step d:** Draw $\theta_d(1)$ from the conditional distribution $p(\theta_d(1) | \theta_1(1), \dots, \theta_{d-1}(1))$
 184

185 In our example we have no additional parameters, so we only need step 0
 186 through to 2. Repeat Steps 1 to d for K = a large number of samples. In terms
 187 of R coding, this means we have to write Gibbs updaters for μ and σ and
 188 embed them into a loop over K iterations. The final code in the form of an R
 189 function is shown in Panel 1.

190 Andy will build the panel environment here soon.

191
 192 Panel 1: R-code for a Gibbs sampler for a Normal model with unknown μ
 193 and σ and conjugate (Normal and Inverse Gamma, respectively) priors
 194 for both parameters.

```
195
196 Normal.Gibbs<-function(y=y,mu0=mu0, sig0=sig0, a=a,b=b,niter=niter) {
197
198   ybar<-mean(y)
199   n<-length(y)
200   mu<-runif(1) #mean initial value
201   sig<-runif(1) #sd initial value
202   an<-n/2 + a
203
204   out<-matrix(nrow=niter, ncol=2)
```

```

205 colnames(out)<-c('mu', 'sig')
206
207 for (i in 1:niter) {
208
209   #update mu
210   mun<- (sig/(sig+n*sig0))*mu0 + (n*sig0/(sig+n* sig0))*ybar
211   sign <- (sig*sig0)/ (sig+n*sig0)
212   mu<-rnorm(1,mun, sqrt(sign))
213
214   #update sig
215   bn<- 0.5 * (sum((y-mu)^2)) +b
216   sig<-1/rgamma(1,shape=an, rate=bn)
217   out[i,]<-c(mu,sqrt(sig))
218
219 }
220 return(out)
221 }

```

222 This is it! You can use the code `NormalGibbs.R` in the **R** package `scrbook`
 223 to simulate some data, $y \sim \text{Normal}(5, 0.5)$ and run your first Gibbs sampler.
 224 Your output will be a table with two columns, one per parameter, and K rows,
 225 one per iteration. For this 2-parameter example you can visualize the joint
 226 posterior by plotting samples of μ against samples of σ (Fig. 2 XXX):

```

227 plot(out[,1], out[,2])

```

228 The marginal distribution of each parameter is approximated by just examining
 229 the samples of this particular parameter you can visualize it by plotting a
 230 histogram of the samples (Fig. 3 a, b XXX):

```

231 par(mfrow=c(1,2))
232 hist(out[,1]); hist (out[,2])

```

233 Finally, recall an important characteristic of Markov chains, namely, that the
 234 chain has to have converged (reached its stationary distribution) for samples to
 235 come from the posterior distribution. In practice, that means you have to throw
 236 out some of the initial samples called the burn-in. We will talk about this in
 237 more when we talk about convergence diagnostics. For now, you can use the
 238 `plot(out[,1])` or `plot(out[,2])` command to make a time series plot of the
 239 samples of each parameter and visually assess how many of the initial samples
 240 you should discard. Figure 3 c and d shows plots for the estimates of μ and
 241 sigma from our simulated data set; you see that in this simple example the
 242 Markov chain apparently reaches its stationary distribution very quickly the
 243 chains look 'grassy' seemingly from the start. It is hard to discern a burn-in
 244 phase visually (but we will see examples further on where the burn-in is clearer)
 245 and you may just discard the first 500 draws to be sure you only use samples
 246 from the posterior distribution. The mean of the remaining samples are your
 247 estimates of μ and σ :




Figure 7.2: Joint posterior distribution of mu and sig from a Normal Model

```

248 > summary(mod[501:10000,])
249      mu              sig
250 Min.   : 4.936      Min.   : 0.4569
251 1st Qu.: 4.984      1st Qu.: 0.4889
252 Median : 4.994      Median : 0.4961
253 Mean   : 4.994      Mean    : 0.4964
254 3rd Qu.: 5.005      3rd Qu.: 0.5037
255 Max.   : 5.062      Max.    : 0.5356

```

256 7.3.2 Metropolis-Hastings sampling

257 Although it is applicable to a wide range of problems, the limitations of Gibbs
 258 sampling are immediately obvious what if we do not want to use conjugate
 259 priors (or what if we cannot recognize the full conditional distribution as a
 260 parametric distribution, or simply do not want to worry about these issues)?
 261 The most general solution is to use the Metropolis-Hastings (MH) algorithm,
 262 which also goes back to the work by ?. You saw the basics of this algorithm in
 263 Chapter 2. In a nutshell, because we do not recognize the posterior $p(\theta|y)$ as
 264 a parametric distribution, the MH algorithm generates samples from a known
 265 proposal distribution, say $h(\theta)$, that depends on θ at t-1. The t^{th} sample is
 266 accepted or rejected based on its joint posterior probability density compared to
 267 the density of the sample at t-1. The original Metropolis algorithm requires $h(\theta)$
 268 to be symmetric so that $h(\theta^t|\theta^{t-1}) = h(\theta^{t-1}|\theta^t)$; but a later development of the

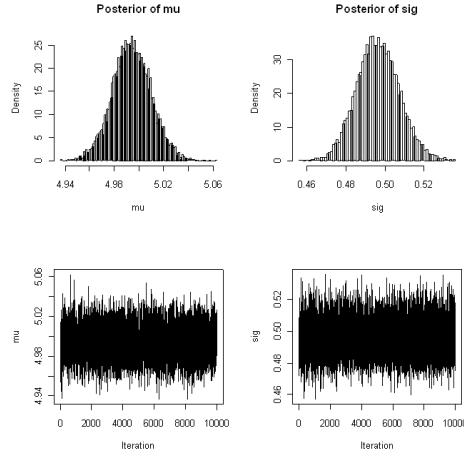


Figure 7.3: Plots of the posterior distributions of μ (a) and σ (b) from a Normal model and time series plots of μ (c) and σ (d).

algorithm by ? lifted this condition. Using a symmetric proposal distribution makes life a little easier and we are going to limit our coverage of the Metropolis-Hastings sampler to this specific case. Specifically, we are going to use a Normal proposal distribution, which is also referred to as 'random walk Metropolis-Hastings sampling'. It is worth knowing that there are alternative formulations of the algorithm. For example, in the independent M-H, θ^t does not depend on θ^{t-1} , while the Langevin algorithm (?) aims at avoiding the random walk by favoring moves towards regions of higher posterior probability density. The interested reader should look up these algorithms in ? or ?.

Building a MH sampler can be broken down into several steps. We are going to demonstrate these steps using a different but still simple and common model the logit-normal or logistic regression model. For simplicity, assume that

$$y \sim \text{Bern}(\exp(\theta)/(1 + \exp(\theta)))$$

and

$$\theta \sim \text{Normal}(\mu_0, \sigma)$$

The following steps are required to set up a random walk MH algorithm:

Step 0: Choose initial values, $\theta(0)$.

Step 1: Generate a proposed value of θ at t from $h(\theta_t - \theta_{t-1})$. We often use a Normal proposal distribution, so we draw θ_1 from $\text{Normal}(\theta_0, \text{sig}^2)$, where sig^2 is the variance of the Normal proposal distribution, a tuning parameter that we have to set.

Step 2: Calculate the ratio of posterior densities for the proposed and the original value for θ :

$$r = p(\theta^t|y)/p(\theta^{t-1}|y)$$

In our example,

$$r = \text{Bern}(y|\theta^t) * \text{Normal}(\theta^t|\mu_0, \sigma_0) / \text{Bernoulli}(y|\theta^{t-1}) * \text{Normal}(\theta^{t-1}|\mu_0, \sigma_0)$$

Step 3: Set

```
\begin{eqnarray*}
\theta^t &= & \theta^{t-1} \text{ with probability } \min(r,1) // \\
&= & \theta^{t-1} \text{ otherwise } \\
\end{eqnarray*}
```

We can do that by drawing a random number u from a $\text{Unif}(0,1)$ and accept θ^t if $u < r$. Repeat for $t = 1, 2, \dots$ a large number of samples. The **R** code for this MH sampler is provided in Panel 2 XXXX.

Panel 2: R code to run a Metropolis sampler on a simple Logit-Normal model.

```
Logreg.MH<-function(y=y, mu0=mu0, sig0=sig0, niter=niter) {
  out<-c()
  theta<-runif(1, -3,3) #initial value
  for (iter in 1:niter){
    theta.cand<-rnorm(1, theta, 0.2)
    loglike<-sum(dbinom(y, 1, exp(theta)/(1+exp(theta)), log=TRUE))
    logprior <- dnorm(theta,mu0 ,sig0, log=TRUE)
    loglike.cand<-sum(dbinom(y, 1, exp(theta.cand)/(1+exp(theta.cand)), log=TRUE))
    logprior.cand <- dnorm(theta.cand, mu0, sig0, log=TRUE)
    if (runif(1)<exp((loglike.cand+logprior.cand)-(loglike+logprior))){
      theta<-theta.cand
    }
    out[iter]<-theta
  }
  return(out)
}
```

The reason we sum the logs of the likelihood and the prior, rather than multiplying the original values, is simply computational. The product of small probabilities can be numbers very close to 0, which computers do not handle well. Thus we add the logarithms, sum, and exponentiate to achieve the desired result. Similarly, in case you have forgotten some elementary math, $x/y = \exp(\log(x) - \log(y))$, with the latter being favored for computational reasons.

Comparing MH sampling to Gibbs sampling, where all draws from the conditional distribution are used, in the MH algorithm we discard a portion of the candidate values, which inherently makes it less efficient than Gibbs sampling the price you pay for its increased generality. In Step 1 of the MH sampler we had to choose a variance for the Normal proposal distribution. Choice of the parameters that define our candidate distribution is also referred to as 'tuning', and it is important since adequate tuning will make your algorithm more efficient, i.e. your Markov chain will converge faster. The variance should be chosen so that (a) each step of drawing a new proposal value for θ can cover a reasonable distance in the parameter space, as otherwise, the random walk moves too slowly; and (b) proposal values are not rejected too often, as otherwise the random walk will 'get stuck' at specific values for too long. As a rule of thumb, your candidate value should be accepted in about 40% of all cases. Acceptance rates of 20–80% are probably ok, but anything below or above may well render your algorithm inefficient (this does not mean that it will give you wrong results only that you will need more iterations to converge to the posterior distribution). In practice, tuning will require some 'trial-and-error' and some common sense. Or, one can use an adaptive phase, where the tuning parameter is automatically adjusted until it reaches a user-defined acceptance rate, at which point the adaptive phase ends and the actual Markov chain begins. This is computationally a little more advanced. ? discuss this in more detail. It is important the samples drawn during the adaptive phase are discarded. You can easily check acceptance rates for the parameters you monitor (that are part of your output) using the `rejectionRate()` function of the package `coda` (we will talk more about this package a little later on). Do not let the term 'rejection rate' confuse you; it is simply 1 - acceptance rate. There may be parameters for example, individual values of a random effect or latent variables that you do not want to save, though, and in our next example we will show you a way to monitor their acceptance rates with a few extra lines of code.

7.3.3 Metropolis-within-Gibbs

One weakness of the MH sampler is that formulating the joint posterior when evaluating whether to accept or reject the candidate values for θ becomes increasingly complex or inefficient as the number of parameters in a model increases. It is probably going to sound like MCMC sampling is too good to be true but in these cases you can simply combine MH sampling and Gibbs sampling. You can use Gibbs sampling to break down your high-dimensional parameter space into easy-to-handle one-dimensional conditional distributions and use MH sampling for these conditional distributions. Better yet if you have some conjugacy in your model, you can use the more efficient Gibbs sampling for these parameters and one-dimensional MH for all the others. You have already seen the basics of how to build both types of algorithms, so we can jump straight into an example here and build a Metropolis-within-Gibbs algorithm.

7.4 GLMMs Poisson regression with a random effect

Let's assume a model that gets us closer to the problem we ultimately want to deal with a GLMM. Here, we assume we have Poisson counts, y , from i plots in j different study sites, and we believe that the counts are influenced by some plot-specific covariate, x , but that there is also a random site effect. So our model is:

$$y_{ij} \sim \text{Poisson}(lami_{ij})$$

$$lami_{ij} = \exp(a_j + b * x_i)$$

Let's use Normal priors on a and b ,

$$a_j \sim \text{Normal}(\mu_a, \sigma_a)$$

and

$$b \sim \text{Normal}(\mu_b, \sigma_b)$$

⁴ Since we want to estimate the random effect in this model, we do not specify μ_a and σ_a , but instead, estimate them as well, so we have to specify hyperpriors for these parameters:

$$\mu_a \sim \text{Normal}(\mu_{a0}, \sigma_{a0})$$

$$\sigma_a \sim \text{InvGamma}(a_0, b_0)$$

With the model fully specified, we can compile the full conditionals, breaking the multi-dimensional parameter space into one-dimensional components:

```
\begin{eqnarray*}
p(a_1|a_2,a_3,a_j,b,y) & \propto & p(y_{i1}|a_1,b) * p(a_1|\mu_a, \sigma_a) \backslash\backslash \\
& \propto & \text{Poisson}(y_{i1} | \exp(a_1 + b*x[j=1])) * \text{Normal}(a_1|\mu_a, \sigma_a) \\
\end{eqnarray*}
\begin{eqnarray*}
p(a_2|a_1,a_3,a_j,b,y) & \propto & p(y_{i2}|a_2,b) * p(a_2|\mu_a, \sigma_a) \backslash\backslash \\
& \propto & \text{Poisson}(y_{i2} | \exp(a_2 + b*x[j=1])) * \text{Normal}(a_2|\mu_a, \sigma_a) \\
\end{eqnarray*}
\text{and so on for all elements of } a.
\begin{eqnarray*}
p(b|a,y) & \propto & p(y|a,b) * p(b) \backslash\backslash \\
& \propto & \text{Poisson}(y | \exp(a + b*x)) * \text{Normal}(b|\mu_b, \sigma_b) \\
\end{eqnarray*}
```

Finally, we need to update the hyperparameters for a :

$$p(\mu_a|a) \propto p(a|\mu_a, \sigma_a) * p(\mu_a)$$

$$p(\sigma_a|a) \propto p(a|\mu_a, \sigma_a) * p(\sigma_a)$$

⁴Why is b a hyperparameter?

Since we assumed a to come from a Normal distribution, the choice of priors for μ_a Normal and σ_a Inverse Gamma leads to the same conjugacy we observed in our initial Normal model, so that both hyperparameters can be updated using Gibbs sampling.

Now let's build the updating steps for these full conditionals. Again, for the MH steps that update a and b we use Normal proposal distributions with standard deviations σ_a and σ_b .

First, we set the initial values $a(0)$ and $b(0)$. Then, starting with $a(1)$, we draw $a(1)$ from Normal($a(0)$, σ_a), calculate the conditional posterior density of $a(0)$ and $a(1)$ and compare their ratios,

$$r = \text{Poisson}(y(j=1)|\exp(a(1)+b*x)) * \text{Normal}(a(1)|\mu_a, \sigma_a) / \text{Poisson}(y(j=1)|\exp(a(0)+b*x)) * \text{Normal}(a(0)|\mu_a, \sigma_a)$$

and accept $a(1)$ with probability $\min(r, 1)$. We repeat this for all a 's.

For b , we draw $b(1)$ from Normal($b(0)$, σ_b), compare the posterior densities of $b(0)$ and $b(1)$,

$$r = \text{Poisson}(y|\exp(a+b(1)*x)) * \text{Normal}(b(1)|\mu_b, \sigma_b) / \text{Poisson}(y|\exp(a+b(0)*x)) * \text{Normal}(b(0)|\mu_b, \sigma_b),$$

and accept $b(1)$ with probability $\min(r, 1)$.

For μ_a and σ_a , we sample directly from the full conditional distributions (Eq XX and Eq XX):

$$\mu_a(1) \sim \text{Normal}(\mu_n, \sigma_n)$$

where $\mu_n = (\sigma_a(0)/\sigma_a(0) + n_a * \sigma_a(0)) * \mu_0 + (n_a * \sigma_a(0)/\sigma_a(0) + n_a * \sigma_a(0)) * \bar{a}$ and $\sigma_n = \sigma_a(0) * \sigma_a(0) / (\sigma_a(0) + n_a * \sigma_a(0))$. Here, \bar{a} is the current mean of the vector a , which we updated before, and n_a is the length of a . For σ_a we use $\sigma_a(1) \sim \text{InvGamma}(a_n, b_n)$, where $a_n = n_a/2 + a_0$, and $b_n = 1/2 \sum (a(1) - \mu_a(1))^2 + b_0$.

We repeat these steps over K iterations of the MCMC algorithm. In this example we may not want to save each value for a , but are only interested in their mean and standard deviation. Since these two parameters will change as soon as the value for one element in a changes, their acceptance rates will always be close to 1 and are not representative of how well your algorithm performs. To monitor the acceptance rates of parameters you do not want to save, you simply need to add a few lines of code into your updater to see how often the individual parameters are accepted. The full code for the MCMC algorithm of our Poisson GLMM in Panel 3 shows one way how to monitor acceptance of individual a 's.

Panel 3: R code for the Metropolis-within-Gibbs sampler for a Poisson regression with random intercepts.

```
Pois.reg<-function(y=y,site=site,mu0=mu0,sig0=sig0,a0=a0,b0=b0,
  mub=mub, sigb=sigb, niter=niter){
  lev<-length(unique(site))    #number of sites
```

```

439 a<-runif(lev,-5,5) #initial values a
440 b<-runif(1,0,5) #initial value b
441 mua<-mean(a)
442 siga<-sd(a)
443
444 out<-matrix(nrow=niter, ncol=3)
445 colnames(out)<-c('mua','siga','b')
446
447 for (iter in 1:niter) {
448
449   #update a
450   aUps<-0 #initiate counter for acceptance rate of a
451   for (j in 1:lev) { #loop over sites
452     a.cand<-rnorm(1, a[j], 0.1) #update intercepts a one at a time
453     loglike<- sum(dpois (y[site==j], exp(a[j] + b*x[site==j]), log=TRUE))
454     logprior<- dnorm(a[j], mua,siga, log=TRUE)
455     loglike.cand<- sum(dpois (y[site==j], exp(a.cand + b *x[site==j]), log=TRUE))
456     logprior.cand<- dnorm(a.cand, mua,siga, log=TRUE)
457     if (runif(1)< exp((loglike.cand+logprior.cand) (loglike+logprior))) {
458       a[j]<-a.cand
459       aUps<-aUps+1
460     }
461   }
462
463   if(iter %% 100 == 0) { #this lets you check the acceptance rate of a at every 100th iteration
464     cat(" Acceptance rates\n")
465     cat(" a =", aUps/lev, "\n")
466   }
467
468   #update b
469   b.cand<-rnorm(1, b, 0.1)
470   avec<-rep(a, times=c(rep(10,10)))
471   loglike<- sum(dpois (y, exp(avec + b*x), log=TRUE))
472   logprior<- dnorm(b, mub,sigb, log=TRUE)
473   loglike.cand<- sum(dpois (y, exp(avec + b.cand *x), log=TRUE))
474   logprior.cand<- dunif(b.cand, mub,sigb, log=TRUE)
475   if (runif(1)< exp((loglike.cand+logprior.cand) (loglike+logprior) )) {
476     b<-b.cand
477   }
478
479   #update mua using Gibbs sampling
480   abar<-mean(a)
481   mun<- (siga/(siga+lev*sig0))*mu0 + (lev*sig0/(siga+lev* sig0))*abar
482   sign <- (siga*sig0)/ (siga+lev*sig0)
483   mua<-rnorm(1,mun, sqrt(sign))
484
485   #update siga using Gibbs sampling
486   a0n<-lev/2 + a0
487   b0n<- 0.5 * (sum((a-mua)^2)) +b0
488   siga<-1/rgamma(1,shape=a0n, rate=b0n)

```

```

489
490 out[iter,]<-c(mua, sqrt(siga), b)
491
492 }
493
494 return(out)
495 }

```

7.4.1 Rejection sampling and slice sampling

While MH and Gibbs sampling are probably the most widely applied algorithms for posterior approximation, there are other options that work under certain circumstances and may be more efficient when applicable. WinBUGS applies these algorithms and we want you to be aware that there is more out there to approximate posterior distributions than Gibbs and MH. One alternative algorithm is rejection sampling. Rejection sampling is not an MCMC method, since each draw is independent of the others. The method can be used when the posterior $p(\theta|y)$ is not a known parametric distribution but can be expressed in closed form. Then, we can use a so-called envelope function, say, $g(\theta)$, that we can easily sample from, with the restriction that $p(\theta|y) < M * g(\theta)$. We then sample a candidate value for θ from $g(\theta)$, calculate $r = p(\theta|y)/M * g(\theta)$ and keep the sample with the probability r . M is a constant that has to be picked so that $r \in [0,1]$, for example by evaluating both $p(\theta|y)$ and $g(\theta)$ at n points and looking at their ratios. Rejection sampling only works well if $g(\theta)$ is similar to $p(\theta|y)$, and packages like WinBUGS use adaptive rejection sampling (?), where a complex algorithm is used to fit an adequate and efficient $g(\theta)$ based on the first few draws. Though efficient in some situations, rejection sampling does not work well with high-dimensional problems, since it becomes increasingly hard to define a reasonable envelope function. For an example of rejection sampling in the context of SCR models, see Chapter 9. Another alternative is slice sampling (?). In slice sampling, we sample uniformly from the area under the plot of $p(\theta|y)$. Considering a single univariate θ . Let's define an auxiliary variable, $U \sim Uniform(0, p(\theta|y))$. Then, θ can be sampled from the vertical slice of $p(\theta|y)$ at U (Figure 4):

```

521 \theta|U \sim \text{Unif}(B),

```

where $B = \theta : U < p(\theta|y)$

523

Slice sampling can be applied in many situations; however, implementing an efficient slice sampling procedure can be complicated. We refer the interested reader to chapter 7 of ? for a simple example. Both rejection sampling and slice sampling can be applied on one-dimensional conditional distributions within a Gibbs sampling setup.

⁵there are supposed to be equations in the caption of figure 4 but it kept causing errors




Figure 7.4: Slice sampling. For...

7.5 MCMC for closed capture-recapture Model

Mh

6

7.6 MCMC algorithm for the basic spatial capture-recapture model

By now you have seen how to build MCMC algorithms for some basic generalized linear models. Now, we'll walk you through the steps of building your own MCMC sampler for the basic SCR model (i.e. without any individual, site or time specific covariates) with both a Poisson and a binomial encounter process. As usual, we will have to go through two general steps before we write the MCMC algorithm:

- (1) Identify your model with all its components (including priors)
- (2) Recognize and express the full conditional distributions for all parameters

It is worthwhile to go through all of step 1 for an SCR model, but you have probably seen enough of step 2 in our previous examples to get the essence of how to express a full conditional distribution. Therefore, we will exemplify step 2 for some parameters and tie these examples directly to the respective R code.

Step 1 Identify your model

Recall the components of the basic SCR model with a Poisson encounter process from Chapter 3: We assume that individuals i , or rather, their activity centers s_i , are uniformly distributed across our state space S ,

$$s_i \sim U(S)$$

⁶Andy could move material from chapter 3 to here.

and that the number of times individual i encounters trap j , y_{ij} , is a random Poisson variable with mean λ_{ij} ,

$$y_{ij} \sim \text{Poisson}(\lambda_{ij})$$

The tie between individual location, movement and trap encounter rates is made by the assumption that λ_{ij} , is a decreasing function of the distance between s_i and j , D_{ij} , of the half-normal form

$$\lambda_{ij} = \lambda_{00} * \exp(-D_{ij}^2 / 2 * \sigma^2),$$

where λ_{00} is the baseline trap encounter rate at $D_{ij} = 0$ and σ controls the shape of the half-normal function.

In order to estimate the number of s_i in S , N , we use data augmentation (sect. 3.XYZ) and create $M-n$ all-0 encounter histories, where n is the number of individuals we observed and M is an arbitrary number that is larger than N . We estimate N by summing over the auxiliary data augmentation variables, z_i , which is 1 if the individual is part of the population and 0 if not, and assume that z_i is a random Bernoulli variable,

$$z_i \sim \text{Bern}(\psi)$$

To link the two model components, we modify our trap encounter model to

$$\lambda_{ij} = \lambda_{00} * \exp(-D_{ij}^2 / 2 * \sigma^2) * z_i.$$

The model has the following structural parameters, for which we need to specify priors ψ the Uniform (0,1) is required as part of the data augmentation procedure and in general is a natural choice of an uninformative prior for a probability; note that this is equivalent to a Beta(1,1) prior, which will come in handy later. s_i since s_i is a pair of coordinates it is two-dimensional and we use a uniform prior limited by the extent of our state-space over both dimensions. σ we can conceive several priors for sigma but let's assume an improper prior one that is Uniform over $(-\infty, \infty)$. We will see why this is convenient when we construct the full conditionals for sigma. λ_0 analogous, we will use a Uniform $(-\infty, \infty)$ improper prior for sigma. The parameter that is the objective of our modeling, N , is a derived parameter that we can simply obtain by summing all z_i 's:

$$N = \sum(z_i)$$

Step 2 - Construct the full conditionals Having completed step 1, let's look at the full conditional distributions for some of these parameters. We find that with improper priors, full conditionals are proportional only to the likelihood of the observations; for example, take the movement parameter sigma:

$$\text{Sig}|s, \lambda_{00}, z, y_{\text{propto}}[y|s, \lambda_{00}, z, \text{sig}] * [\text{sig}]$$

Since the improper prior implies that $[\text{sig}] \propto 1$, we can reduce this further to

$$\text{Sig}|s, \lambda_{00}, z, y_{\text{propto}}[y|s, \lambda_{00}, z, \text{sig}]$$

583 The R code to update sigma is shown in Panel 4.⁷

584 Panel 4: R code to update sigma within an MCMC algorithm for
585 an SCR model when using an improper prior

```
586
587
588 sig.cand <- rnorm(1, sigma, 0.1) #draw candidate value
589 if(sig.cand>0){ #automatically reject sig.cand that are <0
590   lam.cand <- lam0*exp(-(D*D)/(2*sig.cand*sig.cand))
591   ll<- sum(dpois(y, lam*z, log=TRUE))
592   llcand <- sum(dpois(y, lam.cand*z, log=TRUE))
593   if(runif(1) < exp( llcand - ll) ){
594     ll<-llcand
595     lam<-lam.cand
596     sigma<-sig.cand
597   }
598 }
599
```

600 These steps are analogous for lam0 and si and we will use MH steps for all of
601 these parameters. Similar to the random intercepts in our Poisson GLMM, we
602 update each si individually. Note that to be fully correct, the full conditional
603 for si contains both the likelihood and prior component, since we did not specify
604 an improper, but a Uniform prior on si. However, with a Uniform distribution
605 the probability density of any value is 1/(upper limit - lower limit) = constant.
606 Thus, the prior components are identical for both the current and the candidate
607 value and can be ignored (formally, when you calculate the ratio of posterior
608 densities, r, the identical prior component appears both in the numerator and
609 denominator, so that they cancel each other out).

610 We still have to update zi. The full conditional for zi is

$$z_i|y, \sigma, \lambda_0, \text{sprpto}[y|z, \sigma, \lambda_0, s] * [z_i]$$

611 and since $z_i \sim \text{Bernoulli}(\psi_i)$, the term has to be taken into account when
612 updating zi. The R code for updating zi is shown in Panel 5.

613 Panel 5: R code to update z

```
614
615 zUps <- 0 #set counter to monitor acceptance rate
616 for(i in 1:M) {
617   if(seen[i]) #no need to update seen individuals, since their z =1
618     next
619   zcand <- ifelse(z[i]==0, 1, 0)
620   llz <- sum(dpois(y[i,], lam[i,]*z[i], log=TRUE))
621   llcand <- sum(dpois(y[i,], lam[i,]*zcand, log=TRUE))

```

⁷ Somewhere in chapter 2 i added a comment about rejecting parameters outside of the parameter space as being an ok thing to do. Richard said he read something in Robert and Casellas book on that. Hopefully he can remember where and we can cite it back in Ch 2 and again here. It could be mentioned in a sentence or two up in the MCMC section.

```

622      prior <- dbinom(z[i], 1, psi, log=TRUE)
623      prior.cand <- dbinom(zcand, 1, psi, log=TRUE)
624      if(runif(1) < exp( (llcand+prior.cand) - (llz+prior) )) {
625          z[i] <- zcand
626          zUps <- zUps+1
627      }
628  }
629

```

630 ψ itself is a hyperparameter of the model, with an uninformative prior distribution of Unif(0,1) or Beta(1,1), so that

$$Psi|z \propto [z|psi] * Beta(1, 1)$$

632 The Beta distribution is the conjugate prior to the Binomial and Bernoulli distributions (remember that $z \sim Bernoulli(psi)$). The general form of a full conditional of a Beta-Binomial model with $y_i \sim Bernoulli(p)$ and $p \sim Beta(a, b)$ is

$$p(p|y) \propto Beta(a + \text{sum}(yi), b + n - \text{sum}(yi))$$

636 In our case, this means we update psi as follows:

```

637 si<-rbeta(1, 1+sum(z), 1 + M-sum(z))

```

638 These are all the building blocks you need to write the MCMC algorithm for the spatial null model with a Poisson encounter process. You can find the full R code (SCR0pois.R) in the online supplementary material.

641 7.6.1 SCR model with binomial encounter process

642 The equivalent SCR model with a binomial encounter process is very similar. Here, each individual i can only be detected once at any given trap j during a sampling occasion k . Thus

$$y_{ij} \sim Binomial(p_{ij}, K)$$

645 Where p_{ij} is some function of distance between \mathbf{s}_i and trap location \mathbf{x}_j . Here we use:

$$p_{ij} = 1 - \exp(-lam_{ij})$$

647 Recall from Chapter 2 that this is the complementary log-log (cloglog) link function, which constrains p_{ij} to fall between 0 and 1. For our MCMC algorithm that means that, instead of using a Poisson likelihood, $Poisson(y|sigma, lam0, s, z)$, we use a Binomial likelihood, $Binomial(y, K|sigma, lam0, s, z)$, in all the conditional distributions. As an example, Panel 6 shows the updating step for $lam0$ under a binomial encounter model. The full MCMC code for the binomial SCR can be found in the online supplements.

```

654 Panel 6: MCMC updater for lam0 in a SCR model with Binomial encounter
655 process and cloglog link function on detection. Here, pmat =
656 1-exp(-lam).
657
658     lam0.cand <- rnorm(1, lam0, 0.1)
659     if(lam0.cand > 0){ #automatically reject lam0.cand that are < 0
660         lam.cand <- lam0.cand*exp(-(D*D)/(2*sigma*sigma))
661         p.cand <- 1-exp(-lam.cand)
662         ll<- sum(dbinom(y, K, pmat *z, log=TRUE))
663         llcand <- sum(dbinom(y, K, p.cand *z, log=TRUE))
664         if(runif(1) < exp( llcand - ll) ){
665             ll<-llcand
666             pmat<-p.cand
667             lam0<- lam0.cand
668         }
669     }

```

Another possibility is to model variation in the individual and site specific detection probability, p_{ij} , directly, without any transformation, such that

```

672 pij<-p0 * exp(-Dij2/(2*sig^2))

```

and $p_0 = \{0, 1\}$. This formulation is analogous to how detection probability is modeled in distance sampling under a half-normal detection function; however, in distance sampling p_0 - detection of an individual on the transect line - is assumed to be 1 (?). Under this formulation the updater for lam_0 (equivalent to p_0 in Eq XX) becomes:

```

678     lam0.cand <- rnorm(1, lam0, 0.1)
679     if(lam0.cand > 0 & lam0.cand < 1 ){ #automatically reject lam0.cand that are
680         lam.cand <- lam0.cand*exp(-(D*D)/(2*sigma*sigma))
681         ll<- sum(dbinom(y, K, lam *z, log=TRUE)) #no transformation needed
682         llcand <- sum(dbinom(y, K, lam.cand *z, log=TRUE))
683         if(runif(1) < exp( llcand - ll) ){
684             ll<-llcand
685             lam<-lam.cand
686             lam0<- lam0.cand
687         }
688     }

```

7.6.2 Looking at model output

Now that you have an MCMC algorithm to analyze spatial capture-recapture data with, let's run an actual analysis so we can look at the output. As an example, we will use the bear data ... ⁸ You can use the same script provided back in Chapter XX to read in the data and build the augmented encounter history

⁸Does this data set come up before Ch6? If not, introduce data here. Or, Andy, would you rather use simulated data?

7.6. MCMC ALGORITHM FOR THE BASIC SPATIAL CAPTURE-RECAPTURE MODEL 33

array; then source the MCMC code for the binomial encounter model algorithm with the cloglog link and run 5000 iterations. This should take approximately 25 minutes.

```
> source('SCR0binom.txt')
> mod0<-SCR.0(y=bigTrap, X=trapmat, M=M, xl=xl, xu=xu, yl=yl, yu=yu, K=8, niter=5000)
```

Before, we used simple R commands to look at model results. However, there is a specific R package to summarize MCMC simulation output and perform some convergence diagnostics package coda (?). Download and install coda, then convert your model output to an mcmc object

```
> chain<-mcmc(mod0)
```

which can be used by coda to produce MCMC specific output.

Markov chain time series plots

Start by looking at time series plots of your Markov chains using `plot(chain)`. This command produces a time series plot and marginal posterior density plots for each monitored parameter, similar to what we did before using the `hist()` and `plot()` commands (Fig. 5). Time series plots will tell you several things: First, the way the chains move through the parameter space gives you an idea of whether your MH steps are well tuned. If chains were constant over many iterations you would probably need to decrease the tuning parameter of the (Normal) proposal distribution. If a chain moves along some gradient to a stationary state very slowly, you may want to increase the tuning parameter so that the parameter space is explored more efficiently.

Second, you will be able to see if your chains converged and how many initial simulations you have to discard as burn-in. In the case of the chains shown in Figure 5, we would probably consider the first 750 - 1000 iterations as burn-in, as afterwards the chains seem to be fairly stationary.

A word of caution about chain convergence

Since we do not know what the stationary posterior distribution of our Markov chain should look like (this is the whole point of doing an MCMC approximation), we effectively have no means to assess whether it has converged to this desired distribution or not. As mentioned before, the only certainty is that a Markov chain will *eventually* converge to its stationary distribution, but no-one can tell us how long this will take. Also, you only now the part of your posterior distribution that the Markov chain has explored so far for all you know the chain could be stuck in a local maximum, while other maxima remain completely undiscovered. Acknowledging that there is truly nothing we can do to ever proof convergence of our MCMC chains, there are several things we can do to increase the degree of confidence we have about the convergence of our chains. One option, and that advocated by what we will loosely call the WinBUGS community, is to run several Markov chains and to start them off at




Figure 7.5: Time series and posterior density plots for sigma and lam0.

different initial values that are overdispersed relative to the posterior distribution. Such initial values help to explore different areas of the parameter space simultaneously; if after a while all chains oscillate around the same average value, chances are good that they indeed converged to the posterior distribution. Gelman and Rubin came up with a diagnostic statistic that essentially compares within-chain and between-chain variance to check for convergence of multiple chains (?). Of course, running several parallel chains is computationally expensive. Extra computational demands are not the only and by no means the major concern some people voice when it comes to running several parallel MCMC chains to assess convergence. Again, consider the fact that we do not know anything about the true form of the posterior distribution we are trying to approximate. How do we, then, know how to pick overdispersed initial values? We don't all we can do is pick overdispersed values relative to our expectations of what the posterior should look like. To use a quote from the home page of Charlie Geyer, a Bayesian statistician from the University of Minnesota, "If you don't know any good starting points [...], then restarting the sampler at many bad starting points is [...] part of the problem, not part of the solution." (<http://users.stat.umn.edu/~charlie/mcmc/diag.html>). His suggestion is that your only chance to discover a potential problem with your MCMC sampler is to run it for a very long time. But again, there is no way of knowing how long is long enough. It is up to you to decide, which school of thoughts appeals more to you one long versus several parallel Markov chains. Irrespectively, part of developing an MCMC sampler should be to make sure (within reasonable limits) that you are not missing regions of high posterior density because of the way you specify your starting values. Once you have explored the behavior of your chain under a reasonable range of starting values, you may feel comfort-

able enough to run only one long chain. The fact that convergence cannot be proven does not mean that you should not look for potential problems in your MCMC sampler. Some problems are easily detected using simple plots, such as the time series plots we discussed above. If the overall trajectory of your chain at the end of your simulations is still upward or downward, your chain clearly has not converged and you need to run your model much longer. If you run several parallel chains and their stationary distributions look different, you may be looking at a multi-modal posterior or a problem with your sampler. With these words of caution, let's get back to looking at our model output.

7.6.3 Posterior density plots

The `plot()` command also produces posterior density plots and it is worthwhile to look at those carefully. For parameters with priors that have bounds (e.g. Uniform over some interval), you will be able to see if your choice of the prior is truncating the posterior distribution. In the context of SCR models, this will mostly involve our choice of M , the size of the augmented data set. If the posterior of N has a lot of mass concentrated close to M (or equivalently the posterior of ψ has a lot of mass concentrated close to 1), as in the example in Figure 6, we have to re-run the analysis with a larger M . A flat posterior plot shows you that the parameter essentially cannot be identified there may not be enough information in your data to estimate model parameters and you may have to consider a simpler model. Finally, posterior density plots will show you if the posterior distribution is symmetrical or skewed if the distribution has a heavy tail, using the mean as a point estimate of your parameter of interest may be biased and you may want to opt for the median or mode instead.

7.6.4 Serial autocorrelation and effective sample size

Even when we can be relatively confident that our chains have converged, the subsequent samples generated from a Markov chain are not iid samples from the posterior distribution, due to the correlation amongst samples introduced by the Markov process. As a consequence, the variance of the mean cannot simply be derived with the standard variance estimator, which takes into account the sample size (here, number of iterations). Rather, the sample size has to be adjusted to account for the autocorrelation in subsequent samples (see Chapter 8 in ? for more details). This adjusted sample size is referred to as the effective sample size. Checking the degree of autocorrelation in your Markov chains and estimating the effective sample size your chain has generated should be part of evaluating your model output. If you use WinBUGS through the R2WinBUGS package, the `print()` command will automatically return the effective sample size for all monitored parameters. In the coda package there are several functions you can use to do so. `effectiveSize()` will directly give you an estimate of the effective sample size for you parameters:

```
> effectiveSize(chain)
```




Figure 7.6: Time series and posterior density plots of ψ and N for the bear data set truncated by the upper limit of M (500).

```

801      sigma      lam0      psi      N
802      3.930303 78.259159 30.436348 32.047392

```

Alternatively, you can use the `autocorr.diag()` function, which will show you the degree of autocorrelation for different lag values (which you can specify within the function call, we use the defaults below):

```

806 > autocorr.diag(mcmc(mod))
807      sigma      lam0      psi      N
808 Lag 0  1.0000000 1.0000000 1.0000000 1.0000000
809 Lag 1  0.9979948 0.9494134 0.9847503 0.9774201
810 Lag 5  0.9915567 0.8038168 0.9111951 0.9113525
811 Lag 10 0.9836016 0.6714021 0.8462108 0.8509803
812 Lag 50 0.8985337 0.1983780 0.6138516 0.6233994

```

Whichever function you use, if you find that your supposedly long Markov chain has not generated enough pseudo-iid samples, you should consider a longer run. In the present case we see that autocorrelation is especially high for the parameter σ and our effective sample size for this parameter is 4!⁹ This means we would have to run the model for much longer to obtain a reasonable effective sample size. Unfortunately, with many SCR models we observe high degrees of serial autocorrelation, which means we have to run long chains to obtain enough samples that can be considered iid, in order to obtain reasonable estimates of our parameters and their variances. What exactly constitutes a reasonable effective sample size is hard to say, but as a rule of thumb you

⁹Anyone have any idea how the autocorrelation in σ could be reduced?

should probably aim at several hundreds of these pseudo-iid samples. A more meaningful measure of whether you've run your chain for enough iterations is the time-series or Monte Carlo error - the 'noise' introduced into your samples by the stochastic MCMC process - which we introduced in Chapter 2. The MC error decreases with increasing sample size and its magnitude can thus be controlled by adjusting the length of the Markov chain. As a rule of thumb, the MC error should be 1% or less of the parameter estimate. Once you have reached this level, the estimates of the mean, standard error and 95% quantiles should no longer change significantly with additional iterations. For highly correlated samples, it will take more iterations to reduce the MC error. In coda, the MC error is given as part of the summary results (see below). Another option to deal with the serial autocorrelation of samples is to 'thin' Markov chains by some rate r and save only every r -th iteration. But as discussed in Chapter 2, this is not efficient and should only be applied if needed for practical reasons (e.g. a large number of parameters and iterations may force you to thin your samples so you object storing the model output does not become unmanageably large). For now, let's continue using this small set of samples to continue looking at the output.

7.6.5 Summary results

Now that we checked that our chains apparently have converged and pretending that we have generated enough samples from the posterior distribution, we can look at the actual parameter estimates. The `summary()` function will return two sets of results: the mean parameter estimates, with their standard deviation, the naive standard error - i.e. your regular standard error calculated for K (= number of iterations) samples without accounting for serial autocorrelation - and the corrected MC error (Time-series SE), which accounts for autocorrelation. In WinBUGS, this latter value is referred to as MC error and is only given in the log output within BUGS itself. You should adjust the `summary()` call by removing the burn-in from calculating parameter summary statistics. To do so, use the `window()` command, which lets you specify at which iteration to start 'counting'. In contrast to WinBUGS, which requires you to set the burn-in length before you run the model, this command gives us full flexibility to make decisions about the burn-in after we have seen the trajectories of our Markov chains. For our example, `summary(window(chain, start=1001))` returns the following output:

```
Iterations = 1001:5000
Thinning interval = 1
Number of chains = 1
Sample size per chain = 4000

1. Empirical mean and standard deviation for each variable,
   plus standard error of the mean:
```

```

866           Mean          SD   Naive SE Time-series SE
867 sigma    1.9986  0.13805 0.0021827      0.016091
868 lam0     0.1096  0.01523 0.0002407      0.001401
869 psi      0.6113  0.09148 0.0014465      0.010734
870 N        489.8535 71.79695 1.1352094      8.431119

```

```

871
872 2. Quantiles for each variable:

```

```

873
874           2.5%       25%       50%       75%      97.5%
875 sigma    1.75780    1.89847    1.9900    2.0944    2.2772
876 lam0     0.08357    0.09824    0.1087    0.1192    0.1427
877 psi      0.45110    0.54838    0.6052    0.6639    0.8192
878 N        366.00000  440.00000  485.0000  530.0000  654.0000

```

```

879     Looking at the MC errors, we see that in spite of the high autocorrelation, the
880     MC error for sigma is below the 1Our algorithm gives us a posterior distribution
881     of N, but we are usually interested in the density, D. Density itself is not a
882     parameter of our model, but we can derive a posterior distribution for D by
883     dividing each value of N (N at each iteration) by the area of the state-space
884     (here 3032.719 km2) and we can use summary statistics of this distribution to
885     characterize D:

```

```

886 > summary(window(chain[,4]/ 3032.719, start=1001))
887 Iterations = 1001:5000
888 Thinning interval = 1
889 Number of chains = 1
890 Sample size per chain = 4000

```

```

891
892 1. Empirical mean and standard deviation for each variable,
893    plus standard error of the mean:

```

```

894
895           Mean          SD          Naive SE Time-series SE
896    0.1615229    0.0236741    0.0003743      0.0027801

```

```

897
898 2. Quantiles for each variable:

```

```

899
900    2.5%    25%    50%    75%    97.5%
901 0.1207 0.1451 0.1599 0.1748 0.2156

```

```

902 If we compare our mean density of 0.16/km2 (and other parameters) with results
903 from the same model run in secr and WinBUGS in Chapter XX, we see that
904 estimates are almost identical (Table 1).

```

905 7.6.6 Other useful commands

```

906 While inspecting the time series plot gives you a first idea of how well you
907 tuned your MH algorithm, use rejectionRate() to obtain the rejection rates (1
908 acceptance rates) of the parameters that are written to your output:

```

```

909 > rejectionRate(chain)
910      sigma      lam0      psi      N
911 0.44108822 0.77675535 0.00000000 0.01940388

```

Recall that rejection rates should lie between 0.2 and 0.8, so our tuning seems to have been appropriate here. Psi is never rejected since we update it with Gibbs sampling, where all candidate values are kept. And since N is the sum of all z, all it takes for N to change from one iteration to the next are small changes in the z-vector, so the rejection rate of N is always low. If you have run several parallel chains, you can combine them into a single mcmc object using the `mcmc.list()` command on the individual chains (note that each chain has to be converted to an mcmc object before combining them with `mcmc.list()`). You can then easily obtain the Gelman-Rubin diagnostic (?), in WinBUGS called R-hat, using `gelman.diag()`, which will indicate if all chains have converged to the same stationary distribution. For details on these and other functions, see the coda manual, which can be found together with the package on the CRAN mirror.

7.7 Manipulating the state-space

So far, we have constrained the location of the activity centers to fall within the outermost coordinates of our rectangular state space by posing upper and lower bounds for x and y. But what if S has an irregular shape maybe there is a large water body we would like to remove from S, because we know our terrestrial study species does not occur there. Or the study takes place in a clearly defined area such as an island. As mentioned before, this situation is difficult to handle in WinBUGS. In some simple cases we can adjust the state space by setting SX_i to be some function of SY_i or vice versa. In this manner, we can cut off corners of the rectangle to approximate the actual state space. In R, we are much more flexible, as we can use the actual state-space polygon to constrain out si. ¹⁰To illustrate that, let's look at a camera trapping study of Florida panthers (*Puma concolor coryi*) conducted in the Picayune Strand Restoration Project (PSRP) area, southwest Florida (Fig. 7), by XXX, and financed by XXX. In the 1960ies the PSRP area was slated for housing development, but then bought back by the State of Florida and is currently being restored to its original hydrology and vegetation. In an effort to estimate the density of the local Florida panther population, 98 camera traps were operated in the area for 21 months between 2005 and 2007. Florida panthers are wide-ranging animals and in order to account for their wide movements, the state-space was defined as the trapping grid buffered by 15 km around its outermost coordinates. However, the resulting rectangle contained some ocean in its southwestern corner (Fig. 7). In order to precisely describe the state-space, the ocean has to be removed. You can create a precise state-space polygon in ArcGIS and read it into R, or create the polygon directly within R. In the present case we intersected two shape files

¹⁰ Have to check if we can use panther stuff for the book; otherwise, use raccoon example.




Figure 7.7: Rectangular state-space for a Florida panther camera trapping study in the PSRP area (grey outline, red block inset map of Florida) contain some ocean (white) that needs to be removed from the state-space.

one of the state of Florida and one of the rectangle defined by a strip of 15 km around the camera-trapping grid. While you will most likely have to obtain the shapefile describing the landscape of and around your trapping grid (coastlines, water bodies etc.) from some external source, a polygon shapefile buffering your outermost trapping grid coordinates can easily be written in R.

If `xmin`, `xmax`, `ymin` and `ymax`, mark the outermost `x` and `y` coordinates of your trapping grid and `b` is the distance you want to buffer with, load the package `shapefiles` (?) and use:

```

958 xl= xmin-b
959 xu= xmax+b
960 yl= ymin-b
961 yu= ymax+b
962
963 dd <- data.frame(Id=c(1,1,1,1,1),X=c(xl,xu,xu,xl,xl),Y=c(yl,yl,yu,yu,yl)) #create data
964 ddTable <- data.frame(Id=c(1),Name=c("Item1"))
965 ddShapefile <- convert.to.shapefile(dd, ddTable, "Id", 5) #convert #to shapefile, type
966 write.shapefile(ddShapefile, 'c:/, arcgis=T) # save to location of #choice

```

You can read shapefiles into R loading the package `maptools` (?) and using the function `readShapeSpatial()`. Make sure you read in shapefiles in UTM format, so that units of the trap array, the movement parameter `sigma` and the state-space are all identical. Intersection of polygons can be done in R also, using the package `rgeos` (?) and the function `gIntersect()`. The area of your single - polygon can be extracted directly from the state-space object `SSp`:


```
973 > area <- SSp@polygons[[1]]@Polygons[[1]]@area /1000000
```

974 Note that dividing by 1000000 will return the area in km² if your coordi-
 975 nates describing the polygon are in UTM. If your state-space consists of several
 976 disjunct polygons, you will have to sum the areas of all polygons to obtain the
 977 size of the state-space. To include this polygon into our MCMC sampler we
 978 need one last spatial R package `sp` (?), which has a function, `over()`, which
 979 allows us to check if a pair of coordinates falls within a polygon or not. All we
 980 have to do is embed this new check into the updating steps for `s`:

```
981 Scand <- as.matrix(cbind(rnorm(M, S[,1], 2),
982                          rnorm(M, S[,2], 2))) #draw candidate value
983
984 Scoord<-SpatialPoints(Scand*1000) #convert to spatial points on UTM (m) scale
985 SinPoly<-over(Scoord,SSp) # check if scand is within the polygon
986
987 for(i in 1:M) {
988   if(is.na(SinPoly[i])==FALSE) { #if scand falls within polygon, continue update
989     [rest of the updating step remains the same]
```

990 Note that it is much more time-efficient to draw all M candidate values for s
 991 and check once if they fall within the state-space, rather than running the `over()`
 992 command for every individual pair of coordinates. To make sure that our initial
 993 values for s also fall within the polygon of S , we use the function `runifpoint()`
 994 from the package `spatstat` (?), which generates random uniform points within
 995 a specified polygon. You'll find this modified MCMC algorithm in the online
 996 supplementary material (SCR0poisSSp). Finally, observe that we are converting
 997 candidate coordinates of S back to meters to match the UTM polygon. In all
 998 previous examples, for both the trap locations and the activity centers we have
 999 used UTM coordinates divided by 1000 to estimate sigma on a km scale. This is
 1000 adequate for wide ranging individuals like bears. In other cases you may center
 1001 all coordinates on 0. No matter what kind of transformation you use on your
 1002 coordinates, make sure to always convert candidate values for S back to the
 1003 original scale (UTM) before running the `over()` command.

1004 7.8 MCMC software packages

1005 Throughout most of this book we will use WinBUGS and, occasionally, JAGS
 1006 to run MCMC analyses. Here, we will briefly discuss the main pros and cons of
 1007 these two programs as well as WinBUGS successor OpenBUGS. You can find
 1008 scripts to simulate data and run the basic SCR model in all three programs in
 1009 the online supplementary material (simSCR0poisBUGS).

1010 7.8.1 WinBUGS

1011 In a nutshell, WinBUGS (and the other programs) do everything that we just
 1012 went through in this chapter (and quite a bit more). Looking through your

model, WinBUGS determines which parameters it can use standard Gibbs sampling for (i.e. for conjugate full conditional distributions). Then, it determines, in the following hierarchy, whether to use adaptive rejection sampling, slice sampling or in the 'worst' case Metropolis-Hastings sampling for the other full conditionals (?). If it uses MH sampling, it will automatically tune the updater so that it works efficiently. While WinBUGS is a convenient piece of software that is still widely used, its major drawback is that it is no longer being developed, i.e. no new functions or distributions are added and no bugs are fixed.

7.8.2 OpenBUGS

OpenBUGS is essentially the successor of WinBUGS. While the latter is no longer worked on, OpenBUGS is constantly developed further. The name 'OpenBUGS' refers to the software being open source, so users do not need to download a license key, like they have to for WinBUGS (although the license key for WinBUGS is free and valid for life).

Compared to WinBUGS, OpenBUGS has a lot more built-in functions. The method of how to determine the right updater for each model parameter has changed and the user can manually control the MCMC algorithm used to update model parameters. Several other changes have been implemented in OpenBUGS and a detailed list of differences between the two BUGS versions, can be found at <http://www.openbugs.info/w/OpenVsWin>

While OpenBUGS is a useful program for a lot of MCMC sampling applications, for reasons we do not understand, simple SCR models do not converge in OpenBUGS. It is therefore advisable that you check any OpenBUGS SCR model results against result from WinBUGS. Also, currently, the R package BRugs (?) necessary for running OpenBUGS through R has problems with 64-bit machines, so you may have to use the 32-bit version of R and OpenBUGS in order to make it work. The BUGS project site at <http://www.openbugs.info> provides a lot of information on and download links for OpenBUGS.

There is an extensive help archive for both WinBUGS and OpenBUGS and you can subscribe to a mailing list, where people pose and answer questions of how to use these programs at <http://www.mrc-bsu.cam.ac.uk/bugs/overview/list.shtml>

7.8.3 JAGS Just Another Gibbs Sampler

JAGS, currently at Version 3.1.0, is another free program for analysis of Bayesian hierarchical models using MCMC simulation. Originally, JAGS was the only program using the BUGS language that would run on operating systems other than the 32 bit Windows platforms. By now, there are OpenBUGS versions for Linux or Macintosh machines. JAGS 'only' generates samples from the posterior distribution; analysis of the output is done in R either by running JAGS through R using either the packages rjags (?) or R2jags (?), or by using coda on your JAGS output. The program, manuals and rjags can be downloaded at <http://sourceforge.net/projects/mcmc-jags/files/> When run from within R

using the package `rjags` or `R2jags`, writing a JAGS model is virtually identical to writing a WinBUGS model. However, some functions may have slightly different names and you can look up available functions and their use in the JAGS manual. One potential downside is that JAGS can be very particular when it comes to initial values. These may have to be set as close to truth as possible for the model to start. Although JAGS lets you run several parallel Markov chains, this characteristic interferes with the idea of using overdispersed initial values for the different chains. Also, we have occasionally experienced JAGS to crash and take the R GUI with it. Only re-installing both JAGS and `rjags` seemed to solve this problem. On the plus side, JAGS usually runs a little faster than WinBUGS, sometimes considerably faster (see section 4.XYZ), is constantly being developed and improved and it has a variety of functions that are not available in WinBUGS. For example, JAGS allows you to supply observed data for some deterministic functions of unobserved variables. In BUGS we cannot supply data to logical nodes. Another useful feature is that the adaptive phase of the model (the burn-in) is run separately from the sampling from the stationary Markov chains. This allows you to easily add more iterations to the adaptive phase if necessary without the need to start from 0. There are other, more subtle differences and there is an entire manual section on differences between JAGS and OpenBUGS. For questions and problems there is a JAGS forum online at <http://sourceforge.net/projects/mcmc-jags/forums/forum/610037>.¹¹

7.9 Summary and Outlook

While there are a number of flexible and extremely useful software packages to perform MCMC simulations, it sometimes is more efficient to develop your own MCMC algorithm. Building an MCMC code follows three basic steps: Identify your model including priors and express full conditional distributions for each model parameter. If full conditionals are parametric distributions, use Gibbs sampling to draw candidate parameter values from this distributions; otherwise use Metropolis-Hastings sampling to draw candidate values from a proposal distribution and accept or reject them based on their posterior probability densities. These custom-made MCMC algorithms give you more modeling flexibility than existing software packages, especially when it comes to handling the state-space: In BUGS (and JAGS for that matter) we define a continuous rectangular state-space using the corner coordinates to constrain the Uniform priors on the activity centers s . But what if a continuous rectangle isn't an adequate description of the state-space? In this chapter we saw that in R it only takes a few lines of code to use any arbitrary polygon shapefile as the state-space, which is especially useful when you are dealing with coastlines or large bodies of water that need removing from the state-space. Another example is the SCR R package `SPACECAP` (?) that was developed because implementation of an SCR model with a discrete state-space was inefficient in WinBUGS. Another

¹¹As we make progress on the book, let's be sure to add linkages to places where we use JAGS in examples.

1096 situations in which using BUGS/JAGS becomes increasingly complicated or in-
1097 efficient is when using point processes other than the Uniform Poisson point
1098 process which underlies the basic SCR model (see Chapter X). In the Chapters
1099 9 and XX you will see examples of different point processes, implemented using
1100 custom-made MCMC algorithms.¹² Finally, the Chapters XX and XX deal
1101 with unmarked or partially marked populations using hand-made MCMC al-
1102 gorithms to handle the (partially) latent individual encounter histories. While
1103 some of these models can be written in BUGS/JAGS,¹³ they are painstakingly
1104 slow; others cannot be implemented in BUGS/JAGS at all. In conclusion, while
1105 you can certainly get by using BUGS/JAGS for standard SCR models, knowing
1106 how to write your own MCMC sampler allows you to tailor these models to your
1107 specific needs.

¹²Richard, Beth expand on that?

¹³the Poisson one for partially marked we wrote in BUGS and it should work with a known number of marked; the Bernoulli in JAGS with the `dsum()` function should work for the fully unknown; maybe some others? I don't remember. We may have to try writing the others before saying that they don't work in BUGS/JAGS; they are certainly much faster in R, though.

1108 Chapter 8

1109 Goodness of Fit and stuff

1110 Chapter 9

1111 Covariate models

1112 Chapter 10

1113 Inhomogeneous Point 1114 Process

1115 Chapter 11

1116 Open models

Bibliography

- Baddeley, A. and Turner, R. (2005), “Spatstat: an R package for analyzing spatial point patterns,” *Journal of Statistical Software*, 12, 1–42, ISSN 1548-7660.
- Bivand, R. and Rundel, C. (2011), *rgeos: Interface to Geometry Engine - Open Source (GEOS)*, r package version 0.1-8.
- Buckland, S. T. (2001), *Introduction to distance sampling: estimating abundance of biological populations*, Oxford, UK: Oxford University Press.
- Casella, G. and George, E. I. (1992), “Explaining the Gibbs sampler,” *American Statistician*, 46, 167–174.
- Gelfand, A. and Smith, A. (1990), “Sampling-based approaches to calculating marginal densities,” *Journal of the American statistical association*, 85, 398–409.
- Gelman, A., Carlin, J. B., Stern, H. S., and Rubin, D. B. (2004), *Bayesian data analysis, second edition.*, Boca Raton, Florida, USA: CRC/Chapman & Hall.
- Geman, S. and Geman, D. (1984), “Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PAMI-6, 721–741.
- Gilks, W. and Wild, P. (1992), “Adaptive rejection sampling for Gibbs sampling,” *Applied Statistics*, 41, 337–348.
- Gilks, W. R., Thomas, A., and Spiegelhalter, D. J. (1994), “A Language and Program for Complex Bayesian Modelling,” *Journal of the Royal Statistical Society. Series D (The Statistician)*, 43, 169–177, ArticleType: primary_article / Issue Title: Special Issue: Conference on Practical Bayesian Statistics, 1992 (3) / Full publication date: 1994 / Copyright 1994 Royal Statistical Society.
- Gopalaswamy, A. M., Royle, A. J., Hines, J., Singh, P., Jathanna, D., Kumar, N. S., and Karanth, K. U. (2011), *A Program to Estimate Animal Abundance and Density using Spatially-Explicit Capture-Recapture*, r package version 1.0.4.

- Hastings, W. (1970), “Monte Carlo sampling methods using Markov chains and their applications,” *Biometrika*, 57, 97–109.
- Lewin-Koh, N. J., Bivand, R., contributions by Edzer J. Pebesma, Archer, E., Baddeley, A., Bibiko, H.-J., Dray, S., Forrest, D., Friendly, M., Giraudoux, P., Golicher, D., Rubio, V. G., Hausmann, P., Hufthammer, K. O., Jagger, T., Luque, S. P., MacQueen, D., Niccolai, A., Short, T., Stabler, B., and Turner, R. (2011), *maptools: Tools for reading and handling spatial objects*, r package version 0.8-10.
- Link, W. A. and Barker, R. J. (2009), *Bayesian Inference: With Ecological Applications*, London, UK: Academic Press.
- Metropolis, N., Rosenbluth, A., Rosenbluth, M., Teller, A., Teller, E., et al. (1953), “Equation of state calculations by fast computing machines,” *The journal of chemical physics*, 21, 1087–1092.
- Metropolis, N. and Ulam, S. (1949), “The Monte Carlo method,” *Journal of the American Statistical Association*, 44, 335–341.
- Neal, R. (2003), “Slice sampling,” *Annals of Statistics*, 31, 705–741.
- Pebesma, E. and Bivand, R. (2011), *Package ‘sp’*, r package version 0.9-91.
- Plummer, M. (2011), *rjags: Bayesian graphical models using MCMC*, r package version 3-5.
- Plummer, M., Best, N., Cowles, K., and Vines, K. (2006), “CODA: Convergence Diagnosis and Output Analysis for MCMC,” *R News*, 6, 7–11.
- Robert, C. P. and Casella, G. (2004), *Monte Carlo statistical methods*, New York, USA: Springer.
- (2010), *Introducing Monte Carlo Methods with R*, New York, USA: Springer.
- Roberts, G. O. and Rosenthal, J. S. (1998), “Optimal scaling of discrete approximations to Langevin diffusions,” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 60, 255–268.
- Spiegelhalter, D., Thomas, A., Best, N., and Lunn, D. (2003), *WinBUGS User Manual Version 1.4*.
- Stabler, B. (2006), *shapefiles: Read and Write ESRI Shapefiles*, r package version 0.6.
- Su, Y.-S. and Yajima, M. (2011), *R2jags: A Package for Running jags from R*, r package version 0.02-17.
- Thomas, A., O’Hara, B., Ligges, U., and Sturtz, S. (2006), “Making BUGS Open,” *R News*, 6, 12–17.