

₁ Chapter 1

₂ Introduction

³ Chapter 2

⁴ GLMS and WinBUGS

⁵ Chapter 3

⁶ Closed population models

⁷ Chapter 4

⁸ Fully Spatial ⁹ capture-recapture models

¹⁰ **Chapter 5**

¹¹ **Other observation models**

¹² **Chapter 6**

¹³ **MCMC details**

Chapter 7

MCMC for Spatial Capture-Recapture

7.1 Introduction

In this chapter we will dive a little deeper into Markov chain Monte Carlo (MCMC) sampling. We will construct custom MCMC samplers in **R**, starting with easy-to-code GLMs and GLMMs and moving on to simple CR and SCR models. Finally, we will illustrate some alternative ready-to-use software packages for MCMC sampling. We will NOT provide exhaustive background information on the theory and justification of MCMC sampling there are entire books dedicated to that subject and we refer you to Robert and Casella (2004) and Robert and Casella (2010). Rather we aim to provide you with enough background and technical know-how to start building your own MCMC samplers for SCR models in **R**. You will find that quite a few topics that come up in this chapter have already been covered in previous chapters, particularly the introduction into Bayesian analysis in Chapt. 2. To keep you from having to leaf back and forth we will in some places briefly review aspects of Bayesian analysis, but we try to focus on the more technical issues of building MCMC samplers.

7.1.1 Why build your own MCMC algorithm?

The standard program we have used so far to run MCMC analyses is WinBUGS (Gilks et al., 1994). The wonderful thing about WinBUGS is that it will automatically use the most appropriate and efficient form of MCMC sampling for the model specified by the user.

The fact that we have such a Swiss Army knife type of MCMC machine begs the question: Why would anyone want to build their own MCMC algorithm?

For one, there are a limited number of distributions and functions implemented in WinBUGS. While OpenBUGS provides more options, some more complex models may be impossible to build within these programs. A very simple example from spatial capture-recapture that can give you a headache in WinBUGS is when your state-space is an irregular-shaped polygon, rather than an ideal rectangle that can be characterized by four pairs of coordinates. It is easy to restrict activity centers to any arbitrary polygon in R using an ESRI shapefile (and we will show you an example in a little bit), but you cannot use a shape file in a BUGS model.

Sometimes implementing an MCMC algorithm in R may be faster than in WinBUGS - especially if you want to run simulation studies where you have hundreds or more simulated data sets, several years' worth of data or other large models, this can be a big advantage.

Finally, building your own MCMC algorithm is a great exercise to understand how MCMC sampling works. So while using the BUGS language requires you to understand the structure of your model, building an MCMC algorithm requires you to think about the relationship between your data, priors and posteriors, and how these can be efficiently analyzed and characterized. Not to mention that, if you are an R junkie, it can actually be fun. However, if you don't think you will ever sit down and write your own MCMC sampler, consider skipping this chapter - apart from coding it will not cover anything SCR-related that is not covered by other, more model-oriented chapters as well.

7.2 MCMC and posterior distributions

As mentioned in Chapter 2, MCMC is a class of simulation methods for drawing (correlated) random numbers from a target distribution, which in Bayesian inference is the posterior distribution. As a reminder, the posterior distribution is a probability distribution for an unknown parameter, say θ , given a set of observed data and its prior probability distribution (the probability distribution we assign to a parameter before we observe data). The great benefit of computing the posterior distribution of θ is that it can be used to make probability statements about θ , such as the probability that θ is equal to some value, or the probability that θ falls within some range of values. The posterior distribution summarizes all we know about a parameter and thus, is the central object of interest in Bayesian analysis. Unfortunately, in many if not most practical applications, it is nearly impossible to directly compute the posterior. Recall Bayes theorem:

$$p(\theta|y) = p(y|\theta) * p(\theta) / p(y), \quad (7.1)$$

where θ is the parameter of interest, y is the observed data, $p(\theta|y)$ is the posterior, $p(y|\theta)$ the likelihood of the data conditional on θ , $p(\theta)$ the prior probability of θ , and, finally, $p(y)$ is the marginal probability of the data, which can also be written as

$$p(y) = \int p(y|\theta) * p(\theta) d\theta$$

80 This marginal probability is a normalizing constant that ensures that the
 81 posterior integrates to 1. This integral is often hard or impossible to evaluate,
 82 unless you are dealing with a really simple model. For example, consider that
 83 you have a Normal model, with a set of n observations, y that come from a
 84 Normal distribution:

$$y \sim \text{Normal}(\mu, \sigma),$$

85 where σ is known and our objective is to obtain an estimate of μ using Bayesian
 86 statistics. To fully specify the model in a Bayesian framework, we first have
 87 to define a prior distribution for μ . Recall from Chapter 2 that for certain
 88 data models, certain priors lead to conjugacy i.e. if you choose the right prior
 89 for your parameter, your posterior distribution will be of a known parametric
 90 form. The conjugate prior for the mean of a normal model is also a Normal
 91 distribution:

$$\mu \sim \text{Normal}(\mu_0, \sigma_0^2)$$

92 If μ_0 and σ_0^2 are fixed, the posterior for μ has the following form (for the algebraic
 93 proof, see XXX):

$$\mu|y \sim \text{Normal}(\mu_n, \sigma_n^2) \quad (7.2)$$

94 where

$$\mu_n = \frac{\sigma^2}{\sigma^2 + n * \sigma_0^2} * \mu_0 + \frac{n * \sigma_0^2}{\sigma^2 + n * \sigma_0^2} * \bar{y}$$

95 And

$$\sigma_n^2 = \frac{\sigma^2 * \sigma_0^2}{\sigma^2 + n * \sigma_0^2}$$

96 We can directly obtain estimates of interest from this Normal posterior distri-
 97 bution, such as the mean $\hat{\mu}$ and its variance; we do not need to apply MCMC,
 98 since we can recognize the posterior as a parametric distribution, including the
 99 normalizing constant $p(y)$. But generally we will be interested in more complex
 100 models with several, say n , parameters. In this case, computing $p(y)$ from Eq.
 101 7.1 requires n -dimensional integration, which is can be difficult or impossible.
 102 Thus, the posterior distribution is generally only known up to a constant of
 103 proportionality:

$$p(\theta|y) \propto p(y|\theta) * p(\theta)$$

104 The power of MCMC is that it allows us to approximate the posterior using sim-
 105 ulation without evaluating the high dimensional integrals and to directly sample
 106 from the posterior, even when the posterior distribution is unknown! The price
 107 is that MCMC is computationally expensive. Although MCMC first appeared
 108 in the scientific literature in 1949 (Metropolis and Ulam, 1949), widespread use
 109 did not occur until the 1980s when computational power and speed increased
 110 (Gelfand and Smith, 1990). It is safe to say that the advent of practical MCMC
 111 methods is the primary reason why Bayesian inference has become so popular
 112 during the past three decades. In a nutshell, MCMC lets us generate sequential
 113 draws of θ (the parameter(s) of interest) from distributions approximating the
 114 unknown posterior over T iterations. The distribution of the draw at t depends

on the value drawn at $t-1$; hence, the draws from a Markov chain.¹ As T goes to infinity, the Markov chain converges to the desired distribution in our case the posterior distribution for $\theta|y$. Thus, once the Markov chain has reached its stationary distribution, the generated samples can be used to characterize the posterior distribution, $p(\theta|y)$, and point estimates of θ , its standard error and confidence bounds, can be obtained directly from this approximation of the posterior.

7.3 Types of MCMC sampling

There are several MCMC algorithms, the most popular being Gibbs sampling and Metropolis-Hastings sampling, both of which were briefly introduced in Chapt. 2. We will be dealing with these two classes in more detail and use them to construct the MCMC algorithms for SCR models. Also, we will briefly review alternative techniques that are applicable in some situations.

7.3.1 Gibbs sampling

Gibbs sampling was named after the physicist J.W. Gibbs by Geman and Geman (1984), who applied the algorithm to a Gibbs distribution². The roots of Gibbs sampling can be traced back to work of Metropolis et al. (1953), and it is actually closely related to Metropolis sampling (see Chapter 11.5 in Gelman et al. (2004), for the link between the two samplers). We will focus on the technical aspects of this algorithm, but if you find yourself hungry for more background, Casella and George (1992) provide a more in-depth introduction to the Gibbs sampler.

Let's go back to our simple example from above to understand the motivation and functioning of Gibbs sampling. Recall that for a Normal model with known variance and a Normal prior for μ , the posterior distribution of $\mu|y$ is also Normal. Conversely, with a fixed (known) μ , but unknown variance, the conjugate prior for σ^2 is an Inverse-Gamma distribution with shape and scale parameters a and b :

$$\sigma^2 \sim \text{InvGamma}(a, b),$$

With fixed a and b , the posterior $p(\sigma|\mu, y)$ is also an Inverse Gamma distribution, namely:

$$\sigma|\mu, y \sim \text{InvGamma}(a_n, b_n), \quad (7.3)$$

where $a_n = n/2 + a$ and $b_n = 1/2 \sum (y - \mu)^2 + b$. However, what if we know neither μ nor σ , which is probably the more common case? The joint posterior distribution of μ and σ now has the general structure

$$p(\mu, \sigma|y) = \frac{p(y|\mu) * p(\mu) * p(\sigma)}{\int p(y|\mu) * p(\mu) * p(\sigma) d\mu d\sigma}$$

¹In case you are not familiar with Markov chains, for T random samples $\theta^{(1)}, \dots, \theta^{(T)}$ from a Markov chain the distribution of $\theta^{(t)}$ depends only on the immediately preceding value, $\theta^{(t-1)}$.

²a distribution from physics we are not going to worry about, since it has no immediate connection with Gibbs sampling other than giving its name

147 Or

$$p(\mu, \sigma | y) \propto p(y | \mu) * p(\mu) * p(\sigma)$$

148 This cannot easily be reduced to a distribution we recognize. However, we can
 149 condition μ on σ (i.e., we treat σ as fixed) and remove all terms from the joint
 150 posterior distribution that do not involve μ to construct the full conditional
 151 distribution,

$$p(\mu | \sigma, y) \propto p(y | \mu) * p(\mu)$$

152 The full conditional of μ again takes the form of the Normal distribution
 153 shown in Eq. 7.2; similarly, $p(\sigma | \mu, y)$ takes the form of the Inverse Gamma
 154 distribution shown in Eq. 7.3 both distribution we can easily sample from.
 155 And this is precisely what we do when using Gibbs sampling we break down
 156 high-dimensional problems into convenient one-dimensional problems by con-
 157 structing the full conditional distributions for each model parameter separately;
 158 and we sample from these full conditionals, which, if we choose conjugate priors,
 159 are known parametric distributions. Let's put the concept of Gibbs sampling
 160 into the MCMC framework of generating successive samples, using our sim-
 161 ple Normal model with unknown μ and σ and conjugate priors as an example.
 162 These are the steps you need to build a Gibbs sampler:

163 **Step 0:** Begin with some initial values for θ , $\theta^{(0)}$. In our example, we have
 164 to specify initial values for μ and σ , for example by drawing a random number
 165 from some uniform distribution, or by setting them close to what we think they
 166 might be. (Note: This step is required in any MCMC sampling chains have to
 167 start from somewhere. We will get back to these technical details a little later.)

168 **Step 1:** Draw $\theta_1^{(1)}$ from the conditional distribution $p(\theta_1^{(1)} | \theta_2^{(0)}, \dots, \theta_d^{(0)})$. Here,
 169 θ_1 is μ , which we draw from the Normal distribution in Eq. 7.2 using $\sigma^{(0)}$ as
 170 value for σ .

171 **Step 2:** Draw $\theta_2^{(1)}$ from the conditional distribution $p(\theta_2^{(1)} | \theta_1^{(1)}, \theta_3^{(0)}, \dots, \theta_d^{(0)})$.
 172 Here, θ_2 is σ , which we draw from the Inverse Gamma distribution of Eq. 7.3,
 173 using $\mu^{(1)}$ as value for μ .

174 **Step d:** Draw $\theta_d^{(1)}$ from the conditional distribution $p(\theta_d^{(1)} | \theta_1^{(1)}, \dots, \theta_{d-1}^{(1)})$.

175 In our example we have no additional parameters, so we only need step 0
 176 through to 2. Repeat Steps 1 to d for T = a large number of samples. In terms
 177 of R coding, this means we have to write Gibbs updaters for μ and σ and embed
 178 them into a loop over T iterations. The final code in the form of an R function
 179 is shown in Panel 1.

180 Andy will build the panel environment here soon.

181

182 Panel 1: R-code for a Gibbs sampler for a Normal model with unknown mu
 183 and sig and conjugate (Normal and Inverse Gamma, respectively) priors
 184 for both parameters.

185

```

186 Normal.Gibbs<-function(y=y,mu0=mu0, sig0=sig0, a=a,b=b,niter=niter) {
187
188   ybar<-mean(y)
189   n<-length(y)
190   mu<-runif(1) #mean initial value
191   sig<-runif(1) #sd initial value
192   an<-n/2 + a
193
194   out<-matrix(nrow=niter, ncol=2)
195   colnames(out)<-c('mu', 'sig')
196
197   for (i in 1:niter) {
198
199     #update mu
200     mun<- (sig/(sig+n*sig0))*mu0 + (n*sig0/(sig+n* sig0))*ybar
201     sign <- (sig*sig0)/ (sig+n*sig0)
202     mu<-rnorm(1,mun, sqrt(sign))
203
204     #update sig
205     bn<- 0.5 * (sum((y-mu)^2)) +b
206     sig<-1/rgamma(1,shape=an, rate=bn)
207     out[i,]<-c(mu,sqrt(sig))
208
209   }
210   return(out)
211 }

```

212 This is it! You can use the code `NormalGibbs.R` in the **R** package `scrbook`
213 to simulate some data, $y \sim \text{Normal}(5, 0.5)$ and run your first Gibbs sampler.
214 Your output will be a table with two columns, one per parameter, and T rows,
215 one per iteration. For this 2-parameter example you can visualize the joint
216 posterior by plotting samples of μ against samples of σ (Fig. 7.1):

```

217 plot(out[,1], out[,2])

```


218 The marginal distribution of each parameter is approximated by just examining
219 the samples of this particular parameter you can visualize it by plotting a
220 histogram of the samples (Fig. 7.2 a and b):

```

221 par(mfrow=c(1,2))
222 hist(out[,1]); hist (out[,2])

```

223 Finally, recall an important characteristic of Markov chains, namely, that the
224 chain has to have converged (reached its stationary distribution) for samples to
225 come from the posterior distribution. In practice, that means you have to throw
226 out some of the initial samples called the burn-in. We will talk about this in
227 more when we talk about convergence diagnostics. For now, you can use the
228 `plot(out[,1])` or `plot(out[,2])` command to make a time series plot of the

Figure 7.1: Joint posterior distribution of μ and σ from a Normal Model

229 samples of each parameter and visually assess how many of the initial samples
 230 you should discard. Fig. 7.2 c and d shows plots for the estimates of μ and σ
 231 from our simulated data set; you see that in this simple example the Markov
 232 chain apparently reaches its stationary distribution very quickly the chains look
 233 'grassy' seemingly from the start. It is hard to discern a burn-in phase visually
 234 (but we will see examples further on where the burn-in is clearer) and you
 235 may just discard the first 500 draws to be sure you only use samples from the
 236 posterior distribution. The mean of the remaining samples are your estimates
 237 of μ and σ :

```

238 > summary(mod[501:10000,])
239           mu           sig
240 Min.      : 4.936      Min.      : 0.4569
241 1st Qu.: 4.984      1st Qu.: 0.4889
242 Median : 4.994      Median : 0.4961
243 Mean     : 4.994      Mean      : 0.4964
244 3rd Qu.: 5.005      3rd Qu.: 0.5037
245 Max.     : 5.062      Max.      : 0.5356
  
```

246 7.3.2 Metropolis-Hastings sampling

247 Although it is applicable to a wide range of problems, the limitations of Gibbs
 248 sampling are immediately obvious what if we do not want to use conjugate priors

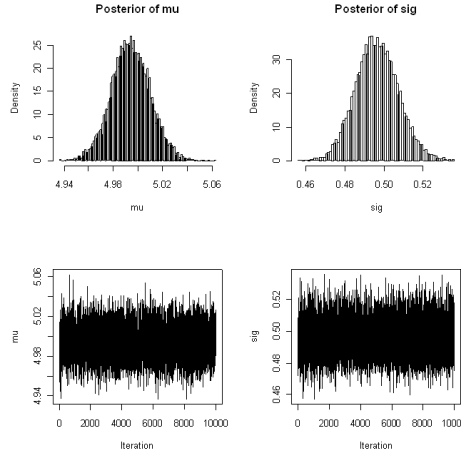


Figure 7.2: Plots of the posterior distributions of μ (a) and σ (b) from a Normal model and time series plots of μ (c) and σ (d).

249 (or what if we cannot recognize the full conditional distribution as a parametric
 250 distribution, or simply do not want to worry about these issues)? The most
 251 general solution is to use the Metropolis-Hastings (MH) algorithm, which also
 252 goes back to the work by Metropolis et al. (1953). You saw the basics of this
 253 algorithm in Chapter 2. In a nutshell, because we do not recognize the posterior
 254 $p(\theta|y)$ as a parametric distribution, the MH algorithm generates samples from
 255 a known proposal distribution, say $h(\theta)$, that depends on θ at $t-1$. The t^{th}
 256 sample is accepted with probability.

$$r = \frac{f(\theta^{(t-1)})h(\theta^{(t)}|\theta^{(t-1)})}{f(\theta^{(t)})h(\theta^{(t-1)}|\theta^{(t)})}$$

257 Proposal distributions can be absolutely anything! You can generate candi-
 258 date values from a *normal*(0, 1) distribution, from a *uniform*(-3455, 3455) distri-
 259 bution, or anything of proper support. Note, however, that good choices of $h()$
 260 are those that approximate the posterior distribution. Obviously if $h() = f(\theta|y)$
 261 (i.e., the posterior) then you always accept the draw, and it stands to reason
 262 that proposals that are more similar to $f(\theta|y)$ will lead to higher acceptance
 263 probabilities.

264 The original Metropolis algorithm required $h(\theta)$ to be symmetric so that
 265 $h(\theta^{(t)}|\theta^{(t-1)}) = h(\theta^{(t-1)}|\theta^{(t)})$. In that case these two terms just cancel out from
 266 the MH acceptance probability and r is then just the ratio of the target density
 267 evaluated at the candidate value to that evaluated at the current value. A later
 268 development of the algorithm by Hastings (1970) lifted this condition. Since
 269 using a symmetric proposal distribution makes life a little easier, we are going
 270 to focus on this specific case. A type of symmetric proposal useful in many

situations is the so-called *random-walk* proposal distribution where candidate values are drawn from a normal distribution with mean equal to the current value and some standard deviation, say δ , which is prescribed by the user (see below for further explanation).

Parameters with bounded support: Many models contain parameters that have bounded support. E.g., variance parameters live on $[0, \infty]$, parameters that represent probabilities live on $[0, 1]$, etc.. In that case it is sometimes convenient to use a random walk proposal distribution that can generate any real number (e.g., a normal random walk proposal). In that case, we can just reject parameters that are outside of the parameter space (XXXX REF FOR THIS XXXX).

It is worth knowing that there are alternatives to the random walk MH algorithm. For example, in the independent M-H, $\theta^{(t)}$ does not depend on $\theta^{(t-1)}$, while the Langevin algorithm (Roberts and Rosenthal, 1998) aims at avoiding the random walk by favoring moves towards regions of higher posterior probability density. The interested reader should look up these algorithms in Robert and Casella (2004) or Robert and Casella (2010).

Building a MH sampler can be broken down into several steps. We are going to demonstrate these steps using a different but still simple and common model the logit-normal or logistic regression model. For simplicity, assume that

$$y \sim \text{Bern}(\exp(\theta)/(1 + \exp(\theta)))$$

and

$$\theta \sim \text{Normal}(\mu, \sigma)$$

The following steps are required to set up a random walk MH algorithm:

Step 0: Choose initial values, $\theta^{(0)}$.

Step 1: Generate a proposed value of θ at t from $h(\theta^{(t)}|\theta^{(t-1)})$. We often use a Normal proposal distribution, so we draw $\theta^{(1)}$ from $\text{Normal}(\theta^{(0)}, \delta)$, where δ is the variance of the Normal proposal distribution, the tuning parameter that we have to set.

Step 2: Calculate the ratio of posterior densities for the proposed and the original value for θ :

$$r = \frac{p(\theta^{(t)}|y)}{p(\theta^{(t-1)}|y)}$$

In our example,

$$r = \frac{\text{Bernoulli}(y|\theta^{(t)}) * \text{Normal}(\theta^{(t)}|\mu, \sigma)}{\text{Bernoulli}(y|\theta^{(t-1)}) * \text{Normal}(\theta^{(t-1)}|\mu, \sigma)}$$

Step 3: Set

$$\begin{aligned} \theta^{(t)} &= \theta^{(t)} \text{ with probability } \min(r, 1) \\ &= \theta^{(t-1)} \text{ otherwise} \end{aligned}$$

302 We can do that by drawing a random number u from a $\text{Unif}(0, 1)$ and accept
 303 $\theta^{(t)}$ if $u < r$. Repeat for $t = 1, 2, \dots$ a large number of samples. The **R** code for
 304 this MH sampler is provided in Panel 2 XXXX.

305 Panel 2: R code to run a Metropolis sampler on a simple Logit-Normal model.

```
306
307 Logreg.MH<-function(y=y, mu0=mu0, sig0=sig0, niter=niter) {
308
309   out<-c()
310
311   theta<-runif(1, -3,3) #initial value
312
313   for (iter in 1:niter){
314     theta.cand<-rnorm(1, theta, 0.2)
315
316     loglike<-sum(dbinom(y, 1, exp(theta)/(1+exp(theta)), log=TRUE))
317     logprior <- dnorm(theta,mu0 ,sig0, log=TRUE)
318     loglike.cand<-sum(dbinom(y, 1, exp(theta.cand)/(1+exp(theta.cand)), log=TRUE))
319     logprior.cand <- dnorm(theta.cand, mu0, sig0, log=TRUE)
320
321     if (runif(1)<exp((loglike.cand+logprior.cand)-(loglike+logprior))){
322       theta<-theta.cand
323     }
324     out[iter]<-theta
325   }
326
327   return(out)
328 }
```

329 The reason we sum the logs of the likelihood and the prior, rather than
 330 multiplying the original values, is simply computational. The product of small
 331 probabilities can be numbers very close to 0, which computers do not handle
 332 well. Thus we add the logarithms, sum, and exponentiate to achieve the desired
 333 result. Similarly, in case you have forgotten some elementary math, $x/y =$
 334 $\exp(\log(x) - \log(y))$, with the latter being favored for computational reasons.

335 Comparing MH sampling to Gibbs sampling, where all draws from the con-
 336 ditional distribution are used, in the MH algorithm we discard a portion of the
 337 candidate values, which inherently makes in less efficient than Gibbs sampling
 338 the price you pay for its increased generality. In Step 1 of the MH sampler
 339 we had to choose a variance, δ , for the Normal proposal distribution. Choice
 340 of the parameters that define our candidate distribution is also referred to as
 341 'tuning', and it is important since adequate tuning will make your algorithm
 342 more efficient. δ should be chosen (a) large enough so that each step of drawing
 343 a new proposal value for θ can cover a reasonable distance in the parameter
 344 space, as otherwise, mixing of the Markov chain is inefficient and chains will
 345 tend to have strong autocorrelation; and (b) small enough so that proposal val-
 346 ues are not rejected too often, as otherwise the random walk will 'get stuck' at
 347 specific values for too long. As a rule of thumb, your candidate value should be
 348 accepted in about 40% of all cases. Acceptance rates of 20 – 80% are proba-

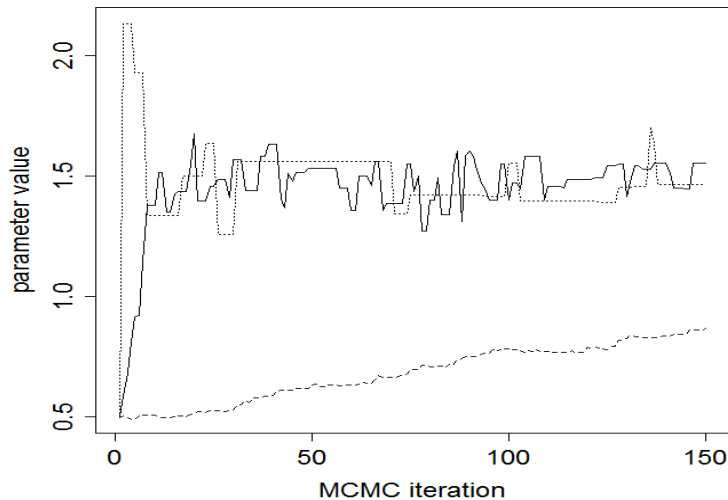


Figure 7.3: Time series plots of θ from a MH algorithm with tuning parameter $\delta = 0.01$ (dashed line), 0.2 (solid line) and 1 (dotted line).

bly ok, but anything below or above may well render your algorithm inefficient (this does not mean that it will give you wrong results only that you will need more iterations to converge to the posterior distribution). In practice, tuning will require some 'trial-and-error' and some common sense. Or, one can use an adaptive phase, where the tuning parameter is automatically adjusted until it reaches a user-defined acceptance rate, at which point the adaptive phase ends and the actual Markov chain begins. This is computationally a little more advanced. Link and Barker (2009) discuss this in more detail. It is important the samples drawn during the adaptive phase are discarded. To illustrate the effects of tuning, we ran the Metropolis-within-Gibbs algorithm in Panel 2 XX with $\delta = 0.01$, $\delta = 0.2$ and $\delta = 1$. The first 150 iterations for θ are shown in Fig. 7.3. We see that for a very small δ (the dashed line) the burn-in is extremely slow - after 150 iterations the chain isn't even half way there, while for the other two values of δ (solid and dotted) the burn-in phase seems to be over after only about 10 iterations. While $\delta = 0.2$ leads to reasonably good mixing, the chain clearly gets stuck on certain values with $\delta = 1$.

Other than graphically, you can easily check acceptance rates for the parameters you monitor (that are part of your output) using the `rejectionRate()` function of the package coda (we will talk more about this package a little later on). Do not let the term 'rejection rate' confuse you; it is simply $1 - \text{acceptance rate}$. There may be parameters for example, individual values of a random effect or latent variables that you do not want to save, though, and in our next

example we will show you a way to monitor their acceptance rates with a few extra lines of code.

7.3.3 Metropolis-within-Gibbs

One weakness of the MH sampler is that formulating the joint posterior when evaluating whether to accept or reject the candidate values for θ becomes increasingly complex or inefficient as the number of parameters in a model increases. As you already saw in Chapter 2, in these cases you can simply combine MH sampling and Gibbs sampling. You can use Gibbs sampling to break down your high-dimensional parameter space into easy-to-handle one-dimensional conditional distributions and use MH sampling for these conditional distributions. Better yet if you have some conjugacy in your model, you can use the more efficient Gibbs sampling for these parameters and one-dimensional MH for all the others. You have already seen the basics of how to build both types of algorithms, so we can jump straight into an example here and build a Metropolis-within-Gibbs algorithm.

7.4 GLMMs Poisson regression with a random effect

Let's assume a model that gets us closer to the problem we ultimately want to deal with a GLMM. Here, we assume we have Poisson counts, y , from i plots in j different study sites, and we believe that the counts are influenced by some plot-specific covariate, x , but that there is also a random site effect. So our model is:

$$y_{ij} \sim \text{Poisson}(\lambda_{ij})$$

$$\lambda_{ij} = \exp(a_j + bx_i)$$

Let's use Normal priors on a and b ,

$$a_j \sim \text{Normal}(\mu_a, \sigma_a)$$

and

$$b \sim \text{Normal}(\mu_b, \sigma_b)$$

.

Since we want to estimate the random effect in this model, we do not specify μ_a and σ_a , but instead, estimate them as well, so we have to specify hyperpriors for these parameters:

$$\begin{aligned} \mu_a &\sim \text{Normal}(\mu_0, \sigma_0) \\ \sigma_a &\sim \text{InvGamma}(a_0, b_0) \end{aligned}$$

With the model fully specified, we can compile the full conditionals, breaking the multi-dimensional parameter space into one-dimensional components:

$$\begin{aligned}
p(a_1|a_2, a_3, \dots, a_j, b, y) &\propto p(y_{i1}|a_1, b) * p(a_1) \\
&\propto \text{Poisson}(y_{i1}|\exp(a_1 + bx_i)) \\
&\quad * \text{Normal}(a_1|\mu_a, \sigma_a)
\end{aligned}$$

```

402 \begin{eqnarray*}
403 p(a_2|a_1, a_3, \ldots, a_j, b, y) &\propto p(y_{i2}|a_2, b) * p(a_2) \\
404 &\propto \text{Poisson}(y_{i2}|\exp(a_2 + bx_i)) * \text{Normal}(a_2|\mu_a, \sigma_a) \\
405 \end{eqnarray*}
406 and so on for all elements of a.
407 \begin{eqnarray*}
408 p(b|a, y) &\propto p(y|a, b) * p(b) \\
409 &\propto \text{Poisson}(y|\exp(a + bx)) * \text{Normal}(b|\mu_b, \sigma_b) \\
410 \end{eqnarray*}

```

411 Finally, we need to update the hyperparameters for a:

$$p(\mu_a|a) \propto p(a|\mu_a, \sigma_a) * p(\mu_a)$$

412

$$p(\sigma_a|a) \propto p(a|\mu_a, \sigma_a) * p(\sigma_a)$$

413 Since we assumed a to come from a Normal distribution, the choice of priors for
414 μ_a (Normal) and σ_a (Inverse Gamma) leads to the same conjugacy we observed
415 in our initial Normal model, so that both hyperparameters can be updated using
416 Gibbs sampling.

417 Now let's build the updating steps for these full conditionals. Again, for
418 the MH steps that update a and b we use Normal proposal distributions with
419 standard deviations δ_a and δ_b .

420 First, we set the initial values $a^{(0)}$ and $b^{(0)}$. Then, starting with a_1 , we draw
421 $a_1^{(1)}$ from $\text{Normal}(a_1^{(0)}, \delta_a)$, calculate the conditional posterior density of $a_1^{(0)}$
422 and $a_1^{(1)}$ and compare their ratios,

$$r = \frac{\text{Poisson}(y_{i1}|\exp(a_1^{(1)} + bx_i)) * \text{Normal}(a_1^{(1)}|\mu_a, \sigma_a)}{\text{Poisson}(y_{i1}|\exp(a_1^{(0)} + bx_i)) * \text{Normal}(a_1^{(0)}|\mu_a, \sigma_a)}$$

423 and accept $a_1^{(1)}$ with probability $\min(r, 1)$. We repeat this for all a's.

424 For b , we draw $b_1^{(1)}$ from $\text{Normal}(b^{(0)}, \delta_b)$, compare the posterior densities
425 of $b^{(0)}$ and $b^{(1)}$,

$$r = \frac{\text{Poisson}(y|\exp(a + b_1^{(1)}x)) * \text{Normal}(b_1^{(1)}|\mu_b, \sigma_b)}{\text{Poisson}(y|\exp(a + b_1^{(0)}x)) * \text{Normal}(b_1^{(0)}|\mu_b, \sigma_b)},$$

426 and accept $b_1^{(1)}$ with probability $\min(r, 1)$.

427 For μ_a and σ_a , we sample directly from the full conditional distributions
428 (Eq. ?? and Eq. ??):

$$\mu_a^{(1)} \sim \text{Normal}(\mu_n, \sigma_n)$$

where

$$\mu_n = \frac{\sigma_a^{(0)}}{\sigma_a^{(0)} + n_a * \sigma_0} * \mu_0 + \frac{n_a * \sigma_0}{\sigma_a^{(0)} + n_a * \sigma_0} * \bar{a}^{(1)}$$

and

$$\sigma_n = \frac{\sigma_a^{(0)} * \sigma_0}{\sigma_a^{(0)} + n * \sigma_0}$$

Here, \bar{a} is the current mean of the vector \mathbf{a} , which we updated before, and n_a is the length of \mathbf{a} . For σ_a we use $\sigma_a^{(1)} \sim \text{InvGamma}(a_n, b_n)$, where $a_n = n_a/2 + a_0$, and $b_n = 0.5(\sum_{j=1}^{n_a} a_j^{(1)} - \mu_a^{(1)})^2 + b_0$.

We repeat these steps over T iterations of the MCMC algorithm. In this example we may not want to save each individual a , but are only interested in their mean and standard deviation. Since these two parameters will change as soon as the value for one element in \mathbf{a} changes, their acceptance rates will always be close to 1 and are not representative of how well your algorithm performs. To monitor the acceptance rates of parameters you do not want to save, you simply need to add a few lines of code into your updater to see how often the individual parameters are accepted. The full code for the MCMC algorithm of our Poisson GLMM in Panel 3 (XXX) shows one way how to monitor acceptance of individual a 's.

```

444 Panel 3: R code for the Metropolis-within-Gibbs sampler for
445 a Poisson regression with random intercepts.
446
447 Pois.reg<-function(y=y,site=site,mu0=mu0,sig0=sig0,a0=a0,b0=b0,
448                   mub=mub, sigb=sigb, niter=niter){
449
450   lev<-length(unique(site))      #number of sites
451   a<-runif(lev,-5,5) #initial values a
452   b<-runif(1,0,5) #initial value b
453   mua<-mean(a)
454   siga<-sd(a)
455
456   out<-matrix(nrow=niter, ncol=3)
457   colnames(out)<-c('mua','siga','b')
458
459   for (iter in 1:niter) {
460
461     #update a
462     aUps<-0 #initiate counter for acceptance rate of a
463     for (j in 1:lev) { #loop over sites
464       a.cand<-rnorm(1, a[j], 0.1) #update intercepts a one at a time
465       loglike<- sum(dpois (y[site==j], exp(a[j] + b*x[site==j]), log=TRUE))
466       logprior<- dnorm(a[j], mua,siga, log=TRUE)
467       loglike.cand<- sum(dpois (y[site==j], exp(a.cand + b *x[site==j]), log=TRUE))

```

```

468 logprior.cand<- dnorm(a.cand, mua,siga, log=TRUE)
469 if (runif(1)< exp((loglike.cand+logprior.cand) (loglike+logprior))) {
470   a[j]<-a.cand
471   aUps<-aUps+1
472 }
473 }
474
475 if(iter %% 100 == 0) { #this lets you check the acceptance rate of a at every 100th iteration
476   cat("   Acceptance rates\n")
477   cat("     a =", aUps/lev, "\n")
478 }
479
480 #update b
481 b.cand<-rnorm(1, b, 0.1)
482 avec<-rep(a, times=c(rep(10,10)))
483 loglike<- sum(dpois (y, exp(avec + b*x), log=TRUE))
484 logprior<- dnorm(b, mub,sigb, log=TRUE)
485 loglike.cand<- sum(dpois (y, exp(avec + b.cand *x), log=TRUE))
486 logprior.cand<- dunif(b.cand, mub,sigb, log=TRUE)
487 if (runif(1)< exp((loglike.cand+logprior.cand) (loglike+logprior) )) {
488   b<-b.cand
489 }
490
491 #update mua using Gibbs sampling
492 abar<-mean(a)
493 mun<- (siga/(siga+lev*sig0))*mu0 + (lev*sig0/(siga+lev* sig0))*abar
494 sign <- (siga*sig0)/ (siga+lev*sig0)
495 mua<-rnorm(1,mun, sqrt(sign))
496
497 #update siga using Gibbs sampling
498 a0n<-lev/2 + a0
499 b0n<- 0.5 * (sum((a-mua)^2)) +b0
500 siga<-1/rgamma(1,shape=a0n, rate=b0n)
501
502 out[iter,]<-c(mua, sqrt(siga), b)
503
504 }
505
506 return(out)
507 }

```

7.4.1 Rejection sampling and slice sampling

While MH and Gibbs sampling are probably the most widely applied algorithms for posterior approximation, there are other options that work under certain circumstances and may be more efficient when applicable. WinBUGS applies these algorithms and we want you to be aware that there is more out there to approximate posterior distributions than Gibbs and MH. One alternative algorithm is rejection sampling. Rejection sampling is not an MCMC method,




Figure 7.4: Slice sampling. For...

since each draw is independent of the others. The method can be used when the
 posterior $p(\theta|y)$ is not a known parametric distribution but can be expressed in
 closed form. Then, we can use a so-called envelope function, say, $g(\theta)$, that we
 can easily sample from, with the restriction that $p(\theta|y) < M * g(\theta)$. We then
 sample a candidate value for θ from $g(\theta)$, calculate $r = p(\theta|y)/M * g(\theta)$ and keep
 the sample with the probability r . M is a constant that has to be picked so
 that r lies between 0 and 1, for example by evaluating both $p(\theta|y)$ and $g(\theta)$ at n
 points and looking at their ratios. Rejection sampling only works well if $g(\theta)$ is
 similar to $p(\theta|y)$, and packages like WinBUGS use adaptive rejection sampling
 (Gilks and Wild, 1992), where a complex algorithm is used to fit an adequate and
 efficient $g(\theta)$ based on the first few draws. Though efficient in some situations,
 rejection sampling does not work well with high-dimensional problems, since
 it becomes increasingly hard to define a reasonable envelope function. For an
 example of rejection sampling in the context of SCR models, see Chapter 9
 XXXX. Another alternative is slice sampling (Neal, 2003). In slice sampling,
 we sample uniformly from the area under the plot of $p(\theta|y)$. Considering a single
 univariate theta. Let's define an auxiliary variable, $U \sim Uniform(0, p(\theta|y))$.
 Then, θ can be sampled from the vertical slice of $p(\theta|y)$ at U (Figure 4):

$$\theta|U \sim \text{Unif}(B),$$

where $B = \{\theta : p(\theta|y) \geq U\}$
³

Slice sampling can be applied in many situations; however, implementing
 an efficient slice sampling procedure can be complicated. We refer the inter-
 ested reader to chapter 7 of Robert and Casella (2010) for a simple example.
 Both rejection sampling and slice sampling can be applied on one-dimensional
 conditional distributions within a Gibbs sampling setup.

³there are supposed to be equations in the caption of figure 4 but it kept causing errors

7.5 MCMC for closed capture-recapture Model Mh

7.5.1 Building your own MCMC algorithm

By now you have seen MCMC samplers for some simple GL(M)M's. Now, to ease you into more complex models, we construct our own MCMC algorithm using a Metropolis-within-Gibbs sampler for the non-spatial Model with individual heterogeneity in capture probability M_h , developed in Chapt. 3.

To recapitulate: Under the non-spatial model, each of the n observed individuals is either detected (1) or not (0) during each of K sampling occasions. We estimate N using data augmentation and have a Bernoulli model for the zero-inflation variables z_i . The binomial observation model is expressed conditional on the latent variables z_i . Further, we prescribe a distribution for the capture probability p_i . Here we assume

$$\text{logit}(p_i) \sim \text{Normal}(\mu, \sigma^2)$$

As usual, we will have to go through two general steps before we write the MCMC algorithm:

- (1) Identify your model with all its components (including priors)
- (2) Recognize and express the full conditional distributions for all parameters

Our model components are as follows: $[y_i|p_i, z_i]$, $[p_i|\mu_p, \sigma_p]$, and $[z_i|\psi]$ for each $i = 1, 2, \dots, M$ and then prior distributions $[\mu_p]$, $[\sigma_p]$ and $[\psi]$. The joint posterior distribution of all unknown quantities in the model is proportional to the joint distribution of all elements y_i, p_i, z_i and also the prior distributions of the prior parameters:

$$\left\{ \prod_{i=1}^M [y_i|p_i, z_i][p_i|\mu_p, \sigma_p][z_i|\psi] \right\} [\mu_p, \sigma_p, \psi]$$

For prior distributions, we assume that μ_p, σ_p, ψ are mutually independent and for μ_p and σ_p we use improper uniform priors, and $\psi \sim \text{Unif}(0, 1)$. Note that the likelihood contribution for each individual, when conditioned on p_i and z_i , does not depend on ψ , μ_p , or σ_p . As such, the full-conditionals for the structural parameters ψ only depends on the collection of data augmentation variables z_i , and that for μ_p and σ_p will only depends on the collection of latent variables $p_i; i = 1, 2, \dots, M$. The full conditionals for all the unknowns are as follows:

- (1) For p_i :

$$[p_i|y_i, \mu_p, \sigma_p, z_i = 1] \propto [y_i|p_i][p_i|\mu_p, \sigma_p^2] \text{ if } z_i = 1$$

$$[p_i|\mu_p, \sigma_p] \text{ if } z_i = 0$$

- (2) for z_i :

$$z_i|\cdot \propto [y_i|z_i * p_i]\text{Bern}(z_i|\psi)$$

571 **(3)** For μ_p :

$$[\mu_p|\cdot] \sim \prod_i [p_i|\cdot] * \text{const}$$

572 **(4)** For σ_p :

$$[\sigma_p|\cdot] \sim \prod_i [p_i|\cdot] * \text{const}$$

573 **(5)** For ψ :

$$\psi|\cdot \sim \text{Beta}(1 + \sum z_i, 1 + M - \sum z_i)$$

574 What we've done here is identify each of the full conditional distributions
 575 in sufficient detail to toss them into our Metropolis-Hastings algorithm. With
 576 the exception of ψ which has a convenient analytic solution – it is a beta dis-
 577 tribution which we can easily sample directly. In truth, we could also sample
 578 μ_p and σ_p^2 directly with certain choices of prior distributions. For example, if
 579 $\mu_p \sim \text{Normal}(0, 1000)$ then the full conditional for μ_p is also normal, etc.. We
 580 implement an MCMC algorithm for this model in the following block of **R** code.

```

581
582   ## obtain the bear data by executing the previous data grabbing
583   ## function
584
585   temp<-getdata()
586   M<-temp$M
587   K<-temp$K
588   ytot<-temp$ytot
589
590
591   ###
592   ### MCMC algorithm for Model Mh
593
594   out<-matrix(NA,nrow=100000,ncol=4)
595   dimnames(out)<-list(NULL,c("mu","sigma","psi","N"))
596   lp<- rnorm(M,-1,1)
597   p<-expit(lp)
598   mu<- -1
599   p0<-exp(mu)/(1+exp(mu))
600   sigma<- 1
601   psi<- .5
602   z<-rbinom(M,1,psi)
603   z[ytot>0]<-1
604
605   for(i in 1:100000){
606
607     ### update the logit(p) parameters
608     lp.cand<- rnorm(M,lp,1) # 0.5 is a tuning parameter

```

```

609 p.cand<-expit(lp.cand)
610 lik.curr<-log(dbinom(ytot,K,z*p)*dnorm(lp,mu,sigma))
611 lik.cand<-log(dbinom(ytot,K,z*pc)*dnorm(lpc,mu,sigma))
612 kp<- runif(M) < exp(lik.cand-lik.curr)
613 p[kp]<-pc[kp]
614 lp[kp]<-lpc[kp]
615
616 p0c<- rnorm(1,p0,.05)
617 if(p0c>0 & p0c<1){
618   muc<-log(p0c/(1-p0c))
619   lik.curr<-sum(dnorm(lp,mu,sigma,log=TRUE))
620   lik.cand<-sum(dnorm(lp,muc,sigma,log=TRUE))
621   if(runif(1)<exp(lik.cand-lik.curr)) {
622     mu<-muc
623     p0<-p0c
624   }
625 }
626
627 sigmac<-rnorm(1,sigma,.5)
628 if(sigmac>0){
629   lik.curr<-sum(dnorm(lp,mu,sigma,log=TRUE))
630   lik.cand<-sum(dnorm(lp,mu,sigmac,log=TRUE))
631   if(runif(1)<exp(lik.cand-lik.curr))
632     sigma<-sigmac
633 }
634
635 ### update the z[i] variables
636 zc<- ifelse(z==1,0,1) # candidate is 0 if current = 1, etc..
637 lik.curr<- dbinom(ytot,K,z*p)*dbinom(z,1,psi)
638 lik.cand<- dbinom(ytot,K,zc*p)*dbinom(zc,1,psi)
639 kp<- runif(M) < (lik.cand/lik.curr)
640 z[kp]<- zc[kp]
641
642 psi<-rbeta(1, sum(z) + 1, M-sum(z) + 1)
643
644 out[i,<- c(mu,sigma,psi,sum(z))
645 }

```

Remarks: (1) for parameters with bounded support, i.e., σ_p and p_0 , we are using a random walk candidate generator but rejecting draws outside of the parameter space. (2) We mostly use Metropolis-Hastings except for the data augmentation parameter ψ which we sample directly from its full-conditional distribution which is a beta distribution. (3) Even the latent data augmentation variables z_i are updated using Metropolis-Hastings although they too can be updated directly from their full-conditional.

7.6 MCMC algorithm for the basic spatial capture-recapture model

Conceptually, but also in terms of MCMC coding, it is only a small step from the non-spatial model Mh to a fully spatial capture-recapture model. Next, we'll walk you through the steps of building your own MCMC sampler for the basic SCR model (i.e. without any individual, site or time specific covariates) with both a Poisson and a binomial encounter process. As usual, we will have to go through two general steps before we write the MCMC algorithm:

- (1) Identify your model with all its components (including priors)
- (2) Recognize and express the full conditional distributions for all parameters

It is worthwhile to go through all of step 1 for an SCR model, but you have probably seen enough of step 2 in our previous examples to get the essence of how to express a full conditional distribution. Therefore, we will exemplify step 2 for some parameters and tie these examples directly to the respective R code.

Step 1 – Identify your model

Recall the components of the basic SCR model with a Poisson encounter process from Chapt. ?? : We assume that individuals i , or rather, their activity centers s_i , are uniformly distributed across our state space S ,

$$s_i \sim U(S)$$

and that the number of times individual i encounters trap j , y_{ij} , is a random Poisson variable with mean λ_{ij} ,

$$y_{ij} \sim \text{Poisson}(\lambda_{ij})$$

The tie between individual location, movement and trap encounter rates is made by the assumption that λ_{ij} , is a decreasing function of the distance between s_i and j , D_{ij} , of the half-normal form

$$\lambda_{ij} = \lambda_0 * \exp(-D_{ij}^2/2\sigma^2),$$

where λ_0 is the baseline trap encounter rate at $D_{ij} = 0$ and σ controls the shape of the half-normal function.

In order to estimate the number of s_i in S , N , we use data augmentation (Sect. ??) and create $M - n$ all-0 encounter histories, where n is the number of individuals we observed and M is an arbitrary number that is larger than N . We estimate N by summing over the auxiliary data augmentation variables, z_i , which is 1 if the individual is part of the population and 0 if not, and assume that z_i is a random Bernoulli variable,

$$z_i \sim \text{Bern}(\psi)$$

To link the two model components, we modify our trap encounter model to

$$\lambda_{ij} = \lambda_0 * \exp(-D_{ij}^2/2\sigma^2) * z_i.$$

The model has the following structural parameters, for which we need to specify priors:

ψ : the *Uniform*(0, 1) is required as part of the data augmentation procedure and in general is a natural choice of an uninformative prior for a probability; note that this is equivalent to a *Beta*(1, 1) prior, which will come in handy later.

s_i : since s_i is a pair of coordinates it is two-dimensional and we use a uniform prior limited by the extent of our state-space over both dimensions.

σ : we can conceive several priors for σ but let's assume an improper prior, one that is Uniform over $(-\text{Inf}, \text{Inf})$. We will see why this is convenient when we construct the full conditionals for σ .

λ_0 : analogous, we will use a *Uniform* $(-\text{Inf}, \text{Inf})$ improper prior for λ_0 .

The parameter that is the objective of our modeling, N , is a derived parameter that we can simply obtain by summing all z_i :

$$N = \sum(z)$$

Step 2 – Construct the full conditionals Having completed step 1, let's look at the full conditional distributions for some of these parameters. We find that with improper priors, full conditionals are proportional only to the likelihood of the observations; for example, take the movement parameter σ :

$$[\sigma|s, \lambda_0, z, y] \propto [y|s, \lambda_0, z, \sigma] * [\sigma]$$

Since the improper prior implies that $[\sigma] \propto 1$, we can reduce this further to

$$[\sigma|s, \lambda_0, z, y] \propto [y|s, \lambda_0, z, \sigma]$$

The R code to update σ is shown in Panel 4. Notice that we automatically reject negative candidate values, since σ cannot be < 0 .

Panel 4: R code to update sigma within an MCMC algorithm for an SCR model when using an improper prior

```

sig.cand <- rnorm(1, sigma, 0.1) #draw candidate value
if(sig.cand>0){ #automatically reject sig.cand that are <0
  lam.cand <- lam0*exp(-(D*D)/(2*sig.cand*sig.cand))
  ll<- sum(dpois(y, lam*z, log=TRUE))
  llcand <- sum(dpois(y, lam.cand*z, log=TRUE))
  if(runif(1) < exp( llcand - ll) ){
    ll<-llcand
    lam<-lam.cand
    sigma<-sig.cand
  }
}
```

These steps are analogous for λ_0 and s_i and we will use MH steps for all of these parameters. Similar to the random intercepts in our Poisson GLMM, we

update each s_i individually. Note that to be fully correct, the full conditional for s_i contains both the likelihood and prior component, since we did not specify an improper, but a Uniform prior on s_i . However, with a Uniform distribution the probability density of any value is $1/(\text{upper limit} - \text{lower limit}) = \text{constant}$. Thus, the prior components are identical for both the current and the candidate value and can be ignored (formally, when you calculate the ratio of posterior densities, r , the identical prior component appears both in the numerator and denominator, so that they cancel each other out).

We still have to update z_i . The full conditional for z_i is

$$[z_i|y, \sigma, \lambda_0, s] \propto [y|z, \sigma, \lambda_0, s] * [z_i]$$

and since $z_i \sim \text{Bernoulli}(\psi)$, the term has to be taken into account when updating z_i . The R code for updating z_i is shown in Panel 5.

Panel 5: R code to update z

```

zUps <- 0 #set counter to monitor acceptance rate
for(i in 1:M) {
  if(seen[i]) #no need to update seen individuals, since their z =1
    next
  zcand <- ifelse(z[i]==0, 1, 0)
  llz <- sum(dpois(y[i,], lam[i,]*z[i], log=TRUE))
  llcand <- sum(dpois(y[i,], lam[i,]*zcand, log=TRUE))

  prior <- dbinom(z[i], 1, psi, log=TRUE)
  prior.cand <- dbinom(zcand, 1, psi, log=TRUE)
  if(runif(1) < exp( (llcand+prior.cand) - (llz+prior) )) {
    z[i] <- zcand
    zUps <- zUps+1
  }
}

```

ψ itself is a hyperparameter of the model, with an uninformative prior distribution of $\text{Unif}(0, 1)$ or $\text{Beta}(1, 1)$, so that

$$\psi|z \propto [z|\psi] * \text{Beta}(1, 1)$$

The Beta distribution is the conjugate prior to the Binomial and Bernoulli distributions (remember that $z \sim \text{Bernoulli}(\psi)$). The general form of a full conditional of a Beta-Binomial model with $y_i \sim \text{Bernoulli}(p)$ and $p \sim \text{Beta}(a, b)$ is

$$p(p|y) \propto \text{Beta}(a + \sum y_i, b + n - \sum y_i)$$

In our case, this means we update psi as follows:

```
si<-rbeta(1, 1+sum(z), 1 + M-sum(z))
```

These are all the building blocks you need to write the MCMC algorithm for the spatial null model with a Poisson encounter process. You can find the full R code (SCR0pois.R) in the R package scrbook XXXXXX.

7.6.1 SCR model with binomial encounter process

The equivalent SCR model with a binomial encounter process is very similar. Here, each individual i can only be detected once at any given trap j during a sampling occasion k . Thus

$$y_{ij} \sim \text{Binomial}(p_{ij}, K)$$

Where p_{ij} is some function of distance between \mathbf{s}_i and trap location \mathbf{x}_j . Here we use:

$$p_{ij} = 1 - \exp(-\lambda_{ij})$$

Recall from Chapter 2 that this is the complementary log-log (cloglog) link function, which constrains p_{ij} to fall between 0 and 1. For our MCMC algorithm that means that, instead of using a Poisson likelihood, $\text{Poisson}(y|\sigma, \lambda_0, s, z)$, we use a Binomial likelihood, $\text{Binomial}(y, K|\sigma, \lambda_0, s, z)$, in all the conditional distributions. As an example, Panel 6 shows the updating step for λ_0 under a binomial encounter model. The full MCMC code for the binomial SCR (SCR0binom.R) can be found in the R package scrbook XXXXXX.

Panel 6: MCMC updater for lam0 in a SCR model with Binomial encounter process and cloglog link function on detection. Here, pmat = 1-exp(-lam).

```

lam0.cand <- rnorm(1, lam0, 0.1)
if(lam0.cand > 0){ #automatically reject lam0.cand that are <0
  lam.cand <- lam0.cand*exp(-(D*D)/(2*sigma*sigma))
  p.cand <- 1-exp(-lam.cand)
  ll<- sum(dbinom(y, K, pmat *z, log=TRUE))
  llcand <- sum(dbinom(y, K, p.cand *z, log=TRUE))
  if(runif(1) < exp( llcand - ll) ){
    ll<-llcand
    pmat<-p.cand
    lam0<- lam0.cand
  }
}
```

Another possibility is to model variation in the individual and site specific detection probability, p_{ij} , directly, without any transformation, such that

$$p_{ij} < -p_0 * \exp(-D_{ij}^2/(2\sigma^2))$$

and $p_0 = \{0, 1\}$. This formulation is analogous to how detection probability is modeled in distance sampling under a half-normal detection function; however, in distance sampling p_0 – detection of an individual on the transect line – is assumed to be 1 (Buckland, 2001). Under this formulation the updater for λ_0 (equivalent to p_0 in Eq XX) becomes:

```

lam0.cand <- rnorm(1, lam0, 0.1)
if(lam0.cand > 0 & lam0.cand < 1 ){ #automatically reject lam0.cand that are not {0,1}
```

```

800     lam.cand <- lam0.cand*exp(-(D*D)/(2*sigma*sigma))
801     ll<- sum(dbinom(y, K, lam *z, log=TRUE)) #no transformation needed
802     llcand <- sum(dbinom(y, K, lam.cand *z, log=TRUE))
803     if(runif(1) < exp( llcand - ll) ){
804         ll<-llcand
805         lam<-lam.cand
806         lam0<- lam0.cand
807     }
808 }

```

809 7.6.2 Looking at model output

810 Now that you have an MCMC algorithm to analyze spatial capture-recapture
811 data with, let's run an actual analysis so we can look at the output. As an
812 example, we will use the Fort Drum bear data set we already analyzed in Chapt.
813 3 with traditional non-spatial models (and that you will see again in Chapt.
814 ??). You can use the same script provided back in Chapt. 3 to read in the
815 trap location (trapmat) and detection data and build the augmented $M \times K$
816 array of individual encounter histories (Xaug). In addition to these data, we
817 need to specify the outermost coordinates of the state-space. Since bears are
818 wide ranging animals we add a 20-km buffer to the maximum and minimum
819 coordinates of the trap array:

```

820 x1<- min(trapmat[,2])- 20
821 y1<- min(trapmat[,3])- 20
822 xu<-max(trapmat[,2])+ 20
823 yu<-max(trapmat[,3])+ 20

```

824 Finally, source the MCMC code for the binomial encounter model algorithm
825 with the cloglog link and run 5000 iterations. This should take approximately
826 25 minutes.

```

827 > source('SCR0binom.txt')
828 > mod0<-SCR.0(y=Xaug, X=trapmat[,2:3], M=M, x1=x1, xu=xu, y1=y1, yu=yu, K=8, niter=5000)

```

829 Before, we used simple R commands to look at model results. However, there
830 is a specific R package to summarize MCMC simulation output and perform
831 some convergence diagnostics – package coda (Plummer et al., 2006). Download
832 and install coda, then convert your model output to an mcmc object

```

833 > chain<-mcmc(mod0)

```

834 which can be used by coda to produce MCMC specific output.

835 Markov chain time series plots

836 Start by looking at time series plots of your Markov chains using `plot(chain)`.
837 This command produces a time series plot and marginal posterior density plots

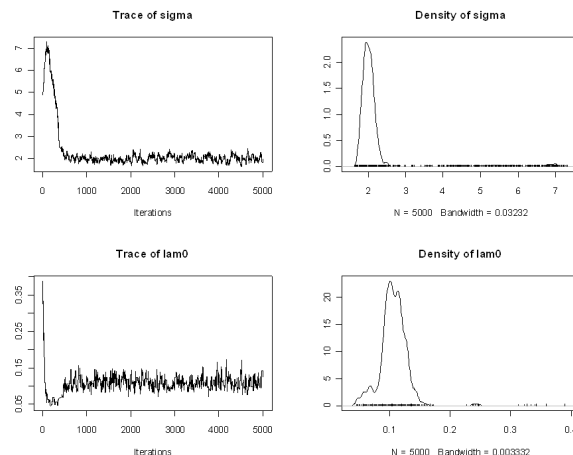


Figure 7.5: Time series and posterior density plots for σ and λ_0 .

for each monitored parameter, similar to what we did before using the `hist()` and `plot()` commands (Fig. 5). Time series plots will tell you several things: First, recall from Sect. 7.1 that the way the chains move through the parameter space gives you an idea of whether your MH steps are well tuned. If chains were constant over many iterations you would need to decrease the tuning parameter of the (Normal) proposal distribution. If a chain moves along some gradient to a stationary state very slowly, you may want to increase the tuning parameter so that the parameter space is explored more efficiently.

Second, you will be able to see if your chains converged and how many initial simulations you have to discard as burn-in. In the case of the chains shown in Fig. 7.5, we would probably consider the first 750 - 1000 iterations as burn-in, as afterwards the chains seem to be fairly stationary.

7.6.3 Posterior density plots

The `plot()` command also produces posterior density plots and it is worthwhile to look at those carefully. For parameters with priors that have bounds (e.g. Uniform over some interval), you will be able to see if your choice of the prior is truncating the posterior distribution. In the context of SCR models, this will mostly involve our choice of M , the size of the augmented data set. If the posterior of N has a lot of mass concentrated close to M (or equivalently the posterior of ψ has a lot of mass concentrated close to 1), as in the example in Fig. 7.6, we have to re-run the analysis with a larger M . A flat posterior plot shows you that the parameter essentially cannot be identified. There may not be enough information in your data to estimate model parameters and you may have to consider a simpler model. Finally, posterior density plots will show you




Figure 7.6: Time series and posterior density plots of ψ and N for the bear data set truncated by the upper limit of M (500).

862 if the posterior distribution is symmetrical or skewed – if the distribution has
 863 a heavy tail, using the mean as a point estimate of your parameter of interest
 864 may be biased and you may want to opt for the median or mode instead.

865 7.6.4 Serial autocorrelation and effective sample size

866 Checking the degree of autocorrelation in your Markov chains and estimating
 867 the effective sample size your chain has generated should be part of evaluating
 868 your model output. If you use WinBUGS through the R2WinBUGS package,
 869 the `print()` command will automatically return the effective sample size for
 870 all monitored parameters. In the coda package there are several functions you
 871 can use to do so. `effectiveSize()` will directly give you an estimate of the
 872 effective sample size for you parameters:

```
873 > effectiveSize(chain)
874      sigma      lam0      psi      N
875 3.930303 78.259159 30.436348 32.047392
```

876 Alternatively, you can use the `autocorr.diag()` function, which will show
 877 you the degree of autocorrelation for different lag values (which you can specify
 878 within the function call, we use the defaults below):

```
879 > autocorr.diag(mcmc(mod))
880      sigma      lam0      psi      N
881 Lag 0 1.0000000 1.0000000 1.0000000 1.0000000
882 Lag 1 0.9979948 0.9494134 0.9847503 0.9774201
```

```

883 Lag 5 0.9915567 0.8038168 0.9111951 0.9113525
884 Lag 10 0.9836016 0.6714021 0.8462108 0.8509803
885 Lag 50 0.8985337 0.1983780 0.6138516 0.6233994

```

886 In the present case we see that autocorrelation is especially high for the param-
887 eter σ and our effective sample size for this parameter is 4! ⁴ This means we
888 would have to run the model for much longer to obtain a reasonable effective
889 sample size. Unfortunately, with many SCR models we observe high degrees of
890 serial autocorrelation. For now, let's continue using this small set of samples to
891 continue looking at the output.

892 7.6.5 Summary results

893 Now that we checked that our chains apparently have converged and pretending
894 that we have generated enough samples from the posterior distribution, we can
895 look at the actual parameter estimates. The `summary()` function will return two
896 sets of results: the mean parameter estimates, with their standard deviation,
897 the naive standard error – i.e. your regular standard error calculated for T (=
898 number of iterations) samples without accounting for serial autocorrelation –
899 and the Time-series SE (in WinBUGS and earlier in this book referred to as
900 MC error), which accounts for autocorrelation. Remember our rule of thumb
901 that this error decreases with increasing chain length and should be 1% or less
902 of the parameter estimate. In WinBUGS the MC error is only given in the log
903 output within BUGS itself. You should adjust the `summary()` call by removing
904 the burn-in from calculating parameter summary statistics. To do so, use the
905 `window()` command, which lets you specify at which iteration to start 'counting'.
906 In contrast to WinBUGS, which requires you to set the burn-in length before
907 you run the model, this command gives us full flexibility to make decisions about
908 the burn-in after we have seen the trajectories of our Markov chains. For our
909 example, `summary(window(chain, start=1001))` returns the following output:

```

910 Iterations = 1001:5000
911 Thinning interval = 1
912 Number of chains = 1
913 Sample size per chain = 4000
914
915 1. Empirical mean and standard deviation for each variable,
916    plus standard error of the mean:
917
918      Mean      SD Naive SE Time-series SE
919 sigma  1.9986 0.13805 0.0021827      0.016091
920 lam0   0.1096 0.01523 0.0002407      0.001401
921 psi    0.6113 0.09148 0.0014465      0.010734
922 N      489.8535 71.79695 1.1352094      8.431119

```

⁴Anyone have any idea how the autocorrelation in sigma could be reduced?

924 2. Quantiles for each variable:

```
925
926           2.5%      25%      50%      75%      97.5%
927 sigma  1.75780    1.89847    1.9900    2.0944    2.2772
928 lam0    0.08357    0.09824    0.1087    0.1192    0.1427
929 psi     0.45110    0.54838    0.6052    0.6639    0.8192
930 N       366.00000  440.00000  485.0000  530.0000  654.0000
```

931 Looking at the MC errors, we see that in spite of the high autocorrelation, the
 932 MC error for σ is below the 1% threshold, whereas for all other parameters, MC
 933 errors are still above, another indication that for a thorough analysis we should
 934 run a longer chain. Our algorithm gives us a posterior distribution of N , but we
 935 are usually interested in the density, D . Density itself is not a parameter of our
 936 model, but we can derive a posterior distribution for D by dividing each value
 937 of N (N at each iteration) by the area of the state-space (here 3032.719 km^2)
 938 and we can use summary statistics of the resulting distribution to characterize
 939 D :

```
940 > summary(window(chain[,4]/ 3032.719, start=1001))
941 Iterations = 1001:5000
942 Thinning interval = 1
943 Number of chains = 1
944 Sample size per chain = 4000
```

945
 946 1. Empirical mean and standard deviation for each variable,
 947 plus standard error of the mean:

```
948
949           Mean           SD      Naive SE Time-series SE
950 0.1615229    0.0236741    0.0003743    0.0027801
```

951
 952 2. Quantiles for each variable:

```
953
954      2.5%   25%   50%   75%  97.5%
955 0.1207 0.1451 0.1599 0.1748 0.2156
```

956 We see that our mean density of $0.16/\text{km}^2$ is very similar to the estimate of
 957 $0.18/\text{km}^2$ obtained under the non-spatial model M0 in Chapt. 3.

958 7.6.6 Other useful commands

959 While inspecting the time series plot gives you a first idea of how well you
 960 tuned your MH algorithm, use `rejectionRate()` to obtain the rejection rates
 961 ($1 - \text{acceptance rates}$) of the parameters that are written to your output:

```
962 > rejectionRate(chain)
963      sigma      lam0      psi      N
964 0.44108822 0.77675535 0.00000000 0.01940388
```


Recall that rejection rates should lie between 0.2 and 0.8, so our tuning seems to have been appropriate here. ψ is never rejected since we update it with Gibbs sampling, where all candidate values are kept. And since N is the sum of all z_i , all it takes for N to change from one iteration to the next are small changes in the z -vector, so the rejection rate of N is always low. If you have run several parallel chains, you can combine them into a single mcmc object using the `mcmc.list()` command on the individual chains (note that each chain has to be converted to an mcmc object before combining them with `mcmc.list()`). You can then easily obtain the Gelman-Rubin diagnostic (Gelman et al., 2004), in WinBUGS called R-hat, using `gelman.diag()`, which will indicate if all chains have converged to the same stationary distribution. For details on these and other functions, see the coda manual, which can be found (together with the package) on the CRAN mirror.

7.7 Manipulating the state-space

So far, we have constrained the location of the activity centers to fall within the outermost coordinates of our rectangular state space by posing upper and lower bounds for x and y . But what if S has an irregular shape – maybe there is a large water body we would like to remove from S , because we know our terrestrial study species does not occur there. Or the study takes place in a clearly defined area such as an island. As mentioned before, this situation is difficult to handle in WinBUGS. In some simple cases we can adjust the state space by setting s_{xi} to be some function of s_{yi} or vice versa. In this manner, we can cut off corners of the rectangle to approximate the actual state space. In R, we are much more flexible, as we can use the actual state-space polygon to constrain out s_i .⁵ To illustrate that, let's look at a camera trapping study of Florida panthers (*Puma concolor coryi*) conducted in the Picayune Strand Restoration Project (PSRP) area, southwest Florida (Fig. 7.7), by XXX, and financed by XXX. In the 1960ies the PSRP area was slated for housing development, but then bought back by the State of Florida and is currently being restored to its original hydrology and vegetation. In an effort to estimate the density of the local Florida panther population, 98 camera traps were operated in the area for 21 months between 2005 and 2007. Florida panthers are wide-ranging animals and in order to account for their wide movements, the state-space was defined as the trapping grid buffered by 15 km around its outermost coordinates. However, the resulting rectangle contained some ocean in its southwestern corner (Fig. 7.7). In order to precisely describe the state-space, the ocean has to be removed. You can create a precise state-space polygon in ArcGIS and read it into R, or create the polygon directly within R. In the present case we intersected two shape files – one of the state of Florida and one of the rectangle defined by a strip of 15 km around the camera-trapping grid. While you will most likely have to obtain the shapefile describing the landscape of and around your trapping grid

⁵ Have to check if we can use panther stuff for the book; otherwise, use raccoon example.




Figure 7.7: Rectangular state-space for a Florida panther camera trapping study in the PSRP area (grey outline, red block inset map of Florida) contain some ocean (white) that needs to be removed from the state-space.

1006 (coastlines, water bodies etc.) from some external source, a polygon shapefile
1007 buffering your outermost trapping grid coordinates can easily be written in R.

1008 If x_{\min} , x_{\max} , y_{\min} and y_{\max} , mark the outermost x and y coordinates
1009 of your trapping grid and b is the distance you want to buffer with, load the
1010 package `shapefiles` (Stabler, 2006) and use:

```
1011 xl= xmin-b
1012 xu= xmax+b
1013 yl= ymin-b
1014 yu= ymax+b
```

1015

```
1016 dd <- data.frame(Id=c(1,1,1,1,1),X=c(xl,xu,xu,xl,xl),Y=c(yl,yl,yu,yu,yl)) #create data
1017 ddTable <- data.frame(Id=c(1),Name=c("Item1"))
1018 ddShapefile <- convert.to.shapefile(dd, ddTable, "Id", 5) #convert #to shapefile, type
1019 write.shapefile(ddShapefile, 'c:/, arcgis=T) # save to location of #choice
```

1020 You can read shapefiles into R loading the package `maptools` (Lewin-Koh
1021 et al., 2011) and using the function `readShapeSpatial()`. Make sure you read
1022 in shapefiles in UTM format, so that units of the trap array, the movement
1023 parameter σ and the state-space are all identical. Intersection of polygons
1024 can be done in R also, using the package `rgeos` (Bivand and Rundel, 2011) and
1025 the function `gIntersect()`. The area of your (single) polygon can be extracted
1026 directly from the state-space object `SSp`:

```
1027 > area <- SSp@polygons[[1]]@Polygons[[1]]@area /1000000
```

Note that dividing by 1000000 will return the area in km^2 if your coordinates describing the polygon are in UTM. If your state-space consists of several disjunct polygons, you will have to sum the areas of all polygons to obtain the size of the state-space. To include this polygon into our MCMC sampler we need one last spatial R package, `sp` (Pebesma and Bivand, 2011), which has a function, `over()`, which allows us to check if a pair of coordinates falls within a polygon or not. All we have to do is embed this new check into the updating steps for `s`:

```
Scand <- as.matrix(cbind(rnorm(M, S[,1], 2),
                        rnorm(M, S[,2], 2)))      #draw candidate value
Scoord<-SpatialPoints(Scand*1000)      #convert to spatial points on UTM (m) scale
SinPoly<-over(Scoord,SSp) # check if scand is within the polygon

for(i in 1:M) {
  if(is.na(SinPoly[i])!=FALSE) { #if scand falls within polygon, continue update
    [rest of the updating step remains the same]
```

Note that it is much more time-efficient to draw all M candidate values for s and check once if they fall within the state-space, rather than running the `over()` command for every individual pair of coordinates. To make sure that our initial values for s also fall within the polygon of S , we use the function `runifpoint()` from the package `spatstat` (Baddeley and Turner, 2005), which generates random uniform points within a specified polygon. You'll find this modified MCMC algorithm (`SCR0poisSSp`) in the R packe `scrbook`. Finally, observe that we are converting candidate coordinates of S back to meters to match the UTM polygon. In all previous examples, for both the trap locations and the activity centers we have used UTM coordinates divided by 1000 to estimate σ on a km scale. This is adequate for wide ranging individuals like bears. In other cases you may center all coordinates on 0. No matter what kind of transformation you use on your coordinates, make sure to always convert candidate values for S back to the original scale (UTM) before running the `over()` command.

7.8 MCMC software packages

Throughout most of this book we will use WinBUGS and, occasionally, JAGS to run MCMC analyses. Here, we will briefly discuss the main pros and cons of these two programs as well as WinBUGS successor OpenBUGS.

7.8.1 WinBUGS

In a nutshell, WinBUGS (and the other programs) do everything that we just went through in this chapter (and quite a bit more). Looking through your

model, WinBUGS determines which parameters it can use standard Gibbs sampling for (i.e. for conjugate full conditional distributions). Then, it determines, in the following hierarchy, whether to use adaptive rejection sampling, slice sampling or – in the ‘worst’ case – Metropolis-Hastings sampling for the other full conditionals (Spiegelhalter et al., 2003). If it uses MH sampling, it will automatically tune the updater so that it works efficiently. While WinBUGS is a convenient piece of software that is still widely used, its major drawback is that it is no longer being developed, i.e. no new functions or distributions are added and no bugs are fixed.

7.8.2 OpenBUGS

OpenBUGS is essentially the successor of WinBUGS. While the latter is no longer worked on, OpenBUGS is constantly developed further. The name ‘OpenBUGS’ refers to the software being open source, so users do not need to download a license key, like they have to for WinBUGS (although the license key for WinBUGS is free and valid for life).

Compared to WinBUGS, OpenBUGS has a lot more built-in functions. The method of how to determine the right updater for each model parameter has changed and the user can manually control the MCMC algorithm used to update model parameters. Several other changes have been implemented in OpenBUGS and a detailed list of differences between the two BUGS versions, can be found at <http://www.openbugs.info/w/OpenVsWin>

While OpenBUGS is a useful program for a lot of MCMC sampling applications, for reasons we do not understand, simple SCR models do not converge in OpenBUGS. It is therefore advisable that you check any OpenBUGS SCR model results against result from WinBUGS. Also, currently, the R package BRugs (Thomas et al., 2006), necessary for running OpenBUGS through R, has problems with 64-bit machines, so you may have to use the 32-bit version of R and OpenBUGS in order to make it work. The BUGS project site at <http://www.openbugs.info> provides a lot of information on and download links for OpenBUGS.

There is an extensive help archive for both WinBUGS and OpenBUGS and you can subscribe to a mailing list, where people pose and answer questions of how to use these programs at <http://www.mrc-bsu.cam.ac.uk/bugs/overview/list.shtml>

7.8.3 JAGS – Just Another Gibbs Sampler

JAGS, currently at Version 3.1.0, is another free program for analysis of Bayesian hierarchical models using MCMC simulation. Originally, JAGS was the only program using the BUGS language that would run on operating systems other than the 32 bit Windows platforms. By now, there are OpenBUGS versions for Linux or Macintosh machines. JAGS ‘only’ generates samples from the posterior distribution; analysis of the output is done in R, either by running JAGS through R using either the packages rjags (Plummer, 2011) or R2jags (Su and Yajima, 2011), or by using coda on your JAGS output. The program, manuals and rjags

can be downloaded at <http://sourceforge.net/projects/mcmc-jags/files/> When run from within R using the package `rjags` or `R2jags`, writing a JAGS model is virtually identical to writing a WinBUGS model. However, some functions may have slightly different names and you can look up available functions and their use in the JAGS manual. One potential downside is that JAGS can be very particular when it comes to initial values. These may have to be set as close to truth as possible for the model to start. Although JAGS lets you run several parallel Markov chains, this characteristic interferes with the idea of using overdispersed initial values for the different chains. Also, we have occasionally experienced JAGS to crash and take the R GUI with it. Only re-installing both JAGS and `rjags` seemed to solve this problem. On the plus side, JAGS usually runs a little faster than WinBUGS, sometimes considerably faster (see Sect. ??), is constantly being developed and improved and it has a variety of functions that are not available in WinBUGS. For example, JAGS allows you to supply observed data for some deterministic functions of unobserved variables. In BUGS we cannot supply data to logical nodes. Another useful feature is that the adaptive phase of the model (the burn-in) is run separately from the sampling from the stationary Markov chains. This allows you to easily add more iterations to the adaptive phase if necessary without the need to start from 0. There are other, more subtle differences and there is an entire manual section on differences between JAGS and OpenBUGS. For questions and problems there is a JAGS forum online at <http://sourceforge.net/projects/mcmc-jags/forums/forum/610037>.⁶

7.9 Summary and Outlook

While there are a number of flexible and extremely useful software packages to perform MCMC simulations, it sometimes is more efficient to develop your own MCMC algorithm. Building an MCMC code follows three basic steps: Identify your model including priors and express full conditional distributions for each model parameter. If full conditionals are parametric distributions, use Gibbs sampling to draw candidate parameter values from this distributions; otherwise use Metropolis-Hastings sampling to draw candidate values from a proposal distribution and accept or reject them based on their posterior probability densities. These custom-made MCMC algorithms give you more modeling flexibility than existing software packages, especially when it comes to handling the state-space: In BUGS (and JAGS for that matter) we define a continuous rectangular state-space using the corner coordinates to constrain the Uniform priors on the activity centers s . But what if a continuous rectangle isn't an adequate description of the state-space? In this chapter we saw that in R it only takes a few lines of code to use any arbitrary polygon shapefile as the state-space, which is especially useful when you are dealing with coastlines or large bodies of water that need removing from the state-space. Another example is

⁶As we make progress on the book, let's be sure to add linkages to places where we use JAGS in examples.

1150 the SCR R package SPACECAP (Gopalaswamy et al., 2011) that was developed
 1151 because implementation of an SCR model with a discrete state-space was inef-
 1152 ficient in WinBUGS. Another situations in which using BUGS/JAGS becomes
 1153 increasingly complicated or inefficient is when using point processes other than
 1154 the Uniform Poisson point process which underlies the basic SCR model (see
 1155 Chapt. ??). In the Chapt. ?? and XX you will see examples of different point
 1156 processes, implemented using custom-made MCMC algorithms.⁷ Finally, the
 1157 Chapt. ?? and XX deal with unmarked or partially marked populations using
 1158 hand-made MCMC algorithms to handle the (partially) latent individual en-
 1159 counter histories. While some of these models can be written in BUGS/JAGS,
 1160 ⁸, they are painstakingly slow; others cannot be implemented in BUGS/JAGS
 1161 at all. In conclusion, while you can certainly get by using BUGS/JAGS for
 1162 standard SCR models, knowing how to write your own MCMC sampler allows
 1163 you to tailor these models to your specific needs.

⁷Richard, Beth expand on that?

⁸the Poisson one for partially marked we wrote in BUGS and it should work with a known number of marked; the Bernoulli in JAGS with the dsum() function should work for the fully unknown; maybe some others? I dont remember. We may have to try writing the others before saying that they dont work in BUGS/JAGS; they are certainly much faster in R, though.

1164 Chapter 8

1165 Goodness of Fit and stuff

1166 Chapter 9

1167 Covariate models

1168 **Chapter 10**

1169 **Inhomogeneous Point**
1170 **Process**

1171 Chapter 11

1172 Open models

Bibliography

- Baddeley, A. and Turner, R. (2005), “Spatstat: an R package for analyzing spatial point patterns,” *Journal of Statistical Software*, 12, 1–42, ISSN 1548-7660.
- Bivand, R. and Rundel, C. (2011), *rgeos: Interface to Geometry Engine - Open Source (GEOS)*, r package version 0.1-8.
- Buckland, S. T. (2001), *Introduction to distance sampling: estimating abundance of biological populations*, Oxford, UK: Oxford University Press.
- Casella, G. and George, E. I. (1992), “Explaining the Gibbs sampler,” *American Statistician*, 46, 167–174.
- Gelfand, A. and Smith, A. (1990), “Sampling-based approaches to calculating marginal densities,” *Journal of the American statistical association*, 85, 398–409.
- Gelman, A., Carlin, J. B., Stern, H. S., and Rubin, D. B. (2004), *Bayesian data analysis, second edition.*, Boca Raton, Florida, USA: CRC/Chapman & Hall.
- Geman, S. and Geman, D. (1984), “Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PAMI-6, 721–741.
- Gilks, W. and Wild, P. (1992), “Adaptive rejection sampling for Gibbs sampling,” *Applied Statistics*, 41, 337–348.
- Gilks, W. R., Thomas, A., and Spiegelhalter, D. J. (1994), “A Language and Program for Complex Bayesian Modelling,” *Journal of the Royal Statistical Society. Series D (The Statistician)*, 43, 169–177, ArticleType: primary_article / Issue Title: Special Issue: Conference on Practical Bayesian Statistics, 1992 (3) / Full publication date: 1994 / Copyright 1994 Royal Statistical Society.
- Gopalaswamy, A. M., Royle, A. J., Hines, J., Singh, P., Jathanna, D., Kumar, N. S., and Karanth, K. U. (2011), *A Program to Estimate Animal Abundance and Density using Spatially-Explicit Capture-Recapture*, r package version 1.0.4.

- 1204 Hastings, W. (1970), “Monte Carlo sampling methods using Markov chains and
1205 their applications,” *Biometrika*, 57, 97–109.
- 1206 Lewin-Koh, N. J., Bivand, R., contributions by Edzer J. Pebesma, Archer, E.,
1207 Baddeley, A., Bibiko, H.-J., Dray, S., Forrest, D., Friendly, M., Giraudoux, P.,
1208 Golicher, D., Rubio, V. G., Hausmann, P., Hufthammer, K. O., Jagger, T.,
1209 Luque, S. P., MacQueen, D., Niccolai, A., Short, T., Stabler, B., and Turner,
1210 R. (2011), *maptools: Tools for reading and handling spatial objects*, r package
1211 version 0.8-10.
- 1212 Link, W. A. and Barker, R. J. (2009), *Bayesian Inference: With Ecological*
1213 *Applications*, London, UK: Academic Press.
- 1214 Metropolis, N., Rosenbluth, A., Rosenbluth, M., Teller, A., Teller, E., et al.
1215 (1953), “Equation of state calculations by fast computing machines,” *The*
1216 *journal of chemical physics*, 21, 1087–1092.
- 1217 Metropolis, N. and Ulam, S. (1949), “The Monte Carlo method,” *Journal of the*
1218 *American Statistical Association*, 44, 335–341.
- 1219 Neal, R. (2003), “Slice sampling,” *Annals of Statistics*, 31, 705–741.
- 1220 Pebesma, E. and Bivand, R. (2011), *Package ‘sp’*, r package version 0.9-91.
- 1221 Plummer, M. (2011), *rjags: Bayesian graphical models using MCMC*, r package
1222 version 3-5.
- 1223 Plummer, M., Best, N., Cowles, K., and Vines, K. (2006), “CODA: Convergence
1224 Diagnosis and Output Analysis for MCMC,” *R News*, 6, 7–11.
- 1225 Robert, C. P. and Casella, G. (2004), *Monte Carlo statistical methods*, New
1226 York, USA: Springer.
- 1227 — (2010), *Introducing Monte Carlo Methods with R*, New York, USA: Springer.
- 1228 Roberts, G. O. and Rosenthal, J. S. (1998), “Optimal scaling of discrete ap-
1229 proximations to Langevin diffusions,” *Journal of the Royal Statistical Society:*
1230 *Series B (Statistical Methodology)*, 60, 255–268.
- 1231 Spiegelhalter, D., Thomas, A., Best, N., and Lunn, D. (2003), *WinBUGS User*
1232 *Manual Version 1.4*.
- 1233 Stabler, B. (2006), *shapefiles: Read and Write ESRI Shapefiles*, r package version
1234 0.6.
- 1235 Su, Y.-S. and Yajima, M. (2011), *R2jags: A Package for Running jags from R*,
1236 r package version 0.02-17.
- 1237 Thomas, A., O’Hara, B., Ligges, U., and Sturtz, S. (2006), “Making BUGS
1238 Open,” *R News*, 6, 12–17.