

1

2

Spatial Capture-Recapture

3

The Four Horsemen (and women)

4

USGS Patuxent Wildlife Research Center
North Carolina State University

5

6

ACADEMIC PRESS

CONTENTS

9	1 Introduction	9
10	2 Bayesian Analysis of GL(M)Ms Using R/WinBUGS	11
11	2.1 Notation	12
12	2.2 GLMs and GLMMs	13
13	2.3 Bayesian Analysis	14
14	2.3.1 Bayes Rule	15
15	2.3.2 Bayesian Inference	16
16	2.3.3 Prior distributions	18
17	2.3.4 Posterior Inference	18
18	2.3.5 Small sample inference	19
19	2.4 Characterizing posterior distributions by MCMC simulation	20
20	2.5 What Goes on Under the MCMC Hood	21
21	2.5.1 Rules for constructing full conditional distributions	22
22	2.5.2 Metropolis-Hastings algorithm	23
23	2.6 Practical Bayesian Analysis and MCMC	24
24	2.6.1 Choice of prior distributions	24
25	2.6.2 Convergence and so-forth	25
26	2.6.3 Bayesian confidence intervals	26
27	2.6.4 Estimating functions of parameters	27
28	2.7 Bayesian Analysis using WinBUGS	27
29	2.7.1 Linear Regression in WinBUGS	27
30	2.7.2 Inference about functions of model parameters	30
31	2.8 Model Checking and Selection	30
32	2.8.1 Goodness-of-fit	31
33	2.8.2 Model Selection	32
34	2.9 Poisson GLMs	33
35	2.9.1 Important properties of the Poisson distribution	34
36	2.9.2 Example: Breeding Bird Survey Data	34
37	2.9.3 Doing it in WinBUGS	36
38	2.9.4 Constructing your own MCMC algorithm	38
39	2.10 Poisson GLM with Random Effects	43
40	2.11 Binomial GLMs	45

41	2.11.1	Binomial regression	45
42	2.11.2	Example: Waterfowl Banding Data	46
43	2.12	Summary and Outlook	48
44	3	Estimating the Size of a Closed Population	51
45	3.1	The Simplest Closed Population Model: Model M0	52
46	3.1.1	The Spatial Context of Capture-Recapture	54
47	3.1.2	Conditional likelihood	54
48	3.2	Data Augmentation	55
49	3.2.1	DA links occupancy models and closed population models . .	56
50	3.2.2	Model M_0 in BUGS	57
51	3.2.3	Formal development of data augmentation	59
52	3.2.4	Remarks on Data Augmentation	60
53	3.2.5	Example: Black Bear Study on Fort Drum	61
54	3.3	Temporally varying and behavioral effects	65
55	3.4	Models with individual heterogeneity	66
56	3.4.1	Analysis of Model Mh	69
57	3.4.2	Analysis of the Fort Drum data	70
58	3.4.3	Building your own MCMC algorithm	72
59	3.4.4	Exercises related to model Mh	76
60	3.5	Individual Covariate Models: Toward Spatial Capture-Recapture . .	76
61	3.5.1	Example: Location of capture as a covariate.	78
62	3.5.2	Extension of the Model	81
63	3.5.3	Invariance of density to maxD	84
64	3.5.4	Toward Fully Spatial Capture-recapture Models	84
65	3.6	DISTANCE SAMPLING: A primitive Spatial Capture-Recapture	
66		Model	85
67	3.6.1	Example: Muntjac deer survey from Nagarahole, India . . .	87
68	3.7	Summary and Outlook	88
69	4	Fully Spatial Capture-Recapture Models	91
70	4.1	Sampling Design and Data Structure	92
71	4.2	The binomial observation model	93
72	4.2.1	Distance as a latent variable	94
73	4.3	The Binomial Point-process Model	95
74	4.3.1	Definition of home range center	96
75	4.3.2	The state-space of the point process	97
76	4.3.3	Invariance and the State-space as a model assumption	98
77	4.3.4	Connection to Model Mh	99
78	4.3.5	Connection to Distance Sampling	100
79	4.4	Simulating SCR Data	100
80	4.4.1	Formatting and manipulating real data sets	102
81	4.5	Fitting an SCR Model in BUGS	102

82	4.6	Unknown N	105
83	4.6.1	Analysis using data augmentation in WinBUGS	106
84	4.6.2	Use of other BUGS engines: JAGS	109
85	4.7	Case Study: Wolverine Camera Trapping Study	110
86	4.7.1	Fitting the model in WinBUGS	113
87	4.7.2	Conclusion of Analysis	116
88	4.8	Constructing Density Maps	117
89	4.8.1	Example: Wolverine density map.	119
90	4.9	Discrete State-Space	120
91	4.9.1	Evaluation of Coarseness of Discrete Approximation	122
92	4.9.2	Analysis of the wolverine camera trapping data	123
93	4.9.3	SCR models as multi-state models	126
94	4.10	Summary and Outlook	127
95	5	Other observation models	131
96	6	Likelihood Analysis of Spatial Capture-Recapture Models	133
97	6.1	Likelihood analysis	134
98	6.1.1	Implementation (simulated data)	136
99	6.2	MLE when N is Unknown	139
100	6.2.1	Exercises	142
101	6.2.2	Integrated Likelihood using the model under data augmenta- tion	143
102	6.2.3	Extensions	143
103	6.3	Classical model selection and assessment	144
104	6.4	Likelihood analysis of the wolverine camera trapping data	144
105	6.4.1	Restricted state-space	147
106	6.4.2	Exercises	147
107	6.5	Program DENSITY and the R package <code>secr</code>	148
108	6.5.1	Analysis using the <code>secr</code> package	149
109	6.6	Summary and Outlook	151
110			
111	7	MCMC for Spatial Capture-Recapture	153
112	7.1	Introduction	153
113	7.1.1	Why build your own MCMC algorithm?	153
114	7.2	MCMC and posterior distributions	154
115	7.3	Types of MCMC sampling	157
116	7.3.1	Gibbs sampling	157
117	7.3.2	Metropolis-Hastings sampling	160
118	7.3.3	Metropolis-within-Gibbs	164
119	7.4	GLMMs Poisson regression with a random effect	164
120	7.4.1	Rejection sampling and slice sampling	168
121	7.5	MCMC for closed capture-recapture Model Mh	168

122	7.6	MCMC algorithm for the basic spatial capture-recapture model . . .	169
123	7.6.1	SCR model with binomial encounter process	172
124	7.6.2	Looking at model output	173
125	7.6.3	Posterior density plots	176
126	7.6.4	Serial autocorrelation and effective sample size	176
127	7.6.5	Summary results	178
128	7.6.6	Other useful commands	180
129	7.7	Manipulating the state-space	180
130	7.8	MCMC software packages	183
131	7.8.1	WinBUGS	183
132	7.8.2	OpenBUGS	183
133	7.8.3	JAGS Just Another Gibbs Sampler	184
134	7.9	Summary and Outlook	185
135	8	Goodness of Fit and stuff	187
136	9	Modeling Encounter Probability	189
137	10	Ecological Distance Models in Spatial Capture-Recapture	191
138	10.1	Cost Distance	191
139	10.1.1	Illustration: Example Good vs. Bad habitat	192
140	10.2	Distance Weighting	192
141	10.3	Hard Boundary	192
142	10.4	Simulation Study	192
143	10.5	Estimating Cost or Resistance Values	192
144	11	State-space Covariates	193
145	11.1	Homogeneous point process revisited	194
146	11.2	Inhomogeneous binomial point process	195
147	11.3	Examples	197
148	11.3.1	Simulation and analysis of inhomogeneous point processes . .	197
149	11.3.2	Fitting inhomogeneous point process SCR models	200
150	11.3.3	The jaguar data	204
151	11.4	Summary	206
152	12	Open models	209
153	13	Spatial Capture-Recapture for Unmarked Populations	211
154	13.1	Existing Models for Inference About Density in Unmarked Populations	212
155	13.2	Spatial Correlation as Information	213
156	13.3	Data Requirements and Survey Designs	214
157	13.4	Encounter Histories as Latent Variables	215
158	13.5	Estimation by MCMC	215

159	13.6 Northern Parula Example	216
160	13.7 On (Im)precision	218
161	13.8 How Much Correlation Is Enough?	218
162	13.9 Mutants	218
163	13.9.1 Other observation models	218
164	13.9.2 Linear designs	218
165	13.10Summary	218
166	13.10.1 Alternative Observation Models	221
167	13.10.2 Spatial point process models	221
168	13.11Conclusion	222
169	Bibliography	223

170
171
172

1

INTRODUCTION

2

BAYESIAN ANALYSIS OF GL(M)MS USING R/WINBUGS

A major theme of this book is that spatial capture-recapture models are, for the most part, just generalized linear models (GLMs) wherein the covariate, distance between trap and home range center, is partially or fully unobserved – and therefore regarded as a random effect. Such models are usually referred to as Generalized Linear Mixed Models (GLMMs) and, therefore, SCR models can be thought of as a specialized type of GLMM. Naturally then, we should consider analysis of these slightly simpler models in order to gain some experience and, hopefully, develop a better understanding of spatial capture-recapture models.

In this chapter, we consider classes of GLM models - Poisson and binomial (i.e., logistic regression) GLMs - that will prove to be enormously useful in the analysis of capture-recapture models of all kinds. Many readers are probably familiar with these models because they represent probably the most generally useful models in all of Ecology and, as such, have received considerable attention in many introductory and advanced texts. We focus on them here in order to introduce the readers to the analysis of such models in **R** and **WinBUGS**, which we will translate directly to the analysis of SCR models in subsequent chapters.

Bayesian analysis is convenient for analyzing GLMMs because it allows us to work directly with the conditional model – i.e., the model that is conditional on the random effects, using computational methods known as Markov chain Monte Carlo (MCMC). Learning how to do Bayesian analysis of GLMs and GLMMs in **WinBUGS** is, in part, the purpose of this chapter. While we use **WinBUGS** to do the Bayesian computations, we organize and summarize our data and execute **WinBUGS** from within **R** using the useful package **R2WinBUGS** (Sturtz et al., 2005). Kéry (2010), and Kéry and Schaub (2011) provide excellent introductions to the

basics of Bayesian analysis and GLMs at an accessible level. We don't want to be too redundant with those books and so we avoid a detailed treatment of Bayesian methodology - instead just providing a cursory overview so that we can move on and attack the problems we're most interested in related to spatial capture-recapture. In addition, there are a number of texts that provide general introductions to Bayesian analysis, MCMC, and their applications in Ecology including McCarthy (2007), Kéry (2010), Link and Barker (2009), and King (2009).

While this chapter is about Bayesian analysis of GLMMS, such models are routinely analyzed using likelihood methods too, as discussed by Royle and Dorazio (2008), and Kéry (2010). Indeed, likelihood analysis of such models is the primary focus of many applied statistics texts, a good one being Zuur et al. (2009). Later in this book, we will use likelihood methods to analyze SCR models but, for now, we concentrate on providing a basic introduction to Bayesian analysis because that is the approach we will use in a majority of cases in later chapters.

2.1 NOTATION

We will sometimes use conventional "bracket notation" to refer to probability distributions. If y is a random variable the $[y]$ indicates its distribution or its probability density/mass function (pdf, pmf) depending on context. If x is another random variable then $[y|x]$ is the conditional distribution of y given x , and $[y, x]$ is the joint distribution of y and x . To differentiate specific distributions in some contexts we might label them $g(y)$, $g(y|\theta)$, $f(x)$, or similar. We will also write $y \sim \text{Normal}(\mu, \sigma^2)$ to indicate that y "is distributed as" a normal random variable with parameters μ and σ^2 . The expected value or mean of a random variable is $E[y] = \mu$, and $\text{Var}[y] = \sigma^2$ is the variance of y . To indicate specific observations we'll use an index such as " i ". So, y_i for $i = 1, 2, \dots, n$ indicates observations for n individuals. Finally, we write $\text{Pr}(y)$ to indicate specific probabilities, i.e., of events " y " or similar.

To illustrate these concepts and notation, suppose z is a binary outcome (e.g., species occurrence) and we might assume the model: $z \sim \text{Bern}(p)$ for observations. Under this model $\text{Pr}(z = 1) = \psi$, which is also the expected value $E[z] = \psi$. The variance is $\text{Var}[z] = \psi * (1 - \psi)$ and the probability mass function (pmf) is $[z] = \psi^z (1 - \psi)^{1-z}$. Sometimes we write $[z|\psi]$ when it is important to emphasize the conditional dependence of z on ψ . As another example, suppose y is a random variable denoting whether or not a species is detected if an occupied site is surveyed. In this case it might be natural to express the pmf of the observations y *conditional* on z . That is, $[y|z]$. In this case, $[y|z = 1]$ is the conditional pmf of y given that a site is occupied, and it is natural to assume that $[y|z = 1] = \text{Bern}(p)$ where p is the "detection probability" - the probability that we detect the species, given that it is present. The model for the observations y is completely specified once we describe the other conditional pmf $[y|z = 0]$. For this conditional distribution it is sometimes reasonable to assume $\text{Pr}(y = 1|z = 0) = 0$ (MacKenzie et al. (2002); see also Royle

and Link (2006)). That is, if the species is absent, the probability of detection is 0. This implies that $\Pr(y = 0|z = 0) = 1$. To allow for situations in which the true state z is unobserved, we assume that $[z]$ is Bernoulli with parameter ψ . In this case, the marginal distribution of y is

$$[y] = [y|z = 1]Pr(z = 1) + [y|z = 0]Pr(z = 0)$$

because $[y|z = 0]$ is a point mass at $y = 0$, by assumption, then

$$\Pr(y = 1) = p\psi$$

And

$$\Pr(y = 0) = (1 - p) * \psi + (1 - \psi)$$

2.2 GLMS AND GLMMS

We have asserted already that SCR models work out most of the time to be variations of GLMs and GLMMS. Some of you might therefore ask: What are GLMs and GLMMS, anyhow? These models are covered extensively in many very good applied statistics books and we refer the reader elsewhere for a detailed introduction. We think Kéry (2010), Kéry and Schaub (2011), and Zuur et al. (2009) are all accessible treatments of considerable merit. Here, we'll give the 1 minute treatment of GLMMS, not trying to be complete but rather only to preserve a coherent organization to the book.

The generalized linear model (GLM) is an extension of standard linear models by allowing the response variable to have some distribution from the exponential family of distributions (i.e., not just normal). This includes the normal distribution but also dozens of others such as the Poisson, binomial, gamma, exponential, and many more. In addition, GLMS allow the response variable to be related to the predictor variables (i.e., covariates) using a link function, which is usually nonlinear. Finally, GLMs typically accommodate a relationship between the mean and variance. The classical reference for GLMs is Nelder and Wedderburn (1972) and also McCullagh and Nelder (1989). The GLM consists of three components:

1. A probability distribution for the dependent variable y , from a class of probability distributions known as the exponential family.
2. A "linear predictor" $\eta = \mathbf{X}\beta$.
3. A link function g that relates $E[y]$ to the linear predictor, $E[y] = \mu = g^{-1}(\eta)$. Therefore $g(E[y]) = \eta$.

The dependent variable y is assumed to be an outcome from a distribution of the exponential family which includes many common distributions including the normal, gamma, Poisson, binomial, and many others. The mean of the distribution of y is assumed to depend on predictor variables x according to

$$g(E[y]) = \mathbf{x}'\beta$$

where $E[y]$ is the expected value of y , and $\mathbf{x}'\beta$ is termed the *linear predictor*, i.e., a linear function of the predictor variables with unknown parameters β to be estimated. The function g is the link function. In standard GLMs, the variance of y is a function V of the mean of y : $\text{Var}(y) = V(\mu)$ (see below for examples).

A Poisson GLM posits that $y \sim \text{Poisson}(\lambda)$ with $E[y] = \lambda$ and usually the model for the mean is specified using the *log link function* by

$$\log(\lambda_i) = \beta_0 + \beta_1 * x_i$$

The variance function is $V(y_i) = \lambda_i$. The binomial GLM posits that $y_i \sim \text{Binomial}(K, p)$ where K is the fixed sample size parameter and $E[y_i] = K * p_i$. Usually the model for the mean is specified using the *logit link function* according to

$$\text{logit}(p_i) = \beta_0 + \beta_1 * x_i$$

Where $\text{logit}(u) = \log(u/(1-u))$. The inverse-logit function, g^{-1} , is a function we will refer to as “expit”, so that $\text{expit}(u) = \exp(u)/(1 + \exp(u))$.

A GLMM is the extension of GLMs to accommodate “random effects”. Often this involves adding a normal random effect to the linear predictor, and so a simple example is:

$$\log(\lambda_i) = \alpha_i + \beta_1 * x_i$$

where

$$\alpha_i \sim \text{Normal}(\mu, \sigma^2)$$

2.3 BAYESIAN ANALYSIS

Bayesian analysis is unfamiliar to many ecological researchers because older cohorts of ecologists were largely educated in the classical statistical paradigm of frequentist inference. But advances in technology and increasing exposure to benefits of Bayesian analysis are fast making Bayesians out of people or at least making Bayesian analysis an acceptable, general, alternative to classical, frequentist inference.

Conceptually, the main thing about Bayesian inference is that it uses probability directly to characterize uncertainty about things we don’t know. “Things”, in this case, are parameters of models and, just as it is natural to characterize uncertain outcomes of stochastic processes using probability, it seems natural also to characterize information about unknown “parameters” using probability. At least this seems natural to us and, we think, most ecologists either explicitly adopt that view or tend to fall into that point of view naturally. Conversely, frequentists use probability in many different ways, but never to characterize uncertainty about parameters¹ Instead, frequentists use probability to characterize the behavior of *procedures* such as estimators or confidence intervals (see below), which can lead to

¹To hear this will be shocking to some readers perhaps.

some inelegant or unnatural interpretations of things. It is paradoxical that people readily adopt a philosophy of statistical inference in which the things you don't know (i.e., parameters) should *not* be regarded as random variables, so that, as a consequence, one cannot use probability to characterize one's state of knowledge about them.

2.3.1 Bayes Rule

As its name suggests, Bayesian analysis makes use of Bayes' rule in order to make direct probability statements about model parameters. Given two random variables z and y , Bayes rule relates the two conditional probability distributions $z|y$ and $y|z$ by the relationship:

$$[z|y] = [y|z][z]/[y]$$

Bayes' rule itself is a mathematical fact and there is no debate in the statistical community as to its validity and relevance to many problems. Generally speaking, these distributions are characterized as follows: $[y|z]$ is the conditional probability distribution of y given z , $[z]$ is the marginal distribution of z and $[y]$ is the marginal distribution of y . In the context of Bayesian inference we usually associate specific meanings in which $[y|z]$ is thought of as "the likelihood", $[z]$ as the "prior" and so on. We leave this for later because here the focus is on this expression of Bayes rule as a basic fact of probability.

As an example of a simple application of Bayes rule, consider the problem of determining species presence at a sample location based on imperfect survey information. Let z be a binary random variable that denotes species presence ($z = 1$) or absence ($z = 0$), let $\Pr(z = 1) = \psi$ where ψ is usually called occurrence probability, "occupancy" (MacKenzie et al., 2002) or "prevalence". Let y be the *observed* presence ($y = 1$) or absence ($y = 0$), and let p be the probability that a species is detected in a single survey at a site given that it is present. Thus, $\Pr(y = 1|z = 1) = p$. The interpretation of this is that, if the species is present, we will only observe presence with probability p . In addition, we assume here that $\Pr(y = 1|z = 0) = 0$. That is, the species cannot be detected if it is not present which is a conventional view adopted in most biological sampling problems (but see Royle and Link (2006)). If we survey a site T times but never detect the species, then this clearly does not imply that the species is not present ($z = 0$) at this site. Rather, our degree of belief in $z = 0$ should be made with a probabilistic statement $\Pr(z = 1|y_1 = 0, \dots, y_T = 0)$. If the T surveys are independent so that we might regard y_t as *iid* Bernoulli trials, then the total number of detections, say y , is Binomial with probability p then we can use Bayes rule to compute the probability that it is present given that it is not detected in T samples. In words, the expression we seek is:

$$\Pr(\text{present}|\text{not detected}) = \frac{\Pr(\text{not detected}|\text{present}) \Pr(\text{present})}{\Pr(\text{detected})}$$

341 Mathematically, this is

$$\begin{aligned}\Pr(z = 1|y = 0) &= \Pr(y = 0|z = 1) \Pr(z = 1) / \Pr(y = 0) \\ &= [(1 - p)^T \psi] / [(1 - p)^T \psi + (1 - \psi)].\end{aligned}$$

342 To apply this, suppose that $T = 2$ surveys are done at a wetland for a species of
343 frog, and the species is not detected there. Suppose further that $\psi = .8$ and $p = .5$
344 are obtained from a prior study. Then the probability that the species is present at
345 this site is $.25 * .8 / (.25 * .8 + .2) = 0.50$. That is, there seems to be about a 50/50
346 chance that the site is occupied despite the fact that the species wasn't observed
347 there.

348 In summary, Bayes' rule provides a simple linkage between the conditional prob-
349 abilities $[y|z]$ and $[z|y]$ which is useful whenever one needs to deduce one from the
350 other. Bayes' rule as a basic fact of probability is not disputed.

351 2.3.2 Bayesian Inference

352 What is controversial to some is the scope and manner in which Bayes rule is ap-
353 plied by Bayesian analysts. Bayesian analysts assert that Bayes rule is relevant,
354 in general, to all statistical problems by regarding all unknown quantities of a
355 model as realizations of random variables - this includes "data", latent variables,
356 and also "parameters". Classical (non-Bayesian) analysts sometimes object to re-
357 garding "parameters" as outcomes of random variables. Classically, parameters are
358 thought of as "fixed but unknown" (using the terminology of classical statistics).
359 Of course, in Bayesian analysis they are also unknown and, in fact, there is a single
360 data-generating value and so they are also fixed. The difference is that this fixed
361 but unknown value is regarded as having been generated from some probability
362 distribution. Specification of that probability distribution is necessary to carryout
363 Bayesian analysis, but it is not required in classical frequentist inference.

364 To see the general relevance of Bayes rule in the context of statistical inference,
365 let y denote observations - i.e., "data" - and let $[y|\theta]$ be the observation model
366 (often colloquially referred to as the "likelihood"). Suppose θ is a parameter of
367 interest having (prior) probability distribution $[\theta]$. These are combined to obtain
368 the posterior distribution using Bayes' rule, which is:

$$[\theta|y] = [y|\theta][\theta]/[y]$$

369 Asserting the general relevance of Bayes rule to all statistical problems, we can con-
370 clude that the two main features of Bayesian inference are that: (1) "parameters"
371 θ are regarded as realizations of a random variable and, as a result, (2) inference is
372 based on the probability distribution of the parameters given the data, $[\theta|y]$, which
373 is called the posterior distribution. This is the result of using Bayes rule to combine
374 "the likelihood" and the prior distribution. The key concept is regarding param-
375 eters as realizations of a random variable because, once you admit this conceptual

view, this leads directly to the posterior distribution, a very natural quantity upon which to base inference about things we don't know - including parameters of statistical models. In particular, $[\theta|y]$ is a probability distribution for θ and therefore we can make direct probability statements to characterize uncertainty about θ .

The denominator of our invocation of Bayes rule, $[y]$, is the marginal distribution of the data y . We note without further remark right now that, in many practical problems, this can be an enormous pain to compute. The main reason that the Bayesian paradigm has become so popular in the last 20 years or so is because methods exist for characterizing the posterior distribution that do not require that we possess a mathematical understanding of $[y]$, i.e., we never have to compute it or know what it looks like, or know anything specific about it.

A common misunderstanding on the distinction between Bayesian and frequentist inference goes something like this "in frequentist inference parameters are fixed but unknown but in a Bayesian analysis parameters are random." At best this is a sad caricature of the distinction and at worst it is downright wrong. What is true is that, to a Bayesian, parameters are random variables. However, a Bayesian assumes, just like a frequentist, that there was a single data-generating value of that parameter - a fixed, and unknown value that produced the given data set. The distinction between Bayesian and frequentist approaches is that Bayesians regard the parameter as a random variable, and its value as the outcome of a random value, on par with the observations. This allows Bayesians to use probability to make direct probability statements about parameters. Frequentist inference procedures do not permit direct probability statements to be made about parameter values - because parameters are not random variables!

While we can understand the conceptual basis of Bayesian inference merely by understanding Bayes rule - that's really all there is to it - it is not so easy to understand the basis of classical "frequentist" inference which is mostly like² a "basket of methods" with little coherent organization. What is mostly coherent in frequentist inference is the manner in which items in this basket of methods are evaluated - the performance of a given procedure is evaluated by "averaging over" hypothetical realizations of y , regarding the *estimator* as a random variable. For example, if $\hat{\theta}$ is an estimator of θ then the frequentist is interested in $E_y[\hat{\theta}|y]$ which is used to characterize bias. If the expected value of $\hat{\theta}$, when averaged over realizations of y , is equal to θ , then $\hat{\theta}$ is unbiased.

The view of parameters as fixed constants and estimators as random variables leads to interpretations that are not so straightforward. For example confidence intervals having the interpretation "95% probability that the interval contains the true value" and p-values being "the probability of observing an outcome as extreme or more than the one observed." These are far from intuitive interpretations to most people. Moreover, this is conceptually problematic to some because the hypothetical realizations that characterize the performance of our procedure we

²Characterization from Sims REF XYZ

will never get to observe.

While we do tend to favor Bayesian inference for the conceptual simplicity (parameters are random, posterior inference), we mostly advocate for a pragmatic non-partisian approach to inference because, frankly, some of these “bucket of methods” are actually very convenient in certain situations as we will see in later chapters.

2.3.3 Prior distributions

The prior distribution $[\theta]$ is an important feature of Bayesian inference. As a conceptual matter, the prior distribution characterizes “prior beliefs” or “prior information” about a parameter. Indeed, an oft-touted benefit of Bayesian analysis is the ease with which prior information can be included in an analysis. However, more commonly, the prior is chosen to express a lack of prior information, even if previous studies have been done and even if the investigator does in fact know quite a bit about a parameter. This is because the manner in which prior information is embodied in a prior (and the amount of information) is usually very subjective and thus the result can wind up being very contentious, e.g., if different investigators might report different results based on subjective assessments of things. Thus it is usually better to “let the data speak” and use priors that reflect absence of information beyond the data set being analyzed.

But still the need occasionally arises to embody prior information or beliefs about a parameter formally into the estimation scheme. In SCR models we often have a parameter that is closely linked to “home range radius” and thus auxiliary information on the home range size of a species can be used as prior information (e.g., see Chandler and Royle (2012) ; also chapter XYZ).

XXXXXXXXXX

noninformative prior on one scale is informative on another scale. e.g., flat prior on $\logit(p)$ is very different from $\text{uniform}(0,1)$ on p ... show graphic.....

reference to non-invariance of prior distributions to transformation.....

XXXXXXXXXX

2.3.4 Posterior Inference

In Bayesian inference, we are not focusing on estimating a single point or interval but rather on characterizing a whole distribution – the posterior distribution – from which one can report any summary of interest. A point estimate might be the posterior mean, median, mode, etc.. In many applications in this book, we will compute 95% Bayesian intervals using the 2.5% and 97.5% quantiles of the posterior distribution. For such intervals, it is correct to say $\Pr(L < \theta < U) = 0.95$. That is, “the probability that θ is between L and U is 0.95”. It is not a subtle thing that this cannot be said using frequentist methods - although people tend to say it anyway and not really understand why it is wrong or even that it is wrong. This is

actually a failing of frequentist ideas and the inability of frequentists to get people to overcome their natural tendency to use probability - which is something that, as a frequentist, you simply cannot do in the manner that you would like to.

Posterior inference is the main practical element of Bayesian analysis. We get to make an inference conditional on the data that we actually observed - i.e., what we actually know. To us, this seems logical - to condition on what we know. Conversely, frequentist inference is based on considering average performance over hypothetical unobserved data sets (i.e., the “relative frequency” interpretation of probability). Frequentists know that their procedures work well when averaged over all hypothetical, unobserved, data sets but no one ever really knows how well they work for the specific data set analyzed. That seems like a relevant question to biologists who oftentimes only have their one, extremely valuable, data set. This distinction comes into play a lot in exposing philosophical biases in the peer review of statistical analyses in ecology in the sense that, despite these opposing conceptual views to inference (i.e. conditional on the data you have, or averaged over hypothetical realizations), those who conduct a Bayesian analysis are often (in ecology, almost always) required to provide a frequentist evaluation of their Bayesian procedure.

2.3.5 Small sample inference

Using Bayesian inference, we obtain an estimate of the posterior distribution which is an exhaustive summary of the state-of-knowledge about an unknown quantity. It is the posterior distribution - not an estimate of that thing. It is also not, usually, an approximation except to within Monte Carlo error (in cases where we use simulation to calculate it). One of the great virtues of Bayesian analysis which is not really appreciated is that it is completely valid for any particular sample size. i.e., it is $[\theta|y]$, as precise as we claim it to be based on our ability to do calculations, for the particular sample size and observations that we have even if we have only a single datum y . The same cannot be said for almost all frequentist procedures in which estimates or variances are very often (almost always in practice) based on “asymptotic approximations” to the procedure which is actually being employed.

There seems to be a prevailing view in statistical ecology that classical likelihood-based procedures are virtuous because of the availability of simple formulas and procedures for carrying out inference, such as calculating standard errors, doing model selection by AIC, and assessing goodness-of-fit. In large samples, this may be an important practical benefit, but the theoretical validity of these procedures cannot be asserted in most situations involving small samples. This is not a minor issue because it is typical in many wildlife sampling problems - especially in surveys of carnivores or rare/endangered species - to wind up with a small, sometimes extremely small, data set. For example, a recent paper on the fossa (*Cryptoprocta ferox*), an endangered carnivore in Madagascar, estimated an adult density of 0.18 adults / km sq based on 20 animals captured over 3 years (Hawkins and Racey,

2005). A similar paper on the endangered southern river otter (*Lontra provocax*)
 estimated a density of 0.25 animals per river km based on 12 individuals captured
 over 3 years (Sepúlveda et al., 2007). Gardner et al. (2010) analyzed data from a
 study of the Pampas cat, a species for which very little is known, wherein only 22
 individual cats were captured during the two year period. Trolle and Kéry (2005)
 reported only 9 individual ocelots captured and Jackson et al. (2006) captured 6
 individual snow leopards using camera trapping. Thus, studies of rare and/or secre-
 tive carnivores necessarily and flagrantly violate one of Le Cam’s Basic Principles,
 that of “If you need to use asymptotic arguments, do not forget to let your number
 of observations tend to infinity.” (Le Cam, 1990).

The biologist thus faces a dilemma with such data. On one hand, these datasets,
 and the resulting inference, are often criticized as being poor and unreliable. Or,
 even worse³, “the data set is so small, this is a poor analysis.” On the other
 hand, such data may be all that is available for species that are extraordinarily
 important for conservation and management. The Bayesian framework for inference
 provides a valid, rigorous, and flexible framework that is theoretically justifiable in
 arbitrary sample sizes. This is not to say that one will obtain precise estimates
 of density or other parameters, just that your inference is coherent and justifiable
 from a conceptual and technical statistical point of view. That is, we report the
 posterior probability $\Pr(D|data)$ which is easily interpretable and just what it is
 advertised to be and we don’t need to do a simulation study to evaluate how well
 some approximate $\Pr(D|data)$ deviates from the actual $\Pr(D|data)$ because they
 are precisely the same quantity.

2.4 CHARACTERIZING POSTERIOR DISTRIBUTIONS BY MCMC SIMULATION

In practice, it is not really feasible to ever compute the marginal probability dis-
 tribution $\Pr(y)$, the denominator resulting from application of Bayes’ rule. For
 decades this impeded the adoption of Bayesian methods by practitioners. Or, the
 few Bayesian analyses done were based on asymptotic normal approximations to the
 posterior distribution. While this was useful stuff from a theoretical and technical
 standpoint and, practically, it allowed people to make the probability statements
 that they naturally would like to make, it was kind of a bad joke around the
 Bayesian water-cooler to, on one hand, criticize classical statistics for being, essen-
 tially, completely ad hoc in their approach to things but then, on the other hand,
 have to devise various approximations to what they were trying to characterize.
 The advent of Markov chain Monte Carlo (MCMC) methods has made it easier to
 calculate posterior distributions for just about any problem to arbitrary levels of
 precision.

Broadly speaking, MCMC is a class of methods for drawing random numbers

³Actual quote from a referee

(sampling or simulating) from the target posterior distribution. Thus, even though we might not recognize the posterior as a named distribution or be able to analyze its features analytically, e.g., devise mathematical expressions for the mean and variance, we can use these MCMC methods to obtain a large sample from the posterior and then use that sample to characterize features of the posterior. What we do with the sample depends on our intentions – typically we obtain the mean or median for use as a point estimate, and take a confidence interval based on Monte Carlo estimates of the quantiles. These are estimates, but not like frequentist estimates. Rather, they are Monte Carlo estimates with an associated Monte Carlo error which is largely determined arbitrarily by the analyst. They are not estimates qualified by a sampling distribution as in classical statistics. If we run our MCMC long enough then our reported value of $E[\theta|y]$ or any feature of the posterior distribution is precisely what we say it is. There is no “sampling variation” in the frequentist sense of the word. In summary, the MCMC samples provide a Monte Carlo characterization of *the* posterior distribution.

2.5 WHAT GOES ON UNDER THE MCMC HOOD

We will develop and apply MCMC methods in some detail for spatial capture-recapture models in chapter 7. Here we provide a simple illustration of some basic ideas related to the practice of MCMC.

A type of MCMC method relevant to most problems is Gibbs sampling (REF XYZ XYZ), which is based on the idea of iterative simulation from the “full conditional” distributions (also called conditional posterior distributions). The full conditional distribution for an unknown quantity is the conditional distribution of that quantity given every other random variable in the model - the data and all other parameters. For example, for a normal regression model with $y \sim \text{Normal}(\alpha + \beta x, 1)$ then the two full conditionals are, in symbolic terms,

$$[\alpha|y, \beta]$$

and

$$[\beta|y, \alpha].$$

We might use our knowledge of probability to identify these mathematically. In particular, by Bayes’ Rule, $[\alpha|y, \beta] = [y|\alpha, \beta][\alpha|\beta]/[y|\beta]$ and similarly for $[\beta|y, \alpha]$. For example, if we have priors for $[\alpha]$ and $[\beta]$ which are also normal distributions, some algebra reveals that XXXX COPY NOTATION FFROM CH. 6 XXXXX

$$[\alpha|y, \beta] = \text{Normal}(ybar, ...weightedvariancehere...).$$

Similarly,

$$[\beta|y, \alpha] \text{ is normal}(\dots\dots\dots)$$

The MCMC algorithm for this model has us simulate in succession, repeatedly, from those two distributions. See Gelman et al. (2004) for more examples of

Gibbs sampling for the normal model. A conceptual representation of the MCMC algorithm for this simple model is therefore: XXXX Check out ALGORITHM environment XXXXX

Algorithm

```

0. Initialize  $\alpha$  and  $\beta$ 

Repeat{
  1. Draw a new value of  $\alpha$  from Eq. \ref{xyz}
  2. Draw a new value of  $\beta$  from Eq. \ref{xyz}
}
```

As we just saw for this simple “normal-normal” model it is sometimes possible to specify the full conditional distributions analytically. In general, when certain so-called conjugate prior distributions are chosen, the form of full conditional distributions is similar to that of the observation model. In this normal-normal case, the normal distribution for the mean parameters is the conjugate prior under the normal model, and thus the full-conditional distributions are also normal. This is convenient because, in such cases, we can simulate directly from them using standard methods (or **R** functions). But, in practice, we don’t really ever need to know such things because most of the time we can get by using a simple algorithm, called the Metropolis-Hastings (henceforth “MH”) algorithm, to obtain samples from these full conditional distributions without having to recognize them as specific, named, distributions. This gives us enormous freedom in developing models and analyzing them without having to resolve them mathematically because to implement the MH algorithm we need only identify the full conditional distribution up to a constant of proportionality, that being the marginal distribution in the denominator (e.g., $[y|\beta]$ above).

We will talk about the Metropolis-Hastings algorithm shortly, and we will use it extensively in the analysis of SCR models (e.g., chapter 7).

2.5.1 Rules for constructing full conditional distributions

The basic strategy for constructing full-conditional distributions for devising MCMC algorithms can be reduced conceptually to a couple of basic steps summarized as follows:

- (step 1) Collect all stochastic components of the model;
- (step 2) Recognize and express the full conditional in question as proportional to the product of all components;
- (step 3) Remove the ones that don’t have the focal parameter in them.
- (step 4) Do some algebra on the result in order to identify the resulting pdf or pmf.

Of the 4 steps, the last of those is the main step that requires quite a bit of statistical experience and intuition because various algebraic tricks can be used to reshape the mess into something noticeable - i.e., a standard, named distribution. But step 4 is not necessary if we decide instead to use the Metropolis-Hastings algorithm as described below.

To illustrate for computing $[\alpha|y, \beta]$ we first apply step 1 and identify the model components as: $[y|\alpha, \beta]$, $[\alpha]$ and $[\beta]$. Step 2 has us write $[\alpha|y, \beta] \propto [y|\alpha, \beta][\alpha][\beta]$. Step 3: We note that $[\beta]$ is not a function of alpha and therefore we remove it to obtain $[\alpha|y, \beta] \propto [y|\alpha, \beta][\alpha]$. Similarly we obtain $[\beta|y, \alpha] \propto [y|\alpha, \beta][\beta]$. We apply step 4 and manipulate these algebraically to arrive at the result or, alternatively, we can sample them indirectly using the Metropolis-Hastings algorithm (see below).

2.5.2 Metropolis-Hastings algorithm

The Metropolis-Hastings (MH) algorithm is a completely generic method for sampling from any distribution, say $f(\theta)$. In our applications, $f(\theta)$ will typically be the full conditional distribution of θ . While we sometimes use Gibbs sampling, we seldom use “pure” Gibbs sampling because we might use MH to sample from one or more of the full conditional distributions. When the MH algorithm is used to sample from full conditional distributions of a Gibbs sampler the resulting hybrid algorithm is called *Metropolized Gibbs sampling* or more commonly *Metropolis-within-Gibbs*. Shortly we will actually construct such an algorithm for a simple class of models.

The MH algorithm generates candidates from some proposal or candidate-generating distribution, that may be conditional on the current value of the parameter, denoted by $h(\theta^*|\theta^t)$. Here, θ^* is the *candidate* or proposed value and θ^t is the current value, i.e., at iteration t of the MCMC algorithm. The proposed value is accepted with probability XXXX check notation with Rahel XXXXXX

$$r = \frac{f(\theta^*)h(\theta^t|\theta^*)}{f(\theta^t)h(\theta^*|\theta^t)}$$

which we call the MH acceptance probability. This ratio can sometimes be > 1 in which case we set it equal to 1. It is useful to note that $h()$ can be anything at all. Absolutely anything! You can generate candidate values from a *normal*(0,1) distribution, from a *uniform*(-3455,3455) distribution, or anything of proper support. Note, however, that good choices of $h()$ are those that approximate the posterior distribution. Obviously if $h() = f(\theta|y)$ (i.e., the posterior) then you always accept the draw, and it stands to reason that proposals that are more similar to $f(\theta|y)$ will lead to higher acceptance probabilities. No matter the choice of $h()$, we can evaluate this ratio numerically because the marginal $f(y)$ cancels from both the numerator and denominator, which is the magic of the MH algorithm.

A special kind of $h()$ are those that are symmetric, which means that $h(a|b) = h(b|a)$ in which case $h(a|b)$ and $h(b|a)$ just cancel out from the MH acceptance

probability and r is then just the ratio of the target density evaluated at the candidate value to that evaluated at the current value. A type of symmetric proposal useful in many situations is the so-called *random-walk* proposal distribution where candidate values are drawn from a normal distribution with mean equal to the current value and some standard deviation, say δ , which is prescribed by the user. For parameters that have support on the real line, e.g., α in our example above, the random walk proposal generator has us generate $\alpha^* \sim \text{Normal}(\alpha^t, \delta)$. If we set δ very small we have a high probability of accepting the proposal and vice versa. In practice, we “tune” delta to achieve a compromise between acceptance rate and efficient mixing of the Markov chains (see below for an example) normally assessed by autocorrelation. Low δ increases the acceptance rate but will tend to produce Markov chains with high autocorrelation, and vice versa.

Parameters with bounded support: Many models contain parameters that have bounded support. E.g., variance parameters live on $[0, \infty]$, parameters that represent probabilities live on $[0, 1]$, etc.. In that case it is sometimes convenient to use a random walk proposal distribution that can generate any real number (e.g., a normal random walk proposal). In that case, we can just reject parameters that are outside of the parameter space (XXXX REF FOR THIS XXXX).

2.6 PRACTICAL BAYESIAN ANALYSIS AND MCMC

There are a number of really important practical issues to be considered in any Bayesian analysis and we cover some of these briefly here.

2.6.1 Choice of prior distributions

**XXX integrate this material with previous section on prior distributions
XXXXXX**

Bayesian analysis requires that we choose prior distributions for all of the structural parameters of the model (we use the term structural parameter to mean all parameters that aren’t customary thought of as latent variables). We will strive to use priors that are meant to express little or no prior information - default or customary “non-informative” or diffuse priors. This will be $\text{Unif}(a, b)$ priors for parameters that have a natural bounded support and, for parameters that live on the real line we use either (1) diffuse normal priors; (2) “improper” uniform priors or (3) sometimes even a bounded $\text{Unif}(a, b)$ prior if that greatly improves the performance of **WinBUGS** or other software doing the MCMC for us. In **WinBUGS** a prior with low “precision”, τ , where $\tau = 1/\sigma^2$, such as $\text{Norm}(0, .01)$ will typically be used. Of course $\tau = 0.01$ ($\sigma^2 = 100$) might be very informative for a regression parameter that has a high variance. Therefore, we recommend that predictor variables *always* be standardized. Clearly there are a lot of choices for ostensibly non-informative priors, and the degree of non-informativeness depends on the parameterization. For example, a natural non-informative prior for the intercept of a

682 logistic regression

$$\text{logit}(p_i) = \alpha + \beta x_i$$

683 Would be $[\alpha] = \text{const}$ which is the same as saying $a \sim \text{Unif}(\infty, \text{infy})$, the custom-
 684 ary improper uniform prior. However, we might also use a prior on the parameter
 685 $p0 = \text{logit}^{-1}(a)$, which is $Pr(y = 1)$ for the value $x = 0$. Since $p0$ is a probability a
 686 natural choice is $p0 \sim \text{Unif}(0, 1)$. These two priors can affect results (see Chapter
 687 3.XYZ), yet they are both sensible non-informative priors. Choice of priors and
 688 parameterization is very much problem-specific and often largely subjective. More-
 689 over, it also affects the behavior of MCMC algorithms and therefore the analyst
 690 needs to pay some attention to this issue and possibly try different things out. XXX
 691 REFS on prior distributions XXXXXX

692 2.6.2 Convergence and so-forth

693 Once we have carried-out an analysis by MCMC, there are many other practical
 694 issues that we have to confront. One of the most important is “have the chains
 695 converged?” Most MCMC algorithms only guarantee that, eventually, the samples
 696 being generated will be from the target posterior distribution. So-called “conver-
 697 gence” of the Markov chain is achieved when that happens. Typically a period of
 698 transience is observed in the early part of the MCMC algorithm, and this is usually
 699 discarded as the “burn-in” period.

700 The quick diagnostic to whether convergence has been achieved is that your
 701 Markov chains look “grassy” – see Fig. 2.5 below. Another way to check conver-
 702 gence is to update the parameters some more and see if the posterior changes. It
 703 is good to confirm convergence using the “R-hat” statistic (\hat{R}) or Brooks-Gelman-
 704 Rubin statistic (Gelman et al., 1996)) which should be close to 1 if the Markov
 705 chains have converged and sufficient posterior samples have been obtained. In
 706 practice, $\hat{R} = 1.2$ is probably good enough for some problems. For some models
 707 you can’t actually realize a low \hat{R} . E.g., if the posterior is a discrete mixture of
 708 distributions then you can be misled into thinking that your Markov chains have
 709 not converged when in fact the chains are just jumping back and forth in the pos-
 710 terior state-space. So, for example, using model selection methods (section XYZ)
 711 sometimes suggests non-convergence. Another situation is when one of the param-
 712 eters is on the boundary of the parameter space which might appear to be very
 713 poor mixing, but all within some extreme region of the parameter space.⁴ This
 714 kind of stuff is normally ok and you need to think really hard about the context
 715 of the model and the problem before you conclude that your MCMC algorithm is
 716 ill-behaved.

717 Some models exhibit “poor mixing” of the Markov chains or what people might
 718 also say “have not covered” (or “slow convergence”) which is a term we would dis-

⁴it would be nice if we could compile examples of this later in the book and reference back to this point

agree with because the samples might well be from the posterior (i.e., the Markov chains have converged to the proper stationary distribution) but simply mix around the posterior rather slowly. Anyway, poor mixing can happen for a huge number of reasons – when parameters are highly correlated (even confounded), or barely identified from the data, or the algorithms are very terrible and probably many other reasons. Slow mixing equates to high autocorrelation in the Markov chain – the successive draws are highly correlated, and thus we need to run the MCMC algorithm much longer to get an effective sample size that is sufficient for estimation – or to reduce the MC error to a tolerable level. A strategy often used to reduce autocorrelation is “thinning” – i.e., keep every m^{th} value of the Markov chain output. However, thinning is necessarily inefficient from the stand point of inference – you can always get more precise posterior estimates by using all of the MCMC output regardless of the level of autocorrelation (MacEachern and Berliner, 1994). Practical considerations might necessitate thinning, even though it is statistically inefficient. For example, in models with many parameters or other unknowns being tabulated, the output files might be enormous and unwieldy to work with. In such cases, thinning is perfectly reasonable. In many cases, how well the Markov chains mix is strongly influenced by parameterization, standardization of covariates, and the prior distributions being used. Some things work better than others, and the investigator should experiment with different settings and remain calm when things don’t work out perfectly. MCMC is an art, and a science.

Is the posterior sample large enough? A good rule of thumb is that you should never report MCMC results to more than 2 decimal places – because they will always be different! Look at the MC error which is printed by default in summaries of BUGS output. You want that to be smallish relative to the magnitude of the parameter and this might depend on the purpose of the analysis. For a preliminary analysis you might settle for a few percent whereas for a final analysis then certainly less than 1% is called for, but you can run your MCMC algorithm as long as it takes. Note that MC error in summaries of the posterior is not the same as having an “approximate” solution in a standard likelihood analysis or similar. The approximate SE in likelihood inference is actually wrong in its actual value.... XYZ.

2.6.3 Bayesian confidence intervals

The 95% Bayesian interval based on percentiles of the posterior is not a unique interval – there are many of them – and the so-called “highest posterior density” (HPD) interval is the narrowest interval. We might compute that frequently because it is easy to do with an integer parameter which N is (See the next chapter). The 95 % HPD is not often exactly 95% but usually slightly more conservative than nominal because it is the narrowest interval that contains at least 95% of the posterior mass.

2.6.4 Estimating functions of parameters

A benefit of analysis by MCMC is that we can seamlessly estimate functions of parameters by simply tabulating the desired function of the simulated posterior draws. For example, if θ is the parameter of interest and let $\theta^{(i)}$ for $i = 1, 2, \dots, M$ be the posterior samples of θ . Let $\eta = \exp(\theta)$, then a posterior sample of η can be obtained simply by computing $\exp(\theta^{(i)})$ for $i = 1, 2, \dots, M$. We give another example in section 2.7.2 below and throughout this book. Almost all SCR models in this book involve at least 1 derived parameter. For example, density D is a derived parameter, being a function of population size N and the area A of the underlying state-space of the point process (see chapter 4).

2.7 BAYESIAN ANALYSIS USING WINBUGS

We won't be too concerned with devising our own MCMC algorithms for every analysis although we will do that a few times for fun. More often, we will rely on the freely available software package **WinBUGS** or **JAGS** for doing this. We will always execute these **BUGS** engines from within **R** using the **R2WinBUGS** (REF XYZ XYZ) or **rjags** packages. **WinBUGS** and **JAGS** are MCMC black boxes that takes a pseudo-code description (i.e., written in the **BUGS** language) of all of the relevant stochastic and deterministic elements of a model and generates an MCMC algorithm for that model. But you never get to see the algorithm. Instead, **WinBUGS/JAGS** will run the algorithm and just return the Markov chain output - the posterior samples of model parameters.

The great thing about using the **BUGS** language is that it forces you to become intimate with your statistical model - you have to write each element of the model down, admit (explicitly) all of the various assumptions, understand what the actual probability assumptions are and how data relate to latent variables and data and latent variables relate to parameters, and how parameters relate to one another.

While we normally use **WinBUGS** or **JAGS** in this book, we note that **OpenBUGS** is the current active development tree of the **BUGS** language. See Kéry (2010, ch.xyz) and Kéry and Schaub (2011, appendix xyz) for more on practical analysis in **WinBUGS**. That book should also be consulted for a more comprehensive introduction to using **WinBUGS**. In this example, we're going to accelerate pretty fast.

2.7.1 Linear Regression in WinBUGS

We provide a brief introductory example of a normal regression model using a small simulated data set. The following commands are executed from within your R workspace, the command line being indicated by '`>`'. First, simulate a covariate x and observations y having prescribed intercept, slope and variance:

```
> x<-rnorm(10)
```

```

795 > mu<- -3.2+ 1.5*x
796 > y<-rnorm(10,mu,sd=4)

```

797 The **BUGS** model specification for a normal regression model is written within **R**
 798 as a character string input to the command `cat()` and then dumped to a text file
 799 named `normal.txt`:

```

800 > cat("
801 model {
802   for (i in 1:10){
803     y[i]~dnorm(mu[i],tau)      # the "likelihood"
804     mu[i]<- beta0 + beta1*x[i] # the linear predictor
805   }
806   beta0~dnorm(0,.01)          # prior distributions
807   beta1~dnorm(0,.01)
808   sigma~dunif(0,100)
809   tau<-1/(sigma*sigma)       # tau is a derived parameter
810 }
811 ",file="normal.txt")

```

812 Alternatively, you can write the model specifications directly within a text file and
 813 save it in your current working directory, but we do not usually take that approach
 814 in this book.

815 **Remarks: 1. WinBUGS** parameterizes the normal in terms of the mean and
 816 inverse-variance, called the precision. Thus, `dnorm(0,.01)` implies a variance of
 817 100; **2.** We typically use diffuse normal priors for mean parameters, β_0 and β_1
 818 in this case, but sometimes we might use uniform priors with suitable bounds $-B$
 819 and $+B$. **3.** We typically use a $\text{Unif}(0, B)$ prior on standard deviation parameters
 820 (Gelman XXX 2006 XXXX). But sometimes we might use a gamma prior on the
 821 precision parameter τ . **4.** In a **WinBUGS** model file, every variable referenced
 822 in the model description has to be either data, which will be input (see below),
 823 a random variable which must have a probability distribution associated with it
 824 using the “~”, or it has to be a derived parameter connected to variables and data
 825 using “<-”.

826 To fit the model, we need to describe various data objects to **WinBUGS**.
 827 In particular, we create an **R** list object called `data` which are the data objects
 828 identified in the BUGS model file. In the example, the data consist of two objects
 829 which exist as y and x in the **R** workspace and also in the **WinBUGS** model
 830 definition. We also have to create an **R** function that produces a list of starting
 831 values `inits` that get sent to **WinBUGS**. Finally, we identify the names of the
 832 parameters (labeled correspondingly in the **WinBUGS** model specification) that
 833 we want **WinBUGS** to save the MCMC output for. In this example, we will
 834 “monitor” the parameters β_0 , β_1 , σ and τ . **WinBUGS** is executed using the **R**
 835 command `bugs()`. We set the option `debug=TRUE` if we want the **WinBUGS** GUI
 836 to stay open (useful for analyzing MCMC output and looking at the **WinBUGS**

error log). Also, we set `working.dir=getwd()` so that **WinBUGS** output files and the log file are saved in the current **R** working directory. All of these activities look like this:

```
library("R2WinBUGS")    # "attach" the R2WinBUGS library
data <- list ( "y","x")
inits <- function()
  list ( beta1=rnorm(1),beta0=rnorm(1),sigma=runif(1,0,2) )
parameters <- c("beta0","beta1","sigma","tau")
out<-bugs (data, inits, parameters, "normal.txt", n.thin=2, n.chains=2,
           n.burnin=2000, n.iter=6000, debug=TRUE,working.dir=getwd())
```

Remarks: A common question is “how should my data be formatted?” That depends on how you describe the model in the **BUGS** language, how your data are input into **R** and subsequently formatted. There is no unique way to describe any particular model and so you have some flexibility. We talk about data format further in the context of capture-recapture models and SCR models in chapter 4 and elsewhere. In general, starting values are optional but we recommend to always provide reasonable starting values for structural parameters, but are not always necessary for random effects. Note that the previously created objects defining data, initial values and parameters to monitor are passed to the function `bugs()`. In addition, various other things are declared: The number of Markov chains (`n.chains`), the thinning rate (`n.thin`), the number of burn-in iterations (`n.burnin`) and the total number of iterations (`n.iter`). To develop a detailed understanding of the various parameters and settings used for MCMC, consult a basic reference such as Kéry (2010).

You should execute all of the commands given above and then look at the resulting output. Kill the **WinBUGS** GUI and the data will be read back into **R** (or specify `debug=FALSE`). We don’t want to give instructions on how to navigate and use the GUI - see XYZ REF (XYZ) for that. The object `out` prints important summaries by default (this is slightly edited):

```
> print(out,digits=2)
Inference for Bugs model at "normal.txt", fit using WinBUGS,
2 chains, each with 6000 iterations (first 2000 discarded), n.thin = 2
n.sims = 4000 iterations saved
```

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
beta0	-2.43	1.84	-6.21	-3.50	-2.42	-1.34	1.27	1	4000
beta1	2.62	1.54	-0.42	1.68	2.62	3.57	5.67	1	4000
sigma	5.29	1.66	3.11	4.14	4.95	6.05	9.39	1	4000
tau	0.05	0.02	0.01	0.03	0.04	0.06	0.10	1	4000
deviance	59.85	3.24	56.18	57.47	59.00	61.37	68.32	1	840

For each parameter, `n.eff` is a crude measure of effective sample size, and `Rhat` is the potential scale reduction factor (at convergence, `Rhat=1`).

880 DIC info (using the rule, $pD = \bar{D} - \hat{D}$)
 881 $pD = 2.6$ and $DIC = 62.4$

882 **Remarks:** (1) convergence is assessed using the \hat{R} statistic – which we might
 883 sometimes write “*Rhat*”. A value of *Rhat* near 1 indicates convergence; (2) DIC
 884 is the “deviance information criterion” (Spiegelhalter et al., 2002) (see section 2.8)
 885 which some people use in a manner similar to AIC although it is recognized to
 886 have some problems in hierarchical models (Millar, 2009). We evaluate this in the
 887 context of SCR models in chapter XYZ XYZ.

888 2.7.2 Inference about functions of model parameters

889 Using the MCMC draws for a given model we can easily obtain the posterior distri-
 890 bution of any function of model parameters. We showed this in the above example
 891 by providing the posterior of τ when the model was parameterized in terms of stan-
 892 dard deviation σ . As another example, suppose that the normal regression model
 893 above had a quadratic response function of the form

$$E(y_i) = \beta_0 + \beta_1 x_i + \beta_2 x_i^2$$

894 Then the optimum value of x , i.e., that corresponding to the optimal expected
 895 response, can be found by setting the derivative of this function to 0 and solving
 896 for x . We find that

$$df/dx = \beta_1 + 2 * \beta_2 x = 0$$

897 yields that $x_{opt} = -\beta_1/(2 * \beta_2)$. We can just take our posterior draws for *beta*₁
 898 and *beta*₂ and obtain a posterior sample of x_{opt} by this simple calculation. As an
 899 exercise, take the normal model above and simulate a quadratic response and then
 900 describe the posterior distribution of x_{opt} .

2.8 MODEL CHECKING AND SELECTION

901 In general terms model checking - or assessing the adequacy of the model - and
 902 model selection are quite thorny issues and, despite contrary and, sometimes,
 903 strongly held belief among practitioners, there are not really definitive, general
 904 solutions to either problem. We’re against dogma on these issues and think people
 905 need to be open-minded about such things and recognize that models can be useful
 906 whether or not they pass certain statistical tests. Some models are intrinsically
 907 better than others because they make more biological sense or foster understanding
 908 or achieve some objective that some bootstrap or other goodness-of-fit test can’t
 909 decide for you. That said, it gives you some confidence if your model seems ade-
 910 quate and we try to provide some fit assessment in most real applications of SCR
 911 models We provide a very brief overview of concepts here, but provide more detailed
 912 coverage in chapter 8. See also Kéry (2010, ch. xyz) and Link and Barker (2009,
 913 ch. xyz) for specific context related to Bayesian model checking and selection.

2.8.1 Goodness-of-fit

Goodness-of-fit testing is an important element of any analysis because our model represents a general set of hypotheses about the ecological and observation processes that generated our data. Thus, if our model “fits” in some statistical or scientific sense, then we believe it to be consistent with the hypotheses that went into the model. More formally, we would conclude that the data are *not inconsistent* with the hypotheses, or that the model appears adequate. If we have enough data, then of course we will reject any set of statistical hypotheses. Conversely, we can always come up with a model that fits by making the model extremely complex. Despite this paradox, it seems to us that simple models that you can understand should usually be preferred even if they don’t fit, for example if they embody essential mechanisms central to our understanding of things, or if we think that some contributing factors to lack-of-fit are minor or irrelevant to the scientific context and intended use of the model. In other words, models can be useful irrespective of whether they fit according to some formal statistical test of fit. Yet the tension is there to obtain fitting models, and this comes naturally at the expense of models that can be easily interpreted and studied and effectively used. Moreover, conducting goodness-of-fit tests is not always so easy to do. Moreover, it is never really easy (or especially convenient) to decide if your goodness-of-fit test is worth anything. It might have 0 power! Despite this, we recommend attempting to assess model fit in real applications, as a general rule, and we provide some basic guidance here and some more specific to SCR models in chapter 8.

To evaluate goodness-of-fit in Bayesian analyses, we will most often use the Bayesian p-value (Gelman et al., 1996). The basic idea is to define a fit statistic or “discrepancy measure” and compare the posterior distribution of that statistic to the posterior predictive distribution of that statistic for hypothetical perfect data sets for which the model is known to be correct. For example, with count frequency data, a standard measure of fit is the sum of squares of the “Pearson residuals”,

$$D(y_i, \theta) = \frac{(y_i - E(y_i))^2}{Var(y_i)}$$

The fit statistic based on the squared residuals is

$$FIT = \sum_i D(y_i, \theta)^2$$

which can be computed at each iteration of a MCMC algorithm given the current values of parameters that determine the response distribution. At the same time (i.e., at each MCMC iteration), the equivalent statistic is computed for a “new” data set, simulated using the current parameter values. The Bayesian p-value is simply the posterior probability $\Pr(\text{Fit} > \text{Fit}_{\text{new}})$ ⁵ which should be close to 0.50 for

⁵Check this definition!

a good model – one that “fits” in the sense that the observed data set is consistent with realizations simulated under the model being fitted to the observed data. In practice we judge “close to 0.50” as being “not too close to 0 or 1” and, as always, closeness is somewhat subjective. We’re happy with anything $> .1$ and $< .9$ but might settle for $> .05$ and < 0.95 . In summary, the Bayesian p-value seems like a bootstrap idea, is easy to compute, and widely used as a result.

Another useful fit statistic is the Freeman-Tukey statistic⁶, in which

$$D(\mathbf{y}, \theta) = \sum_i (\sqrt{y_i} - \sqrt{e_i})^2$$

(Brooks et al., 2000), where y_i is the observed value of observation i and e_i its expected value. In contrast to a chi-square discrepancy, the Freeman-Tukey statistic removes the need to pool cells with small expected values.

2.8.2 Model Selection

For model selection we typically use three different methods: First is, let’s say, common sense. If a parameter has posterior mass concentrated away from 0 then it seems like it should be regarded as important - that is, it is “significant.” This approach seems to have fallen out of favor with all of the interest over the last 10 or 15 years on model selection in ecology. It seems reasonable to us.

For regression problems we sometimes use the factor weighting idea which is to introduce a set of binary variables w_k for variable k , and express the model as, e.g., for a single covariate model:

$$E(y_i) = \alpha + w\beta x_i$$

where w is given a Bernoulli prior distribution with some prescribed probability. E.g., $w \sim \text{Bern}(0.50)$ to provide a prior probability of 0.50 that variable x should be an element of the linear predictor. The posterior probability of the event $w = 1$ is a gauge of the importance of the variable x . i.e., high values of $\text{Pr}(w = 1)$ indicate stronger evidence to support that “ x is in the model” whereas values of $\text{Pr}(w = 1)$ close to 0 suggest that x is less important.

This idea seems to be due to Kuo and Mallick (1998)⁷ and see Royle and Dorazio (2008, ch. XXXX) for an example in the context of logistic regression. This approach seems to even work sometimes with fairly complex hierarchical models of a certain form. E.g., Royle (2008) applied it to a random effects model to evaluate the importance of the random effect component of the model. The main problem with this approach is that its effectiveness and results will typically be highly sensitive to the prior distribution on the structural parameters (e.g., see Royle and

⁶Ref for this?

⁷ Is this also what people call Zellner’s G-priors?

Dorazio (2008, table xyz)). The reason for this is obvious: If $w = 0$ for the current iteration of the MCMC algorithm, so that β is sampled from the prior distribution, and the prior distribution is very diffuse, then extreme values of β are likely. Consequently, when the current value of β is far away from the mass of the posterior when $w = 1$, then the Markov chain may only jump from $w = 0$ to $w = 1$ infrequently. One seemingly reasonable solution to this problem (Aitken XYZ FIND THIS XXXXX⁸) is to fit the full model to obtain posterior distributions for all parameters, and then use those as prior distributions in a “model selection” run of the MCMC algorithm. This seems preferable to more-or-less arbitrary restriction of the prior support to improve the performance of the MCMC algorithm.

A third method that that we advocate is subject-matter context. It seems that there are some situations – some models – where one should not have to do model selection because it is necessitated by the specific context of the problem, thus rendering a formal hypothesis test pointless (Johnson, 1999). SCR models are such an example. In SCR models, we will see that “spatial location” of individuals is an element of the model. The simpler, reduced, model is an ordinary capture-recapture model which is not spatially explicit (i.e., chapter 3), but it seems silly and pointless to think about actually using the reduced model even if we could concoct some statistical test to refute the more complex model. The simpler model is manifestly wrong but, more importantly, not even a plausible data-generating model! Other examples are when effort, area or sample rate is used as a covariate. One might prefer to have such things in models regardless of whether or not they pass some statistical litmus test (although one can always find referees to argue for pedantic procedure over thinking).

Many problems can be approached using one of these methods but there are also broad classes of problems that can’t and, for those, you’re on your own. In later chapters we will address model selection in specific contexts and we hope those will prove useful for a majority of the situations you encounter.

2.9 POISSON GLMS

The Poisson GLM (also known as “Poisson regression”) is probably the most relevant and important class of models in all of ecology. The basic model assumes observations $y_i; i = 1, 2, \dots, n$ follow a Poisson distribution with mean λ which we write

$$y_i \sim \text{Poisson}(\lambda)$$

Commonly y_i is a count of animals or plants at some point in space and λ might depend on i . For example, i might index point count locations in a forest, BBS route centers, or sample quadrats, or similar. If covariates are available it is typical to model them as linear effects on the log mean. If $x(i)$ is some measured

⁸see Royle 2008 paper for reference

1016 covariate associated with observation i . Then,

$$\log(x(i)) = \alpha + \beta * x(i)$$

1017 While we only specify the mean of the Poisson model directly, the Poisson model
 1018 (and all GLMs) has a “built-in” variance which is directly related to the mean. In
 1019 this case, $\text{Var}(y) = E(y) = \lambda$. Thus the model accommodates a linear increase in
 1020 variance with the mean.

1021 2.9.1 Important properties of the Poisson distribution

1022 There are two properties of the Poisson distribution that make it extremely useful
 1023 in ecology. First is the property of *compound additivity*. If y_1 and y_2 are Poisson
 1024 random variables with means λ_1 and λ_2 , then their sum $N = y_1 + y_2$ is Poisson
 1025 with mean $\lambda_1 + \lambda_2$. Thus, if the observations can be viewed as an aggregate of
 1026 counts over some finer unit of measurement, then the mean aggregates in a cor-
 1027 responding manner. Secondly, the Poisson distribution has a direct relationship
 1028 to the multinomial. If y_1 and y_2 are *iid* Poisson then, conditional on their sum
 1029 $N = y_1 + y_2$, their joint distribution is multinomial with sample size N and cell
 1030 probabilities $\lambda_1/(\lambda_1 + \lambda_2)$ and $\lambda_2/(\lambda_1 + \lambda_2)$. As a result of this, most multinomial
 1031 models can be analyzed as a Poisson GLM and *vice versa*.

1032 2.9.2 Example: Breeding Bird Survey Data

1033 As an example we consider a classical situation in ecology where counts of an
 1034 organism are made at a collection of spatial locations. In this particular example,
 1035 we have mourning dove counts made along North American Breeding Bird Survey
 1036 (BBS) routes in Pennsylvania, USA. A route consists of 50 stops separated by 0.5
 1037 mile. For the purposes here we are defining y_i = route total count and the sample
 1038 location will be marked by the center point of the BBS route. The survey is run
 1039 annually and the data set we have is 1966-1998. BBS data can be obtained online at
 1040 <http://www.pwrc.usgs.gov/bbs/>. We will make use of the whole data set shortly
 1041 but for now we’re going to focus on a specific year of counts – 1990 – for the sake of
 1042 building a simple model. For 1990 there were 77 active routes. We have the data
 1043 stored in a .csv file⁹ where rows index the unique route, column 1 is the route ID,
 1044 columns 2-3 are the route coordinates (longitude/latitude), column 4 is a habitat
 1045 covariate “forest cover” (standardized, see below) and the remaining columns are
 1046 the yearly counts. Years for which a route was not run are coded as “NA” in the
 1047 data matrix. We imagine that this will be a typical format for many ecological
 1048 studies, perhaps with more columns representing covariates. To read in the data
 1049 and display the first few elements of this matrix, do this:

```
1050 > a<-read.csv("pa-bbsdovedata-all.csv")
1051 > data[1:2,1:6]
```

⁹check this data format




Figure 2.1. Needs a caption

```

1052      X      lon      lat      habitat X66 X67
1053  1 72002 -80.445 41.501 -0.3871372  NA  24
1054  2 72003 -80.347 41.214 -1.0171629  NA  NA

```

1055 It is useful to display the spatial pattern in the observed counts. For that we
1056 use a spatial dot plot - where we plot the coordinates of the observations and mark
1057 the color of the plotting symbol based on the magnitude of the count. We have a
1058 special plotting function for that which is called `spatial.plot()` and it is available
1059 with the supplemental **R** package. Actually, what we want to do here is plot the
1060 log-count (+1 of course) which (Fig. 2.1) displays a notable pattern that could be
1061 related to something. The **R** commands for obtaining this figure are:

```

1062 data<-read.csv("pa-bbsdovedata-all.csv")
1063 y<-data[,29] # pick out 1990
1064 notna<-!is.na(y)
1065 y<-y[notna]
1066 spatial.plot(data[notna,2:3],y)

```

1067 We can ponder the potential effects that might lead to dove counts being
1068 high....corn fields, telephone wires, barn roofs along with misidentification of pi-
1069 geons, these could all correlated reasonably well with the observed count of mourn-
1070 ing doves. Unfortunately we don't have any of that information.

1071 We do have a measure of forest cover in the vicinity of each point which is
1072 contained in the data set (variable "habitat"). This was derived from a larger GIS




Figure 2.2. Needs a caption

coverage of the state (provided in the data file “pahabdata.csv”) which can be plotted using the `spatial.plot` function using the following commands

```
1075 > map('state',regions="penn",lwd=2)
1076 > spatial.plot(pahabdata[,2:3],pahabdata[, "dfor"],cx=2)
1077 > map('state',regions="penn",lwd=2,add=TRUE)
```

where the result appears in Fig. 2.2. We see a prominent pattern that indicates high forest coverage in the central part of the state and low forest cover in the SE. Inspecting the previous figure of log-counts suggests a relationship between counts and forest cover which is perhaps not surprising.

2.9.3 Doing it in WinBUGS

Here we demonstrate how to fit a Poisson GLM in **WinBUGS** using the covariate x_i = forest cover. It is advisable that x_i be standardized in most cases as this will improve mixing of the Markov chains. Recall that the data we have stored include a standardized covariate (forest cover) and so we don't have to worry about that here. To read the BBS data into **R** and get things set up for **WinBUGS** we issue the following commands:

```
1089 data<-read.csv("pa-bbsdovedata-all.csv")
1090 y<-data[,29] # pick out 1990
1091 notna<-!is.na(y)
```

```

1092 y<-y[notna]                # discard missing
1093 habitat<-data[notna,4]      # get habitat data
1094 library("R2WinBUGS")        # load R2WinBUGS
1095 data <- list ( "y","M","habitat") # bundle data for WinBUGS

```

Now we write out the Poisson model specification in **WinBUGS** pseudo-code, provide initial values, identify parameters to be monitored and then execute **WinBUGS**:

```

1099 cat("
1100 model {
1101     for (i in 1:M){
1102         y[i]~dpois(lam[i])
1103         log(lam[i])<- beta0+beta1*habitat[i]
1104     }
1105     beta0~dunif(-5,5)
1106     beta1~dunif(-5,5)
1107 }
1108 ",file="PoissonGLM.txt")
1109
1110 inits <- function() list ( beta0=rnorm(1),beta1=rnorm(1))
1111 parameters <- c("beta0","beta1")
1112 out<-bugs (data, inits, parameters, "PoissonGLM.txt", n.thin=2,n.chains=2,
1113           n.burnin=2000,n.iter=6000,debug=TRUE,working.dir=getwd())

```

Remarks: (1) Note the close correspondence in how the model is specified here compared with the normal regression model previously. As an exercise you should discuss the specific differences between the **BUGS** model specifications for the normal and Poisson models.

```

1118 > print(out,digits=3)
1119 Inference for Bugs model at
1120 ‘‘PoissonGLM.txt’’, fit using WinBUGS,
1121 2 chains, each with 4000 iterations (first 1000 discarded), n.thin = 2
1122 n.sims = 3000 iterations saved

```

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
beta0	3.151	0.025	3.102	3.135	3.151	3.168	3.199	1.001	2300
beta1	-0.498	0.021	-0.539	-0.512	-0.498	-0.484	-0.457	1.001	3000
fit	869.930	19.856	835.500	855.700	868.600	881.900	913.602	1.002	1600
fitnew	76.709	12.519	54.098	68.107	76.215	84.510	102.602	1.001	3000
deviance	1116.605	2.014	1115.000	1115.000	1116.000	1117.000	1122.000	1.001	3000

We might wonder whether this model provides an adequate fit to our data. To evaluate that, we used a Bayesian p-value analysis with fit statistic based on the Freeman-Tukey residual by replacing the model specification above with this:

```

1132 cat("
1133 model {

```

```

1134   for (i in 1:M){
1135     y[i]~dpois(lam[i])
1136     log(lam[i])<- beta0+beta1*habitat[i]
1137     d[i]<- pow(pow(y[i],0.5)-pow(lam[i],0.5),2)    #
1138
1139     ynew[i]~dpois(lam[i])
1140     dnew[i]<-pow( pow(ynew[i],0.5)-pow(lam[i],0.5),2)
1141
1142   }
1143   fit<-sum(d[])
1144   fitnew<-sum(dnew[])
1145   beta0~dunif(-5,5)
1146   beta1~dunif(-5,5)
1147 }
1148 ",file="PoissonGLM.txt")

```

1149 The Bayesian p-value is the proportion of times $fitnew > fit$ which, for this
 1150 data set, is 0, which was 1.0 in this case (calculation omitted). This suggests that
 1151 the basic Poisson model does not fit well.

1152 2.9.4 Constructing your own MCMC algorithm

1153 It might be helpful to suffer through a couple examples building custom MCMC
 1154 algorithms. Here, we develop an MCMC algorithm for the Poisson regression model,
 1155 using a Metropolis-within-Gibbs sampling framework.

1156 We will assume that the two parameters have diffuse normal priors, say $[\alpha] =$
 1157 $\text{Norm}(0, 100)$ and $[\beta] = \text{Norm}(0, 100)$ where each has *standard deviation* 100 (recall
 1158 that **WinBUGS** parameterizes the normal in terms of $1/\sigma^2$). We need to assemble
 1159 the relevant elements of the model which are these two prior distributions and the
 1160 likelihood $[y|\alpha, \beta] = \prod_i [y_i|\alpha, \beta]$ which is, mathematically, the product of the Poisson
 1161 pmf evaluated at each y_i , given particular values of α and β . Next, we need to
 1162 identify the full conditionals $[\alpha|\beta, \mathbf{y}]$ and $[\beta|\alpha, \mathbf{y}]$. We use the all-purpose rule for
 1163 constructing full conditionals (section 2.5.1) to discover that:

$$[\alpha|\beta, \mathbf{y}] \propto \left\{ \prod_i [y_i|\alpha, \beta] \right\} [\alpha]$$

1164 and

$$[\beta|\alpha, \mathbf{y}] \propto \left\{ \prod_i [y_i|\alpha, \beta] \right\} [\beta]$$

1165 Remember, we could replace the “ \propto ” with “ $=$ ” if we put $[y|\beta]$ or $[y|\alpha]$ in the de-
 1166 nominator. But, in general, $[y|\alpha]$ or $[y|\beta]$ will be quite a pain to compute and, more
 1167 importantly, it is a constant as far as the operative parameters (α or β , respectively)
 1168 are concerned. Therefore, the MH acceptance probability will be the ratio of the

ful-conditional evaluated at a candidate draw to that evaluated at the current draw, and so the denominator required to change \propto to $=$ winds up canceling from the MH acceptance probability. Here we will use the random walk candidate generator so that, for example, $\alpha^* \sim \text{Normal}(\alpha^t, \delta)$ where δ is the standard-deviation of the proposal distribution, which is just a tuning parameter¹⁰. We remark also that calculations are often done on the log-scale to preserve numerical integrity of things when quantities evaluate to small or large numbers, so keep in mind, for example, $a*b = \exp(\log(a) + \log(b))$. The “Metropolis within Gibbs” algorithm for a Poisson regression turns out to be remarkably simple:

```

1178 set.seed(2013)
1179
1180 out<-matrix(NA,nrow=1000,ncol=2)  # matrix to store the output
1181 alpha<- -1                        # starting values
1182 beta <- -.8
1183
1184 # begin the MCMC loop ; do 1000 iterations
1185 for(i in 1:1000){
1186
1187   # update the alpha parameter
1188   lambda<- exp(alpha+beta*habitat)
1189   lik.curr<- sum(log(dpois(y,lambda)))
1190   prior.curr<- log(dnorm(alpha,0,100))
1191   alpha.cand<-rnorm(1,alpha,.25)      # generate candidate
1192   lambda.cand<- exp(alpha.cand + beta*habitat)
1193   lik.cand<- sum(log(dpois(y,lambda.cand)))
1194   prior.cand<- log(dnorm(alpha.cand,0,100))
1195   mhratio<- exp(lik.cand +prior.cand - lik.curr-prior.curr)
1196   if(runif(1)< mhratio)
1197     alpha<-alpha.cand
1198
1199   # update the beta parameter
1200   lik.curr<- sum(log(dpois(y,exp(alpha+beta*habitat))))
1201   prior.curr<- log(dnorm(beta,0,100))
1202   beta.cand<-rnorm(1,beta,.25)
1203   lambda.cand<- exp(alpha+beta.cand*habitat)
1204   lik.cand<- sum(log(dpois(y,lambda.cand)))
1205   prior.cand<- log(dnorm(beta.cand,0,100))
1206   mhratio<- exp(lik.cand + prior.cand - lik.curr - prior.curr)
1207   if(runif(1)< mhratio)
1208     beta<-beta.cand
1209
1210   out[i,]<-c(alpha,beta)              # save the current values
1211 }

```

¹⁰ It would help lots of people out to see a non-symmetric proposal distribution, and the extra step needed to account for it.




Figure 2.3. MCMC output for Poisson regression parameters (top trace: intercept α ; bottom trace: slope β). This is for $\delta = 0.25$.

```

1212
1213
1214 plot(out[,1],ylim=c(-1.5,3.3),type="l",lwd=2,ylab="parameter value",
1215       xlab="MCMC iteration")
1216 lines(out[,2],lwd=2,col="red")

```

1217 The first 300 iterations of the MCMC history of each parameter is shown in
 1218 Fig. 2.3. The appearance of this is not very appealing but a couple of things are
 1219 evident: First, the Markov chains clearly stabilize - “burn-in” - after about 60 or
 1220 70 iterations. They also appear to mix very slowly once convergence is achieved,
 1221 although this is not so clear given the scale of the y -axis. We decreased the standard
 1222 deviation of the candidate generating distribution from $\delta = 0.25$ to $\delta = 0.05$ and
 1223 re-ran re-ran the MCMC algorithm producing the output shown in Fig. 2.4. We see
 1224 that the burn-in takes longer but it seems to mix better although it takes slightly
 1225 longer to burn-in. Using this value of δ we generated 10,000 posterior samples,
 1226 discarding the first 500 as burn-in, and the result is shown in Fig. 2.5, this time
 1227 separate panels for each parameter. The “grassy” look of the MCMC history is
 1228 diagnostic of Markov chains that are well-mixing and we would generally be very
 1229 satisfied with results that look like this.

1230 **Remarks:** We used a specific set of starting values for these simulations. It

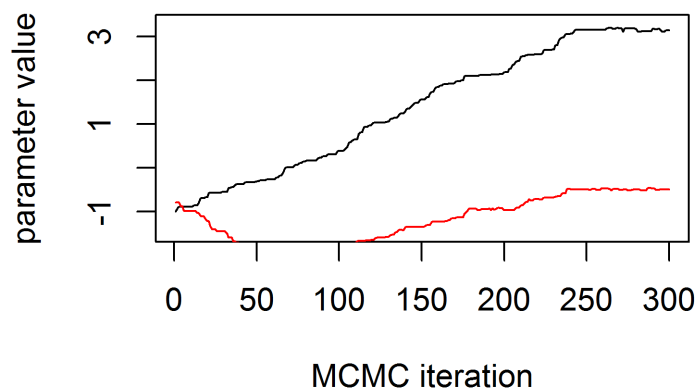


Figure 2.4. Same as previous fig but with $\delta = 0.05$.

1231 should be clear that starting values closer to the mass of the posterior distribution
 1232 might cause burn-in to occur faster. As an exercise, evaluate that. (2) Clearly
 1233 the influence of the proposal standard deviation term is important. Small values
 1234 lead to much better mixing but it should be noted that values that are too small
 1235 will slow down burn-in and also lead to high correlation. This suggests there is an
 1236 optimal value of the Metropolis-Hastings tuning parameter¹¹. As an exercise you
 1237 should contemplate finding that optimal value for this problem¹² (3) For the flat
 1238 normal prior distributions here we could leave the prior contribution out of the full
 1239 conditional evaluation since it is locally constant, i.e., constant in the vicinity of the
 1240 posterior mass, and thus has no practical effect. Removing the prior contribution
 1241 from the MH acceptance probability is equivalent to saying that the parameters
 1242 have an improper uniform prior, i.e., $\alpha \sim \text{const}$, which is commonly used for mean
 1243 parameters in practice. Note also that we have used a different prior than in our
 1244 **WinBUGS** model specification given previously. As an exercise, evaluate whether
 1245 this seems to affect the result.

¹¹Defined previously?????

¹²effective sample size definition?

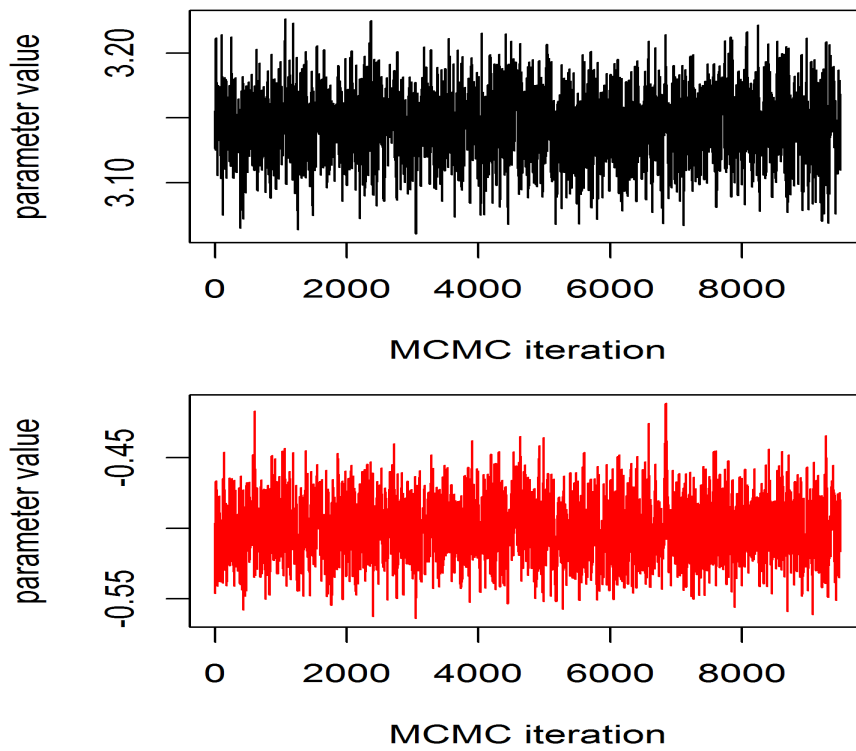


Figure 2.5. nice grassy mcmc output, longer run of previous with $\delta = 0.05$.

2.10 POISSON GLM WITH RANDOM EFFECTS

What we will be doing in most of this book is dealing with random effects in GLM-like models - similar to what are usually referred to as generalized linear mixed models (GLMMs). We provide a brief introduction by way of example, extending our Poisson regression model to include a random effect.

ANDY STOPPED HERE

The Log-Normal mixture: The classical situation involves a GLM with a normally distributed random effect that is additive on the linear predictor. For the Poisson case, we have:

$$\log(\lambda_i) = \alpha + \beta x_i + \eta_i$$

where $\eta_i \sim \text{Normal}(0, \sigma^2)$. A natural alternative is to have multiplicative gamma-distributed noise, $\exp(\eta_i) \sim \text{Gamma}(a, b)$ which would correspond to a negative binomial kind of over-dispersion, implying a different mean/variance relationship to the log-normal mixture (the interested reader should work that out). Choosing between such possibilities is not a topic we will get into here because it doesn't seem possible to provide general guidance on it. For this model we carried-out a goodness-of-fit evaluation using the Bayesian p-value based on a Pearson residual statistic. See also (Kéry, 2010, ch. 18) for an example involving a binomial mixed model¹³. Anyhow, it is really amazingly simple to express this model in **WinBUGS** and have **WinBUGS** draw samples from the posterior distribution using the following code for the BBS dove counts:

```
data<-read.csv("pa-bbsdovedata-all.csv")
locs<-data[,2:3]
habitat<-data[,4]
y<-data[,29]      # grab year 1990
M<-length(y)

set.seed(2013)

cat("
model {
  for (i in 1:M){
    y[i]~dpois(lam[i])
    log(lam[i])<- alpha+ beta*habitat[i] + eta[i]
    frog[i]<-beta*habitat[i] + eta[i]
    eta[i] ~ dnorm(0,tau)
    d[i]<- pow(pow(y[i],0.5)-pow(lam[i],0.5),2)

    ynew[i]~dpois(lam[i])
    dnew[i]<- pow(pow(ynew[i],0.5)-pow(lam[i],0.5),2)
  }
  fit<-sum(d[])
```

¹³Kéry has noticed that such tests probably have 0 power. Should use the marginal frequency of the data

```

1286   fitnew<-sum(dnew[])
1287
1288   alpha~dunif(-5,5)
1289   beta~dunif(-5,5)
1290   sigma~dunif(0,10)
1291   tau<-1/(sigma*sigma)
1292 }
1293
1294 ",file="model.txt")
1295 data <- list ( "y","M","habitat")
1296 inits <- function()
1297   list ( alpha=rnorm(1),beta=rnorm(1),sigma=runif(1,0,4))
1298 parameters <- c("alpha","beta","sigma","tau","fit","fitnew")
1299 library("R2WinBUGS")
1300
1301 out<-bugs (data, inits, parameters, "model.txt", n.thin=2,n.chains=2,
1302   n.burnin=1000,n.iter=5000,debug=TRUE)

```

1303 This produces the following posterior summary statistics:

```

1304 > print(out,digits=2)
1305 Inference for Bugs model at "model.txt", fit using WinBUGS,
1306 2 chains, each with 5000 iterations (first 1000 discarded), n.thin = 2
1307 n.sims = 4000 iterations saved
1308      mean    sd  2.5%   25%   50%   75%  97.5% Rhat n.eff
1309 alpha    2.98 0.08  2.82  2.93  2.98  3.03  3.12 1.00  1400
1310 beta   -0.53 0.07 -0.68 -0.58 -0.53 -0.49 -0.38 1.01   350
1311 sigma    0.60 0.06  0.49  0.56  0.59  0.64  0.73 1.00  2000
1312 tau     2.88 0.57  1.88  2.47  2.86  3.24  4.12 1.00  2000
1313 fit     26.58 3.72 19.87 23.96 26.37 29.01 34.46 1.00  4000
1314 fitnew   26.83 3.90 19.60 24.12 26.68 29.36 35.04 1.00  4000
1315 deviance 445.94 12.18 424.00 437.40 445.20 453.90 471.50 1.00  4000
1316
1317 [... some output deleted ...]

```

1318 The Bayesian p-value for this model is

```

1319 > mean(out$sims.list$fit>out$sims.list$fitnew)
1320 [1] 0.4815

```

1321 indicating a pretty good fit. Given the site-level random effect, it would be surpris-
 1322 ing for this model to not fit! One thing we notice is that the posterior standard
 1323 deviations of the regression parameters are much higher, a result of the excess vari-
 1324 ation. Wwe would also notice much less precise predictions of hypothetical new
 1325 observations.

1326 ANDY STOPPED HERE.

2.11 BINOMIAL GLMS

Another extremely important class of models in ecology are binomial models. We use binomial models for count data whenever the observations are counts or frequencies and it is natural to condition on a “sample size”, say K , the maximum frequency possible in a sample. The random variable, $y \leq K$, is then the frequency of occurrences out of K “trials”. The parameter of the binomial models is p , often called “success probability” which is related to the expected value of y by $E(y) = pK$. Usually we are interested in modeling covariates that affect the parameter p , and such models are called binomial GLMs, binomial regression models or logistic regression, although logistic regression really only applies when the logistic link is used to model the relationship between p and covariates (see below).

One of the most typical binomial GLMs occurs when the sample size equals 1 and the outcome, y , is “presence” ($y = 1$) or “absence” ($y = 0$) of a species. This is a classical “species distribution” modeling situation. A special situation occurs when presence/absence is observed with error (MacKenzie et al., 2002; Tyre et al., 2003). In that case, $K > 1$ samples are usually needed for effective estimation of model parameters.

In standard binomial regression problems the sample size is fixed by design but interesting models also arise when the sample size is itself a random variable. These are the N -mixture models (Royle, 2004; Kéry et al., 2005; Royle and Dorazio, 2008; Kéry, 2010) and related models (in this case, N being the sample size, which we labeled K above)¹⁴. Another situation in which the binomial sample size is “fixed” is closed population capture-recapture models in which a population of individuals is sampled K times. The number of times each individual is encountered is a binomial outcome with parameter - encounter probability - p , based on a sample of size K . In addition, the total number of unique individuals observed, n , is also a binomial random variable based on population size N . We consider such models in the chapter 3.

2.11.1 Binomial regression

In binomial models, covariates are modeled on a suitable transformation (the link function) of the binomial success probability, p . Let x_i denote some measured covariate for sample unit i and let p_i be the success probability for unit i . The standard choice is the “logit” link function which is:

$$\log(p_i/(1 - p_i)) = \alpha + \beta * x_i.$$

¹⁴Some of the jargon is actually a little bit confusing here because the binomial index is customarily referred to as “sample size” but in the context of N -mixture models N is actually the “population size”

1359 The inverse-logit (or “expit”) is

$$p_i = \text{expit}(\alpha + \beta * x_i) = \frac{\exp(\alpha + \beta * x_i)}{1 + \exp(\alpha + \beta * x_i)}$$

1360 There are many other possible link functions. However, ecologists seem to adopt
 1361 the logit link function without question in most applications¹⁵. We sometimes use
 1362 the “complementary log-log” (= “cloglog”) link function in ecological applications
 1363 because it arises naturally in many situations (Royle and Dorazio, 2008, p. 150).
 1364 For example, consider the “probability of observing a count greater than 0” under
 1365 a Poisson model: $\Pr(y > 0) = 1 - \exp(-\lambda)$. In that case,

$$\text{cloglog}(p) = \log(-\log(1 - p)) = \log(\lambda)$$

1366 So that if you have covariates in your linear predictor for $E(y)$ under a Poisson
 1367 model then they are linear on the complementary log-log link of p . In models of
 1368 species occurrence it seems natural to view occupancy as being derived from local
 1369 abundance N (Royle and Nichols, 2003; Royle and Dorazio, 2006; Dorazio, 2007).
 1370 Therefore, models of local abundance in which $N \sim \text{Poisson}(A\lambda)$ for a habitat patch
 1371 of area A implies a model for occupancy ψ of the form

$$\text{cloglog}(\psi) = \log(A) + \log(\lambda).$$

1372 We will use the cloglog link in some analyses of SCR models in chapter 4 and
 1373 elsewhere.

1374 2.11.2 Example: Waterfowl Banding Data

1375 It would be easy to consider a standard “distribution modeling” application where
 1376 $K = 1$ and the outcome is occurrence ($y = 1$) or not ($y = 0$) of some species.
 1377 Such examples abound in books (e.g., Royle and Dorazio (2008, ch. 3); Kéry (2010,
 1378 ch. 21); Kéry and Schaub (2011, ch. 13)) and in the literature. Instead, we will
 1379 consider an example involving band returns of waterfowl which were analyzed by
 1380 Royle and Dubovsky (2001)¹⁶.

1381 For these data, y_i is the number of waterfowl bands recovered out of B_i birds
 1382 banded at some location \mathbf{s}_i . In this case B_i is fixed. Thinking about recovery rate
 1383 as being proportional to harvest rate, we use these data to explore geographic gra-
 1384 dients in recovery rate resulting from variability in harvest pressure experienced by
 1385 populations depending on their migration ecology. As such, we fit a basic binomial
 1386 GLM with a linear response to geographic coordinates (including an interaction
 1387 term). The data are provided with the **R** package **scrbook**. Here we provide the
 1388 part of the script for creating the model and fitting the model in **WinBUGS** using

¹⁵a notable exception is distance sampling, which is all about choosing among link functions

¹⁶I hate this example. Anyone got a better one thats not distribution modeling?

the `bugs` function. There are few structural differences between this model and the Poisson GLM fitted previously. The main things are due to the data structure (we have a matrix here instead of a vector) and otherwise we change the main distributional assumption to binomial (specified with `dbin`) and then use the `logit` function to relate the parameter p_{it} to the covariates. Here is the script:

```

1394 load("mallarddata") # not sure how this will look
1395
1396 sink("model.txt")
1397 cat("
1398 model {
1399   for(t in 1:5){
1400     for (i in 1:nobs){
1401       y[i,t] ~ dbin(p[i,t], B[i,t])
1402       logit(p[i,t]) <- alpha0[t] + alpha1*X[i,1] + alpha2*X[i,2] + alpha3*X[i,1]*X[i,2]
1403     }
1404   }
1405   alpha1~dnorm(0,.001)
1406   alpha2~dnorm(0,.001)
1407   alpha3~dnorm(0,.001)
1408   for(t in 1:5){
1409     alpha0[t] ~ dnorm(0,.001)
1410   }
1411 }
1412 ",fill=TRUE)
1413 sink()
1414
1415 data <- list(B=mallard.bandings, y=mallard.recoveries,
1416             nobs=nrow(banding.locs),X=banding.locs)
1417 inits <- function(){
1418   list(alpha0=rnorm(5),alpha1=0,alpha2=0,alpha3=0) }
1419 parms <- list('alpha0','alpha1','alpha2','alpha3')
1420 out <- bugs(data,inits, parms,"model.txt",n.chains=3,
1421            n.iter=2000,n.burnin=1000, n.thin=2,debug=TRUE)

```

Posterior summaries of model parameters are as follows:

```

1423 > print(out,digits=3)
1424 Inference for Bugs model at "model.txt", fit using WinBUGS,
1425 3 chains, each with 2000 iterations (first 1000 discarded), n.thin = 2
1426 n.sims = 1500 iterations saved

```

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
alpha0[1]	-2.346	0.036	-2.417	-2.370	-2.346	-2.323	-2.277	1.001	1500
alpha0[2]	-2.356	0.032	-2.420	-2.379	-2.356	-2.335	-2.292	1.001	1500
alpha0[3]	-2.220	0.035	-2.291	-2.244	-2.219	-2.197	-2.153	1.001	1500
alpha0[4]	-2.144	0.039	-2.225	-2.169	-2.143	-2.116	-2.068	1.000	1500
alpha0[5]	-1.925	0.034	-1.990	-1.949	-1.924	-1.901	-1.856	1.004	570




Figure 2.6. Predicted recovery rate of bands.

```
1433 alpha1      -0.023 0.003   -0.028  -0.025  -0.023  -0.022  -0.018  1.001  1500
1434 alpha2       0.020 0.006    0.009   0.016   0.020   0.024   0.031  1.001  1500
1435 alpha3       0.000 0.001   -0.002  -0.001   0.000   0.000   0.002  1.001  1500
1436 deviance  1716.001 4.091 1710.000 1713.000 1715.000 1718.000 1726.000 1.001  1500
1437
1438 [... some output deleted ...]
```

1439 The basic result suggests a negative east-west gradient and a positive south to
1440 north gradient but no interaction. A map of the response surface is shown in Fig.
1441 2.6. We did an additional MCMC run where we saved the binomial parameter
1442 p and computed the Bayesian p-value (double use of “p” here is confusing, but I
1443 guess that happens sometimes!) using a fit statistic based on the Freeman-Tukey
1444 statistic (see Section XXX above). The result indicates that the linear response
1445 surface model does not provide an adequate fit of the data. The reader should
1446 contemplate whether this invalidates the basic interpretation of the result.

2.12 SUMMARY AND OUTLOOK

1447 GLMs and GLMMs are the most useful statistical methods in all of ecology. The
1448 principles and procedures underlying these methods are relevant to nearly all mod-
1449 eling and analysis problems in every branch of ecology. Moreover, understanding
1450 how to analyze these models is crucial in a huge number of diverse problems. If you

1451 understand and can conduct classical likelihood and Bayesian analysis of Poisson
1452 and binomial GLM(M)s, then you will be successful analyzing and understanding
1453 more complex classes of models that arise. We will see shortly that spatial capture-
1454 recapture models are a type of GLMM and thus having a basic understanding of
1455 the conceptual origins and formulation of GLM(M)s and their analysis is extremely
1456 useful.

1457 We note that GLM(M)s are routinely analyzed by likelihood methods but we
1458 have focused on Bayesian analysis here in order to develop the tools that are less
1459 familiar to most ecologists. In particular, Bayesian analysis of models with ran-
1460 dom effects is relatively straightforward because the models are easy to analyze
1461 conditional on the random effect, using methods of MCMC. Thus, we will often
1462 analyze SCR models in later chapters by MCMC, explicitly adopting a Bayesian
1463 inference framework. In that regard, the various **BUGS** engines (**WinBUGS**,
1464 **OpenBUGS**, **JAGS**) are enormously useful because they provide an accessible
1465 platform for carrying out analyses by MCMC by just describing the model, and not
1466 having to worry about how to actually build MCMC algorithms. That said, the
1467 **BUGS** language is more important than just to the extent that it enables one to
1468 do MCMC - it is useful as a modeling tool because it fosters understanding, in the
1469 sense that it forces you to become intimate with your model. You have to write
1470 down all of the probability assumptions, the relationships between observations and
1471 latent variables and parameters. This is really a great learning paradigm that you
1472 can grow with.

1473 While we have emphasized Bayesian analysis in this chapter, and make primary
1474 use of it through the book, we will provide an introduction to likelihood analysis
1475 in chapter 6 and use those methods also from time to time. Before getting to that,
1476 however, it will be useful to talk about more basic, conventional closed population
1477 capture-recapture models and these are the topic of the next chapter.

3

ESTIMATING THE SIZE OF A CLOSED POPULATION

In this chapter we will consider ordinary capture-recapture (CR) models for estimating population size in closed populations. We will see that such models are closely related to binomial (or logistic) regression type models. In fact, when N is known, they are precisely such models. We consider some important extensions of ordinary closed population models that accommodate various types of “individual effects” — either in the form of explicit covariates (sex, age, body mass) or unstructured “heterogeneity” in the form of an individual random effect. In general, these models are variations of generalized linear or generalized linear mixed models (GLMMs). Because of the paramount importance of this concept, we focus mainly on fairly simple models in which the observations are individual encounter frequencies, y_i = the number of encounters of individual i out of K replicate samples of the population which, for the models we consider here, is the outcome of a binomial random variable. Along the way, we consider the spatial context of capture-recapture data and models and demonstrate that density cannot be formally estimated when spatial information is ignored. We also review some of the informal methods of estimating density using CR methods, and consider some of their limitations. We will be exposed to our first primitive spatial capture-recapture models which arise as relatively minor variations of so-called “individual covariate models” (of the Huggins (1989) and Alho (1990) variety). In a sense, the point of this chapter is to establish that linkage in a direct and concise manner beginning with the basic “Model M0” and extensions of that model to include individual heterogeneity and also individual covariates. A special type of individual covariate models is distance sampling, which could be thought of as the most primitive spatial capture-recapture model. In later chapters we further develop and extend

ideas introduced in this chapter.

We emphasize Bayesian analysis of capture-recapture models and we accomplish this using a method related to classical “data augmentation” from the statistics literature Tanner and Wong (e.g., 1987)). This is a general concept in statistics but, in the context of capture-recapture models where N is unknown, it has a consistent implementation across classes of capture-recapture models and one that is really convenient from the standpoint of doing MCMC (Royle et al., 2007). We use data augmentation throughout this book and thus emphasize its conceptual and technical origins and demonstrate applications to closed population models. We refer the reader to Kery and Schaub (2011, ch. 6) for an accessible and complimentary development of ordinary closed population models.

3.1 THE SIMPLEST CLOSED POPULATION MODEL: MODEL M0

We suppose that there exists a population of N individuals which we subject to repeated sampling, say over K nights, where individuals are captured, marked, and subsequently recaptured. We suppose that individual encounter histories are obtained, and these are of the form of a sequence of 0’s and 1’s indicating capture ($y = 1$) or not ($y = 0$) during any sampling occasion (“sample”). As an example, suppose $K = 5$ sampling occasions, then an individual captured during sample 2 and 3 but not otherwise would have an encounter history of the form $\mathbf{y} = (0, 1, 1, 0, 0)$. Thus, the observation \mathbf{y}_i for each individual (i) is a vector having elements denoted by y_{ik} for $k = 1, 2, \dots, K$. Usually this is organized as a row of a matrix with elements y_{ik} , see Table 3.1. Except where noted explicitly, we suppose that observations are independent within individuals and among individuals. Formally, this allows us to say that y_{ik} are Bernoulli random variables and we may write $y_{ik} \sim \text{Bern}(p)$. Consequently, for this very simple model in which p is in fact constant, then we can declare that the individual encounter frequencies (total captures), $y_i = \sum_k y_{ik}$, have a binomial distribution based on a sample of size K . That is

$$y_i = \sum_k y_{ik} \sim \text{Bin}(p, K)$$

for every individual in the population. This is a remarkably simple model that forms the cornerstone of almost all of classical capture-recapture models, including most spatial capture-recapture models discussed throughout this book. Evidently, the basic capture-recapture model structure is precisely a simplistic version of a logistic-regression model with only an intercept term ($\text{logit}(p) = \text{constant}$). To say that all capture-recapture models are just logistic regressions is only slightly inaccurate. In fact, we are proceeding here “conditional on N ”, i.e., as if we knew N . In practice we don’t, of course, and that is kind of the point of capture-recapture models as estimating N is the central objective. But, by proceeding conditional on N , we can specify a simple model and then deal with the fact that N is unknown

Table 3.1. a capture-recapture data set with $n = 6$ observed individuals and $K = 5$ samples.

indiv i	Sample occasion					y_i
	1	2	3	4	5	
1	1	0	0	1	0	2
2	0	1	0	0	1	2
3	1	0	0	1	0	2
4	1	0	1	0	1	3
5	0	1	0	0	0	1
$n = 6$	1	0	0	0	0	1

using standard methods that you are already familiar with (i.e., GLMs - see chapter 2).

Assuming individuals of the population are observed independently, the joint probability distribution of the observations is the product of N binomials

$$\begin{aligned} \Pr(y_1, \dots, y_N | p) &= \prod_{i=1}^N \text{Bin}(y_i | K, p) \\ &= \prod_{k=0}^K \pi(k)^{n_k} \end{aligned}$$

where $\pi(k) = \text{Bin}(k | K, p)$ and where $n_k = \sum_{i=1}^N I(y_i = k)$ denotes the number of individuals captured k times in K surveys. We emphasize that this is conditional on N , in which case we get to observe the $y = 0$ observations and the resulting data are just *iid* binomial counts. Because this is a binomial regression model of the variety described in chapter 2, fitting this model using a BUGS engine poses no difficulty.

The essential problem in capture-recapture, however, is that N is not known because the number of uncaptured/missing individuals (i.e., those in the zero cell that occur with probability $\pi(0)$) is unknown. Consequently, the observed capture frequencies n_k are no longer independent. Instead, their joint distribution is multinomial (e.g., see Illian (2008a, p. xyz)):

$$n_1, n_2, \dots, n_K \sim \text{Multin}(N, \pi(1), \pi(2), \dots, \pi(K)) \quad (3.1.1)$$

Note that in our notation the number of uncaptured/missing individuals is denoted by $n_0 = N - n$, where $n = \sum_{k=1}^K n_k$ denotes the total number of distinct individuals seen in the K samples.

To fit the model in which N is *unknown*, we can regard N as a parameter and maximize the multinomial likelihood directly. While direct likelihood analysis of the multinomial model is straightforward, that does not prove to be too useful in practice because we seldom are concerned with models for the aggregated encounter history frequencies. In many instances, including for spatial capture-recapture (SCR)

models, we require a formulation of the model that can accommodate individual level covariates which we address subsequently in this chapter.

3.1.1 The Spatial Context of Capture-Recapture

A common assumption made is that of population “closure” which is really just a colloquial way of saying (in part) the Bernoulli assumptions stated explicitly above. In the biological context, closure means, strictly, no additions or subtractions from the population during study. This is manifest by the statement that the encounters are independent and identically distributed (iid) Bernoulli trials. In practice, closure is usually interpreted by the manner in which potential violations of that assumption arise. In particular, two important elements of the closure assumption are “demographic” and “geographic” closure. If an individual dies then subsequent values of y_{ik} are clearly no longer Bernoulli trials with the same parameter p . If there is no mortality or recruitment in the population, then we say that demographic closure is satisfied. Similarly, animals may emigrate or immigrate. If they do not, then geographic closure is satisfied. Sometimes a distinction is made between temporary and permanent emigration or immigration. That is a relevant distinction in spatial capture-recapture models, because SCR models explicitly accommodate “temporary emigration” of a certain type, due to individuals moving about their home range. The demographic closure assumption can also be relaxed using SCR models, but we will save that discussion for chapter 4.

3.1.2 Conditional likelihood

We saw that a basic closed population model is a simple logistic regression model if N is known and, when N is unknown, the model is multinomial with index or sample size parameter N . This multinomial model, being conditional on N , is sometimes referred to as the “joint likelihood” the “full likelihood” or the “unconditional likelihood” (or model in place of likelihood). This formulation differs from the so-called “conditional likelihood” approach in which the likelihood of the observed encounter histories is devised conditional on the event that an individual is captured at least once. To construct this likelihood, we have to recognize that individuals appear or not in the sample based on the value of the random variable y_i , that is, we capture them if and only if $y_i > 0$. The observation model is therefore based on $\Pr(y|y > 0)$. For the simple case of Model M0, the resulting conditional distribution is a “zero truncated” binomial distribution which accounts for the fact that we cannot observe the value $y = 0$ in the data set (see Royle and Dorazio, 2008, section XYZ). Both the conditional or unconditional models are legitimate modes of analysis in all capture-recapture types of studies, and they provide equally valid descriptions of the data and for many practical purposes provide equivalent inferences, at least in large sample sizes (Sanathanan, 1972).

Mode of analysis	parameters in model	statistical model
Joint likelihood	p, N	multinomial with index N
Conditional likelihood	p	zero-truncated binomial
Data augmentation	p, ψ	zero-inflated binomial

Table 3.2. Modes of analysis of capture-recapture models.

In this book we emphasize Bayesian analysis of capture-recapture models using data augmentation (discussed subsequently), which produces yet a third distinct formulation of capture-recapture-models based on the zero-*inflated* binomial distribution that we describe in the next section. Thus, there are 3 distinct formulations of the model – or models of analysis – for analyzing all capture-recapture models based on the (1) binomial model for the joint or unconditional specification; (2) zero-truncated binomial that arises “conditional on n ”; and (3) the zero-inflated binomial that arises under data augmentation. Each formulation has a distinct complement of model parameters (shown in Table 3.2 for Model M0).

3.2 DATA AUGMENTATION

We consider a method of analyzing closed population models using data augmentation (DA) which is useful for Bayesian analysis and, in particular, analysis of models using the various BUGS engines and other software. Data augmentation is a general statistical concept that is widely used in statistics in many different settings. The classical reference is Tanner and Wong (1987) but see also Liu and Wu (1999). Data augmentation can be adapted to provide a very generic framework for Bayesian analysis of capture-recapture models with unknown N . This idea was introduced for closed populations by Royle et al. (2007), and has subsequently been applied to a number of different contexts including individual covariate models (Royle, 2009), open population models (Royle and Dorazio, 2008, 2010; Gardner et al., 2010), spatial capture-recapture models (Royle and Young, 2008; Royle, 2010; Gardner, 2009), and many others.

Conceptually, data augmentation takes the data you wish you had - that is, the data set with N rows - the known- N data set - and embeds that data set into a larger data set having $M > N$ rows.¹ It is always possible, in practice, to choose M pretty easily for a given problem and context. Then, under data augmentation, analysis is focused on the “augmented data set.” That is, we analyze the bigger data set - the one having M rows - with an appropriate model that accounts for the augmentation. Inference is focused directly on estimating the proportion $\psi = E[N]/M$, instead of directly on N , where ψ is the “data augmentation parameter.”

¹ RC: Might be just me, but I find that formulation a little confusing... I think it's the 'data you wish you had because that's effectively data you don't have. I think it might be easier to grasp if this were explained with the data you do have - based on n .

3.2.1 DA links occupancy models and closed population models

We provide a heuristic description of data augmentation based on the close correspondence between so-called “occupancy” models and closed population models following Royle and Dorazio (2008, sec. xyz).

In occupancy models (MacKenzie et al., 2002; Tyre et al., 2003) the sampling situation is that M sites, or patches, are sampled multiple times to assess whether a species occurs at each site. This yields encounter data such as that illustrated in the left panel of Table 3.3. The important problem is that a species may occur at a site, but go undetected, yielding the “all-zero” encounter histories which are observed. However, some of the all-zeros may well correspond to sites where the species in fact *does not* occur. Thus, while the zeros are observed, there are too many of them and, in a sense, the inference problem is to allocate the zeros into “structural” (fixed) and “sampling” (or stochastic) zeros. More formally, inference is focused on the parameter ψ , the probability that a site is occupied. In contrast, in classical closed population studies, we observe a data set as in the middle panel of Table 3.3 where *no* zeros are observed. The inference problem is, essentially, to estimate how many sampling zeros there are - or should be - in a “complete” data set. The inference objective (how many sampling zeros?) is precisely the same for both types of problems if an upper limit M is specified for the closed population model. The only distinction being that, in occupancy models, M is set by design (i.e., the number of sites to visit) whereas a natural choice of M for capture-recapture models may not be obvious. However, by assuming a uniform prior for N on the integers $[0, M]$, this upper bound is induced (Royle et al., 2007). Then, one can analyze capture-recapture models by adding $M - n$ all-zero encounter histories to the data set and regarding the augmented data set, essentially, as a site-occupancy data set.

Thus, the heuristic motivation of data augmentation is to fix the size of the data set by adding *too many* all-zero encounter histories to create the data set shown in the right panel of Table 3.3 - and then analyze the augmented data set using an occupancy type model which includes both “unoccupied sites” as well as “occupied sites” at which detections did not occur. We call these $M - n$ all-zero histories “potential individuals” because they exist to be recruited (in a non-biological sense) into the population, for example during an analysis by MCMC.

To analyze the augmented data set, we recognize that it is a zero-inflated version of the known- N data set. That is, some of the augmented all-zeros are sampling zeros (corresponding to actual individuals that were missed) and some are “structural” zeros, which do not correspond to individuals in the population. For a basic closed-population model, the resulting likelihood under data augmentation - that is, for the data set of size M - is a simple zero-inflated binomial likelihood. The zero-inflated binomial model can be described “hierarchically”, by introducing a set of binary latent variables, z_1, z_2, \dots, z_M , to indicate whether each individual i is ($z_i = 1$) or is not ($z_i = 0$) a member of the population of N individuals exposed

1675 to sampling. We assume that $z_i \sim \text{Bern}(\psi)$ where ψ is the probability that an
 1676 individual in the data set of size M is a member of the sampled population - in the
 1677 sense that $1 - \psi$ is the probability of realizing a “structural zero” in the augmented
 1678 data set. The zero-inflated binomial model which arises under data augmentation
 1679 can be formally expressed by the following set of assumptions:

$$\begin{aligned} y_i | z_i = 1 &\sim \text{Bin}(K, p) \\ y_i | z_i = 0 &\sim \delta(0) \\ z_i &\stackrel{iid}{\sim} \text{Bern}(\psi) \\ \psi &\sim \text{Unif}(0, 1) \\ p &\sim \text{Unif}(0, 1) \end{aligned}$$

1680 for $i = 1, \dots, M$, where $\delta(0)$ is a point mass at $y = 0$.

1681 We note that N is no longer an explicit parameter of this model. Instead, we
 1682 estimate ψ and functions of the latent variables. In particular, under the assump-
 1683 tions of the zero-inflated model, $z_i \stackrel{iid}{\sim} \text{Bern}(\psi)$; therefore, N is a function of these
 1684 latent variables:

$$N = \sum_{i=1}^M z_i.$$

1685 Further, we note that the latent z_i parameters can be removed from the model by
 1686 integration, in which case the joint probability of the data is

$$\Pr(y_1, \dots, y_M | p, \psi) = \prod_{i=1}^M \psi \text{Bin}(y_i | K, p) + I(y_i = 0)(1 - \psi) \quad (3.2.1)$$

1687 Which can be maximized directly to obtain the MLEs of the structural parameters
 1688 ψ and p or those of other more complex models (e.g., see Royle, 2006). We could
 1689 estimate these parameters and then use them to obtain an estimator of N using
 1690 the so-called “Best unbiased predictor” (see Royle and Dorazio, 2011).

1691 3.2.2 Model M_0 in BUGS

1692 For model M_0 in which we can aggregate the encounter data to individual-specific
 1693 encounter frequencies, the augmented data are given by the vector of frequencies
 1694 $(y_1, \dots, y_n, 0, 0, \dots, 0)$. The zero-inflated model of the augmented data combines
 1695 the model of the latent variables, $z_i \sim \text{Bern}(\psi)$ with the conditional-on- z binomial
 1696 model:

$$\begin{aligned} y_i | z_i = 0 &\sim \delta(0) \\ y_i | z_i = 1 &\sim \text{Bin}(K, p) \end{aligned}$$

Table 3.3. Hypothetical occupancy data set (left), capture-recapture data in standard form (center), and capture-recapture data augmented with all-zero capture histories (right).

Occupancy data				Capture-recapture				Augmented C-R			
site	k=1	k=2	k=3	ind	k=1	k=2	k=3	ind	k=1	k=2	k=3
1	0	1	0	1	0	1	0	1	0	1	0
2	1	0	1	2	1	0	1	2	1	0	1
3	0	1	0	.	0	1	0	3	1	0	1
4	1	0	1	.	1	0	1	4	1	0	1
5	0	1	1	.	0	1	1	5	1	0	1
.	0	1	1	.	0	1	1	.	0	1	1
.	1	1	1	.	1	1	1	.	0	1	1
.	1	1	1	.	1	1	1	.	1	1	1
.	1	1	1	.	1	1	1	.	1	1	1
n	1	1	1	n	1	1	1	n	1	1	1
.	0	0	0					.	0	0	0
.	0	0	0					.	0	0	0
	0	0	0						0	0	0
	0	0	0						0	0	0
	0	0	0						0	0	0
	0	0	0					N	0	0	0
.	0	0	0					.	0	0	0
.	0	0	0						0	0	0
M	0	0	0					.	0	0	0
							
							
							
								M	0	0	0

1697 It is convenient to express the conditional-on- z observation model concisely as:

$$y_i|z_i \sim \text{Bin}(K, pz_i)$$

1698 Thus, if $z_i = 0$ then the success probability of the binomial distribution is identically
 1699 0 whereas, if $z_i = 1$, then the success probability is p . This is useful in describing
 1700 the model in the **BUGS** language, as shown below. Note the last line of the
 1701 model specification here provides the expression for computing N from the data
 1702 augmentation variables z_i .

```

1703 p ~ dunif(0,1)
1704 psi~dunif(0,1)
1705
1706 # nind = number of individuals captured at least once
1707 # nz = number of uncaptured individuals added for PX-DA
1708 for(i in 1:(nind+nz)) {
1709     z[i]~dbern(psi)
1710     mu[i]<-z[i]*p
1711     y[i]~dbin(mu[i],K)
1712 }
1713
1714 N<-sum(z[1:(nind+nz)])
```

1715 Specification of a more general model in terms of the individual encounter obser-
 1716 vations y_{ik} is not much more difficult than for the individual encounter frequencies.
 1717 We define the observation model by a double loop and change the indexing of things
 1718 accordingly, i.e.,

```

1719 for(i in 1:(nind+nz)) {
1720     z[i]~dbern(psi)
1721     for(k in 1:K){
1722         mu[i,k]<-z[i]*p
1723         y[i,k]~dbin(mu[i,k],1)
1724     }
1725 }
```

1726 In this manner, it is straightforward to incorporate covariates on p (see discussion
 1727 of this below and also chapt. 8 (REF XYZ) and consider other extensions.

1728 3.2.3 Formal development of data augmentation

1729 Use of DA for solving inference problems with unknown N can be justified as
 1730 originating from the choice of uniform prior on N . The $\text{Unif}(0, M)$ prior for N is
 1731 innocuous in the sense that the posterior associated with this prior is equal to the
 1732 likelihood for sufficiently large M . One way of inducing the $\text{Unif}(0, M)$ prior on N

1733 is by assuming the following hierarchical prior:

$$\begin{aligned} N &\sim \text{Bin}(M, \psi) \\ \psi &\sim \text{Unif}(0, 1) \end{aligned} \tag{3.2.2}$$

1734 which includes a new model parameter ψ . This parameter denotes the probability
 1735 that an individual in the super-population of size M is a member of the population
 1736 of N individuals exposed to sampling. The model assumptions, specifically the
 1737 multinomial model (eq. XYZ) and eq. 3.2.2, may be combined to yield a reparam-
 1738 eterization of the conventional model that is appropriate for the augmented data
 1739 set of known size M :

$$(n_1, n_2, \dots, n_K) \sim \text{Multin}(M, \psi\pi(1), \psi\pi(2), \dots, \psi\pi(K)) \tag{3.2.3}$$

1740 This arises by removing N from Eq. multinomial XYZ by integrating over the
 1741 binomial prior distribution for N . Thus, the models we analyze under data aug-
 1742 mentation arise formally by removing the parameter N from the ordinary model -
 1743 the model conditional on N - by integrating over a binomial prior distribution for
 1744 N .

1745 Note that the $M - n$ unobserved individuals in the augmented data set have
 1746 probability $\psi\pi(0) + (1 - \psi)$, indicating that these unobserved individuals are a
 1747 mixture of individuals that are sampling zeros ($\psi\pi_0$, and belong to the population
 1748 of size N) and others that are “structural zeros” (occurring in the augmented
 1749 data set with probability $1 - \psi$). In Eq. 3.2.3 N has been eliminated as a formal
 1750 parameter of the model by marginalization (integration) and replaced with the new
 1751 parameter ψ , which we will call the “data augmentation parameter.” However, the
 1752 full likelihood containing both N and ψ can be analyzed (see Royle et al., 2007).

1753 3.2.4 Remarks on Data Augmentation

1754 Data augmentation may seem like a strange and mysterious black-box, and likely
 1755 it is unfamiliar to most people even those with extensive experience with capture-
 1756 recapture models. However, it really is a formal reparameterization of capture-
 1757 recapture models in which N is removed from the ordinary (conditional-on- N)
 1758 model by integration. In the case of Model M0, data augmentation produces the
 1759 zero-inflated binomial which is distinct from the original observation model, but
 1760 only in the sense that it embodies, explicitly, the $\text{Unif}(0, M)$ prior for N . Choice of
 1761 M might be cause for some concern related to potential sensitivity to choice of M .
 1762 The guiding principle is that it should be chosen large enough so that the posterior
 1763 for N is not truncated, but no larger because large values entail more computational
 1764 burden. It seems likely that the properties of the Markov chains should be affected
 1765 by M and so some optimality might exist (Gopalaswamy, 2012), as in occupancy
 1766 models (Mackenzie and Royle, 2005). Formal analysis of this is required.

We emphasize the motivation for data augmentation being that it produces a data set of fixed size, so that the parameter dimension in any capture-recapture model is also fixed. As a result, MCMC is a relatively simple proposition using standard Gibbs Sampling. Consider the simplest context - analyzing Model M0 using the occupancy model. In this case, DA converts Model M0 to a basic occupancy model and the parameters p and ψ have known full-conditional distributions (in fact, beta distributions) that can be sampled from directly. Furthermore, the data augmentation variables - the latent data augmentation variables z , can be sampled from Bernoulli full conditionals. MCMC is not too much more difficult for complicated models - sometimes the hyperparameters need to be sampled using a Metropolis-Hastings step, but nothing more sophisticated than that is required.

There are other approaches to analyzing models with unknown N , using reversible jump MCMC (RJMCMC) or other so-called “trans-dimensional” (TD) algorithms² (Durbin and Elston, 2012; King, missing; Schofield and Barker, missing). What distinguishes DA from RJMCMC and related TD methods is that DA is used to create a distinctly new model that is unconditional on N and we (usually) analyze the unconditional model. The various TD/RJMCMC approaches seek to analyze the conditional-on- N model in which the dimensional of the parameter space is a variable function of N . TD/RJMCMC approaches might appear to have the advantage that one can model N explicitly or consider alternative priors for N . However, despite that N is removed as an explicit parameter in DA, it is possible to develop hierarchical models that involve structure on N (Converse and Royle, 2010; Royle et al., 2011a) which we consider in chapt. XYZ.

3.2.5 Example: Black Bear Study on Fort Drum

To illustrate the analysis of Model M0 using data augmentation, we use a data set collected at Fort Drum Military Installation in upstate New York by the Department of Defense, Cornell University and colleagues. These data have been analyzed in various forms by Gardner (2009); Gardner et al. (2010), and Wegan (2008). The specific data used here are encounter histories on 47 individuals obtained from an array of 38 baited “hair snares” (Fig. 3.1) during June and July 2006. Barbed wire traps were baited and checked for hair samples each week for eight weeks, thus we have $K = 8$ sample intervals. The data are provided on the Web Supplement and the analysis can be set up and run as follows. Here, the data were augmented with $M - n = 128$ ($M = 175$) all-zero encounter histories.

```
# Consider adding comments to your code.
## Good idea. This will be done in final draft
trapmat<-read.csv("FDtrapmat.csv")
bearArray<-source("FDbeararray.R")$value
nind<-dim(bearArray)[1]
```

²Look these citations up in Royle-Dorazio EURING paper

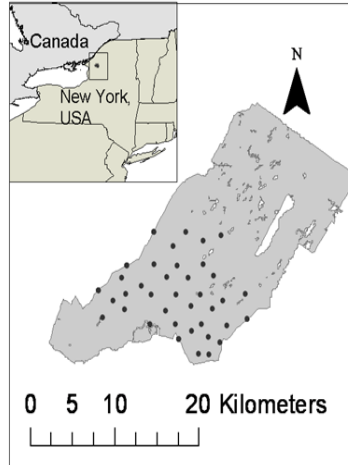


Figure 3.1. Fort Drum study area and hair snare locations.

```

1806 K<-dim(bearArray)[3]
1807 ntraps<-dim(bearArray)[2]
1808
1809 M=175
1810 nz<-M-nind
1811
1812 Xaug <- array(0, dim=c(M,ntraps,K))
1813 Xaug[1:nind,,]<-bearArray
1814 y<- apply(Xaug,c(1,3),sum)
1815 y[y>1]<-1
1816 ytot<-apply(y,1,sum) # total encounters out of K

```

Note that the raw data, \mathbf{y} , is an $M \times K$ array of individual encounter events (i.e., $y_{ik} = 1$ if individual i was encountered in any trap and 0 otherwise). For $i = 48, \dots, 175$, $y_{ik}=0$ as these are augmented observations. For Model M0 it is sufficient to reduce the data to individual encounter frequencies which we have labeled \mathbf{y}_{tot} above. The BUGS model file along with commands to fit the model are as follows:

```

1823 set.seed(2013) # to obtain the same results each time
1824 data0<-list(y=y,M=M,K=K)
1825 params0<-list('psi','p','N')
1826 zst=c(rep(1,nind),rbinom(M-nind, 1, .5))
1827 inits = function() {list(z=zst, psi=runif(1), p=runif(1)) }
1828
1829 cat("

```

```

1830 model {
1831
1832   psi~dunif(0, 1)
1833   p~dunif(0,1)
1834
1835   for (i in 1:M){
1836     z[i]~dbern(psi)
1837     for(k in 1:K){
1838       tmp[i,k]<-p*z[i]
1839       y[i,k]~dbin(tmp[i,k],1)
1840     }
1841   }
1842   N<-sum(z[1:M])
1843 }
1844 ",file="modelM0.txt")
1845
1846 fit0 = bugs(data0, inits, params0, model.file="modelM0.txt",
1847            n.chains=3, n.iter=2000, n.burnin=1000, n.thin=1,
1848            debug=TRUE,working.directory=getwd())

```

1849 The posterior summary statistics from this analysis are as follows:

```

1850 > print(fit0,digits=2)
1851 Inference for Bugs model at "modelM0.txt", fit using WinBUGS,
1852 3 chains, each with 2000 iterations (first 1000 discarded)
1853 n.sims = 3000 iterations saved
1854
1855      mean    sd  2.5%   25%   50%   75%  97.5% Rhat n.eff
1856 psi      0.29 0.04  0.22  0.26  0.29  0.31  0.36   1 3000
1856 p       0.30 0.03  0.25  0.28  0.30  0.32  0.35   1 3000
1857 N       49.94 1.99 47.00 48.00 50.00 51.00 54.00   1 3000
1858 deviance 489.05 11.28 471.00 480.45 488.80 495.40 513.70   1 3000
1859
1860 [.. some output deleted ...]

```

1861 **WinBUGS** did well in choosing an MCMC algorithm for this model – we have
1862 $\hat{R} = 1$ for each parameter, and an effective sample size of 3000, equal to the total
1863 number of posterior samples. We see that the posterior mean of N under this model
1864 is 49.94 and a 95% posterior interval is (48, 54). We revisit these data later in the
1865 context of more complex models.

1866 In order to obtain an estimate of density, D , we need an area to associate with
1867 the estimate of N , and commonly used procedures to conjure up such an area
1868 include buffering the trap array by the home range radius, often estimated by the
1869 mean maximum distance moved (MMDM)³, 1/2 MMDM (Dice, 1938) or directly
1870 from telemetry data (REF XXX NEED REF HERE XXXXX). Typically, the trap
1871 array is defined by the convex hull around the trap locations, and this is what we

³really MMDM? How can this be an estimate of the home range radius? Reference for this?

1872 applied a buffer to. We computed the buffer by using an estimate of the mean female
 1873 home range radius (2.19 km) estimated from telemetry studies (Bales et al., 2005)
 1874 instead of using an estimate based on our relatively more sparse recapture data⁴.
 1875 For the Fort Drum study, the convex hull has area 157.135 km^2 , and the buffered
 1876 convex hull has area 277.011 km^2 . To create this we used functions contained in the
 1877 **R** package **rgeos** and created a utility function **bcharea** which is in our **R** package
 1878 **scrbook**. The commands are as follows:

```
1879 library("rgeos")
1880
1881 bcharea<-function(buff,traplocs){
1882   p1<-Polygon(rbind(traplocs,traplocs[1,]))
1883   p2<-Polygons(list(p1=p1),ID=1)
1884   p3<-SpatialPolygons(list(p2=p2))
1885   p1ch<-gConvexHull(p3)
1886   bp1<-gBuffer(p1ch, width=buff)
1887   plot(bp1, col='gray')
1888   plot(p1ch, border='black', lwd=2, add=TRUE)
1889   gArea(bp1)
1890 }
1891
1892 bcharea(2.19,traplocs=trapmat)
```

1893 The resulting buffered convex hull is shown in Fig. 3.2.

1894 To conjure up a density estimate under model M_0 , we compute the appropriate
 1895 posterior summary of N and the prescribed area (277.011 km^2):

```
1896 > summary(fit0$sims.list$N/277.011)
1897   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
1898 0.1697 0.1733 0.1805 0.1803 0.1841 0.2130
1899
1900 > quantile(fit0$sims.list$N/277.011,c(0.025,0.975))
1901      2.5%      97.5%
1902 0.1696684 0.1949381
```

1903 which yields a density estimate of about 0.18 ind/ km^2 , and a 95% Bayesian confi-
 1904 dence interval of (0.170, 0.195).

1905 The obvious limitation of this estimate and, indeed, of the whole process, is that
 1906 our choice of “area” is completely subjective - which area should we use? MMDM?
 1907 One-half MMDM? Estimated from telemetry data? And, furthermore, how certain
 1908 are we of this area? Can we quantify our uncertainty about this quantity? More
 1909 important, what exactly is the meaning of this area and, in this context, how do
 1910 we gauge bias and/or variance of “estimators” of it? (i.e., what is it estimating?).

⁴BETH: Why?

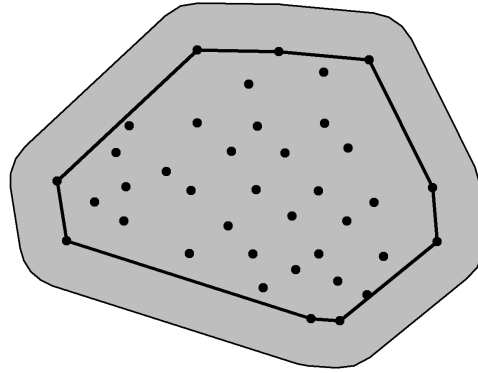


Figure 3.2. buffered convex hull of the bear hair snare array

3.3 TEMPORALLY VARYING AND BEHAVIORAL EFFECTS

The purpose of this chapter is mainly to emphasize the central importance of the binomial model in capture-recapture and so we have considered models for individual encounter frequencies - the number of times individuals are captured out of K samples. Sometimes it is not acceptable to aggregate the encounter data for each individual - such as when encounter probability varies over time among samples. A type of time-varying response that seems relevant in most capture-recapture studies is “effort” such as amount of search time, number of observers, or trap effort or when p depends on date (Kéry et al., 2010; Gardner et al., 2010). A common situation is that in which there exists a “behavioral response” to trapping (even if the animal is not physically trapped).

Behavioral response is an important concept in carnivore studies because individuals might learn to come to baited traps or avoid traps due to trauma related to being encountered. There are a number of ways to parameterize a behavioral response to encounter. The distinction between persistent and ephemeral was made by Yang and Chao (2005) who considered a general behavioral response model of the form:

$$\text{logit}(p_{ik}) = \alpha_0 + \alpha_1 * y_{i,k-1} + \alpha_2 x_{ik}$$

where x_{ik} is a covariate indicator variable of previous capture (i.e., $x_{ik} = 1$ if

captured in any previous period). Therefore, encounter probability changes depending on whether an individual was captured in the immediate previous period (ephemeral behavioral response) or in any previous period (persistent behavioral response). The former probably models a behavioral response due to individuals moving around their territory relatively slowly over time and the latter probably accommodates trap happiness due to baiting or shyness due to trauma. In spatial capture-recapture models it makes sense to consider a local behavioral response that is trap-specific (Royle et al., 2011c) - that is, the encounter probability is modified for individual traps depending on previous capture in specific traps.

Models with temporal effects are easy to describe in the **BUGS** language and analyze and we provide a number of examples in chapt. 8. XXXXX ?? XXXXX

3.4 MODELS WITH INDIVIDUAL HETEROGENEITY

Here we consider models with individual-specific encounter probability parameters, say p_i , which we model according to some probability distribution, $g(\theta)$. We denote this basic model assumption as $p_i \sim g(\theta)$. This type of model is similar in concept to extending a GLM to a GLMM but in the capture-recapture context N is unknown. The basic class of models is often referred to as “Model M_h ” but really this is a broad class of models, each being distinguished by the specific distribution assumed for p_i . There are many different varieties of Model M_h including parametric and various putatively non-parametric approaches (Burnham and Overton, 1978; Norris III and Pollock, 1996; Pledger, 2000). One important practical matter is that estimates of N can be extremely sensitive to the choice of heterogeneity model (Fienberg et al., 1999; Dorazio and Royle, 2003; Link, 2003). Indeed, Link (2003) showed that in some cases it’s possible to find models that yield precisely the same expected data, yet produce wildly different estimates of N . In that sense, N for most practical purposes is not identifiable across classes of mixture models, and this should be understood before fitting any such model. One solution to this problem is to seek to model explicit factors that contribute to heterogeneity, e.g., using individual covariate models (See 3.5 below). Indeed, spatial capture-recapture models seek to do just that, by modeling heterogeneity due to the spatial organization of individuals in relation to traps or other encounter mechanism. For additional background and applications of Model M_h see Royle and Dorazio (2008, chapt. 6) and Kery and Schaub (2011, chapt. 6).

Model M_h has important historical relevance to spatial capture-recapture situations (Karanth, 1995) because investigators recognized that the juxtaposition of individuals with the array of trap locations should yield heterogeneity in encounter probability, and thus it became common to use some version of Model M_h in spatial trapping arrays to estimate N . While this doesn’t resolve the problem of not knowing the area relevant to N , it does yield an estimator that accommodates the heterogeneity in p induced by the spatial aspect of capture-recapture studies.

To see how this juxtaposition induces heterogeneity, we have to understand

the relevance of movement in capture-recapture models. Imagine a quadrat that can be uniformly searched by a crew of biologists for some species of reptile (see Royle and Young (2008)). Figure 3.3 shows a sample quadrat searched repeatedly over a period of time. Further, suppose that species exhibits some sense of spatial fidelity in the form of a home range or territory, and individuals move about their home range (home range centroids are given by the blue dots) in some kind of random fashion. It is natural to think about it in terms of a movement process and sometimes that movement process can be modeled explicitly using hierarchical models (Royle and Young, 2008; Royle et al., 2011b). Heuristically, we imagine that each individual in the vicinity of the study area is liable to experience variable exposure to encounter due to the overlap of its home range with the sampled area - essentially the long-run proportion of times the individual is within the sample plot boundaries, say ϕ . We might model the exposure of an individual to capture by supposing that $z_i = 1$ if individual i is available to be captured (i.e., within the survey plot) during any sample, and 0 otherwise. Then, $\Pr(z_i = 1) = \phi$. In the context of spatial studies, it is natural that ϕ should depend on *where* an individual lives, i.e., it should be individual-specific ϕ_i (Chandler et al., 2011). This system describes, precisely, that of “random temporary emigration” (Kendall, 1997) where ϕ_i is the individual-specific probability of being “available” for capture.

Conceptually, SCR models aim to deal with this problem of variable exposure to sampling due to movement in the proximity of the trapping array explicitly and formally with auxiliary spatial information. If individuals are detected with probability p_0 , *conditional* on $z_i = 1$, then the marginal probability of detection of individual i is

$$p_i = p_0 \phi_i$$

so we see clearly that individual heterogeneity in encounter probability is induced as a result of the juxtaposition of individuals (i.e., their home ranges) with the sample apparatus and the movement of individuals about their home range.

We will work with a specific type of Model M_h here, that in which we extend the basic binomial observation model of Model M_0 so that

$$\text{logit}(p_i) = \mu + \eta_i$$

where

$$\eta_i \sim \text{Normal}(0, \sigma_p^2)$$

We could as well write

$$\text{logit}(p_i) \sim \text{Normal}(\mu, \sigma_p^2)$$

This “logit-normal mixture” was analyzed by Coull and Agresti (1999) and elsewhere. It is a natural extension of the basic model with constant p , as a mixed GLMM, and similar models occur throughout statistics. It is also natural to consider a beta prior distribution for p_i (Dorazio and Royle, 2003) and so-called “finite-mixture” models are also popular (Norris III and Pollock, 1996; Pledger, 2000).

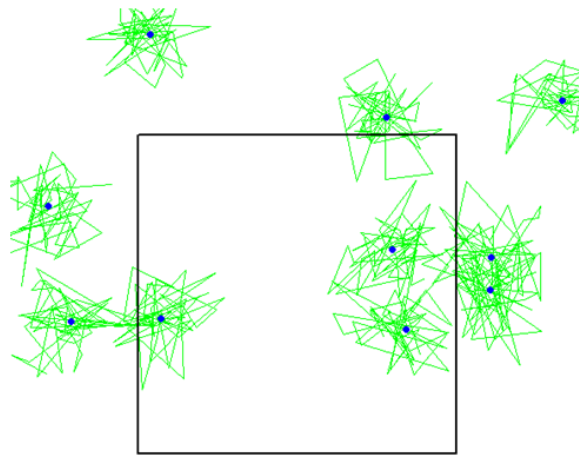


Figure 3.3. A quadrat searched for lizards and the locations of each lizard over some period of time.

3.4.1 Analysis of Model Mh

If N is known, it is worth taking note of the essential simplicity of Model M_h as a binomial GLMM. This is a type of model that is widely applied in just about every scientific discipline and using standard methods of inference based either on integrated likelihood (Laird and Ware, 1982; Berger et al., 1999) which we discuss in chapt. 6 or standard Bayesian methods. However, because N is not known, inference is somewhat more challenging. We address that here using Bayesian analysis based on data augmentation (DA). Although we use data augmentation in the context of Bayesian methods here, we note that heterogeneity models formulated under DA are easily analyzed by conventional likelihood methods as zero-inflated binomial mixtures (Royle, 2006) and more traditional analysis of model M_h based on integrated likelihood, without using data augmentation, has been considered by Coull and Agresti (1999), Dorazio and Royle (2003), and others.

As with model M_0 , we have the Bernoulli model for the zero-inflation variables: $z_i \sim \text{Bern}(\psi)$ and the model of the observations expressed conditional on the latent variables z_i . For $z_i = 1$, we have a binomial model with individual-specific p_i :

$$y_i | z_i = 1 \sim \text{Bin}(K, p_i)$$

and otherwise $y_i | z_i = 0 \sim \delta(0)$. Further, we prescribe a distribution for p_i . Here we assume

$$\text{logit}(p_i) \sim \text{Normal}(\mu, \sigma^2)$$

The basic **BUGS** description for this model, assuming a $\text{Unif}(0, 1)$ prior for $p_0 = \text{logit}^{-1}(\mu)$, is given as follows:

```

model{
  p0 ~ dunif(0,1)          # prior distributions
  mup<- log(p0/(1-p0))
  taup~dgamma(.1,.1)
  psi~dunif(0,1)

  for(i in 1:(nind+nz)){
    z[i]~dbern(psi)        # zero inflation variables
    lp[i] ~ dnorm(mup,taup) # individual effect
    logit(p[i])<-lp[i]
    mu[i]<-z[i]*p[i]
    y[i]~dbin(mu[i],J)     # observation model
  }

  N<-sum(z[1:(nind+nz)])  # N is a derived parameter
}
```

3.4.2 Analysis of the Fort Drum data

The logit-normal heterogeneity model was fitted to the bear data from the Fort Drum study, and we used data augmentation to produce a data set of $M = 500$ individuals. We ran the model using **JAGS** with the instructions given as follows⁵.

```

2045 [... get data as before ....]
2046
2047 set.seed(2013)
2048
2049 cat("
2050 model{
2051   p0 ~ dunif(0,1)          # prior distributions
2052   mup<- log(p0/(1-p0))
2053   sigmap ~ dunif(0,10)
2054   taup<- 1/(sigmap*sigmap)
2055   psi~dunif(0,1)
2056
2057   for(i in 1:(nind+nz)){
2058     z[i]~dbern(psi)        # zero inflation variables
2059     lp[i] ~ dnorm(mup,taup) # individual effect
2060     logit(p[i])<-lp[i]
2061     mu[i]<-z[i]*p[i]
2062     y[i]~dbin(mu[i],K)    # observation model
2063   }
2064
2065   N<-sum(z[1:(nind+nz)])
2066 }
2067 ",file="modelMh.txt")
2068
2069 data1<-list(y=ytot, nz=nz, nind=nind,K=K)
2070 params1= c('p0','sigmap','psi','N')
2071 inits = function() {list(z=as.numeric(ytot>=1), psi=.6, p0=runif(1),
2072   sigmap=runif(1,.7,1.2),lp=rnorm(M,-2)) }
2073
2074 library("rjags")
2075 jm<- jags.model("modelMh.txt", data=data1, inits=inits, n.chains=4,
2076   n.adapt=1000)
2077 jout<- coda.samples(jm, params1, n.iter=200000, thin=1)
2078
2079   This produces the posterior distribution for  $N$  shown in Fig. 3.4. Posterior
2080   summaries of parameters are given as follows:
2081
2082 > summary(jout)
2083
2084 Iterations = 2001:202000

```

⁵For WinBUGS, should provide starts for lp and sigma or sometimes WinBUGS breaks

```

2083 Thinning interval = 1
2084 Number of chains = 4
2085 Sample size per chain = 2e+05
2086
2087 1. Empirical mean and standard deviation for each variable,
2088    plus standard error of the mean:
2089
2090           Mean          SD Naive SE Time-series SE
2091 N      117.7740 56.31633 6.296e-02      1.960115
2092 p0       0.0728  0.05522 6.174e-05      0.001655
2093 psi      0.2366  0.11362 1.270e-04      0.003909
2094 sigmap   2.0795  0.53096 5.936e-04      0.016789
2095
2096 2. Quantiles for each variable:
2097
2098           2.5%      25%      50%      75%      97.5%
2099 N      62.000000 82.00000 102.00000 134.0000 277.0000
2100 p0      0.003143  0.02842  0.06077  0.1066  0.2036
2101 psi     0.117269  0.16377  0.20522  0.2712  0.5560
2102 sigmap  1.211900  1.69434  2.02113  2.4028  3.2694

```

2103 We used $M = 500$ for this analysis and we note that while the posterior mass
 2104 of N is concentrated away from this upper bound (Fig. 3.4), the posterior has an
 2105 extremely long right tail, with some posterior values at the upper bound $N = 500$.
 2106 Maybe or maybe not sufficient data augmentation.⁶ The model runs effectively in
 2107 **WinBUGS** but sometimes with apparently inefficient mixing for reasons that may
 2108 be related to bad starting values. In some cases this was resolved if we supplied
 2109 starting values for the $\text{logit}(p_i)$ parameters and τ .

2110 Because of the skewed posterior we see that the posterior mean ($N = 117$)
 2111 is considerably higher than the posterior mode ($N = 102$). Moreover, posterior
 2112 summaries are estimated with a relatively high error (“Time-series SE” of around
 2113 2.0)⁷. Further, it may be surprising that the posterior mode does not compare well
 2114 with the MLE. To compute the posterior mode we could easily find the posterior
 2115 value of N with the highest mass because N is discrete. But we want to smooth out
 2116 some of the Monte Carlo error a bit so we used a smoothing spline to the posterior
 2117 frequencies of N as follows:

```

2118 > tt<-table(jout[[1]][, "N"])[1:80]
2119 > xg<-as.numeric(names(tt))
2120 > plot(xg, tt)
2121 > sp<- smooth.spline(xg, tt, df=9)
2122 > sp$x[sp$y==max(sp$y)]
2123 [1] 80

```

⁶ to do: insert final results. longer run. more data augmentation. compare with winbugs.

⁷ need to define this somewhere

The `df` argument controls the degree of smoothing and we find in this case that the modal value (i.e., 80) is not too sensitive to the smoothing parameter but this should be checked in any specific instance⁸.

To compare with the MLE, we used the **R** code contained in Panel 6.1 of Royle and Dorazio (2008). The MLE of $\log(n_0)$, the logarithm of the number of uncaptured individuals, is $\log(n_0) = 3.86$ and therefore the MLE is $\hat{N} = \exp(3.86) + 47 = 94.47$ which is not at all consistent with the apparent mode in Fig. 3.4.⁹

Comments: First of all the posterior for this model and data set is very sensitive to prior distributions. While MLEs are invariant to transformation of the parameters, the posterior distribution definitely is *not* invariant. In the present case, the use of a $\text{Unif}(0, 1)$ prior for $p_0 = \expit(\mu)$ is somewhat informative – in particular, it is not at all “flat” on the scale of μ – and this affects the posterior. We generally always recommend use of a $\text{Unif}(0, 1)$ prior for $\expit(\mu)$ in such models. That said, we were surprised at this result, and we experimented with other prior configurations including putting a flat prior on μ directly. This specific prior suggests the possibility that the posterior distribution may be improper for that prior specification. This kind of small sample instability has been widely noted in Model Mh (Fienberg et al., 1999; Dorazio and Royle, 2003) and is not unrelated to sensitivity to model which has also been identified as an important issue in model M_h (Dorazio and Royle, 2003; Link, 2003). Conclusion: The mode is well-defined but the data set is sparse and hence inferences are poor and sensitive to model choices. Get over it.

3.4.3 Building your own MCMC algorithm

For fun, we construct our own MCMC algorithm using a Metropolized Gibbs sampler for Model M_h . In chapter 7 we devise MCMC algorithms for spatial capture-recapture models and the basic conceptual and technical considerations are entirely analogous to Model M_h .

To begin, we first collect all of our model components which are as follows: $[y_i|p_i, z_i]$, $[p_i|\mu_p, \sigma_p]$, and $[z_i|\psi]$ for each $i = 1, 2, \dots, M$ and then prior distributions $[\mu_p]$, $[\sigma_p]$ and $[\psi]$. The joint posterior distribution of all unknown quantities in the model is proportional to the joint distribution of all elements y_i, p_i, z_i and also the prior distributions of the prior parameters:

$$\left\{ \prod_{i=1}^M [y_i|p_i, z_i][p_i|\mu_p, \sigma_p][z_i|\psi] \right\} [\mu_p, \sigma_p, \psi]$$

For prior distributions, we assume that μ_p, σ_p, ψ are mutually independent and for μ_p and σ_p we use improper uniform priors, and $\psi \sim \text{Unif}(0, 1)$. Note that the

⁸we need to give examples of using `density()` to obtain modes

⁹We note that the result is inconsistent with Gardner et al. (2009) who reported an MLE of 104.1 ($\text{density} = 0.437 \text{ inds/km}^2$) although we do not know the reason for this at the present time.




Figure 3.4. Posterior of N for Fort Drum bear study data under the logit-normal version of model M_h . From WinBUGS output. 200k samples.

likelihood contribution for each individual, when conditioned on p_i and z_i , does not depend on ψ , μ_p , or σ_p . As such, the full-conditionals for the structural parameters ψ only depends on the collection of data augmentation variables z_i , and that for μ_p and σ_p will only depends on the collection of latent variables p_i ; $i = 1, 2, \dots, M$. The full conditionals for all the unknowns are as follows:

(1) For p_i :

$$[p_i | y_i, \mu_p, \sigma_p, z_i = 1] \propto [y_i | p_i][p_i | \mu_p, \sigma_p^2] \text{ if } z_i = 1$$

$$[p_i | \mu_p, \sigma_p] \text{ if } z_i = 0$$

(2) for z_i :

$$z_i | \cdot \propto [y_i | z_i * p_i] \text{Bern}(z_i | \psi)$$

(3) For μ_p :

$$[\mu_p | \cdot] \sim \prod_i [p_i | \cdot] * \text{const}$$

(4) For σ_p :

$$[\sigma_p | \cdot] \sim \prod_i [p_i | \cdot] * \text{const}$$

(5) For ψ :

$$\psi | \cdot \sim \text{Beta}(1 + \sum z_i, 1 + M - \sum z_i)$$

We've identified each of the full conditional distributions in sufficient detail to apply the Metropolis-Hastings algorithm. With the exception of ψ which has a convenient analytic solution – it is a beta distribution which we can easily sample directly. In truth, we could also sample μ_p and σ_p^2 directly with certain choices of prior distributions. For example, if $\mu_p \sim \text{Normal}(0, 1000)$ then the full conditional for μ_p is also normal, etc.. We implement an MCMC algorithm for this model in the following block of **R** code. The basic structure is: initialize the parameters and create any required output or intermediate data holders, and then begin the main MCMC loop which, in this case, generates 100000 samples.¹⁰

```
## obtain the bear data by executing the previous data grabbing
## function
temp<-getdata()
M<-temp$M
K<-temp$K
ytot<-temp$ytot

###
### MCMC algorithm for Model Mh
```

¹⁰This data grabbing function is not implemented yet

```

2187
2188 out<-matrix(NA,nrow=100000,ncol=4)
2189 dimnames(out)<-list(NULL,c("mu","sigma","psi","N"))
2190 lp<- rnorm(M,-1,1)
2191 p<-expit(lp)
2192 mu<- -1
2193 p0<-exp(mu)/(1+exp(mu))
2194 sigma<- 1
2195 psi<- .5
2196 z<-rbinom(M,1,psi)
2197 z[ytot>0]<-1
2198
2199 for(i in 1:100000){
2200
2201   ### update the logit(p) parameters
2202   lpc<- rnorm(M,lp,1) # 0.5 is a tuning parameter
2203   pc<-expit(lpc)
2204   lik.curr<-log(dbinom(ytot,K,z*p)*dnorm(lp,mu,sigma))
2205   lik.cand<-log(dbinom(ytot,K,z*pc)*dnorm(lpc,mu,sigma))
2206   kp<- runif(M) < exp(lik.cand-lik.curr)
2207   p[kp]<-pc[kp]
2208   lp[kp]<-lpc[kp]
2209
2210   p0c<- rnorm(1,p0,.05)
2211   if(p0c>0 & p0c<1){
2212     muc<-log(p0c/(1-p0c))
2213     lik.curr<-sum(dnorm(lp,mu,sigma,log=TRUE))
2214     lik.cand<-sum(dnorm(lp,muc,sigma,log=TRUE))
2215     if(runif(1)<exp(lik.cand-lik.curr)) {
2216       mu<-muc
2217       p0<-p0c
2218     }
2219   }
2220
2221   sigmac<-rnorm(1,sigma,.5)
2222   if(sigmac>0){
2223     lik.curr<-sum(dnorm(lp,mu,sigma,log=TRUE))
2224     lik.cand<-sum(dnorm(lp,mu,sigmac,log=TRUE))
2225     if(runif(1)<exp(lik.cand-lik.curr))
2226       sigma<-sigmac
2227   }
2228
2229   ### update the z[i] variables
2230   zc<- ifelse(z==1,0,1) # candidate is 0 if current = 1, etc..
2231   lik.curr<- dbinom(ytot,K,z*p)*dbinom(z,1,psi)
2232   lik.cand<- dbinom(ytot,K,zc*p)*dbinom(zc,1,psi)

```

```

2233 kp<- runif(M) < (lik.cand/lik.curr)
2234 z[kp]<- zc[kp]
2235
2236 psi<-rbeta(1, sum(z) + 1, M-sum(z) + 1)
2237
2238 out[i,]<- c(mu,sigma,psi,sum(z))
2239 }

```

Remarks: (1) for parameters with bounded support, i.e., σ_p and p_0 , we are using a random walk candidate generator but rejecting draws outside of the parameter space. (2) We mostly use Metropolis-Hastings except for the data augmentation parameter ψ which we sample directly from its full-conditional distribution which is a beta distribution. (3) Even the latent data augmentation variables z_i are updated using Metropolis-Hastings although they too can be updated directly from their full-conditional.

3.4.4 Exercises related to model Mh

- (1) Enclose the MCMC algorithm in an R function and provide arguments for some of the parameters of the function that a user might wish to modify.
- (2) Execute the function and compare the results to those generated from WinBUGS in the previous section
- (3) Note that the prior distribution for the “mean” parameter is given on $p_0 = \exp(\mu)/(1 + \exp(\mu))$. Reformulate the algorithm with a flat prior on μ and see what happens. Contemplate this.
- (4) Using Bayes rule, figure out the full conditional for z_i so that you don’t have to use MH for that one. It might be more efficient. Is it?
- (5) Modify the MCMC algorithm so that the prior for μ_p is an improper flat prior. i.e., $[\mu_p] \propto 1$. Describe the posterior distribution of N .

3.5 INDIVIDUAL COVARIATE MODELS: TOWARD SPATIAL CAPTURE-RECAPTURE

A standard situation in capture-recapture models is when an individual covariate is measured, and this covariate is thought to influence encounter probability. As with other closed population models, we begin with the basic binomial observation model:

$$y_i \sim \text{Bin}(K, p_i)$$

and we assume also a model for encounter probability according to:

$$\text{logit}(p_i) = \alpha + \beta x_i \quad (3.5.1)$$

Classical examples of covariates influencing detection probability are type of animal (juvenile/adult or male/female), a continuous covariate such as body mass (Royle

and Dorazio, 2008, ch. 6), or a discrete covariate such as group or cluster size. For example, in models of aerial survey data, it is natural to model detection probabilities as a function of the observation-level individual covariate, “group size” (Royle, 2008, 2009; ?).

Such “individual covariate models” are similar in structure to Model M_h , except that the individual effects are *observed* for the n individuals that appear in the sample. These models are important here because spatial capture-recapture models are precisely a form of individual covariate model, an idea that we will develop here and elsewhere. Specifically, they are such models, but where the individual covariate is a partially observed latent variable similar.. That is, unlike Model M_h , we do have some direct information about the latent variable, which comes from the spatial locations/distribution of individual recaptures.

Traditionally, estimation of N in individual covariate models is achieved using methods based on ideas of unequal probability sampling (i.e., Horwitz-Thompson estimation; see Huggins (1989) and Alho (1990)). An estimator of N is

$$\hat{N} = \sum_i^n \frac{1}{\tilde{p}_i}$$

where \tilde{p}_i is the probability that individual i appeared in the sample. That is, $\tilde{p}_i = \Pr(y_i > 0)$ where, in closed population capture-recapture models,

$$\Pr(y_i > 0) = (1 - (1 - p_i)^K)$$

where p_i is a function of parameters α and β according to Eq. 3.5.1. In practice, parameters are estimated from the conditional-likelihood of the observed encounter histories which is, for observation y_i ,

$$\mathcal{L}(\alpha, \beta | y_i) = \frac{\text{Bin}(y_i | \alpha, \beta)}{\tilde{p}_i}.$$

Here we take a formal model-based approach to Bayesian analysis of such models using data augmentation (Royle, 2009). Classical likelihood analysis of the so-called “full likelihood” is covered by Borchers et al. (2002). For Bayesian analysis of individual covariate models, because the individual covariate is unobserved for the $N - n$ uncaptured individuals, we require a model to describe variation among individuals, essentially allowing the sample to be extrapolated to the population¹¹. For our present purposes, we consider a continuous covariate and we assume that it has a normal distribution:

$$x_i \sim \text{Normal}(\mu, \sigma^2)$$

¹¹weak argument

2294 Data augmentation can be applied directly to this class of models. In particular,
 2295 reformulation of the model under DA yields a basic zero-inflated binomial model of
 2296 the form:

$$\begin{aligned} z_i &\sim \text{Bern}(\psi) \\ y_i|z_i=1 &\sim \text{Bin}(K, p_i(x_i)) \\ y_i|z_i=0 &\sim \delta(0) \\ x_i &\sim \text{Normal}(\mu, \sigma^2) \end{aligned}$$

2297 Fully spatial capture-recapture models use this formulation with a latent covariate
 2298 that is directly related to the individual detection probability (see next Section).
 2299 As with the previous models, implementation is trivial in the **BUGS** language.
 2300 The **BUGS** specification is very similar to that for model M_h , but we require the
 2301 distribution of the covariate to be specified, along with priors for the parameters of
 2302 that distribution.

2303 3.5.1 Example: Location of capture as a covariate.

2304 If we had a regular grid of traps over some closed geographic system then we imagine
 2305 that the average location of capture would be a decent estimate (heuristically) of
 2306 an individual's home range center. Intuitively some measure of typical distance
 2307 from home range center to traps for an individual should be a decent covariate to
 2308 explain heterogeneity in encounter probability, i.e., individuals with more exposure
 2309 to traps should have higher encounter probabilities and vice versa. A version of
 2310 this idea was put forth by Boulanger and McLellan (2001) (see also Ivan (2012)),
 2311 but using the Huggins-Alho estimator and with covariate "distance to edge" of the
 2312 trapping array. A limitation of this approach is that it does not provide a solution
 2313 to the problem that the trap area is fundamentally ill-defined, nor does it readily
 2314 accommodate the inherent and heterogeneous variation in this measured covariate.

2315 Here, we provide an example of this type of heuristically motivated approach
 2316 using the fully model-based individual covariate model described above analyzed
 2317 by data augmentation. We take a slightly different approach than that adopted
 2318 by Boulanger and McLellan (2001). By analyzing the full likelihood and placing a
 2319 prior distribution on the individual covariate, we resolve the problem of having an
 2320 ill-defined area over which the population size is distributed. After you read later
 2321 chapters of this book, it will be apparent that SCR models represent a formalization
 2322 of this heuristic procedure.

2323 For our purposes here, we define $x_i = ||\mathbf{s}_i - \mathbf{x}_0||$ where \mathbf{s}_i is the average encounter
 2324 location of individual i and \mathbf{x}_0 is the centroid of the trap array. Conceptually,
 2325 individuals in the middle of the array should have higher probability of encounter
 2326 and, as x_i increases, p_i should therefore decrease. We note that we have defined
 2327 s_i in terms of a sample quantity - the observed mean - which is ad hoc but maybe
 2328 satisfactory under the circumstances. That said, for an expansive, dense trapping

grid then we might expect the sample mean encounter location to be a good estimate of home range center but, clearly this is biased for individuals that live around the edge (or off) the trapping array. Regardless, it should be good enough for our present purposes of demonstrating this heuristically appealing application of an individual covariate model. A key point is that s_i is missing for each individual that is not encountered and thus so is x_i . Thus, it is a latent variable, or random effect, and we need therefore to specify a probability distribution for it. As a measurement of distance we know it must be positive-valued. Thinking about this like a distance sampling problem lets first try to make x_i uniform from 0 to some large number, say D_{max} , beyond which it would be difficult to imagine an individual being captured. For example, D_{max} should be at a home range diameter past the furthest trap from the center. As such, we use this distribution for the individual covariate “distance from center of the trap array”

$$x_i \sim \text{uniform}(0, D_{max})$$

where D_{max} is a specified constant. In practice, people have used distance from edge of the trap array but that is less easy to define and compute.

ANDY STOPPED RIGHT HERE

Note: I’m a little cautious about stumbling through an analysis that is blatantly wrong – but I’m hoping it will evolve into a learning experience for the reader.

Fort Drum Bear Study

We have to do a little bit of data processing to fit this individual covariate model to the Fort Drum data. We compute the individual covariate \mathbf{x}_i using the **R** script provided in *scrbook*. We picked $D_{max} = 11.5 \text{ km}^2$ which is about the distance from the array center to the furthest trap. Once we specific D_{max} then the implication is that the population size parameter applies to the area within 11.5 units of the trap centroid¹². The **BUGS** model specification and **R** commands to package the data and fit the model are as follows:

```
cat("
model{
  p0 ~ dunif(0,1)          # prior distributions
  mup<- log(p0/(1-p0))
  psi~dunif(0,1)
  beta~dnorm(0,.01)
  for(i in 1:(nind+nz)){
    xcent[i]~dunif(0,maxD)
    z[i]~dbern(psi)         # DA variables
    lp[i] <- mup + beta*xcent[i] # individual effect
```

¹²To be convincing this might need a little bit of hand-holding

```

2367   logit(p[i])<-lp[i]
2368   mu[i]<-z[i]*p[i]
2369   y[i]~dbin(mu[i],K) # observation model
2370 }
2371 N<-sum(z[1:(nind+nz)])
2372 }
2373 ",file="modelMcov.txt")
2374 data2<-list(y=ytot,nz=nz,nind=nind,K=K,xcent=xcent,maxD=11.5)
2375 params2<-list('p0','psi','N','beta')
2376 inits = function() {list(z=z, psi=psi, p0=runif(1),beta=rnorm(1) ) }
2377 fit2 = bugs(data2, inits, params2, model.file="modelMcov.txt",working.directory=getwd(),
2378           debug=T, n.chains=3, n.iter=4000, n.burnin=1000, n.thin=4)

```

Posterior summaries are given in Table ?? XYZ, and the posterior distribution of N is given in Fig. 3.5. It might be perplexing that the estimated N is much lower than obtained by model M_h but there is a good explanation for this, discussed subsequently. That issue notwithstanding, it is worth pondering how this model could be an improvement (conceptually or technically) over some other model/estimator including M_0 and M_h considered previously. Well, for one, we have accounted formally for heterogeneity due to spatial location of individuals relative to exposure to the trap array, characterized by the centroid of the array. Moreover, we have done so using a model that is based on an explicit mechanism, as opposed to a phenomenological one such as Model M_h . Moreover, importantly, using our new model, *the estimated N applies to an explicit area which is defined by our prescribed value of $\max D$* . That is, this area is a fixed component of the model and the parameter N therefore has explicit spatial context, as the number of individuals with home range centers less than $\max D$ from the centroid of the trap array. As such, the implied “effective trap area”¹³ for any $\max D$ is that of a circle with radius $\max D$.

```

2395 %% Not sure whether this should be a table or verbatim print-out
2396 \begin{table}
2397 \tabular{ccccccccc}
2398 Node statistics
2399 node mean sd MC error 2.5% median 97.5% start sample
2400 N 58.89 5.483 0.2199 50.0 58.0 71.0 251 2250
2401 beta -0.246 0.06087 0.003892 -0.3592 -0.2457 -0.126 251 2250
2402 deviance 459.4 13.29 0.4496 435.7 458.4 487.8 251 2250
2403 p0 0.5409 0.06817 0.004052 0.4072 0.544 0.6678 251 2250
2404 psi 0.1706 0.02572 7.759E-4 0.1247 0.1692 0.2242 251 2250
2405 \end{tabular}
2406 \caption{..... xyz .....}
2407 \end{table}

```

¹³This is a bad use of this term. We have never defined ETA or ESA. What is it, exactly?




Figure 3.5. Needs a caption

2408 `\label{tab.maxD}`

2409 We'll remake this figure in R. For now, insert it as is.

2410 3.5.2 Extension of the Model

2411 One important issue in understanding the meaning of estimates produced under
 2412 the individual covariate model is that the uniform distribution on maxD implies
 2413 that density is *not constant* over space. In particular, this model implies that it
 2414 *decreases* as we move away from the centroid of the trap array. This is one reason
 2415 we have a lower estimate of density than that obtained previously and also why,
 2416 if we were to increase maxD , we would see density continue to decrease: $x[i] \sim$
 2417 $\text{Uniform}(0, \text{maxD})$ implies constant N in each distance band from the centroid but
 2418 obviously the *area* of each distance band is increasing. The reader can verify this
 2419 as a homework exercise.

2420 Obviously, the use of an individual covariate model is *not* restricted to use of
 2421 this specific distribution for the individual covariate. Clearly, it is a bad choice and,
 2422 therefore, we should think about whether we can choose a better distribution for
 2423 maxD - one that doesn't imply a decreasing density as distance from the centroid
 2424 increases. Conceptually, what we want to do is impose a prior on distance from
 2425 the centroid, x , such that density is proportional to the amount of area in each
 2426 successive distance band as you move farther away from the centroid. In fact, there
 2427 is theory that exists which tells us what the correct distribution of x is $2x/\text{maxD}^2$.
 2428 This can be derived by noting that $F(x) = \text{Pr}(X < x) = \pi * x * x / \pi * \text{maxD} * \text{maxD}$.
 2429 Then, $f(x) = dF/dx = 2 * x / (\text{maxD} * \text{maxD})$. This might be called a
 2430 triangular distribution, I think, which makes sense because the incremental area
 2431 in each additional distance band increases linearly with radius (i.e., distance from
 2432 centroid). It is sometimes comforting to verify things empirically:

```

2433 > u<-runif(10000,-1,1)
2434 > v<-runif(10000,-1,1)
2435 > d<- sqrt(u*u+v*v)
2436 > hist(d[d<1])
2437 > hist(d[d<1],100)
2438 > hist(d[d<1],100,probability=TRUE)
2439 > abline(0,2)

```

It would be useful if we could describe this distribution in *BUGS but there is not a built-in way to do this. One possibility is to use a discrete version of the pdf. We might also be able to use what is referred to in WinBUGS jargon as the “zeros trick” (see Advanced BUGS tricks) although we haven’t pursued this approach. Instead, we consider using a discrete version and break D_{\max} into L distance classes of width δ , with probabilities proportional to $2 * x$. In particular, if the cut-points are $xg[1] = 0, xg[2], \dots, xg[L + 1] = D_{\max}$ and the interval midpoints are $xm[i] = xg[i + 1] - \delta$. Then, the interval probabilities are $p[i] = 2 * xm[i] * \delta / (D_{\max} * D_{\max})$, which we can compute once and then send them to WinBUGS as data.

The R script is as follows. In the model description the variable x (observed home range center) has been rounded so that the discrete version of the $f(x)$ can be used as described previously. The new variable labeled `xround` is actually then the integer category label in units of δ from 0. Thus, to convert back to distance in the expression for $lp[i]$, `xround[i]` has to be multiplied by δ .

```

2454 delta<-.2
2455 xround<-xcent/%delta + 1
2456 Dgrid<- seq(delta,maxD,delta)
2457 xprobs<- delta*(2*Dgrid/(maxD*maxD))
2458 xprobs<-xprobs/sum(xprobs)
2459
2460 cat("
2461 model{
2462   p0 ~ dunif(0,1)          # prior distributions
2463   mup<- log(p0/(1-p0))
2464   psi~dunif(0,1)
2465   beta~dnorm(0,.01)
2466
2467   for(i in 1:(nind+nz)){
2468     xround[i]~dcat(xprobs[])
2469     z[i]~dbern(psi)          # zero inflation variables
2470     lp[i] <- mup + beta*xround[i]*delta # individual effect
2471     logit(p[i])<-lp[i]
2472     mu[i]<-z[i]*p[i]
2473     y[i]~dbin(mu[i],K)      # observation model
2474   }

```

```

2475
2476 N<-sum(z[1:(nind+nz)])
2477 }
2478 ",file="modelMcov.txt")

```

2479 To fit the model we do this - keeping in mind that the data objects required
 2480 below have been defined in previous analyses of this chapter:

```

2481 data2<-list(y=ytot,nz=nz,nind=nind,K=T,xround=xround,xprobs=xprobs,delta=delta)
2482 params2<-list('p0','psi','N','beta')
2483 inits = function() {list(z=z, psi=psi, p0=runif(1),beta=rnorm(1) ) }
2484 fit = bugs(data2, inits, params2, model.file="modelMcov.txt",working.directory=getwd(),
2485           debug=FALSE, n.chains=3, n.iter=11000, n.burnin=1000, n.thin=2)

```

2486 This is a useful model because it induces a clear definition of area in which
 2487 the population of N individuals reside. Under this model, that area is defined by
 2488 specification of D_{max} . We can apply the model for different values of D_{max} and
 2489 observe that the estimated N varies with D_{max} . Fortunately, we see empirically,
 2490 that while N seems highly sensitive to the prescribed value of D_{max} , density seems
 2491 to be invariant to D_{max} as long as it is chosen to be sufficiently large. We fit the
 2492 model for $D_{max} = 12$ (points in close proximity to the trap arra) to 20 for and the
 2493 results are given in Table ??.

```

2494 \begin{table}
2495 \caption{Table: Analysis of Fort Drum bear hair snare data using the individual covariate }
2496 \begin{tabular}{cccc}
2497     maxD   mn    SD
2498 [1,]   12 0.230 0.038
2499 [2,]   15 0.244 0.041
2500 [3,]   17 0.249 0.044
2501 [4,]   18 0.249 0.043
2502 [5,]   19 0.250 0.043
2503 [6,]   20 0.250 0.044
2504 \end{tabular}
2505 \end{table}

```

2506 We see that the posterior mean and SD of density (individuals per square km)
 2507 appear insensitive to choice of $maxD$ once we get a slight ways away from the
 2508 maximum observed value of about 11.5. The estimated density of 0.250 per km² is
 2509 actually quite a bit lower than we reported using model Mh (0.37, see section XYZ
 2510 above) for which sample area is not an explicit feature of the model. On the other
 2511 hand it is higher than that reported from Model M0 using the buffered area (0.195).
 2512 There is no basis really for comparing or contrasting these various estimates and
 2513 it would be a useful philosophical exercise for the reader to discuss this matter.

In particular, application of model M0 and Mh are distinctly *not* spatially explicit models – the area within which the population¹⁴ resides is not defined under either model. There is therefore no reason at all to think that the estimates produced under either model, using a buffered area, are justifiable based on any theory. In fact, we would get exactly the same estimate of N no matter what we declare the area to be. On the other hand, the individual covariate model explicitly describes a distribution for “distance from centroid” that is a reasonable and standard null model - it posits, in the absence of direct information, that individual home range centers are randomly distributed in space and that probability of detection depends on the distance between home range center and the centroid of the trap array. Under this definition of the system, we see that density is invariant to the choice of sample area which seems like a desirable feature. The individual covariate model is not ideal, however, because it does not make full use of the spatial information in the data set, i.e., the trap locations and the locations of each individual encounter.

3.5.3 Invariance of density to maxD

Under the model above, and also under models that we consider in later chapters, a general property of the estimators is that while N increases with the prescribed trap area (equivalent to maxD in this case), we expect that density estimators should be invariant to this area. In the model used above, we note that $Area(maxD) = \pi * maxD * maxD$ and $E[N(maxD)] = \lambda * A(maxD)$ and thus $E[Density(maxD)] = \lambda$ which is constant. This should be interpreted as the *prior* density. Absent data, then realizations under the model will have density λ regardless of what $maxD$ is prescribed to be. As we verified empirically above, the posterior density is also invariant if $maxD$ as long as the implied area (implied by maxD) is large enough so that the data no longer provide information about density (i.e., “far away”), then our estimator of density should become insensitive.

3.5.4 Toward Fully Spatial Capture-recapture Models

We developed this model for the average observed location and equated it to home range center s_i . Intuitively, taking the average encounter location as an estimate of home range center makes sense but more so when the trapping grid is dense and expansive relative to typical home range sizes. However, our approach also ignored the variable precision with which each $s[i]$ is estimated and also, as noted previously, estimates of $s[i]$ around the “edge” (however we define that) are biased because the observations are truncated (we can only observe locations within the trap array). In the next Chapter we provide a further extension of this individual covariate model that definitively resolves the ad hoc nature of the individual covariate approach we took here. In that model we build a model in which $s[i]$ are regarded as latent

¹⁴We need to look back at Chapter 1 and make sure we quit calling this “sample area” - it really isn’t that at all, but rather the area within which N resides.

2551 variables and the observation locations (i.e., trap specific encounters) are linked
 2552 to those latent variables with an explicit model. We note that the model fitted
 2553 previously could be adapted easily to deal with s_i as a latent variable, simply by
 2554 adding a prior distribution for s_i . The reader should contemplate how to do this
 2555 in WinBUGS.

3.6 DISTANCE SAMPLING: A PRIMITIVE SPATIAL CAPTURE-RECAPTURE MODEL

2556 Distance sampling is one of the most popular methods for estimating animal abun-
 2557 dence. One of the great benefits of distance sampling is that it provides explicit
 2558 estimates of *density*. The distance sampling model is a special case of a closed
 2559 population model with a covariate. The covariate in this case, x_i , is the distance
 2560 between an individual's location “ u ” and the observation location or transect. In
 2561 fact, the model underlying distance sampling is precisely the same model as that
 2562 which applies to the individual-covariate models, except that observations are made
 2563 at only $K = 1$ sampling occasion. In a sense, distance sampling is a spatial capture-
 2564 recapture model, but without the “recapture.” This first and most basic spatial
 2565 capture-recapture model has been used routinely for decades and, formally, it is a
 2566 spatially-explicit model in the sense that it describes, explicitly, the spatial organi-
 2567 zation of individual locations (although this is not always stated explicitly) and, as
 2568 a result, somewhat general models of how individuals are distributed in space can
 2569 be specified (Royle et al., 2004; Johnson, 2010; Sillett, 2011). As before, the dis-
 2570 tance sampling model, under data augmentation, includes a set of M zero-inflation
 2571 variables z_i and the binomial model expressed conditional on z (binomial for $z = 1$,
 2572 and fixed zeros for $z = 0$). In distance sampling we pay for having only a single
 2573 sample (i.e., $K = 1$) by requiring constraints on the model of detection probability.
 2574 A standard model is

$$\log(p_i) = b * x_i^2$$

2575 for $b < 0$, where x_i denotes the distance at which the i th individual is detected
 2576 relative to some reference location where perfect detectability ($p = 1$) is assumed.
 2577 This function corresponds to the “half-normal” detection function (i.e., with $b =$
 2578 $1/\sigma^2$). If $K > 1$ then the intercept alpha is identifiable and such models are
 2579 usually called “capture-recapture distance sampling” (Borchers, missing) and others
 2580 XYZ???.

2581 As with previous examples, we require a distribution for the individual covariate
 2582 x_i . The customary choice is

$$x_i \sim \text{Uniform}(0, B)$$

2583 wherein $B > 0$ is a known constant, being the upper limit of data recording by the
 2584 observer (i.e., the point count radius, or transect half-width). In practice, this is
 2585 sometimes asserted to be infinity, but in such cases the distance data are usually

truncated. Specification of this distance sampling model in the BUGS language is shown in Panel 3.1. Royle and Dorazio (2008), p. xyz) provide a distance sampling example analyzed by DA using the famous Impala data.

```

b~dunif(0,10)
psi~dunif(0,1)

for(i in 1:(nind+nz)){
  z[i]~dbern(psi)      # DA Variables
  x[i]~dunif(0,B)      # B=strip width
  p[i]<-exp(logp[i])    # DETECTION MODEL
  logp[i]<- -((x[i]*x[i])*b)
  mu[i]<-z[i]*p[i]
  y[i]~dbern(mu[i])    # OBSERVATION MODEL
}
N<-sum(z[1:(nind+nz)])
D<- N/striparea # area of transects

```

Panel 3.1: Distance sampling model in WinBUGS, using a “half-normal” detection function.

As with the individual covariate model in the previous section, the distance sampling model can be equivalently specified by putting a prior distribution on individual *location* instead of distance between individual and observation point (or transect). Thus we can write the general distance sampling model as

$$\text{logit}(p[i]) = \alpha + \beta * ||u[i] - x_0||$$

Along with

$$\mathbf{u}_i \sim \text{Uniform}(\mathcal{S})$$

where x_0 is a fixed point (or line) and $u[i]$ is the individual’s location which is observable for n individuals. In practice it is easier to record distance instead of location. Basic math can be used to argue that if individuals have a uniform distribution in space, then the distribution of Euclidean distance is also uniform. In particular, if a transect of length L is used and x is distance to the transect then $F(x) = \text{Pr}(X \leq x) = L * x/L * B = x/B$ and $f(x) = dF/dx = (1/B)$. For measurements of radial distance, see the previous section.

In the context of our general characterization of SCR models (chapter 1.XYZ), we suggested that every SCR model can be described, conceptually, by a hierarchical model of the form:

$$[y|u][u|s][s].$$

Distance sampling ignores s , and treats u as observed data¹⁵. Thus, we are left

¹⁵Formally we could also say that $[u] = \int [y|s][s]ds$

2605 with

$$[y|u][u].$$

2606 In contrast, as we will see in the next chapters, basic SCR models (chapter 4) ignore
2607 u and condition on s , which is not observed:

$$[y|s][s]$$

2608 Since $[u]$ and $[s]$ are both assumed to be uniformly distributed, these are structurally
2609 equivalent models! The main differences have to do with interpretation of model
2610 components and whether or not the latent variables are observable (in distance
2611 sampling they are).

2612 So why bother with SCR models when distance sampling yields density esti-
2613 mates and accounts for spatial heterogeneity in detection? For one, imagine try to
2614 collect distance sampling data on tigers! Clearly, distance sampling requires that
2615 one can collect large quantities of distance data, which is not always possible. For
2616 tigers, it is much easier, efficient, and safer to employ camera traps or tracking
2617 plates and then apply SCR models. Furthermore, as we will see in Ch XYZ, SCR
2618 models can use distance data to estimate all the parameters of our enchilada, al-
2619 lowing us to study distribution, movement, and density. Thus, SCR models are
2620 much more flexible than distance sampling models, and can accommodate data
2621 from virtually all animal survey designs.

2622 3.6.1 Example: Muntjac deer survey from Nagarahole, India

2623 Here we fit distance sampling models to distance sampling data on the muntjac deer
2624 (*Muntiacus muntjak*) collected in the year 2004 from Nagarahole National Park in
2625 southern India (Kumar, missing)(Kumar et al. unpublished data). The muntjac
2626 is a solitary species and distance measurements were made on 57 groups that were
2627 largely singletons with XYZ pairs of individuals. Commands for reading in and
2628 organizing the data for WinBUGS, followed by writing the model to a text file.
2629 Note that the total sampled area of the transects is fed in as “striparea” which is
2630 708 (km of transect) multiplied by the strip width ($B=150 = 0.15$ km) multiplied
2631 by 2.

```
2632 library("R2WinBUGS")
2633 data<- read.csv("Muntjac.csv")
2634 nind<-nrow(data)
2635 y<-rep(1,nind)
2636 nz<-400
2637 y<-c(y,rep(0,nz))
2638 x<-data[,3]
2639 x<-c(x,rep(NA,nz))
2640 z<-y
```

```

2641 data<-list(y=y,x=x,nz=nz,nind=nind,B=150,striparea=708*.15*2)
2642
2643 cat("
2644 model{
2645   b~dunif(0,10)
2646   psi~dunif(0,1)
2647
2648   for(i in 1:(nind+nz)){
2649     z[i]~dbern(psi)    # DA Variables
2650     x[i]~dunif(0,B)    # B=strip width
2651     p[i]<-exp(logp[i])  # DETECTION MODEL
2652     logp[i]<- -((x[i]*x[i])*b)
2653     #logp[i]<- -b*log(x[i]+1)
2654     mu[i]<-z[i]*p[i]
2655     y[i]~dbern(mu[i])  # OBSERVATION MODEL
2656   }
2657   N<-sum(z[1:(nind+nz)])
2658   D<- N/striparea    # area of transects
2659 }
2660 ",file="dsamp.txt")

```

Next, we provide inits, indicate which parameters to monitor, and then pass those things to WinBUGS:

```

2663 params<-list('b','N','D','psi')
2664 inits = function() {list(z=z, psi=runif(1), b=runif(1,0,.02) )}
2665 fit = bugs(data, inits, params, model.file="dsamp.txt",
2666 working.directory=getwd(),debug=T, n.chains=3, n.iter=4000, n.burnin=1000, n.thin=2)

```

Posterior summaries are provided in the following table. Estimated density is pretty low, 1.1 individuals per sq. km.¹⁶

```

2669 node mean sd MC error 2.5% median 97.5% start sample
2670 D 1.096 0.1694 0.009122 0.8098 1.078 1.474 501 4500
2671 N 232.8 35.99 1.938 172.0 229.0 313.0 501 4500
2672 b 5.678E-4 1.05E-4 4.129E-6 3.867E-4 5.616E-4 7.949E-4 501 4500
2673 deviance 681.2 16.72 0.7536 650.8 680.6 716.6 501 4500
2674 psi 0.5099 0.08238 0.004442 0.3681 0.5033 0.6918 501 4500

```

3.7 SUMMARY AND OUTLOOK

Traditional closed population capture-recapture models are closely related to binomial generalized linear models. Indeed, the only real distinction is that in capture-

¹⁶ much lower than Samba's : Observers walked about 708 km from 39 transects in Nagarahole and the muntjac density is about 3 per sq km.. I need to get to the bottom of this.

recapture models, the population size parameter N (corresponding also to the size of a hypothetical “complete” data set) is unknown. This requires special consideration in the analysis of capture-recapture models. The classical approach to inference recognizes that the observations don’t have a standard binomial distribution but, rather, a truncated binomial (from which the so-called “conditional likelihood” derives) since we only have encounter frequency data on observed individuals. If instead we analyze the models using data augmentation, the observations can be modeled using a zero-inflated binomial distribution. In short, when we deal with the unknown- N problem using data augmentation then we are left with zero-inflated GLM and GLMMs instead of ordinary GLM or GLMMs. The analysis of such zero-inflated models is practically convenient, especially using the various Bayesian analysis packages that use the BUGS language.

Spatial capture-recapture models that we will consider in the rest of the chapters of this book are closely related to what have been called individual covariate models. Heuristically, spatial capture-recapture models arise by defining individual covariates based on observed locations of individuals – we can think of using some function of mean encounter location as an individual covariate. We did this in a novel way, by using distance to the centroid of the trapping array as a covariate. We analyzed the “full likelihood” using data augmentation, and placed a prior distribution on the individual covariate which was derived from an assumption that individual locations are, a priori, uniformly distributed in space. This assumption provides for invariance of the density estimator to the choice of population size area (induced by maximum distance from the centroid of the). The model addressed some important problems in the use of closed population models: it allows for heterogeneity in encounter probability due to the spatial context of the problem and it also provides a direct estimate of density because area is a feature of the model (via the prior on the individual covariate). The model is still not completely general because the model does not make use of the fully spatial encounter histories, which provide direct information about the locations and density of individuals. A specific individual covariate model that is in widespread use is classical “distance sampling.” The model underlying distance sampling is precisely a special kind of SCR model - but one without replicate samples. Understanding distance sampling and individual covariate models more broadly provides a solid basis for understanding and analyzing spatial capture-recapture models.

4

FULLY SPATIAL CAPTURE-RECAPTURE MODELS

In previous sections we discussed some classes of models that could be viewed as primitive spatial capture-recapture models. We looked at a basic distance sampling model and we also considered a classical individual covariate modeling approach in which we defined a covariate to be the distance from (estimated) home range center to the center of the trap array. These were spatial in the sense that they included some characterization of where individuals live but, on the other hand, only a primitive or no characterization of trap location. That said, very little distinguishes these two models from spatial capture-recapture models that we consider in this chapter which fully recognize the spatial attribution of both individual animals *and* the locations of encounter devices.

Fully spatial capture-recapture models must accommodate the spatial organization of individuals and the encounter devices because the encounter process occurs at the level of individual traps. Failure to consider the trap-specific collection of data is the key deficiency with classical ad-hoc approaches which aggregate encounter information to the resolution of the entire trap array. We have seen previously some problems that this induces - imbalance in trap-level effort over time is problematic, and not being able to deal with trap-specific behavioral responses. Here, we resolve that by developing what is basically an individual covariate model but operating at the level of traps. That is, we develop our first fully spatial capture-recapture model which turns out to be precisely the model considered in section 3.XXX but instead of defining the individual covariate to be distance to centroid of the array we define J individual covariates - the distance to *each* trap. And, instead of using estimates of individual locations \mathbf{s} , we consider a fully hierarchical model in which we regard \mathbf{s} as a latent variable and impose a prior distribution on

it. We can think of having J independent capture-recapture studies generating one data set for each trap, and applying the individual covariate model with random activity centers, and that is all the basic SCR model is.

In the following sections of this chapter we investigate the basic spatial capture-recapture model and address some important considerations related to its analysis in **WinBUGS**. We also demonstrate how to summarize posterior output for the purposes of producing density maps or spatial predictions of density.

4.1 SAMPLING DESIGN AND DATA STRUCTURE

In our development here, we will assume a standard sampling design in which an array of J traps is operated for K time periods (say, nights) producing encounters of n individuals. Because sampling occurs by traps and also over time, the most general data structure yields encounter histories for *each individual* that are temporally *and* spatially indexed. Thus a typical data set will include an encounter history *matrix* for each individual. For the most basic model, there are no time-varying covariates that influence encounter, there are no explicit individual-specific covariates, and there are no covariates that influence density we will develop models in this chapter for encounter data that are aggregated over the temporal replicates. For example, suppose we observe 6 individuals in sampling at 4 traps over 3 nights of sampling then a plausible data set is the 6×4 matrix of encounters, out of 3, of the form:

	trap1	trap2	trap3	trap4
[1,]	1	0	0	0
[2,]	0	2	0	0
[3,]	0	0	0	1
[4,]	0	1	0	0
[5,]	0	0	1	1
[6,]	1	0	1	0

We develop models in this chapter for devices such as “hair snares” or other DNA sampling methods (Kéry et al., 2010; Gardner et al., 2010) and related types of sampling problems so that we can suppose that “traps” may capture any number of individuals and an individual may be captured in any number of traps during each occasion but individuals can be encountered at most 1 time in a trap during any occasion. Thus, this is a “multi-catch” type of sampling (?, p. xyz). The statistical assumptions are that individual encounters within and among traps are independent. These basic (but admittedly at this point somewhat imprecise) assumptions define the basic spatial capture-recapture model, which we will refer to as “SCR0” henceforth¹ so that we may use that model as a point of reference

¹RC: It would be nice to have a running series of figures to display the various types of models. Each figure could have the same set of traps, use the same symbols, etc... It’s probably worth showing example data (and latent variables) in a table too

without having to provide a long-winded enumeration of assumptions and sampling design each time we do. We will make things more precise as we develop a formal statistical definition of the model shortly.

While the model is mostly directly relevant for hair snares and other DNA sampling methods for which multiple detections of an individual are not distinguishable, we will also make use of the model for data that arise from camera-trapping studies. In practice, with camera trapping, individuals might be photographed several times in a night but we will typically distill such data into a single binary encounter event for reasons discussed later in chapter 6.

4.2 THE BINOMIAL OBSERVATION MODEL

We assume that the individual and trap-specific encounters, y_{ij} , are mutually independent outcomes of a binomial random variable:

$$y_{ij} \sim \text{Bin}(K, p_{ij}) \quad (4.2.1)$$

This is the basic model underlying “logistic regression” (chapter 2) as well as standard closed population models (chapter 3). The key element of the model is that the encounter probability p_{ij} is indexed by (i.e., depends on) both individual and trap. In a sense, then, we can think of each *trap* as producing individual level encounter history data of the classical variety - an $n_{\text{ind}} \times n_{\text{rep}}$ matrix of 0's and 1's (this is the “encountered at most 1 time” assumption).

As we did in section XXX.YYY, we will make explicit the notion that p_{ij} is defined conditional on “where” individual i lives. Naturally, we think about defining an individual home range and then relating p_{ij} explicitly to the centroid of the individual's home range, or its center of activity (Efford, 2004; Borchers and Efford, 2008; Royle and Young, 2008). Therefore, define \mathbf{s}_i , a two-dimensional spatial coordinate, to be the activity center for individual i . Then, the basic SCR model postulates that encounter probability, p_{ij} , is related by a decreasing function to distance between trap j , having location \mathbf{x}_j , and \mathbf{s}_i . Naturally, if we think of modeling binomial counts using logistic regression, we might specify the model according to:

$$\text{logit}(p_{ij}) = \alpha_0 + \theta * ||\mathbf{s}_i - \mathbf{x}_j|| \quad (4.2.2)$$

where, here, $||\mathbf{s}_i - \mathbf{x}_j||$ is the distance between \mathbf{s}_i and \mathbf{x}_j . We sometimes write $||\mathbf{s}_i - \mathbf{x}_j|| = \text{dist}(\mathbf{s}_i, \mathbf{x}_j) = d_{ij}$. Alternatively, if we think about distance sampling then we might use the “half-normal” model of the form:

$$p_{ij} = p_0 * \exp(-\theta * ||\mathbf{s}_i - \mathbf{x}_j||^2)$$

Or any of a large number of standard detection models that are commonly used (we consider more in chapter XYZ). The half-normal model implies

$$\log(p_{ij}) = \log(p_0) - \theta * ||\mathbf{s}_i - \mathbf{x}_j||^2 \quad (4.2.3)$$

Whatever model encounter probability we choose, we should always keep in mind that the model is described conditional on \mathbf{s}_i , which is an unobserved random variable. Thus, to be precise about this, we should write the observation model as

$$y_{ij}|\mathbf{s}_i \sim \text{Bin}(K, p(\mathbf{s}_{ij}; \theta))$$

Note that we probably expect that the parameter θ in Eq. 4.2.2 or 4.2.3 should be negative, so that the probability of encounter decreases with distance between the trap and individual home range center. The joint likelihood for the data, conditional on the collection of individual activity centers, can therefore be expressed as

$$\mathcal{L}(\theta|\{\mathbf{y}_i, \mathbf{s}_i\}_{i=1}^N) = \prod_i \prod_j \text{Bin}(y_{ij}|p_{ij}(\theta))$$

Which, if we switch the indices on the product operators, this shows the SCR likelihood (conditional on \mathbf{s}) to be the product of J *independent* capture-recapture likelihoods - one for each trap. However, the data have a “repeated measures” type of structure, with each of the j likelihood contributions for each individual being grouped by individual. Thus, we cannot analyze the model meaningfully by J trap-specific models. In classical repeated measures types of models, we accommodate the group structure of the data using random effects (random individual or group level variables). For SCR models we take the same basic approach, which we develop subsequently.

4.2.1 Distance as a latent variable

If we knew precisely every \mathbf{s}_i in the population (and how many, N), then the model specified by eqs. 4.2.1 and 4.2.2 or 4.2.3 is just an ordinary logistic regression type of a model which we learned how to fit using **WinBUGS** previously (chapt. 2), with a covariate d_{ij} . However, the activity centers are unobservable even in the best possible circumstances. In that case, d_{ij} is an unobserved variable, analogous to classical “random effects” models. We need to therefore extend the model to accommodate these random variables with an additional model component. A standard, and perhaps not unreasonable, assumption is the so-called “uniformity assumption” which is to say that the \mathbf{s}_i are uniformly distributed over space (the obvious next question “which space?” is addressed below). This uniformity assumption amounts to a uniform prior distribution on \mathbf{s}_i , i.e., the pdf of \mathbf{s}_i is constant, which we may express

$$\Pr(\mathbf{s}_i) \text{propto} \text{const} \quad (4.2.4)$$

To summarize the preceeding model developing, a basic SCR model is defined by 3 essential components:

- (1) Observation model: $y_{ij}|\mathbf{s}_i \sim \text{Bin}(K, p_{ij})$
- (2) Encounter probability: $\text{logit}(p_{ij}) = \alpha_0 + \theta * ||\mathbf{s}_i - \mathbf{x}_j||$

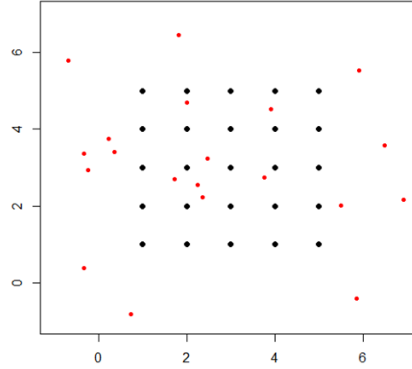


Figure 4.1. Realization of a binomial point process

2840 (3) Point process model: $\Pr[\mathbf{s}_i] \propto \text{const}$

2841 Therefore, the SCR model is little more than an ordinary capture-recapture model
 2842 for closed populations. It is such a model, but augmented with a set of “individual
 2843 effects”, \mathbf{s}_i , which relate some sense of individual location to encounter probability.
 2844 As it turns out, assumption (3) is usually not precise enough to fit a model in
 2845 practice for reasons we discuss in the following section. We will give another way to
 2846 represent this prior distribution that is more concrete, but it depends on specifying
 2847 the “state-space” of the random variable \mathbf{s}_i . The term “state-space” is a technical
 2848 way of saying “possible outcomes”.

4.3 THE BINOMIAL POINT-PROCESS MODEL

2849 The collection of individual activity centers $\mathbf{s}_1, \dots, \mathbf{s}_N$ represent a realization of a
 2850 *binomial point process* (Illian, 2008a, p. xyz). The binomial point process (BPP)
 2851 is analogous to a Poisson point process in the sense that it represents a “random
 2852 scatter” of points in space - except that the total number of points is *fixed*, whereas,
 2853 in a Poisson point process it is random (having a Poisson distribution). As an
 2854 example, we show in Fig. 4.1 locations of 20 individual activity centers (black
 2855 dots) in relation to a grid of 25 traps. For a Poisson point process the number of
 2856 such points in the prescribed state-space would be random whereas often we will
 2857 simulate fixed numbers of points, e.g., for evaluating the performance of procedures
 2858 such as how well does our estimator perform of $N = 50$?

It is natural to consider a binomial point process in the context of capture-recapture models because it preserves N in the model and thus preserves the linkage directly with closed population models. In fact, under the binomial point process model then Model M0 and other closed models are simple limiting cases of SCR models. In addition, use of the BPP model allows us to use data augmentation for Bayesian analysis of the models as in chapter 3, thus yielding a methodologically coherent approach to analyzing the different classes of models. Despite this, making explicit assumptions about N , such as Poisson, is convenient in some cases (see chapt. XYZ).

One consequence of having fixed N , in the BPP model, is that the model is not strictly a model of “complete spatial randomness”. This is because if one forms counts $n(A_1), \dots, n(A_k)$ in any set of disjoint regions say A_1, \dots, A_k , then these counts are *not* independent. In fact, they have a multinomial distribution (see Illian, 2008a, p. XYZ). Thus, the BPP model introduces a slight bit of dependence in the distribution of points. However, in most situations this will have no practical effect on any inference or analysis and, as a practical matter, we will usually regard the BPP model as one of spatial independence among individual activity centers because each activity center is distributed independently of each other activity center. Despite this implicit independence we see in Fig. 4.1 that realizations of randomly distributed points will typically exhibit distinct non-uniformity. Thus, independent, uniformly distributed points will almost never appear regularly, uniformly or systematically distributed. For this reason, the basic binomial (or Poisson) point process models are enormously useful in practical settings. More relevant for SCR models is that we actually have a little bit of data for some individuals and thus the resulting posterior point pattern can deviate strongly from uniformity (we should note this elsewhere too). The uniformity hypothesis is only a *prior* distribution which is directly affected by the quantity and quality of observations.

4.3.1 Definition of home range center

Some will be offended by our use of the concept of “home range center” and thus will have difficulty in believing that the resulting model is really useful for anything. Indeed, the idea of a home range or activity center is a vague concept anyway, a purely phenomenological construct. Despite this, it doesn’t really matter whether or not a home range makes sense for a particular species - individuals of any species inhabit *some* region of space and we can define the “home range center” to be the center of the space that individual was occupying (or using) during the period in which traps were active. Thinking about it in that way, it could even be observable (almost) as the centroid of a very large number of radio fixes over the course of a survey period or a season. Thus, this practical version of a home range center is a well-defined construct regardless of whether one thinks the home range concept is meaningful, even if individuals are not particularly territorial. This is why we usually use the term “activity center” or maybe even “centroid of space usage”

and we recognize that this construct is a transient thing which applies only to a well-defined period of study.

4.3.2 The state-space of the point process

Shortly we will focus on Bayesian analysis of this model with N known so that we can directly apply what we learned in chapter 2 to this situation. To do this, we note that the individual effects $\mathbf{s}_1, \dots, \mathbf{s}_N$ are unknown quantities and we will need to be able to simulate each \mathbf{s}_i in the population from the posterior distribution. It should be self-evident that we cannot simulate the \mathbf{s}_i unless we describe precisely the region over which those \mathbf{s}_i 's are uniformly distributed. This is the quantity referred to above as the state-space, denoted henceforth by \mathcal{S} , which is a region or a set of points comprising the potential values of \mathbf{s}_i . Thus, an equivalent explicit statement of the “uniformity assumption” is

$$\mathbf{s}_i \sim \text{Unif}(\mathcal{S})$$

Prescribing the state-space

Evidently, we need to define the state-space, \mathcal{S} . How can we possibly do this objectively? Prescribing any particular \mathcal{S} seems like the equivalent of specifying a “buffer” which we criticized previously as being ad hoc. How is it that choosing a state-space is *not* ad hoc? As a practical matter, it turns out that estimates of density are insensitive to choice of the state-space. As we observed in chapter 3, it is true that N increases with \mathcal{S} , but only at the same rate as \mathcal{S} under the prior assumption of constant density. As a result, we say that density is invariant to \mathcal{S} as long as \mathcal{S} is sufficiently large. Thus, while choice of \mathcal{S} is (or can be) essentially arbitrary, once \mathcal{S} is chosen, it defines the population being exposed to sampling, which scales appropriately with the size of the state-space.

For our simulated system developed previously in this chapter, we defined the state space to be a square within which our traps were centered perfectly. For many practical situations this might be an acceptable approach to defining the state-space. We provide an example of this in section 4.7 below in which the trap array is irregular and also situated within a realistic landscape that is distinctly irregular. In general, it is most practical to define the state-space as a regular polygon (e.g., rectangle) containing the trap array without differentiating unsuitable habitat. Although defining the state-space to be a regular polygon has computational advantages (e.g., we can implement this more efficiently in **WinBUGS** and cannot for irregular polygons), a regular polygon induces an apparent problem of admitting into the state-space regions that are distinctly non-habitat (e.g., oceans, large lakes, ice fields, etc.). It is difficult to describe complex sets in mathematical terms that can be admitted to this spatial model. As an alternative, we can provide a representation of the state-space as a discrete set of points (section 4.9) that will allow specific points to be deleted or not depending on whether they represent habitat, or we can define the state-space as an intersection of polygons, and analysis of

models with state-space defined in that way can be analyzed easily using MCMC (see section XYZ in chapt. 6). In what follows below we provide an analysis of the camera data defining the state-space to be a regular continuous polygon (a rectangle).

4.3.3 Invariance and the State-space as a model assumption

We will assert for all models we consider in this book that density is invariant to the size and extent of \mathcal{S} , if \mathcal{S} is sufficiently large. In fact, this only holds as long as our model relating p_{ij} to \mathbf{s}_i is a decreasing function of distance. We can prove this thinking about a 1-d case where $E[y]$ for the “last cell” (i.e., for $d > B$ for B large enough) is 0. So it always contributes nothing to the likelihood, i.e., $E[n(\text{lastcell})] = 0$. [sketch out a proof of this], in regular situations in which the detection function decays monotonically with distance and prior density is constant. Sometimes our estimate of density can be influenced if we make \mathcal{S} too small but this might be sensible if \mathcal{S} is naturally well-defined. As we discussed in chapter 1, **choice of \mathcal{S} is part of the model and thus it makes sense that estimates of density might be sensitive to its definition in problems where it is natural to restrict \mathcal{S} .** One could imagine however that in specific cases where you’re studying a small population with well-defined habitat preferences that a problem could arise because changing the state-space around based on differing opinions and GIS layers really changes the estimate of total population size. But this is a real biological problem and a natural consequence of the spatial formalization of capture-recapture models - a feature, not a bug or some statistical artifact - and it should be resolved with better information and research, and not some arbitrary statistical artifact. For situations where there is not a natural choice of \mathcal{S} , we should default to choosing \mathcal{S} to be very large in order to achieve invariance or otherwise evaluate sensitivity of density estimates by trying a couple of different values of \mathcal{S} . This is a standard “sensitivity to prior” argument that Bayesians always have to be conscious of. We demonstrate this in our analysis of section 4.7 below. Note that $area(\mathcal{S})$ affects data augmentation. If you increase $area(\mathcal{S})$ then there are more individuals to account for and therefore the size of the augmented data set M must increase.

We have been told that one can carry-out non-Bayesian analyses of SCR models without having to specify the state-space of the point process or perhaps while only specifying it imprecisely. This assertion is incorrect. We assume people are thinking this because *they* don’t have to specify it explicitly because someone else has done it for them in a package that does integrated likelihood. Even to do integrated likelihood (see chapter 6) we have to integrate the conditional-on- \mathbf{s} likelihood over some 2-dimensional space. It might work that the integration can be done from $-\infty$ to $+\infty$ but that is a mathematical artifact of specific detection functions, and an implicit definition of a state-space that doesn’t make biological sense, even though it may in fact be innocuous;

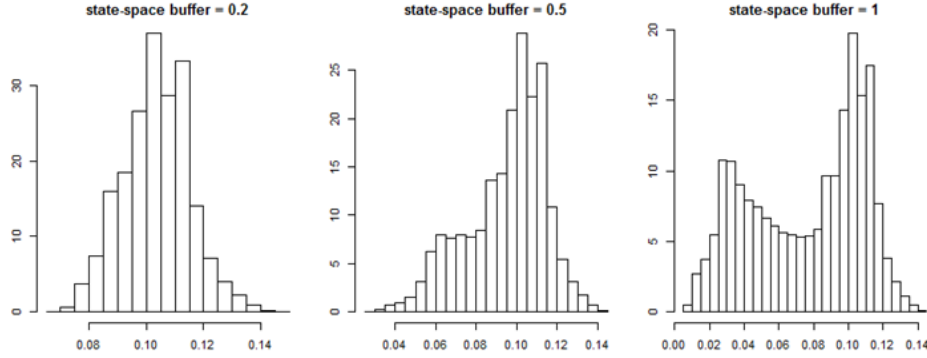


Figure 4.2. Needs a caption

4.3.4 Connection to Model Mh

SCR models are closely related to heterogeneity models. In SCR models, heterogeneity in encounter probability is induced by both the effect of distance in the model for detection probability and also from specification of the state-space. Clearly then the state-space is explicitly part of the model. To understand this, we have a random effect with some prior distribution:

$$\mathbf{s} \sim \text{uniform}(\mathcal{S})$$

And $p(\mathbf{s}) = p(y = 1|\mathbf{s})$ is some function of \mathbf{s} . Therefore, for any specific $g(p)$ and \mathcal{S} we can work out what the implied heterogeneity model is for example, the mean, variance or other moments of the population distribution of p can be evaluated by integrating $p(\mathbf{s})$ over the state-space of \mathbf{s} . Obviously the choice of $p(\mathbf{s})$ and the choice of \mathcal{S} interact to determine the effective heterogeneity in p . We show an illustration in Fig. 4.2 below which shows a histogram of p for a hypothetical population of 100000 individuals on a state-space enclosing our 5×5 trap array above, under the logistic model for distance. **R** code is provided in the **R** package **scrbook** to produce this analysis for the logistic and half-normal models. The histogram shows the encounter probability under buffers of 0.2, 0.5 and 1.0. We see the mass shifts to the left as the buffer increases, implying more individuals in the population but with lower encounter probability as their home range centers increase in distance from the trap array.

Another way to understand this is by representing \mathcal{S} as a set of discrete points on a grid. In the coarsest possible case where \mathcal{S} is a single arbitrary point, then every individual has exactly the same p . As we increase the number of points in \mathcal{S} then more distinct values of p are possible. As such, when \mathcal{S} is characterized by discrete points then SCR models are precisely a type of finite-mixture model (Norris III and

Pollock, 1996; Pledger, 2000), except where we have some information about which group an individual belong (i.e., where their activity center is), as a result of their captures in traps.

This context suggests the problem raised by Link (2003). He showed that in most practical situations N may not be identifiable across classes of mixture distributions which in the context of SCR models is the pair (g, \mathcal{S}) . The difference, however, is that we do obtain some direct information about \mathbf{s} in SCR models and therefore N is identifiable across models characterized by (g, \mathcal{S}) .

4.3.5 Connection to Distance Sampling

It is worth emphasizing that the basic SCR model is a binomial encounter model in which distance is a covariate. As such, it is striking similarity to a classical distance sampling model. Both have distance as a covariate but in classical distance sampling problems the focus is on the distance between the observer and the animal at an instant in time, not the distance between a trap and an animal's home range center. Thus in distance sampling, "distance" is *observed* for those individuals that appear in the sample. Conversely, in SCR problems, it is only imperfectly observed (we have partial information in the form of trap observations). Clearly, it is preferable to observe distance if possible, but as we will discuss in chapter XYZ, distance sampling requires field methods that are often not practical in many situations, e.g. when surveying tigers. Furthermore, SCR models allow us to relax many of the assumption made in classical distance sampling, and SCR models allow for estimates of quantities other than density, such as home range size.

4.4 SIMULATING SCR DATA

It is always useful to simulate data because it allows you to understand the system that you're modeling and also calibrate your understanding with the parameter values of the model. That is, you can simulate data using different parameter values until you obtain data that "looks right" based on your knowledge of the specific situation that you're interested in. Here we provide a simple script to illustrate how to simulate spatial encounter history data. In this exercise we simulate data for 100 individuals and a 25 trap array laid out in a 5×5 grid of unit spacing. The specific encounter model is the half-normal model given above and we used this code to simulate data used in subsequent analyses. The 100 activity centers were simulated on a state-space defined by a 8×8 square within which the trap array was centered (thus the trap array is buffered by 2 units). Therefore, the density of individuals in this system is fixed at $100/64$.

```
set.seed(2013)
# create 5 x 5 grid of trap locations with unit spacing
traplocs<- cbind(sort(rep(1:5,5)),rep(1:5,5))
```

```

3041 Dmat<-e2dist(traplocs,traplocs) # in cases where speed doesn't matter, it might be
3042                                     # clearer to just show the slow for-loop.
3043                                     # Plus, people will want to copy/paste this stuff
3044 ntraps<-nrow(traplocs)
3045
3046 # define state-space of point process. (i.e., where animals live).
3047 # "delta" just adds a fixed buffer to the outer extent of the traps.
3048 delta<-2
3049 Xl<-min(traplocs[,1] - delta)
3050 Xu<-max(traplocs[,1] + delta)
3051 Yl<-min(traplocs[,2] - delta)
3052 Yu<-max(traplocs[,2] + delta)
3053
3054 N<-100    # population size
3055 K<- 20    # number nights of effort
3056
3057 sx<-runif(N,Xl,Xu)    # simulate activity centers
3058 sy<-runif(N,Yl,Yu)
3059 S<-cbind(sx,sy)
3060 D<- e2dist(S,traplocs) # distance of each individual from each trap
3061
3062 alpha0<- -2.5        # define parameters of encounter probability
3063 sigma<- 0.5          #
3064 theta<- 1/(2*sigma*sigma)
3065 probcap<- expit(-2.5)*exp( - theta*D*D)    # probability of encounter
3066 # now generate the encounters of every individual in every trap
3067 Y<-matrix(NA,nrow=N,ncol=ntraps)
3068 for(i in 1:nrow(Y)){
3069     Y[i,<-rbinom(ntraps,K,probcap[i,])
3070 }

```

Subsequently we will generate data using this code packaged in an R function called `simSCR0.fn` which takes a number of arguments including `discard0` which, if `TRUE`, will return only the encounter histories for captured individuals. A second argument is `array3d` which, if `TRUE`, returns the 3-d encounter history array instead of the aggregated `nind × ntraps` encounter frequencies (see below). Finally we provide a random number seed, `sd` which we always set to 2013 in our analyses. Thus we obtain a data set as above using the following command

```

3078 data<-simSCR0.fn(discard0=TRUE,array3d=FALSE,sd=2013)

```

The **R** object `data` is a list, so let's take a look at what's in the list and then harvest some of its elements for further analysis below.

```

3081 > names(data)
3082 [1] "Y"          "traplocs" "xlim"      "ylim"      "N"          "alpha0"    "beta"
3083 [8] "sigma"      "K"

```

```

3084 > Y<-data$Y
3085 > traplocs<-data$traplocs

```

3086 4.4.1 Formatting and manipulating real data sets

3087 Conventional capture-recapture data are easily stored and manipulated as a 2-
 3088 dimensional array, an $nind \times nperiod$ matrix, which is maximally informative for
 3089 any conventional capture-recapture model, but not for spatial capture-recapture
 3090 models. For SCR models we must preserve the spatial information in the encounter
 3091 history information. We will routinely analyze data from 3 standard formats:

- 3092 (1) The basic 2-dimensional data format, which is an $nind \times ntraps$ encounter
 3093 frequency matrix such as that simulated previously;
- 3094 (2) The maximally informative 3-dimensional array which we establish here the
 3095 convention that it has dimensions $nind \times nperiods \times ntraps$ and
- 3096 (3) We use a compact format - the “SCR flat format” - which we describe below
 3097 in section 4.7.

3098 To simulate data in the most informative format - the “3-d array” - we can use the
 3099 **R** commands given previously but replace the last 4 lines with the following:

```

3100 Y<-array(NA,dim=c(N,K,ntraps))
3101 for(i in 1:nrow(Y)){
3102   for(j in 1:ntraps){
3103     Y[i,1:K,j]<-rbinom(K,1,probcap[i,j])
3104   }
3105 }

```

3106 We see that a collection of K binary encounter events are generated for *each*
 3107 individual and for *each* trap. The probabilities have those Bernoulli trials are
 3108 computed based on the distance from each individuals home range center and the
 3109 trap (see calculation above), and those are housed in the matrix probcap. Our
 3110 data simulator function `simSRC0.fn` will return the full 3-d array if `array3d=TRUE`
 3111 is specified in the function call. To recover the 2-d matrix from the 3-d array, and
 3112 subset the 3-d array to individuals that were captured, we do this:

```

3113 Y2d<- apply(Y,c(1,3),sum) # sum over the ‘‘replicates’’ dimension (2nd margin of the array)
3114 ncaps<-apply(Y2d,1,sum)   # compute how many times each individual was captured
3115 Y<-Y[ncaps>0,,]          # keep those individuals that were captured

```

4.5 FITTING AN SCR MODEL IN BUGS

3116 Clearly if we somehow knew the value of N then we could fit this model directly
 3117 because, in that case, it is a special kind of logistic regression model - one with a
 3118 random effect, but that enters into the model in a peculiar fashion - and also with
 3119 a distribution (uniform) which we don’t usually think of as standard for random

effects models. So our aim here is to analyze the known- N problem, using our simulated data, as an incremental step in our progress toward fitting more generally useful models.

To begin, we use our simulator to grab a data set and then harvest the elements of the resulting object for further analysis.

```

3125 data<-simSCR0.fn(discard0=FALSE,sd=2013)
3126 y<-data$Y
3127 traplocs<-data$traplocs
3128 nind<-nrow(y)
3129 X<-data$traplocs
3130 J<-nrow(X)
3131 y<-rbind(y,matrix(0,nrow=(100-nrow(y)),ncol=J ) )
3132 Xl<-data$xlim[1]
3133 Yl<-data$ylim[1]
3134 Xu<-data$xlim[2]
3135 Yu<-data$ylim[2]

```

Note that we specify `discard0 = FALSE` so that we have a "complete" data set, i.e., one with the all-zero encounter histories corresponding to uncaptured individuals. Now, within an **R** session, we can create the **BUGS** model file and fit the model using the following commands. This model describes the half-normal detection model but it would be trivial to modify that to various others including the logistic described above. One consequence of using the half-normal is that we have to constrain the encounter probability to be in $[0, 1]$ which we do here by defining `alpha0` to be the logit of the intercept parameter `p0`. Note that the distance covariate is computed within the **BUGS** model specification given the matrix of trap locations, `X`, which is provided to **WinBUGS** as data.

```

3146 cat("
3147 model {
3148   alpha0~dnorm(0,.1)
3149   logit(p0)<- alpha0
3150   theta~dnorm(0,.1)
3151   for(i in 1:N){
3152     s[i,1]~dunif(Xl,Xu)
3153     s[i,2]~dunif(Yl,Yu)
3154     for(j in 1:J){
3155       d[i,j]<- pow(pow(s[i,1]-X[j,1],2) + pow(s[i,2]-X[j,2],2),0.5)
3156       y[i,j] ~ dbin(p[i,j],K)
3157       p[i,j]<- p0*exp(- theta*d[i,j]*d[i,j])
3158     }
3159   }
3160 }
3161 ",file = "SCR0a.txt")
3162

```

Next we do a number of organizational activities including bundling the data for **WinBUGS**, defining some initial values, the parameters to monitor and some basic MCMC settings. We choose initial values for the activity centers **s** by generating uniform random numbers in the state-space but, for the observed individuals, we replace those values by each individual's mean trap coordinate for all encounters

```

3168 sst<-cbind(runif(nind,Xl,Xu),runif(nind,Yl,Yu)) # starting values for s
3169 for(i in 1:nind){
3170   if(sum(y[i,])==0) next
3171   sst[i,1]<- mean( X[y[i,]>0,1] )
3172   sst[i,2]<- mean( X[y[i,]>0,2] )
3173 }
3174
3175 data <- list (y=y,X=X,K=K,N=nind,J=J,Xl=Xl,Yl=Yl,Xu=Xu,Yu=Yu)
3176 inits <- function(){
3177   list (alpha0=rnorm(1,-4,.4),theta=runif(1,1,2),s=sst)
3178 }
3179
3180 library("R2WinBUGS")
3181 parameters <- c("alpha0","theta")
3182 nthin<-1
3183 nc<-3
3184 nb<-1000
3185 ni<-2000
3186 out <- bugs (data, inits, parameters, "SCROa.txt", n.thin=nthin,
3187   n.chains=nc, n.burnin=nb,n.iter=ni,debug=TRUE,working.dir=getwd())

```

There is little to say about the preceding basic operations other than to suggest that the interested reader explore the output and additional analyses by running the script provided in the **R** package **scrbook**. We ran 1000 burn-in and 1000 after burn-in, 3 chains, to obtain 3000 posterior samples. Because we know N for this particular data set we only have 2 parameters of the detection model to summarize (**alpha0** and **theta**). When the object **out** is produced we print a summary of the results as follows:

```

3195 > print(out,digits=3)
3196 Inference for Bugs model at "SCROa.txt", fit using WinBUGS,
3197 3 chains, each with 2000 iterations (first 1000 discarded)
3198 n.sims = 3000 iterations saved
3199
3200      mean      sd    2.5%    25%    50%    75%   97.5%  Rhat n.eff
3201 alpha0  -2.496  0.224  -2.954  -2.648  -2.48  -2.340  -2.091  1.013   190
3202 theta    2.442  0.419   1.638   2.145   2.44   2.721   3.303  1.005   530
3203 deviance 292.803 21.155 255.597 277.500 291.90 306.000 339.302 1.006   380
3204
3205 For each parameter, n.eff is a crude measure of effective sample size,
3206 and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
3207
3208 DIC info (using the rule, pD = Dbar-Dhat)

```


3208 `pD = -138.8 and DIC = 154.0`
 3209 `DIC is an estimate of expected predictive error (lower deviance is better).`

3210 We know the data were generated with `alpha0 = -2.5` and `theta = -2`. The
 3211 estimates look reasonably close to those data-generating values and we probably feel
 3212 pretty good about the performance of the Bayesian analysis and MCMC algorithm
 3213 that WinBUGS cooked-up based on our sample size of 1 data set. It is worth noting
 3214 that the Rhat statistics indicate reasonable convergence but, as a practical matter,
 3215 we might choose to run the MCMC algorithm for additional time to bring these
 3216 closer to 1.0 and to increase the effective posterior sample size (`n.eff`). Other
 3217 summary output includes “deviance” and related things including the deviance
 3218 information criterion (DIC). We discuss these things in chapter XXXX.

4.6 UNKNOWN N

3219 In all real applications N is unknown and that fact is kind of an important feature
 3220 of the capture-recapture problem! We handled this important issue in chapter 3
 3221 using the method of data augmentation which we apply here to achieve a realistic
 3222 analysis of Model SCR0. As with the basic closed population models considered
 3223 previously, we formulate the problem here by augmenting our observed data set
 3224 with a number of “all zero” encounter histories - what we referred to in Chapter
 3225 3 as potential individuals. If n is the number of observed individuals, then let
 3226 $M - n$ be the number of potential individuals in the data set. For the basic y_{ij}
 3227 data structure (individuals x traps encounter frequencies) we simply add additional
 3228 rows of “all 0” observations to that data set. This is because such “individuals” are
 3229 unobserved, and therefore necessarily have $y_{ij} = 0$ for all j . A data set, say with 4
 3230 traps and 6 individuals, augmented with 4 pseudo-individuals therefore might look
 3231 like this:

	trap1	trap2	trap3	trap4
[1,]	1	0	0	0
[2,]	0	2	0	0
[3,]	0	0	0	1
[4,]	0	1	0	0
[5,]	0	0	1	1
[6,]	1	0	1	0
[7,]	0	0	0	0
[8,]	0	0	0	0
[9,]	0	0	0	0
[10,]	0	0	0	0

3243 We typically have more than 4 traps and, if we’re fortunate, many more indi-
 3244 viduals in our data set.

3245 For the augmented data, we introduce a set of binary latent variables (the data
 3246 augmentation variables), z_i , and the model is extended to describe $\Pr(z_i = 1)$
 3247 which is, in the context of this problem, the probability that an individual in the

augmented data set is a member of the population that was sampled. In other words, if $z_i = 1$ for one of the “all zero” encounter histories, this is implied to be a sampling zero whereas observations for which $z_i = 0$ are “structural zeros” under the model.

How big does the augmented data set have to be? We discussed this issue in chapt. 3 where we noted that the size of the data set is equivalent to the upper limit of a uniform prior distribution on N . Practically speaking, it should be sufficiently large so that the posterior distribution for N is not truncated. On the other hand, if it is too large then unnecessary calculations are being done. An approach to choosing M by trial-and-error is indicated. You can take a ballpark estimate of the probability that an individual is captured (at all during the study), obtain N as $n/pcap$, and then set $M = 2 * N$, as a first guess. Do a short MCMC run and then consider whether you need to do something different. See chapt. 7 for an example of this. Kery and Schaub (2011, ch. 6) provide an assessment of choosing M in closed population models.

Analysis by data augmentation removes N as an explicit parameter of the model. Instead, N is a derived parameter, computed by $N = \sum_{i=1}^M z_i$. Similarly, *density*, D , is also a derived parameter computed as $D = N/area(\mathcal{S})$. For our simulator, we’re using an 8×8 state-space and thus we will compute D as $D = N/64$.

4.6.1 Analysis using data augmentation in WinBUGS

As before we begin by obtaining a data set using our `simSCR0.fn` routine and then harvesting the required data objects from the resulting data list. Note that we use the `discard0=TRUE` option this time so that we get a “real” data set with no all-zero encounter histories. After harvesting the data we produce the **WinBUGS** model specification which now includes M encounter histories including the augmented potential individuals, the data augmentation parameters z_i , and the data augmentation parameter ψ .

```

data<-simSCR0.fn(discard0=TRUE,sd=2013)
y<-data$Y
traplocs<-data$traplocs
nind<-nrow(y)
X<-data$traplocs
J<-nrow(X)
Xl<-data$xlim[1]
Yl<-data$ylim[1]
Xu<-data$xlim[2]
Yu<-data$ylim[2]

cat("
model {
  alpha0~dnorm(0,.1)
  logit(p0)<- alpha0

```

```

3290 theta~dnorm(0,.1)
3291 psi~dunif(0,1)
3292
3293 for(i in 1:M){
3294   z[i] ~ dbern(psi)
3295   s[i,1]~dunif(Xl,Xu)
3296   s[i,2]~dunif(Yl,Yu)
3297   for(j in 1:J){
3298     d[i,j]<- pow(pow(s[i,1]-X[j,1],2) + pow(s[i,2]-X[j,2],2),0.5)
3299     y[i,j] ~ dbin(p[i,j],K)
3300     p[i,j]<- z[i]*p0*exp(- theta*d[i,j]*d[i,j])
3301   }
3302 }
3303 N<-sum(z[])
3304 D<-N/64
3305 }
3306 ",file = "SCR0a.txt")

```

3307 To prepare our data we have to augment the data matrix y with $M - n$ all-
 3308 zero encounter histories, we have to create starting values for the variables z_i and
 3309 also the activity centers s_i of which, for each, we require M values. Otherwise the
 3310 remainder of the code for bundling the data, creating initial values and executing
 3311 **WinBUGS** looks much the same as before except with more or differently named
 3312 arguments.

```

3313 ## Data augmentation stuff
3314 M<-200
3315 y<-rbind(y,matrix(0,nrow=M-nind,ncol=ncol(y)))
3316 z<-c(rep(1,nind),rep(0,M-nind))
3317
3318 sst<-cbind(runif(M,Xl,Xu),runif(M,Yl,Yu)) # starting values for s
3319 for(i in 1:nind){
3320   if(sum(y[i,])==0) next
3321   sst[i,1]<- mean( X[y[i,]>0,1] )
3322   sst[i,2]<- mean( X[y[i,]>0,2] )
3323 }
3324 data <- list (y=y,X=X,K=K,M=M,J=J,Xl=Xl,Yl=Yl,Xu=Xu,Yu=Yu)
3325 inits <- function(){
3326   list (alpha0=rnorm(1,-4,.4),theta=runif(1,1,2),s=sst,z=z)
3327 }
3328
3329 library("R2WinBUGS")
3330 parameters <- c("alpha0","theta","N")
3331 nthin<-1
3332 nc<-3
3333 nb<-1000
3334 ni<-2000

```

```

3335 out <- bugs (data, inits, parameters, "SCR0a.txt", n.thin=nthin,n.chains=nc,
3336   n.burnin=nb,n.iter=ni,debug=TRUE,working.dir=getwd())

```

3337 **Remarks:** (1) Note the differences in this new **WinBUGS** model with that
 3338 appearing in the known- N version. (2) Also the input data has changed - the
 3339 augmented data set has more rows of all-zeros. Previously we knew that $N = 100$
 3340 but in this analysis we pretend not to know N , but think that $N = 200$ is a good
 3341 upper-bound; (3) Population size $N(S)$ is a derived parameter, being computed by
 3342 summing up all of the data augmentation variables z_i (as we've done previously);
 3343 (4) Density, $D \equiv D(S)$, is also a derived parameter. Summarizing the output from
 3344 **WinBUGS** produces:

```

3345 > print(out1,digits=2)
3346 Inference for Bugs model at "SCR0a.txt", fit using WinBUGS,
3347 3 chains, each with 2000 iterations (first 1000 discarded)
3348 n.sims = 3000 iterations saved
3349      mean      sd    2.5%    25%    50%    75%   97.5% Rhat n.eff
3350 alpha0   -2.57  0.23   -3.04  -2.72  -2.56  -2.41  -2.15  1.01   320
3351 theta     2.46  0.42    1.63   2.16   2.46   2.73   3.33  1.02   120
3352 N        113.62 15.73   86.00 102.00 113.00 124.00 147.00 1.01   260
3353 D          1.78  0.25    1.34   1.59   1.77   1.94   2.30  1.01   260
3354 deviance 302.60 23.67 261.19 285.47 301.50 317.90 354.91 1.00  1400
3355
3356 For each parameter, n.eff is a crude measure of effective sample size,
3357 and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
3358
3359 DIC info (using the rule, pD = var(deviance)/2)
3360 pD = 279.9 and DIC = 582.5
3361 DIC is an estimate of expected predictive error (lower deviance is better).

```

3362 The column labeled “MC error” is the Monte Carlo error - the error inherent in
 3363 the attempt to compute these posterior summaries by MCMC. It is desirable to run
 3364 the Markov chain algorithm long enough so as to reduce the MC error to a tolerable
 3365 level. What constitutes tolerable is up to the investigator. Certainly less than 1% is
 3366 called for. As a general rule, Rhat gets closer to 1 and MC error decreases toward 0
 3367 as the number of iterations increases. We see that the estimated parameters (α_0 and
 3368 θ) are comparable to the previous results obtained for the known- N case, and also
 3369 not too different from the data-generating values. The posterior of N overlaps the
 3370 data-generating value substantially with a mean of 113.62. To obtain these results
 3371 we fitted the true data-generating model, that based on the half-normal detection
 3372 model, to a single simulated data set. For fun and excitement we fit the *wrong*
 3373 model - that with the logistic-linear detection model - to the same data set. This is
 3374 easily achieved by modifying the **WinBUGS** model specification above, although
 3375 we provide the **R** script in the **R** package **scrbook**. Those results are given below.
 3376 We see that the estimate of N , the main parameter of interest, is very similar to
 3377 that obtained under the correct model, convergence is worse (as measured by Rhat)
 3378 which probably doesn't have anything to do with the model being wrong, and the

posterior deviance and DIC favor the correct model. We consider the use of DIC for carrying-out model selection in chapter 8.

```

> print(out2,digits=2)
Inference for Bugs model at "SCR0a.txt", fit using WinBUGS,
  3 chains, each with 2000 iterations (first 1000 discarded)
  n.sims = 3000 iterations saved
      mean      sd    2.5%    25%    50%    75%   97.5%  Rhat  n.eff
alpha0  -1.59  0.27   -2.16  -1.77  -1.58  -1.42  -1.07  1.05    60
beta     3.77  0.43    2.92   3.48   3.79   4.05   4.66  1.04    70
N       122.57 18.67   90.00 109.00 122.00 135.00 163.00 1.00  3000
D        1.92  0.29    1.41   1.70   1.91   2.11   2.55  1.00  3000
deviance 312.67 22.43 271.00 297.20 311.50 327.00 359.60 1.02   130

For each parameter, n.eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor (at convergence, Rhat=1).

DIC info (using the rule, pD = var(deviance)/2)
pD = 247.5 and DIC = 560.1
DIC is an estimate of expected predictive error (lower deviance is better).

```

4.6.2 Use of other BUGS engines: JAGS

There are two other popular **BUGS** engines in widespread use: **OpenBUGS** (Thomas et al., 2006) and **JAGS** (Plummer, 2003). Both of these are easily called from **R**. **OpenBUGS** can be used instead of **WinBUGS** by changing the package option in the bugs call to `package=OpenBUGS`. **JAGS** can be called using the function `jags()` in package **R2JAGS** which has nearly the same arguments as `bugs()`. We prefer to use the **R** library `rjags` (Plummer, 2009) which has a slightly different implementation that we demonstrate here as we reanalyze the simulated data set in the previous section (note: the same **R** commands are used to generate the data and package the data, inits and parameters to monitor). The function `jags.model` is used to initialize the model and run the MCMC algorithm for a period in which adaptive rejection (XXXX not sure XXXXX???) sampling is used. Then the Markov chains are updated using `coda.samples()` to obtain posterior samples for analysis, as follows:

```

jm<- jags.model("SCR0a.txt", data=data, inits=inits, n.chains=nc,
               n.adapt=nb))
jm<- coda.samples(jm, parameters, n.iter=ni-nb, thin=nthin)

```

We find that JAGS seems to be 20-30% faster for the basic SCR model which the reader can evaluate using the script `jags.winbugs.R` in the **R** package `scrbook`.

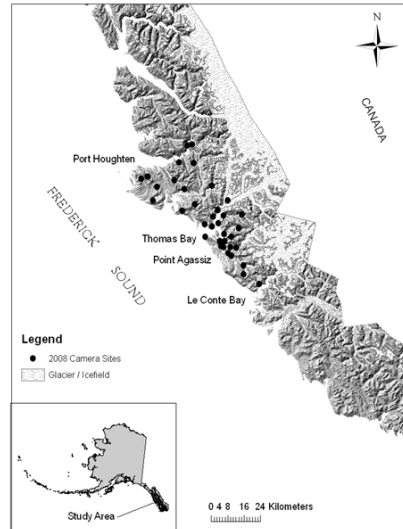


Figure 4.3. Wolverine camera trap locations from Magoun et al. (2011).

4.7 CASE STUDY: WOLVERINE CAMERA TRAPPING STUDY

We provide an analysis here of A. Magoun’s wolverine data (Magoun et al., 2011; Royle et al., 2011c). The study took place in SE Alaska (Fig. 4.3) where 37 cameras were operational for variable periods of time (min = 5 days, max = 108 days, median = 45 days). A consequence of this is that the binomial sample size K (see Eq. 4.2.1) is variable for each camera. Thus, we must provide a matrix of sample sizes as data to BUGS and modify the model specification in sec. 4.6 accordingly. Our treatment of the data here is based on the analysis of Royle et al. (2011c).

To carry-out an analysis of these data, we require the matrix of trap coordinates and the encounter history data. We store data in an the “scr flat format” (see sec. 4.4.1 above), an efficient file format which is easily manipulated and also used as the input file format in our custom **R** script (ch. xxx) and **SPACECAP** (Gopalaswamy, 2012). To illustrate this format, the wolverine data are available as an encounter data **R** object named “**wcaps**” which has 3 columns and 115 rows, each representing a unique encounter event including the trap identity, the individual identity and the sample occasion index (**sample**). The first 10 rows of this matrix are as follows:

```
> wcaps
      trapid individual sample
```

3436	[1,]	1	2	127
3437	[2,]	1	2	128
3438	[3,]	1	2	129
3439	[4,]	1	18	130
3440	[5,]	2	3	106
3441	[6,]	2	18	104
3442	[7,]	5	5	73
3443	[8,]	5	5	89
3444	[9,]	6	18	117
3445	[10,]	6	18	118

3446 This “encounter data file” contains 1 row for each unique individual/trap en-
 3447 counter, and 3 variables (columns): `trapid` is an integer that runs from `1:ntraps`,
 3448 individual runs from `1:nind` and sample runs from `1:nperiods`. Often (as the case
 3449 here) “sample” will correspond to daily sample intervals. The variable `trapid` will
 3450 have to correspond to the row of a matrix containing the trap coordinates - a file
 3451 named `traplocs.csv` available in the **R** package `scrbook`.

3452 Note that these data do not represent a completely informative summary of the
 3453 data. For example, if no individuals were captured in a certain trap or during a
 3454 certain period, then this compact data format will have no record. Thus we will
 3455 need to know `ntraps` and `nperiods` when reformatting this SCR data format into a
 3456 2-d encounter frequency matrix or 3-d array. In addition, the encounter data file
 3457 does not provide information about which periods each trap was operated. This
 3458 additional information is also necessary as the trap-specific sample sizes must be
 3459 passed to **BUGS** as data. We provide this information in a 2nd data file - which
 3460 we call the “trap deployment” file (described below).

3461 The “encounter data file” `wcaps.csv` exists in the **R** package `scrbook` as a .csv
 3462 file that people can read into **R** and do some basic summary statistics on. For our
 3463 purposes we need to convert these data into the “individual x trap” array of binary
 3464 encounter frequencies, although more general models might require an encounter-
 3465 history formulation of the model which requires a full 3-d array. To obtain our `nind`
 3466 x `ntrap` encounter frequency matrix, we do this the hard way by first converting the
 3467 encounter data file into a 3-d array and then summarize to trap totals. We have a
 3468 handy function `SCR23darray.fn` which takes the compact encounter data file with
 3469 optional arguments `ntraps` and `nperiods`, and converts it to a 3-d array, and then
 3470 we use the **R** function `apply` to summarize over the “sample” period dimension (by
 3471 convention here, this is the 2nd dimension):

```

3472 SCR23darray.fn <- function(caps,ntraps=NULL,nperiods=NULL){
3473   nind<-max(caps[,2])
3474   if(is.null(ntraps)) ntraps<-max(caps[,1])
3475   if(is.null(nperiods)) nperiods<- max(caps[,3])
3476
3477   y<-array(0,c(nind,nperiods,ntraps))
3478   tmp<-cbind(caps[,2],caps[,3],caps[,1])
3479   y[tmp]<-1

```

```

3480 y
3481 }
3482
3483 # for the wolverine data do this:
3484
3485 Y3d <-SCR23darray.fn(wcaps,ntraps=37,nperiods=165)
3486 y <- apply(y3d,c(1,3),sum)

```

If `ntraps` and `nperiods` are not specified then they are assumed to be equal to the maximum value provided in the encounter data file. The 3-d array is necessary to fit certain types of models (e.g., behavioral response) and this is why we sometimes will require this maximally informative 3-d data format.

The other data file that we must have is the “trap deployment” file (henceforth “traps file”) which provides the additional information not contained in the encounter data file. The traps file has `nperiods + 3` columns. The first column is assumed to be a trap identifier, columns 2 and 3 are the easting and northing coordinates (assumed to be in a Euclidean coordinate system), and columns 4 to (`nperiods + 3`) are binary indicators of whether each trap was operational in each time period. The first 5 rows (out of 37) and 10 columns (out of 168) of the traps file for the wolverine data (“`wtraps.csv`” in the **R** package `scrbook` are:

	Trap	Easting	Northing	1	2	3	4	5	6	7	<- column names
3499	1	39040	19216	0	0	0	0	0	0	0	
3500	2	41324	19772	1	1	1	1	1	1	1	
3501	3	44957	12985	0	0	0	0	0	0	0	
3502	4	41151	23220	0	0	0	0	0	0	0	
3503	5	44240	17198	0	0	0	0	0	0	0	

This tells us that trap 2 was operated in periods 1-7 but the other traps were not operational during those periods. To extract the relevant information to fit the model in **WinBUGS** we do this:

```

3508 traps<- read.csv("wtraps.csv")
3509 traplocs<- traps[,2:3]
3510 K<- apply(traps[,4:ncol(traps)],1,sum)

```

This results in a matrix `traplocs` which contains the coordinates of each trap and a vector `K` containing the number of days that each trap was operational. We now have all the information required to fit a basic SCR model in **WinBUGS**.

Summarizing these data files for the wolverine study, we see that 21 unique individuals were captured a total of 115 times. Most individuals were captured 1-6 times, with 4, 1, 4, 3, 1, and 2 individuals captured 1-6 times, respectively. In addition, 1 individual was captured each 8 and 14 times and 2 individuals each were captured 10 and 13 times. The number of unique traps that captured a particular individual ranged from 1-6, with 5, 10, 3, 1, 1, and 1 individual captured in each of 1-6 traps, respectively, for a total of 50 unique wolverine-trap encounters. These

numbers might be hard to get your mind around whereas some tabular summary is often more convenient. For that it seems natural to tabulate individuals by trap and total encounter frequencies. The spatial information in SCR data is based on multi-trap captures, and so, it is informative to understand how many unique traps each individual is captured in. At the same, it is useful to understand how many total captures we have of each individual because this is, in an intuitive sense, the effective sample size. So, we reproduce Table 1 from Royle et al. (2011c) which shows the trap and total encounter frequencies:

Table 4.1. Individual frequencies of capture for wolverines captured in camera traps in South-east Alaska in 2008. Rows index unique trap frequencies and columns represent total number of captures (e.g., we captured 4 individuals 1 time, necessarily in only 1 trap; we captured 3 individuals 3 times but in 2 different traps)

	No. of captures									
No. of traps	1	2	3	4	5	6	8	10	13	14
1	4	1	0	0	0	0	0	0	0	0
2	0	0	3	3	0	2	1	2	0	0
3	0	0	1	1	0	0	0	0	0	1
4	0	0	0	0	0	0	0	0	1	0
5	0	0	0	0	1	0	0	0	0	0
6	0	0	0	0	0	0	0	0	1	0

4.7.1 Fitting the model in WinBUGS

For illustrative purposes here we fit the simplest SCR model with the half-normal distance function although we revisit these data with more complex models in later chapters. The model is summarized by the following 3 components:

- (1) $y_{ij} | \mathbf{s}_i \sim \text{Bin}(K, z_i p_{ij})$
- (2) $p_{ij} = p_0 \exp(-\theta \|\mathbf{s}_i - x_j\|^2)$
- (3) $\mathbf{s}_i \sim \text{Unif}(\mathcal{S})$
- (4) $z_i \sim \text{Bern}(\psi)$

We assume customary flat priors on the structural (hyper-) parameters of the model, $\alpha_0 = \text{logit}(p_0)$, θ and ψ . It remains to define the state-space \mathcal{S} . For this, we nested the trap array (Fig. 4.3) in a rectangular state-space extending 20 km beyond the traps in each cardinal direction. We also considered larger state-spaces up to 50 km to evaluate that choice. The buffer of the state space should be larger enough so that individuals beyond the state-space boundary are not likely to be encountered. Thus some knowledge of typical space usage patterns of the species is useful. The coordinate system was scaled so that a unit distance was equal to 10km, producing a rectangular state-space of dimension 9.88x10.5 units ($\text{area} = 10374 \text{km}^2$) within which the trap array was nested. As a general rule, we recommend scaling the state-space so that it is defined near the origin $(x, y) = (0, 0)$. While the scaling of the

coordinate system is theoretically irrelevant, a poorly scaled coordinate system can produce Markov chains that mix poorly. We fitted this model in **WinBUGS** using data augmentation with $M = 300$ potential individuals, using 3 Markov chains each of 12000 total iterations, discarding the first 2000 as burn-in. [R commands for reading in the data and executing the analysis are as follows:

provide those commands here

The output follows (note, we have a parameter “sigma” which we discuss shortly)²:

All based on 3 chains, 12k iters, 2k burn, 30k total

Buffer = 10 km

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
psi	0.13	0.03	0.08	0.11	0.13	0.15	0.20	1	10000
sigma	0.65	0.06	0.55	0.61	0.64	0.68	0.76	1	1800
p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1	20000
N	39.63	6.70	29.00	35.00	39.00	44.00	54.00	1	7100
D	5.92	1.00	4.33	5.22	5.82	6.57	8.06	1	7100
beta	1.23	0.21	0.85	1.08	1.22	1.36	1.66	1	1800
deviance	410.05	12.06	388.70	401.50	409.20	417.80	435.60	1	22000

Buffer = 15 km

n.sims = 30000 iterations saved

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
psi	0.16	0.04	0.10	0.14	0.16	0.19	0.25	1	3800
sigma	0.64	0.06	0.54	0.60	0.64	0.67	0.76	1	510
p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1	17000
N	48.77	9.19	34.00	42.00	48.00	54.00	69.00	1	3300
D	5.78	1.09	4.03	4.98	5.69	6.40	8.18	1	3300
beta	1.25	0.21	0.86	1.10	1.24	1.39	1.70	1	510
deviance	411.00	12.16	389.50	402.40	410.30	418.70	437.00	1	5400

Buffer = 20 km

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
psi	0.20	0.05	0.12	0.17	0.20	0.23	0.30	1	16000
sigma	0.64	0.06	0.54	0.60	0.63	0.67	0.76	1	1200
p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1	1900
N	59.84	11.89	40.00	51.00	59.00	67.00	86.00	1	20000
D	5.77	1.15	3.86	4.92	5.69	6.46	8.29	1	20000
beta	1.26	0.21	0.87	1.11	1.25	1.40	1.71	1	1200
deviance	411.01	12.36	389.10	402.30	410.20	418.80	437.50	1	1500

Buffer = 25 km

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
psi	0.24	0.05	0.15	0.20	0.24	0.28	0.36	1	3400

²Final as of 1/11/2012. output saved in wolv-buffer-study.txt

3590	sigma	0.64	0.05	0.54	0.60	0.63	0.67	0.75	1	3600
3591	p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1	5000
3592	N	72.40	14.72	47.00	62.00	71.00	81.00	105.00	1	2700
3593	D	5.79	1.18	3.76	4.96	5.67	6.47	8.39	1	2700
3594	beta	1.26	0.21	0.88	1.12	1.25	1.40	1.71	1	3600
3595	deviance	411.35	12.23	389.70	402.70	410.55	419.20	437.20	1	30000
3596										
3597	Buffer = 30 km									
3598		mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
3599	psi	0.29	0.06	0.18	0.24	0.28	0.33	0.43	1	3100
3600	sigma	0.63	0.05	0.54	0.60	0.63	0.67	0.75	1	5600
3601	p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1	11000
3602	N	86.42	17.98	56.00	74.00	85.00	97.00	126.02	1	3900
3603	D	5.82	1.21	3.77	4.98	5.72	6.53	8.49	1	3900
3604	beta	1.27	0.21	0.88	1.12	1.26	1.41	1.71	1	5600
3605	deviance	411.06	12.37	389.20	402.50	410.20	418.90	437.60	1	10000
3606										
3607	Buffer = 35 km									
3608		mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
3609	psi	0.34	0.08	0.21	0.29	0.34	0.39	0.50	1	30000
3610	sigma	0.63	0.05	0.54	0.60	0.63	0.67	0.75	1	4500
3611	p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1	24000
3612	N	101.79	21.54	65.00	87.00	100.00	115.00	148.00	1	30000
3613	D	5.85	1.24	3.74	5.00	5.75	6.61	8.51	1	30000
3614	beta	1.27	0.21	0.89	1.12	1.25	1.40	1.70	1	4500
3615	deviance	411.10	12.20	389.50	402.40	410.30	418.90	437.20	1	22000
3616										
3617	Buffer = 40 km									
3618		mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
3619	psi	0.39	0.09	0.24	0.33	0.39	0.45	0.60	1.01	480
3620	sigma	0.64	0.05	0.54	0.60	0.63	0.67	0.75	1.01	410
3621	p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1.00	21000
3622	N	118.05	26.14	75.00	100.00	116.00	133.00	178.00	1.01	450
3623	D	5.87	1.30	3.73	4.97	5.76	6.61	8.84	1.01	450
3624	beta	1.27	0.21	0.89	1.12	1.25	1.40	1.72	1.01	410
3625	deviance	411.37	12.35	389.30	402.60	410.60	419.30	437.50	1.00	9700
3626										
3627	Buffer = 45 km									
3628		mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
3629	psi	0.45	0.10	0.28	0.38	0.44	0.51	0.66	1	3600
3630	sigma	0.64	0.05	0.54	0.60	0.63	0.67	0.75	1	10000
3631	p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1	8100
3632	N	134.43	28.68	85.00	114.00	132.00	153.00	196.00	1	3300
3633	D	5.83	1.24	3.68	4.94	5.72	6.63	8.50	1	3300
3634	beta	1.26	0.21	0.88	1.11	1.24	1.39	1.69	1	10000
3635	deviance	411.36	12.19	389.60	402.70	410.60	419.10	437.30	1	9400

3636										
3637	Buffer = 50 km									
3638		mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
3639	psi	0.51	0.11	0.31	0.43	0.50	0.57	0.74	1	3200
3640	sigma	0.63	0.05	0.54	0.60	0.63	0.67	0.75	1	4700
3641	p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1	3300
3642	N	151.61	31.65	96.00	129.00	149.00	172.00	221.00	1	3400
3643	D	5.79	1.21	3.66	4.92	5.69	6.56	8.43	1	3400
3644	beta	1.27	0.21	0.89	1.12	1.25	1.40	1.70	1	4700
3645	deviance	410.81	12.18	389.20	402.30	410.10	418.50	436.70	1	30000
3646										
3647	Buffer = 55 km									
3648		mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
3649	psi	0.56	0.12	0.35	0.48	0.55	0.64	0.82	1.01	260
3650	sigma	0.64	0.05	0.54	0.60	0.63	0.67	0.76	1.00	1600
3651	p0	0.06	0.01	0.04	0.05	0.06	0.06	0.08	1.00	30000
3652	N	169.28	35.81	108.00	143.00	166.00	192.00	247.00	1.01	260
3653	D	5.73	1.21	3.66	4.84	5.62	6.50	8.36	1.01	260
3654	beta	1.25	0.21	0.88	1.11	1.24	1.39	1.69	1.00	1600
3655	deviance	411.28	12.38	389.40	402.60	410.50	419.10	437.50	1.00	26000

3656 We see that the estimated density is roughly consistent as we increase the state-
 3657 space buffer from 15 to 50 *km*. We do note that the data augmentation parameter
 3658 ψ (and, correspondingly, N) increase with the size of the state space in accordance
 3659 with the deterministic relationship $N = D * A$. However, density is constant more
 3660 or less as we increase the size of the state-space beyond a certain point. For the
 3661 10 *km* state-space buffer, we see a slight effect on the posterior distribution of D .
 3662 This is not a bug but rather a feature. As we noted above, the state-space is part
 3663 of the model.

3664 One thing we haven't talked about yet is that we can calibrate the desired size
 3665 of the state-space by looking at the estimated home range radius of the species. For
 3666 some models it is possible to convert the parameter θ directly into the home range
 3667 radius (section XXX XYZ). For the half-normal model we interpret the half-normal
 3668 scale parameter σ which is related to θ by $\theta = 1/(2\sigma^2)$ as the radius of a bivariate
 3669 normal movement model.

3670 4.7.2 Conclusion of Analysis

3671 Our point estimate of wolverine density from this study, using the posterior mean
 3672 from the state-space based on the 20 *km* buffer, is approximately 5.77 individu-
 3673 als/1000 *km*² with a 95% posterior interval of [3.86, 8.29]. Density is estimated
 3674 imprecisely which might not be surprising given the low sample size ($n = 21$ indi-
 3675 viduals!). This seems to be a basic feature of carnivore studies although it should
 3676 not (in our view) preclude the study of their populations nor attempts to estimate
 3677 density or vital rates.

It is worth thinking about this model, and these estimates, computed under a rectangular state space roughly centered over the trapping array (Fig. 4.3). Does it make sense to define the state-space to include, for example, ocean? What are the possible consequences of this? What can we do about it? There's no reason at all that the state space has to be a regular polygon – we defined it as such here strictly for convenience and for ease of implementation in **WinBUGS** where it enables us to specify the prior for the activity centers as uniform priors for each coordinate. While it would be possible to define a more realistic state-space using some general polygon, it might take some effort to implement that in the **BUGS** language (see chapter XYZXYZ³ for example of a simple case). Alternatively, we recommend using a discrete representation of the state-space – i.e., approximate \mathcal{S} by a grid of G points. We discuss this in the following section.

4.8 CONSTRUCTING DENSITY MAPS

One of the most useful aspects of SCR models is that they are parameterized in terms of individual locations - i.e., *where* each individual lives – and, thus, we can compute many useful or interesting summaries of the activity centers. For example, we can make a spatial density plot by tallying up the number of activity centers \mathbf{s}_i in boxes of arbitrary size and then producing a nice multi-color spatial plot of those which, we find, increases the acceptance probability of your manuscripts by a substantial amount. We discussed in chapter 2 the idea of estimating derived parameters from MCMC output. In SCR models, there are many derived parameters that are functions of the latent point locations $(\mathbf{s}_1, \dots, \mathbf{s}_N)$. In the present context, the number of individuals living in any well-defined polygon is a derived parameter. Specifically, let $B(x)$ indicate a box centered at x then

$$N(x) = \sum_i I(\mathbf{s}_i \in B(x))$$

is the population size of box $B(x)$, and $D(x) = N(x)/|B(x)|$ is the local density. These are just “derived parameters” (see chapter 2) which are estimated from MCMC output using the appropriate Monte Carlo average. One thing to be careful about, in the context of models in which N is unknown, is that, for each MCMC iteration m , we only tabulate those activity centers which correspond to individuals in the sampled population. i.e., for which the data augmentation variable $z_i = 1$. In this case, we take all of the output for MCMC iterations $m = 1, 2, \dots, \text{niter}$ and compute this summary:

$$N(x, m) = \sum_{z_{i,m}=1} I(s_{i,m} \in B(x))$$

³raccoon example or something?

3709 Thus, $N(x, 1), N(x, 2), \dots$, is the Markov chain for parameter $N(x)$. In what follows
 3710 we will provide a set of **R** commands for doing this calculations and making a basic
 3711 image plot from the MCMC output.

3712 **Step 1:** Define the center points of each box, $B(x)$, or point at which local density
 3713 will be estimated:

```
3714 xg<-seq(Xl,Xu,,50)
3715 yg<-seq(Yl,Yu,,50)
```

3716 **Step 2:** Extract the MCMC histories for the activity centers and the data aug-
 3717 mentation variables. Note that these are each $N \times \text{niter}$ matrices:

```
3718 Sxout<-out$sims.list$s[,1]
3719 Syout<-out$sims.list$s[,2]
3720 z<-out$sims.list$z
```

3721 **Step 3:** We associate each coordinate with the proper box using the **R** command
 3722 `cut()`. Note that we keep only the activity centers for which $z = 1$ (i.e., individuals
 3723 that belong to the population of size N):

```
3724 Sxout<-cut(Sxout[z==1],breaks=xg,include.lowest=TRUE)
3725 Syout<-cut(Syout[z==1],breaks=yg,include.lowest=TRUE)
```

3726 **Step 4:** Use the `table()` command to tally up how many activity centers are in
 3727 each $B(x)$:

```
3728 Dn<-table(Sxout,Syout)
```

3729 **Step 5:** Use the `image()` command to display the resulting matrix.

```
3730 image(xg,yg,Dn/nrow(z),col=terrain.colors(10))
```

3731 Praise the Lord! This map is somewhat useful or at least it looks pretty and will
 3732 facilitate the publication of your papers.

3733 It is worth emphasizing here that density maps will not usually appear uniform
 3734 despite that we have assumed that activity centers are uniformly distributed. This is
 3735 because the observed encounters of individuals provide direct information about the
 3736 location of the $i = 1, 2, \dots, n$ activity centers and thus their “estimated” locations
 3737 will be affected by the observations. In a limiting sense, were we to sample space
 3738 intensely enough, every individual would be captured a number of times and we
 3739 would have considerable information about all N point locations. Consequently,
 3740 the uniform prior would have almost no influence at all on the estimated density
 3741 surface in this limiting situation. Thus, in practice, the influence of the uniformity
 3742 assumption increases as the fraction of the population encountered decreases.

3743 **On the non-intuitiveness of `image()`** – the **R** function `image()` might not
 3744 be very intuitive to some – it plots $M[1,1]$ in the lower left corner. If you want $M[]$
 3745 to be plotted “as you look at it” then $M[1,1]$ should be in the upper left corner.
 3746 We have a function `rot()` which does that. If you do `image(rot(M))` then it puts
 3747 it on the monitor as if it was a map you were looking at. You can always specify
 3748 the x and y – labels explicitly as we did above.

3749 **Spatial dot plots** – Now here is a cruder version based on the “spatial
 3750 dot map” function `spatial.plot`. The useful functions in **R** are `image()` and
 3751 `image.scale()` which is a function we grabbed off the web somewhere. Use of
 3752 this function requires arguments of point locations and the resulting value to be
 3753 displayed. The function is defined and applied as follows:

```
3754 spatial.plot<- function(x,y){
3755   nc<-as.numeric(cut(y,20))
3756   plot(x,pch=" ")
3757   points(x,pch=20,col=topo.colors(20)[nc],cex=2)
3758   image.scale(y,col=topo.colors(20))
3759 }
3760 # To execute the function do this:
3761 spatial.plot(cbind(xg,yg), Dn/nrow(z))
```

3762 4.8.1 Example: Wolverine density map.

3763 We used the posterior output from the wolverine model fitted previous to compute
 3764 a relatively coarse version of a density map, using a 10×10 grid (Fig. 4.4) and
 3765 using a 30×30 grid (Fig. 4.5)⁴. In these figures density is expressed in units of
 3766 individuals per 1000 km^2 , while the area of the pixels is about 1037 km^2 and 115
 3767 km^2 , respectively. That calculation is based on⁵:

```
3768 > total.area<- (Yu-Yl)*(Xu-Xl)*1000
3769 > total.area/(10*10)
3770 [1] 1037.427
3771 > total.area/(30*30)
3772 [1] 115.2697
```

3773 A couple of things are worth noting: First is that as we move away from “where
 3774 the data live” - away from the trap array - we see that the density approaches
 3775 the mean density. This is a property of the estimator as long as the “detection
 3776 function” decreases sufficiently rapidly as a function of distance. Relatedly, it is
 3777 also a property of statistical smoothers such as splines, kernel smoothers, and re-
 3778 gression smoothers - predictions tend toward the global mean as the influence of

⁴Final as of 01/11/2012. Note: Not sure if we should use quantiles for color to make equal area slices. ??? Also should we use the same scale?

⁵This is wrong and needs fixed. Move decimal one place over. i.e., 100 instead of 1000.




Figure 4.4. Needs a caption

data diminishes. Another way to think of it is that it is a consequence of the prior - which imposes uniformity, and as you get far away from the data, the predictions tend to the prior. The other thing to note about this map is that density is not 0 over water (although the coastline is not shown). This might be perplexing to some who are fairly certain that wolverines do not like water. However, there is nothing about the model that recognizes water from non-water and so the model predicts over water *as if* it were habitat similar to that within which the array is nested. But, all of this is ok as far as estimating density goes and, furthermore, we can compute valid estimates of N over any well-defined region which presumably wouldn't include water if we so choose.

4.9 DISCRETE STATE-SPACE

The SCR model developed previously in this chapter assumes that individual activity centers are distributed uniformly over the prescribed state-space. Clearly this will not always be a reasonable assumption. In chapter ?? we talk about developing models that allow explicitly for non-uniformity of the activity centers by modeling covariate effects on density. A simpler method of affecting the distribution of activity centers, which we address here, is to modify the shape of the state-space explicitly. For example, we might be able to classify the state-space into distinct




Figure 4.5. Needs a caption

blocks of habitat and non-habitat. In that case we can remove the non-habitat from the state-space and assume uniformity of the activity centers over the remaining portions judged to be suitable habitat. There are two ways to approach this: We can use a regular grid of points to represent the state-space, i.e., by the set of coordinates $\mathbf{s}_1, \dots, \mathbf{s}_G$, and assign a equal probabilities to each possible value, or we can retain the continuous formulation of the state-space but use basic polygon operations to induce constraints on the state-space We focus here on the formulation of our basic SCR model in terms of a discrete state-space but later on (chapter 7 and also Appendix XYZ) we demonstrate the latter approach based on using polygon operations to define an irregular state-space.

Use of a discrete state-space can be computationally expensive in **WinBUGS**. That said, it isn't too difficult to do the MCMC calculations in **R** which we discuss briefly in chapter 7. The **R** package **SPACECAP** (Gopalaswamy et al., 2011) arose from the **R** implementation developed for the application in Royle et al. (2009). As we will see in chapter 6, we must prescribe the state-space by a discrete mesh of points in order to do integrated likelihood and so if we are using a discrete state-space this can be accommodated directly in our code for obtaining MLEs.

While clipping out non-habitat seems like a good idea, its not obvious that we accomplish any biologically reasonable objective by doing so. We might prefer to do it when non-habitat represents a clear-cut restriction on the state-space such as

a reserve boundary or a lake, ocean or river. It makes sense in those situations. Unfortunately, having the capability to do this also causes people to start defining “habitat” vs. “non-habitat” based on their understanding of the system whereas it can’t be known whether the animal being studied has the same understanding. Moreover, differentiating of the landscape by habitat or habitat quality probably affects the geometry and morphology of home ranges much more than the plausible locations of activity centers. That is, a home range centroid could, in actual fact, occur in a walmart parking lot if there is pretty good habitat around walmart, so there is probably no sense to cut out the walmart lot and preclude it as the location for an activity center. It would generally be better to include some definition of habitat quality in the model for the detection probability (see chapter XYZ).

4.9.1 Evaluation of Coarseness of Discrete Approximation

The coarseness of the state-space should not really have much of an effect on estimates if the grain is sufficiently fine relative to typical animal home range sizes. Why is this? We have two analogies that can help us understand this. First is the relationship to Model M_h . As noted in section 4.3.4 above, we can think about SCR models as a type of finite mixture (Norris III and Pollock, 1996; Pledger, 2000) where we are fortunate to be able to obtain direct information about which “group” individuals belong to (group being location of activity center). In the standard finite mixture models we typically find that only 1 or a very small number of groups (e.g., 2 or 3 at the most) can explain really high levels of heterogeneity and are adequate for most data sets of small to moderate sample sizes. We therefore expect a similar effect in SCR models when we discretize the state-space. We can also think about discretizing the state-space as being related to numerical integration where we find (see chapter 6) that we don’t need a very fine grid of support points to evaluate the integral to a reasonable level of accuracy. We demonstrate this here by reanalyzing simulated data using a state-space defined by a different numbers of support points. We provide an R script called `simSCR0discrete.fn` in the **R** package `scrbook`. We note that for this comparison we generated the actual activity centers as a continuous random variable and thus the discrete state-space is, strictly speaking, an approximation to truth. That said, we regard all state-space specifications as approximations to truth because they are all, strictly speaking, models of some unknown truth. Thus the use of any specific discrete state-space is not intrinsically more “wrong” than any specific continuous representation.

We used **JAGS** from the `rjags` function to obtain the results for 6×6 , 9×9 , 12×12 , 15×15 , 20×20 , 25×25 and 30×30 state-space grids. We used 2000 burn, 12000 total iters with 3 chains, therefore a total of 30000 posterior samples. For **WinBUGS** we used 3 chains of 5k total with 1k burnin means 12k total posterior samples. Summary results for these analyses are shown in Table XYZ⁶.

⁶Andy to finish later.

3855 Table XYZ.

3856			Mean	SD	NaiveSE	Time-seriesSE	runtime
3857	6	N	109.7717	15.98959	0.0923160	0.377737	1239
3858	9	N	114.4621	16.72025	0.0965344	0.468659	1267
3859	12	N	115.4309	17.12403	0.098866	0.464830	1576
3860	15	N	114.7699	17.0242	0.0982894	0.425238	1638
3861	20	N	116.0370	17.10686	0.0987665	0.486867	1647
3862	25	N	116.3228	16.98323	0.0980527	0.465527	1661
3863	30	N	116.4252	17.4078	0.100504	0.533735	1806
3864	WinBUGS						
3865			Mean	SD	NaiveSE	Time-seriesSE	runtime
3866	6	N	111.67	16.61			2274
3867	9	N	114.23	17.99			4300
3868	12	N	115.98	17.38			7100
3869	15	N	115.38	17.94			13010

3870
3871 Note: WinBUGS based on fewer samples too!

3872
3873 To get SE and time-series SE do this:

3874 You can use `as.mcmc.list()` to convert to a coda object. Then use `summary`.

3875 The results in terms of the posterior summaries are, as we expect, very similar
3876 using **WinBUGS**. However, it was interesting to note that **WinBUGS** runtime is
3877 much worse (note the number of iterations is lower for **WinBUGS** yet the runtime
3878 is much longer) and, furthermore, it seems to scale with the size of the discrete
3879 state-space grid. While that was expected, it was unexpected that the runtime of
3880 **JAGS** would seem relatively consistent as we increase the grid size. We suspect
3881 that **WinBUGS** is evaluating the full-conditional for each activity center at all
3882 G possible values whereas it may be that **JAGS** is evaluating the full-conditional
3883 only at a subset of values or perhaps using previous calculations more effectively.

3884 While this might suggest that one should always use **JAGS** for this analysis, we
3885 found in our analysis of the wolverine (next section) that **JAGS** could be extremely
3886 sensitive to starting values, producing MCMC algorithms that sometimes simply
3887 did not work.

3888 4.9.2 Analysis of the wolverine camera trapping data

3889 We reanalyzed the wolverine data using discrete state-space grids with points spaced
3890 by 2, 4 and 8 km (depicted in Fig. 4.6). These were constructed from the 40 km
3891 buffered state-space, and deleting the points over water (see Royle et al., 2011c).
3892 Our interest in doing this was to evaluate the relative influence of grid resolution
3893 on estimated density because the coarser grids will be more efficient from a compu-
3894 tational stand-point and so we would prefer to use them, but perhaps not if there




Figure 4.6. 2 km 4 km and 8km wolverine state-space grids extending about 40 km from the vicinity of the trap array.

3895 is a strong influence on estimated density.

3896 **Note:** Results from WinBUGS are given below based on short runs that took
 3897 a long long time. I am rerunning those. I will also show a density map for each
 3898 analysis.

3899 This took about 6 days in WinBUGS. Terrible mixing for the 2km and
 3900 8km. Why is this? We may never know!

3901 > print(out.2km,digits=2)

3902 Inference for Bugs model at "modelfile.txt", fit using WinBUGS,
 3903 3 chains, each with 11000 iterations (first 1000 discarded)

3904 n.sims = 30000 iterations saved

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
psi	0.43	0.09	0.27	0.37	0.43	0.49	0.63	1.00	560
sigma	0.62	0.05	0.54	0.59	0.62	0.65	0.73	1.01	160
lam0	0.05	0.01	0.04	0.04	0.05	0.06	0.07	1.01	320
p0	0.05	0.01	0.03	0.04	0.05	0.05	0.06	1.01	320
N	86.56	16.94	57.00	75.00	85.00	97.00	124.00	1.00	510
D	8.78	1.72	5.78	7.60	8.62	9.83	12.57	1.00	510

3913

3914 For each parameter, n.eff is a crude measure of effective sample size,
 3915 and Rhat is the potential scale reduction factor (at convergence, Rhat=1).

3916 > print(out.4km,digits=2)

3917 Inference for Bugs model at "modelfile.txt", fit using WinBUGS,

```

3918 3 chains, each with 11000 iterations (first 1000 discarded)
3919 n.sims = 30000 iterations saved
3920      mean    sd  2.5%   25%   50%   75%  97.5% Rhat n.eff
3921 psi    0.45  0.09  0.28  0.38  0.44  0.50  0.64    1  1300
3922 sigma 0.61  0.04  0.53  0.58  0.61  0.64  0.71    1  1600
3923 lam0   0.05  0.01  0.04  0.05  0.05  0.06  0.07    1  2500
3924 p0     0.05  0.01  0.03  0.04  0.05  0.05  0.07    1  2500
3925 N      89.25 17.44 59.00 77.00 88.00 100.00 127.00    1  1100
3926 D       9.01  1.76  5.96  7.77  8.88  10.10  12.82    1  1100
3927
3928 For each parameter, n.eff is a crude measure of effective sample size,
3929 and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
3930 > print(out.8km,digits=2)
3931 Inference for Bugs model at "modelfile.txt", fit using WinBUGS,
3932 3 chains, each with 11000 iterations (first 1000 discarded)
3933 n.sims = 30000 iterations saved
3934      mean    sd  2.5%   25%   50%   75%  97.5% Rhat n.eff
3935 psi    0.42  0.09  0.26  0.36  0.41  0.47  0.61  1.00   940
3936 sigma 0.68  0.05  0.59  0.64  0.67  0.71  0.77  1.01   220
3937 lam0   0.05  0.01  0.03  0.04  0.05  0.05  0.06  1.00   560
3938 p0     0.05  0.01  0.03  0.04  0.04  0.05  0.06  1.00   560
3939 N      83.18 16.14 56.00 72.00 82.00 93.00 119.00 1.00   700
3940 D       8.28  1.61  5.57  7.17  8.16  9.26  11.84 1.00   700
3941
3942 For each parameter, n.eff is a crude measure of effective sample size,
3943 and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
3944
3945 We did the analysis in JAGS also. The results are shown below. Note: I am
3946 going to run these again but for longer to finalize the results.
3947
3948 ### 01/10/2012 -- need to rerun these JAGS runs but use more
3949 iterations and check results.
3950
3951 2km
3952 Iterations = 7001:13000
3953 Thinning interval = 1
3954 Number of chains = 3
3955 Sample size per chain = 6000
3956
3957      Mean      SD Naive SE Time-series SE
3958 N      86.28522 16.950626 1.263e-01    0.4878973
3959 lam0    0.04807  0.007512 5.599e-05    0.0002199
3960 p0      0.04581  0.006820 5.083e-05    0.0001996
3961 psi     0.28904  0.062117 4.630e-04    0.0017481
3962 sigma   0.62769  0.043596 3.249e-04    0.0018724

```

3963	4km				
3964		Mean	SD	Naive SE	Time-series SE
3965	N	85.53139	16.998966	1.267e-01	0.5181297
3966	lam0	0.04636	0.007542	5.621e-05	0.0002382
3967	p0	0.04425	0.006867	5.118e-05	0.0002172
3968	psi	0.28650	0.061922	4.615e-04	0.0018276
3969	sigma	0.64281	0.048321	3.602e-04	0.0022911
3970					
3971	8km				
3972		Mean	SD	Naive SE	Time-series SE
3973	N	83.97039	16.508146	1.230e-01	0.4548782
3974	lam0	0.04519	0.006919	5.157e-05	0.0001738
3975	p0	0.04319	0.006319	4.710e-05	0.0001589
3976	psi	0.28146	0.060653	4.521e-04	0.0016555
3977	sigma	0.66956	0.040989	3.055e-04	0.0015070

3978 4.9.3 SCR models as multi-state models

3979 While we invoke a discrete state-space artificially, by gridding the underlying con-
3980 tinuous state-space, sometimes the state-space is more naturally discrete. Consider
3981 a situation in which discrete patches of habitat are searched using some method
3982 and it might be convenient (or occur inadvertently) to associate samples to the
3983 patch level instead of recording observation locations. In this case we might use a
3984 model $\mathbf{s}_i \sim \text{dcat}(\text{probs}[])$ where $\text{probs}[]$ are the probabilities that an individual in-
3985 habits a particular patch. We consider such a case study in chapter XXPoissonXXX
3986 from Mollet et al. (2012) who obtained a population size estimate of a large grouse
3987 species known as the capracaille. Forest patches were searched for scat which was
3988 identified to individual by DNA analysis. Even when space is *not* naturally discrete,
3989 measurements are often made at a fairly coarse grain (e.g., meters or tens of meters
3990 along a stream), or associated with spatial quadrats for scat searches and therefore
3991 the state-space may be effectively discrete in many situations.

3992 This discrete formulation of SCR models suggests that SCR models are related
3993 to ordinary multi-state models (Kery and Schaub, 2011, ch. 9) which are also
3994 parameterized in terms of a discrete state variable which is often defined as a
3995 spatially-indexed state related either to location of capture or breeding location.
3996 While many multi-state models exist in which the state variable is not related to
3997 space, multi-state models have been extremely useful in development models of
3998 movements among geographic states and indeed this type of problem motivated
3999 their early developments by Arnason (1973, 1974) and Hestbeck (1991). We pursue
4000 this connection a little bit more in chapter XXX XYZ.

4.10 SUMMARY AND OUTLOOK

A point we tried to emphasize in this chapter is that the basic SCR model is not much more than an ordinary capture-recapture model for closed populations – it is simply that model but augmented with a set of “individual effects”, \mathbf{s}_i , which relate encounter probability to some sense of individual location. SCR models are therefore a type of individual covariate model (as introduced in chapter 3 – but with imperfect information about the individual covariate. In other words, they are GLMM type models when N is known or, when N is unknown, they are zero-inflated GLMMs (see Royle (2006)). Another class of capture-recapture models that SCR models are closely related to is so-called “Model M_h .” The effect of introducing a spatial location for individuals is that it induces heterogeneity in detection probability, as in Model M_h . However, unlike Model M_h , we obtain some information about the individual effect which is completely latent in Model M_h . If the state-space of the random effect \mathbf{s} is discrete then the SCR model resembles more closely the finite-mixture class of heterogeneity models (Norris III and Pollock, 1996) which parameterizes heterogeneity by assuming that individuals belong to discrete classes or groups (e.g., high, medium, low). In the context of SCR models we obtain some information about the “group membership” in the locations where individuals are captured. Given the direct relationship of SCR models with so many standard classes of models, we find that they are really quite easy to analyze using standard MCMC methods encased in black boxes such as **WinBUGS** or **JAGS** and possibly other packages. They are also easy to analyze using classical likelihood methods, which we address in chapter 6.

Formal consideration of the collection of individual locations $(\mathbf{s}_1, \dots, \mathbf{s}_N)$ in the model is fundamental to all of the models considered in this book. In statistical terminology, we think of the collection of points $\{\mathbf{s}_i\}$ as a realization of a point process and part of the promise, and ongoing challenge, of SCR models is to develop models that reflect interesting biological processes, for example interactions among points or temporal dynamics in point locations. Here we considered the simplest possible point process model – the points are independent and uniformly (“randomly”) distributed over space. Despite the simplicity of this assumption, it should suffice in many applications of SCR models although we do address generalizations of this model in later chapters. Moreover, even though the *prior* distribution on the point locations is uniform, the realized pattern may deviate markedly from uniformity as the observed encounter data provide information to impart deviations from uniformity. Thus, the estimated density map will typically appear distinctly non-uniform. As a general rule, information in the data will govern estimates of individual point locations so even fairly complex patterns of non-independence or non-uniformity will appear in the data. That is, we find in applications of the basic SCR model that this simple *a priori* model can effectively reflect or adapt to complex realizations of the underlying point process. For example, if individuals are highly territorial then the data should indicate this in the form of individuals not

being encountered in the same trap - the resulting posterior distribution of point locations should therefore reflect non-independence. Obviously the complexity of posterior estimates of the point pattern will depend on the quantity of data, both number of individuals and captures per individual. Because the point process is such an integral component of SCR models, the state-space of the point process plays an important role in developing SCR models. As we tried to emphasize in this chapter, the choice of the state-space is part of the model. It can have an influence on parameter estimates and other inferences such as model selection (see chapter 8). We emphasize however that this is not an arbitrary decision like “buffering” because the model induces an explicit interpretation of parameters and statistical effect on estimators.

We showed how to conduct inference about the underlying point process including calculation of density maps from posterior output. We can do other things we normally do with spatial point processes such as compute “K-functions” and test for “complete spatial randomness” (CSR) which we develop in chapter 8. Modifying and applying point process methods to SCR problems seems to us to be a fruitful area of research.

An obvious question that might be floating around in your mind is why should we ever go through all of this trouble when we could just use **MARK** or **CAPTURE** to get an estimate of N and apply 1/2 MMDM methods? The main reason is that these conventional methods are predicated on models that represent explicit misspecifications of both the observation and ecological process - they are wrong! Not just wrong, because of course all models are wrong, but they’re not even *plausible* models! Thus while we might be able to show adequate fit or whatever, we think as a conceptual and philosophical model one should not be using models that are not even plausible data-generating models - even if the plausible ones don’t fit! Perhaps more charitably, these ordinary non-spatial models are models of the wrong system. They do not account for trap identity. They don’t account for spatial organization or “clustering” of individual encounters in space. And, “density” is not a parameter of those models because density has no meaning absent an explicit representation of space. If we do define space explicitly, e.g., as a buffered minimum convex hull, then the normal models (M_0 , M_h , etc..) assume that individual capture-probability is not related to space, no matter how we define the buffer. Conversely, the SCR model is a model for trap-specific encounter data - how individuals are organized in space and interact with traps. SCR models provide a coherent framework for inference about density or population size and also, because of the formality of their derivation, can be extended and generalized to a large variety of different situations, as we demonstrate in subsequent chapters.

In the next few chapters we continue to work with this basic SCR design and model but consider some important extensions of the basic model. For example, we consider extensions to include covariates that vary by individual, trap, or over time (chapter 9), spatial covariates on density (chapter ??), open populations (chapter 12), model assessment and selection (chapter 8) and other topics. We also consider

4085 technical details of Bayesian (chapter 7) and maximum likelihood (chapter 6) esti-
4086 mation so that the interested reader can develop or extend their own methods to
4087 suit their needs.

5

4088
4089
4090

OTHER OBSERVATION MODELS

6

LIKELIHOOD ANALYSIS OF SPATIAL CAPTURE-RECAPTURE MODELS

In this book we mainly focus on Bayesian analysis of spatial capture-recapture models. And, in the previous chapters we learned how to fit some basic spatial capture-recapture models using a Bayesian formulation of the models analyzed in BUGS engines including **WinBUGS** and **JAGS**. Despite our focus on Bayesian analysis, it is instructive to develop the basic conceptual and methodological ideas behind classical analysis based on likelihood methods and frequentist inference. In fact, simple SCR models can be analyzed fairly easily using such methods. This has been the approach taken by Borchers and Efford (2008); Dawson and Efford (2009) and related papers.

This chapter provides some conceptual and technical footing for likelihood-based analysis of spatial capture-recapture models. We recognized earlier (chapt. 4) that SCR models are versions of binomial (or other) GLMs, but with random effects i.e., GLMMs. These models are routinely analyzed by likelihood methods. In particular, likelihood analysis is based on the integrated likelihood in which the random effects are removed by integration from the likelihood. In SCR models, the random effect, \mathbf{s} , i.e., the 2-dimensional coordinate, is a bivariate random effect.

In this chapter, we show that it is straightforward to compute the maximum likelihood estimates (MLE) for SCR models by integrated likelihood. We develop the MLE framework using **R**, and we also provide a basic introduction to an **R** package **secr** (Efford, 2011) which is based on the stand-alone package **DENSITY** (Efford et al., 2004). To set the context we analyze the SCR model here when N is known because, in that case, it is precisely a GLMM and does not pose any difficulty at all. We generalize the model to allow for unknown N using both conventional ideas based on the “joint likelihood” (e.g., Borchers et al., 2002) and also using

4119 a formulation based on data augmentation. We obtain the MLEs for the SCR
 4120 model from the wolverine camera trapping study (Magoun et al., 2011) analyzed
 4121 in previous chapters to compare/contrast the results.

6.1 LIKELIHOOD ANALYSIS

4122 We noted in chapter 4 that, with N known, the basic SCR model is a type of
 4123 binomial regression with a random effect. For such models we can easily obtain
 4124 maximum likelihood estimators of model parameters based on integrated likelihood.
 4125 The integrated likelihood is based on the marginal distribution of the data y in
 4126 which the random effects are removed by integration. Conceptually, our model is a
 4127 specification of the conditional-on- \mathbf{s} model $[y|\mathbf{s}, \theta]$ and we have a “prior distribution”
 4128 for \mathbf{s} , say $[\mathbf{s}]$, and the marginal distribution of the data y is

$$[y|\theta] = \int_{\mathbf{s}} [y|\mathbf{s}, \theta][\mathbf{s}]d\mathbf{s}.$$

4129 When viewed as a function of θ for purposes of estimation, the marginal distribu-
 4130 tion $[y|\theta]$ is often referred to as the *integrated likelihood*.

4131 It is worth analyzing the simplest SCR model with known- N in order to un-
 4132 derstand the underlying mechanics and basic concepts. These are directly relevant
 4133 to the manner in which many capture-recapture models are classically analyzed,
 4134 such as model Mh, and individual covariate models (see chapt. 3 and Royle and
 4135 Dorazio (2008, chapt. 6)). To develop integrated likelihood for SCR models, we
 4136 first identify the conditional likelihood.

4137 The observation model for each encounter observation y_{ij} , specified conditional
 4138 on \mathbf{s}_i , is

$$y_{ij}|\mathbf{s}_i \sim \text{Bin}(K, p_{\theta}(\mathbf{x}_j, \mathbf{s}_i)) \quad (6.1.1)$$

4139 where we have indicated the dependence of p_{ij} on \mathbf{s} and parameters θ explicitly.
 4140 For the random effect we have $\mathbf{s}_i \sim \text{Unif}(S)$. The joint distribution of the data
 4141 for individual i is the product of J such terms (i.e., contributions from each of J
 4142 traps).

$$[\mathbf{y}_i|\mathbf{s}_i, \theta] = \prod_{j=1}^J \text{Bin}(K, p_{\theta}(\mathbf{x}_j, \mathbf{s}_i))$$

4143 We note that this assumes that encounter of individual i in each trap is independent
 4144 of encounter in every other trap, conditional on \mathbf{s}_i , this is the fundamental property
 4145 of SCR0 or “multi-catch” traps.

4146 The so-called marginal likelihood is computed by removing \mathbf{s}_i , by integration
 4147 (hence also *integrated likelihood*), from the conditional-on- \mathbf{s} likelihood and regarding
 4148 the *marginal* distribution of the data as the likelihood. That is, we compute:

$$[y|\theta] = \int_S [\theta|\mathbf{y}_i|\mathbf{s}_i]g(\mathbf{s}_i)d\mathbf{s}_i$$

4149 In most SCR models, $g(\mathbf{s}) = 1/|\mathcal{S}|$ (but see chapt. ??).

4150 The joint likelihood for all N individuals, assuming independence of encounters
4151 among individuals, is the product of N such terms:

$$\mathcal{L}(\theta|\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_N) = \prod_{i=1}^N [\mathbf{y}_i|\theta]$$

4152 We emphasize that two independence assumptions are explicit in this development:
4153 independence of trap-specific encounters within individuals and also independence
4154 among individuals. In particular, this would only be valid when individuals are not
4155 physically restrained or removed upon capture, and when traps do not “fill up”.

4156 The key operation for computing the likelihood is solving a 2-dimensional in-
4157 tegration problem. There are some general purpose **R** packages that implement
4158 a number of multi-dimensional integration routines including **adapt** (Genz et al.,
4159 2007) and **R2cuba** (Hahn et al., 2011). In practice, we won’t rely on these extrane-
4160 ous **R** packages (except see chapt. ?? for an application of **Rcuba**) but instead will
4161 use perhaps less efficient methods in which we replace the integral with a summa-
4162 tion over an equal area mesh of points on the state-space \mathcal{S} and explicitly evaluate
4163 the integrand at each point. We invoke the rectangular rule for integration here¹
4164 in which we evaluate the integrand on a regular grid of points of equal area and
4165 compute the average of the integrand over that grid of points. Let $u = 1, 2, \dots, nG$
4166 index a grid of nG points, \mathbf{s}_u , where the area of grid cell u is constant, say A . In
4167 this case, the integrand, i.e., the marginal pmf of \mathbf{y}_i , is approximated by

$$[\mathbf{y}_i|\theta] = \frac{1}{nG} \sum_{u=1}^{nG} [\mathbf{y}_i|\mathbf{s}_u, \theta] \quad (6.1.2)$$

4168 This is a specific case of the general expression that could be used for approxi-
4169 mating the integral for any arbitrary (bivariate or otherwise) distribution $g(\mathbf{s})$. The
4170 general case is

$$[y] = \frac{A}{nG} \sum_u [y|\mathbf{s}_u][\mathbf{s}_u]$$

4171 In the present context it happens that $[\mathbf{s}] = (1/A)$ and thus the grid-cell area
4172 cancels in the above expression to yield eq. 6.1.2, but we commonly apply this in
4173 the context of normal prior distributions, as we did for likelihood analysis of Model
4174 M_h in Sec. 3.4). The rectangular rule for integration can be seen as an application
4175 of the Law of Total Probability for a discrete random variable \mathbf{s} , having nG unique
4176 values with equal probabilities $1/nG$.

¹e.g., http://en.wikipedia.org/wiki/Rectangle_method

6.1.1 Implementation (simulated data)

Here we will illustrate how to carryout this integration and optimization based on the integrated likelihood using simulated data (i.e., following that from Chapter 4). Using `simSCRO.fn` we simulate data for 100 individuals and a 25 trap array layed out in a 5×5 grid of unit spacing. The specific encounter model is the half-normal model. The 100 activity centers were simulated on a state-space defined by a 8×8 square within which the trap array was centered (thus the trap array is buffered by 2 units). Therefore, the density of individuals in this system is fixed at $100/64$.

In the following set of R commands we generate the data and then harvest the required data objects:

```
4187 data<-simSCRO.fn(discard0=FALSE,sd=2013)
4188 y<-data$Y
4189 traplocs<-data$traplocs
4190 nind<-nrow(y)
4191 X<-data$traplocs
4192 J<-nrow(X)
4193 K<-data$K
4194 Xl<-data$xlim[1]
4195 Yl<-data$ylim[1]
4196 Xu<-data$xlim[2]
4197 Yu<-data$ylim[2]
```

Now we need to define the integration grid, say **G**, which we do with the following set of **R** commands (here, `delta` is the grid spacing):

```
4200 delta<- .2
4201 xg<-seq(Xl+delta/2,Xu-delta/2,by=delta)
4202 yg<-seq(Yl+delta/2,Yu-delta/2,by=delta)
4203 npix<-length(xg)      # assumes xg and yg same dimension here
4204 area<- (Xu-Xl)*(Yu-Yl)/((npix)*(npix)) # dont need area for anything
4205 G<-cbind(rep(xg,npix),sort(rep(yg,npix)))
4206 nG<-nrow(G)
```

In this case, the integration grid is set up as a grid with spacing $\delta = 0.2$ which produces a 40×40 grid of points for evaluating the integrand if the state-space buffer is set at 2.

We next create an **R** function that defines the likelihood as a function of the data objects `y` and `X` which were created above but, in general, you would read these files into **R**, e.g., from a .csv file. In addition to these data objects, we need to have defined the quantities `G` and `nG` associated with the integration grid. However, instead of worrying about making all of these objects and keeping track of them we just put that code above into the likelihood function and pass δ as an additional (optional) argument and a few other things that we need such as the

boundary of the state-space over which the integration (summation) is being done. Here is one reasonably useful variation of a function for estimation based on the integrated likelihood:

```

4220 intlik1<-function(parm,y=y,delta=.2,X=traplocs,ssbuffer=2){
4221
4222   Xl<-min(X[,1]) - ssbuffer
4223   Xu<-max(X[,1]) + ssbuffer
4224   Yu<-max(X[,2]) + ssbuffer
4225   Yl<-min(X[,2]) - ssbuffer
4226
4227   xg<-seq(Xl+delta/2,Xu-delta/2,,length=npix)
4228   yg<-seq(Yl+delta/2,Yu-delta/2,,length=npix)
4229   npix<-length(xg)
4230
4231   G<-cbind(rep(xg,npix),sort(rep(yg,npix)))
4232   nG<-nrow(G)
4233   D<- e2dist(X,G)
4234
4235   alpha<-parm[1]
4236   theta<-parm[2]
4237   probcap<- plogis(alpha)*exp(-theta*D*D)
4238   Pm<-matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
4239           # all zero encounter histories
4240   n0<-sum(apply(y,1,sum)==0)
4241           # encounter histories with at least 1 detection
4242   ymat<-y[apply(y,1,sum)>0,]
4243   ymat<-rbind(ymat,rep(0,ncol(ymat)))
4244   lik.marg<-rep(NA,nrow(ymat))
4245   for(i in 1:nrow(ymat)){
4246     Pm[1:length(Pm)]<- (dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],log=TRUE))
4247     lik.cond<- exp(colSums(Pm))
4248     lik.marg[i]<- sum( lik.cond*(1/nG))
4249   }
4250   nv<-c(rep(1,length(lik.marg)-1),n0)
4251   -1*( sum(nv*log(lik.marg)) )
4252 }

```

The function accepts as input the encounter history matrix, y , the trap locations, X , and the state-space buffer. This allows us to vary the state-space buffer and easily evaluate the sensitivity of the MLE to the size of the state-space. Note that we have a peculiar handling of the encounter history matrix y . In particular, we remove the all-zero encounter histories from the matrix and tack-on a single all-zero encounter history as the last row which then gets weighted by the number of such encounter histories ($n0$). This is a bit long-winded and strictly unnecessary when N is known, but we did it this way because the extension to the unknown- N case

is now transparent (as we demonstrate in the following section). The matrix Pm holds the log-likelihood contributions of each encounter frequency for each possible state-space location of the individual. The log contributions are summed up and the result exponentiated on the next line, producing `lik.cond`, the conditional-on-s likelihood (Eq. 6.1.1 above). The marginal likelihood (`lik.marg`) sums up the conditional elements weighted by $\Pr(\mathbf{s})$ (Eq. 6.1.2 above). This is a fairly primitive function which doesn't allow much flexibility in the data structure. For example, it assumes that K , the number of replicates, is constant for each trap. Further, it assumes that the state-space is a square. We generalize this to some extent later in this chapter.

Here is the **R** command for maximizing the likelihood and saving the results into an object called `frog`. The output is a list of the following structure and these specific estimates are produced using the simulated data set:

```
# should take 15-30 seconds

> starting.values <- c(-2, 2)
> frog<-nlm(intlik1,starting.values,y=y,delta=.1,X=traplocs,ssbuffer=2,hessian=TRUE)
> frog

$minimum
[1] 297.1896

$estimate
[1] -2.504824  2.373343

$gradient
[1] -2.069654e-05  1.968754e-05

$hessian
      [,1]      [,2]
[1,] 48.67898 -19.25750
[2,] -19.25750 13.34114

$code
[1] 1

$iterations
[1] 11
```

Details about this output can be found on the help page for `nlm`. We note briefly that `frog$minimum` is the negative log-likelihood value at the MLEs, which are stored in the `frog$estimate` component of the list. The hessian is the observed Fisher information matrix, which can be inverted to obtain the variance-covariance matrix using the commands:

```
4304 > solve(frog$hessian)
```

4305 It is worth drawing attention to the fact that the estimates are different than
 4306 the Bayesian estimates reported previously in chapt. 4. How can that be?! There
 4307 are several reasons for this. First Bayesian inference is based on the posterior
 4308 distribution and it is not generally the case that the MLE should correspond to any
 4309 particular value of the posterior distribution. If the prior distributions in a Bayesian
 4310 analysis are uniform, then the (multivariate) mode of the posterior is the MLE,
 4311 but note that Bayesians almost always report posterior *means* and so there will
 4312 typically be a discrepancy there. Secondly, we have implemented an approximation
 4313 to the integral here and there might be a slight bit of error induced by that. We
 4314 will evaluate that shortly. Third, the Bayesian analysis by MCMC is subject to
 4315 some amount of Monte Carlo error which the analyst should always be aware of
 4316 in practical situations. All of these different explanations are likely responsible for
 4317 some of the discrepancy. Accounting for these, we see general consistency between
 4318 the two estimates.

4319 To compute the integrated likelihood we used a discrete representation of the
 4320 state-space so that the integral could be approximated as a summation over possible
 4321 values of \mathbf{s} with each value being weighted by its probability of occurring, which
 4322 is $1/nG$ under the assumption that \mathbf{s} is uniform on the state-space \mathcal{S} . Recall in
 4323 chapt. 4 we used a discrete state-space in developing a Bayesian analysis of the
 4324 model in order to be able to modify the state-space in a flexible manner. In that
 4325 case, we could use the discretized state-space as the integration grid and just feed
 4326 it into our integrated likelihood routine.

4327 In summary, we note that, for the basic SCR model, integrated likelihood is
 4328 a really easy calculation when N is known. Even for N unknown it is not too
 4329 difficult, and we will do that shortly. However, if you can solve the known- N
 4330 problem then you should be able to do a real analysis, for example by considering
 4331 different values of N and computing the results for each value and then making a
 4332 plot of the log-likelihood or AIC and choosing the value of N that produces the
 4333 best log likelihood or AIC. As a homework problem we suggest that the reader take
 4334 the code given above and try to estimate N without modifying the code by just
 4335 repeatedly calling that code for different values of N and trying to deduce the best
 4336 value. Nevertheless, we will formalize the unknown- N problem shortly.

4337 The software package **DENSITY** (Efford et al., 2004) implements certain types
 4338 of SCR models using integrated likelihood methods. **DENSITY** has been made
 4339 into an **R** package called **secr** (Efford, 2011) and we provide an analysis of some
 4340 data using **secr** shortly along with a discussion of its capabilities.

6.2 MLE WHEN N IS UNKNOWN

4341 Here we build on the previous introduction to integrated likelihood but we consider
 4342 now the case in which N is unknown. We will see that adapting the analysis based

on the N -known model is really straightforward for the more general problem. The main distinction is that we don't observe the all-zero encounter history so we have to make sure we compute the probability for that encounter history which we do by tacking a row of zeros onto the encounter history matrix. In addition, we include the number of such all-zero encounter histories as an unknown parameter of the model. Call that unknown quantity n_0 . In addition, we have to be sure to include a combinatorial term to account for the fact that of the n observed individuals there are $\binom{N}{n}$ ways to realize a sample of size n . The combinatorial term involves the unknown n_0 and thus it must be included in the likelihood.

Operationally then, things proceed much as before: We compute the marginal probability of each observed \mathbf{y}_i , i.e., by removing the latent \mathbf{s}_i by integration. In addition, we compute the marginal probability of the “all-zero” encounter history \mathbf{y}_{n+1} , and make sure to weight it n_0 times. We accomplish this by “padding” the data set with a single encounter history having $y_{n+1,j} = 0$ for all traps $j = 1, 2, \dots, J$. Then we be sure to include the combinatorial term in the likelihood or log-likelihood computation. We demonstrate this shortly.

To analyze a specific case, we'll read in our fake data set (simulated using the parameters given above). To set some things up in our workspace we do this:

```
data<-simSCRO.fn(discard0=TRUE,sd=2013)
y<-data$Y
nind<-nrow(y)
X<-data$traplocs
J<-nrow(X)
K<-data$K
Xl<-data$xlim[1]
Yl<-data$ylim[1]
Xu<-data$xlim[2]
Yu<-data$ylim[2]
```

Recall that these data were generated with $N = 100$, on an 8×8 unit state-space representing the trap locations (\mathbf{X}) buffered by 2 units. As before, the likelihood is defined in the **R** workspace as an **R** function which takes an argument being the unknown parameters of the model and additional arguments as prescribed. In particular, we provide the encounter history matrix \mathbf{y} , the trap locations `traplocs`, the spacing of the integration grid (δ) and the state-space buffer. Here is the new likelihood function:

```
intlik2<-function(parm,y=y,delta=.3,X=traplocs,ssbuffer=2){
  Xl<-min(X[,1]) -ssbuffer
  Xu<-max(X[,1])+ ssbuffer
  Yu<-max(X[,2])+ ssbuffer
  Yl<-min(X[,2])- ssbuffer
```

```

4384
4385 #delta<- (Xu-Xl)/npix
4386 xg<-seq(Xl+delta/2,Xu-delta/2,delta)
4387 yg<-seq(Yl+delta/2,Yu-delta/2,delta)
4388 npix.x<-length(xg)
4389 npix.y<-length(yg)
4390 area<- (Xu-Xl)*(Yu-Yl)/((npix.x)*(npix.y))
4391 G<-cbind(rep(xg,npix.y),sort(rep(yg,npix.x)))
4392 nG<-nrow(G)
4393 D<- e2dist(X,G)
4394
4395 alpha<-parm[1]
4396 theta<-parm[2]
4397 n0<-exp(parm[3])
4398 probcap<- plogis(alpha)*exp(-theta*D*D)
4399 Pm<-matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
4400 ymat<-rbind(y,rep(0,ncol(y)))
4401
4402 lik.marg<-rep(NA,nrow(ymat))
4403 for(i in 1:nrow(ymat)){
4404   Pm[1:length(Pm)]<- (dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],log=TRUE))
4405   lik.cond<- exp(colSums(Pm))
4406   lik.marg[i]<- sum( lik.cond*(1/nG) )
4407 }
4408 nv<-c(rep(1,length(lik.marg)-1),n0)
4409 part1<- lgamma(nrow(y)+n0+1) - lgamma(n0+1)
4410 part2<- sum(nv*log(lik.marg))
4411 -1*(part1+ part2)
4412 }

```

4413 To execute this function for the data that we created with `simSCRO.fn`, we
 4414 execute the following command (saving the result in our friend `frog`). This re-
 4415 sults in the usual output, including the parameter estimates, the gradient, and the
 4416 numerical Hessian which is useful for obtaining asymptotic standard errors (see
 4417 below):

```

4418 > frog<-nlm(intlik2,c(-2.5,2,log(4)),hessian=TRUE,y=y,X=X,delta=.2,ssbuffer=2)
4419 There were 50 or more warnings (use warnings() to see the first 50)
4420 >
4421 >
4422 > frog
4423 $minimum
4424 [1] 113.5004
4425

```

```

4426 $estimate
4427 [1] -2.538334  2.466515  4.232810
4428
4429 [. Additional output deleted .]

```

While this produces some **R** warnings, these happen to be harmless in this case, and we will see from the **nlm** output that the algorithm performed satisfactory in minimizing the objective function. The estimate of population size for the state-space (using the default state-space buffer) is

```

4434 > nrow(y)+exp(4.2328)
4435 [1] 110.9099

```

Which differs from the data-generating value ($N = 100$) as we might expect. We usually will present an estimate of uncertainty associated with this MLE which we can obtain by inverting the Hessian. Note that $Var(\hat{N}) = n + Var(\hat{n}_0)$. Since we have parameterized the model in terms of $\log(n_0)$ we use a delta approximation to obtain the variance on the scale of n_0 as follows:

```

4441 > (exp(4.2328)^2)*solve(frog$hessian)[3,3]
4442 [1] 260.2033
4443 > sqrt(260)
4444 [1] 16.12452

```

Therefore, the asymptotic “Wald-type” confidence interval for N is $110.91 + / - 1.96 \times 16.125 = (79.305, 142.515)$. To report this in terms of density, we scale appropriately by the area of the prescribed state-space which is 64 units of area (i.e., an 8×8 square).

6.2.1 Exercises

1. Run the analysis with different state-space buffers and comment on the result.
2. Conduct a brief simulation study using this code by simulating 100 data sets and obtain the MLEs for each data set. Do things seem to be working as you expect?
3. Further extensions: It should be straightforward to generalize the integrated likelihood function to accommodate many different situations. For examples, if we want to include more covariates in the model we can just add stuff to the object **probcap**, and add the relevant parameters to the argument that gets passed to the main function. For the simulated data, make up a covariate by generating a Bernoulli covariate (“trap type” perhaps baited or not baited) randomly and try to modify the likelihood to accommodate that.
4. We would probably be interested in devising the integrated likelihood for the full 3-d encounter history array so that we could include temporally varying covariates. This is not difficult but naturally will slow down the execution substantially. The interested reader should try to expand the capabilities of this basic **R** function.

4464 6.2.2 Integrated Likelihood using the model under data augmentation

4465 Note that this likelihood analysis is based on the standard likelihood in which N
 4466 (or n_0) is an explicit parameter. This is usually called the “joint likelihood” or
 4467 “unconditional likelihood”. We could also express the joint likelihood using data
 4468 augmentation, replacing the parameter N with ψ (e.g., see Sec. 7.1.6 Royle and
 4469 Dorazio, 2008, for an example). We don’t go into detail here, but we note that the
 4470 likelihood under data augmentation is a zero-inflated binomial mixture precisely
 4471 an occupancy type model (Royle, 2006). Thus, while it is possible to carryout
 4472 likelihood analysis of models under data augmentation, we primarily advocate data
 4473 augmentation for Bayesian analysis.

4474 6.2.3 Extensions

4475 We have only considered basic SCR models with no additional covariates. However,
 4476 in practice, we are interested in other types of covariate effects including “behavioral
 4477 response”, sex-specificity of parameters, and potentially other effects. Some of these
 4478 can be added directly to the likelihood if the covariate is fixed and known for all
 4479 individuals captured or not. An example is a behavioral response, which amounts
 4480 to having a covariate $x_{ik} = 1$ if individual i was captured prior to occasion k and
 4481 $x_{ik} = 0$ otherwise. For uncaptured individuals, $x_{ik} = 0$ for all k . Royle et al.
 4482 (2011c) called this a global behavioral response because the covariate is defined
 4483 for all traps, no matter the trap in which an individual was captured. We could
 4484 also define a *local* behavioral response which occurs at the level of the trap, i.e.,
 4485 $x_{ijk} = 1$ if individual i was captured in trap j prior to occasion k , etc.. Trap-
 4486 specific covariates such as trap type or status, or time-specific covariates such as
 4487 date, are easily accommodated as well. As an example, Kéry et al. (2010) develop
 4488 a model for the European wildcat in which traps are either baited or not (a trap-
 4489 specific covariate with only 2 values), and also encounter probability varies over time
 4490 in the form of a quadratic seasonal response. We consider models with behavioral
 4491 response or fixed covariates in chapter XXXX, although the integrated likelihood
 4492 routines we provided above can be modified directly for such cases, which we leave
 4493 to the interested reader.

4494 Sex-specificity is more difficult to deal with since sex is not known for uncaptured
 4495 individuals (and sometimes not even for all captured individuals). To analyze
 4496 such models, we do Bayesian analysis of the joint likelihood facilitated by the use of
 4497 data augmentation (Gardner et al., 2010; R.E. et al., 2012). For covariates that are
 4498 not fixed and known for all individuals, it is somewhat more challenging to do MLE
 4499 for these based on the joint likelihood as we have developed above. Instead it is
 4500 more conventional to use what is colloquially referred to as the “Huggins-Alho” type
 4501 model which is one of the approaches taken in the software package `secr` (Efford,
 4502 2011, see Sec. 6.5). This idea is motivated by thinking about unequal probability
 4503 sampling methods known as Horvitz-Thompson sampling (e.g., see Overton and

4504 Stehman, 1995). We don't use that method anywhere in this book because it rep-
 4505 represents a paradigm shift in the inference framework which is done historically only
 4506 for convenience (i.e., ease of constructing an estimator) and not for philosophical
 4507 or theoretical reasons.

6.3 CLASSICAL MODEL SELECTION AND ASSESSMENT

4508 In most analyses, one is interested in choosing from among various potential mod-
 4509 els. A good thing about classical analysis based on likelihood is we can apply AIC
 4510 methods without difficulty (Burnham and Anderson, 2002). There are two distinct
 4511 contexts for model-selection that we think are relevant to SCR models. First is se-
 4512 lecting among models that represent distinct biological hypotheses (e.g., covariates
 4513 affecting encounter probability or density), and AIC is convenient for assessing the
 4514 relative merits of these different models although if there are only a few models
 4515 it is not objectionable to use hypothesis tests or confidence intervals to determine
 4516 importance of effects. The second context is selecting among various detection func-
 4517 tions. Indeed, when distance is used as a covariate (e.g., distance sampling), AIC
 4518 is usually applied to some large and arbitrary selection of distance functions with
 4519 no biological motivation. As a general rule, we don't recommend this given there
 4520 is hardly ever (if at all) a rational subject-matter based reason motivating specific
 4521 distance functions. As a result, we believe that doing too much model selection
 4522 will invariably lead to over-fitting and thus over-statement of precision. This is the
 4523 main reason that we haven't loaded you down with a basket of models for detection
 4524 probability so far, although we discuss many possibilities in chapter XYZ.

4525 Goodness-of-fit: For many standard capture-recapture models, it is possible to
 4526 identify goodness-of-fit statistics based on the multinomial likelihood and evaluate
 4527 model adequacy using formal statistical tests. Similar strategies can be applied
 4528 to SCR models using expected cell-frequencies based on the marginal distribution
 4529 of the observations. Also, because computing MLEs is somewhat more efficient in
 4530 many cases compared to Bayesian analysis, it is also sometimes easy to use boot-
 4531 strap methods². Bayesian goodness-of-fit is almost always addressed with Bayesian
 4532 p-values or some other posterior predictive check (chapter 2 REF XXX and see
 4533 chapter 8). Royle et al. (2011b) suggested checking model fit by decomposing
 4534 fit into two components: an evaluation of the encounter process model based on
 4535 expected encounter frequencies computed *conditional* on \mathbf{s} , and then independent
 4536 evaluation of the "spatial randomness" hypothesis. We discuss this in chapter 8.

6.4 LIKELIHOOD ANALYSIS OF THE WOLVERINE CAMERA TRAPPING DATA

4537 Here we compute the MLEs for the wolverine data using an expanded version of
 4538 the function we developed in the previous section. To accommodate that each trap

²I could use some references in the context of SCR models for this stuff

might be operational a variable number of nights, we provided an additional argument to the likelihood function (allowing for a vector K), which requires also a modification to the construction of the likelihood. In addition, we had to accommodate that the state-space is a general rectangle, and we included a line in the code to compute the state-space area which we apply below for computing density. The more general function (`intlik3`) is given in the **R** package. It has a general purpose wrapper named `scr`³ which has other capabilities too.

The data were read into our R session and manipulated using the following commands. Note that we use the utility **R** function `SCR23darray.fn` which we defined in chapt. 4.

```

4539 > wcaps<-source("wcaps.R")$value
4540 > wtraps<-source("wtraps.R")$value
4541 > K.wolv<-apply(wtraps[,4:ncol(wtraps)],1,sum)
4542 >
4543 > xx<-SCR23darray.fn(wcaps,ntraps=37,nperiods=165)
4544 > y.wolv<- apply(xx,c(1,3),sum)
4545 > traplocs.wolv<-wtraps[,2:3]
4546 > traplocs.wolv<-traplocs.wolv/10000
4547 >
4548 > frog<-nlm(intlik3,c(-1.5,1.2,log(4)),hessian=TRUE,y=y.wolv,K=K.wolv,X=traplocs.wolv,delt
4549 There were 23 warnings (use warnings() to see them)
4550 > frog
4551
4552 $minimum
4553 [1] 220.4355
4554
4555 $estimate
4556 [1] -2.817570  1.255112  3.599040
4557
4558 $gradient
4559 [1] -6.274309e-06  2.146722e-05 -1.045566e-05
4560
4561 $hessian
4562      [,1]      [,2]      [,3]
4563 [1,] 37.687931 -11.852236  4.688911
4564 [2,] -11.852236 30.846144 -9.199113
4565 [3,]  4.688911 -9.199113 13.050428
4566
4567 $code
4568 [1] 1
4569

```

³Not written yet

```

4580 $iterations
4581 [1] 12
4582
4583 > exp(3.599)*sqrt(solve(frog$hessian)[3,3])
4584 [1] 11.41059
4585 >
4586

```

4587 We obtained the MLEs for a state-space buffer of 2 (standardized units) and
 4588 for integration grid with spacing $\delta = .3, .2, .1, .05$. The MLEs for these 4 cases
 4589 including the relative runtime are given in Table 6.1.

Table 6.1. Run time and MLEs for different integration grid resolutions.

δ	Estimates			
	runtime	α_0	θ	$\log(n_0)$
0.30	9.9	-2.819786	1.258468	3.569731
0.20	32.3	-2.817610	1.254757	3.583690
0.10	115.1	-2.817570	1.255112	3.599040
0.05	407.3	-2.817559	1.255281	3.607158

4590 We see the results change only slightly as the fineness of the integration grid
 4591 increases. Conversely, the runtime on the platform of the day for the 4 cases
 4592 increases rapidly. As we have suggested previously these runtimes could be regarded
 4593 in relative terms, across platforms, for gaging the decrease in speed as the fineness
 4594 of the integration grid increases. The effect of this is that we anticipate some
 4595 numerical error in approximating the integral on a mesh of points, and that error
 4596 increases as the coarseness of the mesh increases.

4597 We studied the effect of the state-space buffer on the MLEs, using a fixed $\delta = .2$
 4598 for all analyses. We used state-space buffers of 1 to 4 units stepped by .5. This
 4599 produced the following results, given here are the state-space buffer, area of the
 4600 state-space, the MLE of N for the prescribed state-space and the corresponding
 4601 MLE of density:

	ssbuff	Ass	Nhat	Dhat
4602 [1,]	1.0	66.98212	37.73338	0.5633352
4603 [2,]	1.5	84.36242	46.21008	0.5477567
4604 [3,]	2.0	103.74272	57.00617	0.5494956
4605 [4,]	2.5	125.12302	69.03616	0.5517463
4606 [5,]	3.0	148.50332	82.17550	0.5533580
4607 [6,]	3.5	173.88362	96.44018	0.5546249
4608 [7,]	4.0	201.26392	111.83524	0.5556646

4610 The estimates of D stabilize rapidly and the incremental difference is within the
 4611 numerical error associated with approximating the integral. The results suggest

that wolverine density is around 0.55 individuals per 100 km^2 (recall that a state-space unit is $10 \times 10 km$). This is about 5.5 individuals per thousand km^2 which compares closely with 5.77 reported in Sec. 4.7 based on Bayesian analysis of the model.

6.4.1 Restricted state-space

In Sec. 4.9 we used a discrete representation of the state-space in order to have control over its extent and shape, for example so that we could clip out “non-habitat”. Clearly that formulation of the model is relevant to the use of integrated likelihood in the sense that such a representation of the state-space underlies the computation of the integral. Thus, for example, we could easily compute the MLE of parameters under some model with a restricted state-space merely by creating the required state-space at whatever grid resolution is desired, and then feed that state-space into the likelihood evaluation above. The **R** function `scr` which comes with the **R** package for this book accommodates an arbitrary state-space fashioned in this manner, as well as state-spaces created by polygons or GIS shapefiles which we demonstrate here. Our approach here is to create the integration grid (or state-space grid) outside of the likelihood evaluation, and then determine which points of the grid lie in the polygon defined by the shapefile using functions in the **R** libraries `sp` and `maptools`.

```
library(maptools)
library(sp)
SSp<-readShapeSpatial('Sim_Polygon.shp')
Pcoord<-SpatialPoints(G)
PinPoly<-over(Pcoord,SSp)
Pin<-as.numeric(!is.na(PinPoly[,1]))
G<-G[Pin==1,]
```

We modified the function `intlik3` to accept the integration grid as an argument and so we pass `G` directly. The **R** script is in package `scrbook`.

XYZ-lookup-XYZ reported in Royle et al. (2011c) based on a clipped state-space as described in section XYZ (XYZ chapter 4 XYZ).

TO BE COMPLETED

6.4.2 Exercises

1. Compute the 95% confidence interval for wolverine density, somehow. Comment on the practical implication of this level of precision.

2. Compute the AIC of this model and modify `intlik3` to consider alternative link functions (at least one additional) and compare the AIC of the different models and the estimates. Comment.

6.5 PROGRAM DENSITY AND THE R PACKAGE SECR

DENSITY is a software program developed by Efford (2004) for fitting spatial capture-recapture models based mostly on classical maximum likelihood estimation and related inference methods. Efford (2011) has also released an **R** package named **secr**, that contains much of the functionality of **DENSITY** but also incorporates new models and features. Here, we will focus on **secr** as it will continue to be developed, contains more functionality and is based in **R**.

To install and run models in **secr**, you must download the package and load it in **R**.

```
> install.packages(secr)
> library(secr)
```

secr allows the user to simulate data and fit a suite of models with various detection functions and covariate responses. **secr** uses the standard **R** model specification framework using tildes. E.g., the model command is **secr.fit** and is generally written as

```
> secr.fit(capturedata, model = list(D~1, g0~1, sigma~1), buffer = 20000)
```

where we have **g0~1** indicating the intercept model. Possible predictors for detection probability include both pre-defined variables (e.g., **t** and **b** corresponding to “time” and “behavior”), and user-defined covariates of several kinds. For example, to include a behavioral response, this would be written as **g0~b**. The discussion of covariates is developed more in chapter XX(8)⁴

Before we can fit the models, the data must first be packaged properly for **secr**. Two input files are required: trap layout (location and identification information for each trap) and capture data (e.g., sampling session, animal identification, trap day, and trap location). **secr** requires that you specify the trap type, the two most common for camera trapping/hair snares are proximity detectors and count detectors. The ‘proximity’ detector type allows, at most, one detection of each individual at a particular detector on any occasion. The count detector designation allows repeat encounters of each individual at a particular detector on any occasion. There are other detector types that one can select such as: ‘polygon’ detector type which allows for a trap to be a sampled polygon, e.g., scat surveys, and ‘signal’ detector which allows for traps that have a strength indicator, e.g., acoustic arrays. The detector types single and multi can be confusing as multi seems like it would be appropriate for something like a camera trap, but instead these two designations refer to traps that retain individuals, thus precluding the ability for animals to be captured in other traps during the sampling occasion. The single type indicates trap that can only catch one animal at a time, while multi indicates traps that may catch more than one animal at a time. For a full review of the detector types, one

⁴Beth: does secr fit a local trap-specific response or just a global behavioral response?

4686 should look at the help manual, which can be accessed in **R** after installing the
 4687 **secr** package by using the command:

```
4688 > RShowDoc("secr-manual", package = "secr")
```

4689 As with all of the **scr** models, **secr** fits a detection function relating the proba-
 4690 bility of detection to the distance of a detector from an individual activity center.
 4691 **secr** allows the user to specify one of a variety of detection functions including the
 4692 commonly used half-normal, hazard rate, and exponential. There are 12 different
 4693 functions, but some are only available for simulating data, and one should take
 4694 caution when using different detection functions as the interpretation of the pa-
 4695 rameters, such as sigma, may not be consistent across formulations. The different
 4696 detection functions are defined in the **secr** manual and can be found by calling the
 4697 help function for the detection function:

```
4698 > ?detectfn
```

4699 It is useful to note that **secr** requires the buffer distance to be defined in meters
 4700 and density will be returned as number of animals per hectare. Thus to make
 4701 comparisons between **secr** and other models, we will often have to convert the
 4702 density to the same units. Also, note that sigma is returned in units of meters.

5

4704 6.5.1 Analysis using the **secr** package

4705 To demonstrate the use of the **secr** package, we will show how to do the same
 4706 analysis on the wolverine study as shown in section 4.6. To use the **secr** package,
 4707 the data need to be formatted in a similar but slightly different manner than we use
 4708 in **WinBUGS**⁶. After installing the **secr** package, we first have to read in the trap
 4709 locations and other related information, such as if the trap is operational during a
 4710 sampling occasion. The **secr** package reads in the trap data through a command
 4711 called “**read.traps**”, which requires the detector type as input. The detector type
 4712 is important because it will determine the likelihood that **secr** will use to fit the
 4713 model. Here, we have selected proximity since individuals are captured at most
 4714 once in each trap during each sampling occasion.

```
4715 > traps= read.csv(wtraps.csv)
4716 > #name the first 3 columns to match the secr nomenclature
4717 > colnames(traps)[1:3]<- c("trapID","x", "y")
4718
4719 > trapfile <- read.traps(data = traps, detector = "proximity")
```

⁵One question: SECR only ever reports sigma. What exactly is sigma? It is a scale parameter of a detection function and all detection functions have a scale parameter. But in what sense is this sigma parameter related to home range diameter? Efford doesn't explain this, does he? In some sections in chapter 4 or possibly 6 we get into this issue.

⁶Elaborate on this point – and how is this different than introduced in chapter 4?

After reading in the data, we now need to create the encounter matrix or array. The `secr` package does this through the use of the `make.caphist` command, where we provide the capture histories in raw data format (each line contains the session, identification number, occasion, and trap id for only 1 individual). This is the format that was shown in the data input file “`wcaps`”, and we only need a line or two to organize the data into the order that the `make.caphist` command wants. In creating the capture history, we provide also the trapfile with the trap information, and the format (e.g., here `fmt= ‘‘trapID’’`) so that `secr` knows how to match the encounters to the trap, and finally, we provide the number of occasions:⁷

```

4729 > wolv.dat <- wcaps[,c(2, 3, 1)]
4730         #NEED TO UPDATE THIS WHEN I GET THE FILES,
4731         ### I JUST GUESSED AT THE CODE, BUT WOULD LIKE TO TRY IT.
4732 > wolv.dat <- cbind(rep(1, dim(wolv.dat)[1]), wolv.dat)
4733 > colnames(wolv.dat) <- c("Session", "ID", "Occasion", "trapID")
4734
4735 > wolvcapt=make.caphist(wolv.dat, trapfile, fmt = "trapID", noccasions = 165)

```

Calling the `secr.fit` command, will run the model. We are using the basic model (SCR0), so we do not need to make any specifications in the command line except for the providing the buffer size (in *m*). To specify different models, you can change the default `D~1`, `g0~1`, `sigma~1`, which the interested reader can do with very little difficulty.

```

4741 > wolv.secr=secr.fit(wolvcapt, model = list(D~1, g0~1, sigma~1), buffer = 20000)
4742
4743 > wolv.secr
4744
4745 secr.fit( caphist = wolvcapt, buffer = 20000, binomN = 1 )
4746 secr 2.0.0, 18:26:39 05 Jul 2011
4747
4748 Detector type      proximity
4749 Detector number    37
4750 Average spacing    4415.693 m
4751 x-range            593498 652294 m
4752 y-range            6296796 6361803 m
4753 N animals          : 21
4754 N detections        : 115
4755 N occasions         : 165
4756 Mask area          : 1037069 ha
4757
4758 Model              : D~1 g0~1 sigma~1
4759 Fixed (real)        : none
4760 Detection fn        : halfnormal
4761 Distribution         : poisson

```

⁷Beth: Do you need to update this?

```

4762 N parameters      : 3
4763 Log likelihood     : -746.754
4764 AIC                 : 1499.508
4765 AICc                : 1500.920
4766
4767 Beta parameters (coefficients)
4768           beta      SE.beta      lcl      ucl
4769 D      -9.749576 0.23027860 -10.200913 -9.298238
4770 g0     -4.275736 0.15846104 -4.586313 -3.965158
4771 sigma  8.699202 0.07868944  8.544973  8.853430
4772
4773 Variance-covariance matrix of beta parameters
4774           D           g0          sigma
4775 D      0.053028233 0.000546922 -0.005226926
4776 g0     0.000546922 0.025109900 -0.005885213
4777 sigma -0.005226926 -0.005885213 0.006192027
4778
4779 Fitted (real) parameters evaluated at base levels of covariates
4780           link      estimate SE.estimate      lcl      ucl
4781 D      log 5.831941e-05 1.360973e-05 3.713638e-05 9.158548e-05
4782 g0     logit 1.371121e-02 2.142902e-03 1.008756e-02 1.861207e-02
4783 sigma  log 5.998124e+03 4.727205e+02 5.140849e+03 6.998355e+03

```

4784 Under the fitted (real) parameters, we find D , the density, given in units of
4785 individuals/hectare (1 hectare = 10000 m^2). To convert this into individuals/1000
4786 km^2 , we multiply by 100000, thus our density estimate is 5.83 individuals/1000
4787 km^2 . σ is given in units of meters, to convert to kilometers, we divide by 1000,
4788 which puts sigma at 5.99 km . Both of these estimates are very similar to those
4789 provided in section 4.6 for the buffer size equal to 20 km XXXX How similar?
4790 XXXXX.

4791 As an exercise, run this analysis for 30 and 40 km buffers and compare those
4792 found in section 4.6 under **WinBUGS**. NOTE: The function `seccr.fit` will return
4793 a warning when the buffer size appears to be too small. This is useful particularly
4794 with the different units being used between programs and packages.

6.6 SUMMARY AND OUTLOOK

4795 In this chapter, we showed that classical analysis of SCR models based on likeli-
4796 hood methods is a relatively simple proposition. Analysis is based on the so-called
4797 integrated likelihood in which the individual activity centers (random effects) are
4798 removed from the conditional-on-s likelihood by integration. We showed how to
4799 construct the integrated likelihood and fit some simple models in the **R** program-
4800 ming language. In addition, likelihood analysis for some broad classes of SCR
4801 models can be accomplished in the software package **DENSITY** or the **R** library

`secr` which we provided an illustration of here. In later chapters we provide more detailed analyses of SCR data using the `secr` package.

To compute the integrated likelihood we have to precisely describe the state-space of the underlying point process. In practice, this leads to a “buffer” around the trap array. We note that this is not really a “buffer strip” in the sense of Wilson and Anderson (1985) which is a feature of the analysis but it is somewhat more general here. In particular, it establishes the support of the integrand which we generally require to compute any integral. It might be that the integrand itself is finite even if the support is infinity but that may or may not be the case depending on the choice of detection function. As a practical matter then, it will typically be the case that, while estimates of N increase with the size of the buffer, estimates of density stabilize. This is not a feature of the classical methods based on using model M_0 or model M_h and buffering the trap array.

Why or why not use likelihood inference exclusively? For certain specific models, it is probably more computationally efficient to produce MLEs. However, **BUGS** is extremely flexible in terms of describing models, although it sometimes can be quite slow. We can devise models in the **BUGS** language easily that we cannot fit in `secr`. E.g., random individual effects of various types (see next chapter), we can handle missing covariates in complete generality and seamlessly, and impose arbitrary distributions on random variables. Moreover, models can easily be adapted to include auxiliary data types. For example, we might have camera trapping and genetic data and we can describe the models directly in **BUGS** and fit a joint model. For the MLE we have to write a custom new piece of code for each model or hope someone has done it for us. Later we consider open population models which are straightforward to develop in **BUGS** but, so far, there is no available platform for doing MLE although we imagine one could develop this. Another thing that is more conceptual here is non-CSR point processes (see chapter XXXX) and generating predictions of how many individuals have home range centers in any particular polygon. Basic benefits of Bayesian analysis have been discussed elsewhere (XXXXXXXXX Chapter 2? BPA book? Link and Barker?) and we believe these are compelling. On the other hand, likelihood analysis makes it easy to do model-selection by AIC and in some cases compute standard errors or carry-out goodness-of-fit evaluations.

In summary, basic SCR models are easy to implement by either likelihood or Bayesian methods but we feel that the typical user will realize much more flexibility in model development using existing platforms for Bayesian analysis. While these tend to be slow (sometimes excruciatingly slow), this will probably not be an impediment in most problems, especially at some near point in the future. Since we spent a lot of time here talking about specific technical details on how to implement likelihood analysis of SCR models, we provided a corresponding treatment in the next chapter on how to devise MCMC algorithms for SCR models. This is a bit more tedious and requires more coding, but is not technically challenging (except perhaps to develop highly efficient algorithms which we don't excel at).

7

MCMC FOR SPATIAL CAPTURE-RECAPTURE

7.1 INTRODUCTION

In this chapter we will dive a little deeper into Markov chain Monte Carlo (MCMC) sampling. We will construct custom MCMC samplers in R, starting with easy-to-code GLMs and GLMMs and moving on to simple SCR models. We will also demonstrate some tricks and simple extensions to the 'spatial null model'. Finally, we will illustrate some alternative ready-to-use software packages for MCMC sampling. We will NOT provide exhaustive background information on the theory and justification of MCMC sampling there are entire books dedicated to that subject and we refer you to Robert and Casella (2004) and Robert and Casella (2010). Rather we aim to provide you with enough background and technical know-how to start building your own MCMC samplers for SCR models in R.

7.1.1 Why build your own MCMC algorithm?

The standard program we have used so far to run MCMC analyses is WinBUGS (Gilks et al., 1994). The wonderful thing about WinBUGS is that it will automatically use the most appropriate and efficient form of MCMC sampling for the model specified by the user.

The fact that we have such a Swiss Army knife type of MCMC machine begs the question: Why would anyone want to build their own MCMC algorithm? For one, there are a limited number of distributions and functions implemented in WinBUGS. While OpenBUGS provides more options, some more complex models may be impossible to build within these programs. A very simple example from spatial capture-recapture that can give you a headache in WinBUGS is when your

state-space is an irregular-shaped polygon, rather than an ideal rectangle that can be characterized by four pairs of coordinates. It is easy to restrict activity centers to any arbitrary polygon in R using an ESRI shapefile (and we will show you an example in a little bit), but you cannot use a shape file in a BUGS model.

Sometimes implementing an MCMC algorithm in R may be faster than in WinBUGS - especially if you want to run simulation studies where you have hundreds or more simulated data sets, several years' worth of data or other large models, this can be a big advantage.

Finally, building your own MCMC algorithm is a great exercise to understand how MCMC sampling works. So while using the BUGS language requires you to understand the structure of your model, building an MCMC algorithm requires you to think about the relationship between your data, priors and posteriors, and how these can be efficiently analyzed and characterized. Not to mention that, if you are an R junkie, it can actually be fun. However, if you don't think you will ever sit down and write your own MCMC sampler, consider skipping this chapter - apart from coding it will not cover anything SCR-related that is not covered by other, more model-oriented chapters as well.

7.2 MCMC AND POSTERIOR DISTRIBUTIONS

As mentioned in Chapter 2, MCMC is a class of simulation methods for drawing (correlated) random numbers from a target distribution, which in Bayesian inference is the posterior distribution. As a reminder, the posterior distribution is a probability distribution for an unknown parameter, say θ , given a set of observed data and its prior probability distribution (the probability distribution we assign to a parameter before we observe data). The great benefit of computing the posterior distribution of θ is that it can be used to make probability statements about θ , such as the probability that θ is equal to some value, or the probability that θ falls within some range of values. As an example, suppose we conducted a Bayesian analysis to estimate detection probability of some species at a study site (p), and we obtained a posterior distribution of $\text{beta}(20,10)$ for the parameter p . The following R commands demonstrate how we make inferences based upon summaries of the posterior distribution. Fig 1 shows the posterior along with the summary statistics.

```
> (post.median <- qbeta(0.5, 20, 10))
[1] 0.6704151
> (post.95ci <- qbeta(c(0.025, 0.975), 20, 10))
[1] 0.4916766 0.8206164
```

Thus, we can state that there is a 95% probability that θ lies between 0.49 and 0.82.

The posterior distribution summarizes all we know about a parameter and thus, is the central object of interest in Bayesian analysis. Unfortunately, in many if not




Figure 7.1. Probability density plot of a hypothetical posterior distribution of $\text{beta}(20,10)$; dashed lines indicate mean and upper and lower 95% interval

most practical applications, it is nearly impossible to directly compute the posterior.
Recall Bayes theorem:

$$p(\theta|y) = p(y|\theta) * p(\theta) / p(y), \quad (7.2.1)$$

where θ is the parameter of interest, y is the observed data, $p(\theta|y)$ is the posterior, $p(y|\theta)$ the likelihood of the data conditional on θ , $p(\theta)$ the prior probability of θ , and, finally, $p(y)$ is the marginal probability of the data, which can also be written as

$$p(y) = \int p(y|\theta) * p(\theta) d\theta$$

This marginal probability is a normalizing constant that ensures that the posterior integrates to 1. You read in Chapter 2 that this integral is often hard or impossible to evaluate, unless you are dealing with a really simple model. For example, consider that you have a Normal model, with a set of n observations, y that come from a Normal distribution:

$$y \sim \text{Normal}(\mu, \sigma),$$

where σ is known and our objective is to obtain an estimate of μ using Bayesian statistics. To fully specify the model in a Bayesian framework, we first have to define a prior distribution for μ . Recall from Chapter 2 that for certain data models, certain priors lead to conjugacy i.e. if you choose the right prior for your

parameter, your posterior distribution will be of a known parametric form. The conjugate prior for the mean of a normal model is also a Normal distribution:

$$\mu \sim \text{Normal}(\mu_0, \sigma_0^2)$$

If μ_0 and σ_0^2 are fixed, the posterior for μ has the following form (for the algebraic proof, see XXX):

$$\mu|y \sim \text{Normal}(\mu_n, \sigma_n^2) \quad (7.2.2)$$

where

$$\mu_n = (\text{sig}^2 / \text{sig}^2 + n * \text{sig}0^2) * \text{mu}0 + (n * \text{sig}0^2 / \text{sig}^2 + n * \text{sig}0^2) * y - \text{bar}$$

And

$$\text{sign}^2 = \text{sig}^2 * \text{sig}0^2 / (\text{sig}^2 + n * \text{sig}0^2)$$

We can directly obtain estimates of interest from this Normal posterior distribution, such as the mean μ -hat and its variance; we do not need to apply MCMC, since we can recognize the posterior as a parametric distribution, including the normalizing constant $p(y)$. But generally we will be interested in more complex models with several, say n , parameters. In this case, computing $p(y)$ from Eq. 7.2.1 requires n -dimensional integration, which is can be difficult or impossible. Thus, the posterior distribution is generally only known up to a constant of proportionality:

$$p(\theta|y) \propto p(y|\theta) * p(\theta)$$

The power of MCMC is that it allows us to approximate the posterior using simulation without evaluating the high dimensional integrals and to directly sample from the posterior, even when the posterior distribution is unknown! The price is that MCMC is computationally expensive. Although MCMC first appeared in the scientific literature in 1949 (Metropolis and Ulam, 1949), widespread use did not occur until the 1980s when computational power and speed increased (Gelfand and Smith, 1990). It is safe to say that the advent of practical MCMC methods is the primary reason why Bayesian inference has become so popular during the past three decades. In a nutshell, MCMC lets us generate sequential draws of θ (the parameter(s) of interest) from distributions approximating the unknown posterior over T iterations. The distribution of the draw at t depends on the value drawn at $t-1$; hence, the draws form a Markov chain.¹ As T goes to infinity, the Markov chain converges to the desired distribution. In our case the posterior distribution for θ — y . Thus, once the Markov chain has reached its stationary distribution, the generated samples can be used to characterize the posterior distribution, $p(\theta|y)$, and point estimates of θ , its standard error and confidence bounds, can be obtained directly from this approximation of the posterior. In practice, although we know that

¹In case you are not familiar with Markov chains, for t random samples $\theta(1), \dots, \theta(t)$ from a Markov chain the distribution of $\theta(t)$ depends only on the most recent value, $\theta(t-1)$.

a Markov chain will eventually converge, we can only generate a limited number of samples a process that depending on the model can be quite time consuming. Assessing whether our Markov chain has indeed converged is an important part of MCMC sampling and we will speak about some common diagnostics in Section XX.

7.3 TYPES OF MCMC SAMPLING

There are several MCMC algorithms, the most popular being Gibbs sampling and Metropolis-Hastings sampling. We will be dealing with these two classes in more detail and use them to construct the MCMC algorithms for SCR models. Also, we will briefly review alternative techniques that are applicable in some situations.

7.3.1 Gibbs sampling

Gibbs sampling was named after the physicist J.W. Gibbs by Geman and Geman (1984), who applied the algorithm to a Gibbs distribution². The roots of Gibbs sampling can be traced back to work of Metropolis et al. (1953), and it is actually closely related to Metropolis sampling (see Chapter 11.5 in Gelman et al. (2004), for the link between the two samplers). We will focus on the technical aspects of this algorithm, but if you find yourself hungry for more background, Casella and George (1992) provide a more in-depth introduction to the Gibbs sampler.

In Chapter 2 you already heard about the basic principles of Gibbs sampling³. But as a refresher, let's go back to our simple example from above to understand the motivation and functioning of Gibbs sampling. Recall that for a Normal model with known variance and a Normal prior for μ , the posterior distribution of $\mu|y$ is also Normal. Conversely, with a fixed (known) μ , but unknown variance, the conjugate prior for σ^2 is an Inverse-Gamma distribution with shape and scale parameters a and b :

$$\sigma^2 \sim \text{Inv-Gamma}(a, b),$$

With fixed a and b , the posterior $p(\text{sig}|\mu, y)$ is also an Inverse Gamma distribution, namely:

$$\text{sig}|\mu, y \sim \text{InvGamma}(an, bn), \quad (7.3.1)$$

where $an = n/2 + a$ and $bn = 1/2\sigma(y_i - \mu)^2 + b$. However, what if we know neither μ nor sig , which is probably the more common case? The joint posterior distribution of μ and sig now has the general structure

$$p(\mu, \text{sig}|y) = \frac{p(y|\mu) * p(\mu) * p(\text{sig})}{\int p(y|\mu) * p(\mu) * p(\text{sig}) d\mu d\text{sig}}$$

²a distribution from physics we are not going to worry about, since it has no immediate connection with Gibbs sampling other than giving its name

³maybe we should think out chapter 2 and concentrate that material here?

4983 Or

$$p(\mu, \sigma|y) \propto p(y|\mu) * p(\mu) * p(\sigma)$$

4984 This cannot easily be reduced to a distribution we recognize. However, we can
 4985 condition μ on σ (i.e., we treat σ as fixed) and remove all terms from the
 4986 joint posterior distribution that do not involve μ to construct the full conditional
 4987 distribution,

$$p(\mu|\sigma, y) \propto p(y|\mu) * p(\mu)$$

4988 The full conditional of μ again takes the form of the Normal distribution shown
 4989 in Eq. ??; similarly, $p(\sigma|\mu, y)$ takes the form of the Inverse Gamma distribution
 4990 shown in Eq. Eq. 7.3.1 both distribution we can easily sample from. And this is
 4991 precisely what we do when using Gibbs sampling we break down high-dimensional
 4992 problems into convenient one-dimensional problems by constructing the full con-
 4993 ditional distributions for each model parameter separately; and we sample from
 4994 these full conditionals, which, if we choose conjugate priors, are known parametric
 4995 distributions. Let's put the concept of Gibbs sampling into the MCMC framework
 4996 of generating successive samples, using our simple Normal model with unknown μ
 4997 and σ and conjugate priors as an example. These are the steps you need to build
 4998 a Gibbs sampler:

4999 **Step 0:** Begin with some initial values for θ , $\theta(0)$. In our example, we have to
 5000 specify initial values for μ and σ , for example by drawing a random number from
 5001 some uniform distribution, or by setting them close to what we think they might
 5002 be. (Note: This step is required in any MCMC sampling chains have to start from
 5003 somewhere. We will get back to these technical details a little later.)

5004 **Step 1:** Draw $\theta_1(1)$ from the conditional distribution $p(\theta_1(1) | \theta_2(0), \dots, \theta_d(0))$ Here,
 5005 θ_1 is μ , which we draw from the Normal distribution in Eq. ?? using $\sigma(0)$ as
 5006 value for σ .

5007 Step 2: Draw $\theta_2(1)$ from the conditional distribution $p(\theta_2(1) | \theta_1(1), \theta_3(0), \dots, \theta_d(0))$
 5008 Here, θ_2 is σ , which we draw from the Inverse Gamma distribution of Eq. 7.3.1,
 5009 using $\mu(1)$ as value for μ ...

5010 **Step d:** Draw $\theta_d(1)$ from the conditional distribution $p(\theta_d(1) | \theta_1(1), \dots, \theta_{d-1}(1))$

5011 In our example we have no additional parameters, so we only need step 0 through
 5012 to 2. Repeat Steps 1 to d for $K =$ a large number of samples. In terms of R coding,
 5013 this means we have to write Gibbs updaters for μ and σ and embed them into
 5014 a loop over K iterations. The final code in the form of an R function is shown in
 5015 Panel 1.

5016 Andy will build the panel environment here soon.

5017

5018 Panel 1: R-code for a Gibbs sampler for a Normal model with unknown μ

5019 and sig and conjugate (Normal and Inverse Gamma, respectively) priors
5020 for both parameters.

```
5021
5022 Normal.Gibbs<-function(y=y,mu0=mu0, sig0=sig0, a=a,b=b,niter=niter) {
5023
5024   ybar<-mean(y)
5025   n<-length(y)
5026   mu<-runif(1) #mean initial value
5027   sig<-runif(1) #sd initial value
5028   an<-n/2 + a
5029
5030   out<-matrix(nrow=niter, ncol=2)
5031   colnames(out)<-c('mu', 'sig')
5032
5033   for (i in 1:niter) {
5034
5035     #update mu
5036     mun<- (sig/(sig+n*sig0))*mu0 + (n*sig0/(sig+n* sig0))*ybar
5037     sign <- (sig*sig0)/ (sig+n*sig0)
5038     mu<-rnorm(1,mun, sqrt(sign))
5039
5040     #update sig
5041     bn<- 0.5 * (sum((y-mu)^2)) +b
5042     sig<-1/rgamma(1,shape=an, rate=bn)
5043     out[i,<-c(mu,sqrt(sig))
5044
5045   }
5046   return(out)
5047 }
```

5048 This is it! You can use the code `NormalGibbs.R` in the **R** package `scrbook` to
5049 simulate some data, $y \sim \text{Normal}(5, 0.5)$ and run your first Gibbs sampler. Your
5050 output will be a table with two columns, one per parameter, and K rows, one per
5051 iteration. For this 2-parameter example you can visualize the joint posterior by
5052 plotting samples of μ against samples of σ (Fig. 2 XXX):

```
5053 plot(out[,1], out[,2])
```

5054 The marginal distribution of each parameter is approximated by just examining the
5055 samples of this particular parameter you can visualize it by plotting a histogram
5056 of the samples (Fig. 3 a, b XXX):

```
5057 par(mfrow=c(1,2))
5058 hist(out[,1]); hist (out[,2])
```

Finally, recall an important characteristic of Markov chains, namely, that the chain has to have converged (reached its stationary distribution) for samples to come from the posterior distribution. In practice, that means you have to throw out some of the initial samples called the burn-in. We will talk about this in more when we talk about convergence diagnostics. For now, you can use the `plot(out[,1])` or `plot(out[,2])` command to make a time series plot of the samples of each parameter and visually assess how many of the initial samples you should discard. Figure 3 c and d shows plots for the estimates of mu and sigma from our simulated data set; you see that in this simple example the Markov chain apparently reaches its stationary distribution very quickly the chains look 'grassy' seemingly from the start. It is hard to discern a burn-in phase visually (but we will see examples further on where the burn-in is clearer) and you may just discard the first 500 draws to be sure you only use samples from the posterior distribution. The mean of the remaining samples are your estimates of mu and sig:

```
> summary(mod[501:10000,])
      mu              sig
Min.   : 4.936      Min.   : 0.4569
1st Qu.: 4.984      1st Qu.: 0.4889
Median : 4.994      Median : 0.4961
Mean   : 4.994      Mean   : 0.4964
3rd Qu.: 5.005      3rd Qu.: 0.5037
Max.   : 5.062      Max.   : 0.5356
```

7.3.2 Metropolis-Hastings sampling

Although it is applicable to a wide range of problems, the limitations of Gibbs sampling are immediately obvious what if we do not want to use conjugate priors (or what if we cannot recognize the full conditional distribution as a parametric distribution, or simply do not want to worry about these issues)? The most general solution is to use the Metropolis-Hastings (MH) algorithm, which also goes back to the work by Metropolis et al. (1953). You saw the basics of this algorithm in Chapter 2. In a nutshell, because we do not recognize the posterior $p(\theta|y)$ as a parametric distribution, the MH algorithm generates samples from a known proposal distribution, say $h(\theta)$, that depends on θ at $t-1$. The t^{th} sample is accepted or rejected based on its joint posterior probability density compared to the density of the sample at $t-1$. The original Metropolis algorithm requires $h(\theta)$ to be symmetric so that $h(\theta^t|\theta^{t-1}) = h(\theta^{t-1}|\theta^t)$; but a later development of the algorithm by Hastings (1970) lifted this condition. Using a symmetric proposal distribution makes life a little easier and we are going to limit our coverage of the Metropolis-Hastings sampler to this specific case. Specifically, we are going to use a Normal proposal distribution, which is also referred to as 'random walk Metropolis-Hastings




Figure 7.2. Joint posterior distribution of μ and σ from a Normal Model

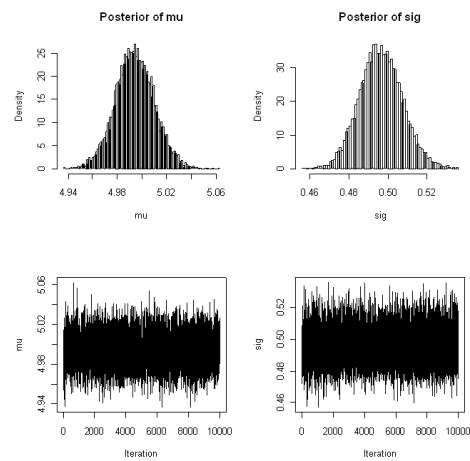


Figure 7.3. Plots of the posterior distributions of μ (a) and σ (b) from a Normal model and time series plots of μ (c) and σ (d).

sampling'. It is worth knowing that there are alternative formulations of the algorithm. For example, in the independent M-H, θ^t does not depend on θ^{t-1} , while the Langevin algorithm (Roberts and Rosenthal, 1998) aims at avoiding the random walk by favoring moves towards regions of higher posterior probability density. The interested reader should look up these algorithms in Robert and Casella (2004) or Robert and Casella (2010).

Building a MH sampler can be broken down into several steps. We are going to demonstrate these steps using a different but still simple and common model the logit-normal or logistic regression model. For simplicity, assume that

$$y \sim \text{Bern}(\exp(\theta)/(1 + \exp(\theta)))$$

and

$$\theta \sim \text{Normal}(\mu_0, \sigma)$$

The following steps are required to set up a random walk MH algorithm:

Step 0: Choose initial values, $\theta(0)$.

Step 1: Generate a proposed value of θ at t from $h(\theta^t - \theta^{t-1})$. We often use a Normal proposal distribution, so we draw θ_1 from $\text{Normal}(\theta^t, \text{sig}^2)$, where sig^2 is the variance of the Normal proposal distribution, a tuning parameter that we have to set.

Step 2: Calculate the ratio of posterior densities for the proposed and the original value for θ :

$$r = p(\theta^t|y)/p(\theta^{t-1}|y)$$

In our example,

$$r = \text{Bern}(y|\theta^t) * \text{Normal}(\theta^t|\mu_0, \sigma_0) / \text{Bernoulli}(y|\theta^{t-1}) * \text{Normal}(\theta^{t-1}|\mu_0, \text{sig}_0)$$

Step 3: Set

```
\begin{eqnarray*}
\theta(t) &= & \theta(t) \text{ with probability } \min(r,1) //
&= & \theta(t-1) \text{ otherwise }
\end{eqnarray*}
```

We can do that by drawing a random number u from a $\text{Unif}(0, 1)$ and accept θ^t if $u < r$. Repeat for $t = 1, 2, \dots$ a large number of samples. The **R** code for this MH sampler is provided in Panel 2 XXXX.

Panel 2: R code to run a Metropolis sampler on a simple Logit-Normal model.

```
Logreg.MH<-function(y=y, mu0=mu0, sig0=sig0, niter=niter) {
```

```

5129 out<-c()
5130
5131 theta<-runif(1, -3,3) #initial value
5132
5133 for (iter in 1:niter){
5134   theta.cand<-rnorm(1, theta, 0.2)
5135
5136   loglike<-sum(dbinom(y, 1, exp(theta)/(1+exp(theta)), log=TRUE))
5137   logprior <- dnorm(theta,mu0 ,sig0, log=TRUE)
5138   loglike.cand<-sum(dbinom(y, 1, exp(theta.cand)/(1+exp(theta.cand)), log=TRUE))
5139   logprior.cand <- dnorm(theta.cand, mu0, sig0, log=TRUE)
5140
5141   if (runif(1)<exp((loglike.cand+logprior.cand)-(loglike+logprior))){
5142     theta<-theta.cand
5143   }
5144   out[iter]<-theta
5145 }
5146
5147 return(out)
5148 }

```

5149 The reason we sum the logs of the likelihood and the prior, rather than multiply-
5150 ing the original values, is simply computational. The product of small probabilities
5151 can be numbers very close to 0, which computers do not handle well. Thus we add
5152 the logarithms, sum, and exponentiate to achieve the desired result. Similarly, in
5153 case you have forgotten some elementary math, $x/y = \exp(\log(x) - \log(y))$, with
5154 the latter being favored for computational reasons.

5155 Comparing MH sampling to Gibbs sampling, where all draws from the condi-
5156 tional distribution are used, in the MH algorithm we discard a portion of the
5157 candidate values, which inherently makes it less efficient than Gibbs sampling the
5158 price you pay for its increased generality. In Step 1 of the MH sampler we had to
5159 choose a variance for the Normal proposal distribution. Choice of the parameters
5160 that define our candidate distribution is also referred to as 'tuning', and it is im-
5161 portant since adequate tuning will make your algorithm more efficient, i.e. your
5162 Markov chain will converge faster. The variance should be chosen so that (a) each
5163 step of drawing a new proposal value for θ can cover a reasonable distance in the
5164 parameter space, as otherwise, the random walk moves too slowly; and (b) proposal
5165 values are not rejected too often, as otherwise the random walk will 'get stuck' at
5166 specific values for too long. As a rule of thumb, your candidate value should be ac-
5167 cepted in about 40% of all cases. Acceptance rates of 20–80% are probably ok, but
5168 anything below or above may well render your algorithm inefficient (this does not
5169 mean that it will give you wrong results only that you will need more iterations to
5170 converge to the posterior distribution). In practice, tuning will require some 'trial-
5171 and-error' and some common sense. Or, one can use an adaptive phase, where the
5172 tuning parameter is automatically adjusted until it reaches a user-defined accep-

5173 tance rate, at which point the adaptive phase ends and the actual Markov chain
 5174 begins. This is computationally a little more advanced. Link and Barker (2009)
 5175 discuss this in more detail. It is important the samples drawn during the adaptive
 5176 phase are discarded. You can easily check acceptance rates for the parameters you
 5177 monitor (that are part of your output) using the `rejectionRate()` function of the
 5178 package `coda` (we will talk more about this package a little later on). Do not let
 5179 the term 'rejection rate' confuse you; it is simply $1 - \text{acceptance rate}$. There may
 5180 be parameters for example, individual values of a random effect or latent variables
 5181 that you do not want to save, though, and in our next example we will show you a
 5182 way to monitor their acceptance rates with a few extra lines of code.

5183 7.3.3 Metropolis-within-Gibbs

5184 One weakness of the MH sampler is that formulating the joint posterior when
 5185 evaluating whether to accept or reject the candidate values for θ becomes increas-
 5186 ingly complex or inefficient as the number of parameters in a model increases. It
 5187 is probably going to sound like MCMC sampling is too good to be true but in
 5188 these cases you can simply combine MH sampling and Gibbs sampling. You can
 5189 use Gibbs sampling to break down your high-dimensional parameter space into
 5190 easy-to-handle one-dimensional conditional distributions and use MH sampling for
 5191 these conditional distributions. Better yet if you have some conjugacy in your
 5192 model, you can use the more efficient Gibbs sampling for these parameters and
 5193 one-dimensional MH for all the others. You have already seen the basics of how to
 5194 build both types of algorithms, so we can jump straight into an example here and
 5195 build a Metropolis-within-Gibbs algorithm.

7.4 GLMMS POISSON REGRESSION WITH A RANDOM EFFECT

5196 Let's assume a model that gets us closer to the problem we ultimately want to
 5197 deal with a GLMM. Here, we assume we have Poisson counts, y , from i plots
 5198 in j different study sites, and we believe that the counts are influenced by some
 5199 plot-specific covariate, x , but that there is also a random site effect. So our model
 5200 is:

$$y_{ij} \sim \text{Poisson}(l_{amij})$$

5201

$$l_{amij} = \exp(a_j + b * x_i)$$

5202 Let's use Normal priors on a and b ,

$$a_j \sim \text{Normal}(m_{ua}, s_{iga})$$

5203 and

$$b \sim \text{Normal}(m_{ub}, s_{igb})$$

.⁴ Since we want to estimate the random effect in this model, we do not specify μ_a and σ_a , but instead, estimate them as well, so we have to specify hyperpriors for these parameters:

$$\begin{aligned}\mu_a &\sim \text{Normal}(\mu_0, \text{sig}_0) \\ \sigma_a &\sim \text{InvGamma}(a_0, b_0)\end{aligned}$$

With the model fully specified, we can compile the full conditionals, breaking the multi-dimensional parameter space into one-dimensional components:

```
\begin{eqnarray*}
p(a_1|a_2,a_3,a_j,b,y) &\propto p(y_{i1}|a_1,b) * p(a_1|mua, \text{sig}_a) \\
&\propto \text{Poisson}(y_{i1} | \exp(a_1 + b*x[j=1])) * \text{Normal}(a_1|mua, \text{sig}_a)
\end{eqnarray*}
\begin{eqnarray*}
p(a_2|a_1,a_3,a_j,b,y) &\propto p(y_{i2}|a_2,b) * p(a_2|mua, \text{sig}_a) \\
&\propto \text{Poisson}(y_{i2} | \exp(a_2 + b*x[j=1])) * \text{Normal}(a_2|mua, \text{sig}_a)
\end{eqnarray*}
and so on for all elements of a.
\begin{eqnarray*}
p(b|a,y) &\propto p(y|a,b) * p(b) \\
&\propto \text{Poisson}(y | \exp(a + b*x)) * \text{Normal}(b|mub, \text{sig}_b)
\end{eqnarray*}
```

Finally, we need to update the hyperparameters for a:

$$\begin{aligned}p(mua|a) &\propto p(a|mua, \text{sig}_a) * p(mua) \\ p(\text{sig}_a|a) &\propto p(a|mua, \text{sig}_a) * p(\text{sig}_a)\end{aligned}$$

Since we assumed a to come from a Normal distribution, the choice of priors for mua Normal and sig_a Inverse Gamma leads to the same conjugacy we observed in our initial Normal model, so that both hyperparameters can be updated using Gibbs sampling.

Now let's build the updating steps for these full conditionals. Again, for the MH steps that update a and b we use Normal proposal distributions with standard deviations sig_ha and sig_hb.

First, we set the initial values a(0) and b(0). Then, starting with a1, we draw a1(1) from Normal(a1(0), sig_ha), calculate the conditional posterior density of a1(0) and a1(1) and compare their ratios,

$$r = \text{Poisson}(y(j=1) | \exp(a1(1) + b*x)) * \text{Normal}(a1(1) | mua, \text{sig}_a) / \text{Poisson}(y(j=1) | \exp(a1(0) + b*x)) * \text{Normal}(a1(0) | mua, \text{sig}_a)$$

and accept a1(1) with probability min(r,1). We repeat this for all a's.

⁴Why is b a hyperparameter?

5235 For b, we draw $b(1)$ from $\text{Normal}(b(0), \text{sigbh})$, compare the posterior densities
5236 of $b(0)$ and $b(1)$,

$$r = \text{Poisson}(y|\exp(a+b(1)*x)) * \text{Normal}(b(1)|\text{mub}, \text{sigb}) / \text{Poisson}(y|\exp(a+b(0)*x)) * \text{Normal}(b(0)|\text{mub}, \text{sigb}),$$

5237 and accept $b(1)$ with probability $\min(r, 1)$.

5238 For mua and sig , we sample directly from the full conditional distributions (Eq
5239 XX and Eq XX):

$$\text{mua}(1) \sim \text{Normal}(\text{mun}, \text{sign})$$

5240 where $\text{mun} = (\text{sig}(0)/\text{sig}(0) + n_a * \text{sig}) * \text{mu0} + (n_a * \text{sig0}/\text{sig}(0) + n_a * \text{sig0}) * \text{abar}(1)$ and $\text{sign} = \text{sig}(0) * \text{sig0}/(\text{sig}(0) + n * \text{sig0})$. Here, abar is the
5241 current mean of the vector a , which we updated before, and n_a is the length of
5242 a . For sig we use $\text{sig}(1) \sim \text{InvGamma}(an, bn)$, where $an = n_a/2 + a0$, and
5243 $bn = 1/2 \sum (\text{a}(1) - \text{mua}(1))^2 + b0$.
5244

5245 We repeat these steps over K iterations of the MCMC algorithm. In this example
5246 we may not want to save each value for a , but are only interested in their mean and
5247 standard deviation. Since these two parameters will change as soon as the value for
5248 one element in a changes, their acceptance rates will always be close to 1 and are
5249 not representative of how well your algorithm performs. To monitor the acceptance
5250 rates of parameters you do not want to save, you simply need to add a few lines
5251 of code into your updater to see how often the individual parameters are accepted.
5252 The full code for the MCMC algorithm of our Poisson GLMM in Panel 3 shows
5253 one way how to monitor acceptance of individual a 's.

5254 Panel 3: R code for the Metropolis-within-Gibbs sampler for
5255 a Poisson regression with random intercepts.

```
5256
5257 Pois.reg<-function(y=y,site=site,mu0=mu0,sig0=sig0,a0=a0,b0=b0,
5258                   mub=mub, sigb=sigb, niter=niter){
5259
5260   lev<-length(unique(site))      #number of sites
5261   a<-runif(lev,-5,5) #initial values a
5262   b<-runif(1,0,5) #initial value b
5263   mua<-mean(a)
5264   sig<-sd(a)
5265
5266   out<-matrix(nrow=niter, ncol=3)
5267   colnames(out)<-c('mua','sig','b')
5268
5269   for (iter in 1:niter) {
5270
5271     #update a
5272     aUps<-0 #initiate counter for acceptance rate of a
5273     for (j in 1:lev) { #loop over sites
5274       a.cand<-rnorm(1, a[j], 0.1) #update intercepts a one at a time
```

```

5275 loglike<- sum(dpois (y[site==j], exp(a[j] + b*x[site==j]), log=TRUE))
5276 logprior<- dnorm(a[j], mua,siga, log=TRUE)
5277 loglike.cand<- sum(dpois (y[site==j], exp(a.cand + b *x[site==j]), log=TRUE))
5278 logprior.cand<- dnorm(a.cand, mua,siga, log=TRUE)
5279 if (runif(1)< exp((loglike.cand+logprior.cand) (loglike+logprior))) {
5280   a[j]<-a.cand
5281   aUps<-aUps+1
5282 }
5283 }
5284
5285 if(iter %% 100 == 0) { #this lets you check the acceptance rate of a at every 100th iteration
5286   cat("   Acceptance rates\n")
5287   cat("     a =", aUps/lev, "\n")
5288 }
5289
5290 #update b
5291 b.cand<-rnorm(1, b, 0.1)
5292 avec<-rep(a, times=c(rep(10,10)))
5293 loglike<- sum(dpois (y, exp(avec + b*x), log=TRUE))
5294 logprior<- dnorm(b, mub,sigb, log=TRUE)
5295 loglike.cand<- sum(dpois (y, exp(avec + b.cand *x), log=TRUE))
5296 logprior.cand<- dunif(b.cand, mub,sigb, log=TRUE)
5297 if (runif(1)< exp((loglike.cand+logprior.cand) (loglike+logprior) )) {
5298   b<-b.cand
5299 }
5300
5301 #update mua using Gibbs sampling
5302 abar<-mean(a)
5303 mun<- (siga/(siga+lev*sig0))*mu0 + (lev*sig0/(siga+lev* sig0))*abar
5304 sign <- (siga*sig0)/ (siga+lev*sig0)
5305 mua<-rnorm(1,mun, sqrt(sign))
5306
5307 #update siga using Gibbs sampling
5308 a0n<-lev/2 + a0
5309 b0n<- 0.5 * (sum((a-mua)^2)) +b0
5310 siga<-1/rgamma(1,shape=a0n, rate=b0n)
5311
5312 out[iter,]<-c(mua, sqrt(siga), b)
5313 }
5314 }
5315
5316 return(out)
5317 }

```

7.4.1 Rejection sampling and slice sampling

While MH and Gibbs sampling are probably the most widely applied algorithms for posterior approximation, there are other options that work under certain circumstances and may be more efficient when applicable. WinBUGS applies these algorithms and we want you to be aware that there is more out there to approximate posterior distributions than Gibbs and MH. One alternative algorithm is rejection sampling. Rejection sampling is not an MCMC method, since each draw is independent of the others. The method can be used when the posterior $p(\theta|y)$ is not a known parametric distribution but can be expressed in closed form. Then, we can use a so-called envelope function, say, $g(\theta)$, that we can easily sample from, with the restriction that $p(\theta|y) < M * g(\theta)$. We then sample a candidate value for θ from $g(\theta)$, calculate $r = p(\theta|y) / M * g(\theta)$ and keep the sample with the probability r . M is a constant that has to be picked so that $r \in [0,1]$, for example by evaluating both $p(\theta|y)$ and $g(\theta)$ at n points and looking at their ratios. Rejection sampling only works well if $g(\theta)$ is similar to $p(\theta|y)$, and packages like WinBUGS use adaptive rejection sampling (Gilks and Wild, 1992), where a complex algorithm is used to fit an adequate and efficient $g(\theta)$ based on the first few draws. Though efficient in some situations, rejection sampling does not work well with high-dimensional problems, since it becomes increasingly hard to define a reasonable envelope function. For an example of rejection sampling in the context of SCR models, see Chapter 9. Another alternative is slice sampling (Neal, 2003). In slice sampling, we sample uniformly from the area under the plot of $p(\theta|y)$. Considering a single univariate θ . Let's define an auxiliary variable, $U \sim \text{Uniform}(0, p(\theta|y))$. Then, θ can be sampled from the vertical slice of $p(\theta|y)$ at U (Figure 4):

$$\theta|U \sim \text{Unif}(B),$$

where $B = \{\theta : U < p(\theta|y)\}$

⁵

Slice sampling can be applied in many situations; however, implementing an efficient slice sampling procedure can be complicated. We refer the interested reader to chapter 7 of Robert and Casella (2010) for a simple example. Both rejection sampling and slice sampling can be applied on one-dimensional conditional distributions within a Gibbs sampling setup.

7.5 MCMC FOR CLOSED CAPTURE-RECAPTURE MODEL MH

⁶

⁵there are supposed to be equations in the caption of figure 4 but it kept causing errors

⁶Andy could move material from chapter 3 to here.




Figure 7.4. Slice sampling. For...

7.6 MCMC ALGORITHM FOR THE BASIC SPATIAL CAPTURE-RECAPTURE MODEL

By now you have seen how to build MCMC algorithms for some basic generalized linear models. Now, we'll walk you through the steps of building your own MCMC sampler for the basic SCR model (i.e. without any individual, site or time specific covariates) with both a Poisson and a binomial encounter process. As usual, we will have to go through two general steps before we write the MCMC algorithm:

- (1) Identify your model with all its components (including priors)
 - (2) Recognize and express the full conditional distributions for all parameters
- It is worthwhile to go through all of step 1 for an SCR model, but you have probably seen enough of step 2 in our previous examples to get the essence of how to express a full conditional distribution. Therefore, we will exemplify step 2 for some parameters and tie these examples directly to the respective R code.

Step 1 Identify your model

Recall the components of the basic SCR model with a Poisson encounter process from Chapter 3: We assume that individuals i , or rather, their activity centers s_i , are uniformly distributed across our state space S ,

$$s_i \sim U(S)$$

and that the number of times individual i encounters trap j , y_{ij} , is a random Poisson variable with mean λ_{mij} ,

$$y_{ij} \sim \text{Poisson}(\lambda_{mij})$$

The tie between individual location, movement and trap encounter rates is made by the assumption that λ_{mij} , is a decreasing function of the distance between s_i

5370 and j , D_{ij} , of the half-normal form

$$Lam_{ij} = lam0 * \exp(-D_{ij}^2/2 * sig^2),$$

5371 where $lam0$ is the baseline trap encounter rate at $D_{ij} = 0$ and sig controls the
5372 shape of the half-normal function.

5373 In order to estimate the number of s_i in S , N , we use data augmentation (sect.
5374 3.XYZ) and create M - n all-0 encounter histories, where n is the number of individ-
5375 uals we observed and M is an arbitrary number that is larger than N . We estimate
5376 N by summing over the auxiliary data augmentation variables, z_i , which is 1 if the
5377 individual is part of the population and 0 if not, and assume that z_i is a random
5378 Bernoulli variable,

$$z_i \sim \text{Bern}(\psi)$$

5379 To link the two model components, we modify our trap encounter model to

$$Lam_{ij} = lam0 * \exp(-D_{ij}^2/2 * sig^2) * z_i.$$

5380 The model has the following structural parameters, for which we need to specify
5381 priors ψ the Uniform (0,1) is required as part of the data augmentation procedure
5382 and in general is a natural choice of an uninformative prior for a probability; note
5383 that this is equivalent to a Beta(1,1) prior, which will come in handy later. s_i since
5384 s_i is a pair of coordinates it is two-dimensional and we use a uniform prior limited
5385 by the extent of our state-space over both dimensions. σ we can conceive several
5386 priors for sigma but let's assume an improper prior one that is Uniform over (-Inf,
5387 Inf). We will see why this is convenient when we construct the full conditionals for
5388 sigma. λ_0 analogous, we will use a Uniform (-Inf, Inf) improper prior for sigma.
5389 The parameter that is the objective of our modeling, N , is a derived parameter that
5390 we can simply obtain by summing all z 's:

$$N = \text{sum}(z)$$

5391 **Step 2 - Construct the full conditionals** Having completed step 1, let's
5392 look at the full conditional distributions for some of these parameters. We find
5393 that with improper priors, full conditionals are proportional only to the likelihood
5394 of the observations; for example, take the movement parameter sigma:

$$Sig|s, lam0, z, ypropto[y|s, lam0, z, sig] * [sig]$$

5395 Since the improper prior implies that $[sig] \propto 1$, we can reduce this further to

$$Sig|s, lam0, z, ypropto[y|s, lam0, z, sig]$$

5396 The R code to update sigma is shown in Panel 4. ⁷

⁷ Somewhere in chapter 2 i added a comment about rejecting parameters outside of the parameter space as being an ok thing to do. Richard said he read something in Robert and Casellas book on that. Hopefully he can remember where and we can cite it back in Ch 2 and again here. It could be mentioned in a sentence or two up in the MCMC section.

5397 Panel 4: R code to update sigma within an MCMC algorithm for
 5398 an SCR model when using an improper prior

```
5399
5400
5401 sig.cand <- rnorm(1, sigma, 0.1) #draw candidate value
5402 if(sig.cand>0){ #automatically reject sig.cand that are <0
5403   lam.cand <- lam0*exp(-(D*D)/(2*sig.cand*sig.cand))
5404   ll<- sum(dpois(y, lam*z, log=TRUE))
5405   llcand <- sum(dpois(y, lam.cand*z, log=TRUE))
5406   if(runif(1) < exp( llcand - ll) ){
5407     ll<-llcand
5408     lam<-lam.cand
5409     sigma<-sig.cand
5410   }
5411 }
5412
```

5413 These steps are analogous for lam0 and si and we will use MH steps for all
 5414 of these parameters. Similar to the random intercepts in our Poisson GLMM, we
 5415 update each si individually. Note that to be fully correct, the full conditional for
 5416 si contains both the likelihood and prior component, since we did not specify an
 5417 improper, but a Uniform prior on si. However, with a Uniform distribution the
 5418 probability density of any value is 1/(upper limit - lower limit) = constant. Thus,
 5419 the prior components are identical for both the current and the candidate value
 5420 and can be ignored (formally, when you calculate the ratio of posterior densities, r,
 5421 the identical prior component appears both in the numerator and denominator, so
 5422 that they cancel each other out).

5423 We still have to update zi. The full conditional for zi is

$$z_i | y, \sigma, \lambda_0, \text{sprpto}[y | z, \sigma, \lambda_0, s] * [z_i]$$

5424 and since $z_i \sim \text{Bernoulli}(\psi_i)$, the term has to be taken into account when updating
 5425 zi. The R code for updating zi is shown in Panel 5.

5426 Panel 5: R code to update z

```
5427
5428 zUps <- 0 #set counter to monitor acceptance rate
5429 for(i in 1:M) {
5430   if(seen[i]) #no need to update seen individuals, since their z =1
5431     next
5432   zcand <- ifelse(z[i]==0, 1, 0)
5433   llz <- sum(dpois(y[i,], lam[i,]*z[i], log=TRUE))
5434   llcand <- sum(dpois(y[i,], lam[i,]*zcand, log=TRUE))
5435
5436   prior <- dbinom(z[i], 1, psi, log=TRUE)
5437   prior.cand <- dbinom(zcand, 1, psi, log=TRUE)

```

```

5438         if(runif(1) < exp( (llcand+prior.cand) - (llz+prior) )) {
5439             z[i] <- zcand
5440             zUps <- zUps+1
5441         }
5442     }

```

5443 ψ itself is a hyperparameter of the model, with an uninformative prior distribu-
5444 tion of Unif(0,1) or Beta(1,1), so that

$$Psi|z \propto [z|psi] * Beta(1, 1)$$

5445 The Beta distribution is the conjugate prior to the Binomial and Bernoulli distribu-
5446 tions (remember that $z \sim Bernoulli(psi)$). The general form of a full conditional
5447 of a Beta-Binomial model with $y_i \sim Bernoulli(p)$ and $p \sim Beta(a, b)$ is

$$p(p|y) \propto Beta(a + sum(yi), b + n - sum(yi))$$

5448 In our case, this means we update psi as follows:

```

5449 si<-rbeta(1, 1+sum(z), 1 + M-sum(z))

```

5450 These are all the building blocks you need to write the MCMC algorithm for
5451 the spatial null model with a Poisson encounter process. You can find the full R
5452 code (SCR0pois.R) in the online supplementary material.

5453 7.6.1 SCR model with binomial encounter process

5454 The equivalent SCR model with a binomial encounter process is very similar. Here,
5455 each individual i can only be detected once at any given trap j during a sampling
5456 occasion k . Thus

$$y_{ij} \sim Binomial(p_{ij}, K)$$

5457 Where p_{ij} is some function of distance between s_i and trap location x_j . Here we
5458 use:

$$p_{ij} = 1 - \exp(-lam_{ij})$$

5459 Recall from Chapter 2 that this is the complementary log-log (cloglog) link func-
5460 tion, which constrains p_{ij} to fall between 0 and 1. For our MCMC algorithm that
5461 means that, instead of using a Poisson likelihood, $Poisson(y|sigma, lam0, s, z)$, we
5462 use a Binomial likelihood, $Binomial(y, K|sigma, lam0, s, z)$, in all the conditional
5463 distributions. As an example, Panel 6 shows the updating step for $lam0$ under
5464 a binomial encounter model. The full MCMC code for the binomial SCR can be
5465 found in the online supplements.

5466 Panel 6: MCMC updater for $lam0$ in a SCR model with Binomial encounter
5467 process and cloglog link function on detection. Here, `pmat =`

```

5468 1-exp(-lam).
5469
5470     lam0.cand <- rnorm(1, lam0, 0.1)
5471     if(lam0.cand > 0){ #automatically reject lam0.cand that are <0
5472         lam.cand <- lam0.cand*exp(-(D*D)/(2*sigma*sigma))
5473         p.cand <- 1-exp(-lam.cand)
5474         ll<- sum(dbinom(y, K, pmat *z, log=TRUE))
5475         llcand <- sum(dbinom(y, K, p.cand *z, log=TRUE))
5476         if(runif(1) < exp( llcand - ll) ){
5477             ll<-llcand
5478             pmat<-p.cand
5479             lam0<- lam0.cand
5480         }
5481     }

```

Another possibility is to model variation in the individual and site specific detection probability, p_{ij} , directly, without any transformation, such that

```

5484 pij<-p0 * exp(-Dij2/(2*sig^2))

```

and $p_0 = \{0,1\}$. This formulation is analogous to how detection probability is modeled in distance sampling under a half-normal detection function; however, in distance sampling p_0 - detection of an individual on the transect line - is assumed to be 1 (Buckland, 2001). Under this formulation the updater for lam_0 (equivalent to p_0 in Eq XX) becomes:

```

5490     lam0.cand <- rnorm(1, lam0, 0.1)
5491     if(lam0.cand > 0 & lam0.cand < 1 ){ #automatically reject lam0.cand that are not
5492         lam.cand <- lam0.cand*exp(-(D*D)/(2*sigma*sigma))
5493         ll<- sum(dbinom(y, K, lam *z, log=TRUE)) #no transformation needed
5494         llcand <- sum(dbinom(y, K, lam.cand *z, log=TRUE))
5495         if(runif(1) < exp( llcand - ll) ){
5496             ll<-llcand
5497             lam<-lam.cand
5498             lam0<- lam0.cand
5499         }
5500     }

```

5501 7.6.2 Looking at model output

5502 Now that you have an MCMC algorithm to analyze spatial capture-recapture data
 5503 with, let's run an actual analysis so we can look at the output. As an example,
 5504 we will use the bear data ...⁸ You can use the same script provided back in

⁸Does this data set come up before Ch6? If not, introduce data here. Or, Andy, would you rather use simulated data?

Chapter XX to read in the data and build the augmented encounter history array; then source the MCMC code for the binomial encounter model algorithm with the cloglog link and run 5000 iterations. This should take approximately 25 minutes.

```
> source('SCR0binom.txt')
> mod0<-SCR.0(y=bigTrap, X=trapmat, M=M, xl=xl, xu=xu, yl=yl, yu=yu, K=8, niter=5000)
```

Before, we used simple R commands to look at model results. However, there is a specific R package to summarize MCMC simulation output and perform some convergence diagnostics package coda (Plummer et al., 2006). Download and install coda, then convert your model output to an mcmc object

```
> chain<-mcmc(mod0)
```

Markov chain time series plots

Start by looking at time series plots of your Markov chains using `plot(chain)`. This command produces a time series plot and marginal posterior density plots for each monitored parameter, similar to what we did before using the `hist()` and `plot()` commands (Fig. 5). Time series plots will tell you several things: First, the way the chains move through the parameter space gives you an idea of whether your MH steps are well tuned. If chains were constant over many iterations you would probably need to decrease the tuning parameter of the (Normal) proposal distribution. If a chain moves along some gradient to a stationary state very slowly, you may want to increase the tuning parameter so that the parameter space is explored more efficiently.

Second, you will be able to see if your chains converged and how many initial simulations you have to discard as burn-in. In the case of the chains shown in Figure 5, we would probably consider the first 750 - 1000 iterations as burn-in, as afterwards the chains seem to be fairly stationary.

A word of caution about chain convergence

Since we do not know what the stationary posterior distribution of our Markov chain should look like (this is the whole point of doing an MCMC approximation), we effectively have no means to assess whether it has converged to this desired distribution or not. As mentioned before, the only certainty is that a Markov chain will *eventually* converge to its stationary distribution, but no-one can tell us how long this will take. Also, you only now the part of your posterior distribution that the Markov chain has explored so far for all you know the chain could be stuck in a local maximum, while other maxima remain completely undiscovered. Acknowledging that there is truly nothing we can do to ever proof convergence of our MCMC chains, there are several things we can do to increase the degree of confidence we have about the convergence of our chains. One option, and that advocated by what we will loosely call the WinBUGS community, is to run several Markov chains and to start them off at different initial values that are overdispersed relative to the posterior distribution. Such initial values help to explore different

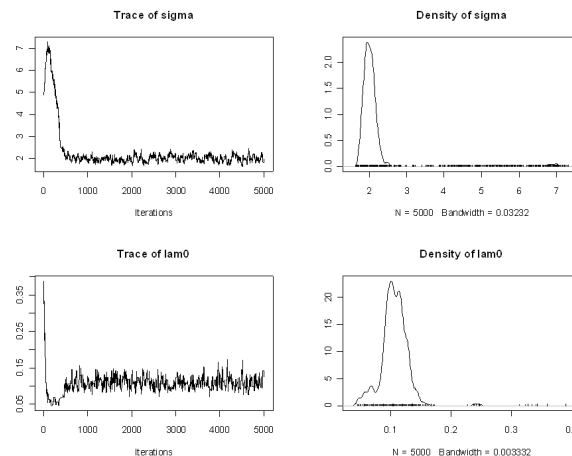


Figure 7.5. Time series and posterior density plots for sigma and lam0.

5545 areas of the parameter space simultaneously; if after a while all chains oscillate
 5546 around the same average value, chances are good that they indeed converged to
 5547 the posterior distribution. Gelman and Rubin came up with a diagnostic statistic
 5548 that essentially compares within-chain and between-chain variance to check for
 5549 convergence of multiple chains (Gelman et al., 2004). Of course, running several
 5550 parallel chains is computationally expensive. Extra computational demands are not
 5551 the only and by no means the major concern some people voice when it comes to
 5552 running several parallel MCMC chains to assess convergence. Again, consider the
 5553 fact that we do not know anything about the true form of the posterior distribution
 5554 we are trying to approximate. How do we, then, know how to pick overdispersed
 5555 initial values? We don't all we can do is pick overdispersed values relative to our
 5556 expectations of what the posterior should look like. To use a quote from the home
 5557 page of Charlie Geyer, a Bayesian statistician from the University of Minnesota,
 5558 "If you don't know any good starting points [...], then restarting the sampler at
 5559 many bad starting points is [...] part of the problem, not part of the solution."
 5560 (<http://users.stat.umn.edu/~charlie/mcmc/diag.html>). His suggestion is that your
 5561 only chance to discover a potential problem with your MCMC sampler is to run it
 5562 for a very long time. But again, there is no way of knowing how long is long enough.
 5563 It is up to you to decide, which school of thoughts appeals more to you one long
 5564 versus several parallel Markov chains. Irrespectively, part of developing an MCMC
 5565 sampler should be to make sure (within reasonable limits) that you are not missing
 5566 regions of high posterior density because of the way you specify your starting values.
 5567 Once you have explored the behavior of your chain under a reasonable range of

starting values, you may feel comfortable enough to run only one long chain. The fact that convergence cannot be proven does not mean that you should not look for potential problems in your MCMC sampler. Some problems are easily detected using simple plots, such as the time series plots we discussed above. If the overall trajectory of your chain at the end of your simulations is still upward or downward, your chain clearly has not converged and you need to run your model much longer. If you run several parallel chains and their stationary distributions look different, you may be looking at a multi-modal posterior or a problem with your sampler. With these words of caution, let's get back to looking at our model output.

7.6.3 Posterior density plots

The `plot()` command also produces posterior density plots and it is worthwhile to look at those carefully. For parameters with priors that have bounds (e.g. Uniform over some interval), you will be able to see if your choice of the prior is truncating the posterior distribution. In the context of SCR models, this will mostly involve our choice of M , the size of the augmented data set. If the posterior of N has a lot of mass concentrated close to M (or equivalently the posterior of ψ has a lot of mass concentrated close to 1), as in the example in Figure 6, we have to re-run the analysis with a larger M . A flat posterior plot shows you that the parameter essentially cannot be identified there may not be enough information in your data to estimate model parameters and you may have to consider a simpler model. Finally, posterior density plots will show you if the posterior distribution is symmetrical or skewed if the distribution has a heavy tail, using the mean as a point estimate of your parameter of interest may be biased and you may want to opt for the median or mode instead.

7.6.4 Serial autocorrelation and effective sample size

Even when we can be relatively confident that our chains have converged, the subsequent samples generated from a Markov chain are not iid samples from the posterior distribution, due to the correlation amongst samples introduced by the Markov process. As a consequence, the variance of the mean cannot simply be derived with the standard variance estimator, which takes into account the sample size (here, number of iterations). Rather, the sample size has to be adjusted to account for the autocorrelation in subsequent samples (see Chapter 8 in Robert and Casella (2010) for more details). This adjusted sample size is referred to as the effective sample size. Checking the degree of autocorrelation in your Markov chains and estimating the effective sample size your chain has generated should be part of evaluating your model output. If you use WinBUGS through the R2WinBUGS package, the `print()` command will automatically return the effective sample size for all monitored parameters. In the coda package there are several functions you




Figure 7.6. Time series and posterior density plots of ψ and N for the bear data set truncated by the upper limit of M (500).

5606 can use to do so. `effectiveSize()` will directly give you an estimate of the effective
5607 sample size for your parameters:

```
5608 > effectiveSize(chain)
5609     sigma     lam0     psi      N
5610 3.930303 78.259159 30.436348 32.047392
```

5611 Alternatively, you can use the `autocorr.diag()` function, which will show you the
5612 degree of autocorrelation for different lag values (which you can specify within the
5613 function call, we use the defaults below):

```
5614 > autocorr.diag(mcmc(mod))
5615     sigma     lam0     psi      N
5616 Lag 0  1.0000000 1.0000000 1.0000000 1.0000000
5617 Lag 1  0.9979948 0.9494134 0.9847503 0.9774201
5618 Lag 5  0.9915567 0.8038168 0.9111951 0.9113525
5619 Lag 10 0.9836016 0.6714021 0.8462108 0.8509803
5620 Lag 50 0.8985337 0.1983780 0.6138516 0.6233994
```

5621 Whichever function you use, if you find that your supposedly long Markov chain
5622 has not generated enough pseudo-iid samples, you should consider a longer run. In
5623 the present case we see that autocorrelation is especially high for the parameter
5624 `sigma` and our effective sample size for this parameter is 4!⁹ This means we would

⁹Anyone have any idea how the autocorrelation in `sigma` could be reduced?

have to run the model for much longer to obtain a reasonable effective sample size. Unfortunately, with many SCR models we observe high degrees of serial autocorrelation, which means we have to run long chains to obtain enough samples that can be considered iid, in order to obtain reasonable estimates of our parameters and their variances. What exactly constitutes a reasonable effective sample size is hard to say, but as a rule of thumb you should probably aim at several hundreds of these pseudo-iid samples. A more meaningful measure of whether you've run your chain for enough iterations is the time-series or Monte Carlo error the 'noise' introduced into your samples by the stochastic MCMC process which we introduced in Chapter 2. The MC error decreases with increasing sample size and its magnitude can thus be controlled by adjusting the length of the Markov chain. As a rule of thumb, the MC error should be 1% or less of the parameter estimate. Once you have reached this level, the estimates of the mean, standard error and 95% quantiles should no longer change significantly with additional iterations. For highly correlated samples, it will take more iterations to reduce the MC error. In coda, the MC error is given as part of the summary results (see below). Another option to deal with the serial autocorrelation of samples is to 'thin' Markov chains by some rate r and save only every r -th iteration. But as discussed in Chapter 2, this is not efficient and should only be applied if needed for practical reasons (e.g. a large number of parameters and iterations may force you to thin your samples so you object storing the model output does not become unmanageably large). For now, let's continue using this small set of samples to continue looking at the output.

7.6.5 Summary results

Now that we checked that our chains apparently have converged and pretending that we have generated enough samples from the posterior distribution, we can look at the actual parameter estimates. The `summary()` function will return two sets of results: the mean parameter estimates, with their standard deviation, the naive standard error - i.e. your regular standard error calculated for K (= number of iterations) samples without accounting for serial autocorrelation - and the corrected MC error (Time-series SE), which accounts for autocorrelation. In WinBUGS, this latter value is referred to as MC error and is only given in the log output within BUGS itself. You should adjust the `summary()` call by removing the burn-in from calculating parameter summary statistics. To do so, use the `window()` command, which lets you specify at which iteration to start 'counting'. In contrast to WinBUGS, which requires you to set the burn-in length before you run the model, this command gives us full flexibility to make decisions about the burn-in after we have seen the trajectories of our Markov chains. For our example, `summary(window(chain, start=1001))` returns the following output:

```
Iterations = 1001:5000
Thinning interval = 1
```

```

5665 Number of chains = 1
5666 Sample size per chain = 4000
5667
5668 1. Empirical mean and standard deviation for each variable,
5669    plus standard error of the mean:

```

```

5670
5671      Mean      SD Naive SE Time-series SE
5672 sigma  1.9986  0.13805 0.0021827      0.016091
5673 lam0   0.1096  0.01523 0.0002407      0.001401
5674 psi    0.6113  0.09148 0.0014465      0.010734
5675 N      489.8535 71.79695 1.1352094      8.431119

```

```

5676
5677 2. Quantiles for each variable:

```

```

5678
5679      2.5%      25%      50%      75%      97.5%
5680 sigma  1.75780  1.89847  1.9900  2.0944  2.2772
5681 lam0   0.08357  0.09824  0.1087  0.1192  0.1427
5682 psi    0.45110  0.54838  0.6052  0.6639  0.8192
5683 N      366.00000 440.00000 485.0000 530.0000 654.0000

```

5684 Looking at the MC errors, we see that in spite of the high autocorrelation, the
 5685 MC error for sigma is below the 1Our algorithm gives us a posterior distribution of
 5686 N, but we are usually interested in the density, D. Density itself is not a parameter
 5687 of our model, but we can derive a posterior distribution for D by dividing each
 5688 value of N (N at each iteration) by the area of the state-space (here 3032.719 km²)
 5689 and we can use summary statistics of this distribution to characterize D:

```

5690 > summary(window(chain[,4]/ 3032.719, start=1001))
5691 Iterations = 1001:5000
5692 Thinning interval = 1
5693 Number of chains = 1
5694 Sample size per chain = 4000
5695
5696 1. Empirical mean and standard deviation for each variable,
5697    plus standard error of the mean:
5698
5699      Mean      SD      Naive SE Time-series SE
5700  0.1615229  0.0236741  0.0003743      0.0027801
5701
5702 2. Quantiles for each variable:
5703
5704      2.5%      25%      50%      75%      97.5%
5705 0.1207 0.1451 0.1599 0.1748 0.2156

```

If we compare our mean density of 0.16/km² (and other parameters) with results from the same model run in secr and WinBUGS in Chapter XX, we see that estimates are almost identical (Table 1).

7.6.6 Other useful commands

While inspecting the time series plot gives you a first idea of how well you tuned your MH algorithm, use `rejectionRate()` to obtain the rejection rates (1 - acceptance rates) of the parameters that are written to your output:

```
> rejectionRate(chain)
      sigma      lam0      psi      N
0.44108822 0.77675535 0.00000000 0.01940388
```

Recall that rejection rates should lie between 0.2 and 0.8, so our tuning seems to have been appropriate here. Psi is never rejected since we update it with Gibbs sampling, where all candidate values are kept. And since N is the sum of all z, all it takes for N to change from one iteration to the next are small changes in the z-vector, so the rejection rate of N is always low. If you have run several parallel chains, you can combine them into a single mcmc object using the `mcmc.list()` command on the individual chains (note that each chain has to be converted to an mcmc object before combining them with `mcmc.list()`). You can then easily obtain the Gelman-Rubin diagnostic (Gelman et al., 2004), in WinBUGS called R-hat, using `gelman.diag()`, which will indicate if all chains have converged to the same stationary distribution. For details on these and other functions, see the coda manual, which can be found together with the package on the CRAN mirror.

7.7 MANIPULATING THE STATE-SPACE

So far, we have constrained the location of the activity centers to fall within the outermost coordinates of our rectangular state space by posing upper and lower bounds for x and y. But what if S has an irregular shape maybe there is a large water body we would like to remove from S, because we know our terrestrial study species does not occur there. Or the study takes place in a clearly defined area such as an island. As mentioned before, this situation is difficult to handle in WinBUGS. In some simple cases we can adjust the state space by setting `SXi` to be some function of `SYi` or vice versa. In this manner, we can cut off corners of the rectangle to approximate the actual state space. In R, we are much more flexible, as we can use the actual state-space polygon to constrain out si.¹⁰ To illustrate that, let's look at a camera trapping study of Florida panthers (*Puma concolor coryi*) conducted in the Picayune Strand Restoration Project (PSRP) area, southwest Florida (Fig. 7), by XXX, and financed by XXX. In the 1960ies the PSRP area was slated for

¹⁰ Have to check if we can use panther stuff for the book; otherwise, use raccoon example.

housing development, but then bought back by the State of Florida and is currently being restored to its original hydrology and vegetation. In an effort to estimate the density of the local Florida panther population, 98 camera traps were operated in the area for 21 months between 2005 and 2007. Florida panthers are wide-ranging animals and in order to account for their wide movements, the state-space was defined as the trapping grid buffered by 15 km around its outermost coordinates. However, the resulting rectangle contained some ocean in its southwestern corner (Fig. 7). In order to precisely describe the state-space, the ocean has to be removed. You can create a precise state-space polygon in ArcGIS and read it into R, or create the polygon directly within R. In the present case we intersected two shape files one of the state of Florida and one of the rectangle defined by a strip of 15 km around the camera-trapping grid. While you will most likely have to obtain the shapefile describing the landscape of and around your trapping grid (coastlines, water bodies etc.) from some external source, a polygon shapefile buffering your outermost trapping grid coordinates can easily be written in R.

If `xmin`, `xmax`, `ymin` and `ymax`, mark the outermost `x` and `y` coordinates of your trapping grid and `b` is the distance you want to buffer with, load the package `shapefiles` (Stabler, 2006) and use:

```

5759 x1= xmin-b
5760 xu= xmax+b
5761 y1= ymin-b
5762 yu= ymax+b
5763
5764 dd <- data.frame(Id=c(1,1,1,1,1),X=c(x1,xu,xu,x1,x1),Y=c(y1,y1,yu,yu,y1)) #create data fra
5765 ddTable <- data.frame(Id=c(1),Name=c("Item1"))
5766 ddShapefile <- convert.to.shapefile(dd, ddTable, "Id", 5) #convert #to shapefile, type pol
5767 write.shapefile(ddShapefile, 'c:/', arcgis=T) # save to location of #choice

```

You can read shapefiles into R loading the package `maptools` (Lewin-Koh et al., 2011) and using the function `readShapeSpatial()`. Make sure you read in shapefiles in UTM format, so that units of the trap array, the movement parameter `sigma` and the state-space are all identical. Intersection of polygons can be done in R also, using the package `rgeos` (Bivand and Rundel, 2011) and the function `gIntersect()`. The area of your single - polygon can be extracted directly from the state-space object `SSp`:

```

5775 > area <- SSp@polygons[[1]]@Polygons[[1]]@area /1000000

```

Note that dividing by 1000000 will return the area in km² if your coordinates describing the polygon are in UTM. If your state-space consists of several disjunct polygons, you will have to sum the areas of all polygons to obtain the size of the state-space. To include this polygon into our MCMC sampler we need one last spatial R package `sp` (Pebesma and Bivand, 2011), which has a function, `over()`,




Figure 7.7. Rectangular state-space for a Florida panther camera trapping study in the PSRP area (grey outline, red block inset map of Florida) contain some ocean (white) that needs to be removed from the state-space.

```

5781 which allows us to check if a pair of coordinates falls within a polygon or not. All
5782 we have to do is embed this new check into the updating steps for s:

5783         Scand <- as.matrix(cbind(rnorm(M, S[,1], 2),
5784                                 rnorm(M, S[,2], 2)))          #draw candidate value
5785
5786 Scoord<-SpatialPoints(Scand*1000)      #convert to spatial points on UTM (m) scale
5787 SinPoly<-over(Scoord,SSp) # check if scand is within the polygon
5788
5789         for(i in 1:M) {
5790 if(is.na(SinPoly[i])==FALSE) { #if scand falls within polygon, continue update
5791   [rest of the updating step remains the same]

```

```

5792 Note that it is much more time-efficient to draw all M candidate values for s and
5793 check once if they fall within the state-space, rather than running the over() com-
5794 mand for every individual pair of coordinates. To make sure that our initial values
5795 for s also fall within the polygon of S, we use the function runifpoint() from the
5796 package spatstat (Baddeley and Turner, 2005), which generates random uniform
5797 points within a specified polygon. You'll find this modified MCMC algorithm in
5798 the online supplementary material (SCR0poisSSp). Finally, observe that we are
5799 converting candidate coordinates of S back to meters to match the UTM polygon.
5800 In all previous examples, for both the trap locations and the activity centers we

```

5801 have used UTM coordinates divided by 1000 to estimate sigma on a km scale. This
5802 is adequate for wide ranging individuals like bears. In other cases you may center
5803 all coordinates on 0. No matter what kind of transformation you use on your co-
5804 ordinates , make sure to always convert candidate values for S back to the original
5805 scale (UTM) before running the `over()` command.

7.8 MCMC SOFTWARE PACKAGES

5806 Throughout most of this book we will use WinBUGS and, occasionally, JAGS to
5807 run MCMC analyses. Here, we will briefly discuss the main pros and cons of these
5808 two programs as well as WinBUGS successor OpenBUGS. You can find scripts to
5809 simulate data and run the basic SCR model in all three programs in the online
5810 supplementary material (`simSCR0poisBUGS`).

7.8.1 WinBUGS

5812 In a nutshell, WinBUGS (and the other programs) do everything that we just went
5813 through in this chapter (and quite a bit more). Looking through your model, Win-
5814 BUGS determines which parameters it can use standard Gibbs sampling for (i.e.
5815 for conjugate full conditional distributions). Then, it determines, in the following
5816 hierarchy, whether to use adaptive rejection sampling, slice sampling or in the
5817 'worst' case Metropolis-Hastings sampling for the other full conditionals (Spiegel-
5818 halter et al., 2003). If it uses MH sampling, it will automatically tune the updater
5819 so that it works efficiently. While WinBUGS is a convenient piece of software that
5820 is still widely used, its major drawback is that it is no longer being developed, i.e.
5821 no new functions or distributions are added and no bugs are fixed.

7.8.2 OpenBUGS

5823 OpenBUGS is essentially the successor of WinBUGS. While the latter is no longer
5824 worked on, OpenBUGS is constantly developed further. The name 'OpenBUGS'
5825 refers to the software being open source, so users do not need to download a license
5826 key, like they have to for WinBUGS (although the license key for WinBUGS is free
5827 and valid for life).

5828 Compared to WinBUGS, OpenBUGS has a lot more built-in functions. The
5829 method of how to determine the right updater for each model parameter has
5830 changed and the user can manually control the MCMC algorithm used to update
5831 model parameters. Several other changes have been implemented in OpenBUGS
5832 and a detailed list of differences between the two BUGS versions, can be found at
5833 <http://www.openbugs.info/w/OpenVsWin>

5834 While OpenBUGS is a useful program for a lot of MCMC sampling applications,
5835 for reasons we do not understand, simple SCR models do not converge in Open-
5836 BUGS. It is therefore advisable that you check any OpenBUGS SCR model results

against result from WinBUGS. Also, currently, the R package BRugs (Thomas et al., 2006) necessary for running OpenBUGS through R has problems with 64-bit machines, so you may have to use the 32-bit version of R and OpenBUGS in order to make it work. The BUGS project site at <http://www.openbugs.info> provides a lot of information on and download links for OpenBUGS.

There is an extensive help archive for both WinBUGS and OpenBUGS and you can subscribe to a mailing list, where people pose and answer questions of how to use these programs at <http://www.mrc-bsu.cam.ac.uk/bugs/overview/list.shtml>

7.8.3 JAGS Just Another Gibbs Sampler

JAGS, currently at Version 3.1.0, is another free program for analysis of Bayesian hierarchical models using MCMC simulation. Originally, JAGS was the only program using the BUGS language that would run on operating systems other than the 32 bit Windows platforms. By now, there are OpenBUGS versions for Linux or Macintosh machines. JAGS 'only' generates samples from the posterior distribution; analysis of the output is done in R either by running JAGS through R using either the packages `rjags` (Plummer, 2011) or `R2jags` (Su and Yajima, 2011), or by using `coda` on your JAGS output. The program, manuals and `rjags` can be downloaded at <http://sourceforge.net/projects/mcmc-jags/files/> When run from within R using the package `rjags` or `R2jags`, writing a JAGS model is virtually identical to writing a WinBUGS model. However, some functions may have slightly different names and you can look up available functions and their use in the JAGS manual. One potential downside is that JAGS can be very particular when it comes to initial values. These may have to be set as close to truth as possible for the model to start. Although JAGS lets you run several parallel Markov chains, this characteristic interferes with the idea of using overdispersed initial values for the different chains. Also, we have occasionally experienced JAGS to crash and take the R GUI with it. Only re-installing both JAGS and `rjags` seemed to solve this problem. On the plus side, JAGS usually runs a little faster than WinBUGS, sometimes considerably faster (see section 4.XYZ), is constantly being developed and improved and it has a variety of functions that are not available in WinBUGS. For example, JAGS allows you to supply observed data for some deterministic functions of unobserved variables. In BUGS we cannot supply data to logical nodes. Another useful feature is that the adaptive phase of the model (the burn-in) is run separately from the sampling from the stationary Markov chains. This allows you to easily add more iterations to the adaptive phase if necessary without the need to start from 0. There are other, more subtle differences and there is an entire manual section on differences between JAGS and OpenBUGS. For questions and problems there is a JAGS forum online at <http://sourceforge.net/projects/mcmc-jags/forums/forum/610037>.

¹¹

¹¹As we make progress on the book, lets be sure to add linkages to places where we use JAGS in examples.

7.9 SUMMARY AND OUTLOOK

While there are a number of flexible and extremely useful software packages to perform MCMC simulations, it sometimes is more efficient to develop your own MCMC algorithm. Building an MCMC code follows three basic steps: Identify your model including priors and express full conditional distributions for each model parameter. If full conditionals are parametric distributions, use Gibbs sampling to draw candidate parameter values from this distributions; otherwise use Metropolis-Hastings sampling to draw candidate values from a proposal distribution and accept or reject them based on their posterior probability densities. These custom-made MCMC algorithms give you more modeling flexibility than existing software packages, especially when it comes to handling the state-space: In BUGS (and JAGS for that matter) we define a continuous rectangular state-space using the corner coordinates to constrain the Uniform priors on the activity centers s . But what if a continuous rectangle isn't an adequate description of the state-space? In this chapter we saw that in R it only takes a few lines of code to use any arbitrary polygon shapefile as the state-space, which is especially useful when you are dealing with coastlines or large bodies of water that need removing from the state-space. Another example is the SCR R package SPACECAP (Gopalaswamy et al., 2011) that was developed because implementation of an SCR model with a discrete state-space was inefficient in WinBUGS. Another situations in which using BUGS/JAGS becomes increasingly complicated or inefficient is when using point processes other than the Uniform Poisson point process which underlies the basic SCR model (see Chapter X). In the Chapters 9 and XX you will see examples of different point processes, implemented using custom-made MCMC algorithms.¹² Finally, the Chapters XX and XX deal with unmarked or partially marked populations using hand-made MCMC algorithms to handle the (partially) latent individual encounter histories. While some of these models can be written in BUGS/JAGS,¹³ they are painstakingly slow; others cannot be implemented in BUGS/JAGS at all. In conclusion, while you can certainly get by using BUGS/JAGS for standard SCR models, knowing how to write your own MCMC sampler allows you to tailor these models to your specific needs.

¹²Richard, Beth expand on that?

¹³the Poisson one for partially marked we wrote in BUGS and it should work with a known number of marked; the Bernoulli in JAGS with the `dsum()` function should work for the fully unknown; maybe some others? I don't remember. We may have to try writing the others before saying that they don't work in BUGS/JAGS; they are certainly much faster in R, though.

8

5906

5907

5908

GOODNESS OF FIT AND STUFF

9

5909

5910

5911

MODELING ENCOUNTER PROBABILITY

ECOLOGICAL DISTANCE MODELS IN SPATIAL CAPTURE-RECAPTURE

Previously we have only considered stationary and symmetric models for detection probability. While these will often be sufficient for practical purposes, especially in small data sets, there will sometimes be interest in developing more complex models of the detection process as it relates to space usage of individuals.

These models are simple because they are based on Euclidean distance – which are convenient because we know this distance precisely conditional on individual activity center \mathbf{s} . However, animals may not judge distance in terms of euclidean distance but, rather, according to local habitat conditions and so forth.

In this section we develop models for detection probability based on a distance metric that tries to mimic ecological distance.

10.1 COST DISTANCE

We use a cost-weighted distance metric in the package `gdistance` which computes the distance between points by accumulating pixel-specific costs assigned by the user (but we consider estimating these in Sec. XYZ). The idea is due to XYZ XYZ REF XYZ. The basic idea is that distance bewtween x_i and x_j is

$$wd_{ij} = w_{ij} \dots \text{formula for this?}$$

where and so on....XXX XYZ. To be consistent with the functioning of `costdistance` the incremental cost is assessed for *leaving* a pixel, not arriving. As a consequence, the terminal value is not counted. As an example of the cost-weighted distance calculation consider the following landscape comprised of 9 pixels identified as follows, along with their costs:

pixel ID	cost
1 2 3	100 1 1
4 5 6	100 100 1
7 8 9	100 100 1

Then we assign low cost of 1 to “good habitat” pixels (or pixels we think of as “highly connected” by virtue of being in good habitat) and, conversely, we assign high cost (100) to “bad habitat”. So the cost-weighted distance between pixels 2 and 3 in this example is just 1 unit, the distance between pixels 2 and 6 is $\sqrt{2}$ units, and the distance between pixels 4 and 5 is 100 units, while the cost-distance between 4 and 6 is 101.

Do the calculation here:

Alternatively we might label our pixels by transposing the id’s so that pixels 1, 4 and 7 represent the first row:

1 4 7
2 5 8
3 6 9

In this case we have to transpose or rotate something along the way. Be careful.

10.1.1 Illustration: Example Good vs. Bad habitat

We made up a polygon which we might think of as representing a habitat corridor or park unit or something which has been preserved. It is surrounded by a suburban wasteland of McDonalds and Wal-Marts, much less hospitable habitat for most things. Figure here.

10.2 DISTANCE WEIGHTING

10.3 HARD BOUNDARY

10.4 SIMULATION STUDY

evaluate effect of ignoring ecological distance

10.5 ESTIMATING COST OR RESISTANCE VALUES

STATE-SPACE COVARIATES

Underlying all spatial capture recapture models is a point process model describing the distribution of individual activity centers (\mathbf{s}_i) within the state space (\mathcal{S}). So far we have focused our discussion on the homogeneous binomial point process, $\mathbf{s}_i \sim \text{Uniform}(\mathcal{S}), i = 1, 2, \dots, N$, where N is the size of the population. This is a model of “spatial-randomness”¹ because the intensity of the activity centers is constant across the study area and the activity centers are distributed independently of each other.

The spatial-randomness assumption is often viewed as restrictive because ecological processes such as territoriality and habitat selection can result in non-random distributions of organisms. We have argued, however, that this assumption is less restrictive than may be recognized because the homogeneous point process actually allows for infinite possible configurations of activity centers. Furthermore, given enough data, the uniform prior will have very little influence on the estimated locations of activity centers. Nonetheless, the homogeneous point process model does not allow one to model population density using covariates—a central objective of much ecological research. For example, a homogeneous point process model may result in a density surface map indicating that individuals were more abundant in one habitat than another, but it does not do so explicitly. A more direct approach would be to model density using covariates as is done in generalized linear models (GLMs).

In this chapter we will present a method for fitting inhomogeneous binomial point process models using covariates in much the same way as is done with GLMs. The covariates we consider differ from those covered in previous chapters, which were typically attributes of the animal (*e.g.* sex, age) and were used to model movement or encounter rate. In contrast, here we wish to model covariates that

¹The phrase “complete spatial-randomness” is reserved for the homogeneous Poisson point process

are defined for all points in \mathcal{S} , which we will refer to as state-space, or density, covariates. These may include continuous covariates such as elevation, or discrete covariates such as habitat type.

Borchers and Efford (2008) were the first to propose an inhomogeneous point process model for SCR models, and our approach is similar to theirs with the exception that we will use a binomial rather than a Poisson model because the binomial model is easily integrated into our data augmentation scheme and is consistent with the objective of determining how a *fixed* number of activity centers are distributed with respect to covariates.

The method we use to accommodate inhomogeneous binomial point process models within our MCMC algorithm is simple—we replace the uniform prior with a prior describing the distribution of the N activity centers conditional on the covariates. Development of this prior, which does not have a standard form, is a central component of this chapter. First we will begin with a review of homogeneous point process models.

11.1 HOMOGENEOUS POINT PROCESS REVISITED

The homogeneous Poisson point process is *the* model of “complete spatial randomness” and is often used in ecology as a null model to test for departures from randomness. Given its central role in the analysis of point processes, it is helpful to compare it with the binomial model that we use in our SCR models. The primary descriptor of the homogeneous point process model is the “intensity” parameter, μ which describes the expected number of points in an infinitesimally small area. The intensity parameter can also be used to determine the expected number of points in any region of the state-space \mathcal{S} . To denote this, we say that the expected number of points in region $B \in \mathcal{S}$ is $n(B) = A(B)\mu$ where $A(B)$ is the area of region B . In words, the expected number of points in B is simply the area of B multiplied by the intensity parameter. One property of the Poisson model is that if we divide the entire state-space into $k = 1, \dots, K$ disjunct regions, the counts $\mathbf{n}(\mathbf{B})$ are independent and identically distributed, (*i.i.d.*). This is one of the distinctions between the Poisson model and the binomial model, for which the counts $n(B_k)$ are not *i.i.d.* as we will explain shortly. This difference is also related to another distinction between the two models, namely that the binomial model conditions on the number of points to be simulated N ; whereas under the Poisson model N is random. Here is some simple **R** code to illustrate this point.

```

6019 mu <- 4                                # intensity
6020 Np <- rpois(1, mu)                     # Np is random
6021 PPP <- cbind(runif(Np), runif(Np)) # Poisson point process
6022
6023 Nb <- 4                                # Nb is fixed
6024 BPP <- cbind(runif(Nb), runif(Nb)) # Binomial point process

```

Note that in both models, the N points are independent of one another and distributed uniformly throughout \mathcal{S} . Thus, the intensity at any point $x \in \mathcal{S}$ is $\mu = 1/A(\mathcal{S})$ where $A(\mathcal{S})$ denotes the area of the state-space. In the **R** code above, the area of the state-space is 1 unit, and thus the intensity is $\mu = 1/1$.

Although the Poisson model is typically described in terms of μ , the binomial model is not; rather, it is more common to consider a discrete state space, such as a grid with K pixels. Under the binomial model, the number of points in each region is $n(B_k) \sim \text{Bin}(N, p_k)$ where $p_k = A(B_k)/A(\mathcal{S})$, ie p_k is simply the fraction of the state-space area in B_k . This discrete space representation of the binomial point process is shown in Fig. 11.1. The state-space in this case is the unit square, and thus the probability of a point falling in each of the 25 disjunct regions is $p_k = 1/25$ and thus the expected counts are simply $\mathbb{E}(n(B_k)) = Np_k$. In the figure $N = 50$ and thus we would expect 2 points per pixel, which happens to be the empirical mean of the data in Fig. 11.1. Note also that these counts are not independent realizations from a binomial distribution since $\sum_k n(B_k) = N$. Instead, the model for the entire vector is $\mathbf{n}(\mathbf{B}) \sim \text{Multinomial}(N, \pi = (p_1, p_2, \dots, p_K))$ (Illian, 2008b). The dependence among counts has virtually no practical consequence when the number of pixels is large. For example, if we have 100 pixels, the number of counts in one pixels tells you very little about the expected count in another pixel. However, if there are only 2 pixels, then clearly the number of points in one pixel tells you exactly how many will occur in the remaining pixel. To gain familiarity with the multinomial distribution and the discrete representation of space, use the `rmultinom` function in **R** to simulate counts similar to those shown in Fig. 11.1, for example using a command such as:

```
n.B_k <- rmultinom(1, size=50, prob=rep(1/25, 25))
matrix(n.B_k, 5, 5)
```

The discrete space representation of the binomial point process is of practical importance when fitting SCR models because spatial covariates are almost always represented in a discrete format, often called “rasters” in GIS-speak. In such cases, we often need to change our definition of the prior for an activity center from $s_i \sim \text{Uniform}(\mathcal{S})$ to $s_i \sim \text{Multinomial}(1, \pi)$. In the latter case, the activity center is simply defined as an integer representing pixel “id”. Note also that the multinomial distribution with an index of 1 (i.e. `size=1` in `rmultinom`) is referred to as the categorical distribution, which we will frequently use in the BUGS language.

11.2 INHOMOGENEOUS BINOMIAL POINT PROCESS

As with the homogeneous model, the inhomogeneous binomial point process model is developed conditional on N . The primary distinction is that the uniform distribution is replaced with another distribution allowing for the intensity parameter to vary spatially. To arrive at this new distribution, define $\mu(x, \beta)$ to be a function of




Figure 11.1. Homogeneous binomial point process with $N=50$ points represented in continuous and discrete space.

6063 spatially-referenced covariates (β) available at all points of the state space. To be
 6064 concise we will subsequently drop the vector of coefficients from our notation, and
 6065 simply use $\mu(x)$. Since an intensity must be strictly positive, it is natural to model
 6066 $\mu(x)$ using the log-link.

$$\log(\mu(x)) = \sum_{j=1}^J \beta_j v_j(x), \quad x \in \mathcal{S}$$

6067 where β_j is the regression coefficient for covariate $v_j(x)$. To be clear, $v(x)$ is the
 6068 value of any covariate, such as habitat type or elevation, at location x . This equa-
 6069 tion should look familiar because it is the standard linear model used in log-linear
 6070 GLMs. Note, however, that we have no need for an intercept because it would be
 6071 confounded with N . This should be intuitive since an intercept would represent the
 6072 expected value of N when $\beta = 0$, but we already have a parameter in the model for
 6073 expected abundance, namely $\mathbb{E}[N] = \psi M$. Thus an intercept would be redundant,
 6074 and without it we are still able to achieve our goal of describing the distribution of
 6075 N activity centers as a function of spatial covariates.

6076 Now that we have a model of the intensity parameter $\mu(x)$, we need to develop
 6077 the associated probability density function to use in place of the uniform prior.
 6078 Remembering that the integral of a pdf must be unity, we can create a pdf by
 6079 dividing $\mu(x)$ by a normalizing constant, which in this case is the integral of $\mu(x)$
 6080 evaluated over the entire state-space. **ANDY, is there a better justification for this?**

The probability density function is therefore

$$f(x) = \frac{\mu(x)}{\int_{x \in S} \mu(x) dx} \quad (11.2.1)$$

Substituting this distribution for the uniform prior allows us to fit inhomogeneous binomial point process models to spatial capture-recapture data. We can also use this distribution to obtain the expected number of individuals in any given region. Specifically, the proportion of N expected to occur in any region B when heterogeneity in density is present is $p(B) = \int_B f(x) dx$. These are also the multinomial cell probabilities if the regions are disjoint and compose the entire state-space.

As a practical matter, note that the integral in the denominator of $f(x)$ is evaluated over space, and since we almost always regard space as two-dimensional, this is a two-dimensional integral that can be approximated using the methods discussed in refChXXX. These methods include Monte Carlo integration, Gaussian quadrature, etc... Alternatively, if our state-space covariates are in raster format, *i.e* they are in discrete space, the integral can be replaced with a sum over all pixels, which is much more efficient computationally.

We now have all the tools needed to fit inhomogeneous point process (IPP) models. Before doing so, we note that the IPP for the activity centers results in another IPP for the observation process, $\lambda(x)$. As a reminder, $\lambda(x)$ is the expected number of captures for a trap at point x . As was true for the homogeneous model, this intensity function is a product of the point process intensity and the encounter rate function, $\lambda(x) = \mu(x)g(x)$.

In the next section we walk through a few examples, building up from the simplest case where we actually observe the activity centers as though they were data. In the second example, we fit our new model to simulated data in which density is a function of a single continuous covariate. Example three shows an analysis in discrete space using both **secr** (Efford, 2011) and **JAGS** (Plummer, 2003). In the last example, we model the intensity of activity centers for a real dataset collected on jaguars (*Panthera onca*) in Argentina.

11.3 EXAMPLES

11.3.1 Simulation and analysis of inhomogeneous point processes

In SCR models, the point process is not directly observed, but in other contexts it is. Examples include the locations of disease outbreaks, the locations of trees in a forest, or the locations of radio-tracked animals. Indeed Eq. 11.2.1 has been used extensively in the radio-telemetry literature to model so-called “resource selection functions” (??). In such cases where the point locations are directly observed, estimating the parameters β is straight-forward as demonstrated in the following example. This example also illustrates the fundamental process that we will later embed in our MCMC algorithm used to fit SCR models.

Suppose we knew the locations of 100 animals' activity centers. To estimate the intensity surface $\mu(x)$ underlying these points, we need to derive the likelihood for our data under this model. Given the pdf $f(x)$ (Eq. 11.2.1) and assuming that the points are mutually independent of one another, we may write the likelihood as the product of R such terms, where $R = 100$ is the sample size in this case, *ie* the observed number of activity centers.

$$\mathcal{L}(\beta|\mathbf{x}_i) = \prod_{i=1}^R f(x_i)$$

Having defined the likelihood we could choose a prior and obtain the posterior for β using Bayesian methods, or we can find the maximum likelihood estimates (MLEs) using standard numerical methods as is demonstrated below.

First, let's simulate some data. Simulating data under an inhomogeneous point process model is often accomplished using indirect methods such as rejection sampling. Rejection sampling proceeds by simulating data from a standard distribution and then accepting or rejecting each sample using probabilities defined by the distribution of interest. For more information, readers should consult an accessible text such as Robert and Casella (2004). In our example, we simulate from a uniform distribution and then accept or reject using the (scaled) probability density function $f(x)$. Note that we first define a spatial covariate (elevation) that is a simple function of the spatial coordinates increasing from the southwest to the northeast of our state-space.²

The following **R** commands demonstrate the use of rejection sampling to simulate an inhomogeneous point process for the covariate depicted in Fig. 11.3.1. The code uses the **cuhre** function in the **R2Cuba** package to integrate the intensity function over space (Hahn et al., 2011).

```
# spatial covariate (with mean 0)
elev.fn <- function(x) x[1]+x[2]-1
# intensity function
mu <- function(x, beta) exp(beta*elev.fn(x=x))

# Simulate PP using rejection sampling
set.seed(300225)
N <- 100
count <- 1
s <- matrix(NA, N, 2)
beta <- 2 # parameter of interest
int.mu <- R2Cuba::cuhre(2, 1, mu, beta=beta)$value
elev.min <- elev.fn(c(0,0)) #elev.fn(cbind(0,0))
elev.max <- elev.fn(c(1,1)) #elev.fn(cbind(1,1))
```

²Such functional forms of covariates are rarely available, which is why continuous spatial covariates are more often measured on a discrete grid.




Figure 11.2. An example of a spatial covariate, say elevation, and a realization of an inhomogeneous binomial point process with $N=100$ and $\mu(x) = \exp(\beta \text{Elev})$ where $\beta = 2$.

```

6154 Q <- max(c(exp(beta*elev.min) / int.mu, #2d(beta),
6155             exp(beta*elev.max) / int.mu)) #2d(beta))
6156 while(count <= 100) {
6157   x.c <- runif(1, 0, 1); y.c <- runif(1, 0, 1)
6158   s.cand <- c(x.c,y.c)
6159   pr <- exp(beta*elev.fn(s.cand)) / int.mu #2d(beta)
6160   if(runif(1) < pr/Q) {
6161     s[count,] <- s.cand
6162     count <- count+1
6163   }
6164 }

```

6165 The simulated data are shown in Fig 11.3.1. High elevations are represented by
 6166 light green and low elevations by dark green. The activity centers of one hundred
 6167 animals are shown as points, and it is clear that these simulated animals prefer the
 6168 high elevations. Perhaps they are mountain goats. The underlying model describing
 6169 this preference is $\log(\mu(x)) = \exp(\beta \times \text{Elevation}(x))$ where $\beta = 2$ is the parameter
 6170 to be estimated and $\text{Elevation}(x)$ is a function of the coordinates at x , as displayed
 6171 on the map.

6172 Given these points, we will now estimate β by minimizing the negative-log-
 6173 likelihood using R's `optim` function.

```

6174 # Negative log-likelihood

```

```

6175 nll <- function(beta) {
6176   int.mu <- cuhre(2, 1, mu, beta=beta)$value
6177   -sum(beta*elev.fn(s) - log(int.mu))
6178 }
6179 starting.value <- 0
6180 fm <- optim(starting.value, nll, method="Brent",
6181            lower=-5, upper=5, hessian=TRUE)
6182 c(Est=fm$par, SE=sqrt(1/fm$hessian)) # estimates and SEs

```

6183 Maximizing the likelihood took a small fraction of a second, and we obtained
 6184 an estimate of $\hat{\beta} = 1.99$. We could plug in this estimate to our linear model at each
 6185 point in the state-space to obtain the MLE for the intensity surface.

6186 This example demonstrates that if we had the data we wish we had, *i.e.* if we
 6187 knew the coordinates of the activity centers, we could easily estimate the parameters
 6188 governing the underlying point process. Unfortunately, in SCR models, the activity
 6189 centers cannot be directly observed, but spatial re-captures, that is captures of
 6190 individuals at multiple locations in space, provide us with the information needed
 6191 to estimate these latent parameters.

6192 11.3.2 Fitting inhomogeneous point process SCR models

6193 Continuous space

6194 One of the nice things about hierarchical models is that they allow us to break a
 6195 problem up into a series of simple conditional relationships. Thus, we can simply
 6196 add the methods described above into our existing MCMC algorithm to simulate
 6197 the posteriors of β conditional on the simulated values of \mathbf{s}_i . To demonstrate, we
 6198 will continue with the previous example. Specifically, we will overlay a grid of
 6199 traps upon the map shown in Fig. 11.3.1. We will then simulate capture histories
 6200 conditional upon the activity centers shown on the map. Then, we will attempt to
 6201 estimate the activity center locations as though we did not know where they were,
 6202 as is the case in real applications.

6203 Here is some **R** code to simulate the encounter histories under a Poisson ob-
 6204 servation model, which would be appropriate if animals could be detected multiple
 6205 times at a trap during a single occasion.

```

6206 # Create trap locations
6207 xsp <- seq(-0.8, 0.8, by=0.2)
6208 len <- length(xsp)
6209 X <- cbind(rep(xsp, each=len), rep(xsp, times=len))
6210
6211 # Simulate capture histories, and augment the data
6212 ntraps <- nrow(X)
6213 T <- 5
6214 y <- array(NA, c(N, ntraps, T))
6215

```



```

6216 nz <- 50 # augmentation
6217 M <- nz+nrow(y)
6218 yz <- array(0, c(M, ntraps, T))
6219
6220 sigma <- 0.1 # half-normal scale parameter
6221 lam0 <- 0.5 # basal encounter rate
6222 lam <- matrix(NA, N, ntraps)
6223
6224 set.seed(5588)
6225 for(i in 1:N) {
6226   for(j in 1:ntraps) {
6227     distSq <- (s[i,1]-X[j,1])^2 + (s[i,2] - X[j,2])^2
6228     lam[i,j] <- exp(-distSq/(2*sigma^2)) * lam0
6229     y[i,j,] <- rpois(T, lam[i,j])
6230   }
6231 }
6232 yz[1:nrow(y),,] <- y # Fill

```

Now that we have a simulated capture-recapture dataset y , and we have augmented it to create the new data object yz , we are ready to begin sampling from the posteriors. A commented Gibbs sampler written in **R** is available in the accompanying **R** package **scrbook** (see ?scrIPP). There are two small parts of the **R** code that distinguish it from previous code we have shown to fit homogeneous point processes. First, we need to update the parameter β conditional on all other parameters in the model. The code to do so is:

```

6240 D1 <- cuhre(2, 1, mu, lower=c(xlims[1], ylims[1]),
6241             upper=c(xlims[2], ylims[2]), beta=beta1)$value
6242 beta1.cand <- rnorm(1, beta1, tune[3])
6243 D1.cand <- cuhre(2, 1, mu, lower=c(xlims[1], ylims[1]),
6244                 upper=c(xlims[2], ylims[2]), beta=beta1.cand)$value
6245 ll.beta1 <- sum( beta1*elev.fn.v(S) - log(D1) )
6246 ll.beta1.cand <- sum( beta1.cand*elev.fn.v(S) - log(D1.cand) )
6247 if(runif(1) < exp(ll.beta1.cand - ll.beta1) ) {
6248   beta1<-beta1.cand
6249 }

```

Next, we need to put the new prior on the activity centers:

```

6251 #ln(prior), denominator is constant
6252 prior.S <- beta1*cov(S[i,1], S[i,2]) # - log(D1)
6253 prior.S.cand <- beta1*(Scand[1] + Scand[2]) # - log(D1)
6254 if(runif(1)< exp((ll.S.cand+prior.S.cand) - (ll.S+prior.S))) {
6255   S[i,] <- Scand
6256   lam <- lam.cand
6257   D[i,] <- dtmp
6258 }

```

We can apply this modified sampler to our data using the code shown in the help file for `scrIPP`. We obtain posterior distributions summarized in Table 11.1. Mixing is good, and as usual, life is very nice when we are working with simulated data.

Fitting continuous space IPP models is somewhat difficult in **BUGS** because our prior $f(x)$ is not one of the available distributions that come with the software³ `secr` allows users to fit continuous space using polynomials of the x - and y - coordinates, but not for truly continuous covariates. However, these are not really important limitations because discrete space versions are straight-forward, and virtually all spatial covariates are defined as such.

Discrete space

To fit discrete space models, we follow the same steps as outlined in Chapter XXX—we define s_i as pixel ID, and we use the categorical distribution as a prior. A good example of this is in `+citeKery capricallie`. Here we present an analysis of the simulated data shown in the right panel of Fig. 11.3.1. The spatial covariate, let's call it elevation again, was simulated from a kriging type of model as shown on the help page `ch9simData` in `scrbook`. The points are the number of activity centers in each pixel, generated from a single realization of the IPP $\mu(x) = 2elev$.

The **BUGS** code to fit an IPP model to these data is shown in the following panel.

```

model{
  sigma ~ dunif(0, 1)
  lam0 ~ dunif(0, 5)
  beta ~ dnorm(0,0.1)
  psi ~ dbeta(1,1)

  for(j in 1:nPix) {
    theta[j] <- exp(beta*elevation[j])
    probs[j] <- theta[j]/sum(theta[])
  }

  for(i in 1:M) {

```

³It is possible, if somewhat cumbersome, to add new distributions in **BUGS**.

Table 11.1. Posterior summaries from inhomogeneous point proces model

	Mean	SD	2.5%	50%	97.5%
$\sigma = 0.10$	0.1026	0.0048	0.0935	0.1025	0.1123
$\lambda_0 = 0.50$	0.4419	0.0493	0.3496	0.4400	0.5390
$\psi = 0.66$	0.6826	0.0554	0.5762	0.6820	0.7923
$\beta = 2.00$	2.1601	0.3390	1.5193	2.1583	2.8043
$N = 100$	102.7696	6.2689	92.0000	102.0000	117.0000




Figure 11.3. Simulated activity centers in discrete space. The spatial covariate, elevation, is highest in the higher areas. Density of activity centers (circles) increases with elevation. Trap locations are shown as crosses.

```

6291   w[i] ~ dbern(psi)
6292   s[i] ~ dcat(probs[])
6293   x0g[i] <- Sgrid[s[i],1]
6294   y0g[i] <- Sgrid[s[i],2]
6295   for(j in 1:ntraps) {
6296     dist[i,j] <- sqrt(pow(x0g[i]-grid[j,1],2) + pow(y0g[i]-grid[j,2],2))
6297     lambda[i,j] <- lam0*exp(-dist[i,j]*dist[i,j]/(2*sigma*sigma)) * w[i]
6298     y[i,j] ~ dpois(lambda[i,j])
6299   }
6300 }
6301
6302 N <- sum(w[])
6303 Density <- N/1 # unit square
6304 }

```

6305 This model can also be fit in **secr**, which refers to the pixel locations as a
 6306 “mask”. **R** code to fit the models using **secr** and **JAGS** is available in **scrbook**,
 6307 see **help(ch9secrYjags)**. Results of the comparison are shown in Table 11.2 and
 6308 are very similar as expected.

6309 Density surface maps can be created for fun, and of course to inform manage-
 6310 ment decisions. [describe how to do this]

6311 11.3.3 The jaguar data

6312 Estimating density of large felines has been a priority for many conservation orga-
 6313 nizations, but no robust methodologies existed before the advent of SCR. Distance
 6314 sampling is not feasible for such rare and cryptic species, and traditional capture-
 6315 recapture methods yield estimates that are highly sensitive to the subjective choice
 6316 of the effective survey area. In this example, we demonstrate how readily density
 6317 can be estimated for a globally imperilled species using SCR. Furthermore, we show
 6318 how inhomogeneous point process models can be used to test important hypotheses
 6319 regarding the factors affecting density.

6320 In this example, we make use a single year of data from an 8-year camera-

Table 11.2. Comparison of **secr** and **JAGS** results

Software	Par	Est.	SD	lower	upper
secr	N	49.2803	5.7535	41.0087	64.3879
	β	2.1772	0.5628	1.0741	3.2804
	λ_0	0.9203	0.0764	0.7824	1.0825
	σ	0.0990	0.0038	0.0918	0.1068
JAGS	N	48.2072	5.4053	39.0000	60.0000
	β	2.1026	0.5323	1.0889	3.1506
	λ_0	0.9328	0.0766	0.7898	1.0921
	σ	0.1004	0.0041	0.0929	0.1089

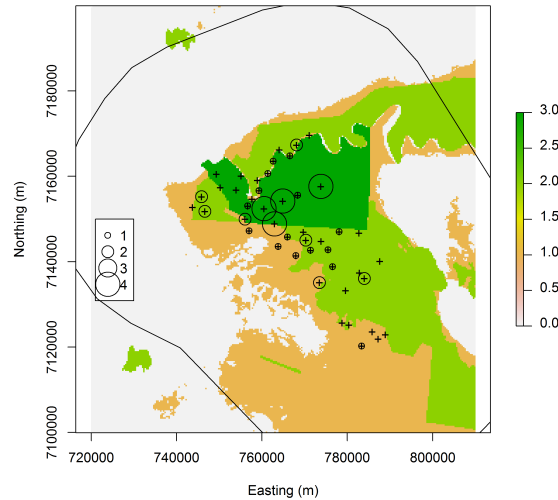


Figure 11.4. Jaguar detections at 46 camera trap stations. The three levels of protection status are no protection (beige), some protection (light green), and national park (dark green). Non-habitat is shown in gray and represents large soybean monocultures.

trapping study of jaguars (*Panthera onca*) in Argentina, along the borders with Brazil and Paraguay. The data consist of 46 camera stations, each consisting of a camera on either side of a road or trail. Forty-five detections of 16 jaguars were made over a 95 day sampling period. The mean number of sampling days at each camera station was 48.2. Eight females and 8 males were detected.

Estimating density is a central objective of this study because ultimately, an estimate of the total population size for the “green corridor” is needed. A second, and related, objective was to assess the influence of poaching on jaguar density. Although jaguars themselves are occasionally killed by poachers, the bigger influence is the effect of poaching on prey species such as peccaries (*Pecari tajacu*, *Tayassu pecari* ASK AGUSTIN WHICH ARE MORE COMMON). To protect jaguars and related species, protected areas have been established (DETAILS) and three levels of protection are recognized as depicted in Fig. 11.3.3.

MENTION SEX-SPECIFIC SIGMAS

MENTION THE RIVER

To assess the influence of protection on jaguar density, we treated protection status as an ordinal variable with 3 levels: no protection, some protection???, national park. Clearly these are ordered, and our hypothesis is that density should

Table 11.3. "Jaguar density estimates and associated parameters"

	Mean	SD	2.5%	50%	97.5%
sigmaF	7340.1547	1998.5926	4736.1803	6933.9760	12440.3079
sigmaM	8154.5222	1581.5224	5801.0509	7916.7572	11890.7957
rho	0.5163	0.1170	0.2874	0.5172	0.7388
lam0	0.0071	0.0024	0.0033	0.0067	0.0126
psi	0.3145	0.0699	0.1908	0.3100	0.4638
beta	4.3238	1.5115	2.4148	4.0295	8.0894
N	20.3828	2.8671	16.0000	20.0000	27.0000
N1	0.2759	0.6360	0.0000	0.0000	2.0000
N2	2.6285	1.8589	0.0000	2.0000	7.0000
N3	17.4784	2.8458	12.0000	17.0000	24.0000

increase with the level of protection. Thus, β in this example is a single “slope” parameter describing the degree to which protection status affects jaguar density. `°` code to fit the model are available in `scrbook`. Parameter estimates are shown in Table XXX. Our results indicate that efforts to protect jaguars by reducing poaching are working. Density was X times higher in the national park than in the unprotected areas. Fig. 11.3.3 shows the estimated density surface.

We note that there is room for improvement in our analysis. The political boundaries use to demarcate protected areas are not as concrete as we might like. In reality poaching pressure is likely to be higher near remote park boundaries than in well-guarded park interiors. One option for addressing this would be to use a continuous measure of poaching pressure such as distance from the nearest town, or some other accessibility metric. It would also be interesting to model density seperately for each sex. Many of the detections outside of the park were of males, and thus it is possible that the sexes use habitat differently.

11.4 SUMMARY

When state-space covariates are available, we can model density by replacing the uniform prior on the activity centers with a prior based on a normalized log-linear function of covariates. This distribution has been widely used in ecology to model point processes as well as resource selection probability functions. In our SCR context, use of this new prior results in a model for the inhomogeneous point process describing the location of activity centers, which can be used to test hypotheses about covariates affecting density. In rare cases, these covariates are truly continuous in the sense that they are defined as a function of space. More often, covariates are represented on rasters, which simplifies the analysis. Fitting these models can be accomplished using **BUGS**, **secr**, or the custom **R** code presented in this chapter and found in the package `scrbook`.

Note that density cannot be modeled using traditional CR methods.

All the examples in this section included a single state-space covariate, but this

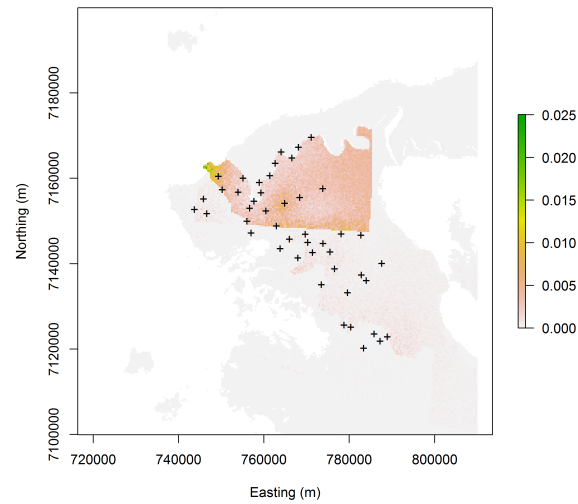


Figure 11.5. Estimated density surface for the jaguar dataset

6366 was for simplicity only. Including multiple covariates poses no additional challenges.
6367 Likewise, additional model structure such sex-specific encounter rate parameters or
6368 behavioral responses can be accommodated.

6369
6370
6371

12

OPEN MODELS

SPATIAL CAPTURE-RECAPTURE FOR UNMARKED POPULATIONS

Traditional capture-recapture models share the fundamental assumption that each individual in a population can be uniquely identified when captured. This can often be accomplished by marking individuals with color bands, ear tags, or some other artificial mark that can be subsequently read in the field. For other species, such as tigers or marbled salamanders, individuals can be easily identified using only their natural markings. Many species, however, do not possess adequate natural markings and are difficult to capture, making it impractical to use standard capture-recapture techniques.

Estimating density when individuals are unmarked can be accomplished using a variety of alternatives to capture-recapture. However, many of these methods have important limitations that warrant the exploration of new approaches. In this chapter we highlight the work of Chandler and Royle (2012) who demonstrated that the “individual recognition” assumption of capture-recapture models is not a requirement of spatial capture-recapture models. They showed that spatial correlation alone is sufficient for making inference about animal distribution and density. That is, if we simply have spatially-correlated count data at a collection of survey points, we can estimate density even if all individuals are unmarked, assuming that the underlying SCR model is valid. This “spatial count model” is virtually identical to other SCR models except that the encounter histories are not directly observed.

The ability to fit SCR models to data from unmarked populations has important consequences. For one, it means that SCR models can be applied to data collected using methods like point counts in which observers record simple counts of animals at an array of survey points. This development also has important implications for traditional SCR studies because many SCR datasets include some individuals that cannot be identified due to poor photo quality or indistinguishable natural

markings.

It's also interesting to note that by disregarding individual identity, we wind up with a model that closely resembles another large class of spatial models, known as convolution models (Wolpert and Ickstadt, 1998; Higdon, 1998). These models have been used for purposes as varied as describing oceanic surface temperatures and correlation in tree locations within managed forests. The Chandler-Royle model offers an improvement in some respects over existing convolution models because it does not require arbitrary decisions about the location and number of "support points". We will clarify this later in the chapter, and briefly mention how this model can be used outside of SCR contexts for general purpose spatial modeling of correlated count data.

13.1 EXISTING MODELS FOR INFERENCE ABOUT DENSITY IN UNMARKED POPULATIONS

When capture-recapture methods are not a viable option, researchers often collect simple count data or even detection/non-detection data to estimate population parameters. These data are often analyzed using Poisson regression or logistic regression, perhaps with random effects. When detection is imperfect, as it almost always is, these methods cannot be used to obtain unbiased estimates of population size or occurrence probability. Even when these data are used as an index of abundance or occurrence, standard models may yield unreliable results when covariates affect both the state variable and detection probability. A classic example is the finding by Bibby and Buckland (1987) who reported that the probability of detecting songbirds in restocked conifer plantations decreased with vegetation height; whereas population density was positively related to vegetation height. This intuitive and common phenomenon has led to the development of a vast number of models to estimate population size and detection probability when individuals are unmarked. A review of these models is beyond the scope of this chapter, but we mention a few deficiencies of existing methods that warrant the exploration of alternatives for robust inference when standard capture-recapture methods do not apply.

Distance sampling (Buckland, 2001), which we briefly introduced in chapter XXXX, is perhaps the most widely used method for estimating population density when individuals are unmarked and detection probability is less than one. This class of methods is known to work impeccably when estimating the number of stakes in a field or the number of duck nests in a wetland. Distance sampling can also work very well in more interesting situations, and is an extremely powerful method when the assumptions can be met. However, the assumptions that distance data can be recorded without error and that animals are distributed randomly with respect to the transect can be easily violated by common processes such as animal movement and measurement error may result in substantial bias. Although numerous methods have been proposed to relax some of these assumptions Royle et al. (2004); ?;

Johnson (2010); Chandler et al. (2011), another issue is that distance sampling is simply not practical in many settings. For example, many species are so rare and elusive that they can only be reliably surveyed using methods such as camera traps.

Other common sampling methods used to estimate density when individuals are unmarked include double-observer sampling, removal sampling, and repeated counts. Models for each of these data structures have been developed (Royle, 2004; ?). To obtain reliable density estimates using these methods, the area surveyed must be well defined and closed with respect to movement and demographic processes. Given a short enough sampling interval, such as a 5-min point-count, the closure assumption may be reasonable. However, short sampling intervals can reduce the number of detections, so observers generally visit each survey point multiple times during a season. But then you can expect animal movement to invalidate the closure assumption, so there is this kind of endless morass of shit to worry about and you might as well jump off a cliff—especially if you start thinking about false positives!

MENTION SIMILARITY BETWEEN THE TEMPORARY EMIGRATION PROBLEM FACED HERE AND THE ONE THAT MOTIVATES SCR

Deep breath. Anyhow, we mention these issues not to suggest that existing models do not have value—indeed we believe that they can be used to obtain reliable density estimates in many situations—rather our aim is to highlight the need for alternative methods when the assumptions of existing methods cannot be met. Additionally, the model we develop in this chapter serves as the foundation for a broad class of SCR models in which all or some of the individuals cannot be uniquely identified.

13.2 SPATIAL CORRELATION AS INFORMATION

All of the previous methods require some sort of auxiliary information to separately model the state and observation processes. That is, we need multiple observers or distance data or repeated visits to ensure that model parameters are identifiable¹. The same is true for Chandler-Royle model, but the auxiliary information comes in the form of spatial correlation (Chandler and Royle, 2012).

It is natural to be suspicious of the claim that spatial correlation is a good thing. Indeed, elaborate methods have been devised to deal with spatial correlation as a nuisance parameter (?), and ecologists have been admonished for failing to obtain “real” replicates uncontaminated by spatial correlation (Hurlbert, 1984). The following heuristic may be helpful.

Imagine a 10×10 grid of camera traps and a single individual exposed to capture whose home range center lies in the center of the trapping grid. If the individual has a small home range size relative to the extent of the trapping grid, we can imagine what the spatial correlation structure of the encounters might look like.

¹Or we can make very strong model assumptions and get away without any auxiliary data (?)

If the animal's movement is symmetric around the activity center then the number of times the individual is detected at each trap (the trap counts) is a function of the distance between the home range center and the trap, *i.e.* traps with the same distance from the activity center will yield counts that are more highly correlated with one another than traps located at different distances from the activity center. Thus, the correlation in counts tells us something about the location of the activity center. That is, correlation carries information about distribution. What about density?

Imagine now that there are two activity centers located in our trapping grid. Using trap counts alone, can our model tell us both where the activity centers are and how many exist in the population exposed to capture? The answer is yes, at least under certain circumstances. EXPLAIN.

Thus, intuition alone suggests that we might be able to estimate the number and location of activity centers using simple count data that are spatially-correlated. The results of Chandler and Royle (2012) provide additional support for this assertion.

13.3 DATA REQUIREMENTS AND SURVEY DESIGNS

One of the important benefits of Chandler-Royle model is that it can be applied to data collected using an enormous variety of survey methods. Whereas traditional SCR models require spatially-referenced encounter histories, this model requires simple count data. Once again, suppose that we have J traps surveyed on K time periods during which no births or deaths occur. Suppose also that an individual can be encountered at multiple traps during a single time period, say one night during a camera-trapping study. If individuals can also be encountered multiple times at a single trap during a night, then the encounter history data could be regarded as

$$z_{ijk} \sim \text{Poisson}(\lambda_{ij}). \quad (13.3.1)$$

if the animals were marked. Here λ_{ij} is the encounter rate for individual i at trap j . A common form of this parameter is

$$\lambda_{ij} = \lambda_0 \exp(\|\mathbf{x}_j - \mathbf{s}_i\|/2\sigma^2)$$

where λ_0 is the baseline encounter rate and σ is the scale parameter describing the distance-related decay in encounter rate.

When individuals cannot be uniquely identified, the z_{ijk} cannot be directly observed. So just when you thought we ran out of things to treat as latent variables, we are now going to regard even the encounter histories as latent. The data are now just a reduced-information summary of the latent encounter histories. That is, they are the sample- and trap-specific totals, aggregated over all individuals:

$$n_{jk} = \sum_{i=1}^N z_{ijk}.$$

This data structure, a matrix of counts made at a collection of sampling locations on one or more occasions is extremely common in ecology. Note also that the reason why we can get by with a single occasion of data ($J \equiv 1$) is that, under the Poisson model,

$$n_{jk} \sim \text{Poisson}(\Lambda_j) \quad (13.3.2)$$

where

$$\Lambda_j = \lambda_0 \sum_i k_{ij},$$

and because Λ_r does not depend on t , we can aggregate the replicated counts, defining $n_{j\cdot} = \sum_k n_{jk}$ and then

$$n_{j\cdot} \sim \text{Poisson}(K\Lambda_j)$$

As such, K and λ_0 serve equivalent roles as affecting baseline encounter rate. This has been noted elsewhere (?).

We imagine that other observation models might be possible (see Discussion) although we focus here on the Poisson encounter model because it has considerable relevance to animal surveys, and has additional methodological context related to point process models which we address in the Discussion.

13.4 ENCOUNTER HISTORIES AS LATENT VARIABLES

State model is the same as other SCR models.

This formulation of the model in terms of the aggregate count simplifies computations as the latent variables z_{irt} do not need to be updated in the MCMC estimation scheme (see below). However, retaining z_{irt} in the formulation of the model is important if some individuals are uniquely marked, in which case modifying the MCMC algorithm (see below) to include both types of data is trivial. This is because uniquely identifiable individuals produce observations of some of the z_{irt} variables.

13.5 ESTIMATION BY MCMC

We adopt a Bayesian framework for inference allowing estimation of N while retaining the formulation of the model that is conditional on the latent activity centers \mathbf{s}_i . Specifically, we employ Markov chain Monte Carlo (MCMC) to simulate posterior distributions of the parameters. However, the fact that N is unknown presents a technical challenge because the size of the parameter space can change with each MC iteration. To resolve this, we adopt the formulation of data augmentation in Royle et al. (2007) who used a specific prior construction for N in terms of individual level Bernoulli trials. In particular, we assume $N \sim \text{Unif}(0, M)$ for some large integer M . We construct this prior by assuming $N|M, \phi \sim \text{Bin}(M, \phi)$ and $\phi \sim \text{DUnif}(0, 1)$ which implies, marginally, that N has the requisite $\text{DUnif}(0, M)$

distribution. However the hierarchical formulation of the prior suggests an implementation in which we introduce a set of latent indicator variables $w_i \sim \text{Bern}(\phi)$ and, furthermore, the model implies that z_{irt} are obtained from the specified distribution (Eq. 13.3.1) if $w_i = 1$, or if $w_i = 0$, $z_{irt} = 0$ with probability 1. In effect, extending the model in this way induces a reparameterization for the latent counts that is a zero-inflated version of the original conditional-on- N model. Specifically, the model under data augmentation becomes

$$\begin{aligned} z_{irt} | w_i &\sim \text{Poisson}(\lambda_{ir} w_i) \\ w_i &\sim \text{Bern}(\phi) \end{aligned}$$

Under this formulation $N = \sum_{i=1}^M w_i$, and population density is simply $D = N/A(\mathcal{S})$ where $A(\mathcal{S})$ is the area of the point process state-space \mathcal{S} .

We developed two distinct MCMC implementations for this model. In the first, we devised an algorithm for the model conditional on the latent variables z_{irt} . This formulation is useful for problems in which one or more individual identities are available, in which case the z_{irt} are observable for those individuals. The unobserved z_{irt} are easily updated using their full-conditional distribution which is multinomial with sample size n_{rt} . The remaining parameters are updated using Metropolis-Hastings steps. In the second formulation of the algorithm we applied the Metropolis-Hastings algorithm to the model *unconditional* on the z_{irt} variables. In that case, the marginal distribution for n_{rt} is precisely Eq. 13.3.2. This algorithm is slightly more convenient because it avoids having to update the z_{irt} variables of which there are many.

13.6 NORTHERN PARULA EXAMPLE

To apply our model to data collected in the field, we designed a point count study of the northern parula (*Parula americana*), a Neotropical-Nearctic migratory passerine. This species defends well-defined territories during the breeding season (?), and thus our modeling effort was focused on estimating the number and location of territory centers. Points were located on a 50-m grid to ensure spatial correlation. This small grid spacing contrasts with the conventional practice of spacing points by > 200 m to obtain *i.i.d.* counts. Figure 13.1 depicts the spatially-correlated counts ($n_{r.}$) from the 105 point count locations surveyed three times each during June 2006 at the Patuxent Wildlife Research Center in Laurel Maryland, USA. A total of 226 detections were made with a maximum count of 4 during a single survey. At 38 points, no warblers were detected. All but one of the detections were of singing males, and this one observation was not included in the analysis.

In our analysis of the parula data, we defined the point process state-space by buffering the grid of point count locations by 250 m and used $M = 300$. We simulated posterior distributions using three Markov chains, each consisting of 300000 iterations after discarding the initial 10000 draws. Convergence was satisfactory, as indicated by an \hat{R} statistic of < 1.02 (Gelman and Rubin, 1992).




Figure 13.1. Spatially-correlated counts of northern parula on a 50-m grid. The size of the circle represents the total number of detections at each point.

One benefit of a Bayesian analysis is that it can accommodate prior information on the home range size and encounter rate parameters, which are readily available for many species. To illustrate, we analyzed the parula data using two sets of priors. In the first set, all priors were improper, customary non-informative priors (see Table 13.1). Uniform priors were also used in the second set, with the exception of an informative prior for the scale parameter $\sigma \sim \text{Gamma}(13, 10)$. We arrived at this prior using the methods described by Royle et al. (2011b) and published information on the warbler's home range size and detection probability (Simons et al., 2009). More details on this derivation are found in. We briefly note here that this prior includes the biologically-plausible range of values from σ suggested by the published literature.

The posterior distribution for N was highly skewed with a long right tail resulting in a wide 95% credible interval (Table 13.1). Nonetheless, the interval for density, D , includes estimates reported from more intensive field studies (?). This was true when considering both sets of priors, although posterior precision was higher under the informative set of priors. Specifically, the use of prior information reduced posterior density at high, biologically implausible, values of σ , and hence decreased the posterior mass for low values of N (Fig. 13.2).

In addition to estimating density, our model can be used to produce density surface maps, which are often used in applied ecological research to direct management efforts and develop hypotheses regarding the factors influencing abundance. Density surface maps can be produced by discretized the state-space and tallying the number of activity centers occurring in each pixel during each MCMC iteration. Parula density was highest near the northeastern corner of the study plot, which

Table 13.1. Posterior summary statistics for spatial Poisson-count model applied to the northern parula data. Two sets of priors were considered. $M = 300$ was used in both cases. Parulas/ha, D , is a derived parameter.

Par	Prior	Mean	SD	Mode	q0.025	q0.50	q0.975
σ	$U(0, \infty)$	2.154	1.222	1.230	0.896	1.665	5.170
λ_0	$U(0, \infty)$	0.284	0.149	0.212	0.084	0.256	0.665
N	$U(0, M)$	40.953	38.072	4.000	3.000	31.000	143.000
D	–	0.427	0.397	0.0417	0.0313	0.323	1.490
σ	$G(13, 10)$	1.301	0.258	1.230	0.889	1.266	1.908
λ_0	$U(0, \infty)$	0.298	0.132	0.240	0.098	0.279	0.603
N	$U(0, M)$	59.321	36.489	36.000	18.000	50.000	157.000
D	–	0.618	0.380	0.375	0.188	0.521	1.635

may correspond to important habitat features such as suitable nest site locations (Fig. 13.3). We anticipate future model extensions to directly model the point process intensity using habitat covariates.

13.7 ON (IM)PRECISION

13.8 HOW MUCH CORRELATION IS ENOUGH?

13.9 MUTANTS

13.9.1 Other observation models

13.9.2 Linear designs

13.10 SUMMARY

In this paper, we confronted one of the most difficult challenges faced in wildlife sampling — estimation of density in the absence of data to distinguish among individuals. To do so, we developed a novel class of spatially-explicit models that applies to spatially organized counts, where the count locations or devices are located sufficiently close together so that individuals are exposed to encounter at multiple devices. This design yields correlation in the observed counts, and this correlation proves to be informative about encounter probability parameters and hence density. We note that sample locations in count-based studies are typically *not* organized close together in space because conventional wisdom and standard practice dictate that independence of sample units is necessary (Hurlbert, 1984). Our model suggests that in some cases it might be advantageous to deviate from the conventional wisdom if one is interested in direct inference about density. Of course, this is also known in the application of standard spatial capture-recapture models (Borchers and Efford, 2008) where individual identity is preserved across trap encounters, but it is seldom, if ever, considered in the design of more traditional count surveys.




Figure 13.2. Effects of $\sigma \sim \text{Gamma}(13,10)$ prior on the posterior distributions from the northern parula model. Posteriors from model with uniform priors are shown in black, and posteriors from the informative prior model are shown in gray. The prior itself is shown as dotted line in the upper panel.




Figure 13.3. Estimated density surface of northern parula activity centers. The grid of point count locations with count totals is superimposed. See Fig. 1 for additional details.

Our model has broad relevance to an incredible number of animal sampling problems. Our motivating problem involved bird point counts where individual identity is typically not available. The model also applies to other standard methods used to sample unmarked populations, such as camera traps or even methods that yield sign (*e.g.* scat, track) counts indexed by space. However, results of our simulation study reveal some important limitations of the basic estimator applied to situations in which none of the individuals can be uniquely identified. In particular, posterior distributions are highly skewed in typical small to moderate sample size situations and posterior precision is low.

Several modifications of the model can lead to improved performance of the estimator. Our simulation results demonstrate that marking a subset of individuals can yield substantial increases in posterior precision. Marking a subset of individuals is commonplace in animal studies such as when a small number of individuals are radio-collared in conjunction with a count-based survey (Bartmann et al., 1987). In many other situations a subset of individuals can be identified by natural marks alone, and thus our model could be applied to data from camera-trapping studies of species such as mountain lions, deer, coyotes for which traditional SCR methods are not effective (Kelly et al., 2008). Thus, the ability to study partially-marked populations adds flexibility to existing SCR methods, and also creates new opportunities for designing efficient SCR studies since the costs of marking all individuals in a population can be prohibitive.

We note the existence of traditional approaches to combining data on marked and unmarked animals based on either the Lincoln-Peterson estimator or so-called “mark-resight” methods. (Bartmann et al., 1987; ?; ?). In their simplest form, mark-resight methods involve fitting standard closed-population mark-recapture models to the data on marked individuals, and the resultant estimate of detection probability (\hat{p}) is used to estimate population size as $\hat{N} = m + u/\hat{p}$ where m and u are the number of marked and unmarked individual, respectively. In this case, the unmarked individuals provide no information about the encounter rate parameters, and thus mark-resight methods cannot be used unless a large sample of marked individuals is available. This contrasts with our approach which can be used even when all individuals are unmarked.

In some cases, such as in point counts of birds, it may not be practical to mark individuals. An alternative to increasing posterior precision is to utilize prior information on home range size. Indeed, extensive information on home range size has been compiled for many species in diverse habitats (*e.g.*, DeGraaf and Yamasaki, 2001). It is easy to embody this information in a prior distribution as we demonstrated for the parula data.

An additional design extension that could increase precision is to use multiple sampling methods, in which one method generates encounter frequencies and the other method generates individuality. For example, camera traps are now commonly used with surveys for sign (scat or tracks), or hair snares for sampling bear populations. These distinct methods would have different basal detection rates

but share an underlying spatial model describing the organization of individuals in space. Our models show promise for using these disparate data types efficiently for estimating density.

13.10.1 Alternative Observation Models

Several aspects of our “spatial N -mixture model” can be modified to accommodate alternative sampling designs or parametric distributions. We considered situations where an individual can be detected more than once at a trap during a single occasion, but under some designs this is not possible. When collecting DNA samples, for instance, an individual can often be detected at most once during an occasion, because multiple samples of biological material cannot be attributed to distinct episodes. Therefore, rather than $z_{irt} \sim \text{Poisson}(\lambda_{ir})$ we have $z_{irt} \sim \text{Bernoulli}(p_{ir})$ where, for example, $p_{ir} = p_0 \exp(-d_{ir}^2/(2\sigma^2))$, and p_0 is the probability of detecting an individual whose home range is centered on trap r . This Bernoulli model is a focus of ongoing investigations.

Both the Poisson and the Bernoulli models produce count observations when aggregated over individuals to form trap-specific totals; however, ecologists often collect so-called “detection/non-detection” data because it can be easier to determine if “at least one” individual was present rather than enumerating all individuals in a location. In this case, the underlying z_{irt} array is the same as the above cases, but we observe $y_{rt} = I(\sum_{i=1}^N z_{irt} > 0)$ where I is the indicator function. This “Poisson-binary model” is a spatially explicit extension of the model of Royle and Nichols (2003) in which the underlying abundance state is inferred from binary data. We have investigated this model to a limited extent but do not report on those results here.

13.10.2 Spatial point process models

Our model has some direct linkages to existing point process models. We note that the observation intensity function (i.e., corresponding to the observation locations) is a compound Gaussian kernel similar to that of the Thomas process (??, pp. 61-62). Also, the Poisson-Gamma Convolution models (Wolpert and Ickstadt, 1998) are structurally similar (see also Higdon (1998) and Best et al. (2000)). In particular, our model is such a model but with a *constant* basal encounter rate λ_0 and *unknown* number and location of “support points”, which in our case are the animal activity centers, \mathbf{s}_i . We can thus regard our model as a model for *estimating* the location and local density of support points in such models, which we believe could be useful in the application of convolution models. Best et al. (2000) devise an MCMC algorithm for the Poisson-Gamma model based on data augmentation, which is similar to the component of our algorithm for updating the z variables in the conditional-on- z formulation of the model. We emphasize that our model

is distinct from these Poisson-Gamma models in that the number *and* location of such support points are estimated.

If individuals were perfectly observable then the resulting point process of locations is clearly a standard Poisson or Binomial (fixed N) cluster process or Neyman-Scott process. If detection is uniform over space but imperfect, then the basic process is unaffected by this random thinning. Our model can therefore be viewed formally as a Poisson (or Binomial) cluster process model but one in which the thinning is non-uniform, governed by the encounter model which dictates that thinning rate increases with distance from the observation points. In addition, our inference objective is, essentially, to estimate the number of parents in the underlying Poisson cluster process, where the observations are biased by an incomplete sampling apparatus (points in space).

As a model of a thinned point process, our model has much in common with classical distance sampling models (Buckland, 2001). The main distinction is that our data structure does *not* include observed distances, although the underlying observation model is fundamentally the same as in distance sampling if there is only a single replicate sample and \mathbf{s}_i is defined as an individual's location at an instant in time. For replicate samples, our model preserves (latent) individuality across samples and traps which is not a feature of distance sampling. We note that error in measurement of distance is not a relevant consideration in our model, and we explicitly do not require the standard distance sampling assumption that the probability of detection is 1 if an individual occurs at the survey point. More importantly, distance sampling models cannot be applied to data from many of the sampling designs for which our model is relevant. For example, many rare and endangered species can only be effectively surveyed using methods such as hair snares and camera traps that do not produce distance data (O'Connell et al., 2010).

13.11 CONCLUSION

Concerns about “statistical independence” have prompted ecologists to design count-based studies such that observed random variables can be regarded as *i.i.d.* outcomes (Hurlbert, 1984). Interestingly, this often proves impossible in practice, and elaborate methods have been devised to model spatial dependence as a nuisance parameter. Our paper presents a modeling framework that directly confronts this view by demonstrating that spatial correlation carries information about the locations of individuals, which can be used to estimate density even when individuals are unmarked and distance-related heterogeneity exists in encounter probability.

BIBLIOGRAPHY

- Alho, J. (1990), “Logistic regression in capture-recapture models,” *Biometrics*, 623–635.
- Arnason (1973), “Missing,” *Missing*, Missing, Missing.
- (1974), “Missing,” *Missing*, Missing, Missing.
- Baddeley, A. and Turner, R. (2005), “Spatstat: an R package for analyzing spatial point patterns,” *Journal of Statistical Software*, 12, 1–42, ISSN 1548-7660.
- Bales, S. L., Hellgren, E. C., Leslie Jr., D. M., and Hemphill Jr., J. (2005), “Dynamics of recolonizing populations of black bears in the Ouachita Mountains of Oklahoma,” *Wildlife Society Bulletin*, 1342–1351.
- Bartmann, R. M., White, G. C., Carpenter, L. H., and Garrott, R. A. (1987), “Aerial Mark-Recapture Estimates of Confined Mule Deer in Pinyon-Juniper Woodland,” *The Journal of Wildlife Management*, 51, 41–46.
- Berger, J. O., Liseo, B., and Wolpert, R. L. (1999), “Integrated likelihood methods for eliminating nuisance parameters,” *Statistical Science*, 1–22.
- Best, N. G., Ickstadt, K., and Wolpert, R. L. (2000), “Spatial Poisson Regression for Health and Exposure Data Measured at Disparate Resolutions,” *Journal of the American Statistical Association*, 95, 1076.
- Bibby, C. and Buckland, S. (1987), “Bias of bird census results due to detectability varying with habitat,” *Acta Ecologica*, 8, 103–112.
- Bivand, R. and Rundel, C. (2011), *rgeos: Interface to Geometry Engine - Open Source (GEOS)*, r package version 0.1-8.
- Borchers, D. and Efford, M. (2008), “Spatially explicit maximum likelihood methods for capture-recapture studies,” *Biometrics*, 64, 377–385.
- Borchers, D. L. (missing), “missing,” *Missing*, missing.
- Borchers, D. L., Buckland, S. T., and Zucchini, W. (2002), *Estimating animal abundance: closed populations*, vol. 13, Springer Verlag.
- Boulanger, J. and McLellan, B. (2001), “Closure violation in DNA-based mark-recapture estimation of grizzly bear populations,” *Canadian Journal of Zoology*, 79, 642–651.
- Brooks, S. P., Catchpole, E. A., and Morgan, B. J. T. (2000), “Bayesian Animal Survival Estimation,” *Statistical Science*, 15, 357–376.
- Buckland, S. T. (2001), *Introduction to distance sampling: estimating abundance of biological populations*, Oxford, UK: Oxford University Press.

- 6774 Burnham, K. P. and Anderson, D. R. (2002), *Model selection and multimodel in-*
6775 *ference: a practical information-theoretic approach*, Springer Verlag.
- 6776 Burnham, K. P. and Overton, W. S. (1978), “Estimation of the size of a closed
6777 population when capture probabilities vary among animals,” *Biometrika*, 65,
6778 625.
- 6779 Casella, G. and George, E. I. (1992), “Explaining the Gibbs sampler,” *American*
6780 *Statistician*, 46, 167–174.
- 6781 Chandler, R. and Royle, J. (2012), “Spatially-explicit models for inference about
6782 density in unmarked populations,” *Biometrics (in review)*.
- 6783 Chandler, R., Royle, J., and King, D. (2011), “Inference about density and tempo-
6784 rary emigration in unmarked populations,” *Ecology*, 92, 1429–1435.
- 6785 Converse, S. J. and Royle, J. A. (2010), “Missing,” *Missing*, Missing.
- 6786 Coull, B. A. and Agresti, A. (1999), “The Use of Mixed Logit Models to Reflect
6787 Heterogeneity in Capture-Recapture Studies,” *Biometrics*, 55, 294–301.
- 6788 Dawson, D. and Efford, M. (2009), “Bird population density estimated from acous-
6789 tic signals,” *Journal of Applied Ecology*, 46, 1201–1209.
- 6790 DeGraaf, R. M. and Yamasaki, M. (2001), *New England wildlife: habitat, natural*
6791 *history, and distribution*, University Press of New England.
- 6792 Dice, L. R. (1938), “Some census methods for mammals,” *Journal of Wildlife Man-*
6793 *agement*, 2, 119–130.
- 6794 Dorazio, R. and Royle, J. (2003), “Mixture models for estimating the size of a closed
6795 population when capture rates vary among individuals,” *Biometrics*, 351–364.
- 6796 Dorazio, R. M. (2007), “On the choice of statistical models for estimating occurrence
6797 and extinction from animal surveys,” *Ecology*, 88, 2773–2782.
- 6798 Durbin and Elston (2012), “Missing,” *Missing*, Missing.
- 6799 Efford, M. (2004), “Density estimation in live-trapping studies,” *Oikos*, 106, 598–
6800 610.
- 6801 — (2011), “secre-spatially explicit capture-recapture in R,” .
- 6802 Efford, M., Dawson, D., and Robbins, C. (2004), “DENSITY: software for analysing
6803 capture-recapture data from passive detector arrays,” *Animal Biodiversity and*
6804 *Conservation*, 217–228.
- 6805 Fienberg, S. E., Johnson, M. S., and Junker, B. W. (1999), “Classical multilevel and
6806 Bayesian approaches to population size estimation using multiple lists,” *Journal*
6807 *of the Royal Statistical Society of London A*, 163, 383–405.
- 6808 Gardner, B. (2009), “missing,” *Missing*, missing.
- 6809 Gardner, B., Royle, J., Wegan, M., Rainbolt, R., and Curtis, P. (2010), “Estimating
6810 black bear density using DNA data from hair snares,” *The Journal of Wildlife*
6811 *Management*, 74, 318–325.
- 6812 Gelfand, A. and Smith, A. (1990), “Sampling-based approaches to calculating
6813 marginal densities,” *Journal of the American statistical association*, 85, 398–409.
- 6814 Gelman, A., Carlin, J. B., Stern, H. S., and Rubin, D. B. (2004), *Bayesian data*
6815 *analysis, second edition.*, Boca Raton, Florida, USA: CRC/Chapman & Hall.
- 6816 Gelman, A., Meng, X. L., and Stern, H. (1996), “Posterior predictive assessment

- of model fitness via realized discrepancies,” *Statistica Sinica*, 6, 733–759.
- Gelman, A. and Rubin, D. B. (1992), “Inference from iterative simulation using multiple sequences,” *Statistical Science*, 7, 457–511.
- Geman, S. and Geman, D. (1984), “Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PAMI-6, 721–741.
- Genz, A. S., Meyer, M. R., Lumley, T., and Maechler, M. (2007), “The adapt Package. R package version 1.0-4,” .
- Gilks, W. and Wild, P. (1992), “Adaptive rejection sampling for Gibbs sampling,” *Applied Statistics*, 41, 337–348.
- Gilks, W. R., Thomas, A., and Spiegelhalter, D. J. (1994), “A Language and Program for Complex Bayesian Modelling,” *Journal of the Royal Statistical Society. Series D (The Statistician)*, 43, 169–177, ArticleType: primary_article / Issue Title: Special Issue: Conference on Practical Bayesian Statistics, 1992 (3) / Full publication date: 1994 / Copyright 1994 Royal Statistical Society.
- Gopalaswamy (2012), “Missing,” *Missing*, missing.
- Gopalaswamy, A. M., Royle, A. J., Hines, J., Singh, P., Jathanna, D., Kumar, N. S., and Karanth, K. U. (2011), *A Program to Estimate Animal Abundance and Density using Spatially-Explicit Capture-Recapture*, r package version 1.0.4.
- Hahn, T., Bouvier, A., and Kieu, K. (2011), “Package ‘R2Cuba’ R package version 1.0-6,” .
- Hastings, W. (1970), “Monte Carlo sampling methods using Markov chains and their applications,” *Biometrika*, 57, 97–109.
- Hawkins, C. and Racey, P. (2005), “Low population density of a tropical forest carnivore, *Cryptoprocta ferox*: implications for protected area management,” *Oryx*, 39, 35–43.
- Hestbeck (1991), “Missing,” *Missing*, Missing, Missing.
- Higdon, D. (1998), “A process-convolution approach to modelling temperatures in the North Atlantic Ocean,” *Environmental and Ecological Statistics*, 5, 173–190.
- (2002), “Space and Space-Time Modeling using Process Convolutions,” in *Quantitative methods for current environmental issues*, eds. Anderson, C., Barnett, V., Chatwin, P., and El-Shaarawi, A., Springer Verlag, p. 37.
- Huggins, R. M. (1989), “On the statistical analysis of capture experiments,” *Biometrika*, 76, 133.
- Hurlbert, S. (1984), “Pseudoreplication and the design of ecological field experiments,” *Ecological monographs*, 54, 187–211.
- Illian (2008a), “Missing,” *Missing*, Missing.
- Illian, J. (2008b), *Statistical analysis and modelling of spatial point patterns*, Wiley-Interscience.
- Ivan, J. (2012), “Density, demography, and seasonal movements of snowshoe hares in central Colorado,” Ph.D. thesis, COLORADO STATE UNIVERSITY.
- Jackson, R., Roe, J., Wangchuk, R., and Hunter, D. (2006), “Estimating Snow Leopard Population Abundance Using Photography and Capture-Recapture

- Techniques,” *Wildlife Society Bulletin*, 34, 772–781.
- Johnson (2010), “A Model-Based Approach for Making Ecological Inference from Distance Sampling Data,” *Biometrics*, 66, 310318.
- Johnson, D. (1999), “The insignificance of statistical significance testing,” *The journal of wildlife management*, 763–772.
- Karanth, K. U. (1995), “Estimating tiger *Panthera tigris* populations from camera-trap data using capture–recapture models,” *Biological Conservation*, 71, 333–338.
- Kelly, M., Noss, A., Di Bitetti, M., Maffei, L., Arispe, R., Paviolo, A., De Angelo, C., and Di Blanco, Y. (2008), “Estimating puma densities from camera trapping across three study sites: Bolivia, Argentina, and Belize,” *Journal of Mammalogy*, 89, 408418.
- Kendall (1997), “Missing,” *Missing*, missing.
- Kéry, M. (2010), *Introduction to WinBUGS for Ecologists: Bayesian Approach to Regression, ANOVA, Mixed Models and Related Analyses*, Academic Press.
- Kéry, M., Gardner, B., Stoeckle, T., Weber, D., and Royle, J. A. (2010), “Use of Spatial Capture-Recapture Modeling and DNA Data to Estimate Densities of Elusive Animals,” *Conservation Biology*, 25, 356–364.
- Kéry, M., Royle, J., and Schmid, H. (2005), “Modeling avian abundance from replicated counts using binomial mixture models,” *Ecological Applications*, 15, 1450–1461.
- Kery, M. and Schaub, M. (2011), *Bayesian Population Analysis Using WinBugs*, Academic Press.
- King, R. (2009), “Missing,” *missing*, Missing.
- (missing), “Missing,” *missing*, Missing.
- Kumar (missing), “Unpublished data,” .
- Kuo, L. and Mallick, B. (1998), “Variable selection for regression models,” *Sankhyā*, 60, 65–81.
- Laird, N. M. and Ware, J. H. (1982), “Random-effects models for longitudinal data,” *Biometrics*, 963–974.
- Langtimm (2010), “Missing,” *Missing*, missing.
- Le Cam, L. (1990), “Maximum likelihood: an introduction,” *International Statistical Review/Revue Internationale de Statistique*, 153–171.
- Lewin-Koh, N. J., Bivand, R., contributions by Edzer J. Pebesma, Archer, E., Baddeley, A., Bibiko, H.-J., Dray, S., Forrest, D., Friendly, M., Giraudoux, P., Golicher, D., Rubio, V. G., Hausmann, P., Hufthammer, K. O., Jagger, T., Luque, S. P., MacQueen, D., Niccolai, A., Short, T., Stabler, B., and Turner, R. (2011), *maptools: Tools for reading and handling spatial objects*, r package version 0.8-10.
- Link, W. A. (2003), “Missing,” *missing*, missing.
- Link, W. A. and Barker, R. J. (2009), *Bayesian Inference: With Ecological Applications*, London, UK: Academic Press.
- Liu and Wu (1999), “Parameter expansion for data augmentation,” *J Am Stat Assoc*, 94, 1264–1274.

- MacEachern, S. and Berliner, L. (1994), "Subsampling the Gibbs sampler," *American Statistician*, 188–190.
- MacKenzie, D. I., Nichols, J. D., Lachman, G. B., Droege, S., Royle, J. A., and Langtimm, C. A. (2002), "Estimating site occupancy rates when detection probabilities are less than one," *Ecology*, 83, 2248–2255.
- Mackenzie, D. I. and Royle, J. (2005), "Designing occupancy studies: general advice and allocating survey effort," *Journal of Applied Ecology*, 42, 1105–1114.
- Magoun, A. J., Long, C. D., Schwartz, M. K., Pilgrim, K. L., Lowell, R. E., and Valkenburg, P. (2011), "Integrating motion-detection cameras and hair snags for wolverine identification," *The Journal of Wildlife Management*, 75, 731–739.
- McCarthy, M. A. (2007), *Bayesian Methods for Ecology*, Cambridge: Cambridge University Press.
- McCullagh, P. and Nelder, J. (1989), *Generalized linear models*, Chapman & Hall/CRC.
- Metropolis, N., Rosenbluth, A., Rosenbluth, M., Teller, A., Teller, E., et al. (1953), "Equation of state calculations by fast computing machines," *The journal of chemical physics*, 21, 1087–1092.
- Metropolis, N. and Ulam, S. (1949), "The Monte Carlo method," *Journal of the American Statistical Association*, 44, 335–341.
- Millar, R. (2009), "Comparison of hierarchical Bayesian models for overdispersed count data using DIC and Bayes' Factors," *Biometrics*, 65, 962–969.
- Mollet, P., Kery, M., Gardner, B., Pasinelli, G., and A, R. J. (2012), "Population size estimation for capercaillie (*Tetrao urogallus* L.) using DNA-based individual recognition and spatial capture-recapture models," *missing*, missing, missing.
- Neal, R. (2003), "Slice sampling," *Annals of Statistics*, 31, 705–741.
- Nelder, J. and Wedderburn, R. (1972), "Generalized linear models," *Journal of the Royal Statistical Society. Series A (General)*, 370–384.
- Norris III, J. L. and Pollock, K. H. (1996), "Nonparametric MLE under two closed capture-recapture models with heterogeneity," *Biometrics*, 639–649.
- O'Connell, A. F., Nichols, J. D., and Karanth, U. K. (2010), *Camera traps in animal ecology: Methods and analyses*, Springer.
- Overton, W. and Stehman, S. (1995), "The Horvitz-Thompson theorem as a unifying perspective for probability sampling: with examples from natural resource sampling," *American Statistician*, 261–268.
- Pebesma, E. and Bivand, R. (2011), *Package 'sp'*, r package version 0.9-91.
- Pledger, S. (2000), "Unified maximum likelihood estimates for closed capture-recapture models using mixtures," *Biometrics*, 434–442.
- Plummer, M. (2003), "JAGS: A program for analysis of Bayesian graphical models using Gibbs sampling," in *Proceedings of the 3rd International Workshop on Distributed Statistical Computing (DSC 2003)*. March, pp. 20–22.
- (2009), "rjags: Bayesian graphical models using mcmc. R package version 1.0.3-12," .
- (2011), *rjags: Bayesian graphical models using MCMC*, r package version 3-5.

- Plummer, M., Best, N., Cowles, K., and Vines, K. (2006), “CODA: Convergence Diagnosis and Output Analysis for MCMC,” *R News*, 6, 7–11.
- R.E., R., Royle, J., Desimone, R., Schwartz, M., Edwards, V., Pilgrim, K., and McKelvey, K. (2012), “Estimating abundance of mountain lions from unstructured spatial samples,” .
- Robert, C. P. and Casella, G. (2004), *Monte Carlo statistical methods*, New York, USA: Springer.
- (2010), *Introducing Monte Carlo Methods with R*, New York, USA: Springer.
- Roberts, G. O. and Rosenthal, J. S. (1998), “Optimal scaling of discrete approximations to Langevin diffusions,” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 60, 255–268.
- Royle, J. (2006), “Site occupancy models with heterogeneous detection probabilities,” *Biometrics*, 62, 97–102.
- (2009), “Analysis of capture-recapture models with individual covariates using data augmentation,” *Biometrics*, 65, 267–274.
- Royle, J. and Dorazio, R. (2006), “Hierarchical models of animal abundance and occurrence,” *Journal of Agricultural, Biological, and Environmental Statistics*, 11, 249–263.
- (2008), *Hierarchical modeling and inference in ecology: the analysis of data from populations, metapopulations and communities*, Academic Press.
- (2010), “Missing,” *Missing*, missing, missing.
- (2011), “Missing,” *Missing*, missing, missing.
- Royle, J., Dorazio, R., and Link, W. (2007), “Analysis of multinomial models with unknown index using data augmentation,” *Journal of Computational and Graphical Statistics*, 16, 67–85.
- Royle, J. and Dubovsky, J. (2001), “Modeling spatial variation in waterfowl band-recovery data,” *The Journal of wildlife management*, 726–737.
- Royle, J., Karanth, K., Gopalaswamy, A., and Kumar, N. (2009), “Bayesian inference in camera trapping studies for a class of spatial capture-recapture models,” *Ecology*, 90, 3233–3244.
- Royle, J. and Link, W. (2006), “Generalized site occupancy models allowing for false positive and false negative errors,” *Ecology*, 87, 835–841.
- Royle, J. and Nichols, J. (2003), “Estimating abundance from repeated presence-absence data or point counts,” *Ecology*, 84, 777–790.
- Royle, J. A. (2004), “ N -mixture models for estimating population size from spatially replicated counts,” *Biometrics*, 60, 108–115.
- (2008), “Modeling individual effects in the Cormack–Jolly–Seber model: a state-space formulation,” *Biometrics*, 64, 364–370.
- (2010), “missing,” *missing*, missing.
- Royle, J. A., Converse, S., and Link, W. (2011a), “Data augmentation for structured populations,” *unpublished*.
- Royle, J. A., Dawson, D. K., and Bates, S. (2004), “Modeling abundance effects in distance sampling,” *Ecology*, 85, 1591–1597.

- 6989 Royle, J. A., Kéry, M., and Guélat, J. (2011b), “Spatial capture-recapture models
6990 for search-encounter data,” *Methods in Ecology and Evolution*, 1–10.
- 6991 Royle, J. A., Magoun, A. J., Gardner, B., Valkenburg, P., and Lowell, R. E. (2011c),
6992 “Density estimation in a wolverine population using spatial capture–recapture
6993 models,” *The Journal of Wildlife Management*, 75, 604–611.
- 6994 Royle, J. A. and Young, K. V. (2008), “A Hierarchical Model For Spatial Capture-
6995 Recapture Data,” *Ecology*, 89, 2281–2289.
- 6996 Sanathanan, L. (1972), “Estimating the size of a multinomial population,” *The*
6997 *Annals of Mathematical Statistics*, 142–152.
- 6998 Schofield and Barker (missing), “Missing,” *Missing*, missing.
- 6999 Sepúlveda, M., Bartheld, J., Monsalve, R., Gómez, V., and Medina-Vogel, G.
7000 (2007), “Habitat use and spatial behaviour of the endangered Southern river ot-
7001 ter (*Lontra provocax*) in riparian habitats of Chile: conservation implications,”
7002 *Biological Conservation*, 140, 329–338.
- 7003 Sillett (2011), “Missing,” *Missing*, missing.
- 7004 Simons, T. R., Pollock, K. H., Wettroth, J. M., Alldredge, M. W., Pacifici, K., and
7005 Brewster, J. (2009), “Sources of Measurement Error, Misclassification Error, and
7006 Bias in Auditory Avian Point Count Data,” in *Modeling Demographic Processes*
7007 *In Marked Populations*, eds. Thomson, D. L., Cooch, E. G., and Conroy, M. J.,
7008 Boston, MA: Springer US, vol. 3, pp. 237–254.
- 7009 Spiegelhalter, D., Thomas, A., Best, N., and Lunn, D. (2003), *WinBUGS User*
7010 *Manual Version 1.4*.
- 7011 Spiegelhalter, D. J., Best, N. G., Carlin, B. P., and Van Der Linde, A. (2002),
7012 “Bayesian measures of model complexity and fit,” *Journal of the Royal Statistical*
7013 *Society. Series B, Statistical Methodology*, 583–639.
- 7014 Stabler, B. (2006), *shapefiles: Read and Write ESRI Shapefiles*, r package version
7015 0.6.
- 7016 Sturtz, S., Ligges, U., and Gelman, A. (2005), “R2WinBUGS: A Package for Run-
7017 ning WinBUGS from R,” *Journal of Statistical Software*, 12, 1–16.
- 7018 Su, Y.-S. and Yajima, M. (2011), *R2jags: A Package for Running jags from R*, r
7019 package version 0.02-17.
- 7020 Tanner, M. A. and Wong, W. H. (1987), “The calculation of posterior distributions
7021 by data augmentation,” *J Am Stat Assoc*, 82, 528–540.
- 7022 Thomas, A., O’Hara, B., Ligges, U., and Sturtz, S. (2006), “Making BUGS Open,”
7023 *R News*, 6, 12–17.
- 7024 Trolle, M. and Kéry, M. (2005), “Camera-trap study of ocelot and other secretive
7025 mammals in the northern Pantanal,” *Mammalia*, 69, 409–416.
- 7026 Tyre, A. J., Tenhumberg, B., Field, S. A., Niejalke, D., Parris, K., and Possing-
7027 ham, H. P. (2003), “Improving precision and reducing bias in biological surveys:
7028 estimating false-negative error rates,” *Ecological Applications*, 13, 1790–1801.
- 7029 Wegan, M. (2008), “Aversive conditioning, population estimation, and habitat pref-
7030 erence of black bears (*Ursus americanus*) on Fort Drum Military Installation in
7031 northern New York,” Ph.D. thesis, Cornell University, Jan.

- 7032 Wilson, K. R. and Anderson, D. R. (1985), “Evaluation of Two Density Estimators
7033 of Small Mammal Population Size,” *Journal of Mammalogy*, 66, 13–21.
- 7034 Wolpert, R. L. and Ickstadt, K. (1998), “Poisson/gamma random field models for
7035 spatial statistics,” *Biometrika*, 85, 251 –267.
- 7036 Yang, H. C. and Chao, A. (2005), “Modeling Animals’ Behavioral Response by
7037 Markov Chain Models for Capture–Recapture Experiments,” *Biometrics*, 61,
7038 1010–1017.
- 7039 Zuur, A., Ieno, E., Walker, N., Saveliev, A., and Smith, G. (2009), *Mixed effects*
7040 *models and extensions in ecology with R*, Springer Verlag.

INDEX

- 7041 bracket notation, 12
- 7042 R package
- 7043 sp, 147
- 7044 R package
- 7045 maptools, 147
- 7046 shapefile, 147