

₁ Chapter 1

₂ Introduction

³ Chapter 2

⁴ GLMS and WinBUGS

⁵ Chapter 3

⁶ Closed population models

⁷ Chapter 4

⁸ Fully Spatial ⁹ capture-recapture models

¹⁰ **Chapter 5**

¹¹ **Other observation models**

¹² **Chapter 6**

¹³ **MCMC details**

Chapter 7

MCMC Details

7.1 Introduction

In this chapter we will dive a little deeper into Markov chain Monte Carlo (MCMC) sampling. We will construct custom MCMC samplers in R, starting with easy-to-code GLMs and GLMMs and moving on to simple SCR models. We will also demonstrate some tricks and simple extensions to the 'spatial null model'. Finally, we will illustrate some alternative ready-to-use software packages for MCMC sampling. We will NOT provide exhaustive background information on the theory and justification of MCMC sampling there are entire books dedicated to that subject and we refer you to Robert and Casella (2004) and Robert and Casella (2010). Rather we aim to provide you with enough background and technical know-how to start building your own MCMC samplers for SCR models in R.

7.1.1 Why build your own MCMC algorithm?

The standard program we have used so far to run MCMC analyses is WinBUGS (Gilks et al., 1994). The wonderful thing about WinBUGS is that it will automatically use the most appropriate and efficient form of MCMC sampling for the model specified by the user.

The fact that we have such a Swiss Army knife type of MCMC machine begs the question: Why would anyone want to build their own MCMC algorithm? For one, there are a limited number of distributions and functions implemented in WinBUGS. While OpenBUGS provides more options, some more complex models may be impossible to build within these programs. A very simple example from spatial capture-recapture that can give you a headache in WinBUGS is when your state-space is an irregular-shaped polygon, rather than an ideal rectangle that can be characterized by four pairs of coordinates. It is easy to restrict activity centers to any arbitrary polygon in R using an ESRI shapefile

(and we will show you an example in a little bit), but you cannot use a shape file in a BUGS model.

Sometimes implementing an MCMC algorithm in R may be faster than in WinBUGS - especially if you want to run simulation studies where you have hundreds or more simulated data sets, several years' worth of data or other large models, this can be a big advantage.

Finally, building your own MCMC algorithm is a great exercise to understand how MCMC sampling works. So while using the BUGS language requires you to understand the structure of your model, building an MCMC algorithm requires you to think about the relationship between your data, priors and posteriors, and how these can be efficiently analyzed and characterized. Not to mention that, if you are an R junkie, it can actually be fun. However, if you don't think you will ever sit down and write your own MCMC sampler, consider skipping this chapter - apart from coding it will not cover anything SCR-related that is not covered by other, more model-oriented chapters as well.

7.2 MCMC and posterior distributions

As mentioned in Chapter 2, MCMC is a class of simulation methods for drawing (correlated) random numbers from a target distribution, which in Bayesian inference is the posterior distribution. As a reminder, the posterior distribution is a probability distribution for an unknown parameter, say θ , given a set of observed data and its prior probability distribution (the probability distribution we assign to a parameter before we observe data). The great benefit of computing the posterior distribution of θ is that it can be used to make probability statements about θ , such as the probability that θ is equal to some value, or the probability that θ falls within some range of values. As an example, suppose we conducted a Bayesian analysis to estimate detection probability of some species at a study site (p), and we obtained a posterior distribution of $\text{beta}(20,10)$ for the parameter p . The following R commands demonstrate how we make inferences based upon summaries of the posterior distribution. Fig 1 shows the posterior along with the summary statistics.

```
> (post.median <- qbeta(0.5, 20, 10))
[1] 0.6704151
> (post.95ci <- qbeta(c(0.025, 0.975), 20, 10))
[1] 0.4916766 0.8206164
```

Thus, we can state that there is a 95% probability that θ lies between 0.49 and 0.82.

The posterior distribution summarizes all we know about a parameter and thus, is the central object of interest in Bayesian analysis. Unfortunately, in many if not most practical applications, it is nearly impossible to directly compute the posterior. Recall Bayes theorem:

$$p(\theta|y) = p(y|\theta) * p(\theta)/p(y), \quad (7.1)$$




Figure 7.1: Probability density plot of a hypothetical posterior distribution of $\text{beta}(20,10)$; dashed lines indicate mean and upper and lower 95% interval

where θ is the parameter of interest, y is the observed data, $p(\theta|y)$ is the posterior, $p(y|\theta)$ the likelihood of the data conditional on θ , $p(\theta)$ the prior probability of θ , and, finally, $p(y)$ is the marginal probability of the data, which can also be written as

$$p(y) = \int p(y|\theta) * p(\theta) d\theta$$

This marginal probability is a normalizing constant that ensures that the posterior integrates to 1. You read in Chapter 2 that this integral is often hard or impossible to evaluate, unless you are dealing with a really simple model. For example, consider that you have a Normal model, with a set of n observations, y that come from a Normal distribution:

$$y \sim \text{Normal}(\mu, \sigma),$$

where σ is known and our objective is to obtain an estimate of μ using Bayesian statistics. To fully specify the model in a Bayesian framework, we first have to define a prior distribution for μ . Recall from Chapter 2 that for certain data models, certain priors lead to conjugacy i.e. if you choose the right prior for your parameter, your posterior distribution will be of a known parametric form. The conjugate prior for the mean of a normal model is also a Normal distribution:

$$\mu \sim \text{Normal}(\mu_0, \sigma_0^2)$$

If μ_0 and σ_0^2 are fixed, the posterior for μ has the following form (for the algebraic proof, see XXX):

$$\mu|y \sim \text{Normal}(\mu_n, \sigma_n^2) \quad (7.2)$$

100 where

$$\mu_n = (sig^2/sig^2 + n * sig0^2) * mu0 + (n * sig0^2/sig^2 + n * sig0^2) * y - bar$$

101 And

$$sign^2 = sig^2 * sig0^2 / (sig^2 + n * sig0^2)$$

102 We can directly obtain estimates of interest from this Normal posterior distri-
 103 bution, such as the mean mu-hat and its variance; we do not need to apply
 104 MCMC, since we can recognize the posterior as a parametric distribution, in-
 105 cluding the normalizing constant $p(y)$. But generally we will be interested in
 106 more complex models with several, say n , parameters. In this case, computing
 107 $p(y)$ from Eq. 7.1 requires n -dimensional integration, which is can be difficult
 108 or impossible. Thus, the posterior distribution in generally only known up to a
 109 constant of proportionality:

$$p(\theta|y)proptop(y|\theta) * p(\theta)$$

110 The power of MCMC is that it allows us to approximate the posterior using sim-
 111 ulation without evaluating the high dimensional integrals and to directly sample
 112 from the posterior, even when the posterior distribution is unknown! The price
 113 is that MCMC is computationally expensive. Although MCMC first appeared
 114 in the scientific literature in 1949 (Metropolis and Ulam, 1949), widespread use
 115 did not occur until the 1980s when computational power and speed increased
 116 (Gelfand and Smith, 1990). It is safe to say that the advent of practical MCMC
 117 methods is the primary reason why Bayesian inference has become so popular
 118 during the past three decades. In a nutshell, MCMC lets us generate sequential
 119 draws of θ (the parameter(s) of interest) from distributions approximating the
 120 unknown posterior over T iterations. The distribution of the draw at t depends
 121 on the value drawn at $t-1$; hence, the draws from a Markov chain.¹ As T goes
 122 to infinity, the Markov chain converges to the desired distribution in our case
 123 the posterior distribution for $\theta|y$. Thus, once the Markov chain has reached
 124 its stationary distribution, the generated samples can be used to characterize
 125 the posterior distribution, $p(\theta|y)$, and point estimates of θ , its standard error
 126 and confidence bounds, can be obtained directly from this approximation of the
 127 posterior. In practice, although we know that a Markov chain will eventually
 128 converge, we can only generate a limited number of samples a process that
 129 depending on the model can be quite time consuming. Assessing whether our
 130 Markov chain has indeed converged is an important part of MCMC sampling
 131 and we will speak about some common diagnostics in Section XX.

132 7.3 Types of MCMC sampling

133 There are several MCMC algorithms, the most popular being Gibbs sampling
 134 and Metropolis-Hastings sampling. We will be dealing with these two classes in

¹In case you are not familiar with Markov chains, for t random samples $\theta(1), \dots, \theta(t)$ from a Markov chain the distribution of $\theta(t)$ depends only on the most recent value, $\theta(t-1)$.

more detail and use them to construct the MCMC algorithms for SCR models. Also, we will briefly review alternative techniques that are applicable in some situations.

7.3.1 Gibbs sampling

Gibbs sampling was named after the physicist J.W. Gibbs by Geman and Geman (1984), who applied the algorithm to a Gibbs distribution². The roots of Gibbs sampling can be traced back to work of Metropolis et al. (1953), and it is actually closely related to Metropolis sampling (see Chapter 11.5 in Gelman et al. (2004), for the link between the two samplers). We will focus on the technical aspects of this algorithm, but if you find yourself hungry for more background, Casella and George (1992) provide a more in-depth introduction to the Gibbs sampler.

In Chapter 2 you already heard about the basic principles of Gibbs sampling³. But as a refresher, let's go back to our simple example from above to understand the motivation and functioning of Gibbs sampling. Recall that for a Normal model with known variance and a Normal prior for μ , the posterior distribution of $\mu|y$ is also Normal. Conversely, with a fixed (known) μ , but unknown variance, the conjugate prior for σ^2 is an Inverse-Gamma distribution with shape and scale parameters a and b :

$$\sigma^2 \sim \text{Inv-Gamma}(a, b),$$

With fixed a and b , the posterior $p(\text{sig}|\mu, y)$ is also an Inverse Gamma distribution, namely:

$$\text{sig}|\mu, y \sim \text{InvGamma}(an, bn), \quad (7.3)$$

where $an = n/2 + a$ and $bn = 1/2\sigma(y_i - \mu)^2 + b$. However, what if we know neither μ nor sig , which is probably the more common case? The joint posterior distribution of μ and sig now has the general structure

$$p(\mu, \text{sig}|y) = \frac{p(y|\mu) * p(\mu) * p(\text{sig})}{\int p(y|\mu) * p(\mu) * p(\text{sig}) d\mu d\text{sig}}$$

Or

$$p(\mu, \text{sig}|y) \propto p(y|\mu) * p(\mu) * p(\text{sig})$$

This cannot easily be reduced to a distribution we recognize. However, we can condition μ on sig (i.e., we treat sig as fixed) and remove all terms from the joint posterior distribution that do not involve μ to construct the full conditional distribution,

$$p(\mu|\text{sig}, y) \propto p(y|\mu) * p(\mu)$$

The full conditional of μ again takes the form of the Normal distribution shown in Eq. ??; similarly, $p(\text{sig}|\mu, y)$ takes the form of the Inverse Gamma

²a distribution from physics we are not going to worry about, since it has no immediate connection with Gibbs sampling other than giving its name

³maybe we should think out chapter 2 and concentrate that material here?

165 distribution shown in Eq. 7.3 both distribution we can easily sample from.
 166 And this is precisely what we do when using Gibbs sampling we break down
 167 high-dimensional problems into convenient one-dimensional problems by con-
 168 structing the full conditional distributions for each model parameter separately;
 169 and we sample from these full conditionals, which, if we choose conjugate priors,
 170 are known parametric distributions. Let's put the concept of Gibbs sampling
 171 into the MCMC framework of generating successive samples, using our simple
 172 Normal model with unknown μ and σ and conjugate priors as an example.
 173 These are the steps you need to build a Gibbs sampler:

174 **Step 0:** Begin with some initial values for θ , $\theta(0)$. In our example, we have to
 175 specify initial values for μ and σ , for example by drawing a random number
 176 from some uniform distribution, or by setting them close to what we think they
 177 might be. (Note: This step is required in any MCMC sampling chains have to
 178 start from somewhere. We will get back to these technical details a little later.)

179 **Step 1:** Draw $\theta_1(1)$ from the conditional distribution $p(\theta_1(1) | \theta_2(0), \dots, \theta_d(0))$
 180 Here, θ_1 is μ , which we draw from the Normal distribution in Eq. 7.3 using
 181 $\sigma(0)$ as value for σ .

182 **Step 2:** Draw $\theta_2(1)$ from the conditional distribution $p(\theta_2(1) | \theta_1(1), \theta_3(0), \dots, \theta_d(0))$
 183 Here, θ_2 is σ , which we draw from the Inverse Gamma distribution of
 184 Eq. 7.3, using $\mu(1)$ as value for μ ...

185 **Step d:** Draw $\theta_d(1)$ from the conditional distribution $p(\theta_d(1) | \theta_1(1), \dots, \theta_{d-1}(1))$
 186

187 In our example we have no additional parameters, so we only need step 0
 188 through to 2. Repeat Steps 1 to d for K = a large number of samples. In terms
 189 of R coding, this means we have to write Gibbs updaters for μ and σ and
 190 embed them into a loop over K iterations. The final code in the form of an R
 191 function is shown in Panel 1.

192 Andy will build the panel environment here soon.

193
 194 Panel 1: R-code for a Gibbs sampler for a Normal model with unknown μ
 195 and σ and conjugate (Normal and Inverse Gamma, respectively) priors
 196 for both parameters.

```
197
198 Normal.Gibbs<-function(y=y,mu0=mu0, sig0=sig0, a=a,b=b,niter=niter) {
199
200   ybar<-mean(y)
201   n<-length(y)
202   mu<-runif(1) #mean initial value
203   sig<-runif(1) #sd initial value
204   an<-n/2 + a
205
206   out<-matrix(nrow=niter, ncol=2)
```

```

207 colnames(out)<-c('mu', 'sig')
208
209 for (i in 1:niter) {
210
211   #update mu
212   mun<- (sig/(sig+n*sig0))*mu0 + (n*sig0/(sig+n* sig0))*ybar
213   sign <- (sig*sig0)/ (sig+n*sig0)
214   mu<-rnorm(1,mun, sqrt(sign))
215
216   #update sig
217   bn<- 0.5 * (sum((y-mu)^2)) +b
218   sig<-1/rgamma(1,shape=an, rate=bn)
219   out[i,]<-c(mu,sqrt(sig))
220
221 }
222 return(out)
223 }

```

224 This is it! You can use the code `NormalGibbs.R` in the **R** package `scrbook`
 225 to simulate some data, $y \sim \text{Normal}(5, 0.5)$ and run your first Gibbs sampler.
 226 Your output will be a table with two columns, one per parameter, and K rows,
 227 one per iteration. For this 2-parameter example you can visualize the joint
 228 posterior by plotting samples of μ against samples of σ (Fig. 2 XXX):

```

229 plot(out[,1], out[,2])

```

230 The marginal distribution of each parameter is approximated by just examining
 231 the samples of this particular parameter you can visualize it by plotting a
 232 histogram of the samples (Fig. 3 a, b XXX):

```

233 par(mfrow=c(1,2))
234 hist(out[,1]); hist (out[,2])

```

235 Finally, recall an important characteristic of Markov chains, namely, that the
 236 chain has to have converged (reached its stationary distribution) for samples to
 237 come from the posterior distribution. In practice, that means you have to throw
 238 out some of the initial samples called the burn-in. We will talk about this in
 239 more when we talk about convergence diagnostics. For now, you can use the
 240 `plot(out[,1])` or `plot(out[,2])` command to make a time series plot of the
 241 samples of each parameter and visually assess how many of the initial samples
 242 you should discard. Figure 3 c and d shows plots for the estimates of μ and
 243 sigma from our simulated data set; you see that in this simple example the
 244 Markov chain apparently reaches its stationary distribution very quickly the
 245 chains look 'grassy' seemingly from the start. It is hard to discern a burn-in
 246 phase visually (but we will see examples further on where the burn-in is clearer)
 247 and you may just discard the first 500 draws to be sure you only use samples
 248 from the posterior distribution. The mean of the remaining samples are your
 249 estimates of μ and σ :




Figure 7.2: Joint posterior distribution of mu and sig from a Normal Model

```

250 > summary(mod[501:10000,])
251      mu              sig
252 Min.   : 4.936      Min.   : 0.4569
253 1st Qu.: 4.984      1st Qu.: 0.4889
254 Median : 4.994      Median : 0.4961
255 Mean   : 4.994      Mean   : 0.4964
256 3rd Qu.: 5.005      3rd Qu.: 0.5037
257 Max.   : 5.062      Max.   : 0.5356

```

258 7.3.2 Metropolis-Hastings sampling

259 Although it is applicable to a wide range of problems, the limitations of Gibbs
 260 sampling are immediately obvious what if we do not want to use conjugate priors
 261 (or what if we cannot recognize the full conditional distribution as a parametric
 262 distribution, or simply do not want to worry about these issues)? The most
 263 general solution is to use the Metropolis-Hastings (MH) algorithm, which also
 264 goes back to the work by Metropolis et al. (1953). You saw the basics of this
 265 algorithm in Chapter 2. In a nutshell, because we do not recognize the posterior
 266 $p(\theta|y)$ as a parametric distribution, the MH algorithm generates samples from a
 267 known proposal distribution, say $h(\theta)$, that depends on θ at $t-1$. The t^{th} sample
 268 is accepted or rejected based on its joint posterior probability density compared
 269 to the density of the sample at $t-1$. The original Metropolis algorithm requires
 270 $h(\theta)$ to be symmetric so that $h(\theta^t|\theta^{t-1}) = h(\theta^{t-1}|\theta^t)$; but a later development




Figure 7.3: Plots of the posterior distributions of μ (a) and sig (b) from a Normal model and time series plots of μ (c) and sig (d).

of the algorithm by Hastings (1970) lifted this condition. Using a symmetric proposal distribution makes life a little easier and we are going to limit our coverage of the Metropolis-Hastings sampler to this specific case. Specifically, we are going to use a Normal proposal distribution, which is also referred to as 'random walk Metropolis-Hastings sampling'. It is worth knowing that there are alternative formulations of the algorithm. For example, in the independent M-H, θ^t does not depend on θ^{t-1} , while the Langevin algorithm (Roberts and Rosenthal, 1998) aims at avoiding the random walk by favoring moves towards regions of higher posterior probability density. The interested reader should look up these algorithms in Robert and Casella (2004) or Robert and Casella (2010).

Building a MH sampler can be broken down into several steps. We are going to demonstrate these steps using a different but still simple and common model the logit-normal or logistic regression model. For simplicity, assume that

$$y \sim \text{Bern}(\exp(\theta)/(1 + \exp(\theta)))$$

and

$$\theta \sim \text{Normal}(\mu_0, \sigma)$$

The following steps are required to set up a random walk MH algorithm:

Step 0: Choose initial values, $\theta(0)$.

Step 1: Generate a proposed value of θ at t from $h(\theta_t - \theta_{t-1})$. We often use a Normal proposal distribution, so we draw θ_1 from $\text{Normal}(\theta_0, \text{sig}^2)$, where sig^2 is the variance of the Normal proposal distribution, a tuning parameter that we have to set.

Step 2: Calculate the ratio of posterior densities for the proposed and the original value for θ :

$$r = p(\theta^t|y)/p(\theta^{t-1}|y)$$

In our example,

$$r = \text{Bern}(y|\theta^t) * \text{Normal}(\theta^t|\mu_0, \sigma_0) / \text{Bernoulli}(y|\theta^{t-1}) * \text{Normal}(\theta^{t-1}|\mu_0, \sigma_0)$$

Step 3: Set

```
\begin{eqnarray*}
\theta^t &= & \theta^{t-1} \text{ with probability } \min(r,1) // \\
&= & \theta^{t-1} \text{ otherwise } \\
\end{eqnarray*}
```

We can do that by drawing a random number u from a $\text{Unif}(0,1)$ and accept θ^t if $u < r$. Repeat for $t = 1, 2, \dots$ a large number of samples. The **R** code for this MH sampler is provided in Panel 2 XXXX.

Panel 2: R code to run a Metropolis sampler on a simple Logit-Normal model.

```
Logreg.MH<-function(y=y, mu0=mu0, sig0=sig0, niter=niter) {
  out<-c()
  theta<-runif(1, -3,3) #initial value
  for (iter in 1:niter){
    theta.cand<-rnorm(1, theta, 0.2)
    loglike<-sum(dbinom(y, 1, exp(theta)/(1+exp(theta)), log=TRUE))
    logprior <- dnorm(theta,mu0 ,sig0, log=TRUE)
    loglike.cand<-sum(dbinom(y, 1, exp(theta.cand)/(1+exp(theta.cand)), log=TRUE))
    logprior.cand <- dnorm(theta.cand, mu0, sig0, log=TRUE)
    if (runif(1)<exp((loglike.cand+logprior.cand)-(loglike+logprior))){
      theta<-theta.cand
    }
    out[iter]<-theta
  }
  return(out)
}
```

The reason we sum the logs of the likelihood and the prior, rather than multiplying the original values, is simply computational. The product of small probabilities can be numbers very close to 0, which computers do not handle well. Thus we add the logarithms, sum, and exponentiate to achieve the desired result. Similarly, in case you have forgotten some elementary math, $x/y = \exp(\log(x) - \log(y))$, with the latter being favored for computational reasons.

Comparing MH sampling to Gibbs sampling, where all draws from the conditional distribution are used, in the MH algorithm we discard a portion of the candidate values, which inherently makes it less efficient than Gibbs sampling the price you pay for its increased generality. In Step 1 of the MH sampler we had to choose a variance for the Normal proposal distribution. Choice of the parameters that define our candidate distribution is also referred to as 'tuning', and it is important since adequate tuning will make your algorithm more efficient, i.e. your Markov chain will converge faster. The variance should be chosen so that (a) each step of drawing a new proposal value for θ can cover a reasonable distance in the parameter space, as otherwise, the random walk moves too slowly; and (b) proposal values are not rejected too often, as otherwise the random walk will 'get stuck' at specific values for too long. As a rule of thumb, your candidate value should be accepted in about 40% of all cases. Acceptance rates of 20–80% are probably ok, but anything below or above may well render your algorithm inefficient (this does not mean that it will give you wrong results only that you will need more iterations to converge to the posterior distribution). In practice, tuning will require some 'trial-and-error' and some common sense. Or, one can use an adaptive phase, where the tuning parameter is automatically adjusted until it reaches a user-defined acceptance rate, at which point the adaptive phase ends and the actual Markov chain begins. This is computationally a little more advanced. Link and Barker (2009) discuss this in more detail. It is important the samples drawn during the adaptive phase are discarded. You can easily check acceptance rates for the parameters you monitor (that are part of your output) using the `rejectionRate()` function of the package `coda` (we will talk more about this package a little later on). Do not let the term 'rejection rate' confuse you; it is simply $1 - \text{acceptance rate}$. There may be parameters for example, individual values of a random effect or latent variables that you do not want to save, though, and in our next example we will show you a way to monitor their acceptance rates with a few extra lines of code.

7.3.3 Metropolis-within-Gibbs

One weakness of the MH sampler is that formulating the joint posterior when evaluating whether to accept or reject the candidate values for θ becomes increasingly complex or inefficient as the number of parameters in a model increases. It is probably going to sound like MCMC sampling is too good to be true but in these cases you can simply combine MH sampling and Gibbs sampling. You can use Gibbs sampling to break down your high-dimensional parameter space into easy-to-handle one-dimensional conditional distributions and use MH sampling for these conditional distributions. Better yet if you have some conjugacy in your model, you can use the more efficient Gibbs sampling for these parameters and one-dimensional MH for all the others. You have already seen the basics of how to build both types of algorithms, so we can jump straight into an example here and build a Metropolis-within-Gibbs algorithm.

376 7.4 GLMMs Poisson regression with a random 377 effect

378 Let's assume a model that gets us closer to the problem we ultimately want to
379 deal with a GLMM. Here, we assume we have Poisson counts, y , from i plots
380 in j different study sites, and we believe that the counts are influenced by some
381 plot-specific covariate, x , but that there is also a random site effect. So our
382 model is:

$$383 \quad y_{ij} \sim \text{Poisson}(lami_{ij})$$

$$lami_{ij} = \exp(a_j + b * x_i)$$

384 Let's use Normal priors on a and b ,

$$a_j \sim \text{Normal}(\mu_a, \sigma_a)$$

385 and

$$b \sim \text{Normal}(\mu_b, \sigma_b)$$

386 .⁴ Since we want to estimate the random effect in this model, we do not specify
387 μ_a and σ_a , but instead, estimate them as well, so we have to specify hyperpriors
388 for these parameters:

$$\mu_a \sim \text{Normal}(\mu_{a0}, \sigma_{a0})$$

$$\sigma_a \sim \text{InvGamma}(a_0, b_0)$$

389 With the model fully specified, we can compile the full conditionals, breaking
390 the multi-dimensional parameter space into one-dimensional components:

```
391 \begin{eqnarray*}
392 p(a_1|a_2,a_3,a_j,b,y) & \propto & p(y_{i1}|a_1,b) * p(a_1|\mu_a, \sigma_a) \backslash\backslash \\
393 & \propto & \text{Poisson}(y_{i1}|\exp(a_1 + b*x[j=1])) * \text{Normal}(a_1|\mu_a, \sigma_a) \\
394 \end{eqnarray*}
395 \begin{eqnarray*}
396 p(a_2|a_1,a_3,a_j,b,y) & \propto & p(y_{i2}|a_2,b) * p(a_2|\mu_a, \sigma_a) \backslash\backslash \\
397 & \propto & \text{Poisson}(y_{i2}|\exp(a_2 + b*x[j=1])) * \text{Normal}(a_2|\mu_a, \sigma_a) \\
398 \end{eqnarray*}
399 \text{and so on for all elements of } a.
400 \begin{eqnarray*}
401 p(b|a,y) & \propto & p(y|a,b) * p(b) \backslash\backslash \\
402 & \propto & \text{Poisson}(y|\exp(a + b*x)) * \text{Normal}(b|\mu_b, \sigma_b) \\
403 \end{eqnarray*}
```

404 Finally, we need to update the hyperparameters for a :

$$p(\mu_a|a) \propto p(a|\mu_a, \sigma_a) * p(\mu_a)$$

$$405 \quad p(\sigma_a|a) \propto p(a|\mu_a, \sigma_a) * p(\sigma_a)$$

⁴Why is b a hyperparameter?

Since we assumed a to come from a Normal distribution, the choice of priors for μ_a Normal and σ_a Inverse Gamma leads to the same conjugacy we observed in our initial Normal model, so that both hyperparameters can be updated using Gibbs sampling.

Now let's build the updating steps for these full conditionals. Again, for the MH steps that update a and b we use Normal proposal distributions with standard deviations σ_a and σ_b .

First, we set the initial values $a(0)$ and $b(0)$. Then, starting with $a(1)$, we draw $a(1)$ from Normal($a(0)$, σ_a), calculate the conditional posterior density of $a(0)$ and $a(1)$ and compare their ratios,

$$r = \text{Poisson}(y(j=1)|\exp(a(1)+b*x)) * \text{Normal}(a(1)|\mu_a, \sigma_a) / \text{Poisson}(y(j=1)|\exp(a(0)+b*x)) * \text{Normal}(a(0)|\mu_a, \sigma_a)$$

and accept $a(1)$ with probability $\min(r, 1)$. We repeat this for all a 's.

For b , we draw $b(1)$ from Normal($b(0)$, σ_b), compare the posterior densities of $b(0)$ and $b(1)$,

$$r = \text{Poisson}(y|\exp(a+b(1)*x)) * \text{Normal}(b(1)|\mu_b, \sigma_b) / \text{Poisson}(y|\exp(a+b(0)*x)) * \text{Normal}(b(0)|\mu_b, \sigma_b),$$

and accept $b(1)$ with probability $\min(r, 1)$.

For μ_a and σ_a , we sample directly from the full conditional distributions (Eq XX and Eq XX):

$$\mu_a(1) \sim \text{Normal}(\mu_n, \sigma_n)$$

where $\mu_n = (\sigma_a(0)/\sigma_a(0) + n_a * \sigma_a(0)) * \mu_0 + (n_a * \sigma_a(0)/\sigma_a(0) + n_a * \sigma_a(0)) * \bar{a}$ and $\sigma_n = \sigma_a(0) * \sigma_a(0) / (\sigma_a(0) + n_a * \sigma_a(0))$. Here, \bar{a} is the current mean of the vector a , which we updated before, and n_a is the length of a . For σ_a we use $\sigma_a(1) \sim \text{InvGamma}(a_n, b_n)$, where $a_n = n_a/2 + a_0$, and $b_n = 1/2 \sum (a(1) - \mu_a(1))^2 + b_0$.

We repeat these steps over K iterations of the MCMC algorithm. In this example we may not want to save each value for a , but are only interested in their mean and standard deviation. Since these two parameters will change as soon as the value for one element in a changes, their acceptance rates will always be close to 1 and are not representative of how well your algorithm performs. To monitor the acceptance rates of parameters you do not want to save, you simply need to add a few lines of code into your updater to see how often the individual parameters are accepted. The full code for the MCMC algorithm of our Poisson GLMM in Panel 3 shows one way how to monitor acceptance of individual a 's.

Panel 3: R code for the Metropolis-within-Gibbs sampler for a Poisson regression with random intercepts.

```
Pois.reg<-function(y=y,site=site,mu0=mu0,sig0=sig0,a0=a0,b0=b0,
  mub=mub, sigb=sigb, niter=niter){
  lev<-length(unique(site))    #number of sites
```

```

444 a<-runif(lev,-5,5) #initial values a
445 b<-runif(1,0,5) #initial value b
446 mua<-mean(a)
447 siga<-sd(a)
448
449 out<-matrix(nrow=niter, ncol=3)
450 colnames(out)<-c('mua','siga','b')
451
452 for (iter in 1:niter) {
453
454   #update a
455   aUps<-0 #initiate counter for acceptance rate of a
456   for (j in 1:lev) { #loop over sites
457     a.cand<-rnorm(1, a[j], 0.1) #update intercepts a one at a time
458     loglike<- sum(dpois (y[site==j], exp(a[j] + b*x[site==j]), log=TRUE))
459     logprior<- dnorm(a[j], mua,siga, log=TRUE)
460     loglike.cand<- sum(dpois (y[site==j], exp(a.cand + b *x[site==j]), log=TRUE))
461     logprior.cand<- dnorm(a.cand, mua,siga, log=TRUE)
462     if (runif(1)< exp((loglike.cand+logprior.cand) (loglike+logprior))) {
463       a[j]<-a.cand
464       aUps<-aUps+1
465     }
466   }
467
468   if(iter %% 100 == 0) { #this lets you check the acceptance rate of a at every 100th iteration
469     cat(" Acceptance rates\n")
470     cat(" a =", aUps/lev, "\n")
471   }
472
473   #update b
474   b.cand<-rnorm(1, b, 0.1)
475   avec<-rep(a, times=c(rep(10,10)))
476   loglike<- sum(dpois (y, exp(avec + b*x), log=TRUE))
477   logprior<- dnorm(b, mub,sigb, log=TRUE)
478   loglike.cand<- sum(dpois (y, exp(avec + b.cand *x), log=TRUE))
479   logprior.cand<- dunif(b.cand, mub,sigb, log=TRUE)
480   if (runif(1)< exp((loglike.cand+logprior.cand) (loglike+logprior) )) {
481     b<-b.cand
482   }
483
484   #update mua using Gibbs sampling
485   abar<-mean(a)
486   mun<- (siga/(siga+lev*sig0))*mu0 + (lev*sig0/(siga+lev* sig0))*abar
487   sign <- (siga*sig0)/ (siga+lev*sig0)
488   mua<-rnorm(1,mun, sqrt(sign))
489
490   #update siga using Gibbs sampling
491   a0n<-lev/2 + a0
492   b0n<- 0.5 * (sum((a-mua)^2)) +b0
493   siga<-1/rgamma(1,shape=a0n, rate=b0n)

```

```

494
495 out[iter,]<-c(mua, sqrt(siga), b)
496
497 }
498
499 return(out)
500 }

```

7.4.1 Rejection sampling and slice sampling

While MH and Gibbs sampling are probably the most widely applied algorithms for posterior approximation, there are other options that work under certain circumstances and may be more efficient when applicable. WinBUGS applies these algorithms and we want you to be aware that there is more out there to approximate posterior distributions than Gibbs and MH. One alternative algorithm is rejection sampling. Rejection sampling is not an MCMC method, since each draw is independent of the others. The method can be used when the posterior $p(\theta|y)$ is not a known parametric distribution but can be expressed in closed form. Then, we can use a so-called envelope function, say, $g(\theta)$, that we can easily sample from, with the restriction that $p(\theta|y) < M * g(\theta)$. We then sample a candidate value for θ from $g(\theta)$, calculate $r = p(\theta|y)/M * g(\theta)$ and keep the sample with the probability r . M is a constant that has to be picked so that $r \in [0,1]$, for example by evaluating both $p(\theta|y)$ and $g(\theta)$ at n points and looking at their ratios. Rejection sampling only works well if $g(\theta)$ is similar to $p(\theta|y)$, and packages like WinBUGS use adaptive rejection sampling (Gilks and Wild, 1992), where a complex algorithm is used to fit an adequate and efficient $g(\theta)$ based on the first few draws. Though efficient in some situations, rejection sampling does not work well with high-dimensional problems, since it becomes increasingly hard to define a reasonable envelope function. For an example of rejection sampling in the context of SCR models, see Chapter 9. Another alternative is slice sampling (Neal, 2003). In slice sampling, we sample uniformly from the area under the plot of $p(\theta|y)$. Considering a single univariate θ . Let's define an auxiliary variable, $U \sim Uniform(0, p(\theta|y))$. Then, θ can be sampled from the vertical slice of $p(\theta|y)$ at U (Figure 4):

$\theta|U \sim \text{Unif}(B)$,

where $B = \{\theta : U < p(\theta|y)\}$

5

Slice sampling can be applied in many situations; however, implementing an efficient slice sampling procedure can be complicated. We refer the interested reader to chapter 7 of ? for a simple example. Both rejection sampling and slice sampling can be applied on one-dimensional conditional distributions within a Gibbs sampling setup.

⁵there are supposed to be equations in the caption of figure 4 but it kept causing errors




Figure 7.4: Slice sampling. For...

7.5 MCMC for closed capture-recapture Model

Mh

6

7.6 MCMC algorithm for the basic spatial capture-recapture model

By now you have seen how to build MCMC algorithms for some basic generalized linear models. Now, we'll walk you through the steps of building your own MCMC sampler for the basic SCR model (i.e. without any individual, site or time specific covariates) with both a Poisson and a binomial encounter process. As usual, we will have to go through two general steps before we write the MCMC algorithm:

- (1) Identify your model with all its components (including priors)
- (2) Recognize and express the full conditional distributions for all parameters

It is worthwhile to go through all of step 1 for an SCR model, but you have probably seen enough of step 2 in our previous examples to get the essence of how to express a full conditional distribution. Therefore, we will exemplify step 2 for some parameters and tie these examples directly to the respective R code.

Step 1 Identify your model

Recall the components of the basic SCR model with a Poisson encounter process from Chapter 3: We assume that individuals i , or rather, their activity centers s_i , are uniformly distributed across our state space S ,

$$s_i \sim U(S)$$

⁶Andy could move material from chapter 3 to here.

and that the number of times individual i encounters trap j , y_{ij} , is a random Poisson variable with mean lam_{ij} ,

$$y_{ij} \sim \text{Poisson}(\text{lam}_{ij})$$

The tie between individual location, movement and trap encounter rates is made by the assumption that lam_{ij} , is a decreasing function of the distance between s_i and j , D_{ij} , of the half-normal form

$$\text{Lam}_{ij} = \text{lam}_0 * \exp(-D_{ij}^2/2 * \text{sig}^2),$$

where lam_0 is the baseline trap encounter rate at $D_{ij} = 0$ and sig controls the shape of the half-normal function.

In order to estimate the number of s_i in S , N , we use data augmentation (sect. 3.XYZ) and create M - n all-0 encounter histories, where n is the number of individuals we observed and M is an arbitrary number that is larger than N . We estimate N by summing over the auxiliary data augmentation variables, z_i , which is 1 if the individual is part of the population and 0 if not, and assume that z_i is a random Bernoulli variable,

$$z_i \sim \text{Bern}(\psi)$$

To link the two model components, we modify our trap encounter model to

$$\text{Lam}_{ij} = \text{lam}_0 * \exp(-D_{ij}^2/2 * \text{sig}^2) * z_i.$$

The model has the following structural parameters, for which we need to specify priors ψ the Uniform (0,1) is required as part of the data augmentation procedure and in general is a natural choice of an uninformative prior for a probability; note that this is equivalent to a Beta(1,1) prior, which will come in handy later. s_i since s_i is a pair of coordinates it is two-dimensional and we use a uniform prior limited by the extent of our state-space over both dimensions. σ we can conceive several priors for sigma but let's assume an improper prior one that is Uniform over $(-\text{Inf}, \text{Inf})$. We will see why this is convenient when we construct the full conditionals for sigma. λ_0 analogous, we will use a Uniform $(-\text{Inf}, \text{Inf})$ improper prior for sigma. The parameter that is the objective of our modeling, N , is a derived parameter that we can simply obtain by summing all z 's:

$$N = \text{sum}(z)$$

Step 2 - Construct the full conditionals Having completed step 1, let's look at the full conditional distributions for some of these parameters. We find that with improper priors, full conditionals are proportional only to the likelihood of the observations; for example, take the movement parameter sigma:

$$\text{Sig}|s, \text{lam}_0, z, y \propto [y|s, \text{lam}_0, z, \text{sig}] * [\text{sig}]$$

Since the improper prior implies that $[\text{sig}] \propto 1$, we can reduce this further to

$$\text{Sig}|s, \text{lam}_0, z, y \propto [y|s, \text{lam}_0, z, \text{sig}]$$

588 The R code to update sigma is shown in Panel 4.⁷

589 Panel 4: R code to update sigma within an MCMC algorithm for
590 an SCR model when using an improper prior

```
591
592
593 sig.cand <- rnorm(1, sigma, 0.1) #draw candidate value
594 if(sig.cand>0){ #automatically reject sig.cand that are <0
595     lam.cand <- lam0*exp(-(D*D)/(2*sig.cand*sig.cand))
596     ll<- sum(dpois(y, lam*z, log=TRUE))
597     llcand <- sum(dpois(y, lam.cand*z, log=TRUE))
598     if(runif(1) < exp( llcand - ll) ){
599         ll<-llcand
600         lam<-lam.cand
601         sigma<-sig.cand
602     }
603 }
604
```

605 These steps are analogous for lam0 and si and we will use MH steps for all of
606 these parameters. Similar to the random intercepts in our Poisson GLMM, we
607 update each si individually. Note that to be fully correct, the full conditional
608 for si contains both the likelihood and prior component, since we did not specify
609 an improper, but a Uniform prior on si. However, with a Uniform distribution
610 the probability density of any value is 1/(upper limit - lower limit) = constant.
611 Thus, the prior components are identical for both the current and the candidate
612 value and can be ignored (formally, when you calculate the ratio of posterior
613 densities, r, the identical prior component appears both in the numerator and
614 denominator, so that they cancel each other out).

615 We still have to update zi. The full conditional for zi is

$$z_i|y, \sigma, \lambda_0, \text{sprpto}[y|z, \sigma, \lambda_0, s] * [z_i]$$

616 and since $z_i \sim \text{Bernoulli}(\psi_i)$, the term has to be taken into account when
617 updating zi. The R code for updating zi is shown in Panel 5.

618 Panel 5: R code to update z

```
619
620 zUps <- 0 #set counter to monitor acceptance rate
621 for(i in 1:M) {
622     if(seen[i]) #no need to update seen individuals, since their z =1
623         next
624     zcand <- ifelse(z[i]==0, 1, 0)
625     llz <- sum(dpois(y[i,], lam[i,]*z[i], log=TRUE))
626     llcand <- sum(dpois(y[i,], lam[i,]*zcand, log=TRUE))

```

⁷ Somewhere in chapter 2 i added a comment about rejecting parameters outside of the parameter space as being an ok thing to do. Richard said he read something in Robert and Casellas book on that. Hopefully he can remember where and we can cite it back in Ch 2 and again here. It could be mentioned in a sentence or two up in the MCMC section.

```

627      prior <- dbinom(z[i], 1, psi, log=TRUE)
628      prior.cand <- dbinom(zcand, 1, psi, log=TRUE)
629      if(runif(1) < exp( (llcand+prior.cand) - (llz+prior) )) {
630          z[i] <- zcand
631          zUps <- zUps+1
632      }
633  }
634

```

635 ψ itself is a hyperparameter of the model, with an uninformative prior dis-
636 tribution of Unif(0,1) or Beta(1,1), so that

$$Psi|z \propto [z|psi] * Beta(1, 1)$$

637 The Beta distribution is the conjugate prior to the Binomial and Bernoulli
638 distributions (remember that $z \sim Bernoulli(psi)$). The general form of a full
639 conditional of a Beta-Binomial model with $y_i \sim Bernoulli(p)$ and $p \sim Beta(a, b)$
640 is

$$p(p|y) \propto Beta(a + sum(yi), b + n - sum(yi))$$

641 In our case, this means we update psi as follows:

```

642 si<-rbeta(1, 1+sum(z), 1 + M-sum(z))

```

643 These are all the building blocks you need to write the MCMC algorithm
644 for the spatial null model with a Poisson encounter process. You can find the
645 full R code (SCR0pois.R) in the online supplementary material.

646 7.6.1 SCR model with binomial encounter process

647 The equivalent SCR model with a binomial encounter process is very similar.
648 Here, each individual i can only be detected once at any given trap j during a
649 sampling occasion k . Thus

$$y_{ij} \sim Binomial(p_{ij}, K)$$

650 Where p_{ij} is some function of distance between \mathbf{s}_i and trap location \mathbf{x}_j . Here
651 we use:

$$p_{ij} = 1 - \exp(-lam_{ij})$$

652 Recall from Chapter 2 that this is the complementary log-log (cloglog) link func-
653 tion, which constrains p_{ij} to fall between 0 and 1. For our MCMC algorithm that
654 means that, instead of using a Poisson likelihood, $Poisson(y|sigma, lam0, s, z)$,
655 we use a Binomial likelihood, $Binomial(y, K|sigma, lam0, s, z)$, in all the con-
656 ditional distributions. As an example, Panel 6 shows the updating step for $lam0$
657 under a binomial encounter model. The full MCMC code for the binomial SCR
658 can be found in the online supplements.

```

659 Panel 6: MCMC updater for lam0 in a SCR model with Binomial encounter
660 process and cloglog link function on detection. Here, pmat =
661 1-exp(-lam).
662
663     lam0.cand <- rnorm(1, lam0, 0.1)
664     if(lam0.cand > 0){ #automatically reject lam0.cand that are <0
665         lam.cand <- lam0.cand*exp(-(D*D)/(2*sigma*sigma))
666         p.cand <- 1-exp(-lam.cand)
667         ll<- sum(dbinom(y, K, pmat *z, log=TRUE))
668         llcand <- sum(dbinom(y, K, p.cand *z, log=TRUE))
669         if(runif(1) < exp( llcand - ll) ){
670             ll<-llcand
671             pmat<-p.cand
672             lam0<- lam0.cand
673         }
674     }

```

Another possibility is to model variation in the individual and site specific detection probability, p_{ij} , directly, without any transformation, such that

```

677 pij<-p0 * exp(-Dij2/(2*sig^2))

```

and $p_0 = \{0, 1\}$. This formulation is analogous to how detection probability is modeled in distance sampling under a half-normal detection function; however, in distance sampling p_0 - detection of an individual on the transect line - is assumed to be 1 (Buckland, 2001). Under this formulation the updater for lam_0 (equivalent to p_0 in Eq XX) becomes:

```

683     lam0.cand <- rnorm(1, lam0, 0.1)
684     if(lam0.cand > 0 & lam0.cand < 1 ){ #automatically reject lam0.cand that are
685         lam.cand <- lam0.cand*exp(-(D*D)/(2*sigma*sigma))
686         ll<- sum(dbinom(y, K, lam *z, log=TRUE)) #no transformation needed
687         llcand <- sum(dbinom(y, K, lam.cand *z, log=TRUE))
688         if(runif(1) < exp( llcand - ll) ){
689             ll<-llcand
690             lam<-lam.cand
691             lam0<- lam0.cand
692         }
693     }

```

7.6.2 Looking at model output

Now that you have an MCMC algorithm to analyze spatial capture-recapture data with, let's run an actual analysis so we can look at the output. As an example, we will use the bear data ... ⁸ You can use the same script provided back in Chapter XX to read in the data and build the augmented encounter history

⁸Does this data set come up before Ch6? If not, introduce data here. Or, Andy, would you rather use simulated data?

7.6. MCMC ALGORITHM FOR THE BASIC SPATIAL CAPTURE-RECAPTURE MODEL33

699 array; then source the MCMC code for the binomial encounter model algorithm
700 with the cloglog link and run 5000 iterations. This should take approximately
701 25 minutes.

```
702 > source('SCR0binom.txt')  
703 > mod0<-SCR.0(y=bigTrap, X=trapmat, M=M, xl=xl, xu=xu, yl=yl, yu=yu, K=8, niter=5000)
```

704 Before, we used simple R commands to look at model results. However, there
705 is a specific R package to summarize MCMC simulation output and perform
706 some convergence diagnostics package coda (Plummer et al., 2006). Download
707 and install coda, then convert your model output to an mcmc object

```
708 > chain<-mcmc(mod0)
```

709 which can be used by coda to produce MCMC specific output.

710 Markov chain time series plots

711 Start by looking at time series plots of your Markov chains using `plot(chain)`.
712 This command produces a time series plot and marginal posterior density plots
713 for each monitored parameter, similar to what we did before using the `hist()`
714 and `plot()` commands (Fig. 5). Time series plots will tell you several things:
715 First, the way the chains move through the parameter space gives you an idea
716 of whether your MH steps are well tuned. If chains were constant over many
717 iterations you would probably need to decrease the tuning parameter of the
718 (Normal) proposal distribution. If a chain moves along some gradient to a
719 stationary state very slowly, you may want to increase the tuning parameter so
720 that the parameter space is explored more efficiently.

721 Second, you will be able to see if your chains converged and how many initial
722 simulations you have to discard as burn-in. In the case of the chains shown in
723 Figure 5, we would probably consider the first 750 - 1000 iterations as burn-in,
724 as afterwards the chains seem to be fairly stationary.

725 A word of caution about chain convergence

726 Since we do not know what the stationary posterior distribution of our Markov
727 chain should look like (this is the whole point of doing an MCMC approxima-
728 tion), we effectively have no means to assess whether it has converged to this
729 desired distribution or not. As mentioned before, the only certainty is that a
730 Markov chain will *eventually* converge to its stationary distribution, but no-one
731 can tell us how long this will take. Also, you only now the part of your pos-
732 terior distribution that the Markov chain has explored so far for all you know
733 the chain could be stuck in a local maximum, while other maxima remain com-
734 pletely undiscovered. Acknowledging that there is truly nothing we can do to
735 ever proof convergence of our MCMC chains, there are several things we can do
736 to increase the degree of confidence we have about the convergence of our chains.
737 One option, and that advocated by what we will loosely call the WinBUGS com-
738 munity, is to run several Markov chains and to start them off at different initial




Figure 7.5: Time series and posterior density plots for sigma and lam0.

values that are overdispersed relative to the posterior distribution. Such initial
 values help to explore different areas of the parameter space simultaneously;
 if after a while all chains oscillate around the same average value, chances are
 good that they indeed converged to the posterior distribution. Gelman and Ru-
 bin came up with a diagnostic statistic that essentially compares within-chain
 and between-chain variance to check for convergence of multiple chains (Gel-
 man et al., 2004). Of course, running several parallel chains is computationally
 expensive. Extra computational demands are not the only and by no means
 the major concern some people voice when it comes to running several parallel
 MCMC chains to assess convergence. Again, consider the fact that we do not
 know anything about the true form of the posterior distribution we are trying to
 approximate. How do we, then, know how to pick overdispersed initial values?
 We don't all we can do is pick overdispersed values relative to our expectations
 of what the posterior should look like. To use a quote from the home page
 of Charlie Geyer, a Bayesian statistician from the University of Minnesota, "If
 you don't know any good starting points [...], then restarting the sampler at
 many bad starting points is [...] part of the problem, not part of the solution."
 (<http://users.stat.umn.edu/~charlie/mcmc/diag.html>). His suggestion is that
 your only chance to discover a potential problem with your MCMC sampler is
 to run it for a very long time. But again, there is no way of knowing how long
 is long enough. It is up to you to decide, which school of thoughts appeals more
 to you one long versus several parallel Markov chains. Irrespectively, part of
 developing an MCMC sampler should be to make sure (within reasonable lim-
 its) that you are not missing regions of high posterior density because of the
 way you specify your starting values. Once you have explored the behavior of
 your chain under a reasonable range of starting values, you may feel comfort-

able enough to run only one long chain. The fact that convergence cannot be proven does not mean that you should not look for potential problems in your MCMC sampler. Some problems are easily detected using simple plots, such as the time series plots we discussed above. If the overall trajectory of your chain at the end of your simulations is still upward or downward, your chain clearly has not converged and you need to run your model much longer. If you run several parallel chains and their stationary distributions look different, you may be looking at a multi-modal posterior or a problem with your sampler. With these words of caution, let's get back to looking at our model output.

7.6.3 Posterior density plots

The `plot()` command also produces posterior density plots and it is worthwhile to look at those carefully. For parameters with priors that have bounds (e.g. Uniform over some interval), you will be able to see if your choice of the prior is truncating the posterior distribution. In the context of SCR models, this will mostly involve our choice of M , the size of the augmented data set. If the posterior of N has a lot of mass concentrated close to M (or equivalently the posterior of ψ has a lot of mass concentrated close to 1), as in the example in Figure 6, we have to re-run the analysis with a larger M . A flat posterior plot shows you that the parameter essentially cannot be identified there may not be enough information in your data to estimate model parameters and you may have to consider a simpler model. Finally, posterior density plots will show you if the posterior distribution is symmetrical or skewed if the distribution has a heavy tail, using the mean as a point estimate of your parameter of interest may be biased and you may want to opt for the median or mode instead.

7.6.4 Serial autocorrelation and effective sample size

Even when we can be relatively confident that our chains have converged, the subsequent samples generated from a Markov chain are not iid samples from the posterior distribution, due to the correlation amongst samples introduced by the Markov process. As a consequence, the variance of the mean cannot simply be derived with the standard variance estimator, which takes into account the sample size (here, number of iterations). Rather, the sample size has to be adjusted to account for the autocorrelation in subsequent samples (see Chapter 8 in ? for more details). This adjusted sample size is referred to as the effective sample size. Checking the degree of autocorrelation in your Markov chains and estimating the effective sample size your chain has generated should be part of evaluating your model output. If you use WinBUGS through the R2WinBUGS package, the `print()` command will automatically return the effective sample size for all monitored parameters. In the coda package there are several functions you can use to do so. `effectiveSize()` will directly give you an estimate of the effective sample size for you parameters:

```
> effectiveSize(chain)
```




Figure 7.6: Time series and posterior density plots of ψ and N for the beaver data set truncated by the upper limit of M (500).

```

806      sigma      lam0      psi      N
807      3.930303 78.259159 30.436348 32.047392

```

Alternatively, you can use the `autocorr.diag()` function, which will show you the degree of autocorrelation for different lag values (which you can specify within the function call, we use the defaults below):

```

811 > autocorr.diag(mcmc(mod))
812      sigma      lam0      psi      N
813 Lag 0  1.0000000 1.0000000 1.0000000 1.0000000
814 Lag 1  0.9979948 0.9494134 0.9847503 0.9774201
815 Lag 5  0.9915567 0.8038168 0.9111951 0.9113525
816 Lag 10 0.9836016 0.6714021 0.8462108 0.8509803
817 Lag 50 0.8985337 0.1983780 0.6138516 0.6233994

```

Whichever function you use, if you find that your supposedly long Markov chain has not generated enough pseudo-iid samples, you should consider a longer run. In the present case we see that autocorrelation is especially high for the parameter σ and our effective sample size for this parameter is 4!⁹ This means we would have to run the model for much longer to obtain a reasonable effective sample size. Unfortunately, with many SCR models we observe high degrees of serial autocorrelation, which means we have to run long chains to obtain enough samples that can be considered iid, in order to obtain reasonable estimates of our parameters and their variances. What exactly constitutes a reasonable effective sample size is hard to say, but as a rule of thumb you

⁹Anyone have any idea how the autocorrelation in σ could be reduced?

should probably aim at several hundreds of these pseudo-iid samples. A more meaningful measure of whether you've run your chain for enough iterations is the time-series or Monte Carlo error - the 'noise' introduced into your samples by the stochastic MCMC process - which we introduced in Chapter 2. The MC error decreases with increasing sample size and its magnitude can thus be controlled by adjusting the length of the Markov chain. As a rule of thumb, the MC error should be 1% or less of the parameter estimate. Once you have reached this level, the estimates of the mean, standard error and 95% quantiles should no longer change significantly with additional iterations. For highly correlated samples, it will take more iterations to reduce the MC error. In coda, the MC error is given as part of the summary results (see below). Another option to deal with the serial autocorrelation of samples is to 'thin' Markov chains by some rate r and save only every r -th iteration. But as discussed in Chapter 2, this is not efficient and should only be applied if needed for practical reasons (e.g. a large number of parameters and iterations may force you to thin your samples so you object storing the model output does not become unmanageably large). For now, let's continue using this small set of samples to continue looking at the output.

7.6.5 Summary results

Now that we checked that our chains apparently have converged and pretending that we have generated enough samples from the posterior distribution, we can look at the actual parameter estimates. The `summary()` function will return two sets of results: the mean parameter estimates, with their standard deviation, the naive standard error - i.e. your regular standard error calculated for K (= number of iterations) samples without accounting for serial autocorrelation - and the corrected MC error (Time-series SE), which accounts for autocorrelation. In WinBUGS, this latter value is referred to as MC error and is only given in the log output within BUGS itself. You should adjust the `summary()` call by removing the burn-in from calculating parameter summary statistics. To do so, use the `window()` command, which lets you specify at which iteration to start 'counting'. In contrast to WinBUGS, which requires you to set the burn-in length before you run the model, this command gives us full flexibility to make decisions about the burn-in after we have seen the trajectories of our Markov chains. For our example, `summary(window(chain, start=1001))` returns the following output:

```
Iterations = 1001:5000
Thinning interval = 1
Number of chains = 1
Sample size per chain = 4000

1. Empirical mean and standard deviation for each variable,
   plus standard error of the mean:
```

```

871           Mean          SD   Naive SE Time-series SE
872 sigma    1.9986  0.13805 0.0021827      0.016091
873 lam0     0.1096  0.01523 0.0002407      0.001401
874 psi      0.6113  0.09148 0.0014465      0.010734
875 N        489.8535 71.79695 1.1352094      8.431119

```

```

876
877 2. Quantiles for each variable:
878

```

```

879           2.5%      25%      50%      75%      97.5%
880 sigma    1.75780    1.89847    1.9900    2.0944    2.2772
881 lam0     0.08357    0.09824    0.1087    0.1192    0.1427
882 psi      0.45110    0.54838    0.6052    0.6639    0.8192
883 N        366.00000 440.00000 485.0000 530.0000 654.0000

```

```

884     Looking at the MC errors, we see that in spite of the high autocorrelation, the
885     MC error for sigma is below the 1Our algorithm gives us a posterior distribution
886     of N, but we are usually interested in the density, D. Density itself is not a
887     parameter of our model, but we can derive a posterior distribution for D by
888     dividing each value of N (N at each iteration) by the area of the state-space
889     (here 3032.719 km2) and we can use summary statistics of this distribution to
890     characterize D:

```

```

891 > summary(window(chain[,4]/ 3032.719, start=1001))
892 Iterations = 1001:5000
893 Thinning interval = 1
894 Number of chains = 1
895 Sample size per chain = 4000
896

```

```

897 1. Empirical mean and standard deviation for each variable,
898    plus standard error of the mean:
899

```

```

900           Mean          SD          Naive SE Time-series SE
901    0.1615229    0.0236741    0.0003743      0.0027801
902

```

```

903 2. Quantiles for each variable:
904

```

```

905    2.5%    25%    50%    75%    97.5%
906 0.1207 0.1451 0.1599 0.1748 0.2156

```

```

907 If we compare our mean density of 0.16/km2 (and other parameters) with results
908 from the same model run in secr and WinBUGS in Chapter XX, we see that
909 estimates are almost identical (Table 1).

```

910 7.6.6 Other useful commands

```

911 While inspecting the time series plot gives you a first idea of how well you
912 tuned your MH algorithm, use rejectionRate() to obtain the rejection rates (1
913 acceptance rates) of the parameters that are written to your output:

```

```

914 > rejectionRate(chain)
915      sigma      lam0      psi      N
916 0.44108822 0.77675535 0.00000000 0.01940388

```

Recall that rejection rates should lie between 0.2 and 0.8, so our tuning seems to have been appropriate here. Psi is never rejected since we update it with Gibbs sampling, where all candidate values are kept. And since N is the sum of all z, all it takes for N to change from one iteration to the next are small changes in the z-vector, so the rejection rate of N is always low. If you have run several parallel chains, you can combine them into a single mcmc object using the `mcmc.list()` command on the individual chains (note that each chain has to be converted to an mcmc object before combining them with `mcmc.list()`). You can then easily obtain the Gelman-Rubin diagnostic (Gelman et al., 2004), in WinBUGS called R-hat, using `gelman.diag()`, which will indicate if all chains have converged to the same stationary distribution. For details on these and other functions, see the coda manual, which can be found together with the package on the CRAN mirror.

7.7 Manipulating the state-space

So far, we have constrained the location of the activity centers to fall within the outermost coordinates of our rectangular state space by posing upper and lower bounds for x and y. But what if S has an irregular shape maybe there is a large water body we would like to remove from S, because we know our terrestrial study species does not occur there. Or the study takes place in a clearly defined area such as an island. As mentioned before, this situation is difficult to handle in WinBUGS. In some simple cases we can adjust the state space by setting SX_i to be some function of SY_i or vice versa. In this manner, we can cut off corners of the rectangle to approximate the actual state space. In R, we are much more flexible, as we can use the actual state-space polygon to constrain out si. ¹⁰To illustrate that, let's look at a camera trapping study of Florida panthers (*Puma concolor coryi*) conducted in the Picayune Strand Restoration Project (PSRP) area, southwest Florida (Fig. 7), by XXX, and financed by XXX. In the 1960ies the PSRP area was slated for housing development, but then bought back by the State of Florida and is currently being restored to its original hydrology and vegetation. In an effort to estimate the density of the local Florida panther population, 98 camera traps were operated in the area for 21 months between 2005 and 2007. Florida panthers are wide-ranging animals and in order to account for their wide movements, the state-space was defined as the trapping grid buffered by 15 km around its outermost coordinates. However, the resulting rectangle contained some ocean in its southwestern corner (Fig. 7). In order to precisely describe the state-space, the ocean has to be removed. You can create a precise state-space polygon in ArcGIS and read it into R, or create the polygon directly within R. In the present case we intersected two shape files

¹⁰ Have to check if we can use panther stuff for the book; otherwise, use raccoon example.




Figure 7.7: Rectangular state-space for a Florida panther camera trapping study in the PSRP area (grey outline, red block inset map of Florida) contain some ocean (white) that needs to be removed from the state-space.

one of the state of Florida and one of the rectangle defined by a strip of 15 km around the camera-trapping grid. While you will most likely have to obtain the shapefile describing the landscape of and around your trapping grid (coastlines, water bodies etc.) from some external source, a polygon shapefile buffering your outermost trapping grid coordinates can easily be written in R.

If `xmin`, `xmax`, `ymin` and `ymax`, mark the outermost `x` and `y` coordinates of your trapping grid and `b` is the distance you want to buffer with, load the package `shapefiles` (Stabler, 2006) and use:

```

955 xl= xmin-b
956 xu= xmax+b
957 yl= ymin-b
958 yu= ymax+b
959
960 dd <- data.frame(Id=c(1,1,1,1,1),X=c(xl,xu,xu,xl,xl),Y=c(yl,yl,yu,yu,yl)) #create data
961 ddTable <- data.frame(Id=c(1),Name=c("Item1"))
962 ddShapefile <- convert.to.shapefile(dd, ddTable, "Id", 5) #convert #to shapefile, type
963 write.shapefile(ddShapefile, 'c:/', arcgis=T) # save to location of #choice

```

You can read shapefiles into R loading the package `maptools` (Lewin-Koh et al., 2011) and using the function `readShapeSpatial()`. Make sure you read in shapefiles in UTM format, so that units of the trap array, the movement parameter `sigma` and the state-space are all identical. Intersection of polygons can be done in R also, using the package `rgeos` (Bivand and Rundel, 2011) and the function `gIntersect()`. The area of your single - polygon can be extracted

978 directly from the state-space object SSp:

```
979 > area <- SSp@polygons[[1]]@Polygons[[1]]@area /1000000
```

980 Note that dividing by 1000000 will return the area in km² if your coordi-
 981 nates describing the polygon are in UTM. If your state-space consists of several
 982 disjunct polygons, you will have to sum the areas of all polygons to obtain the
 983 size of the state-space. To include this polygon into our MCMC sampler we
 984 need one last spatial R package `sp` (Pebesma and Bivand, 2011), which has a
 985 function, `over()`, which allows us to check if a pair of coordinates falls within a
 986 polygon or not. All we have to do is embed this new check into the updating
 987 steps for `s`:

```
988         Scand <- as.matrix(cbind(rnorm(M, S[,1], 2),
989                                rnorm(M, S[,2], 2)))          #draw candidate value
990
991 Scoord<-SpatialPoints(Scand*1000)      #convert to spatial points on UTM (m) scale
992 SinPoly<-over(Scoord,SSp) # check if scand is within the polygon
993
994         for(i in 1:M) {
995 if(is.na(SinPoly[i])==FALSE) { #if scand falls within polygon, continue update
996   [rest of the updating step remains the same]
```

997 Note that it is much more time-efficient to draw all M candidate values for s
 998 and check once if they fall within the state-space, rather than running the `over()`
 999 command for every individual pair of coordinates. To make sure that our initial
 1000 values for s also fall within the polygon of S , we use the function `runifpoint()`
 1001 from the package `spatstat` (Baddeley and Turner, 2005), which generates random
 1002 uniform points within a specified polygon. You'll find this modified MCMC al-
 1003 gorithm in the online supplementary material (SCR0poisSSp). Finally, observe
 1004 that we are converting candidate coordinates of S back to meters to match the
 1005 UTM polygon. In all previous examples, for both the trap locations and the
 1006 activity centers we have used UTM coordinates divided by 1000 to estimate
 1007 sigma on a km scale. This is adequate for wide ranging individuals like bears.
 1008 In other cases you may center all coordinates on 0. No matter what kind of
 1009 transformation you use on your coordinates, make sure to always convert can-
 1010 didate values for S back to the original scale (UTM) before running the `over()`
 1011 command.

1012 7.8 MCMC software packages

1013 Throughout most of this book we will use WinBUGS and, occasionally, JAGS
 1014 to run MCMC analyses. Here, we will briefly discuss the main pros and cons of
 1015 these two programs as well as WinBUGS successor OpenBUGS. You can find
 1016 scripts to simulate data and run the basic SCR model in all three programs in
 1017 the online supplementary material (simSCR0poisBUGS).

1018 7.8.1 WinBUGS

1019 In a nutshell, WinBUGS (and the other programs) do everything that we just
 1020 went through in this chapter (and quite a bit more). Looking through your
 1021 model, WinBUGS determines which parameters it can use standard Gibbs sam-
 1022 pling for (i.e. for conjugate full conditional distributions). Then, it determines,
 1023 in the following hierarchy, whether to use adaptive rejection sampling, slice
 1024 sampling or in the 'worst' case Metropolis-Hastings sampling for the other full
 1025 conditionals (Spiegelhalter et al., 2003). If it uses MH sampling, it will auto-
 1026 matically tune the updater so that it works efficiently. While WinBUGS is a
 1027 convenient piece of software that is still widely used, its major drawback is that
 1028 it is no longer being developed, i.e. no new functions or distributions are added
 1029 and no bugs are fixed.

1030 7.8.2 OpenBUGS

1031 OpenBUGS is essentially the successor of WinBUGS. While the latter is no
 1032 longer worked on, OpenBUGS is constantly developed further. The name
 1033 'OpenBUGS' refers to the software being open source, so users do not need
 1034 to download a license key, like they have to for WinBUGS (although the license
 1035 key for WinBUGS is free and valid for life).

1036 Compared to WinBUGS, OpenBUGS has a lot more built-in functions. The
 1037 method of how to determine the right updater for each model parameter has
 1038 changed and the user can manually control the MCMC algorithm used to update
 1039 model parameters. Several other changes have been implemented in OpenBUGS
 1040 and a detailed list of differences between the two BUGS versions, can be found
 1041 at <http://www.openbugs.info/w/OpenVsWin>

1042 While OpenBUGS is a useful program for a lot of MCMC sampling appli-
 1043 cations, for reasons we do not understand, simple SCR models do not converge
 1044 in OpenBUGS. It is therefore advisable that you check any OpenBUGS SCR
 1045 model results against result from WinBUGS. Also, currently, the R package
 1046 BRugs (Thomas et al., 2006) necessary for running OpenBUGS through R
 1047 has problems with 64-bit machines, so you may have to use the 32-bit version
 1048 of R and OpenBUGS in order to make it work. The BUGS project site at
 1049 <http://www.openbugs.info> provides a lot of information on and download links
 1050 for OpenBUGS.

1051 There is an extensive help archive for both WinBUGS and OpenBUGS and
 1052 you can subscribe to a mailing list, where people pose and answer questions of
 1053 how to use these programs at <http://www.mrc-bsu.cam.ac.uk/bugs/overview/list.shtml>

1054 7.8.3 JAGS Just Another Gibbs Sampler

1055 JAGS, currently at Version 3.1.0, is another free program for analysis of Bayesian
 1056 hierarchical models using MCMC simulation. Originally, JAGS was the only
 1057 program using the BUGS language that would run on operating systems other
 1058 than the 32 bit Windows platforms. By now, there are OpenBUGS versions for

Linux or Macintosh machines. JAGS 'only' generates samples from the posterior distribution; analysis of the output is done in R either by running JAGS through R using either the packages `rjags` (Plummer, 2011) or `R2jags` (Su and Yajima, 2011), or by using `coda` on your JAGS output. The program, manuals and `rjags` can be downloaded at <http://sourceforge.net/projects/mcmc-jags/files/>. When run from within R using the package `rjags` or `R2jags`, writing a JAGS model is virtually identical to writing a WinBUGS model. However, some functions may have slightly different names and you can look up available functions and their use in the JAGS manual. One potential downside is that JAGS can be very particular when it comes to initial values. These may have to be set as close to truth as possible for the model to start. Although JAGS lets you run several parallel Markov chains, this characteristic interferes with the idea of using overdispersed initial values for the different chains. Also, we have occasionally experienced JAGS to crash and take the R GUI with it. Only re-installing both JAGS and `rjags` seemed to solve this problem. On the plus side, JAGS usually runs a little faster than WinBUGS, sometimes considerably faster (see section 4.XYZ), is constantly being developed and improved and it has a variety of functions that are not available in WinBUGS. For example, JAGS allows you to supply observed data for some deterministic functions of unobserved variables. In BUGS we cannot supply data to logical nodes. Another useful feature is that the adaptive phase of the model (the burn-in) is run separately from the sampling from the stationary Markov chains. This allows you to easily add more iterations to the adaptive phase if necessary without the need to start from 0. There are other, more subtle differences and there is an entire manual section on differences between JAGS and OpenBUGS. For questions and problems there is a JAGS forum online at <http://sourceforge.net/projects/mcmc-jags/forums/forum/610037>.¹¹

7.9 Summary and Outlook

While there are a number of flexible and extremely useful software packages to perform MCMC simulations, it sometimes is more efficient to develop your own MCMC algorithm. Building an MCMC code follows three basic steps: Identify your model including priors and express full conditional distributions for each model parameter. If full conditionals are parametric distributions, use Gibbs sampling to draw candidate parameter values from these distributions; otherwise use Metropolis-Hastings sampling to draw candidate values from a proposal distribution and accept or reject them based on their posterior probability densities. These custom-made MCMC algorithms give you more modeling flexibility than existing software packages, especially when it comes to handling the state-space: In BUGS (and JAGS for that matter) we define a continuous rectangular state-space using the corner coordinates to constrain the Uniform priors on the activity centers s . But what if a continuous rectangle isn't an ad-

¹¹As we make progress on the book, let's be sure to add linkages to places where we use JAGS in examples.

1100 equate description of the state-space? In this chapter we saw that in R it only
 1101 takes a few lines of code to use any arbitrary polygon shapefile as the state-
 1102 space, which is especially useful when you are dealing with coastlines or large
 1103 bodies of water that need removing from the state-space. Another example is
 1104 the SCR R package SPACECAP (Gopalaswamy et al., 2011) that was developed
 1105 because implementation of an SCR model with a discrete state-space was inef-
 1106 ficient in WinBUGS. Another situations in which using BUGS/JAGS becomes
 1107 increasingly complicated or inefficient is when using point processes other than
 1108 the Uniform Poisson point process which underlies the basic SCR model (see
 1109 Chapter X). In the Chapters 9 and XX you will see examples of different point
 1110 processes, implemented using custom-made MCMC algorithms.¹² Finally,
 1111 the Chapters XX and XX deal with unmarked or partially marked populations
 1112 using hand-made MCMC algorithms to handle the (partially) latent individual
 1113 encounter histories. While some of these models can be written in BUGS/JAGS,
 1114 ¹³, they are painstakingly slow; others cannot be implemented in BUGS/JAGS
 1115 at all. In conclusion, while you can certainly get by using BUGS/JAGS for
 1116 standard SCR models, knowing how to write your own MCMC sampler allows
 1117 you to tailor these models to your specific needs.

¹²Richard, Beth expand on that?

¹³the Poisson one for partially marked we wrote in BUGS and it should work with a known number of marked; the Bernoulli in JAGS with the `dsum()` function should work for the fully unknown; maybe some others? I dont remember. We may have to try writing the others before saying that they dont work in BUGS/JAGS; they are certainly much faster in R, though.

1118 Chapter 8

1119 Goodness of Fit and stuff

1120 Chapter 9

1121 Covariate models

1122 Chapter 10

1123 Inhomogeneous Point 1124 Process

1125 Chapter 11

1126 Open models

Bibliography

- Baddeley, A. and Turner, R. (2005), “Spatstat: an R package for analyzing spatial point patterns,” *Journal of Statistical Software*, 12, 1–42, ISSN 1548-7660.
- Bivand, R. and Rundel, C. (2011), *rgeos: Interface to Geometry Engine - Open Source (GEOS)*, r package version 0.1-8.
- Buckland, S. T. (2001), *Introduction to distance sampling: estimating abundance of biological populations*, Oxford, UK: Oxford University Press.
- Casella, G. and George, E. I. (1992), “Explaining the Gibbs sampler,” *American Statistician*, 46, 167–174.
- Gelfand, A. and Smith, A. (1990), “Sampling-based approaches to calculating marginal densities,” *Journal of the American statistical association*, 85, 398–409.
- Gelman, A., Carlin, J. B., Stern, H. S., and Rubin, D. B. (2004), *Bayesian data analysis, second edition.*, Boca Raton, Florida, USA: CRC/Chapman & Hall.
- Geman, S. and Geman, D. (1984), “Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PAMI-6, 721–741.
- Gilks, W. and Wild, P. (1992), “Adaptive rejection sampling for Gibbs sampling,” *Applied Statistics*, 41, 337–348.
- Gilks, W. R., Thomas, A., and Spiegelhalter, D. J. (1994), “A Language and Program for Complex Bayesian Modelling,” *Journal of the Royal Statistical Society. Series D (The Statistician)*, 43, 169–177, ArticleType: primary_article / Issue Title: Special Issue: Conference on Practical Bayesian Statistics, 1992 (3) / Full publication date: 1994 / Copyright 1994 Royal Statistical Society.
- Gopalaswamy, A. M., Royle, A. J., Hines, J., Singh, P., Jathanna, D., Kumar, N. S., and Karanth, K. U. (2011), *A Program to Estimate Animal Abundance and Density using Spatially-Explicit Capture-Recapture*, r package version 1.0.4.

- 1158 Hastings, W. (1970), “Monte Carlo sampling methods using Markov chains and
1159 their applications,” *Biometrika*, 57, 97–109.
- 1160 Lewin-Koh, N. J., Bivand, R., contributions by Edzer J. Pebesma, Archer, E.,
1161 Baddeley, A., Bibiko, H.-J., Dray, S., Forrest, D., Friendly, M., Giraudoux, P.,
1162 Golicher, D., Rubio, V. G., Hausmann, P., Hufthammer, K. O., Jagger, T.,
1163 Luque, S. P., MacQueen, D., Niccolai, A., Short, T., Stabler, B., and Turner,
1164 R. (2011), *maptools: Tools for reading and handling spatial objects*, r package
1165 version 0.8-10.
- 1166 Link, W. A. and Barker, R. J. (2009), *Bayesian Inference: With Ecological*
1167 *Applications*, London, UK: Academic Press.
- 1168 Metropolis, N., Rosenbluth, A., Rosenbluth, M., Teller, A., Teller, E., et al.
1169 (1953), “Equation of state calculations by fast computing machines,” *The*
1170 *journal of chemical physics*, 21, 1087–1092.
- 1171 Metropolis, N. and Ulam, S. (1949), “The Monte Carlo method,” *Journal of the*
1172 *American Statistical Association*, 44, 335–341.
- 1173 Neal, R. (2003), “Slice sampling,” *Annals of Statistics*, 31, 705–741.
- 1174 Pebesma, E. and Bivand, R. (2011), *Package ‘sp’*, r package version 0.9-91.
- 1175 Plummer, M. (2011), *rjags: Bayesian graphical models using MCMC*, r package
1176 version 3-5.
- 1177 Plummer, M., Best, N., Cowles, K., and Vines, K. (2006), “CODA: Convergence
1178 Diagnosis and Output Analysis for MCMC,” *R News*, 6, 7–11.
- 1179 Robert, C. P. and Casella, G. (2004), *Monte Carlo statistical methods*, New
1180 York, USA: Springer.
- 1181 — (2010), *Introducing Monte Carlo Methods with R*, New York, USA: Springer.
- 1182 Roberts, G. O. and Rosenthal, J. S. (1998), “Optimal scaling of discrete ap-
1183 proximations to Langevin diffusions,” *Journal of the Royal Statistical Society:*
1184 *Series B (Statistical Methodology)*, 60, 255–268.
- 1185 Spiegelhalter, D., Thomas, A., Best, N., and Lunn, D. (2003), *WinBUGS User*
1186 *Manual Version 1.4*.
- 1187 Stabler, B. (2006), *shapefiles: Read and Write ESRI Shapefiles*, r package version
1188 0.6.
- 1189 Su, Y.-S. and Yajima, M. (2011), *R2jags: A Package for Running jags from R*,
1190 r package version 0.02-17.
- 1191 Thomas, A., O’Hara, B., Ligges, U., and Sturtz, S. (2006), “Making BUGS
1192 Open,” *R News*, 6, 12–17.