

AUTHOR QUERY FORM

	Book: Spatial Capture-Recapture Chapter: 6	Please e-mail your responses and any corrections to: E-mail: a.kittusamyramasamy@elsevier.com
---	--	---

Dear Author,

Any queries or remarks that have arisen during the processing of your manuscript are listed below and are highlighted by flags in the proof. (AU indicates author queries; ED indicates editor queries; and TS/TY indicates typesetter queries.) Please check your proof carefully and answer all AU queries. Mark all corrections and query answers at the appropriate place in the proof (e.g., by using on-screen annotation in the PDF file http://www.elsevier.com/framework_authors/tutorials/ePDF_voice_skin.swf) or compile them in a separate list, and tick off below to indicate that you have answered the query.
Please return your input as instructed by the project manager.

Uncited references: References that occur in the reference list but are not cited in the text. Please position each reference in the text or delete it from the reference list.

Missing references: References listed below were noted in the text but are missing from the reference list. Please make the reference list complete or remove the references from the text.

Location in Chapter	Query / remark

CHAPTER

6

Likelihood Analysis of Spatial Capture-Recapture Models

p0005 We have so far mainly focused on Bayesian analysis of spatial capture-recapture models. And, in the previous chapters we learned how to fit some basic spatial capture-recapture models using a Bayesian formulation of the models analyzed in **BUGS** engines including **WinBUGS** and **JAGS**. Despite our focus on Bayesian analysis, it is instructive to develop the basic concepts and ideas behind classical analysis based on likelihood methods and frequentist inference for SCR models. We recognized earlier (Chapter 5) that SCR models are versions of binomial (or other) GLMs, but with random effects (i.e., GLMMs). Throughout statistics, such models are routinely analyzed by likelihood methods. In particular, likelihood analysis is based on the integrated or marginal likelihood, in which the random effects are removed, by integration, from the conditional-on- s likelihood (s being the individual activity center). This has been the approach taken by Borchers and Efford (2008), Dawson and Efford (2009), and related papers. Therefore, in this chapter, we provide some conceptual and technical foundation for likelihood-based analysis of spatial capture-recapture models.

p0010 We will show here that it is straightforward to compute the maximum likelihood estimates (MLE) for SCR models by integrated likelihood. We develop the MLE framework using **R**, and we also provide a basic introduction to the **R** package **secr** (Efford, 2011a) which does likelihood analysis of SCR models (see also the standalone program **DENSITY** (Efford et al., 2004)). To set the context for likelihood analysis of SCR models, we first analyze the SCR model when N is known because, in that case, analysis is no different at all than a standard GLMM. We generalize the model to allow for unknown N using both conventional ideas based on the “full likelihood” (e.g., Borchers et al., 2002) and also using a formulation based on data augmentation. We obtain the MLEs for the SCR model from the wolverine camera trapping study (Magoun et al., 2011) analyzed in previous chapters to compare/contrast the results.

This sentence starts out as general (GLMMs) and then becomes about SCR. I'd separate those two.

Maybe use a general description, not specific to s , e.g., random effect, say η , ... conditional-on- n

of SCR

6.1 MLE with known N

p0015 We noted in Chapter 5 that, with N known, the basic SCR model is a type of binomial model with a random effect. For such models we can obtain maximum likelihood

estimators of model parameters based on integrated likelihood. The integrated likelihood is based on the marginal distribution of the data y , in which the random effects are removed by integration from the conditional-on- s distribution of the observations. See Chapter 2 for a review of marginal, conditional, and joint distributions. Conceptually, any SCR model begins with a specification of the conditional-on- s model $[y|s, \alpha]$ and we have a “prior distribution” for s , say $[s]$. Then, the marginal distribution of the data y is

$$[y|\alpha] = \int_S [y|s, \alpha][s]ds.$$

When viewed as a function of α for purposes of estimation, the marginal distribution $[y|\alpha]$ is often referred to as the *integrated likelihood*.

It is worth analyzing the simplest SCR model with known- N in order to understand the underlying mechanics and basic concepts. These are directly relevant to the manner in which many capture-recapture models are classically analyzed, such as model M_h , and individual covariate models (see Chapter 4).

To develop the integrated likelihood for SCR models, we first identify the conditional-on- s likelihood. The observation model for each encounter observation y_{ij} , for individual i and trap j , specified conditional-on- s_i , is

$$y_{ij}|s_i \sim \text{Binomial}(K, p_\alpha(x_j, s_i)), \quad (6.1.1)$$

where we have indicated the dependence of encounter probability, p_{ij} , on s and parameters α explicitly. For example, p_{ij} might be the Gaussian model given by

$$p_{ij} = \text{logit}^{-1}(\alpha_0) \exp(-\alpha_1 \|x_j - s_i\|^2),$$

where $\alpha_1 = 1/(2\sigma^2)$. The joint distribution of the data for individual i is the product of J such terms (i.e., contributions from each of J traps).

$$[y_i|s_i, \alpha] = \prod_{j=1}^J \text{Binomial}(K, p_\alpha(x_j, s_i)).$$

We note this assumes that encounter of individual i in each trap is independent of encounter in every other trap, conditional-on- s_i . This is the fundamental property of the basic model SCRO. The marginal likelihood is computed by removing s_i , by integration from the conditional-on- s likelihood, so we compute:

$$[y_i|\alpha] = \int_S [y_i|s_i, \alpha][s_i]ds_i.$$

In most SCR models, $[s] = 1/A(S)$ where $A(S)$ is the area of the prescribed state-space S (but see Chapter 11 for alternative specifications of $[s]$).

The joint likelihood for all N individuals, assuming independence of encounters among individuals, is the product of N such terms:

$$\mathcal{L}(\alpha|y_1, y_2, \dots, y_N) = \prod_{i=1}^N [y_i|\alpha].$$

We emphasize that two independence assumptions are explicit in this development: independence of trap-specific encounters within individuals and independence among individuals. In particular, this would only be valid when individuals are not physically restrained or removed upon capture, and when traps do not “fill up.”

p0035 The key operation for computing the likelihood is solving a two-dimensional integration problem. There are some general-purpose R packages that implement a number of multi-dimensional integration routines including `adapt` (Genz et al., 2007) and `R2cuba` (Hahn et al., 2010). In practice, we won’t rely on these extraneous R packages (except see Chapter 11 for an application of `R2cuba`), but instead will use perhaps less efficient methods in which we replace the integral with a summation over an equal-area mesh of points on the state-space \mathcal{S} and explicitly evaluate the integrand at each point. We invoke the rectangular rule for integration here¹ in which we evaluate the integrand on a regular grid of points of equal area and compute the average of the integrand over that grid of points. Let $u = 1, 2, \dots, nG$ index a grid of nG points, s_u , where the area of grid cells is constant, say A . In this case, the integrand, i.e., the marginal pmf of y_i , is approximated by

$$[y_i|\alpha] = \frac{1}{nG} \sum_{u=1}^{nG} [y_i|s_u, \alpha]. \quad (6.1.2)$$

p0040 This is a specific case of the general expression that could be used for approximating the integral for any arbitrary distribution $[s]$. The general case is

subscript \leftarrow $[y|\alpha] = \frac{A(\mathcal{S})}{nG} \sum_{u=1}^{nG} [y|s_u, \alpha][s_u].$

bold, with subscript

Under the uniformity assumption, $[s] = 1/A(\mathcal{S})$ and thus the grid-cell area cancels in the above expression to yield Eq. (6.1.2). The rectangular rule for integration can be seen as an application of the Law of Total Probability for a discrete random variable s , having nG unique values with equal probabilities $1/nG$.

s0010 6.1.1 Implementation (simulated data)

p0045 Here¹ we will illustrate how to carry out this integration and optimization based on the integrated likelihood using simulated data (i.e., see Section 5.5). Using `simSCR()`, we simulate data for 100 individuals and an array of 25 traps laid out in a 5×5 grid of traps having unit spacing. The specific encounter model is the Gaussian model. The 100 activity centers were simulated on a state-space defined by an 8×8 square within which the trap array was centered (thus the trap array is buffered by 2 units). Therefore, the density of individuals in this system is fixed at 100/64. In the following set of R commands we generate the data and then harvest the required data objects:

np005 ¹e.g., http://en.wikipedia.org/wiki/Rectangle_method.

```
## simulate a complete data set (perfect detection) o0005
> data <- simSCR0(discard0=FALSE, rnd=2013) o0010
  ## extract the objects that we need for analysis o0015
> y <- data$Y o0020
> traplocs <- data$traplocs o0025
> nind <- nrow(y)    ## in this case nind=N o0030
> J <- nrow(traplocs) o0035
> K <- data$K o0040
> xlim <- data$xlim o0045
> ylim <- data$ylim o0050
```

Now, we need to define the integration grid, say G , which we do with the following set of **R** commands (here, δ is the grid spacing):

```
> delta <- .2 o0055
> xg <- seq(xlim[1]+delta/2,xlim[2]-delta/2,by=delta) o0060
> yg <- seq(ylim[1]+delta/2,ylim[2]-delta/2,by=delta) o0065
> npix <- length(xg)      # valid for square state-space only o0070
> G <- cbind(rep(xg,npix),sort(rep(yg,npix))) o0075
> nG <- nrow(G) o0080
```

In this case, the integration grid is set up as a grid with spacing $\delta = 0.2$, which produces, for our example, a 40×40 grid of points for evaluating the integrand if the state-space buffer is set at 2. We note that the integration grid is set up here to correspond exactly to the state-space used in simulating the data. However, in practice, we wouldn't know this, and our estimate of N (for the unknown case, see below) would be sensitive to choice of the extent of the integration grid. As we've discussed previously, density, which is N standardized by the area of the state-space, will not be so sensitive in most cases.

We are now ready to compute the conditional-on- s likelihood and carry out the marginalization described by Eq. (6.1.2). We need to do this by defining an **R** function that computes the likelihood for the integration grid, as a function of the data objects y and $traplocs$, which were created above. However, it is a bit untidy to store the grid information in your workspace, and define the likelihood function in a way that depends on these things that exist in your workspace. Therefore, we build the **R** function so that it computes the integration grid *within* the function, thereby avoiding potential problems if our trapping grid locations change, or if we want to modify the state-space buffer easily. We therefore define the function, called `intlik1`, to which we pass the data objects and other information necessary to compute the marginal likelihood. This function is available in the `scrbook` package (use `?intlik1` at the **R** prompt). The code is reproduced here:

tiny! /
attitudes
used

```
intlik1 <- function(parm,y=y,X=traplocs, delta=.2, ssbuffer=2) { o0085
  Xl <- min(X[,1]) - ssbuffer    ## These lines of code are setting up the o0090
  Xu <- max(X[,1]) + ssbuffer    ## support for the integration which is o0095
```

```
o0100  Yu <- max(X[,2]) + ssbuffer ## the same as the state-space of "s"
o0105  Yl <- min(X[,2]) - ssbuffer
o0110  xg <- seq(Xl+delta/2,Xu-delta/2,length=npix)
o0115  yg <- seq(Yl+delta/2,Yu-delta/2,length=npix)
o0120  npix<- length(xg)

o0125  G <- cbind(rep(xg,npix),sort(rep(yg,npix)))
o0130  nG <- nrow(G)
o0135  D <- e2dist(X,G)

o0140  alpha0 <- parm[1]
o0145  alpha1 <- exp(parm[2]) # alpha1 restricted to be positive here

o0150  probcap <- plogis(alpha0)*exp(-alpha1*D*D)
o0155  Pm <- matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
o0160      # Frequency of all-zero encounter histories
o0165  n0 <- sum(apply(y,1,sum)==0)
o0170      # Encounter histories with at least 1 detection
o0175  ymat <- y[apply(y,1,sum)>0,]
o0180  ymat <- rbind(ymat,rep(0,nrow(ymat)))
o0185  lik.marg <- rep(NA,nrow(ymat))

o0190  for(i in 1:nrow(ymat)){
o0195      ## Nextline: log conditional likelihood for ALL possible values of s
o0200      Pm[1:length(Pm)] <- dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],
o0205          log=TRUE)
o0210      ## Next line: sum the log conditional likelihoods, exp() result
o0215      ## same as taking the product
o0220  lik.cond <- exp(colSums(Pm))
o0225      ## Take the average value == computing marginal
o0230  lik.marg[i] <- sum(lik.cond*(1/nG))
o0235  }
o0240  ## n0 = number of all-0 encounter histories
o0245  nv <- c(rep(1,length(lik.marg)-1),n0)
o0250  return( -1*( sum(nv*log(lik.marg)) ) )
o0255 }
```

p0055 We emphasize that this function (and subsequent) are not meant to be general-purpose routines for solving all of your SCR problems but, rather, they are meant for illustrative purposes—so you can see how the integrated likelihood is constructed and how we connect it to data and other information that is needed.

p0060 The function `intlik1` accepts as input the encounter history matrix, y , the trap locations, X , and the state-space buffer. This allows us to vary the state-space buffer and easily evaluate the sensitivity of the MLE to the size of the state-space. Note that we have a peculiar handling of the encounter history matrix y . In particular, we remove the all-zero encounter histories from the matrix and tack-on a single all-zero encounter history as the last row, which then gets weighted by the number of such encounter histories ($n0$). This is a bit long-winded and strictly unnecessary when N is known, but we did it this way because the extension to the unknown- N case is now transparent (as we demonstrate in the following section). The matrix Pm holds the log-likelihood contributions of each encounter frequency for each possible state-space location of the individual. The log contributions are summed up and the result exponentiated

on the next line, producing `lik.cond`, the conditional-on-s likelihood (Eq. (6.1.1)). The marginal likelihood (`lik.marg`) sums up the conditional elements weighted by the probabilities [`s`] (Eq. (6.1.2)).

This is a fairly primitive function which doesn't allow much flexibility in the data structure. For example, it assumes that K , the number of replicates, is constant for each trap. Further, it assumes that the state-space is a square. We generalize this to some extent later in this chapter.

Here is the R command for maximizing the likelihood using `nlm` (the function `optim` could also be used) and saving the results into an object called `frog`. The output is a list of the following structure, and these specific estimates are produced using the simulated data set:

```
# should take 15-30 seconds          o0260  
  
> starts <- c(-2,2)                  o0265  
> frog <- nlm(intlik1,starts,y=y,X=traplocs,delta=.1,ssbuffer=2,hessian=TRUE) o0270  
> frog                                o0275  
  
$minimum                               o0280  
[1] 297.1896                            o0285  
  
$estimate                             o0290  
[1] -2.504824 2.373343                 o0295  
  
$gradient                            o0300  
[1] -2.069654e-05 1.968754e-05       o0305  
  
$hessian  
[1] 48.67898 -19.25750  
[2] -19.25750 13.34114  
  
$code                                 o0310  
[1] 1                                    o0315  
  
$iterations                          o0320  
[1] 11                                  o0325
```

Caution

Details about this output can be found on the help page for `nlm`. We note briefly that `frog$minimum` is the negative log-likelihood value at the MLEs, which are stored in the `frog$estimate` component of the list. The order of the parameters is as they are defined in the likelihood function so, in this case, the first element (value = -2.504824) is the logit transform of p_0 and the second element (value = 2.373343) is the value of α_1 , the “coefficient” on distance-squared. The Hessian is the observed Fisher information matrix, which can be inverted to obtain the variance-covariance matrix using the command:

```
> solve(frog$hessian)                o0330
```

p0080 It is worth drawing attention to the fact that the estimates are slightly different than the Bayesian estimates reported previously in Section 5.6. There are several reasons for this. First, Bayesian inference is based on the posterior distribution and it is not generally the case that the MLE should correspond to any particular value of the posterior distribution. If the prior distributions in a Bayesian analysis are uniform, then the (multivariate) mode of the posterior is the MLE, but note Bayesians almost always report posterior means, and so there will typically be a discrepancy there. Secondly, we have implemented an approximation to the integral here and there might be a slight bit of error induced by that. We will evaluate that shortly. Third, the Bayesian analysis by MCMC is itself subject to some amount of Monte Carlo error which the analyst should always be aware of in practical situations. All of these different explanations are likely responsible for some of the discrepancy. Accounting for these, we see general consistency between the two estimates.

I don't think at
any point do we
say that the analysis
here is equivalent
to that in Ch 5, so
this comparison
comes out of the
blue.

p0085 In summary, for the basic SCR model, computing the integrated likelihood is a simple task when N is known. Even for N unknown it is not too difficult, and we will do that shortly. However, if you can solve the known- N problem then you should be able to do a real analysis, for example by considering different values of N , and computing the results for each value and then making a plot of the log-likelihood or AIC and choosing the value of N that produces the best log-likelihood or AIC. As a homework problem, we suggest that you can take the code given above and try to estimate N without modifying the code by just repeatedly applying it for different values of N in attempt to deduce the best value. We will formalize the unknown- N problem next.

s0015 6.2 MLE when N is unknown

Here, we build on the previous introduction to integrated likelihood but we consider now the case in which N is unknown. We will see that adapting the analysis based on the known- N model is straightforward for the more general problem. The main distinction is that we don't observe the all-zero encounter history, so we have to make sure we compute the probability for that encounter history, which we do by tacking a row of zeros onto the encounter history matrix. In addition, we include the number of such all-zero encounter histories (that is, the number of individuals not encountered) as an unknown parameter of the model. Call that unknown quantity n_0 , so that $N = n_0 + n$, where n is the number of unique individuals encountered. We will usually parameterize the likelihood in terms of n_0 , because optimization over a parameter space in which $\log(n_0)$ is unconstrained is preferred to a parameter space in which N must be constrained $N \geq n$. With n_0 unknown, we have to be sure to include a combinatorial term to account for the fact that, of the n observed individuals, there are $\binom{N}{n}$ ways to realize a sample of size n . The combinatorial term involves the unknown n_0 and thus it must be included in the likelihood. In evaluating the log-likelihood, we have to compute terms such as the log-factorial, $\log(N!) = \log((n_0 + n)!)$. We do

* remove; this is
whether we estimate
anything, plus we
did that in the known-
 N case, too.
* "We do that by
including..."

this in R by making use of the log-gamma function (`lgamma`) and the identity

$$\log(N!) = \underline{\text{lgamma}}(N + 1).$$

Is this supposed

Therefore, to compute the likelihood, we require the following three components: p0095 to be +?

(1) The marginal probability of each y_i as before,

$$[y_i | \alpha] = \int_S [y_i | s_i, \alpha] [s_i] ds_i.$$

(2) We compute the probability of an all-0 encounter history:

$$\pi_0 = [y = 0 | \alpha] = \int_S \text{Binomial}(0 | s_i, \alpha) [s_i] ds_i.$$

(3) The combinatorial term: $\binom{N}{n}$. Put on new line Then, the marginal likelihood has this form:

$$\mathcal{L}(\alpha, n_0 | y) = \frac{N!}{n! n_0!} \left\{ \prod_{i=1}^n [y_i | \alpha] \right\} \pi_0^{n_0}. \quad (6.2.1)$$

This is discussed in Borchers and Efford (2008, p. 379) as the conditional-on- N form of the likelihood—we also call it the “binomial form” of the likelihood because of its appearance.

Operationally, things proceed much as before: We compute the marginal probability of each observed y_i , i.e., by removing the latent s_i by integration. In addition, we compute the marginal probability of the “all-zero” encounter history y_{n+1} , and make sure to weight it n_0 times. We accomplish this by “padding” the data set with a single encounter history having $y_{n+1,j} = 0$ for all traps $j = 1, 2, \dots, J$. Then we be sure to include the combinatorial term in the likelihood or log-likelihood computation. We demonstrate this shortly. To analyze a specific case, we'll simulate our fake data set (simulated using the parameters given above). To set some things up in our workspace we do this:

```
## Obtain a simulated data set
> data <- simSCRO(discard0=TRUE, rnd=2013)

## Extract the items we need for analysis
> y <- data$Y
> nind <- nrow(y)
> traplocs <- data$traplocs
> J <- nrow(traplocs)
> K <- data$K
```

p0100

make

analyze ?

00335 | Isn't that
00340 the same
00345 data set?
00350 Maybe re-
00355 move this
00360 command
00365
00370

Recall that these data are simulated by default with $N = 100$, on an 8×8 unit state-space representing the trap locations buffered by two units, although you can modify the simulation script easily.

p0105 As before, the likelihood is defined in the R workspace as an R function, `intlik2`, which takes an argument being the unknown parameters of the model and additional arguments as prescribed. In particular, we provide the encounter history matrix `y`, the trap locations `traplocs`, the spacing of the integration grid (argument `delta`), and the state-space buffer. Here is the new likelihood function:

```
o0375 intlik2 <- function(parm,y=y,X=traplocs,delta=.3,ssbuffer=2){  
o0380   Xl <- min(X[,1]) - ssbuffer  
o0385   Xu <- max(X[,1]) + ssbuffer  
o0390   Yu <- max(X[,2]) + ssbuffer  
o0395   Yl <- min(X[,2]) - ssbuffer  
o0400   xg <- seq(Xl+delta/2,Xu-delta/2,delta)  
o0405   yg <- seq(Yl+delta/2,Yu-delta/2,delta)  
o0410   npix.x <- length(xg)  
o0415   npix.y <- plength(yg)  
o0420   area <- (Xu-Xl)*(Yu-Yl)/((npix.x)*(npix.y))  
o0425   G <- cbind(rep(xg,npix.y),sort(rep(yg,npix.x)))  
o0430   nG <- nrow(G)  
o0435   D <- e2dist(X,G)  
o0440   # extract the parameters from the input vector  
o0445   alpha0 <- parm[1]  
o0450   alpha1 <- exp(parm[2])  
o0455   n0 <- exp(parm[3]) # note parm[3] lives on the real line  
o0460   probcap <- plogis(alpha0)*exp(-alpha1*D*D)  
o0465   Pm <- matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))  
o0470   ymat <- rbind(y,rep(0,ncol(y)))  
  
o0475   lik.marg <- rep(NA,nrow(ymat))  
o0480   for(i in 1:nrow(ymat)){  
o0485     Pm[1:length(Pm)] <- dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],  
o0490       log=TRUE)  
o0495     lik.cond <- exp(colSums(Fm))  
o0500     lik.marg[i] <- sum(lik.cond*(1/nG))  
o0505   }  
o0510   nv <- c(rep(1,length(lik.marg)-1),n0)  
o0515   ## part1 here is the combinatorial term.  
o0520   ## math: log(factorial(N)) = lgamma(N+1)  
o0525   part1 <- lgamma(nrow(y)+n0+1) - lgamma(n0+1)  
o0530   part2 <- sum(nv*log(lik.marg))  
o0535   return( -1*(part1+part2))  
o0540 }
```

p0110 To execute this function for the data that we created with `simSCR0`, we execute the following command (saving the result in our friend `frog`). This results in the usual output, including the parameter estimates, the gradient, and the numerical Hessian, which is useful for obtaining asymptotic standard errors (see below):

```
o0545 > starts <- c(-2.5,0,4)  
o0550 > frog <- nlm(intlik2,starts,hessian=TRUE,y=y,X=traplocs,delta=.2,ssbuffer=2)  
  
o0555 Warning message:  
o0560 In nlm(intlik2, starts, hessian = TRUE, y = y, X = traplocs, delta = 0.2, :  
o0565 NA/Inf replaced by maximum positive value
```

```
> frog  
$minimum  
[1] 113.5004  
  
$estimate  
[1] -2.538333 0.902807 4.232810  
  
[... additional output deleted ...]
```

00570
00575
00580

00585
00590

00595

Executing `nls` here usually produces one or more R warnings due to numerical calculations happening on extremely small or large numbers (calculation of p near the edge of the state-space), and they also happen if a poor parameterization is used which produces evaluations of the objective function beyond the boundary of the parameter space (e.g., $n_0 < 0$). Such numerical warnings can often be minimized or avoided altogether by picking judicious starting values of parameters or properly transforming or scaling the parameters but, in general, they can be ignored. You will see from the `nls` output that the algorithm performed satisfactorily in minimizing the objective function. The estimate of population size, \hat{N} , for the state-space (using the default state-space buffer) is

```
> Nhat <- nrow(y) + exp(4.2328) #### This is n + MLE of n0  
> Nhat  
[1] 110.9099
```

p0115

column a

column b

This estimate which differs from the data-generating value ($N = 100$), as we might expect for a single realization. We usually will present an estimate of uncertainty associated with this MLE, which we can obtain by inverting the Hessian. Note that $\text{Var}(\hat{N}) = n + \text{Var}(\hat{n}_0)$. Since we have parameterized the model in terms of $\log(n_0)$, we use the delta method² described in Williams et al. (2002, Appendix F4) (see also Ver Hoef, 2012) to obtain the variance on the scale of n_0 as follows:

```
> (exp(4.2328)^2) * solve(frog$hessian)[3,3]  
[1] 260.2033
```

00600

00605

00610

```
> sqrt(260)  
[1] 16.12452
```

00625

00630

Therefore, the asymptotic “Wald-type” confidence interval for N is $110.91 \pm 1.96 \times 16.125 = (79.305, 142.515)$. To report this in terms of density, we scale appropriately by the area of the prescribed state-space, which is 64 units of area (i.e., an 8×8 square). Our MLE of D is $\hat{D} = 110.91/64 = 1.733$ individuals per square unit. To get the standard error for \hat{D} we need to divide the SE for \hat{N} by the area of the state-space, and so $\text{SE}(\hat{D}) = (1/64) * 16.12452 = 0.252$.

²We found a good set of notes on the delta approximation on Dr. David Patterson’s ST549 notes: <http://www.math.umt.edu/patterson/549/Delta.pdf>.

s0020 6.2.1 Integrated likelihood under data augmentation

p0120 The likelihood analysis developed in the previous sections is based on the likelihood in which N (or n_0) is an explicit parameter. This is usually called the “full likelihood” or sometimes “unconditional likelihood” (Borchers et al., 2002), because it is the likelihood for all individuals in the population, not just those which have been captured, i.e., not that which is *conditional on capture*. It is also possible to express an alternative unconditional likelihood using data augmentation, replacing the parameter N with ψ (e.g., see Section 7.1.6 [Royle and Dorazio, 2008], for an example). We don't go into detail here, but we note that the likelihood under data augmentation is a zero-inflated binomial mixture—precisely an occupancy-type model (Royle, 2006). Thus, while it is possible to carry out likelihood analysis of models under data augmentation, we primarily advocate data augmentation for Bayesian analysis.

s0025 6.2.2 Extensions

p0125 So far, we have only considered basic SCR models with no additional covariates. However, in practice, we are interested in covariate effects including “behavioral response,” sex-specificity of parameters, and potentially others. Some of these can be added directly to the likelihood if the covariate is fixed and known for all individuals, captured or not. An example is a behavioral response, which amounts to having a covariate $x_{ik} = 1$ if individual i was captured prior to occasion k and $x_{ik} = 0$ otherwise. For uncaptured individuals, $x_{ik} = 0$ for all k . Royle et al. (2011b) called this a global behavioral response, because the covariate is defined for all traps, no matter the trap in which an individual was captured. We could also define a local behavioral response which occurs at the level of the trap, i.e., $x_{ijk} = 1$ if individual i was captured in trap j prior to occasion k , etc. Trap-specific covariates such as trap type or status, or time-specific covariates such as date, are easily accommodated as well. As an example, Kéry et al. (2010) develop a model for the European wildcat (*Felis silvestris*) in which traps are either baited or not (a trap-specific covariate with only two values), and encounter probability varies over time in the form of a quadratic seasonal response. We consider models with behavioral response or fixed covariates in Chapter 7. The integrated likelihood routines we provided above can be modified directly for such cases, which we leave to the interested reader to investigate.

p0130 Sex specificity is more difficult to deal with, since sex is not known for uncaptured individuals (and sometimes not even for all captured individuals). To analyze such models, we do Bayesian analysis of the joint likelihood using data augmentation (Gardner et al., 2010b; Russell et al., 2012), discussed further in Chapter 7. For such covariates (i.e., that are not fixed and known for all individuals), it is somewhat more challenging to do MLE based on the joint likelihood as we have developed above.

Instead, it is more conventional to use what is colloquially referred to as the “Huggins-Alho”-type model, which is one of the approaches taken in the software package *secr* (Efford, 2011a). We introduce the *secr* package in Section 6.5.

184 CHAPTER 6 Likelihood Analysis of Spatial Capture-Recapture Models

I thought that were issues in determining the # of parameters in random effects models. How does scr do that?

6.3 Classical model selection and assessment

s0030

In most analyses, one is interested in choosing from among various potential models, or ranking models, or something else to do with assessing the relative merits of a set of models. A good thing about classical analysis based on likelihood is that we can apply Akaike Information Criterion (AIC) methods (Burnham and Anderson, 2002) without difficulty. AIC is convenient for assessing the relative merits of these different models, although if there are only a few models it is not objectionable to use hypothesis tests or confidence intervals to determine importance of effects. A second model selection context has to do with choosing among various detection models, although, as a general rule, we don't recommend this application of model selection. This is because there is hardly ever (if at all) a rational subject-matter based reason motivating specific distance functions. As a result, we believe that doing too much model selection will invariably lead to overfitting and thus overstatement of precision. This is the main reason that we haven't loaded you down with a basket of models for detection probability so far, although we discuss many possibilities in Chapter 7.

Comment Goodness-of-fit or model-checking—For many standard capture-recapture models, it is possible to identify goodness-of-fit statistics based on the multinomial likelihood (Cooch and White, 2006, Chapter 5), and evaluate model adequacy using formal statistical tests. Similar strategies can be applied to SCR models using expected cell frequencies based on the marginal distribution of the observations. Also, because computing MLEs is somewhat more efficient in many cases compared to Bayesian analysis, it is sometimes feasible to use bootstrap methods. At the present time, we don't know of any applications of goodness-of-fit testing for SCR models based on likelihood inference, although we discuss the use of Bayesian p-values for assessing model fit in Chapter 8. An important practical problem in trying to evaluate goodness-of-fit is that, in realistic sample sizes, fit tests often lack the power to detect departures from the model under consideration and so they may not be generally useful in practice.

p0140

6.4 Likelihood analysis of the wolverine camera trapping data

s0035

Comment Here we compute the MLEs for the wolverine data using an expanded version of the function we developed in the previous section. To accommodate that each trap might be operational a variable number of nights, we provided an additional argument to the likelihood function (allowing for a vector $\mathbf{K} = (K_1, \dots, K_J)$), which requires also a modification to the construction of the likelihood. In addition, we accommodate the state-space is a general rectangle, and we included a line in the code to compute the state-space area, which we apply below for computing density. The more general function (`intlik3`) is given in the R package `scrbook`. Incidentally, this function also returns the area of the state-space for a given set of parameter values, as an

p0145

something's missing here
already say
that two
sentences
earlier

attribute to the function value, which will be used in converting \hat{N} to \hat{D} . To use this function to obtain the MLEs for the wolverine camera trap study, we execute the following commands (note: these are in the help file and will execute if you type example(intlik3):

o0635 > library(scrbook)
o0640 > data(wolverine)

o0645 > traps <- wolverine\$wtraps
o0650 > traplocs <- traps[,2:3]/10000
o0655 > K.wolv <- apply(traps[,4:ncol(traps)],1,sum)

o0660 > y3d <- SCR23darray(wolverine\$wcaps,traps)
o0665 > y2d <- apply(y3d,c(1,2),sum)

o0670 > starts <- c(-1.5,0,3)

o0675 > wolv <- nlm(intlik3,starts,hessian=TRUE,y=y2d,K=K.wolv,X=traplocs,
o0680 delta=.2,ssbuffer=2)

o0685 > wolv
o0690 \$minimum
o0695 [1] 220.4313

o0700 \$estimate
o0705 [1] -2.8176120 0.2269395 3.5836875

o0710 [.... output deleted]

p0150 Of course we're interested in obtaining an estimate of population size for the prescribed state-space, or density, and associated measures of uncertainty which we do using the delta method (Williams et al., 2002, Appendix F4). To do all of that we need to manipulate the output of nlm, since we have our estimate in terms of $\log(n_0)$. *Comment*
We execute the following commands:

o0715 > wolv <- nlm(intlik3,starts,hessian=TRUE,y=y2d,K=K.wolv,X=traplocs,delta=.2,
o0720 ssbuffer=2)
o0725 > Nhat <- nrow(y2d)*exp(wolv\$estimate[3])
o0730 > area <- attr(intlik3(starts,y=y2d,K=K.wolv,X=traplocs,delta=.2,ssbuffer=2),
o0735 "SSarea")
o0740 > Dhat <- Nhat/area

o0745 > Dhat
o0750 [1] 0.5494947

o0755 > SE <- (1/area)*exp(wolv\$estimate[3])*sqrt(solve(wolv\$hessian)[3,3])

o0760 > SE
o0765 [1] 0.1087073

p0155 Our estimate of density is 0.55 individuals per "standardized unit" which is 100 km^2 , because we divided UTM coordinates by 10,000. So this is about 5.5 individuals per 1000 km^2 , with a SE of around 1.09 individuals. This compares closely with 5.77 reported in Section 5.9 based on Bayesian analysis of the model.
Comment

6.4.1 Sensitivity to integration grid and state-space buffer

s0040

The effect of approximating the integral by a discrete mesh of points is that it induces some numerical error in evaluation of the integral and, further, that error increases as the coarseness of the mesh increases. To evaluate the effect (or sensitivity) of the integration grid spacing, we obtained the MLEs for a state-space buffer of 2 (standardized units) and for integration grid with spacing $\delta = .3, .2, .1, .05$. The MLEs for these four cases including the relative runtime are given in Table 6.1. We see the results change only slightly as the integration grid changes. Conversely, the runtime on the platform of the day for the four cases increases rapidly. These runtimes could be regarded in relative terms, across platforms, for gaging the decrease in speed as the fineness of the integration grid increases.

p0160

We studied the effect of the state-space buffer on the MLEs, using a fixed $\delta = .2$ for all analyses. We used state-space buffers of 1–4 units stepped by .5. As we can see (Table 6.2), the estimates of D stabilize rapidly and the incremental difference is within the numerical error associated with approximating the integral.

p0165

Table 6.1 Runtime and MLEs for different integration grid resolutions for the wolverine camera trapping data.

sp005

δ	Estimates			
	Runtime(s)	$\hat{\alpha}_0$	$\hat{\alpha}_1$	$\log(n_0)$
0.30	9.9	-2.819786	1.258468	3.569731
0.20	32.3	-2.817610	1.254757	3.583690
0.10	115.1	-2.817570	1.255112	3.599040
0.05	407.3	-2.817559	1.255281	3.607158

fewer digits?

center column

Table 6.2 Results of the effect of the state-space buffer on the MLE. Given here are the state-space buffer, area of the state-space (area), the MLE of N (\hat{N}) for the prescribed state-space, and the corresponding MLE of density (\hat{D}).

sp010

Buffer	Area	\hat{N}	\hat{D}
1.0	66.98212	37.73338	0.5633352
1.5	84.36242	46.21008	0.5477567
2.0	103.74272	57.00617	0.5494956
2.5	125.12302	69.03616	0.5517463
3.0	148.50332	82.17550	0.5533580
3.5	173.88362	96.44018	0.5546249
4.0	201.26392	111.83524	0.5556646

center column

s0045 **6.4.2 Using a habitat mask (restricted state-space)**

p0170 In Section 5.10 we used a discrete representation of the state-space in order to have control over its extent and shape. This makes it easy to do things like clip out non-habitat, or create a *habitat mask* which defines suitable habitat. Clearly that formulation of the model is relevant to the calculation of the marginal likelihood in the sense that the discrete state-space is equivalent to the integration grid. Thus, for example, we could easily compute the MLE of parameters under some model with a restricted state-space merely by creating the required state-space at whatever grid resolution is desired, and then inputting that state-space into the likelihood function above, instead of computing it within the function. We can easily create an explicit state-space grid for integration from arbitrary polygons or GIS shapefiles which we demonstrate here. Our approach is to create the integration grid (or state-space grid) outside of the likelihood evaluation, and then determine which points of the grid lie in the polygon defined by the shapefile using functions in the R packages sp and maptools. For each point in the state-space grid (object G in the code below which is assumed to exist), we determine whether it is inside the polygon,³ identifying such points with a value of mask=1, and mask=0 for points that are *not* in the polygon. We load the shapefile, which originates by an application of the readShapeSpatial function. We have saved the result into an R data object called SSp, which is in the scrbook package. Here are the R commands for doing this (see the helpfile ?intlik4):

o0770 > library(maptools)
o0775 > library(sp)
o0780 > library(scrbook)

o0785 ##### If we have the .shp file in place, we would use this command:
o0790 ##### SSp <- readShapeSpatial('Sim_Polygon.shp')
o0795 ##### The object SSp is in data(fakeshapefile)
o0800 > data(fakeshapefile)
o0805 > Pcoord <- SpatialPoints(G)
o0810 > PinPoly <- over(Pcoord,SSp) #### determine if each point is in polygon
o0815 > mask <- as.numeric(!is.na(PinPoly[,1])) ## convert to binary 0/1
o0820 > G <- G[mask==1,]

semicolon

We created the function intlik4 which accepts the integration grid as an explicit argument; and this function is also available in the package scrbook.

p0175 We apply this modification to the wolverine camera trapping study. Royle et al. (2011b) created 2, 4, and 8-km state-space grids so as to remove “non-habitat” (mostly ocean, bays, and large lakes). We previously analyzed the model using JAGS and

np015 ³We perform this check using the over function. This function takes as its second argument (among others) an object of the class “SpatialPolygons” or “SpatialPolygonsDataFrame,” which can hold additional information for each polygon, and the output value of the function differs slightly for these two classes: if using a “SpatialPolygons” object, the function returns a vector of length equal to the number of points (e.g., in the example above), but if using a “SpatialPolygonsDataFrame” it returns a data frame (e.g., see Section 17.7 in Chapter 17). If you use the over function, make sure you know the class of your second argument so that when processing the function output you index it correctly.

→ since there are several, maybe give Section?

RETS

commu

wanna

applying

commu

also

modified function

dash

In this info should go in 6.4., or remove.

with the halftext mask

WinBUGS in Chapter 5 To set up the wolverine data and fit the model using maximum likelihood we execute the following commands:

This block is identical to the one in 6.4; remove and only refer here to intlik4?

```
> library(scrbook) o0825
> data(wolverine) o0830

> traps <- wolverine$wtraps o0835
> traplocs <- traps[,2:3]/10000 o0840
> K.wolv <- apply(traps[,4:ncol(traps)],1,sum) o0845

> y3d <- SCR23darray(wolverine$wcaps,traps) o0850
> y2d <- apply(y3d,c(1,2),sum) o0855
> G <- wolverine$grid2/10000 o0860

> starts <- c(-1.5,0,3) o0865
> wolv <- nlm(intlik4, starts, y=y2d, K=K.wolv, X=traplocs, G=G) o0870

> wolv o0875

$minimum
[1] 225.8355 o0880
o0885

$estimate
[1] -2.9955424 0.2350885 4.1104757 o0890
o0895

[... some output deleted ...] o0900
```

As before

commu Next we convert the parameter estimates to estimates of total population size for the prescribed state-space, and then obtain an estimate of density (per 1,000 km²) using the area computed as the number of pixels in the state-space grid, G, multiplied by the area per grid cell. In the present case (the calculation above) we used a state-space grid with 2 km × 2 km pixels. Finally, we compute a standard error using the delta approximation:

```
> area <- nrow(G)*4 o0905
# Nhat = n (observed) + MLE of n0 (not observed)
> Nhat <- 21 + exp(wolv$estimate[3]) o0910
> SE <- exp(wolv$estimate[3])*sqrt(solve(wolv$hessian)[3,3]) o0915
> D <- (Nhat/(nrow(G)*area))*1000 o0920
> SE.D <- (SE/(nrow(G)*area))*1000 o0925
o0930
```

We did this for each 2 km, 4 km, and 8 km state-space grid which produced the estimates summarized in Table 6.3. These estimates compare with the 8.6 (2 km grid) and 8.2 (8 km grid) reported in Royle et al. (2011b) based on a clipped state-space as described in Section 5.10.

commu

p0185 commu
dash
dash

sp015

Table 6.3 Maximum likelihood estimates (MLEs) and asymptotic standard errors (SE) for the wolverine camera trapping data using 2, 4, and 8 km state-space grids. → and a habitat mask.

Grid	α_0	α_1	$\log(n_0)$	N	SE	D(1000)	SE
2	-3.00	1.27	4.11	81.98	16.31	8.31	1.65
4	-2.99	1.34	4.16	84.88	16.76	8.57	1.69
8	-3.05	1.08	4.06	78.89	15.31	7.85	1.52

s0050

6.5 DENSITY and the R package secr

p0190

(e.g. Chapter 7) DENSITY is a software program developed by Efford (2004) for fitting spatial capture-recapture models based mostly on classical maximum likelihood estimation and related inference methods. Efford (2011a) has also released an R package called secr, that contains much of the functionality of DENSITY but also incorporates new models and features. Here, we briefly introduce the secr package, which we prefer to use over DENSITY, because it allows us to remain in the R environment for data processing and summarization. We provide a brief introduction to secr and some of its capabilities here, and we also use it for doing some analysis in other parts of this book! We believe that secr will be sufficient for many (if not most) of the SCR problems that one might encounter. It provides a flexible analysis platform, with a large number of summary features, and “publication ready” output. Its user interface is clean and intuitive to R users, and it has been stable, efficient, and reliable in the (fairly extensive) evaluations that we have done.

p0195

To install and run models in secr, you must download the package and load it in R.

o0935 > install.packages("secr")
o0940 > library(secr)

secr allows the user to simulate data and fit a suite of models with various detection functions and covariate responses. It also contains a number of helpful constructor functions for creating objects of the proper class that are recognized by other secr functions. We provide a brief overview of the capabilities here, but the secr help manual can be accessed with the command:

o0945 > RShowDoc("secr-manual", package = "secr")

We note that secr has many capabilities that we will not cover or do so only sparingly. We encourage you to read through the manual, the extensive documentation, and the vignettes, in order to get a better understanding of what the package is capable of. We also cover certain capabilities of secr in other chapters!

The main model-fitting function in `secr` is called `secr.fit`, which makes use p0200 of the standard R model specification framework with tildes. As an example, the equivalent of the basic model SCRO is fitted as follows:

```
> secr.fit(capturedata, model = list(D ~ 1, g0 ~ 1, sigma ~ 1),  
           buffer = 20000) 00950  
00955
```

where `capturedata` is the object created by `secr` containing the encounter history data and the trap information, and the model expression $\text{g0} \sim 1$ indicates the intercept-only (i.e., constant) model. Note that we use p_0 for the baseline encounter probability parameter, which is `g0` in `secr` notation. A number of possible models for encounter probability can be fitted including both pre-defined variables (e.g., `t` and `b` corresponding to “time” and “behavior”), and user-defined covariates of several kinds. For example, to include a global behavioral response, this would be written as $\text{g0} \sim \text{b}$. The discussion of this (global versus local trap-specific behavioral response) and other covariates is developed more in Chapter 7. We can also model covariates on density in `secr`, which we discuss in Chapter 11. It is important to note that `secr` requires the buffer distance to be defined in meters and density will be returned as number of animals per hectare. Thus to make comparisons between `secr` and output from other programs, we will often have to convert the density to the same units.

Before we can fit the models, the data must first be packaged properly for `secr`. p0205 We require data files that contain two types of information: trap layout (location and identification information for each trap), which is equivalent to the trap deployment file (TDF) described in Section 5.9 and the capture data file containing sampling session, animal identification, trap occasion, and trap location, equivalent in information content to the encounter data file (EDF). Sample session can be thought of as primary period identifier in a robust design-like framework—it could represent a yearly sample or multiple sample periods within a year, each of them producing data on a closed population. We discuss “multi-session” models in more detail below in Section 6.5.4 and Chapter 14.

There are three important constructor functions that help package-up your data p0210 no dash for use in `secr`: `read.traps`, `make.capthist`, and `read.mask`. We provide a brief description of each here, but apply them to our wolverine camera trapping data in Section 6.5.1.

1. `read.traps`: This function points to an external file or R data object containing the trap coordinates, and other information, and requires specification of the type of encounter devices (described in Section 6.5.1). A typical application of this function looks like the following, invoking the `data=` option when there is an existing R object containing the trap information:

```
> trapfile <- read.traps(data=traps, detector="proximity") 00960
```

2. `make.capthist`: This function takes the EDF and combines it with trap information, and the number of sampling occasions. A typical application looks like this:

o0965 > capturedata <- make.capthist(enc.data, trapfile, fmt="trapID",
noccasions=165)

See `?make.capthist` for definition of distinct file formats. Specifying `fmt = trapID` is equivalent to our EDF format.

- o1200 3. `read.mask`: If there is a habitat mask available (as described in Section 6.4.2), then this function will organize it so that `secr.fit` knows what to do with it. The function accepts either an external file name (see `?read.mask` for details of the structure) or a $nG \times 2$ **R** object, say `mask.coords`, containing the coordinates of the mask. A typical application looks like the following:

o0970 > grid <- read.mask(data=mask.coords)

These constructor functions produce output that can then be used in the fitting of models using `secr.fit`.

s0055 6.5.1 Encounter device types and detection models

p0215 The `secr` package requires that you specify the type of encounter device. Instead of describing models by their statistical distribution (Bernoulli, Poisson, etc.), `secr` uses certain operational classifications of detector types including “proximity,” “multi,” “single,” “polygon,” and “signal.” For camera trapping/hair snares we might consider “proximity” detectors or “count” detectors. The “proximity” detector type allows, at most, one detection of each individual at a particular detector on any occasion (i.e., it is equivalent to what we call the Bernoulli or binomial encounter process model, or model SCR0). The “count” detector designation allows repeat encounters of each individual at a particular detector on any occasion. There are other detector types that one can select such as: “polygon” detector type, which allows for a trap to be a sampled polygon (Royle and Young, 2008), which we discuss further in Chapter 15, and “signal” detector, which allows for traps that have a strength indicator, e.g., acoustic arrays (Dawson and Efford, 2009). The detector types “single” and “multi” refer to traps that retain individuals, thus precluding the ability for animals to be captured in other traps during the sampling occasion. The “single” type indicates trap that can only catch one animal at a time (single-catch traps), while “multi” indicates traps that may catch more than one animal at a time (multi-catch). These are both variations of the multinomial encounter models described in Chapter 9.

p0220 As with all SCR models, `secr` fits an encounter probability model (“detection function” in `secr` terminology) relating the probability of encounter to the distance of a detector from an individual activity center. `secr` allows the user to specify one of a variety of detection functions including the commonly used half-normal (“Gaussian”), hazard rate (“Gaussian hazard”), and (negative) exponential models. There are 12 different functions as of version 2.3.1 (see Table 7.1 in Chapter 7), but some are only available for simulating data. The different detection functions are defined in the `secr` manual and can be found by calling the help function for the detection function:

o0975 > ?detectfn

Most of the detection functions available in `secr` contain some kind of a scale parameter which is usually labeled σ . The units of this parameter default to meters in the `secr` output. We caution that the meaning of this parameter depends on the specific detection model being used, and it should not be directly compared as a measure of home range size across models. Instead, as we noted in Section 5.4, most encounter probability models imply a model of space usage and fitted encounter models should be converted to a common currency such as "area used." p0225

6.5.2 Analysis using the `secr` package

To demonstrate the use of the `secr` package, we will show how to do the same analysis on the wolverine study as shown in Section 5.9. To use the `secr` package, the data need to be formatted in a similar but slightly different manner than we use in WinBUGS. p0230

For example, in Section 5.9 we introduced a standard data format for the encounter data file (EDF) and trap deployment file (TDF). The EDF shares the same format as that used by the `secr` package, with 1 row for every encounter observation and 4 columns representing trap session ("Session"), individual identity ("ID"), sample occasion ("Occasion"), and trap identity ("trapID"). For a standard closed population study that takes place during a single season, the "Session" column is all 1's, to indicate a single primary sampling occasion. In addition to providing the encounter data file (EDF), we must tell `secr` information about the traps, which is formatted as a matrix with column labels "trapID," "x," and "y," the last two being the coordinates of each trap, with additional columns representing the operational state of each trap during each occasion (1 = operational, 0 = not). p0235

We demonstrate these differences now by walking through an analysis of the wolverine camera trapping data using `secr`. To read in the trap locations and other related information, we make use of the constructor function `read.traps`, which also requires that we specify the detector type. The detector type is important because it will determine the likelihood that `secr` will use to fit the model. Here, we have selected "proximity", which corresponds to the Bernoulli encounter model in which individuals are captured at most once in each trap during each sampling occasion: p0240

```
> library(secr) 00980  
> library(scrbook) 00985  
> data(wolverine) 00990  
  
> traps <- as.matrix(wolverine$traps) 00995  
> dimnames(traps) <- list(NULL,c("trapID","x","y",paste("day",1:165,sep=""))) 01000  
> traps1 <- as.data.frame(traps[,1:3]) 01005  
> trapfile1 <- read.traps(data=traps1,detector="proximity") 01010
```

Here we note that trap coordinates are extracted from the wolverine data but we do not scale them. This is because `secr` defaults to coordinate scaling of meters, which is the extant scaling of the wolverine trap coordinates. Note that we add a "trapID" column to the trap coordinates and provide appropriate column labels to the "traps" matrix. An important aspect of the wolverine study is that while the camera traps

data frame

were operated over a 165-day period, each trap was operational during only a portion of that period. We need to provide the trap operation information, which is contained in the columns to the right of the trap coordinates in our standard trap deployment file (TDF). Unfortunately, this is less easy to do in `secr`,⁴ which requires an external file with a single long string of 1's and 0's indicating the days in which each trap was operational (1) or not (0). The `read.traps` function will not allow for this information on trap operation if the data exists as an **R** object—instead, we can create this external file and then read it back in with `read.traps` using these commands:

```
o1015 > hold <- rep(NA,nrow(traps))
o1020 > for(i in 1:nrow(traps)){
o1025 >   hold[i] <- paste(traps[i,4:ncol(traps)],collapse="")
o1030 > }
o1035 > traps1 <- cbind(traps[,1:3],"usage"=hold)

o1040 > write.table(traps1, "traps.txt", row.names=FALSE, col.names=FALSE)
o1045 > trapfile2 <- read.traps("traps.txt", detector="proximity")
```

These operations can be accomplished using the function `scr2secr`, which is provided in the **R** package `scrbook`.

p0250 After reading in the trap data, we now need to create the encounter matrix or array using the `make.capthist` command, where we provide the capture histories in EDF format, which is the existing format of the data input file `wcaps`. In creating the capture history, we provide also the trapfile created previously, the format (e.g., here EDF format is `fmt = "trapID"`), and finally, we provide the number of occasions.

```
o1050 #
o1055 # Grab the encounter data file and format it:
o1060 #
o1065 > wolv.dat <- wolverine$wcaps
o1070 > dimnames(wolv.dat) <- list(NULL,c("Session","ID","Occasion","trapID"))
o1075 > wolv.dat <- as.data.frame(wolv.dat)
o1080 > wolvcapt2 <- make.capthist(wolv.dat,trapfile2,fmt="trapID",noccasions=165)
```

p0255 We also set up a habitat mask using the 2×2 km grid which we used previously in the analysis of the wolverine data and then pass the relevant objects to `secr.fit` as follows:

```
o1085 #
o1090 # Grab the habitat mask ( $2 \times 2$  km) and format it:
o1095 #
o1100 > gr2 <- (as.matrix(wolverine$grid2))
o1105 > dimnames(gr2) <- list(NULL,c("x","y"))
o1110 > gr2 <- read.mask(data=gr2)
o1115 #
o1120 # To fit the model we use secr.fit:
o1125 #
```

np020 ⁴as of v. 2.3.1.

```
>wolv.secr2 <- secr.fit(wolvcapt2,model=list(D ~ 1, g0 ~ 1, sigma ~ 1),  
buffer=20000,mask=gr2)  
model  
o1130  
o1135
```

We are using the "proximity-detector" model (SCR0), so we do not need to make any specifications in the command line, because we have specified the detector type using the constructor function `read.traps`, except to provide the buffer size (in meters). To specify different models, you can change the default model `D~1, g0~1, sigma~1`. We provide all of these commands and additional analyses in the `scrbook` package with the function called `secr_wolverine`. Printing the output object produces the following (slightly edited):

```
> wolv.secr2  
  
secr 2.3.1, 15:52:45 29 Aug 2012  
  
Detector type proximity  
Detector number 37  
Average spacing 4415.693 m  
x-range 593498 652294 m  
y-range 6296796 6361803 m  
N animals : 21  
N detections : 115  
N occasions : 165  
Mask area : 987828.1 ha  
  
Model : D ~ 1 g0 ~ 1 sigma ~ 1  
Fixed (real) : none  
Detection fn : halfnormal  
Distribution : poisson  
N parameters : 3  
Log likelihood : -602.9207  
AIC : 1211.841  
AICc : 1213.253  
  
Beta parameters (coefficients)  
  
beta SE.beta lcl ucl  
D -9.390124 0.22636698 -9.833795 -8.946452  
g0 -2.995611 0.16891982 -3.326688 -2.664535  
sigma 8.745547 0.07664648 8.595323 8.895772  
  
Variance-covariance matrix of beta parameters  
D g0 sigma  
D 0.0512420110 -0.0004113326 -0.003945371  
g0 -0.0004113326 0.0285339045 -0.006269477  
sigma -0.0039453711 -0.0062694767 0.005874683
```

```
Fitted (real) parameters evaluated at base levels of covariates
      link      estimate    SE.estimate      lcl      ucl
D      log     8.354513e-05  1.915674e-05  5.360894e-05  1.301982e-04
g0     logit   4.762453e-02  7.661601e-03  3.466689e-02  6.509881e-02
sigma   log    6.282651e+03  4.822512e+02  5.406315e+03  7.301037e+03
```

p0265 The object returned by `secr.fit` provides extensive default output when printed.

Much of this is basic descriptive information about the model, the traps, or the encounter data. We focus here on the parameter estimates. Under the fitted (real) parameters, we find D , the density, given in units of individuals/hectare (1 hectare = 10,000 m²). To convert this into individuals/1000 km², we multiply by 100,000, thus, our density estimate is 8.35 individuals/1000 km². The parameter σ is given in units of meters, and so this corresponds to 6.283 km. Both of these estimates are very similar to those obtained in our likelihood analysis summarized in Table 6.3 which, for the 2×2 km grid, we obtained $D = 8.31$ with a SE of $100,000 \times 1.915674e-05 = 1.9156$ and, accounting for the scale difference (1 unit = 10,000 m in the previous analysis), $\hat{\sigma} = \sqrt{1/(2\hat{\alpha}_1)} * 10,000 = 6.289$ km. The difference in the MLE between Table 6.3 and those produced by `secr` could be due to subtle differences in internal tuning of optimization algorithms, starting values, or other numerical settings. In addition, the likelihood is based on a Poisson prior for N (see Section 6.5.3). On the other hand, the SE is slightly larger based on `secr` which is due to a subtle difference in the interpretation of D under the `secr` model (see Section 6.5.3).

s0065 6.5.3 Likelihood analysis in the secr package

p0270 The `secr` package does likelihood analysis of SCR models for most classes of models as developed by Borchers and Efford (2008). Their formulation deviates slightly from the binomial form we presented in Section 6.2 above (though Borchers and Efford (2008) also mention the binomial form). Specifically, the likelihood that `secr` implements is that based on removing N from the likelihood by integrating the binomial likelihood (Eq. (6.2.1)) over a Poisson prior for N —what we will call the *Poisson-integrated likelihood* as opposed to the conditional-on- N (*binomial-form*) considered previously.

p0275 To develop the Poisson-integrated likelihood we compute the marginal probability of each y_i and the probability of an all-zero encounter history, π_0 , as before, to arrive at the marginal likelihood in the binomial-form:

$$\mathcal{L}(\boldsymbol{\alpha}, n_0 | \mathbf{y}) = \frac{N!}{n! n_0!} \left\{ \prod_i [y_i | \boldsymbol{\alpha}] \right\} \pi_0^{n_0}. \quad \text{for consistency, use } \prod_{i=1}^N$$

Now, what Borchers and Efford (2008) do is assume that $N \sim \text{Poisson}(\Lambda)$ and they do a further level of marginalization over this prior distribution:

$$\sum_{n_0=0}^{\infty} \frac{N!}{n! n_0!} \left\{ \prod_i [y_i | \boldsymbol{\alpha}] \right\} \pi_0^{n_0} \frac{\exp(-\Lambda) \Lambda^{n_0}}{N!}.$$

In Chapter 11 we write $\Lambda = \mu \|\mathcal{S}\|$ where $\|\mathcal{S}\|$ is the area of the state-space and μ p0280 is the density (“intensity”) of the point process. Carrying out the summation above produces exactly this marginal likelihood:

$$\mathcal{L}_2(\alpha, \Lambda | y) = \left\{ \prod_i [y_i | \alpha] \right\} \Lambda^n \exp(-\Lambda(1 - \pi_0)), \quad | \text{ see above}$$

which is Eq. (6.1.2) of Borchers and Efford (2008) except for notational differences. It also resembles the binomial form of the likelihood in Eq. (6.2.1) with $\Lambda^n \exp(-\Lambda\pi_0)$ replacing the combinatorial term and the $\pi_0^{n_0}$ term. We emphasize there are two marginalizations going on here: (1) the integration to remove the latent variables s ; and (2) summation to remove the parameter N . We provide a function for computing this in the `scrbook` package called `intlik3Poisson`. The help file for that function shows how to conduct a small simulation study to compare the MLE under the Poisson-integrated likelihood with that from the binomial form. (2008)

The essential distinction between our MLE and Borchers and Efford as implemented in `secr` is whether you keep N in the model or remove it by integration over a Poisson prior. If you have prescribed a state-space explicitly with a sufficiently large buffer, then we imagine there should be hardly any difference at all between the MLEs obtained by either the Poisson-integrated likelihood or the binomial-form of the likelihood which retains N . There is a subtle distinction in the sense that under the binomial form, we estimate the realized population size N for the state-space, whereas for the Poisson-integrated form we estimate the *prior* expected value which would apply to a hypothetical new study of a similar population (see Section 5.7.3).

Both models (likelihoods) assume s is uniformly distributed over space, but for the binomial model we make no additional assumption about N whereas we assume N is Poisson using the formulation in `secr` from (Borchers and Efford, 2008). Using data augmentation we could do a similar kind of integration but integrate N over a binomial (M, ψ) prior—which we referred to as the binomial-integrated likelihood in Section 4.2.4. So obviously the two approaches (data augmentation and Poisson-integrated likelihood) are approximately the same as M gets large. However, doing a Bayesian analysis by MCMC, we obtain an estimate of both N , the *realized population size*, and the parameter controlling its expected value ψ , which are, in fact, both identifiable from the data even using likelihood analysis (Royle et al., 2007). That said we can integrate N out completely and just estimate ψ as we noted in Section 6.2.1.

6.5.4 Multi-session models in `secr`

In practice we will often deal with SCR data that have some meaningful stratification or group structure. For example, we might conduct mist netting of birds on K consecutive days, repeated, say, T times during a year, or perhaps over T years. Or we might collect data from R distinct trapping grids. In these cases, we have T or R groups which we might reasonably regard as being samples of independent populations. While the groups might be distinct sites, years, or periods within years, they could also be other

pi

underlying biological groups such as sex or age. Conveniently, `secr` fits a specific model for stratified populations—referred to as *multi-session models*. These models build on the Poisson assumption which *underlies* the integrated likelihood used in `secr` (as described in Section 6.5.3). To understand the technical framework, let N_g be the population size of group g and *assume*

$$N_g \sim \text{Poisson}(\Lambda_g).$$

Naturally, we model group-specific covariates on Λ_g :

$$\log(\Lambda_g) = \beta_0 + \beta_1 z_g,$$

where z_g is some group-specific covariate such as a categorical index to the group, or a trend variable, or a spatial covariate, such as treatment effect or habitat structure, if the groups represent spatial units. Under this model, we can marginalize *all* N_g parameters out of the likelihood to concentrate the likelihood on the parameters β_0 and β_1 precisely as discussed in Section 6.5.3. This Poisson hierarchical model is the basis of the multi-session models in `secr`.

p0300 To implement a multi-session model (or stratified population model) in `secr`, we provide the relevant stratification information in the “Session” variable of the input encounter data file (EDF). If “Session” has multiple values, then a “multi-session” object is created by default and session-specific variables can be described in the model. For example, if the session has two values for males and females then we have sex-specific densities, and baseline encounter probability p_0 (g_0 in `secr`) by just doing this (see Chapter 8 for the R code to set this up):

```
o1140 > out <- secr.fit(capdata, model=list(D ~ session, g0 ~ session, sigma ~ 1),
o1145           buffer=20000)
```

More detailed analysis is given in Section 8.1 where we fit a number of different models and apply methods of model selection to obtain model-averaged estimates of density.

p0305 We can also easily implement stratified population models in the various **BUGS** engines using data augmentation (Converse and Royle, 2012; Royle and Converse, *COMING IN REVIEW*, which we discuss, with examples, in Chapter 14.

s0075 6.5.5 Some additional capabilities of secr

p0310 The `secr` package has capabilities to do a complete analysis of SCR data sets, including model fitting, selection, and many summary analyses. In the previous sections, we've given a basic overview, and we do more in later chapters of this book. Here we mention a few of these other capabilities that you should know about as you use `secr`. Of course, you should skim through the associated documentation (`?secr`) to see more of what's available.

6.5.5.1 Alternative observation models

secr fits a wide range of alternative observation models besides the Bernoulli encounter model, including multinomial encounter models for “multi-catch” and “single catch” traps, models for sound attenuation from acoustic detection devices, and many others. We discuss many of these other methods in Chapter 9 and elsewhere in the book.

s0080

p0315

6.5.5.2 Summary statistics

secr provides a useful default summary of the data, but it also has summary statistics about animal movement including mean-maximum distance moved (the function MMDM). For example, see the help page ?MMDM which lists a number of other summary functions which take a capthist object.

s0085

p0320

For
> moves(capthist) o1150
> dbar(capthist) o1155
> RPSV(capthist) o1160
> MMDM(capthist, min.recapt = 1, full = FALSE) o1165
> ARL(capthist, min.recapt = 1, plt = FALSE, full = FALSE) o1170

The function moves returns the observed distances moved, dbar returns the average distance moved, RPSV produces a measure of dispersion about the home range center, and ARL gives the *Asymptotic Range Length* which is the asymptote of an exponential model fit to the observed range length vs. the number of detections of each individual (Jett and Nichols, 1987).

6.5.5.3 State-space buffer

secr will produce a warning if the state-space buffer is chosen too small. For example, in fitting the wolverine data as in Section 6.5.2 but with a 1000 m buffer, and we see the following warning message:

s0090

p0325

commu

Warning message:
In secr.fit(wolvcapt2, model=list(D ~ 1, g0 ~ 1, sigma ~ 1), buffer=1000):
predicted relative bias exceeds 0.01 with buffer = 1000 o1175
o1180
o1185

This should cause you to contemplate modifying the state-space buffer if that is a reasonable thing to do in the specific application.

increasing

commu

6.5.5.4 Model selection and averaging

secr does likelihood ratio tests to compare nested models using the function LR.test. You can create model selection tables based on AIC or AICc, using the function AIC, and obtain model-averaged parameter estimates using the function model.average (See Chapter 8 for examples).

s0095

p0330

s0100 **6.5.5.5 Population closure test**

p0335 **secr** has a population closure test with the function `closure.test`, which implements the tests of Stanley and Burnham (1999) or Otis et al. (1978). The function is used like this: `closure.test(object, SB = FALSE)`. Here `object` is a capthist object and `SB` is a logical variable that, if TRUE, produces the Stanley and Burnham (1999) test.

commma
communc

s0105 **6.5.5.6 Density mapping and effective sample area**

p0340 **secr** produces likelihood versions of the various summaries of posterior density and effective sample area that we discussed in Chapter 5. For example, while **secr** reports estimates of the expected value of N or density directly in the summary output from fitting a model, you can use the function `region.N` to produce estimates of N for any given region. In addition, **secr** has functions for creating maps of detection contours for individuals traps, or for the entire trap array. See the function `pdot.contour`, and `fxi.contour` for computing the two-dimensional pdf of the locations of one or more individual activity centers (as in Section 5.11.3). In the context of likelihood analysis, estimation of a random effect s is based on a plug-in application of Bayes' Rule. When s has a uniform distribution, and we use a discrete evaluation of the integral, it can be computed simply by renormalizing the likelihood:

$$[s|y, \theta] = \frac{[y|s, \theta]}{\sum_s [y|s, \theta]}.$$

sounds like these
confusing, having
these functions
listed under one
"See"; sounds like
they do the same

Any of the `intlik` functions given previously in this chapter can be easily modified to return the posterior distribution of s for any, or all, individuals, or an individual that is not encountered.

p0345 Effective sample area (see Section 5.12) can be calculated in **secr** using the functions `esa` and `esa.plot`.

s0110 **6.5.5.7 Covariate models**

p0350 **secr** has many capabilities for modeling covariates. It has a number of built-in models that allow certain covariates on encounter probability, which we cover to a large extent in Chapter 7, and also see Chapter 8 for more examples. **secr** also allows covariates to be built into the density model (see Chapter 11). It has some built-in response surface models, allowing for the fitting of linear or quadratic response surfaces. This is done by modifying the density model in `secr.fit`. For example, $D \sim 1$ is a constant density surface, and $D \sim x + y$ fits a linear response surface, etc. See the manual `secr-densitysurfaces.pdf` for the details.

→ sounds like we
can only use covs
on density, but
we can use custom
covs on g5, too,
right?

some word missing?

p0355 There are a number of ways to model your own "custom" covariates (as opposed to pre-specified models). One way is to use the `addCovariates` function and supply it a `mask` or `traps` object along with some "spatialdata." Or, if you have covariates at each trap location then it will extrapolate to all points on the habitat mask. There's also a method by which the user can create a function of geographic coordinates, `userDfn`, which seems to provide additional flexibility, although we haven't used

this method. There is a handy function `predictDsurface` for producing density maps under the specified model for density.

6.6 Summary and outlook

s0115

In this chapter, we discussed basic concepts related to classical analysis of SCR models based on likelihood methods. Analysis is based on the so-called integrated or marginal likelihood, in which the individual activity centers (random effects) are removed from the conditional-on-s likelihood by integration. We showed how to construct the integrated likelihood and fit some simple models in the R programming language. In addition, likelihood analysis for some broad classes of SCR models can be accomplished using the R library `secr` (Efford, 2011a), which we provided a brief introduction to. In later chapters we provide more detailed analyses of SCR data using likelihood methods and the `secr` package.

p0360

Why or why not use likelihood inference exclusively? For certain specific models, it is/may be more computationally efficient to produce MLEs (for an example see Chapter 12). And, likelihood analysis makes it easy to do model selection by AIC and compute standard errors or confidence intervals. However, **BUGS** is extremely flexible in terms of describing models and we can devise models in the **BUGS** language easily that we cannot fit in `secr`. For example, in Chapter 16 we consider open population models, which are straightforward to develop in **BUGS** but, so far, there is no available platform for doing MLE of such models. We can also fit models in **BUGS** that accommodate missing covariates in complete generality (e.g., unobserved sex of individuals), and we can adopt SCR models to include auxiliary data types. For example, we might have camera trapping and genetic data and we can describe the models directly in **BUGS** and fit a joint model (Gopalaswamy et al., 2012b). To do maximum likelihood estimation, we have to write a custom new piece of code for each model⁵ or hope someone has done it for us. You should have some capability to develop your own MLE routines with the tools we provided in this chapter.

p0365

⁵Although we may be able to handle multiple survey methods together in `secr` using the multi-session models.

np025

Non Print Items

- sp005 **Abstract:** In this chapter we show how to compute maximum likelihood estimates of SCR model parameters both directly in R and using the R package secr (Efford, 2011a,b). We walk through a simulation example in which the number of animals, N , is known in order to provide a clear development of concepts and implementation, and show that this approach is readily extended to the more general case where N is unknown. Using a discrete state space gives us the flexibility to clip out areas of non-habitat (for example, bodies of water) and carry out our likelihood analysis only across habitat deemed suitable to the species. Likelihood functions for all these scenarios are provided in the accompanying R package scrbook and we apply these models to a wolverine camera trapping data set. Alternatively, we can use the R package secr, which provides a range of functionality for preparing data, fitting different models (e.g., with behavioral, time of site-specific effects) and summarizing and visualizing the results. We guide the reader through all steps associated with an analysis in secr and re-analyze the wolverine data set. While the Bayesian implementation of SCR models in the BUGS language gives us the flexibility to fit models we cannot readily fit in secr, for certain models a likelihood-based analysis is straightforward and can be more practical and efficient. We cover some of these cases in later chapters of the book.
- st075 **Keywords:** Akaike Information Criterion (AIC), Binomial integrated likelihood, Conditional likelihood, Effective sample area, Habitat mask, Integrated likelihood, Marginal likelihood, Numerical integration, Numerical optimization, Poisson integrated likelihood, Program DENSITY, secr, Unconditional likelihood