

1

INTRODUCTION

2

GLMS AND WINBUGS

3

7
8
9

CLOSED POPULATION MODELS

4

FULLY SPATIAL CAPTURE-RECAPTURE MODELS

5

OTHER OBSERVATION MODELS

6

LIKELIHOOD ANALYSIS OF SPATIAL CAPTURE-RECAPTURE MODELS

We have so far mainly focused on Bayesian analysis of spatial capture-recapture models. And, in the previous chapters we learned how to fit some basic spatial capture-recapture models using a Bayesian formulation of the models analyzed in BUGS engines including **WinBUGS** and **JAGS**. Despite our focus on Bayesian analysis, it is instructive to develop the basic concepts and ideas behind classical analysis based on likelihood methods and frequentist inference for SCR models. This has been the approach taken by Borchers and Efford (2008); Dawson and Efford (2009) and related papers. Therefore, in this chapter, we provide some conceptual and technical foundation for likelihood-based analysis of spatial capture-recapture models. We recognized earlier (Chapt. 4) that SCR models are versions of binomial (or other) GLMs, but with random effects, i.e., GLMMs. These models are routinely analyzed by likelihood methods. In particular, likelihood analysis is based on the integrated likelihood in which the random effects are removed by integration from the likelihood. In SCR models, the 2-dimensional coordinate, \mathbf{s} , is a bivariate random effect. Beyond that, there is little difference between likelihood analysis of SCR models and ordinary GLMMs.

We will show here that it is straightforward to compute the maximum likelihood estimates (MLE) for SCR models by integrated or marginal likelihood. We develop the MLE framework using **R**, and we also provide a basic introduction to an **R** package **secr** (Efford, 2011) which mostly does likelihood analysis of SCR models (see also the stand-alone package **DENSITY** (Efford et al., 2004)). To set the context for likelihood analysis of SCR models, we first analyze the SCR model here when N is known because, in that case, it is precisely a GLMM and does not pose any difficulty at all. We generalize the model to allow for unknown N using both

conventional ideas based on the “joint likelihood” (e.g., Borchers et al., 2002) and also using a formulation based on data augmentation. We obtain the MLEs for the SCR model from the wolverine camera trapping study (Magoun et al., 2011) analyzed in previous chapters to compare/contrast the results.

6.1 MLE WITH KNOWN N

We noted in Chapt. 4 that, with N known, the basic SCR model is a type of binomial regression with a random effect. For such models we can obtain maximum likelihood estimators of model parameters based on integrated likelihood. The integrated likelihood is based on the marginal distribution of the data y in which the random effects are removed by integration from the conditional-on- \mathbf{s} distribution of the observations. See Chapt. ?? for a review of marginal, conditional and joint distributions. Conceptually, any SCR model begins with a specification of the conditional-on- \mathbf{s} model $[y|\mathbf{s}, \boldsymbol{\alpha}]$ and we have a “prior distribution” for \mathbf{s} , say $[\mathbf{s}]$. Then, the marginal distribution of the data y is

$$[y|\boldsymbol{\alpha}] = \int_{\mathbf{s}} [y|\mathbf{s}, \boldsymbol{\alpha}][\mathbf{s}]d\mathbf{s}.$$

When viewed as a function of $\boldsymbol{\alpha}$ for purposes of estimation, the marginal distribution $[y|\boldsymbol{\alpha}]$ is often referred to as the *integrated likelihood*.

It is worth analyzing the simplest SCR model with known- N in order to understand the underlying mechanics and basic concepts. These are directly relevant to the manner in which many capture-recapture models are classically analyzed, such as model M_h , and individual covariate models (see Chapt. 3).

To develop the integrated likelihood for SCR models, we first identify the conditional-on- \mathbf{s} likelihood. The observation model for each encounter observation y_{ij} , specified conditional on \mathbf{s}_i , is

$$y_{ij}|\mathbf{s}_i \sim \text{Binomial}(K, p_{\alpha}(\mathbf{x}_j, \mathbf{s}_i)) \quad (6.1.1)$$

where we have indicated the dependence of p_{ij} on \mathbf{s} and parameters $\boldsymbol{\alpha}$ explicitly. For example, p_{ij} might be the Gaussian model given by

$$p_{ij} = \text{logit}^{-1}(\alpha_0) \exp(\alpha_1 \|\mathbf{x}_j - \mathbf{s}_i\|)$$

where $\alpha_1 = -1/(2\sigma^2)$. The joint distribution of the data for individual i is the product of J such terms (i.e., contributions from each of J traps).

$$[\mathbf{y}_i|\mathbf{s}_i, \boldsymbol{\alpha}] = \prod_{j=1}^J \text{Binomial}(K, p_{\alpha}(\mathbf{x}_j, \mathbf{s}_i))$$

We note this assumes that encounter of individual i in each trap is independent of encounter in every other trap, conditional on \mathbf{s}_i , this is the fundamental property of the basic model SCR0.

The marginal likelihood is computed by removing \mathbf{s}_i , by integration (hence also *integrated* likelihood), from the conditional-on- \mathbf{s} likelihood and regarding the *marginal* distribution of the data as the likelihood. That is, we compute:

$$[y|\boldsymbol{\alpha}] = \int_{\mathcal{S}} [y_i|\mathbf{s}_i, \boldsymbol{\alpha}][\mathbf{s}_i] d\mathbf{s}_i$$

In most SCR models, $[\mathbf{s}] = 1/A(\mathcal{S})$ where $A(\mathcal{S})$ is the area of the prescribed state-space \mathcal{S} (but see Chapt. ?? for alternative specifications of $[\mathbf{s}]$).

The joint likelihood for all N individuals, assuming independence of encounters among individuals, is the product of N such terms:

$$\mathcal{L}(\boldsymbol{\alpha}|\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_N) = \prod_{i=1}^N [y_i|\boldsymbol{\alpha}]$$

We emphasize that two independence assumptions are explicit in this development: independence of trap-specific encounters within individuals and also independence among individuals. In particular, this would only be valid when individuals are not physically restrained or removed upon capture, and when traps do not “fill up” (i.e., this is model SCR0, from Chapt. 4).

The key operation for computing the likelihood is solving a 2-dimensional integration problem. There are some general purpose **R** packages that implement a number of multi-dimensional integration routines including **adapt** (Genz et al., 2007) and **R2cuba** (Hahn et al., 2010). In practice, we won’t rely on these extraneous **R** packages (except see Chapt. ?? for an application of **R2cuba**) but instead will use perhaps less efficient methods in which we replace the integral with a summation over an equal area mesh of points on the state-space \mathcal{S} and explicitly evaluate the integrand at each point. We invoke the rectangular rule for integration here¹ in which we evaluate the integrand on a regular grid of points of equal area and compute the average of the integrand over that grid of points. Let $u = 1, 2, \dots, nG$ index a grid of nG points, \mathbf{s}_u , where the area of grid cells is constant, say A . In this case, the integrand, i.e., the marginal pmf of \mathbf{y}_i , is approximated by

$$[y_i|\boldsymbol{\alpha}] = \frac{1}{nG} \sum_{u=1}^{nG} [y_i|\mathbf{s}_u, \boldsymbol{\alpha}] \quad (6.1.2)$$

This is a specific case of the general expression that could be used for approximating the integral for any arbitrary distribution $[\mathbf{s}]$. The general case is

$$[y|\boldsymbol{\alpha}] = \frac{A(\mathcal{S})}{nG} \sum_{u=1}^{nG} [y|\mathbf{s}_u, \boldsymbol{\alpha}][\mathbf{s}_u]$$

¹e.g., http://en.wikipedia.org/wiki/Rectangle_method

Under the uniformity assumption, $[s] = 1/A(\mathcal{S})$ and thus the grid-cell area cancels in the above expression to yield eq. 6.1.2. The rectangular rule for integration can be seen as an application of the Law of Total Probability for a discrete random variable s , having nG unique values with equal probabilities $1/nG$.

6.1.1 Implementation (simulated data)

Here we will illustrate how to carryout this integration and optimization based on the integrated likelihood using simulated data (i.e., see sec. ??). Using `simSCR0` we simulate data for 100 individuals and a 25 trap array laid out in a 5×5 grid of unit spacing. The specific encounter model is the Gaussian model. The 100 activity centers were simulated on a state-space defined by a 8×8 square within which the trap array was centered (thus the trap array is buffered by 2 units). Therefore, the density of individuals in this system is fixed at $100/64$. In the following set of **R** commands we generate the data and then harvest the required data objects:

```
## simulate a complete data set (perfect detection)
data<-simSCR0(discard0=FALSE,rnd=2013)
## extract the objects that we need for analysis
y<-data$Y
traplocs<-data$traplocs
nind<-nrow(y) ## in this case nind=N
J<-nrow(traplocs)
K<-data$K
xlim<-data$xlim
ylim<-data$ylim
```

Now we need to define the integration grid, say G , which we do with the following set of **R** commands (here, `delta` is the grid spacing):

```
delta<- .2
xg<-seq(xlim[1]+delta/2,xlim[2]-delta/2,by=delta)
yg<-seq(ylim[1]+delta/2,ylim[2]-delta/2,by=delta)
npix<-length(xg) # valid for square state-space only
G<-cbind(rep(xg,npix),sort(rep(yg,npix)))
nG<-nrow(G)
```

In this case, the integration grid is set up as a grid with spacing $\delta = 0.2$ which produces, for our example, a 40×40 grid of points for evaluating the integrand if the state-space buffer is set at 2. We note that the integration grid is set-up here to correspond exactly to the state-space used in simulating the data. However, in practice, we wouldn't know this, and our estimate of N (for the unknown case, see below) would be sensitive to choice of the extent of the integration grid. As we've discussed previously, density, which is N standardized by the area of the state-space, will not be so sensitive in most cases.

We are now ready to compute the conditional-on- s likelihood and carry-out the marginalization described by Eq. 6.1.2. We need to do this by defining an **R**

function that computes the likelihood for the integration grid, as a function of the data objects `y` and `X` which were created above. However, it is a bit untidy to store the grid information in your workspace, and define the likelihood function in a way that depends on these things that exist in your workspace. Therefore, we build the **R** function so that it computes the integration grid *within* the function, thereby avoiding potential problems if our trapping grid locations change, or if we want to modify the state-space buffer easily. We therefore define the function, called `intlik1`, to which we pass the data objects and other information necessary to compute the marginal likelihood. This function is available in the `scrbook` package (use `?intlik1` at the **R** prompt). The code is reproduced here:

```

151 intlik1<-function(parm,y=y,delta=.2,X=traplocs,ssbuffer=2){
152
153   Xl<-min(X[,1]) - ssbuffer    ## these lines of code are setting up the
154   Xu<-max(X[,1]) + ssbuffer    ## support for the integration which is
155   Yu<-max(X[,2]) + ssbuffer    ## the same as the state-space of "s"
156   Yl<-min(X[,2]) - ssbuffer
157   xg<-seq(Xl+delta/2,Xu-delta/2,,length=npix)
158   yg<-seq(Yl+delta/2,Yu-delta/2,,length=npix)
159   npix<-length(xg)
160
161   G<-cbind(rep(xg,npix),sort(rep(yg,npix)))
162   nG<-nrow(G)
163   D<- e2dist(X,G)
164
165   alpha0<-parm[1]
166   alpha1<-exp(parm[2]) # alpha1 restricted to be positive here
167                        # for convenience (it is negated below)
168   probcap<- plogis(alpha0)*exp(-alpha1*D*D)
169   Pm<-matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
170                        # all zero encounter histories
171   n0<-sum(apply(y,1,sum)==0)
172                        # encounter histories with at least 1 detection
173   ymat<-y[apply(y,1,sum)>0,]
174   ymat<-rbind(ymat,rep(0,ncol(ymat)))
175   lik.marg<-rep(NA,nrow(ymat))
176   for(i in 1:nrow(ymat)){
177     ## next line: log conditional likelihood for ALL possible values of s
178     Pm[1:length(Pm)]<- (dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],log=TRUE))
179     ## next line: sum the log conditional likelihoods, exp() result
180     ## same as taking the product
181     lik.cond<- exp(colSums(Pm))
182     ## take the average value == computing marginal
183     lik.marg[i]<- sum( lik.cond*(1/nG))
184   }
185   ## n0 = number of all-0 encounter histories

```

```

186 nv<-c(rep(1,length(lik.marg)-1),n0)
187 -1*( sum(nv*log(lik.marg)) )
188 }

```

The function `intlik1` accepts as input the encounter history matrix, `y`, the trap locations, `X`, and the state-space buffer. This allows us to vary the state-space buffer and easily evaluate the sensitivity of the MLE to the size of the state-space. Note that we have a peculiar handling of the encounter history matrix `y`. In particular, we remove the all-zero encounter histories from the matrix and tack-on a single all-zero encounter history as the last row which then gets weighted by the number of such encounter histories (`n0`). This is a bit long-winded and strictly unnecessary when N is known, but we did it this way because the extension to the unknown- N case is now transparent (as we demonstrate in the following section). The matrix `Pm` holds the log-likelihood contributions of each encounter frequency for each possible state-space location of the individual. The log contributions are summed up and the result exponentiated on the next line, producing `lik.cond`, the conditional-on-`s` likelihood (Eq. 6.1.1 above). The marginal likelihood (`lik.marg`) sums up the conditional elements weighted by the probabilities `[s]` (Eq. 6.1.2 above). This is a fairly primitive function which doesn't allow much flexibility in the data structure. For example, it assumes that K , the number of replicates, is constant for each trap. Further, it assumes that the state-space is a square. We generalize this to some extent later in this chapter.

Here is the **R** command for maximizing the likelihood using `nlm` (the function `optim` could also be used) and saving the results into an object called `frog`. The output is a list of the following structure and these specific estimates are produced using the simulated data set:

```

211 # should take 15-30 seconds
212
213 starts<-c(-2,2)
214 frog<-nlm(intlik1,starts,y=y,delta=.1,X=traplocs,ssbuffer=2,hessian=TRUE)
215 frog
216
217 $minimum
218 [1] 297.1896
219
220 $estimate
221 [1] -2.504824  2.373343
222
223 $gradient
224 [1] -2.069654e-05  1.968754e-05
225
226 $hessian
227      [,1]      [,2]
228 [1,] 48.67898 -19.25750

```



```

229 [2,] -19.25750  13.34114
230
231 $code
232 [1] 1
233
234 $iterations
235 [1] 11

```

236 Details about this output can be found on the help page for `nlm`. We note
 237 briefly that `frog$minimum` is the negative log-likelihood value at the MLEs, which
 238 are stored in the `frog$estimate` component of the list. The Hessian is the observed
 239 Fisher information matrix, which can be inverted to obtain the variance-covariance
 240 matrix using the command:

```

241 > solve(frog$hessian)

```

242 It is worth drawing attention to the fact that the estimates are slightly different
 243 than the Bayesian estimates reported previously in sec. ???. There are several
 244 reasons for this. First Bayesian inference is based on the posterior distribution
 245 and it is not generally the case that the MLE should correspond to any particular
 246 value of the posterior distribution. If the prior distributions in a Bayesian analysis
 247 are uniform, then the (multivariate) mode of the posterior is the MLE, but note
 248 that Bayesians almost always report posterior *means* and so there will typically
 249 be a discrepancy there. Secondly, we have implemented an approximation to the
 250 integral here and there might be a slight bit of error induced by that. We will
 251 evaluate that shortly. Third, the Bayesian analysis by MCMC is itself subject to
 252 some amount of Monte Carlo error which the analyst should always be aware of
 253 in practical situations. All of these different explanations are likely responsible for
 254 some of the discrepancy. Accounting for these, we see general consistency between
 255 the two estimates.

256 In summary, for the basic SCR model, integrated likelihood is a really easy
 257 calculation when N is known. Even for N unknown it is not too difficult, and
 258 we will do that shortly. However, if you can solve the known- N problem then you
 259 should be able to do a real analysis, for example by considering different values of N
 260 and computing the results for each value and then making a plot of the log-likelihood
 261 or AIC and choosing the value of N that produces the best log-likelihood or AIC.
 262 As a homework problem we suggest that the reader take the code given above and
 263 try to estimate N without modifying the code by just repeatedly applying it for
 264 different values of N in attempt to deduce the best value. We will formalize the
 265 unknown- N problem next.

6.2 MLE WHEN N IS UNKNOWN

266 Here we build on the previous introduction to integrated likelihood but we consider
 267 now the case in which N is unknown. We will see that adapting the analysis based

on the known- N model is straightforward for the more general problem. The main distinction is that we don't observe the all-zero encounter history so we have to make sure we compute the probability for that encounter history which we do by tacking a row of zeros onto the encounter history matrix. In addition, we include the number of such all-zero encounter histories (that is, the number of individuals *not* encountered) as an unknown parameter of the model. Call that unknown quantity n_0 . Then, $N = n_0 + n$. We will usually parameterize the likelihood in terms of n_0 because optimization over a parameter space in which $\log(n_0)$ is unconstrained is preferred to a parameter space in which N must be constrained $N \geq n$. With n_0 unknown, we have to be sure to include a combinatorial term to account for the fact that of the n observed individuals there are $\binom{N}{n}$ ways to realize a sample of size n . The combinatorial term involves the unknown n_0 and thus it must be included in the likelihood. In evaluating the *log*-likelihood, we have to compute terms such as the log-factorial $\log(N!) = \log((n_0 + n)!)$. We do this in **R** by making use of the log-gamma function (`lgamma`) and the identity

$$\log(N!) = \text{lgamma}(N + 1).$$

Therefore, to compute the likelihood, we require the following 3 components:
(1) the marginal probability of each \mathbf{y}_i as before,

$$[\mathbf{y}_i | \boldsymbol{\alpha}] = \int_S [\mathbf{y}_i | \mathbf{s}_i, \boldsymbol{\alpha}] [\mathbf{s}_i] d\mathbf{s}_i.$$

(2) We compute the probability of an all-0 encounter history:

$$\pi_0 = [\mathbf{y} = \mathbf{0} | \boldsymbol{\alpha}] = \int_S \text{Binomial}(\mathbf{0} | \mathbf{s}_i, \boldsymbol{\alpha}) [\mathbf{s}_i] d\mathbf{s}_i$$

(3) The combinatorial term: $\binom{N}{n}$. Then, the marginal likelihood has this form:

$$\mathcal{L}(\boldsymbol{\alpha}, n_0 | \mathbf{y}) = \frac{N!}{n!n_0!} \left\{ \prod_{i=1}^n [\mathbf{y}_i | \boldsymbol{\alpha}] \right\} (\pi_0)^{n_0}. \quad (6.2.1)$$

This is discussed in Borchers and Efford (2008, p. 379) as the conditional-on- N form of the likelihood – we might also call it the “binomial form” because of its appearance.

Operationally, things proceed much as before: We compute the marginal probability of each observed \mathbf{y}_i , i.e., by removing the latent \mathbf{s}_i by integration. In addition, we compute the marginal probability of the “all-zero” encounter history \mathbf{y}_{n+1} , and make sure to weight it n_0 times. We accomplish this by “padding” the data set with a single encounter history having $y_{n+1,j} = 0$ for all traps $j = 1, 2, \dots, J$. Then we be sure to include the combinatorial term in the likelihood or log-likelihood computation. We demonstrate this shortly. To analyze a specific case, we'll read in our fake data set (simulated using the parameters given above). To set some things up in our workspace we do this:

```

299 data<-simSCRO(discard0=TRUE,rnd=2013) # obtain a simulated data set
300 ## extract the items we need for analysis
301 y<-data$Y
302 nind<-nrow(y)
303 traplocs<-data$traplocs
304 J<-nrow(traplocs)
305 K<-data$K

306 Recall that these data were generated with  $N = 100$ , on an  $8 \times 8$  unit state-space
307 representing the trap locations buffered by 2 units.
308 As before, the likelihood is defined in the R workspace as an R function,
309 intlik2, which takes an argument being the unknown parameters of the model
310 and additional arguments as prescribed. In particular, we provide the encounter
311 history matrix y, the trap locations traplocs, the spacing of the integration grid
312 (argument delta) and the state-space buffer. Here is the new likelihood function:

313 intlik2<-function(parm,y=y,delta=.3,X=traplocs,ssbuffer=2){
314
315   Xl<-min(X[,1]) -ssbuffer
316   Xu<-max(X[,1]) + ssbuffer
317   Yl<-min(X[,2]) - ssbuffer
318   Yl<-min(X[,2]) - ssbuffer
319
320   xg<-seq(Xl+delta/2,Xu-delta/2,delta)
321   yg<-seq(Yl+delta/2,Yu-delta/2,delta)
322   npix.x<-length(xg)
323   npix.y<-length(yg)
324   area<- (Xu-Xl)*(Yu-Yl)/((npix.x)*(npix.y))
325   G<-cbind(rep(xg,npix.y),sort(rep(yg,npix.x)))
326   nG<-nrow(G)
327   D<- e2dist(X,G)
328   # extract the parameters from the input vector
329   alpha0<-parm[1]
330   alpha1<-exp(parm[2])
331   n0<-exp(parm[3]) # note parm[3] lives on the real line
332   probcap<- plogis(alpha0)*exp(-alpha1*D*D)
333   Pm<-matrix(NA,nrow=nrow(probcap),ncol=ncol(probcap))
334   ymat<-rbind(y,rep(0,ncol(y)))
335
336   lik.marg<-rep(NA,nrow(ymat))
337   for(i in 1:nrow(ymat)){
338     Pm[1:length(Pm)]<- (dbinom(rep(ymat[i,],nG),K,probcap[1:length(Pm)],log=TRUE))
339     lik.cond<- exp(colSums(Pm))
340     lik.marg[i]<- sum( lik.cond*(1/nG) )
341   }
342   nv<-c(rep(1,length(lik.marg)-1),n0)
343   ## part1 here is the combinatorial term.

```

```

344  ## math: log(factorial(N)) = lgamma(N+1)
345  part1<- lgamma(nrow(y)+n0+1) - lgamma(n0+1)
346  part2<- sum(nv*log(lik.marg))
347  -1*(part1+ part2)
348  }

```

To execute this function for the data that we created with `simSCRO`, we execute the following command (saving the result in our friend `frog`). This results in the usual output, including the parameter estimates, the gradient, and the numerical Hessian which is useful for obtaining asymptotic standard errors (see below):

```

353  > starts<-c(-2.5,0,4)
354  > frog<-nlm(intlik2,starts,hessian=TRUE,y=y,X=traplocs,delta=.2,ssbuffer=2)
355
356  Warning message:
357  In nlm(intlik2, starts, hessian = TRUE, y = y, X = traplocs, delta = 0.2,  :
358    NA/Inf replaced by maximum positive value
359
360  > frog
361  $minimum
362  [1] 113.5004
363
364  $estimate
365  [1] -2.538333  0.902807  4.232810
366
367  [... additional output deleted ...]

```

While usually produces one or more **R** warnings due to numerical calculations happening on extremely small or large numbers (calculation of p near the edge of the state-space), and they also happen if a poor parameterization is used which produces evaluations of the objective function beyond the boundary of the parameter space (e.g., $n_0 < 0$). You will see from the `nlm` output that can be reproduced that the algorithm performed satisfactory in minimizing the objective function. The estimate of population size for the state-space (using the default state-space buffer) is

```

376  nrow(y)+exp(4.2328)
377  [1] 110.9099

```

Which differs from the data-generating value ($N = 100$) as we might expect for a single realization. We usually will present an estimate of uncertainty associated with this MLE which we can obtain by inverting the Hessian. Note that $\text{Var}(\hat{N}) = n + \text{Var}(\hat{n}_0)$. Since we have parameterized the model in terms of $\log(n_0)$ we use the delta method² (Williams et al., 2002, Appendix F4) to obtain the variance on the

² We found a good set of notes on the delta approximation on Dr. David Patterson's ST549 notes: <http://www.math.umd.edu/patterson/549/Delta.pdf>

383 scale of n_0 as follows:

```
384 (exp(4.2328)^2)*solve(frog$hessian)[3,3]
385 [1] 260.2033
386 > sqrt(260)
387 [1] 16.12452
```

388 Therefore, the asymptotic “Wald-type” confidence interval for N is $110.91 \pm 1.96 \times$
 389 $16.125 = (79.305, 142.515)$. To report this in terms of density, we scale appropri-
 390 ately by the area of the prescribed state-space which is 64 units of area (i.e., an
 391 8×8 square).

392 6.2.1 Integrated Likelihood using the model under data augmentation

393 The likelihood analysis developed in the previous sections is based on the likeli-
 394 hood in which N (or n_0) is an explicit parameter. This is usually called the “full
 395 likelihood” or sometimes “unconditional likelihood” (Borchers et al., 2002) because
 396 it is the likelihood for all individuals in the population, not just those which have
 397 been captured, i.e., not that which is *conditional on capture*. It is also possible to
 398 express an alternative unconditional likelihood using data augmentation, replacing
 399 the parameter N with ψ (e.g., see Sec. 7.1.6 Royle and Dorazio, 2008, for an ex-
 400 ample). We don’t go into detail here, but we note that the likelihood under data
 401 augmentation is a zero-inflated binomial mixture – precisely an occupancy type
 402 model (Royle, 2006). Thus, while it is possible to carryout likelihood analysis of
 403 models under data augmentation, we primarily advocate data augmentation for
 404 Bayesian analysis.

405 6.2.2 Extensions

406 We have only considered basic SCR models with no additional covariates. However,
 407 in practice, we are interested in other types of covariate effects including “behavioral
 408 response”, sex-specificity of parameters, and potentially other effects. Some of these
 409 can be added directly to the likelihood if the covariate is fixed and known for all
 410 individuals captured or not. An example is a behavioral response, which amounts
 411 to having a covariate $x_{ik} = 1$ if individual i was captured prior to occasion k and
 412 $x_{ik} = 0$ otherwise. For uncaptured individuals, $x_{ik} = 0$ for all k . Royle et al.
 413 (2011) called this a global behavioral response because the covariate is defined
 414 for all traps, no matter the trap in which an individual was captured. We could
 415 also define a *local* behavioral response which occurs at the level of the trap, i.e.,
 416 $x_{ijk} = 1$ if individual i was captured in trap j prior to occasion k , etc.. Trap-
 417 specific covariates such as trap type or status, or time-specific covariates such as
 418 date, are easily accommodated as well. As an example, Kéry et al. (2010) develop
 419 a model for the European wildcat *Felis silvestris* in which traps are either baited
 420 or not (a trap-specific covariate with only 2 values), and also encounter probability

varies over time in the form of a quadratic seasonal response. We consider models with behavioral response or fixed covariates in Chapt. 9. The integrated likelihood routines we provided above can be modified directly for such cases, which we leave to the interested reader to investigate.

Sex-specificity is more difficult to deal with since sex is not known for uncaptured individuals (and sometimes not even for all captured individuals). To analyze such models, we do Bayesian analysis of the joint likelihood using data augmentation (Gardner et al., 2010; Russell et al., 2012), discussed further in Chapt. 9. For such covariates (i.e., that are not fixed and known for all individuals), it is somewhat more challenging to do MLE for these based on the joint likelihood as we have developed above. Instead it is more conventional to use what is colloquially referred to as the “Huggins-Alho” type model which is one of the approaches taken in the software package `secr` (Efford, 2011) which we describe in sec. 6.5 below. This idea is motivated by thinking about unequal probability sampling methods known as Horvitz-Thompson sampling (e.g., see Overton and Stehman, 1995).

6.3 CLASSICAL MODEL SELECTION AND ASSESSMENT

In most analyses, one is interested in choosing from among various potential models, or ranking models, or something else to do with assessing the relative merits of a set of models. A good thing about classical analysis based on likelihood is we can apply AIC methods (Burnham and Anderson, 2002) without difficulty. There are two distinct contexts for model-selection that we think are relevant to SCR models. First is, and AIC selecting among models that represent distinct biological hypotheses (e.g., covariates affecting encounter probability or density). AIC is convenient for assessing the relative merits of these different models although if there are only a few models it is not objectionable to use hypothesis tests or confidence intervals to determine importance of effects. The second model selection context has to do with choosing among various detection functions although, as a general rule, we don’t recommend this application of model selection. This is because there is hardly ever (if at all) a rational subject-matter based reason motivating specific distance functions. As a result, we believe that doing too much model selection will invariably lead to over-fitting and thus over-statement of precision. This is the main reason that we haven’t loaded you down with a basket of models for detection probability so far, although we discuss many possibilities in Chapt. 9.

Goodness-of-fit – For many standard capture-recapture models, it is possible to identify goodness-of-fit statistics based on the multinomial likelihood and evaluate model adequacy using formal statistical tests. Similar strategies can be applied to SCR models using expected cell-frequencies based on the marginal distribution of the observations. Also, because computing MLEs is somewhat more efficient in many cases compared to Bayesian analysis, it is also sometimes easy to use bootstrap methods although, at the present time, we don’t know of any applications of

460 goodness-of-fit testing for SCR models based on likelihood inference³.

6.4 LIKELIHOOD ANALYSIS OF THE WOLVERINE CAMERA TRAPPING DATA

461 Here we compute the MLEs for the wolverine data using an expanded version of
 462 the function we developed in the previous section. To accommodate that each
 463 trap might be operational a variable number of nights, we provided an additional
 464 argument to the likelihood function (allowing for a vector $\mathbf{K} = (K_1, \dots, K_J)$),
 465 which requires also a modification to the construction of the likelihood. In addition,
 466 we accommodate the state-space is a general rectangle, and we included a line in
 467 the code to compute the state-space area which we apply below for computing
 468 density. The more general function (`intlik3`) is given in the **R** package `scrbook`.
 469 Incidentally, this function also returns the area of the state-space for a given set
 470 of parameter values, as an attribute to the function value, which will be used in
 471 converting \hat{N} to \hat{D} . To use this function to obtain the MLEs for the wolverine
 472 camera trap study, we execute the following commands (note: these are in the help
 473 file and will execute if you type `example(intlik3)`:

```
474 library("scrbook")
475 data("wolverine")
476
477 traps<-wolverine$wtraps
478 traplocs<-traps[,2:3]/10000
479 K.wolv<-apply(traps[,4:ncol(traps)],1,sum)
480
481 y3d<-SCR23darray(wolverine$wcaps,traps)
482 y2d<-apply(y3d,c(1,3),sum)
483
484 starts<-c(-1.5,0,3)
485
486 frog<-nlm(intlik3,starts,hessian=TRUE,y=y2d,K=K.wolv,X=traplocs,delta=.2,ssbuffer=2)
487
488 frog
489 $minimum
490 [1] 220.4313
491
492 $estimate
493 [1] -2.8176120 0.2269395 3.5836875
494
495 [.... output deleted ....]
```

496 Of course we're interested in obtaining an estimate of population size for the
 497 prescribed state-space, or density, and associated measures of uncertainty which we
 498 do using the delta method (Williams et al., 2002, Appendix F4). To do all of that

³WE NEED TO LOOK INTO THIS!!

```

499 we need to manipulate the output of nlm since we have our estimate in terms of
500 log(n0). We execute the following commands:

501 frog<-nlm(intlik3,starts,hessian=TRUE,y=y2d,K=K.wolv,X=traplocs,delta=.2,ssbuffer=2)
502 Nhat<-nrow(y2d)*exp(frog$estimate[3])
503 area<-attr(intlik3(starts,y=y2d,K=K.wolv,X=traplocs,delta=.2,ssbuffer=2),"SSarea")
504 Dhat<- Nhat/area
505
506 Dhat
507 [1] 0.5494947
508
509 SE<- (1/area)*exp(frog$estimate[3])*sqrt(solve(frog$hessian)[3,3])
510
511 SE
512 [1] 0.1087073

```

So our estimate of density is 0.55 individuals per “standardized unit” which is 100 km^2 , because we divided UTM coordinates by 10000. So this is about 5.5 individuals per 1000 km^2 , with a SE of around 1.09 individuals. This compares closely with 5.77 reported in sec. ?? based on Bayesian analysis of the model.

To evaluate the effect of the integration grid density, we obtained the MLEs for a state-space buffer of 2 (standardized units) and for integration grid with spacing $\delta = .3, .2, .1, .05$. The MLEs for these 4 cases including the relative runtime are given in Table 6.1. We see the results change only slightly as the fineness of the integration grid increases. Conversely, the runtime on the platform of the day for the 4 cases increases rapidly. These runtimes could be regarded in relative terms, across platforms, for gaging the decrease in speed as the fineness of the integration grid increases. The effect of this is that we anticipate some numerical error in approximating the integral on a mesh of points, and that error increases as the coarseness of the mesh increases.

Table 6.1. Run time and MLEs for different integration grid resolutions for the wolverine camera trapping data.

δ	Estimates			
	runtime	$\hat{\alpha}_0$	$\hat{\alpha}_1$	$\widehat{\log(n_0)}$
0.30	9.9	-2.819786	1.258468	3.569731
0.20	32.3	-2.817610	1.254757	3.583690
0.10	115.1	-2.817570	1.255112	3.599040
0.05	407.3	-2.817559	1.255281	3.607158

We studied the effect of the state-space buffer on the MLEs, using a fixed $\delta = .2$ for all analyses. The results are show in Table 6.2. We used state-space buffers of 1 to 4 units stepped by .5. As we can see in Table 6.2, the estimates of D stabilize rapidly and the incremental difference is within the numerical error associated with approximating the integral.

Table 6.2. Results of the effect of the state-space buffer on the MLE. Given here are the state-space buffer (buff), area of the state-space (area), the MLE of N (\hat{N}) for the prescribed state-space and the corresponding MLE of density (\hat{D}).

buff	area	\hat{N}	\hat{D}
1.0	66.98212	37.73338	0.5633352
1.5	84.36242	46.21008	0.5477567
2.0	103.74272	57.00617	0.5494956
2.5	125.12302	69.03616	0.5517463
3.0	148.50332	82.17550	0.5533580
3.5	173.88362	96.44018	0.5546249
4.0	201.26392	111.83524	0.5556646

6.4.1 Using a habitat mask (Restricted state-space)

In sec. ?? we used a discrete representation of the state-space in order to have control over its extent and shape. This makes it easy to do things like clip out non-habitat, or create a *habitat mask* which defines suitable habitat. Clearly that formulation of the model is relevant to the calculation of the marginal likelihood in the sense that the discrete state-space is equivalent to the integration grid. Thus, for example, we could easily compute the MLE of parameters under some model with a restricted state-space merely by creating the required state-space at whatever grid resolution is desired, and then inputting that state-space into the likelihood function above, instead of computing it in the function itself. We can easily create an explicit state-space grid for integration from arbitrary polygons or GIS shapefiles which we demonstrate here. Our approach is to create the integration grid (or state-space grid) outside of the likelihood evaluation, and then determine which points of the grid lie in the polygon defined by the shapefile using functions in the **R** packages **sp** and **maptools**. For each point in the state-space grid (object **G** in the code below which is assumed to exist), we determine whether it is inside the polygon⁴, identifying such points with a value of **mask=1** and **mask=0** for points that are *not* in the polygon. We load the shapefile which originates by an application of the **readShapeSpatial** function. We have saved the result into an **R** data object called **SSp** which is in the **scrbook** package. Here are the **R** commands for doing this (see the helpfile **?intlik4**):

```
library(maptools)
library(sp)
library(scrbook)
```

⁴We perform this check using the **over** function. This function takes as its second argument (among others) an object of the class “SpatialPolygons” or “SpatialPolygonsDataFrame”, which can hold additional information for each polygon, and the output value of the function differs slightly for these two classes: if using a “SpatialPolygons” object, the function returns a vector of length equal to the number of points (e.g., in the example above), but if using a SpatialPolygons-DataFrame it returns a data frame (e.g., see sec. ?? in Chapt. 7). If you use the **over** function, make sure you know the class of your second argument so that when processing the function output you index it correctly.

```

556
557 data("fakeshapefile")
558 ##### replaces this:
559 #####SSp<-readShapeSpatial('Sim_Polygon.shp')
560 Pcoord<-SpatialPoints(G)
561 PinPoly<-over(Pcoord,SSp) ### determine if each point is in polygon
562 mask<-as.numeric(!is.na(PinPoly[,1])) ## convert to binary 0/1
563 G<-G[mask==1,]

```

564 We created the function `intlik4` which accepts the integration grid as an explicit argument, and this function is also available in the package `scrbook`.

565 We apply this modification to the wolverine camera trapping study. Royle et al. (2011) created 2, 4 and 8 km state-space grids so as to remove “non-habitat” (mostly ocean, bays, and large lakes). We previously analyzed the model using **JAGS** and **WinBUGS** in Chapt. 4. To set up the wolverine data and fit the model we execute the following commands

```

571 library("scrbook")
572 data("wolverine")
573
574 traps<-wolverine$wtraps
575 traplocs<-traps[,2:3]/10000
576 K.wolv<-apply(traps[,4:ncol(traps)],1,sum)
577
578 y3d<-SCR23darray(wolverine$wcaps,traps)
579 y2d<-apply(y3d,c(1,3),sum)
580 G<-wolverine$grid2/10000
581
582 starts<-c(-1.5,0,3)
583 frog<-nlm(intlik4,starts,hessian=TRUE,y=y2d,K=K.wolv,X=traplocs,G=G)
584
585 frog
586 $minimum
587 [1] 225.8355
588
589 $estimate
590 [1] -2.9955424 0.2350885 4.1104757
591
592 [... some output deleted ...]

```

593 Next we convert the parameter estimates to estimates of total population size for the prescribed state-space, and then obtain an estimate of density (per 1000 km²) using the area computed as the number of pixels in the state-space grid `G` multiplied by the area per grid cell. In the present case (the calculation above) we used a state-space grid with 2km × 2km pixels. Finally, we compute a standard errors using the delta approximation: **XXXXX Check these commands and compare with whats in the table XXXXXX**

Table 6.3. MLEs for the wolverine camera trapping data using 2, 4 and 8 km state-space grids.

grid	α_0	α_1	$\log(n_0)$	N	SE	D(1000)	SE
2	-3.00	1.27	4.11	81.98	16.31	8.31	1.65
4	-2.99	1.34	4.16	84.88	16.76	8.57	1.69
8	-3.05	1.08	4.06	78.89	15.31	7.85	1.52

```

600 area<- nrow(G)*4
601 Nhat<- 21*exp(frog$estimate[3])
602 SE<- exp(frog$estimate[3])*sqrt(solve(frog$hessian)[3,3])
603 D<- (Nhat/(nrow(G)*area))*1000
604 SE.D<- (SE/(nrow(G)*area))*1000

```

605 We did this for each the 2 km, 4 km and 8 km state-space grids which produced
606 the estimates summarized in Tab. 6.3. These estimates compare with the 8.6 (2
607 km grid) and 8.2 (8 km grid) reported in Royle et al. (2011) based on a clipped
608 state-space as described in sec. ???.

6.5 DENSITY AND THE R PACKAGE SECR

609 **DENSITY** is a software program developed by Efford (2004) for fitting spatial
610 capture-recapture models based mostly on classical maximum likelihood estima-
611 tion and related inference methods. Efford (2011) has also released an **R** package
612 called **secr**, that contains much of the functionality of **DENSITY** but also in-
613 corporates new models and features. Here, we briefly introduce the **secr** package
614 which we prefer to use instead of **DENSITY** because it allows us to remain in the
615 **R** environment for data processing and summarization.

616 To install and run models in **secr**, you must download the package and load it
617 in **R**.

```

618 install.packages("secr")
619 library(secr)

```

620 **secr** allows the user to simulate data and fit a suite of models with various detection
621 functions and covariate responses. It also contains a number of helpful constructor
622 functions for creating objects of the proper class that are recognized by other **secr**
623 functions. We provide a brief overview of the capabilities here, but the **secr** help
624 manual can be accessed with the command:

```

625 RShowDoc("secr-manual", package = "secr")

```

626 We note that **secr** has many capabilities that we will not cover or do so only spar-
627 ingly. We encourage you to read through the manual to get a better understanding
628 of what the package is capable of.

The main model-fitting function in **secr** is called **secr.fit**, which makes use of the standard **R** model specification framework with tildes. As an example, the equivalent of the basic model SCR0 is fitted as follows: **XXXX need centered tildes here XXXXX**

```
secr.fit(capturedata, model = list(D~1, g0~1, sigma~1), buffer = 20000)
```

where **capturedata** is the object created by **secr** containing the encounter history data and the trap information, and the model expression **g0~1** indicates the intercept-only (i.e., constant) model. Possible predictors for detection probability include both pre-defined variables (e.g., **t** and **b** corresponding to “time” and “behavior”), and user-defined covariates of several kinds. For example, to include a global behavioral response, this would be written as **g0~b**. The discussion of this (global versus local trap-specific behavioral response) and other covariates is developed more in Chapt. 9.

Before we can fit the models, the data must first be packaged properly for **secr**. We require data files that contain two types of information: trap layout (location and identification information for each trap), which is equivalent to our trap deployment file (TDF) described in sec. ?? and the capture data file containing sampling *session*, animal identification, trap day, and trap location, equivalent in information content to our encounter data file (EDF). Sample session can be thought of as primary period identifier in a robust design like framework – it could represent a yearly sample or multiple sample periods within a year, each of them producing data on a closed population. We discuss “multi-session” models in more detail below, in sec. 6.5.5.

There are three important constructor functions that help package-up your data for use in **secr**: **read.traps**, **make.caphist** and **read.mask**. We provide a brief description of each here, but apply them to our wolverine camera trapping data in the next section:

(1) **read.traps**: This function points to an external file *or* **R** data object containing the trap coordinates, and other information, and also requires specification of the type of encounter devices (described in the next section). A typical application of this function looks like the following, invoking the **data=** option when there is an existing **R** object containing the trap information:

```
trapfile<-read.traps(data=traps,detector="proximity")
```

(2) **make.caphist**: This function takes the EDF and combines it with trap information, and the number of sampling occasions. A typical application looks like this:

```
capturedata<-make.caphist(enc.data,trapfile,fmt="trapID",noccasions=165)
```

See **?make.caphist** for definition of distinct file formats. Specifying **fmt = trapID** is equivalent to our EDF format.

(3) `read.mask`: If there is a habitat mask available (as described in sec. 6.4.1), then this function will organize it so that `secr.fit` knows what to do with it. The function accepts either an external file name (see `?read.mask` for details of the structure) or a $nG \times 2$ **R** object, say `mask.coords`, containing the coordinates of the mask. A typical application looks like the following:

```
grid<-read.mask(data=mask.coords)
```

These constructor functions produce output that can then be used in the fitting of models using `secr.fit`.

6.5.1 Encounter device types and detection models

The `secr` package requires that you specify the type of encounter device. Instead of describing models by their statistical distribution (Bernoulli, Poisson, etc.), `secr` uses certain operational classifications of detector types including 'proximity', 'multi', 'single', 'polygon' and 'signal'. For camera trapping/hair snares we might consider 'proximity' detectors or 'count' detectors. The 'proximity' detector type allows, at most, one detection of each individual at a particular detector on any occasion (i.e., it is equivalent to the Bernoulli or binomial encounter process model, or model SCR0). The 'count' detector designation allows repeat encounters of each individual at a particular detector on any occasion. There are other detector types that one can select such as: 'polygon' detector type which allows for a trap to be a sampled polygon (Royle and Young, 2008) which we discuss further in Chapt. ??, and 'signal' detector which allows for traps that have a strength indicator, e.g., acoustic arrays (Dawson and Efford, 2009). The detector types 'single' and 'multi' refer to traps that retain individuals, thus precluding the ability for animals to be captured in other traps during the sampling occasion. The 'single' type indicates trap that can only catch one animal at a time (single-catch traps), while 'multi' indicates traps that may catch more than one animal at a time (multi-catch). These are both variations of the multinomial encounter models described in Chapt. ??.

As with all SCR models, `secr` fits a detection function relating the probability of detection to the distance of a detector from an individual activity center. `secr` allows the user to specify one of a variety of detection functions including the commonly used half-normal, hazard rate, and exponential. There are 12 different functions as of version 2.3.1 (see Tab. ?? in Chapt. 9), but some are only available for simulating data. The different detection functions are defined in the `secr` manual and can be found by calling the help function for the detection function:

```
?detectfn
```

It is useful to note that `secr` requires the buffer distance to be defined in meters and density will be returned as number of animals per hectare. Thus to make comparisons between `secr` and other models, we will often have to convert the density to the same units.

Most of the detection functions available in `secr` contain some kind of a scale parameter which is usually labeled σ . The units of this parameter default to meters in the `secr` output. We caution that the meaning of this parameter depends on the specific model being used and it should not be directly compared as a measure of home-range size across models. Instead, as we noted in sec. ?? every encounter probability model implies a model of space-usage and fitted encounter models should be converted to a common currency such as “area used.”

6.5.2 Analysis using the `secr` package

To demonstrate the use of the `secr` package, we will show how to do the same analysis on the wolverine study as shown in sec. ?. To use the `secr` package, the data need to be formatted in a similar but slightly different manner than we use in **WinBUGS**.

For example, in sec. ?? we introduced a standard data format for the encounter data file (EDF) and trap deployment file (TDF). The EDF shares the same format as that used by the `secr` package with 1 row for every encounter observation and 4 columns representing trap session ('Session'), individual identity ('ID'), sample occasion ('Occasion'), and trap identity ('trapID'). For a standard closed population study that takes place during a single season, the 'Session' column in our case is all 1s, to indicate a single primary sampling occasion. In addition to providing the encounter data file (EDF), we must tell `secr` information about the traps, which is formatted as a matrix with column labels 'trapID', 'x' and 'y', the last two being the coordinates of each trap, with additional columns representing the operational state of each trap during each occasion (1=operational, 0=not).

We demonstrate these differences now by walking through an analysis of the wolverine camera trapping data using `secr`. To read in the trap locations and other related information, we make use of the constructor function `read.traps` which also requires that we specify the detector type. The detector type is important because it will determine the likelihood that `secr` will use to fit the model. Here, we have selected “proximity” which corresponds to the Bernoulli encounter model in which individuals are captured at most once in each trap during each sampling occasion:

```
library("secr")
library("scrbook")
data("wolverine")

traps<-as.matrix(wolverine$wtraps)
dimnames(traps)<-list(NULL,c("trapID","x","y",paste("day",1:165,sep="")))
traps1<-as.data.frame(traps[,1:3])
trapfile1<-read.traps(data=traps1,detector="proximity")
```

Here we note that trap coordinates are extracted from the wolverine data but we do *not* standardize them. This is because `secr` defaults to coordinate scaling of meters which is the extant scaling of the wolverine trap coordinates. Note that

we add a 'trapID' column to the trap coordinates and provide appropriate column labels to the 'traps' matrix. An important aspect of the wolverine study is that while the camera traps were operated over a 165 day period, each trap was operational during only a portion of that period. We need to provide the trap operation information which is contained in the columns to the right of the trap coordinates in our standard trap deployment file (TDF). Unfortunately, this is less easy to do in **secr**, which requires an external file with a single long string of 1's and 0's indicating the days in which each trap was operational (1) or not (0). The **read.traps** function will not allow for this information on trap operation if the data exists as an **R** object – instead, we can create this external file and then read it back in with **read.traps** using these commands:

```

759 hold<-rep(NA,nrow(traps))
760 for(i in 1:nrow(traps)){
761   hold[i]<-paste(traps[i,4:ncol(traps)],collapse="")
762 }
763 traps1<- cbind(traps[,1:3], "usage"=hold)
764
765 write.table(traps1, "traps.txt", row.names=FALSE, col.names=FALSE)
766 trapfile2<-read.traps("traps.txt",detector="proximity")

```

These operations can be accomplished using the function **scr2secr** which is provided in the **R** package **scrbook**.

After reading in the trap data, we now need to create the encounter matrix or array using the **make.capthist** command, where we provide the capture histories in EDF format, which is the existing format of the data input file **wcaps**. In creating the capture history, we provide also the trapfile created previously, the format (e.g., here EDF format is **fmt= 'trapID'**), and finally, we provide the number of occasions. We also set up a habitat mask using the 2×2 km grid which we used previously in the analysis of the wolverine data and then pass the relevant objects to **secr.fit** as follows:

```

777 #
778 # grab the encounter data file and format it:
779 #
780 wolv.dat<-wolverine$wcaps
781 dimnames(wolv.dat)<-list(NULL,c("Session","ID","Occasion","trapID"))
782 wolv.dat<-as.data.frame(wolv.dat)
783 wolvcapt2<-make.capthist(wolv.dat,trapfile2,fmt="trapID",noccasions=165)
784
785 # grab the habitat mask (2 x 2 km) and format it:
786 #
787 gr2<-(as.matrix(wolverine$grid2))
788 dimnames(gr2)<-list(NULL,c("x","y"))
789
790 # To fit the model we use secr.fit:

```

```

791 #
792 wolv.secr2<-secr.fit(wolvcapt2,model=list(D~1, g0~1, sigma~1), buffer=20000,mask=gr2)

```

We are using the basic “proximity detector” model (SCR0), so we do not need to make any specifications in the command line because we have specified the detector type using the constructor function `read.traps`, except to provide the buffer size (in *m*). To specify different models, you can change the default `D~1`, `g0~1`, `sigma~1`, which the interested reader can do with very little difficulty. We provide all of these commands and additional analyses in the `scrbook` package with the function called `secr.wolverine`. Printing the output object produces the following (slightly edited):

```

801 wolv.secr2
802
803 secr 2.3.1, 15:52:45 29 Aug 2012
804
805 Detector type      proximity
806 Detector number   37
807 Average spacing    4415.693 m
808 x-range           593498 652294 m
809 y-range           6296796 6361803 m
810 N animals         : 21
811 N detections      : 115
812 N occasions       : 165
813 Mask area         : 987828.1 ha
814
815 Model              : D~1 g0~1 sigma~1
816 Fixed (real)      : none
817 Detection fn       : halfnormal
818 Distribution       : poisson
819 N parameters       : 3
820 Log likelihood     : -602.9207
821 AIC                : 1211.841
822 AICc               : 1213.253
823
824 Beta parameters (coefficients)
825      beta      SE.beta      lcl      ucl
826 D      -9.390124 0.22636698 -9.833795 -8.946452
827 g0     -2.995611 0.16891982 -3.326688 -2.664535
828 sigma  8.745547 0.07664648  8.595323  8.895772
829
830 Variance-covariance matrix of beta parameters
831      D      g0      sigma
832 D      0.0512420110 -0.0004113326 -0.003945371
833 g0     -0.0004113326  0.0285339045 -0.006269477
834 sigma -0.0039453711 -0.0062694767  0.005874683

```



```

835
836 Fitted (real) parameters evaluated at base levels of covariates
837      link      estimate SE.estimate      lcl      ucl
838 D      log 8.354513e-05 1.915674e-05 5.360894e-05 1.301982e-04
839 g0     logit 4.762453e-02 7.661601e-03 3.466689e-02 6.509881e-02
840 sigma  log 6.282651e+03 4.822512e+02 5.406315e+03 7.301037e+03

```

The object returned by `secr.fit` provides extensive default output when printed. Much of this is basic descriptive information about the model, the traps, or the encounter data. We focus here on the parameter estimates. Under the fitted (real) parameters, we find D , the density, given in units of individuals/hectare (1 hectare = 10000 m^2). To convert this into individuals/1000 km^2 , we multiply by 100000, thus our density estimate is 8.35 individuals/1000 km^2 . The parameter σ is given in units of meters, and so this corresponds to 6.283 km . Both of these estimates are very similar to those obtained in our likelihood analysis summarized in Tab. ?? which, for the 2×2 km grid, we obtained $\hat{D} = 8.31$ with a SE of $100000 \times 1.915674e - 05 = 1.9156$ and, accounting for the scale difference (1 unit = 10000 m in the previous analysis), $\hat{\sigma} = \sqrt{1/(2\hat{\alpha}_1)} * 10000 = 6.289$ km . The difference in the MLE between Tab. ?? and those produced by `secr` are likely due to subtle differences in internal tuning of optimization algorithms, starting values or other numerical settings. In addition, see the next section. On the other hand, the SE is slightly larger based on `secr` which is due to a subtle difference in the interpretation of D under the `secr` model (See below).

6.5.3 Likelihood Analysis in the `secr` Package

The `secr` package does likelihood analysis of SCR models for most classes of models as developed by Borchers and Efford (2008). Their formulation deviates slightly from the binomial form we presented in sec. 6.2 above (though Borchers and Efford (2008) mention the binomial form). Specifically, the likelihood that `secr` implements is that based on removing N from the likelihood by integrating the binomial likelihood (Eq. 6.2.1 above) over a Poisson prior for N – what we will call the *Poisson-integrated likelihood* as opposed to the conditional-on- N (*binomial-form*) considered previously.

To develop the Poisson-integrated likelihood we compute the marginal probability of each \mathbf{y}_i and the probability of an all-0 encounter history, π_0 , as before, to arrive at the marginal likelihood in the binomial-form:

$$\mathcal{L}(\boldsymbol{\alpha}, n_0 | \mathbf{y}) = \frac{N!}{n!n_0!} \left\{ \prod_i [\mathbf{y}_i | \boldsymbol{\alpha}] \right\} \pi_0^{n_0}$$

Now, what Borchers and Efford (2008) do is assume that $N \sim \text{Poisson}(\Lambda)$ and they

do a further level of marginalization over this prior distribution:

$$\sum_{n_0=0}^{\infty} \frac{N!}{n!n_0!} \left\{ \prod_i [\mathbf{y}_i | \boldsymbol{\alpha}] \right\} \pi_0^{n_0} \frac{\exp(-\Lambda) \Lambda^N}{N!}$$

Carrying-out the summation above produces exactly this marginal likelihood:

$$\mathcal{L}_2(\boldsymbol{\alpha}, \Lambda | \mathbf{y}) = \left\{ \prod_i [\mathbf{y}_i | \boldsymbol{\alpha}] \right\} \Lambda^n \exp(-\Lambda \pi_0)$$

which is Eq. 2 of Borchers and Efford (2008) except for notational differences. It also resembles the binomial-form of the likelihood in Eq. 6.2.1 except with $\Lambda^n \exp(-\Lambda \pi_0)$ replacing the combinatorial term and the $\pi_0^{n_0}$ term. We emphasize there are two marginalizations going on here: (1) the integration to remove the latent variables \mathbf{s} ; and, (2) summation to remove the parameter N . We provide a function for computing this in the `scrbook` package called `intlik3Poisson`. The help file for that function shows how to conduct a small simulation study to compare the MLE under the Poisson-integrated likelihood with that from the binomial form.

The essential distinction between our MLE and Borchers and Efford as implemented in `secr` is whether you keep N in the model or remove it by integration over a Poisson prior. If you have prescribed a state-space explicitly with a sufficiently larger buffer, then we imagine there should be hardly any difference at all between the MLEs obtained by either the Poisson-integrated likelihood or the binomial-form of the likelihood which retains N . There is a subtle distinction in the sense that under the binomial form, we estimate the realized population size N for the state-space whereas, for the Poisson-integrated form we estimate the *prior* expected value which would apply to a hypothetical new study of a similar population.

Both models (likelihoods) assume \mathbf{s} is uniformly distributed over space, but for the binomial model we make no additional assumption about N whereas we assume N is Poisson using the formulation in `secr` from (Borchers and Efford, 2008). Using data augmentation we could do a similar kind of integration but integrate N over a binomial (M, ψ) prior – which we referred to as the binomial-integrated likelihood in sec. ???. So obviously the two approaches (data augmentation and Poisson-integrated likelihood) are approximately the same as M gets large. However, doing a Bayesian analysis by MCMC, we obtain an estimate of both N , the *realized population size*, and the parameter controlling its expected value ψ which are, in fact, both identifiable from the data even using likelihood analysis (Royle et al., 2007). That said we can integrate N out completely and just estimate ψ as we noted in sec. 6.2.1 above. And we could make a prediction for a new study which would be based on the posterior distribution of $M\psi$ which, we imagine, should have slightly larger uncertainty associated with it.

6.5.4 Other stuff secr does

secr has a function called `region.N` for estimating the realized value of N and not its expectation.

We can also model covariates on density, allowing that the activity centers \mathbf{s} are a realization of an inhomogeneous point process. We cover such models in Chapt. 10.

mapping, and other things???

6.5.5 Multi-Session Models in secr

In practice we will often deal with SCR data that have some meaningful stratification or group structure. For example, we might conduct mist-netting of birds on K consecutive days, repeated, say, T times during a year, or perhaps over T years. Or we might collect data from R distinct trapping grids. In these cases, we have T or R groups which we might reasonably regard as being samples of independent populations. While the groups might be distinct sites, year, or periods within years, they could also be other biological groups such as sex or age. Conveniently, `secr` fits a specific model for stratified populations – referred to as *multi-session* models. These models build on the Poisson assumption which underlies the integrated likelihood used in `secr` (as described in the previous section). To understand the technical framework, let N_g be the population size of group g and *assume*

$$N_g \sim \text{Poisson}(\Lambda_g).$$

Naturally, we model group-specific covariates on Λ_g :

$$\log(\Lambda_g) = \beta_0 + \beta_1 z_g$$

where z_g is some group-specific covariate such as a categorical index to the group, or a trend variable, or a spatial covariate, such as treatment effect or habitat structure, if the groups represent spatial units. Under this model, we can marginalize *all* N_g parameters out of the likelihood to concentrate the likelihood on the parameters β_0 and β_1 precisely as discussed in the previous section. This Poisson hierarchical model is the basis of the multi-session models in `secr`.

To implement a multi-session model (or stratified population model) in `secr`, we provide the relevant stratification information in the 'Session' variable of the input encounter data file (EDF). If 'Session' has multiple values then a "multi-session" object is created by default and session-specific variables can be described in the model. For example, if the session has 2 values for males and females then we have sex-specific densities, and baseline encounter probability p_0 (named g_0 in `secr`) by just doing this: **XXXXXX need centered tildes here XXXXX**

```
out<-secr.fit(wolvcapt,model=list(D~session, g0~session, sigma~1), buffer=20000)
```

More detailed analysis is given in sec. ?? where we fit a number of different models and apply methods of model selection to obtain model-averaged estimates of density.

We can also easily implement stratified population models in the various **BUGS** engines using data augmentation (Converse and Royle, 2012, 2013) which we address, with examples, in Chapt. ??.

6.6 SUMMARY AND OUTLOOK

In this chapter, we discussed basic concepts related to classical analysis of SCR models based on likelihood methods. Analysis is based on the so-called integrated or marginal likelihood in which the individual activity centers (random effects) are removed from the conditional-on-s likelihood by integration. We showed how to construct the integrated likelihood and fit some simple models in the **R** programming language. In addition, likelihood analysis for some broad classes of SCR models can be accomplished using the **R** library **secr** (Efford, 2011) which we provided a brief introduction to. In later chapters we provide more detailed analyses of SCR data using likelihood methods and the **secr** package.

Why or why not use likelihood inference exclusively? For certain specific models, it is probably more computationally efficient to produce MLEs (for an example see Chapt. ??). However, **BUGS** is extremely flexible in terms of describing models, although it sometimes can be quite inefficient. We can devise models in the **BUGS** language easily that we cannot fit in **secr**. E.g., random individual effects of various types (Chapt. 9), we can handle missing covariates in complete generality and seamlessly, and impose arbitrary distributions on random variables. Moreover, models can easily be adapted to include auxiliary data types. For example, we might have camera trapping and genetic data and we can describe the models directly in **BUGS** and fit a joint model (?). For the MLE we have to write a custom new piece of code for each model or hope someone has done it for us, although you should be able to do this with the tools we have provided here. Later we consider open population models which are straightforward to develop in **BUGS** but, so far, there is no available platform for doing MLE of such models, although we imagine one could develop this. On the other hand, likelihood analysis makes it easy to do model-selection by AIC and in some cases compute standard errors or carry-out goodness-of-fit evaluations.

7

MCMC DETAILS

8

971

972

973

GOODNESS OF FIT AND STUFF

9

974
975

COVARIATE MODELS

976

10

977
978

INHOMOGENEOUS POINT PROCESS

979

980
981
982

11

OPEN MODELS

BIBLIOGRAPHY

- Borchers, D. L., Buckland, S. T., and Zucchini, W. (2002), *Estimating animal abundance: closed populations*, vol. 13, Springer Verlag.
- Borchers, D. L. and Efford, M. G. (2008), “Spatially explicit maximum likelihood methods for capture–recapture studies,” *Biometrics*, 64, 377–385.
- Burnham, K. P. and Anderson, D. R. (2002), *Model selection and multimodel inference: a practical information-theoretic approach*, Springer Verlag.
- Converse, S. J. and Royle, J. A. (2012), “Dealing with incomplete and variable detectability in multi-year, multi-site monitoring of ecological populations,” in *Design and Analysis of Long-term Ecological Monitoring Studies*, ed. XXXXXXXX, XXXXXXXX, p. XXXX.
- (2013), “Hierarchical Spatial Capture-Recapture Models: Modeling population density based on replicated capture-recapture experiments,” XXXX.
- Dawson, D. K. and Efford, M. G. (2009), “Bird population density estimated from acoustic signals,” *Journal of Applied Ecology*, 46, 1201–1209.
- Efford, M. (2004), “Density estimation in live-trapping studies,” *Oikos*, 106, 598–610.
- (2011), *secr: Spatially explicit capture-recapture models*, r package version 2.3.1.
- Efford, M. G., Dawson, D. K., and Robbins, C. S. (2004), “DENSITY: software for analysing capture-recapture data from passive detector arrays,” *Animal Biodiversity and Conservation*, 217–228.
- Gardner, B., Royle, J. A., Wegan, M. T., Rainbolt, R. E., and Curtis, P. D. (2010), “Estimating black bear density using DNA data from hair snares,” *The Journal of Wildlife Management*, 74, 318–325.
- Genz, A. S., Meyer, M. R., Lumley, T., and Maechler, M. (2007), “The adapt Package. R package version 1.0-4,” .
- Hahn, T., Bouvier, A., and Kiêu, K. (2010), *R2Cuba: Multidimensional Numerical Integration*, r package version 1.0-6.
- Kéry, M., Gardner, B., Stoeckle, T., Weber, D., and Royle, J. A. (2010), “Use of Spatial Capture-Recapture Modeling and DNA Data to Estimate Densities of Elusive Animals,” *Conservation Biology*, 25, 356–364.
- Magoun, A. J., Long, C. D., Schwartz, M. K., Pilgrim, K. L., Lowell, R. E., and Valkenburg, P. (2011), “Integrating motion-detection cameras and hair snags for wolverine identification,” *The Journal of Wildlife Management*, 75, 731–739.
- Overton, W. S. and Stehman, S. V. (1995), “The Horvitz-Thompson theorem as

- 1019 a unifying perspective for probability sampling: with examples from natural
1020 resource sampling,” *American Statistician*, 261–268.
- 1021 Royle, J. A. (2006), “Site occupancy models with heterogeneous detection proba-
1022 bilities,” *Biometrics*, 62, 97–102.
- 1023 Royle, J. A. and Dorazio, R. M. (2008), *Hierarchical modeling and inference in*
1024 *ecology: the analysis of data from populations, metapopulations and communities*,
1025 Academic Press.
- 1026 Royle, J. A., Dorazio, R. M., and Link, W. A. (2007), “Analysis of multinomial
1027 models with unknown index using data augmentation,” *Journal of Computational*
1028 *and Graphical Statistics*, 16, 67–85.
- 1029 Royle, J. A., Magoun, A. J., Gardner, B., Valkenburg, P., and Lowell, R. E. (2011),
1030 “Density estimation in a wolverine population using spatial capture–recapture
1031 models,” *The Journal of Wildlife Management*, 75, 604–611.
- 1032 Royle, J. A. and Young, K. V. (2008), “A Hierarchical Model For Spatial Capture-
1033 Recapture Data,” *Ecology*, 89, 2281–2289.
- 1034 Russell, R. E., Royle, J. A., Desimone, R., Schwartz, M. K., Edwards, V. L.,
1035 Pilgrim, K. P., and McKelvey, K. S. (2012), “Estimating abundance of mountain
1036 lions from unstructured spatial samples,” *Journal of Wildlife Management*.
- 1037 Williams, B. K., Nichols, J. D., and Conroy, M. J. (2002), *Analysis and management*
1038 *of animal populations: modeling, estimation, and decision making*, Academic Pr.